

Recommending Products Based on the Latent Factors Model



Swetha Kolalapudi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Understand the latent factors model for collaborative filtering

Contrast the latent factors model and the nearest neighbors model

Use optimization techniques to solve for latent factors

- Stochastic gradient descent

Collaborative Filtering Techniques

Nearest Neighbors Model



Use the ratings of
“most similar” users

Latent Factor Analysis



Solve for underlying
factors that drive the
ratings

Latent Factor Analysis



**Analogous to PCA
(Principal Components
Analysis)
in matrix algebra**

Latent Factor Analysis



Why are some products rated high and others low?

Why do some users like certain products?

Are there some underlying factors that influence all users' ratings?


	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5

Once underlying factors are known


Predict any rating for a product by a user

Latent Factor Analysis

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5




	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			



	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					

Latent Factor Analysis



Users described by some
underlying factors

Latent Factor Analysis

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	2	4	3
U ₂	3	2	5	1	5
U ₃	4	2	3	5	4
U ₄	1	2	4	3	5
U ₅	1	4	2	5	3
U ₆	3	4	3	1	5




	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			



	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					


Products described by
the same underlying
factors

Latent Factor Analysis



Solve using known ratings

Latent Factor Analysis



Use the solution to find unknown ratings

Latent Factor Analysis

This representation is
analogous to content
based filtering

	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			

*

	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					

Content Based Filtering

**Rate every product
against the relevant
attributes**

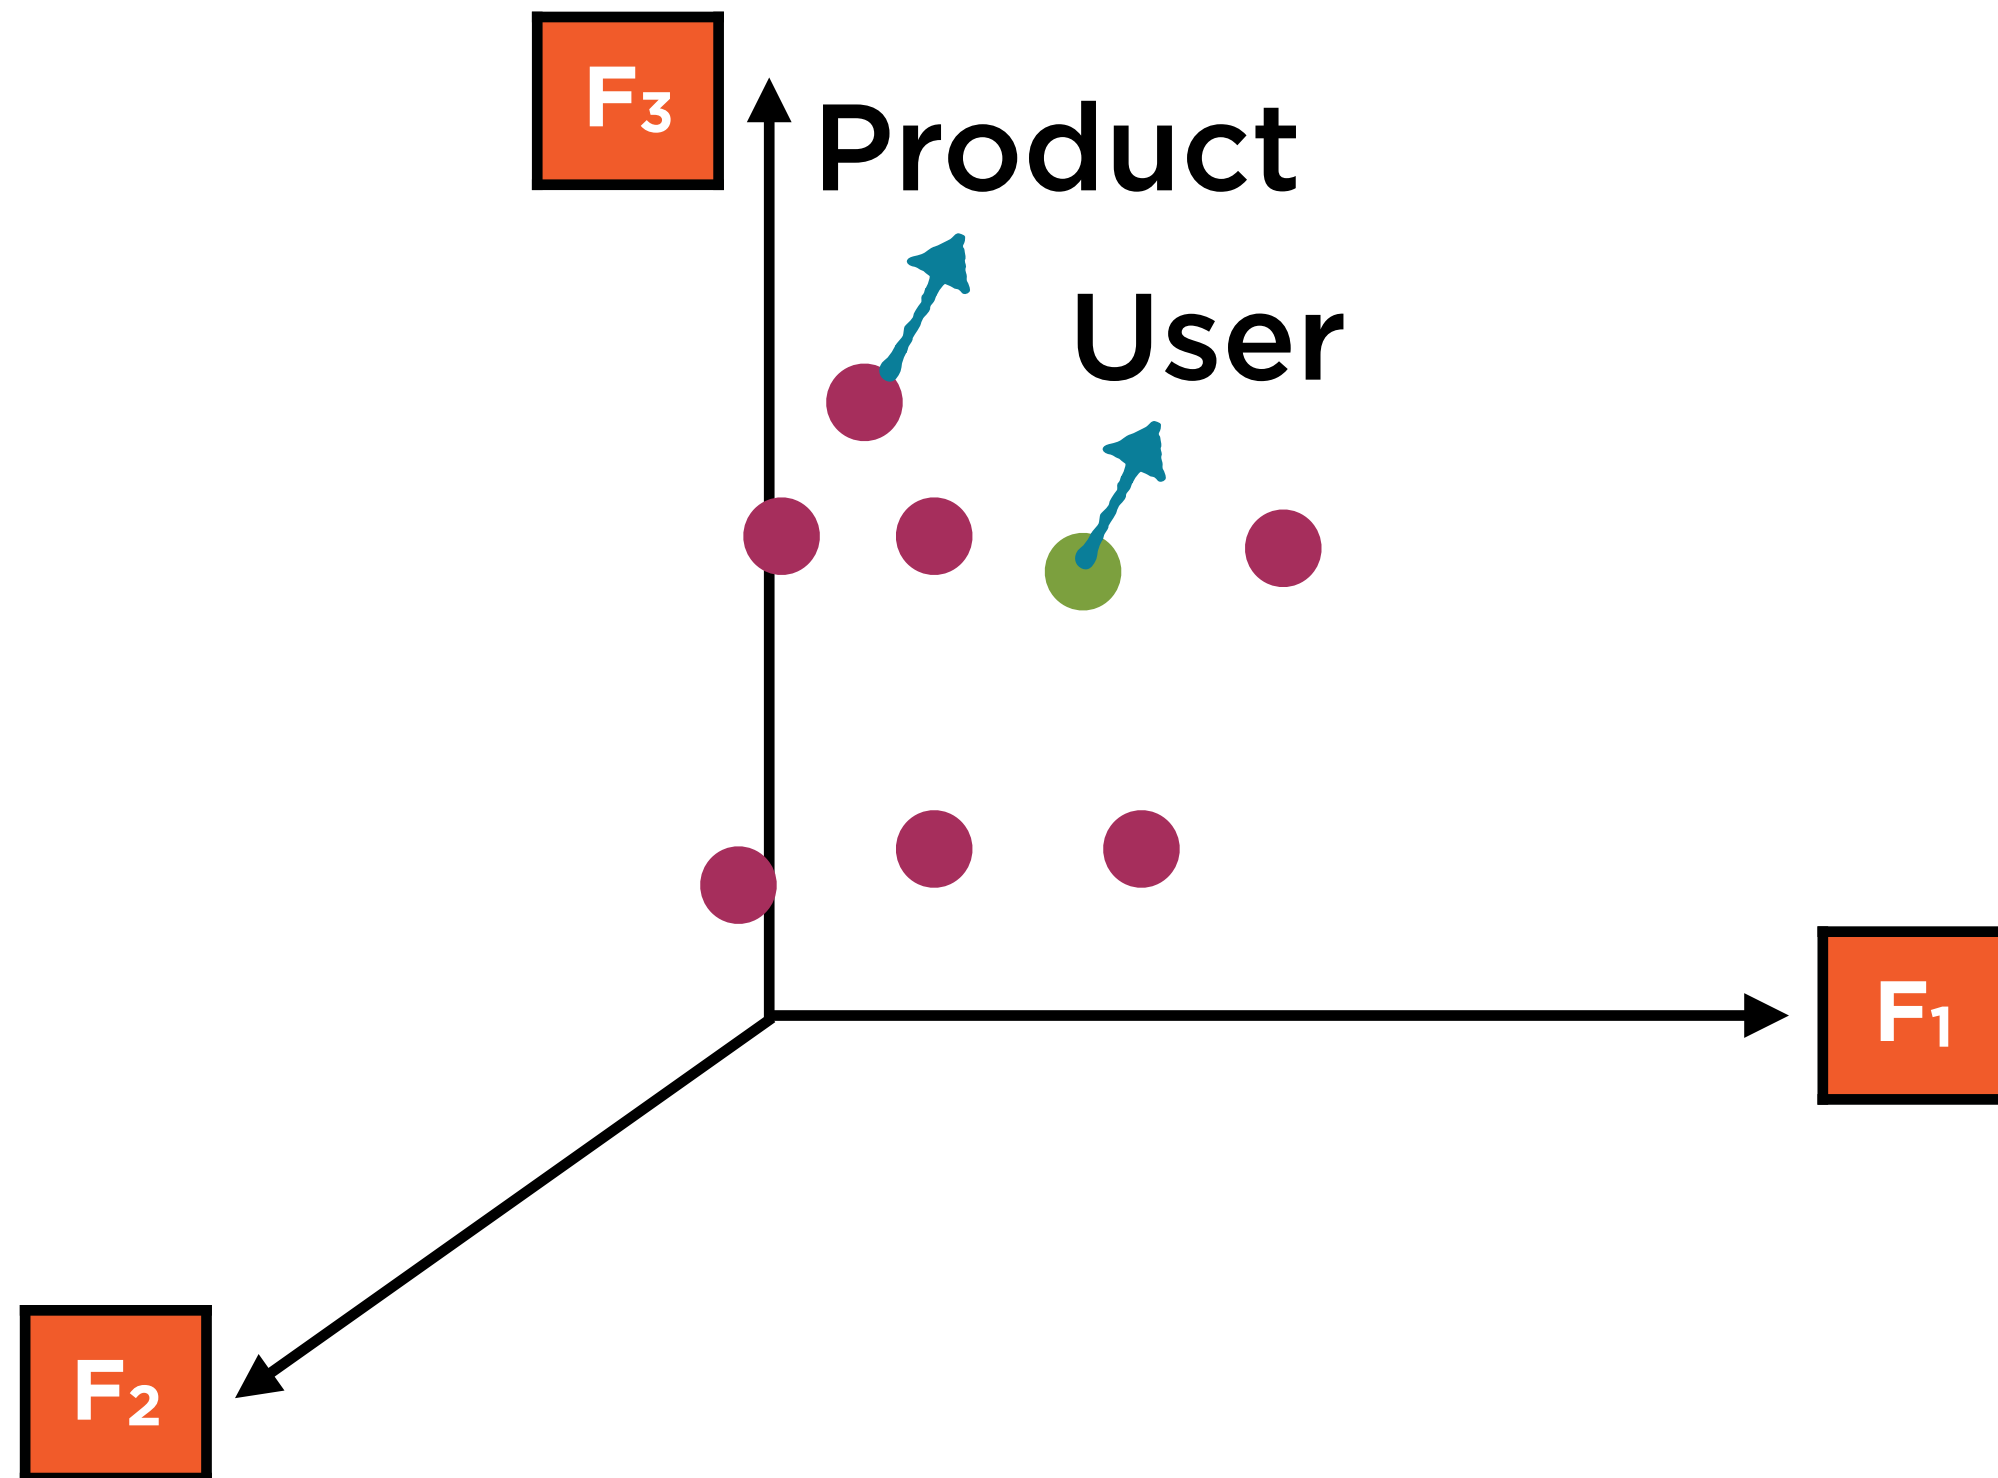
Product	F1	F2	F3	F4
A	0	3	2	5
B	5	2	3	4
C	4	5	2	1
D	3	4	5	2

**Rate the user on the
importance he/she
gives to these factors**

User	F1	F2	F3	F4
A	0	3	2	5

**Ex : Average of ratings of
products that the user
already likes**

Data Representation



Content Based Filtering vs Latent Factor Analysis




Factors are identified by experts

Factors are product attributes

Factors are derived using machine learning techniques

Factors may be related to product attributes or may be abstract



Contrasting the Nearest Neighbors Model and Latent Factor Analysis


Nearest Neighbors Model vs Latent Factor Analysis



Data Representation



Model Updates



Memory Usage

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation




Model Updates




Memory Usage


Nearest Neighbors Model



Latent Factor Analysis



Latent Factor Analysis



Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors



Model Updates



Memory Usage

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors




Model Updates




Memory Usage

Nearest Neighbors Model

Active user



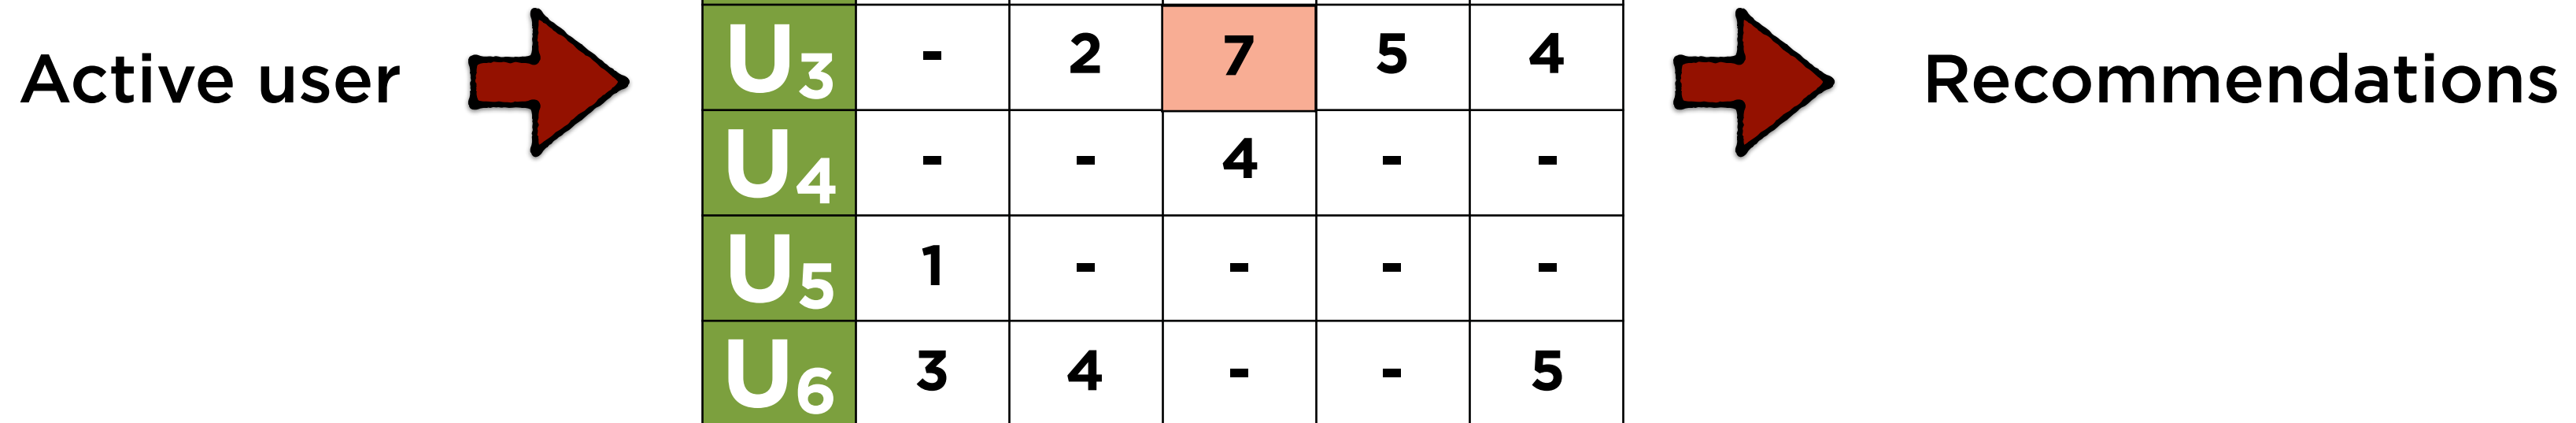
	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5



Recommendations


Nearest Neighbors Model

Online updates




Latent Factor Analysis

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5



	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			



	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					


Latent Factor Analysis

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5


	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			

*

	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					



Recommendations



Active user

Latent Factor Analysis


Offline updates

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	7	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5


	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			

*

	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					



Recommendations



Active user

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors



Model Updates

Online vs offline
updates



Memory Usage

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors



Model Updates

Online vs offline
updates



Memory Usage

Nearest Neighbors Model

	P ₁	P ₂	P ₃	P ₄	P ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5

Ratings kept in-memory at all times

User recommendations computed on-demand

Latent Factor Analysis

	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			

*

	P ₁	P ₂	P ₃	P ₄	P ₅
F ₁					
F ₂					
F ₃					

**Rating matrix is
decomposed offline**

**All recommendations can
easily be pre-computed
as a one off**

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors



Model Updates

Online vs offline
updates



Memory Usage

**In-memory vs pre-
computed**

Nearest Neighbors Model vs Latent Factor Analysis



Data Representation

Observable attributes vs
hidden driving factors



Model Updates

Online vs offline
updates



Memory Usage


In-memory vs pre-
computed

Decomposing the Rating Matrix

Using is Easy, Building is Hard



Driving a car



Building an engine

Solving for Factors

R

	I_1	I_2	I_3	I_4	I_5
U_1	3	4	-	-	-
U_2	3	2	-	-	5
U_3		r_{32}	-	5	4
U_4	-	-	4	-	-
U_5	1	-	-	-	-
U_6	3	4	-	-	5

P

	F_1	F_2	F_3
U_1			
U_2			
U_3	p_3		
U_4			
U_5			
U_6			

Q


	I_1	I_2	I_3	I_4	I_5
F_1		q_2			
F_2					
F_3					

*

Solving for Factors


R

	I ₁	I ₂	I ₃	I ₄	I ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5



P

	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			



Q

	I ₁	I ₂	I ₃	I ₄	I ₅
F ₁					
F ₂					
F ₃					

$$r_{ui} = p_u \cdot q_i$$



Solve the set of equations


$$r_{ui} = p_u \cdot q_i$$

Equations = # Known ratings

Solving for Factors


R

	I ₁	I ₂	I ₃	I ₄	I ₅
U ₁	3	4	-	-	-
U ₂	3	2	-	-	5
U ₃	-	2	-	5	4
U ₄	-	-	4	-	-
U ₅	1	-	-	-	-
U ₆	3	4	-	-	5



P

	F ₁	F ₂	F ₃
U ₁			
U ₂			
U ₃			
U ₄			
U ₅			
U ₆			



Q


	I ₁	I ₂	I ₃	I ₄	I ₅
F ₁					
F ₂					
F ₃					

$$r_{ui} = p_u \cdot q_i$$

Solving for Factors

R

	I₁	I₂	I₃	I₄	I₅
U₁	3	4	-	-	-
U₂	3	2	-	-	5
U₃	-	2	-	5	4
U₄	-	-	4	-	-
U₅	1	-	-	-	-
U₆	3	4	-	-	5



P

	F₁	F₂	F₃
U₁			
U₂			
U₃			
U₄			
U₅			
U₆			

Q

	I₁	I₂	I₃	I₄	I₅
F₁					
F₂					
F₃					

$$r_{ui} - p_u \cdot q_i = \epsilon_{ui}$$



Find p_u and q_i such that

$$\sum (r_{ui} - p_u \cdot q_i)^2$$

Total error is minimized



$$\min \sum (r_{ui} - p_u \cdot q_i)^2$$

Add a term to penalize
the model for the number
of factors



$$\min \sum (r_{ui} - p_u \cdot q_i)^2 + \lambda (||p_u||^2 + ||q_i||^2)$$

Add a term to penalize
the model for the number
of factors



$$\min \sum (r_{ui} - p_u \cdot q_i)^2 + \lambda (||p_u||^2 + ||q_i||^2)$$

Regularization term



$$\min \sum (r_{ui} - p_u \cdot q_i)^2 + \lambda (||p_u||^2 + ||q_i||^2)$$

Regularization factor



$$\min E(p_u, q_i)$$


A standard optimization problem

Demo

Solve for latent factors using the Stochastic Gradient Descent method (SGD)

- Set up a function to compute error
- Set up a function to implement SGD

Solving for Latent Factors



Solving for Latent Factors

Set up the data


Functions to access
relevant information

Construct a rating matrix

The representation needed
for collaborative filtering

Setting Up the Rating Matrix

User	ISBN	Rating




	I ₁	I ₂	I ₃	I ₄	I ₅
U ₁	3	4	-	-	4
U ₂	3	5	3	4	5
U ₃	4	2	-	5	4
U ₄	3	-	4	5	2
U ₅	1	-	4	2	1
U ₆	3	4	-	2	5

`pandas.pivot_table`

Setting Up the Rating Matrix

User	ISBN	Rating



	I ₁	I ₂	I ₃	I ₄	I ₅
U ₁	3	4	-	-	4
U ₂	3	5	3	4	5
U ₃	4	2	-	5	4
U ₄	3	-	4	5	2
U ₅	1	-	4	2	1
U ₆	3	4	-	2	5

`scipy.coo_matrix`

```
coo_matrix((values, (rowsource, columnsource)))
```

Creating a Rating Matrix

Takes a tuple with rating data, row names and column names

User	ISBN	Rating



```
coo_matrix((values, (rowsource, columnsource)))
```

Creating a Rating Matrix

Takes a tuple with rating data, row names and column names



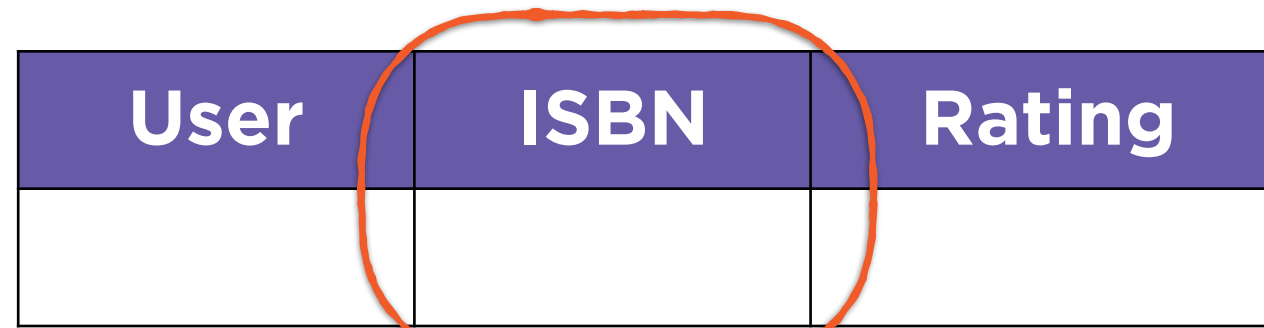
User	ISBN	Rating

```
coo_matrix((values, (rowsource, columnsource)))
```

Creating a Rating Matrix

Takes a tuple with rating data, row names and column names

User	ISBN	Rating



```
coo_matrix((values, (rowsource, columnsource)))
```

Creating a Rating Matrix

Takes a tuple with rating data, row names and column names

	..	3053
..	4		-	-
..	2		-	5
10633	2	0	5	4
..	-	4	-	-
..	-	-	-	-
..	4	-	-	5



Solving for Latent Factors

Set up the data

Functions to access
relevant information

Initialize factor matrices

Construct a rating matrix

The representation needed
for collaborative filtering

Solving for Latent Factors

Set up the data

Functions to access
relevant information

Initialize factor matrices

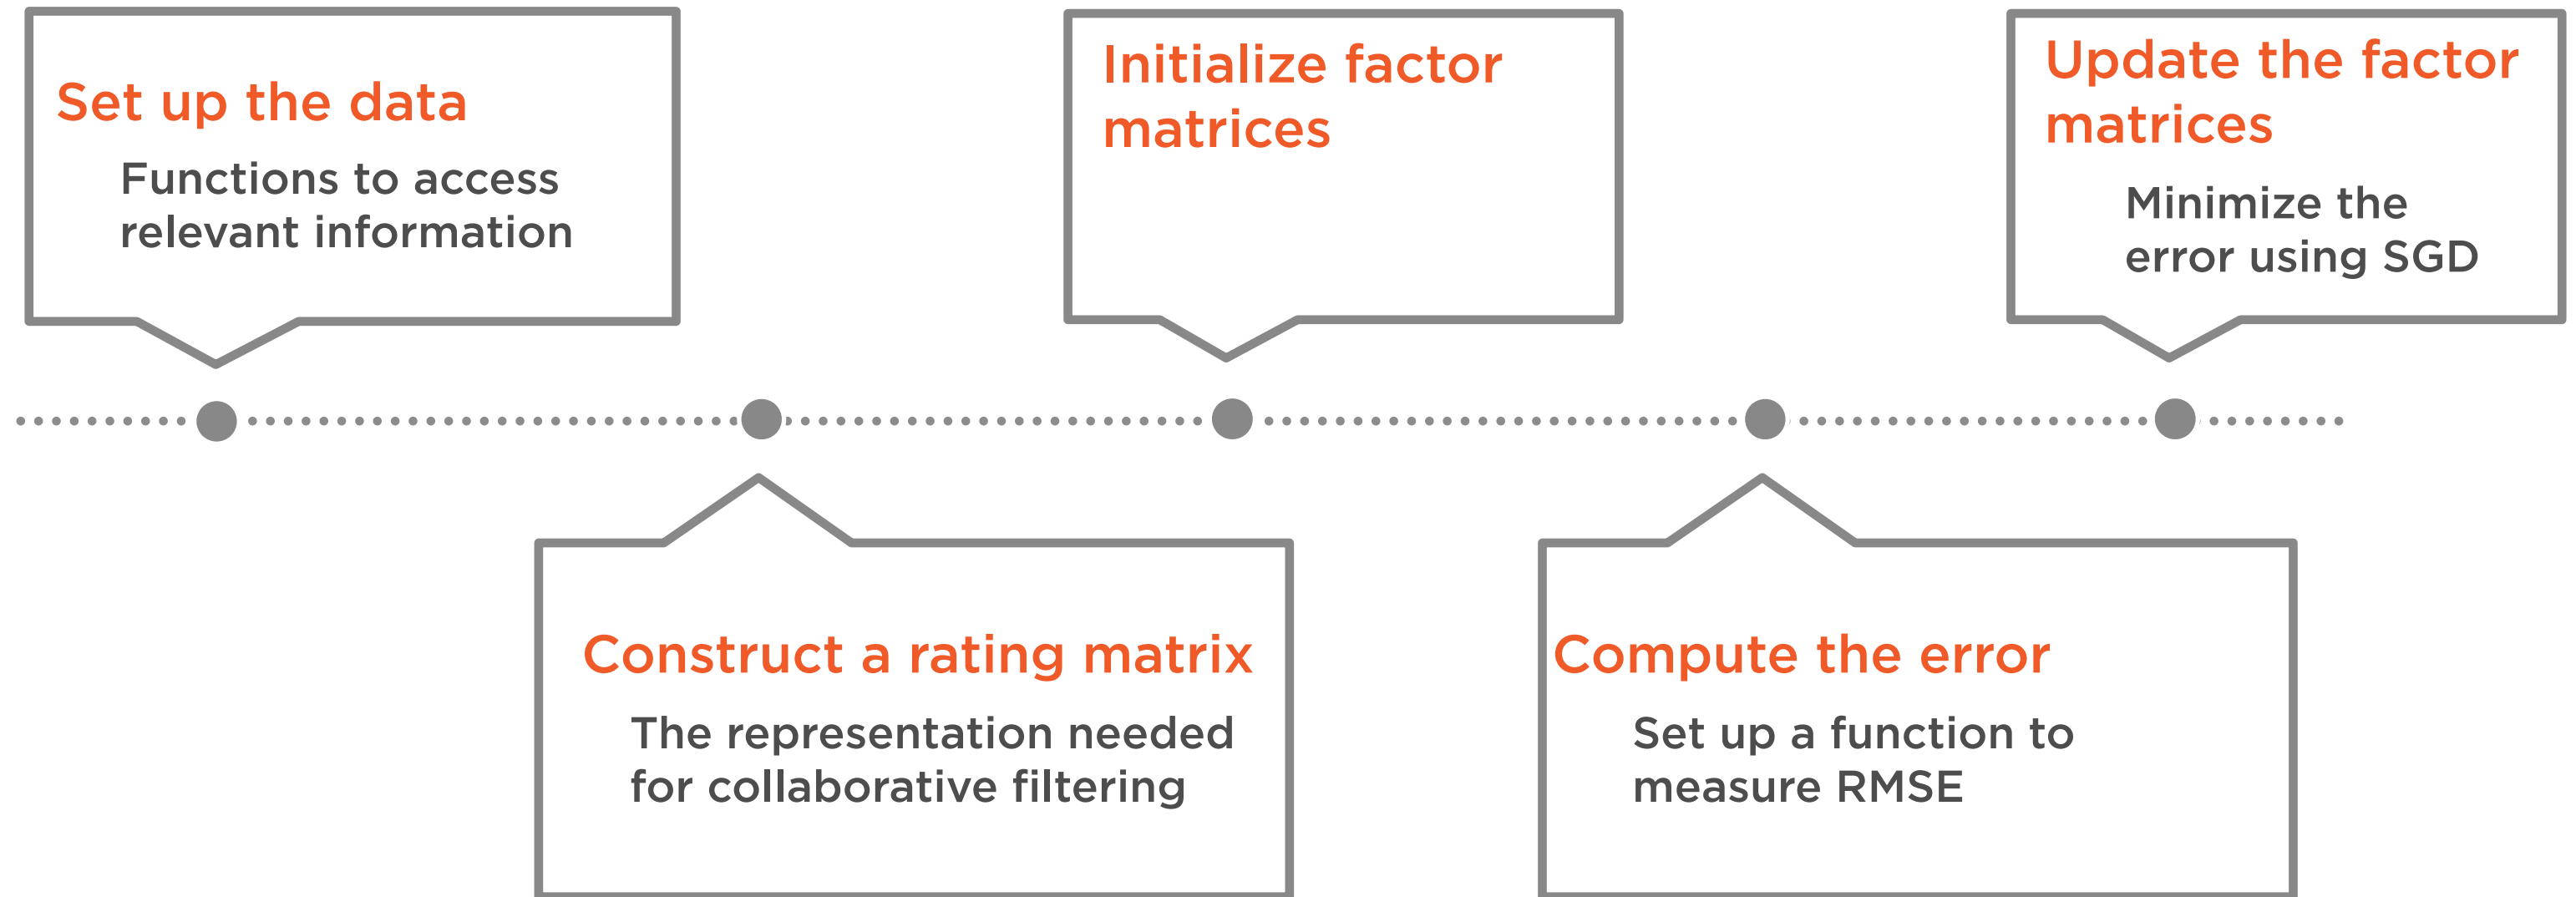
Construct a rating matrix

The representation needed
for collaborative filtering

Compute the error

Set up a function to
measure RMSE

Solving for Latent Factors





Stochastic Gradient
Descent

$$\sum (r_{ui} - p_u \cdot q_i)^2 + \lambda (||p_u||^2 + ||q_i||^2)$$

Total error across all ratings



Stochastic Gradient
Descent

$$\sum (r_{ui} - p_u \cdot q_i)^2 + \lambda (||p_u||^2 + ||q_i||^2)$$

Look at the error for 1 rating




Stochastic Gradient
Descent

$$e(p_u, q_i) = (r_{ui} - p_u \cdot q_i)^2 + \lambda(||p_u||^2 + ||q_i||^2)$$

Look at the error for 1 rating




Stochastic Gradient
Descent






Stochastic Gradient
Descent



Compute the initial value of
the error




Stochastic Gradient
Descent



Compute the slope at that
point




Stochastic Gradient
Descent



Take a small step in the
downward direction of the
slope




Stochastic Gradient
Descent



Take a small step i.e. update
 p_u and q_i by a small value




Stochastic Gradient
Descent



Repeat until you reach a
minimum




Stochastic Gradient
Descent



Take a small step i.e. update
 p_u and q_i by a small value




Stochastic Gradient
Descent






Stochastic Gradient
Descent






Stochastic Gradient
Descent





Stochastic Gradient
Descent

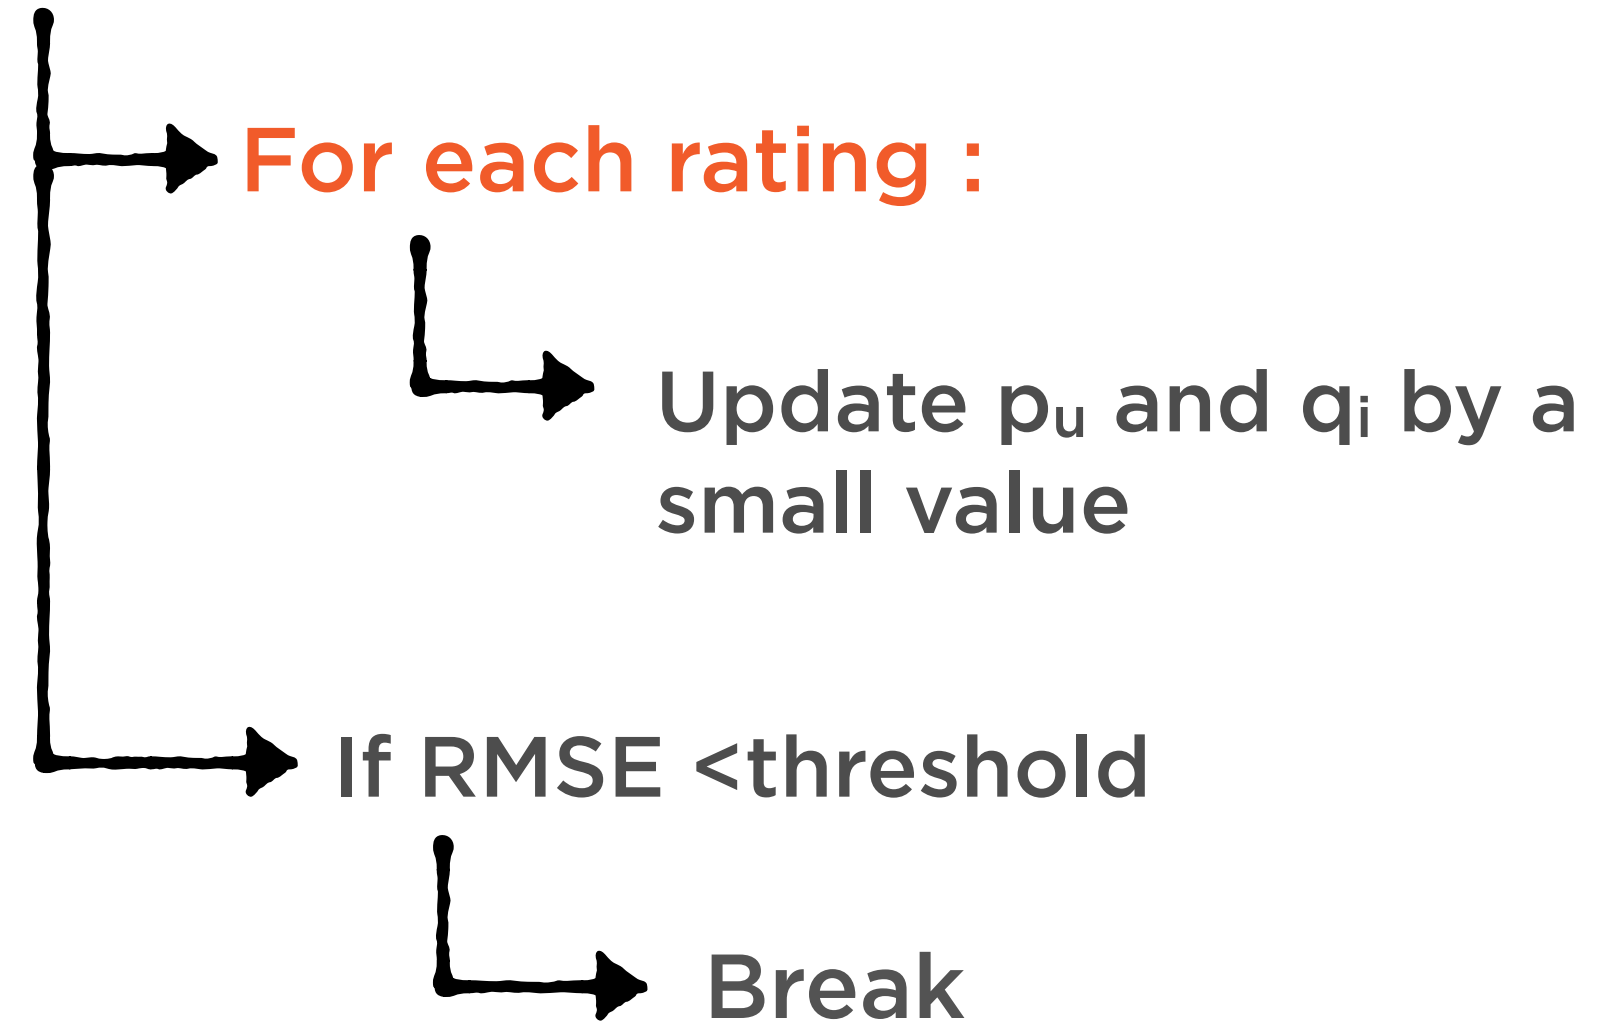


$$q_i = q_i + \gamma \left(- \frac{\partial E(p_u, q_i)}{\partial q_i} \right)$$




Stochastic Gradient Descent

While max # allowed
steps is not reached :





Stochastic Gradient Descent



$$p_u = p_u + \gamma \left(- \frac{\partial E(p_u, q_i)}{\partial p_u} \right)$$

$$q_i = q_i + \gamma \left(- \frac{\partial E(p_u, q_i)}{\partial q_i} \right)$$

```
eui=rui-np.dot(P[u,:],Q[:,i])  
P[u,:]=P[u,:]+gamma*2*(eui*Q[:,i]-lamda*P[u,:])  
Q[:,i]=Q[:,i]+gamma*2*(eui*P[u,:]-lamda*Q[:,i])
```

Updating p_u and q_i

$$\epsilon_{ui} = r_{ui} - \mathbf{p}_u \cdot \mathbf{q}_i$$

```
eui=rui-np.dot(P[u,:],Q[:,i])
```

```
P[u,:]=P[u,:]+gamma*2*(eui*Q[:,i]-lamda*P[u,:])
```

```
Q[:,i]=Q[:,i]+gamma*2*(eui*P[u,:]-lamda*Q[:,i])
```

Updating \mathbf{p}_u and \mathbf{q}_i

$$p_u = p_u + \gamma \left(\frac{-\partial E(p_u, q_i)}{\partial p_u} \right)$$

```
eui=rui-np.dot(P[u, :], Q[:, i])
```

```
P[u, :]=P[u, :]+gamma*2*(eui*Q[:, i]-lamda*P[u, :])
```

```
Q[:, i]=Q[:, i]+gamma*2*(eui*P[u, :]-lamda*Q[:, i])
```

Updating p_u and q_i

$$p_u = p_u + \gamma (2(\epsilon_{ui}q_i - \lambda p_u))$$

```
eui=rui-np.dot(P[u,:],Q[:,i])
```

```
P[u,:]=P[u,:]+gamma*(2*(eui*Q[:,i]-lamda*P[u,:]))
```

```
Q[:,i]=Q[:,i]+gamma*2*(eui*P[u,:]-lamda*Q[:,i])
```

Updating p_u and q_i

$$q_i = q_i + \gamma \left(- \frac{\partial E(p_u, q_i)}{\partial q_i} \right)$$

```
eui=rui-np.dot(P[u,:],Q[:,i])
```

```
P[u,:]=P[u,:]+gamma*2*(eui*Q[:,i]-lamda*P[u,:])
```

```
Q[:,i]=Q[:,i]+gamma*2*(eui*P[u,:]-lamda*Q[:,i])
```

Updating p_u and q_i

$$q_i = q_i + \gamma (2(\epsilon_{ui}p_u - \lambda q_i))$$

```
eui=rui-np.dot(P[u,:],Q[:,i])  
P[u,:]=P[u,:]+gamma*2*(eui*Q[:,i]-lamda*P[u,:])  
Q[:,i]=Q[:,i]+gamma*2*(eui*P[u,:]-lamda*Q[:,i])
```

Updating p_u and q_i

Summary

Understand the latent factors model for collaborative filtering

Contrast the latent factors model and the nearest neighbors model

Use optimization techniques to solve for latent factors

- Stochastic gradient descent