

Building Regression Models using TensorFlow

LEARNING USING NEURONS



Vitthal Srinivasan

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

The most common use of TensorFlow is in deep learning

Neural networks are the most common type of deep learning algorithm

The basic building block of a neural network is a neuron

Linear regression can be “learnt” using a single neuron

Deep learning extends this idea to more complex, non-linear functions

Prerequisites and Course Outline

Course Outline



- Learning using Neurons**
- Linear Regression in TensorFlow**
- Logistic Regression in TensorFlow**
- Estimators**

Related Courses

Understanding the Foundations of TensorFlow

Understanding and Applying Linear Regression

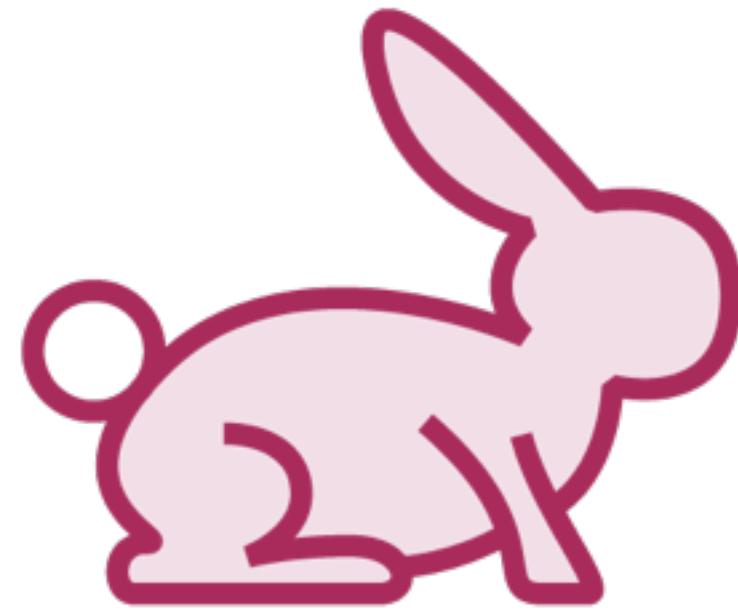
Understanding and Applying Logistic Regression

Software and Skills

Have TensorFlow installed
Know some Python programming

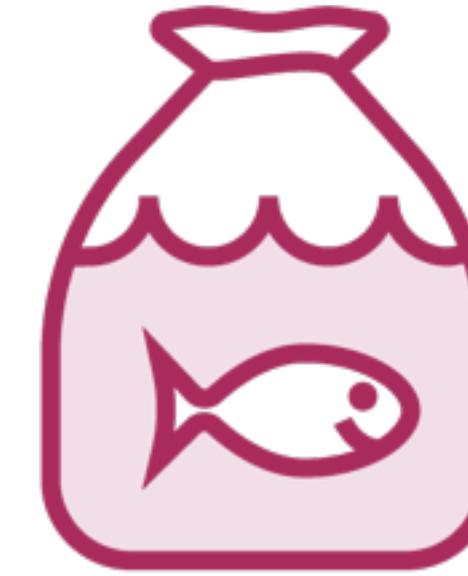
Understanding Deep Learning

Whales: Fish or Mammals?



Mammals

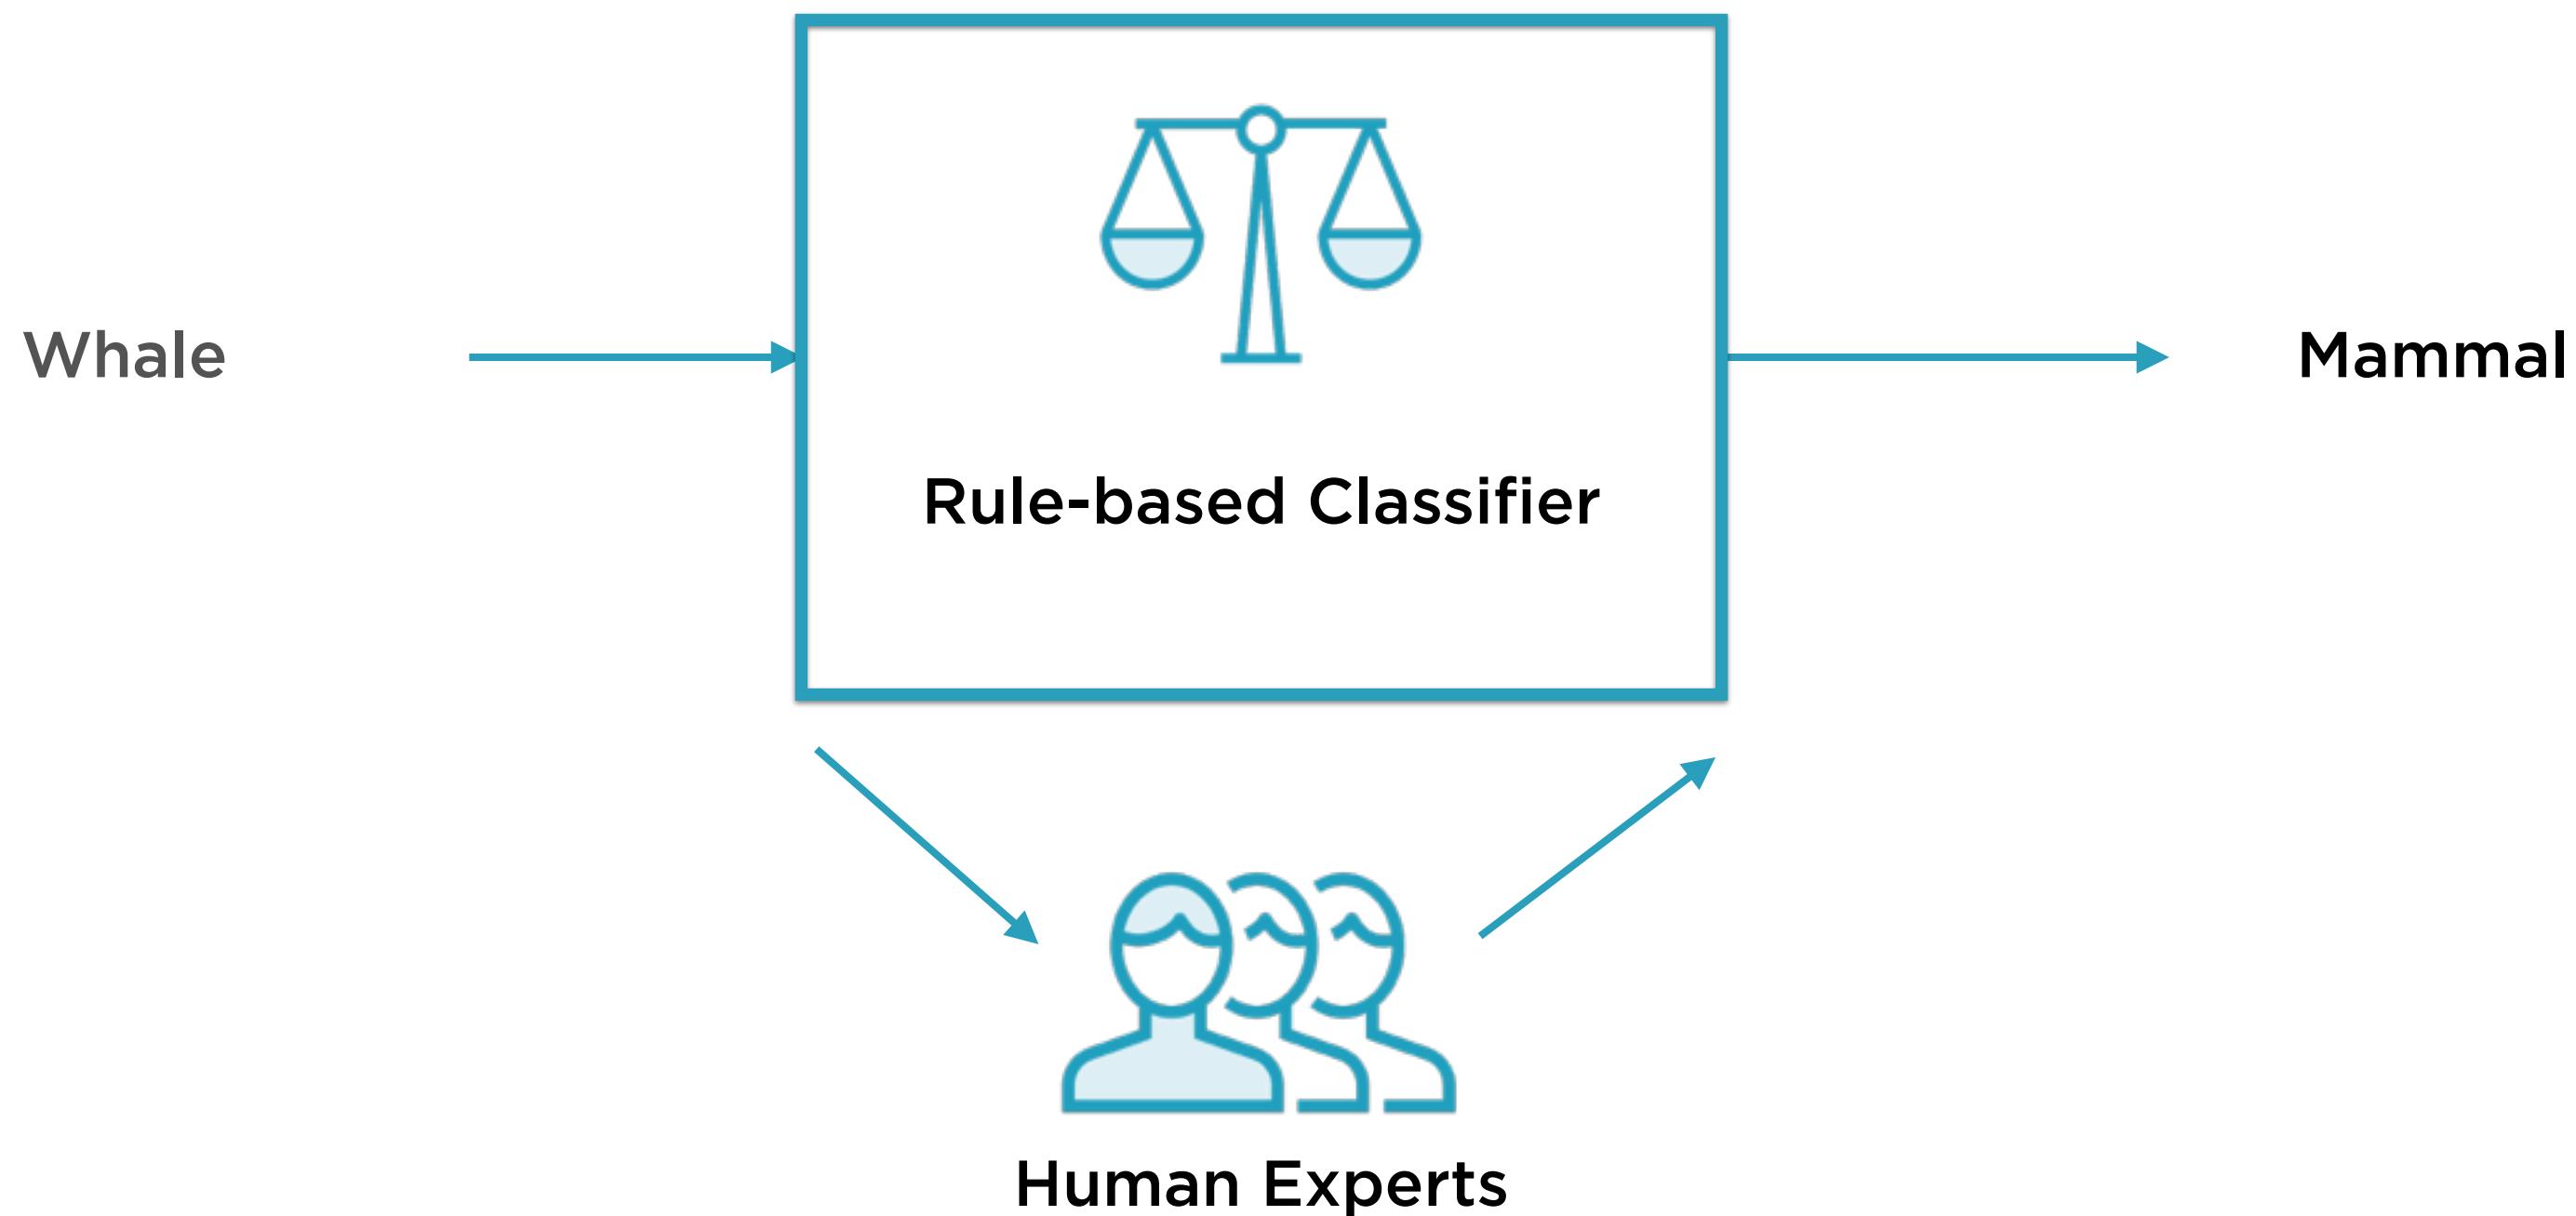
**Members of the infraorder
*Cetacea***



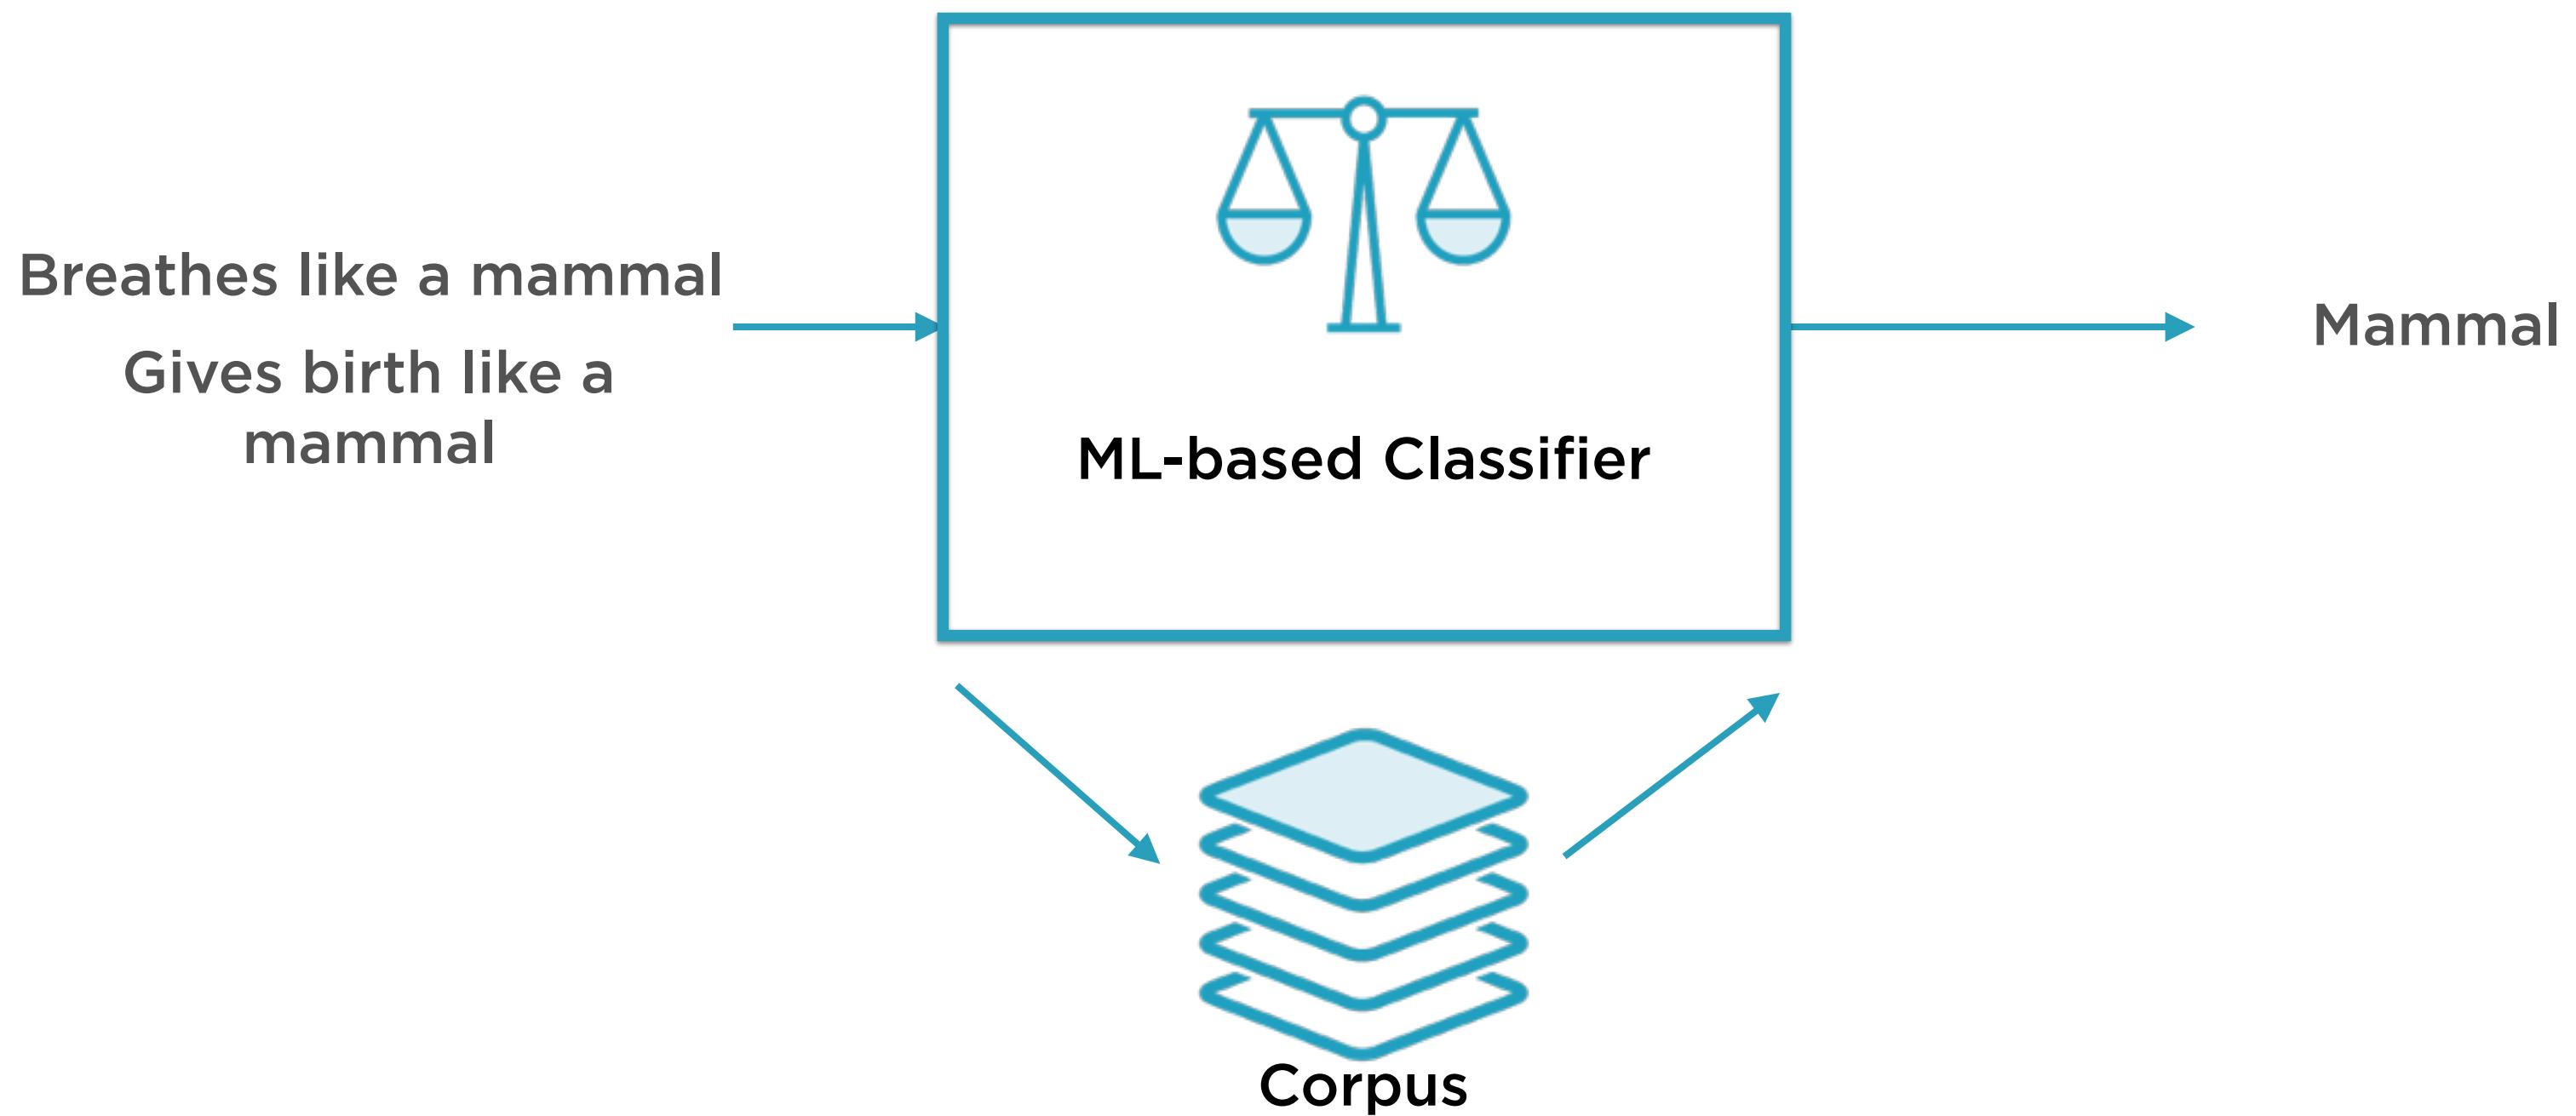
Fish

**Look like fish, swim like fish,
move with fish**

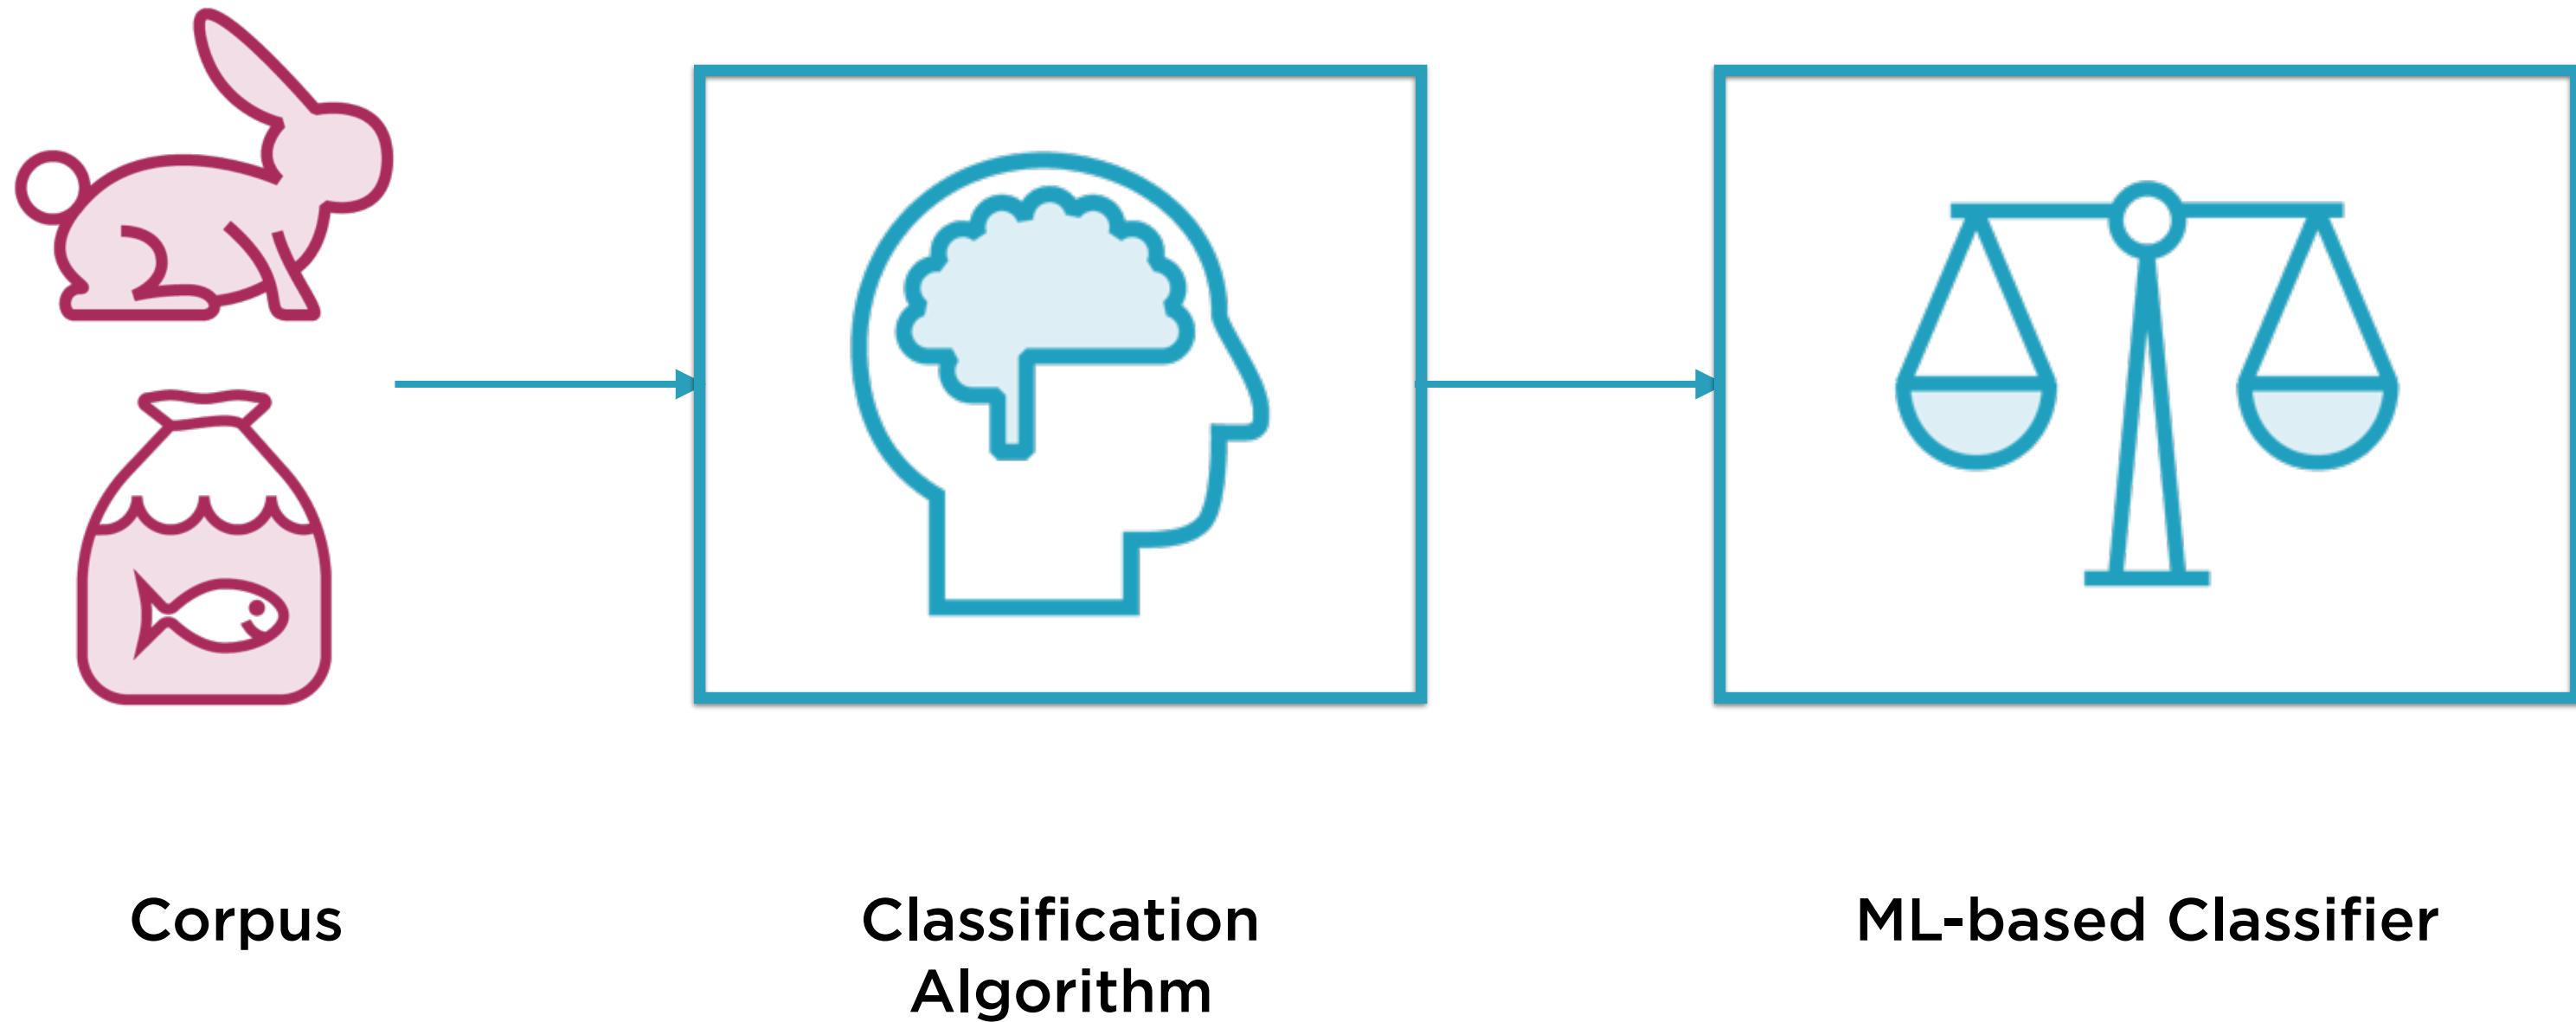
Rule-based Binary Classifier



“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



“Traditional” ML-based vs. Rule-based

“Traditional” ML-based

Dynamic

Experts optional

Corpus required

Training step

Rule-based

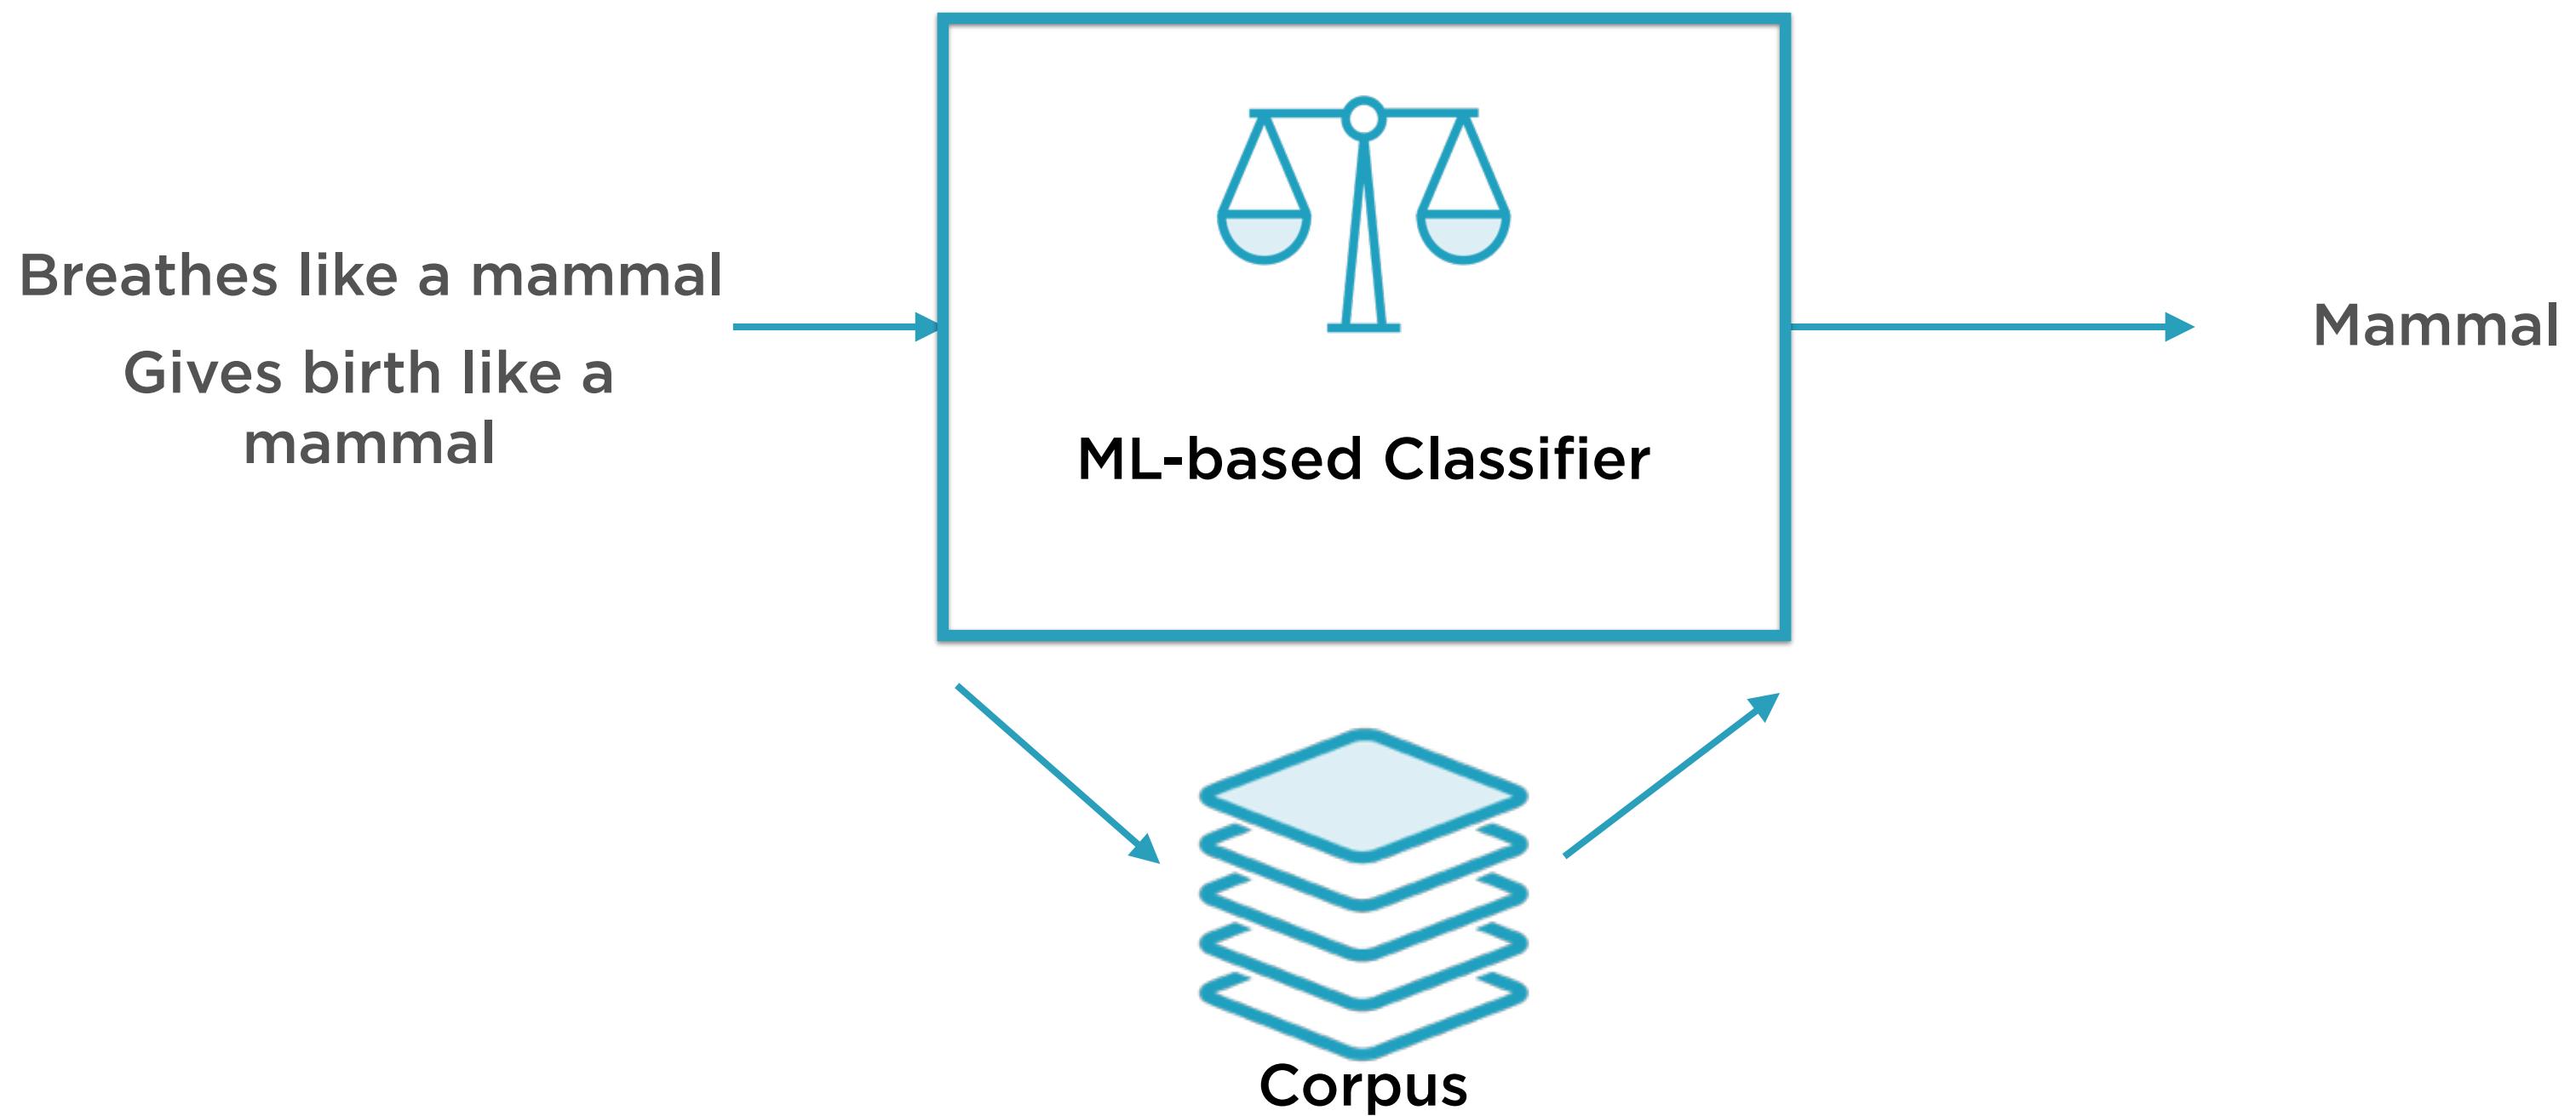
Static

Experts required

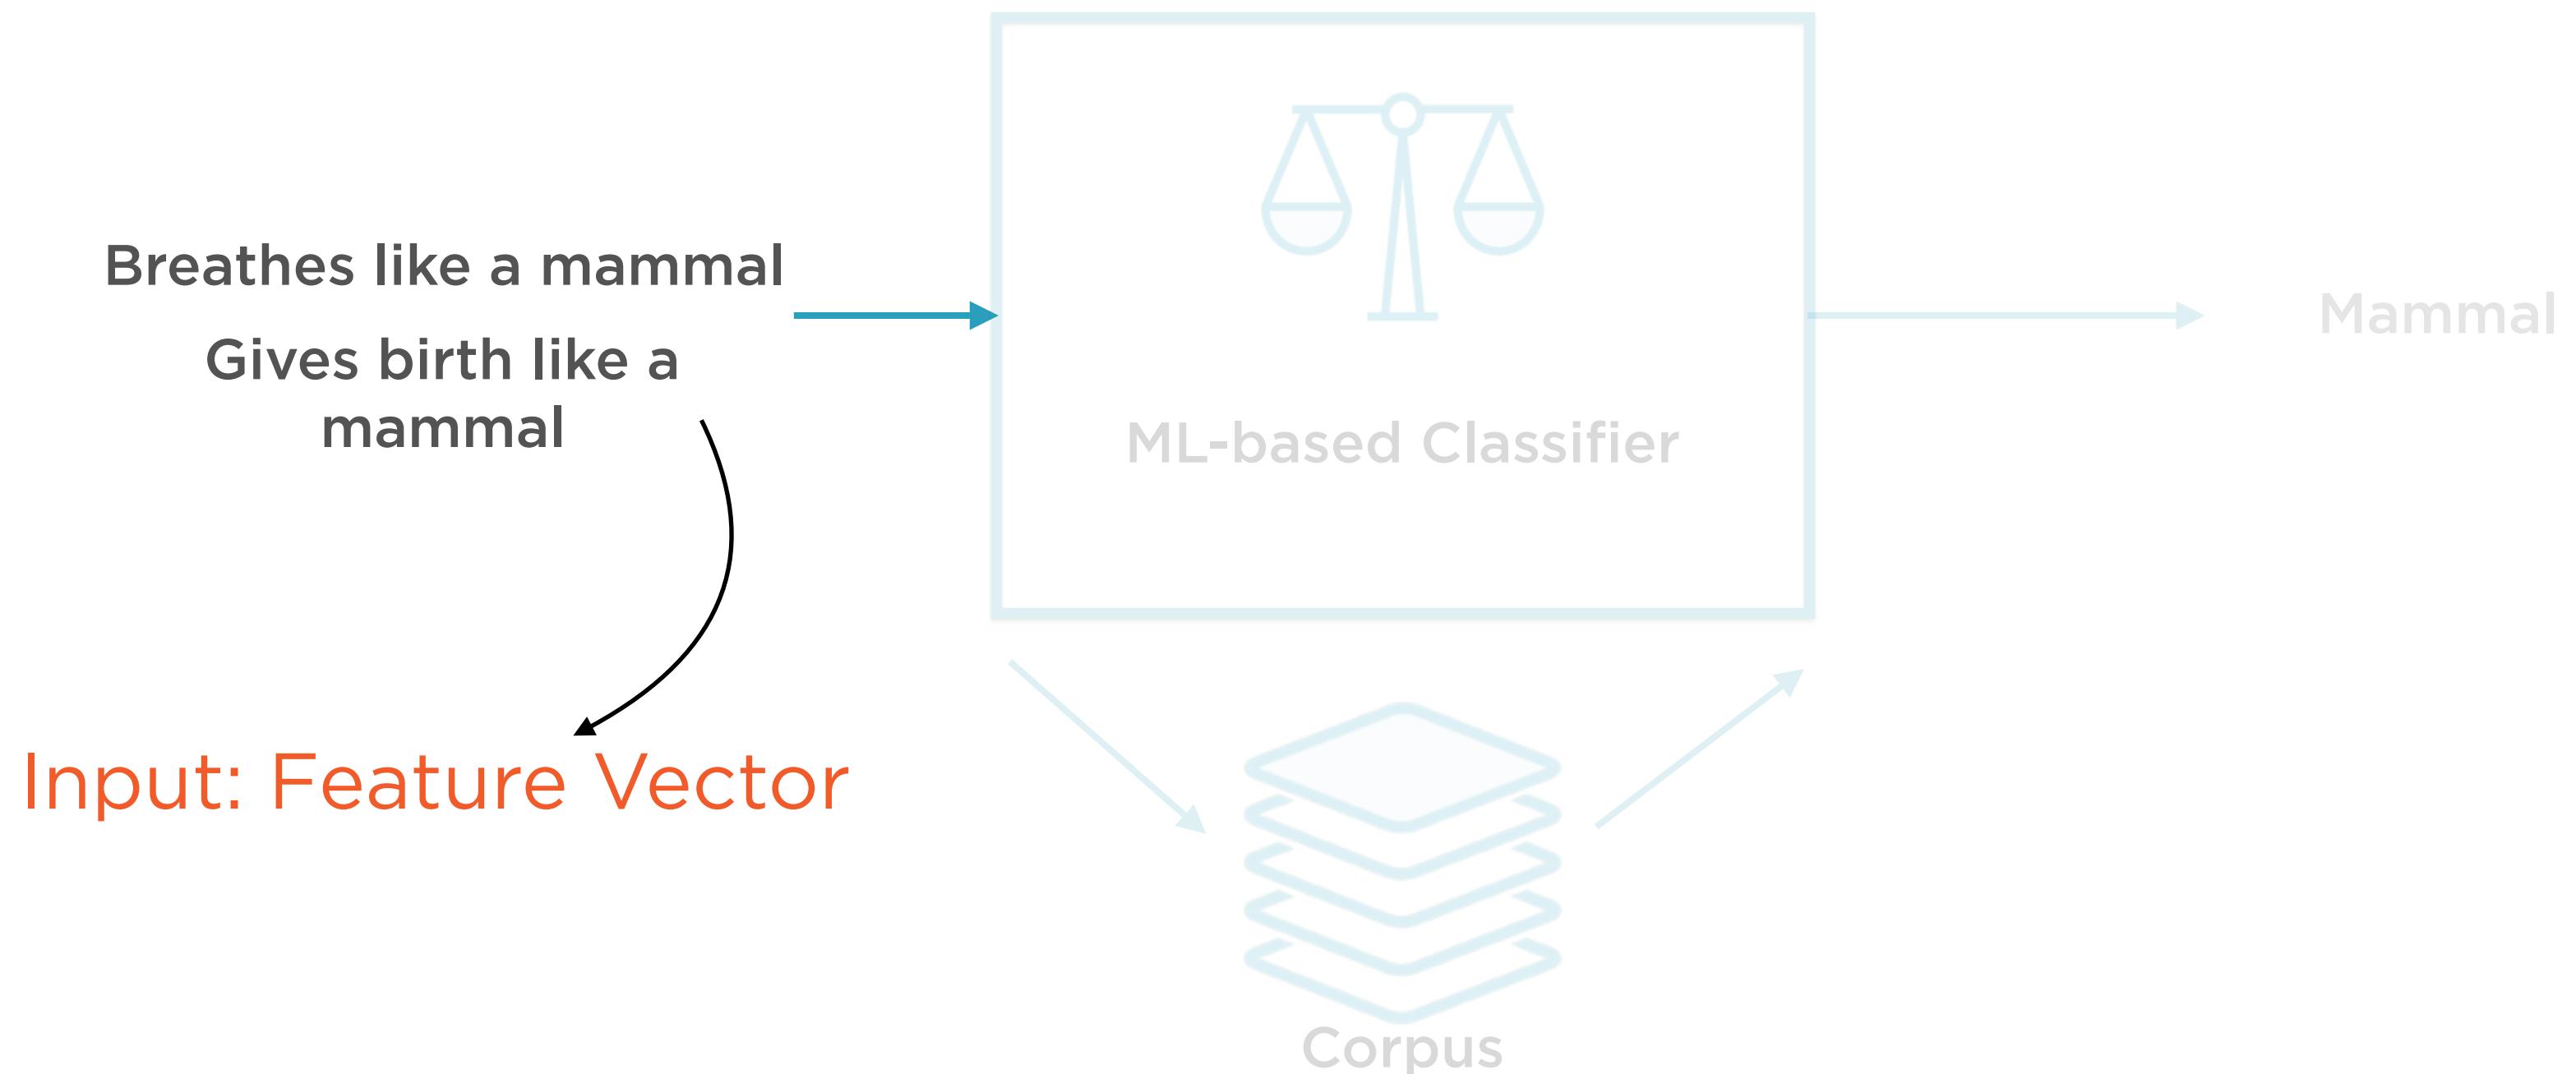
Corpus optional

No training step

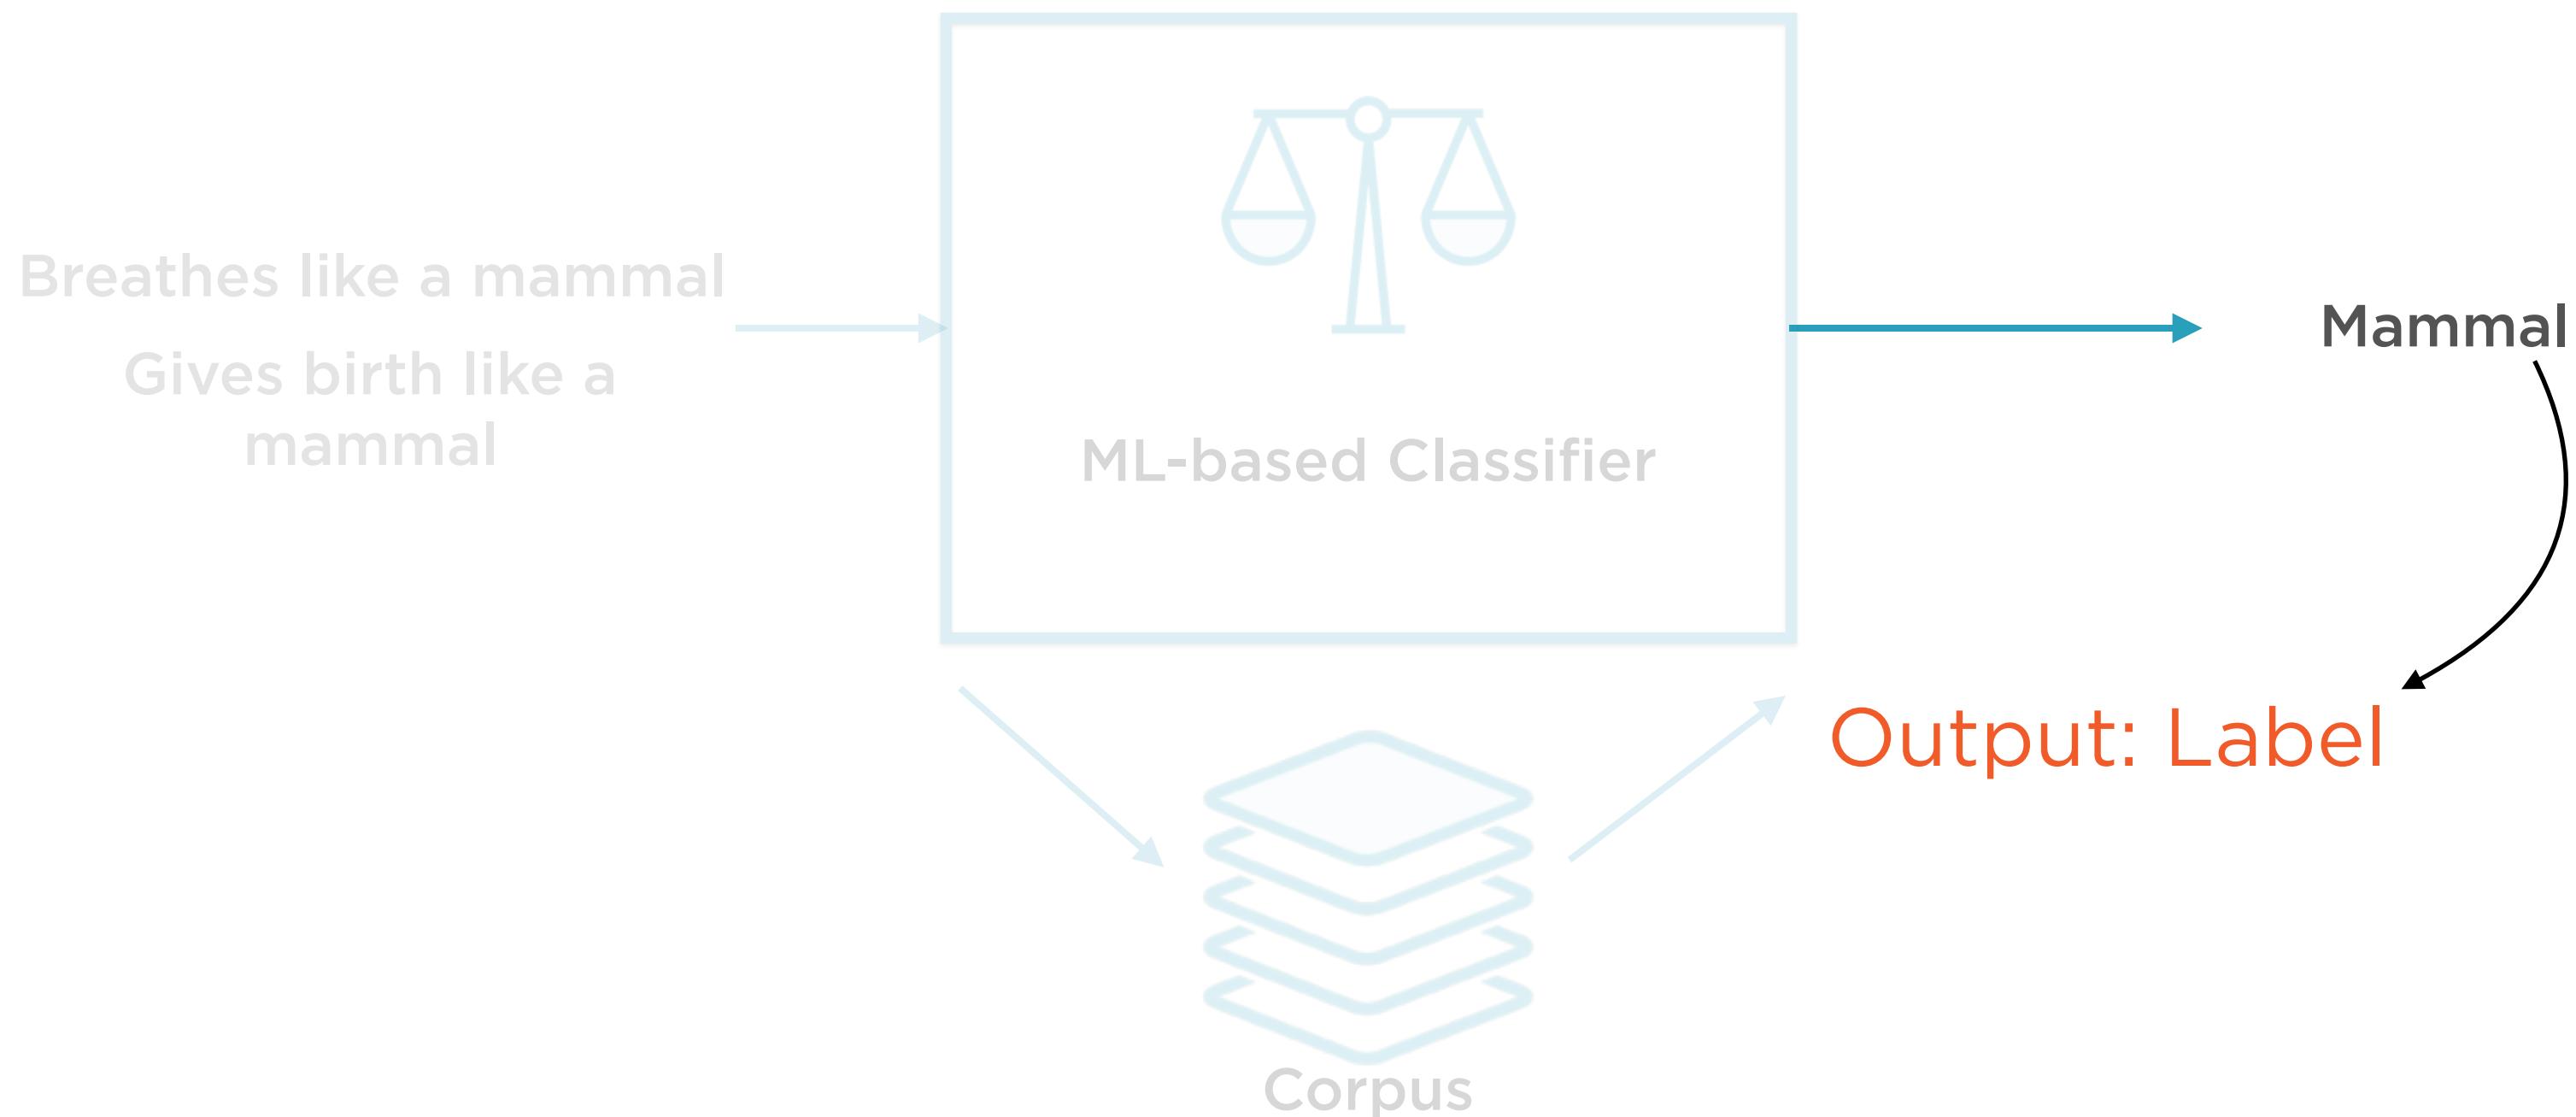
“Traditional” ML-based Binary Classifier



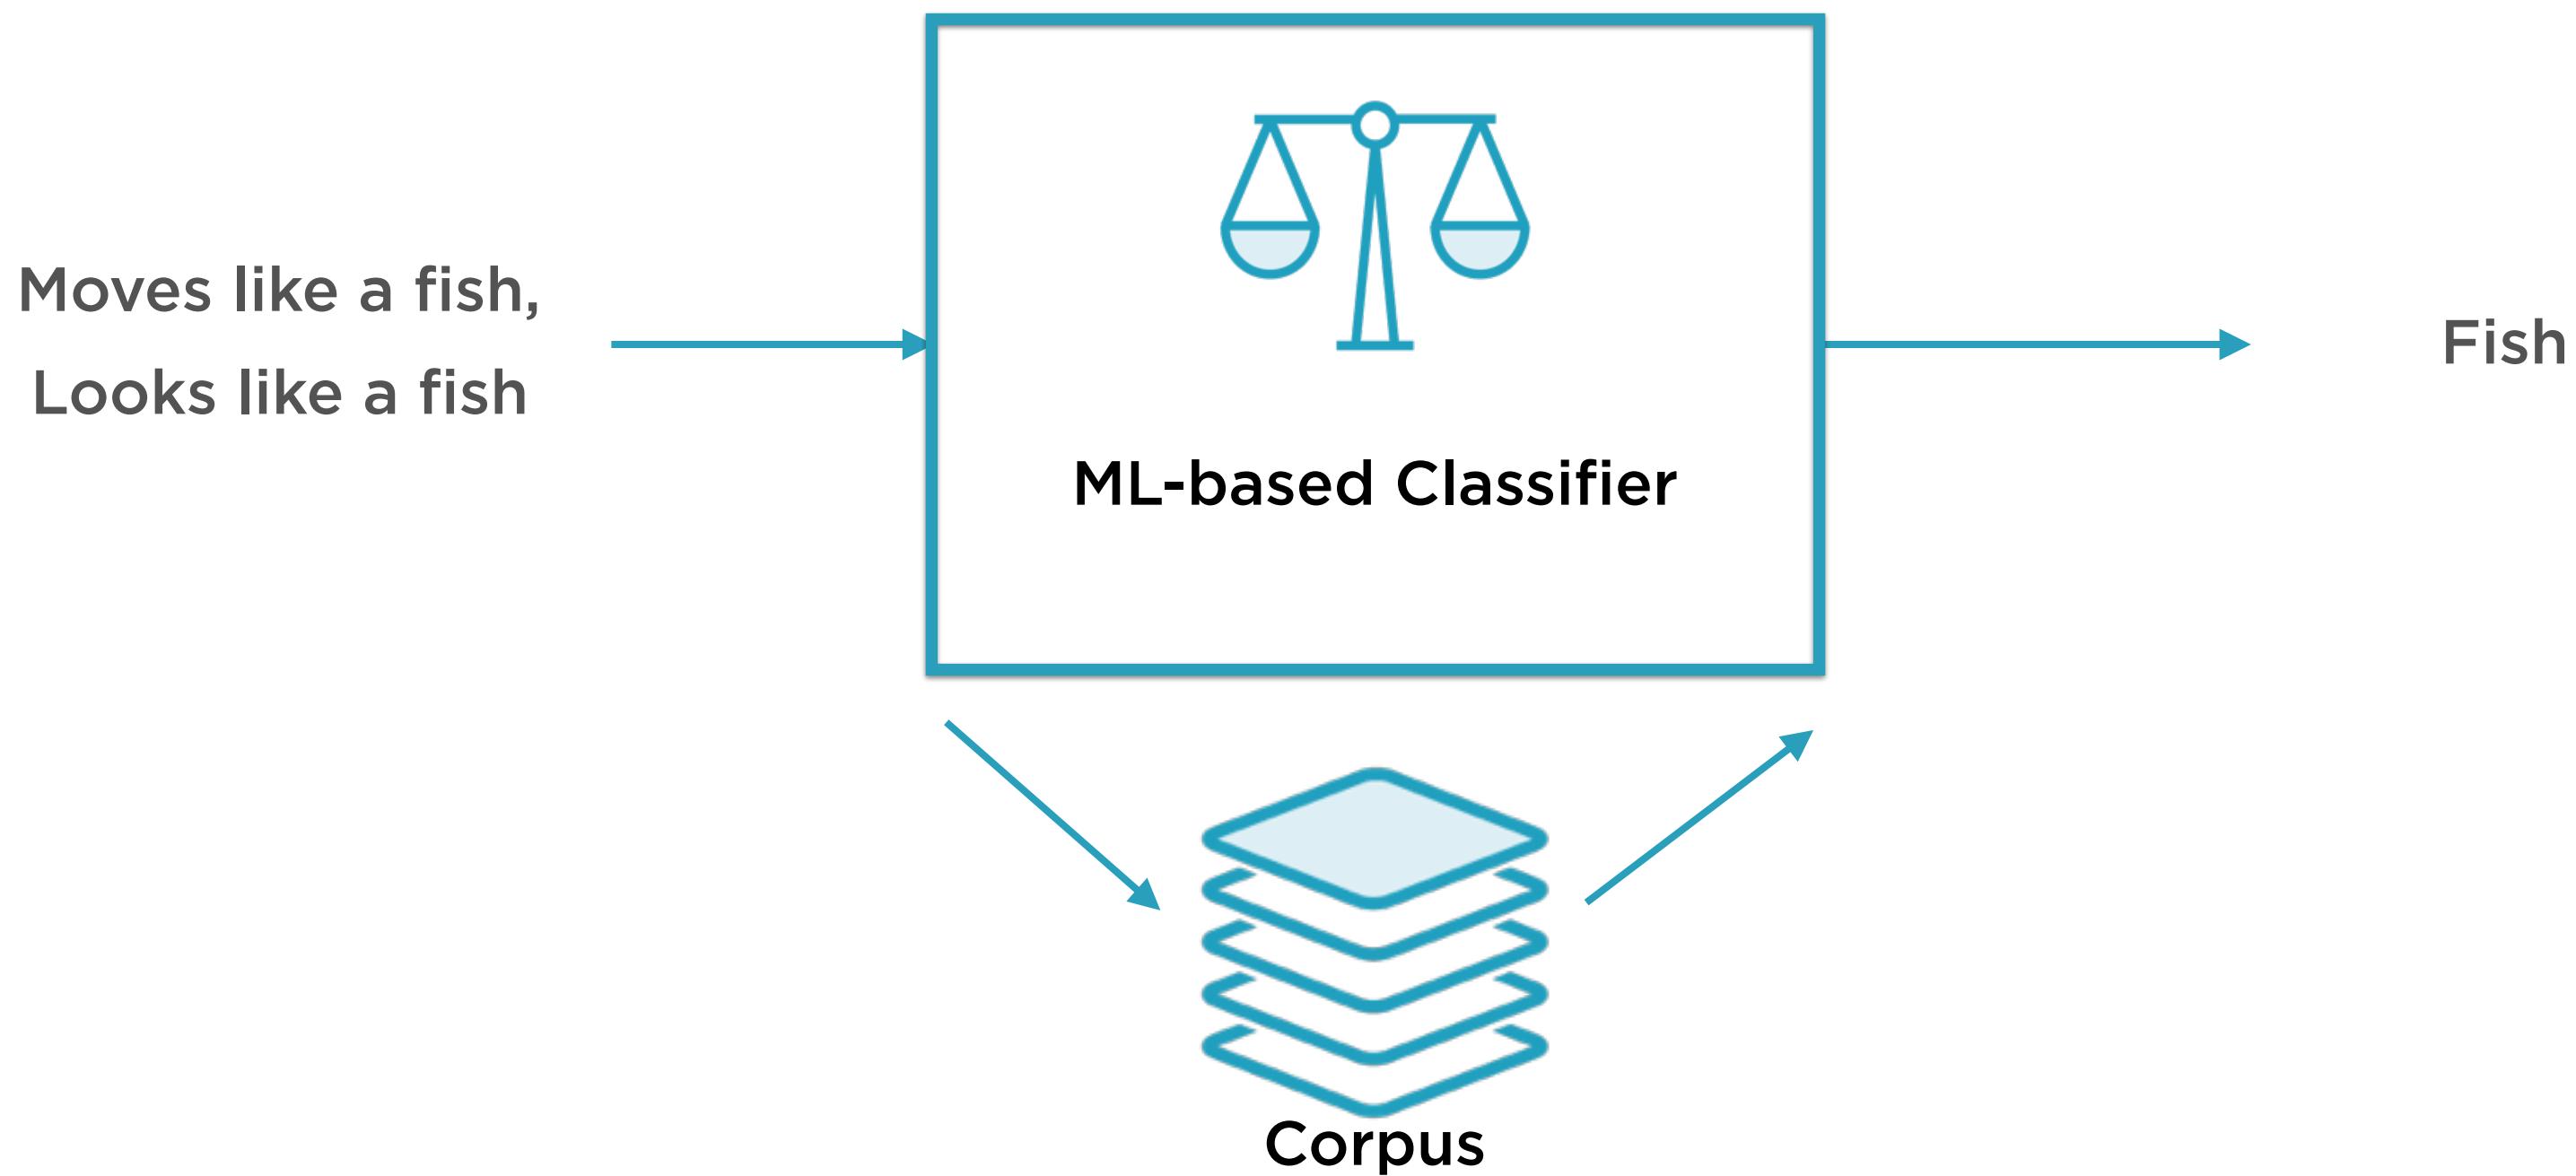
“Traditional” ML-based Binary Classifier



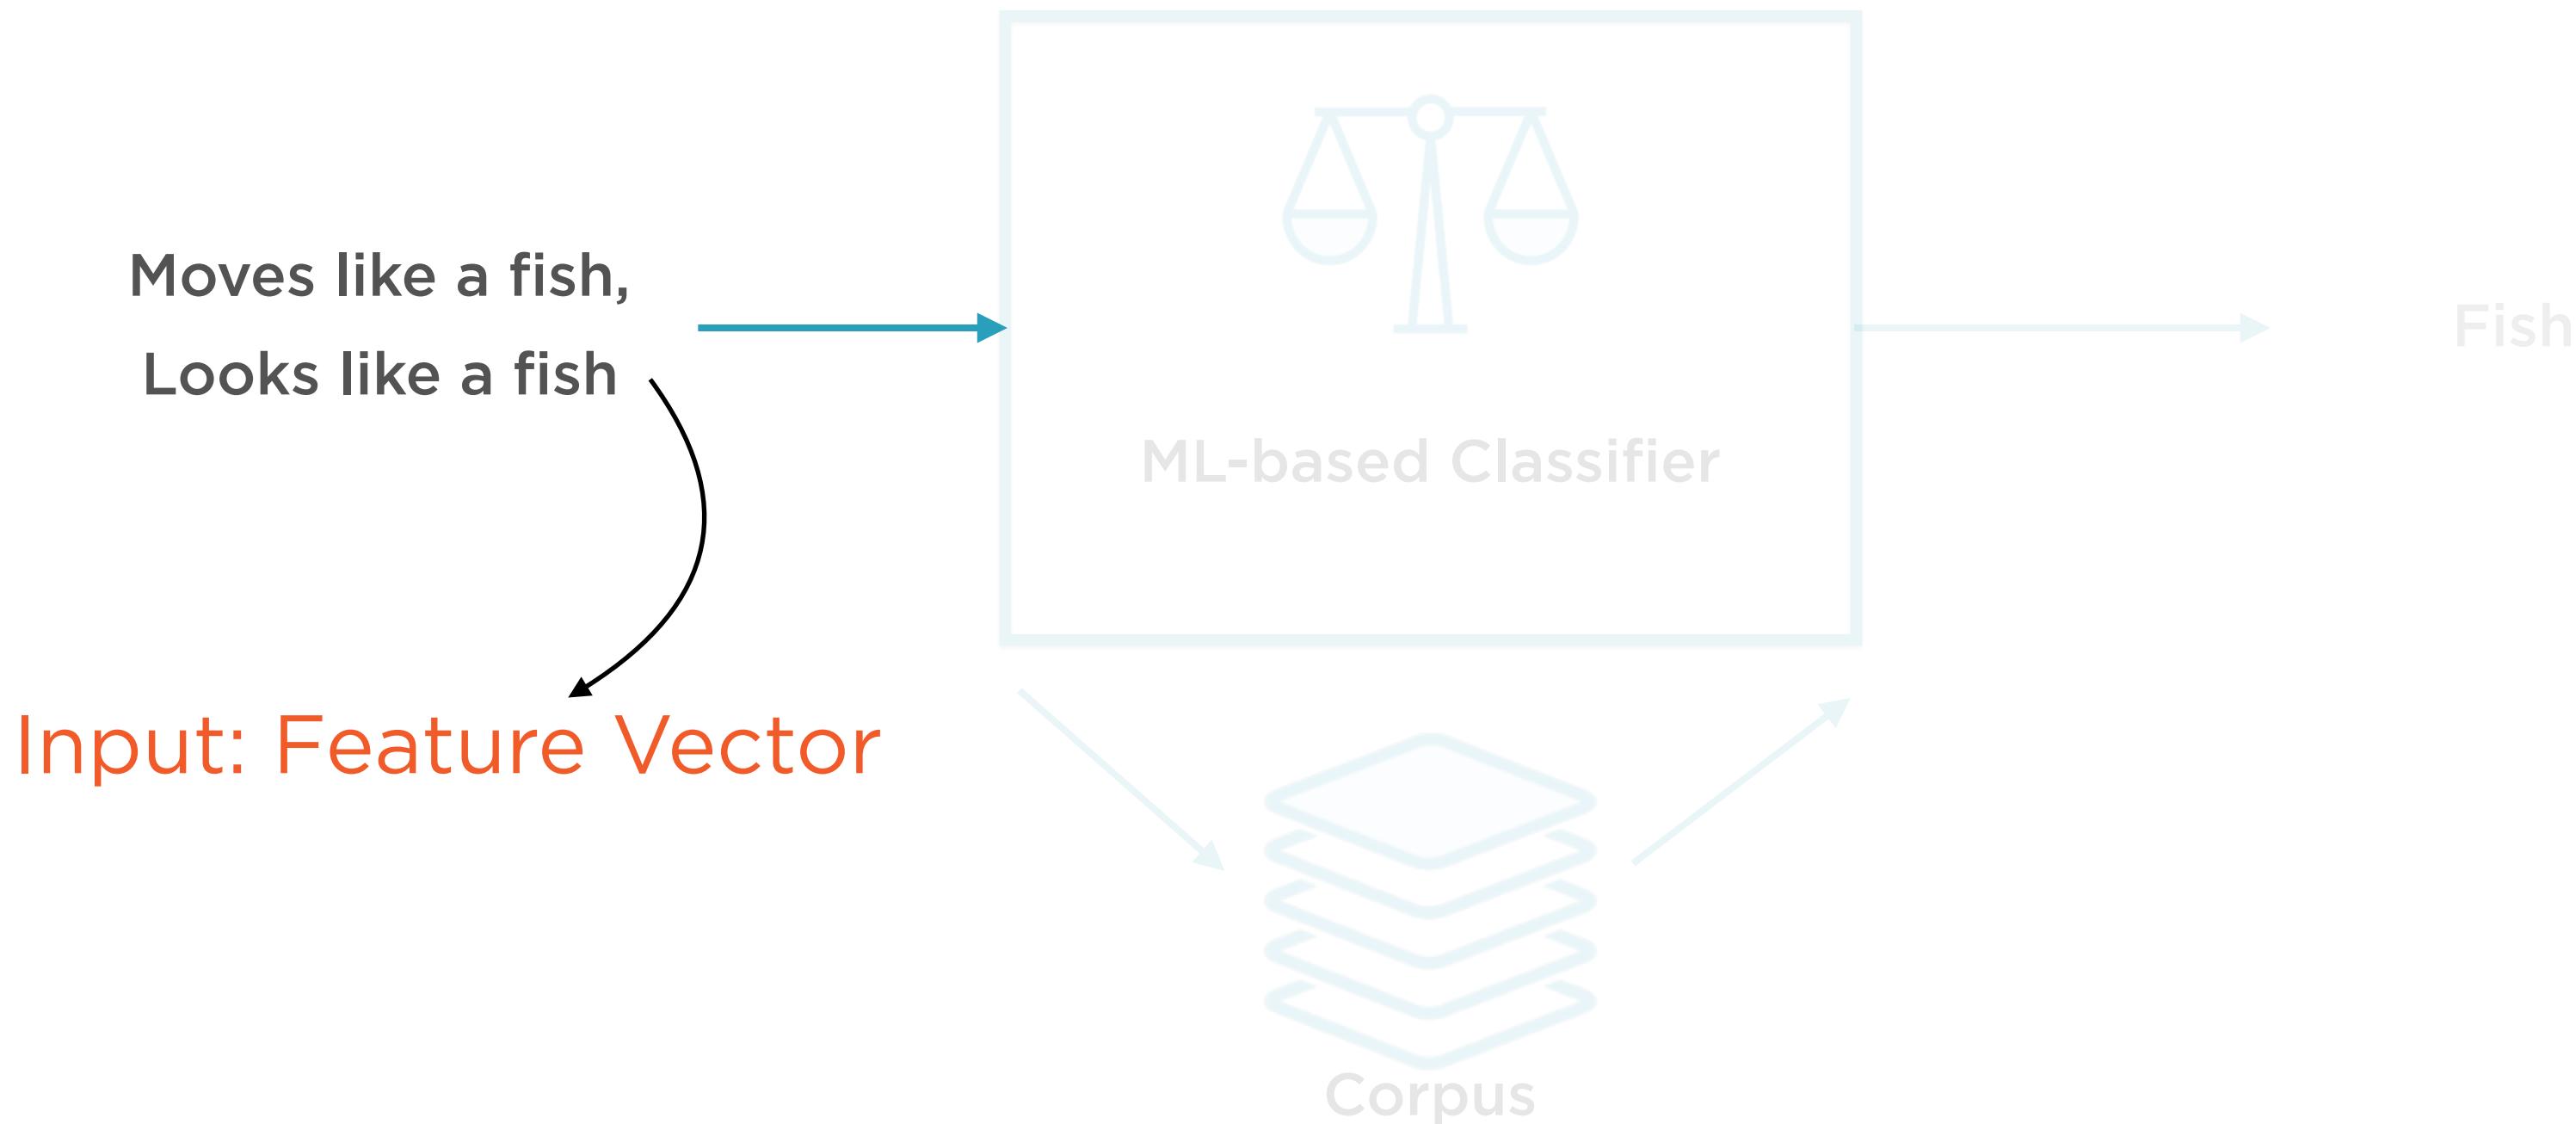
“Traditional” ML-based Binary Classifier



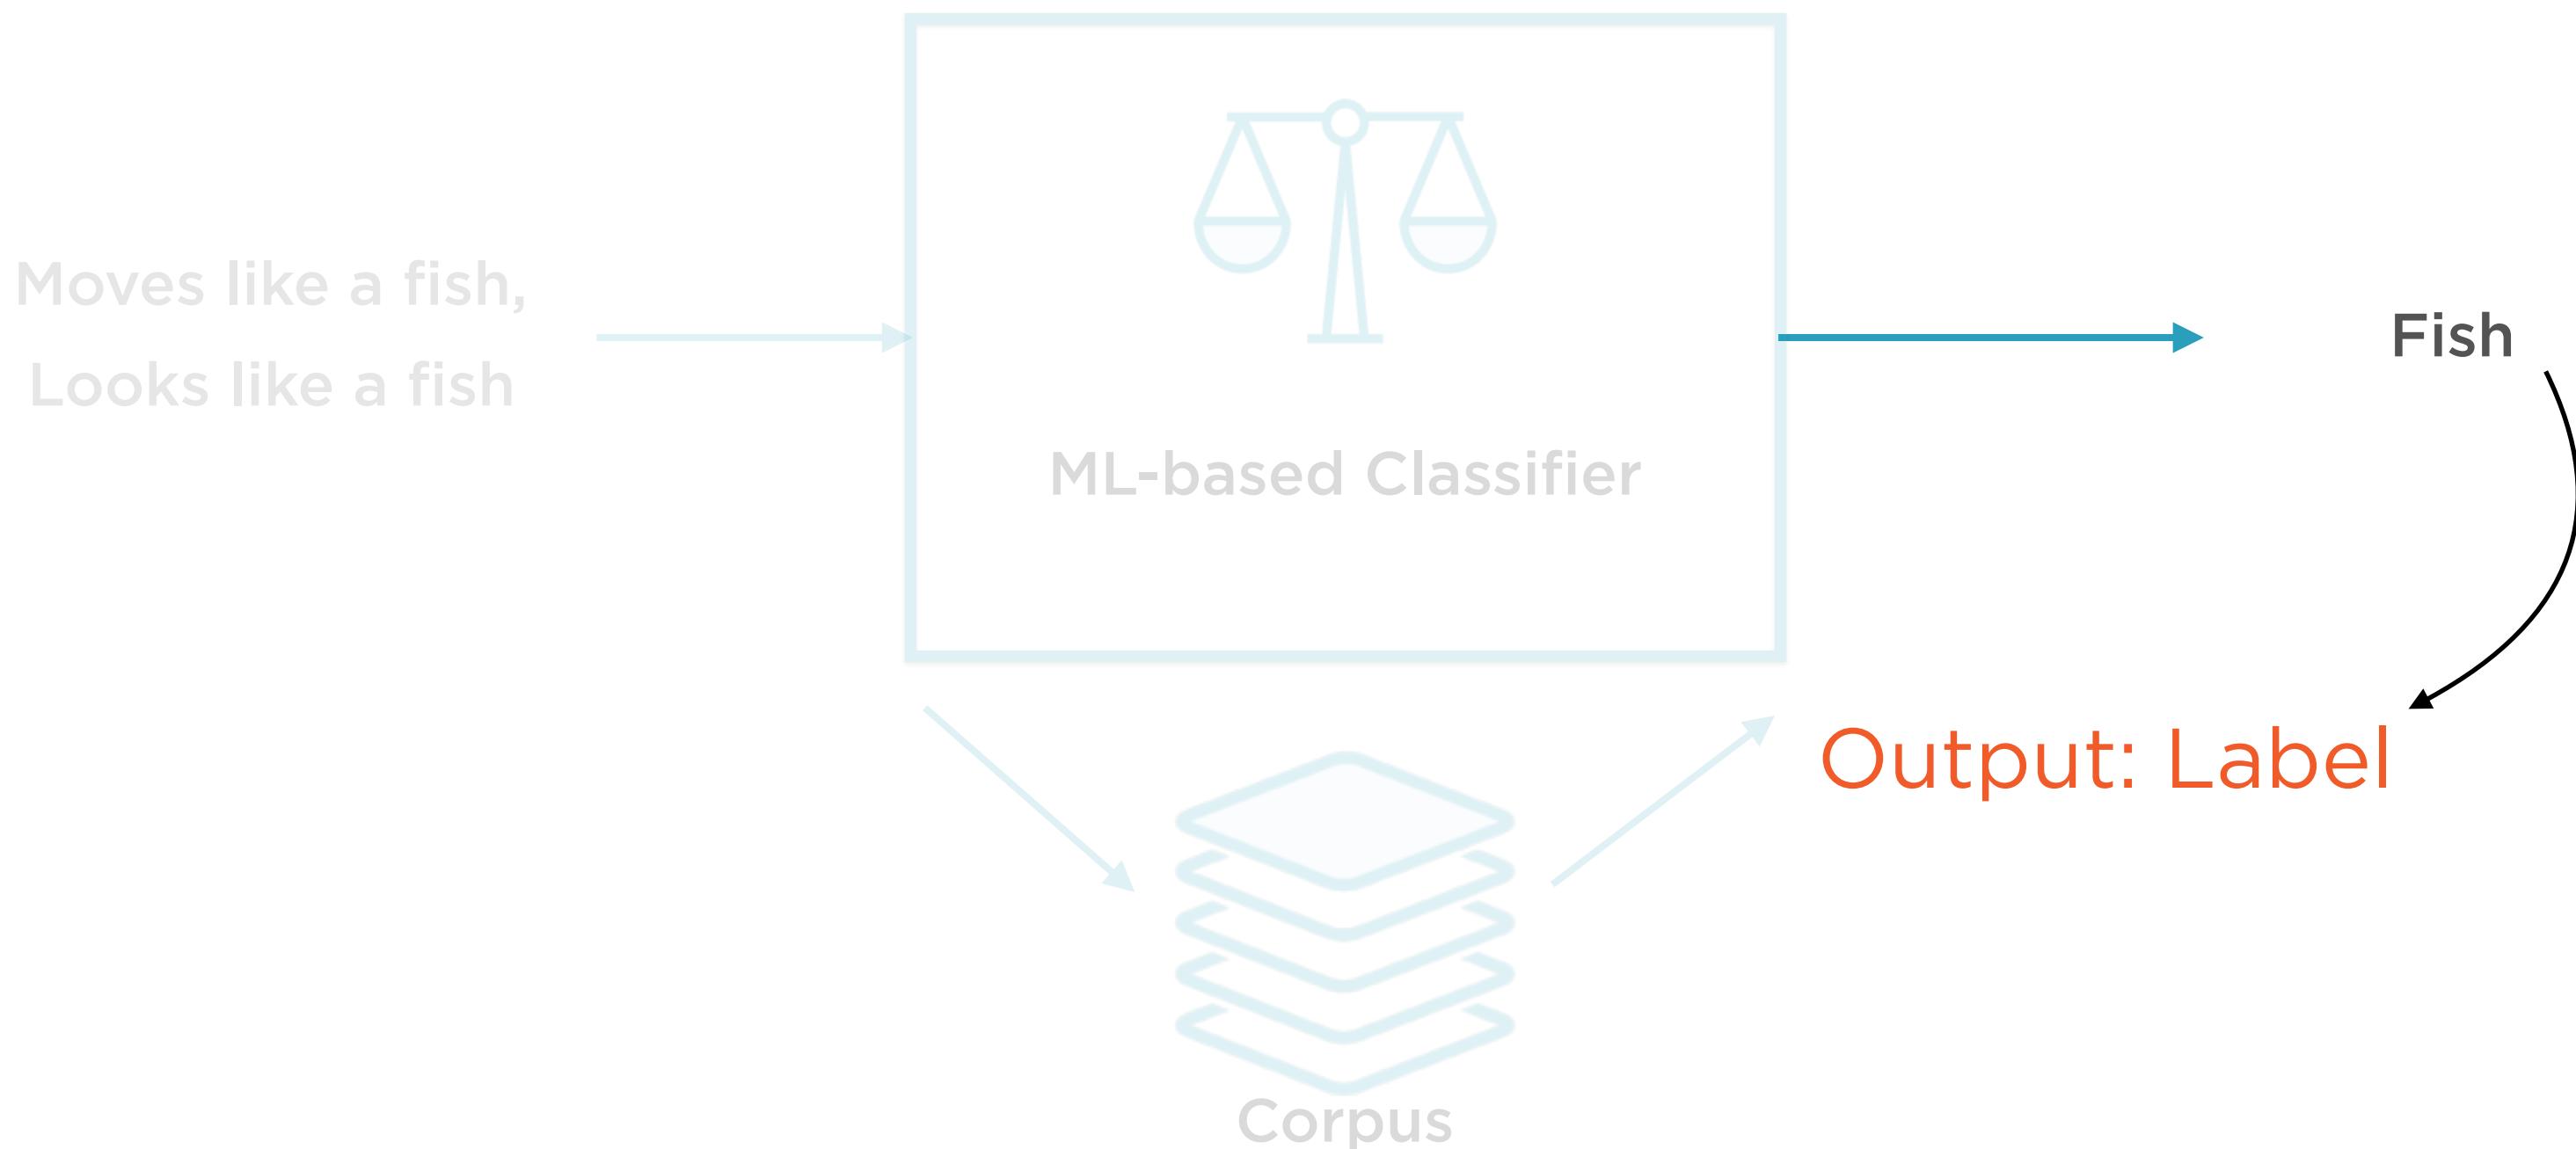
“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



“Traditional” ML-based Binary Classifier



Feature Vectors

The attributes that the ML algorithm focuses on are called **features**

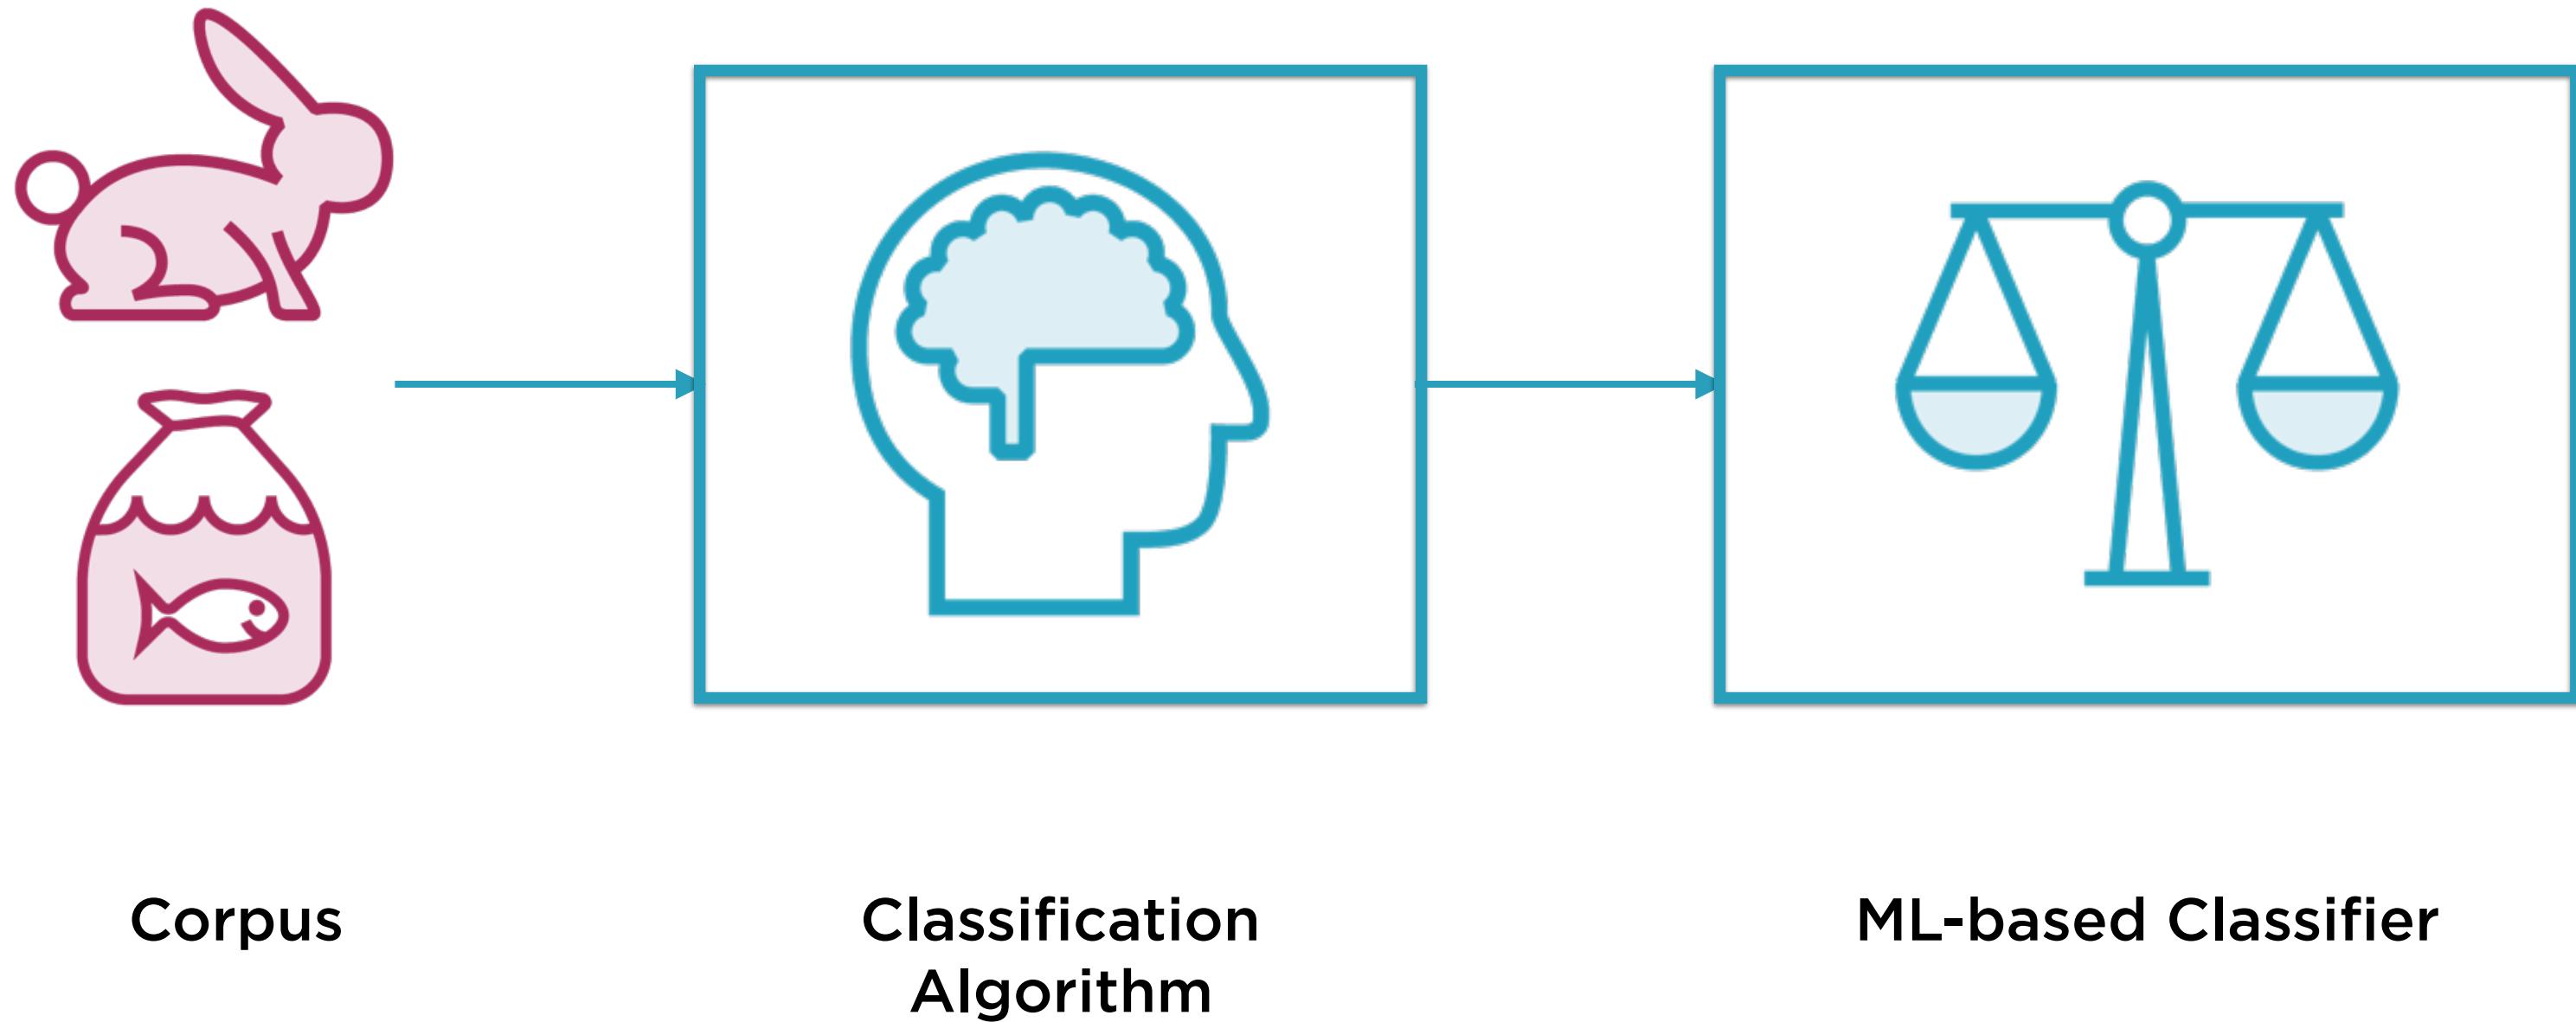
Each data point is a list - or **vector** - of such features

Thus, the input into an ML algorithm is a **feature vector**

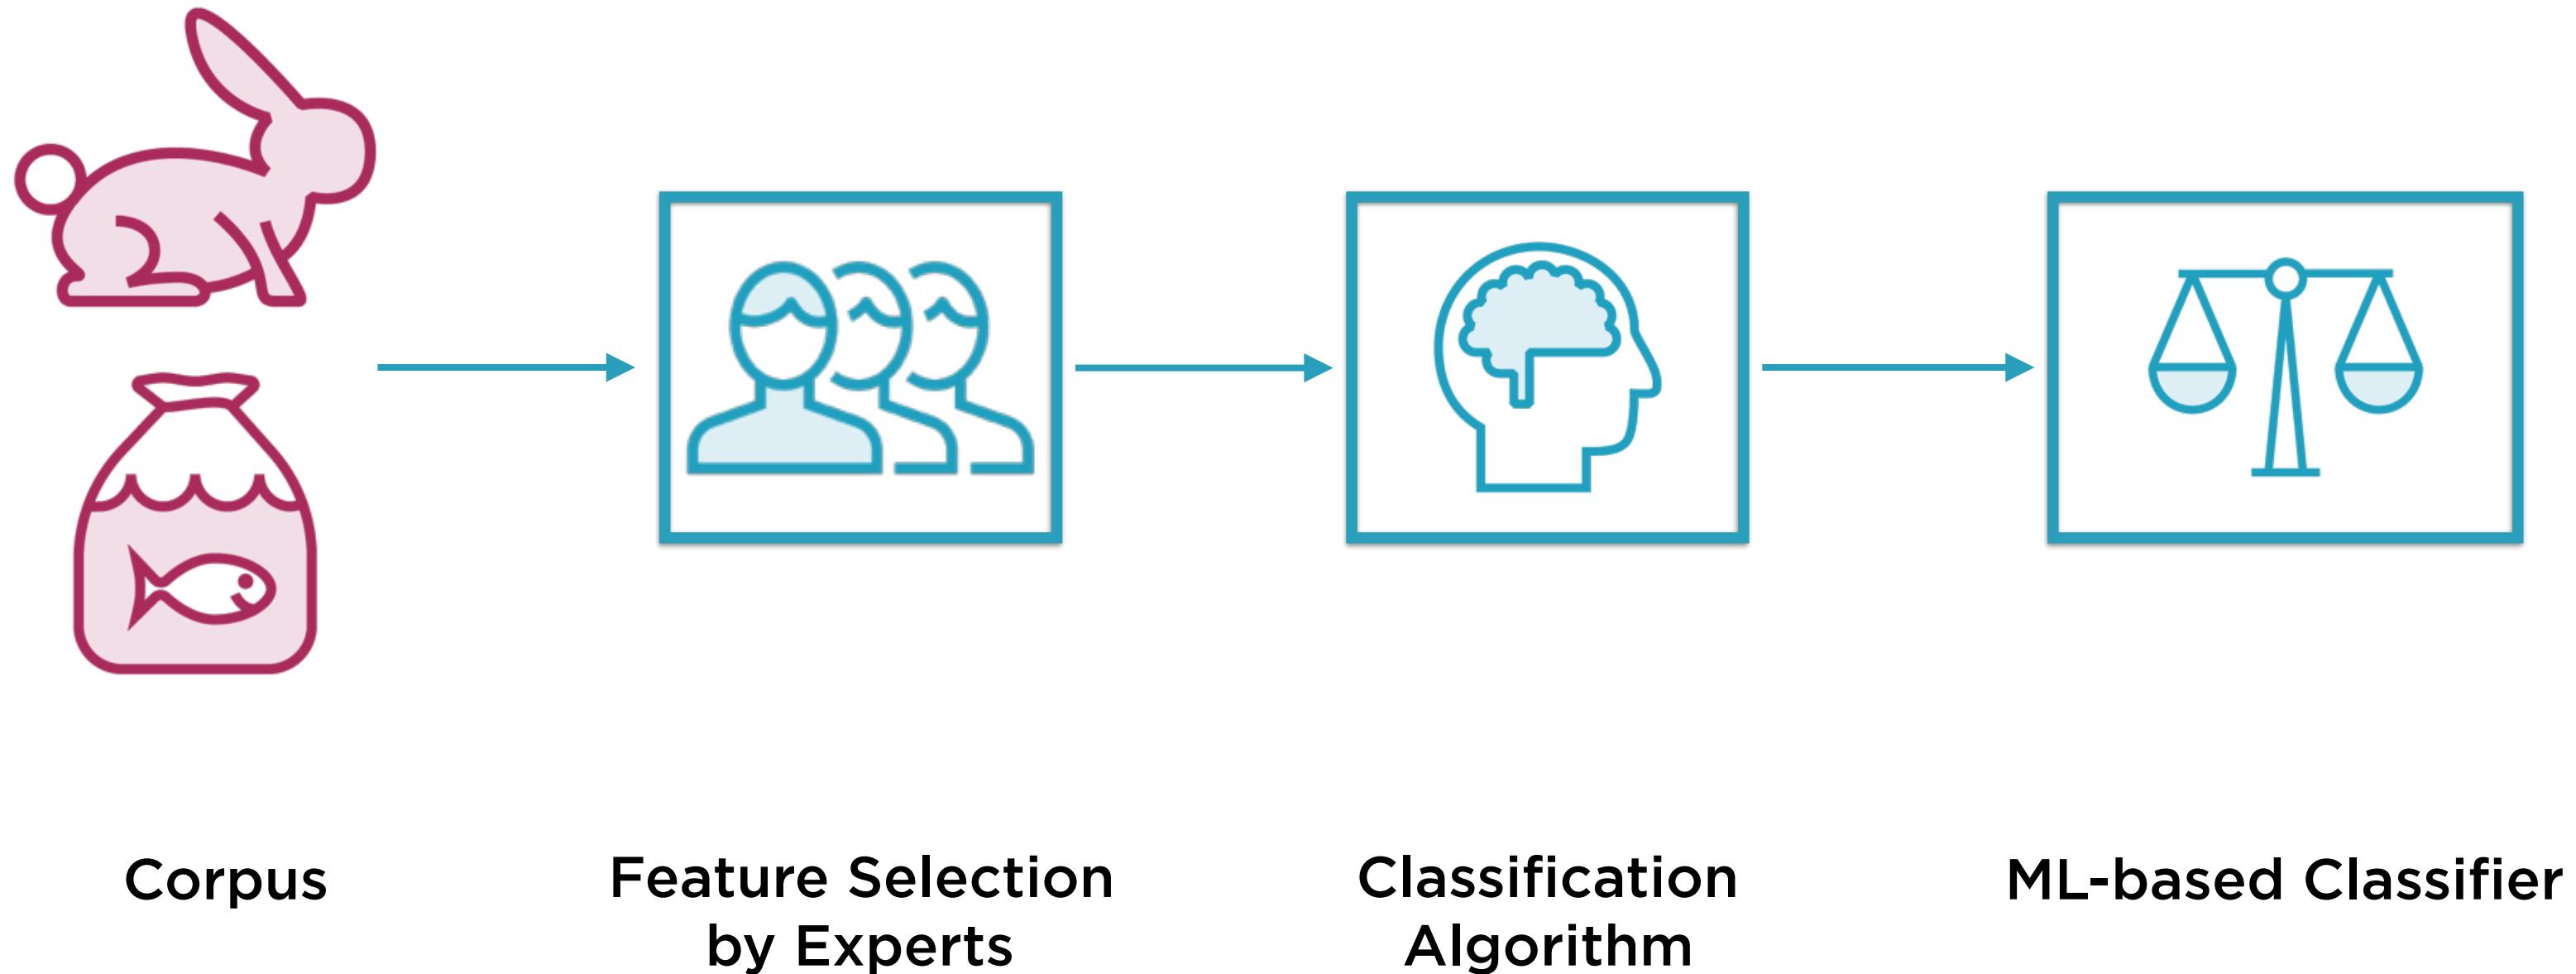
“Traditional” ML-based systems still
rely on experts to decide what
features to pay attention to

“Representation” ML-based systems
figure out by themselves what
features to pay attention to

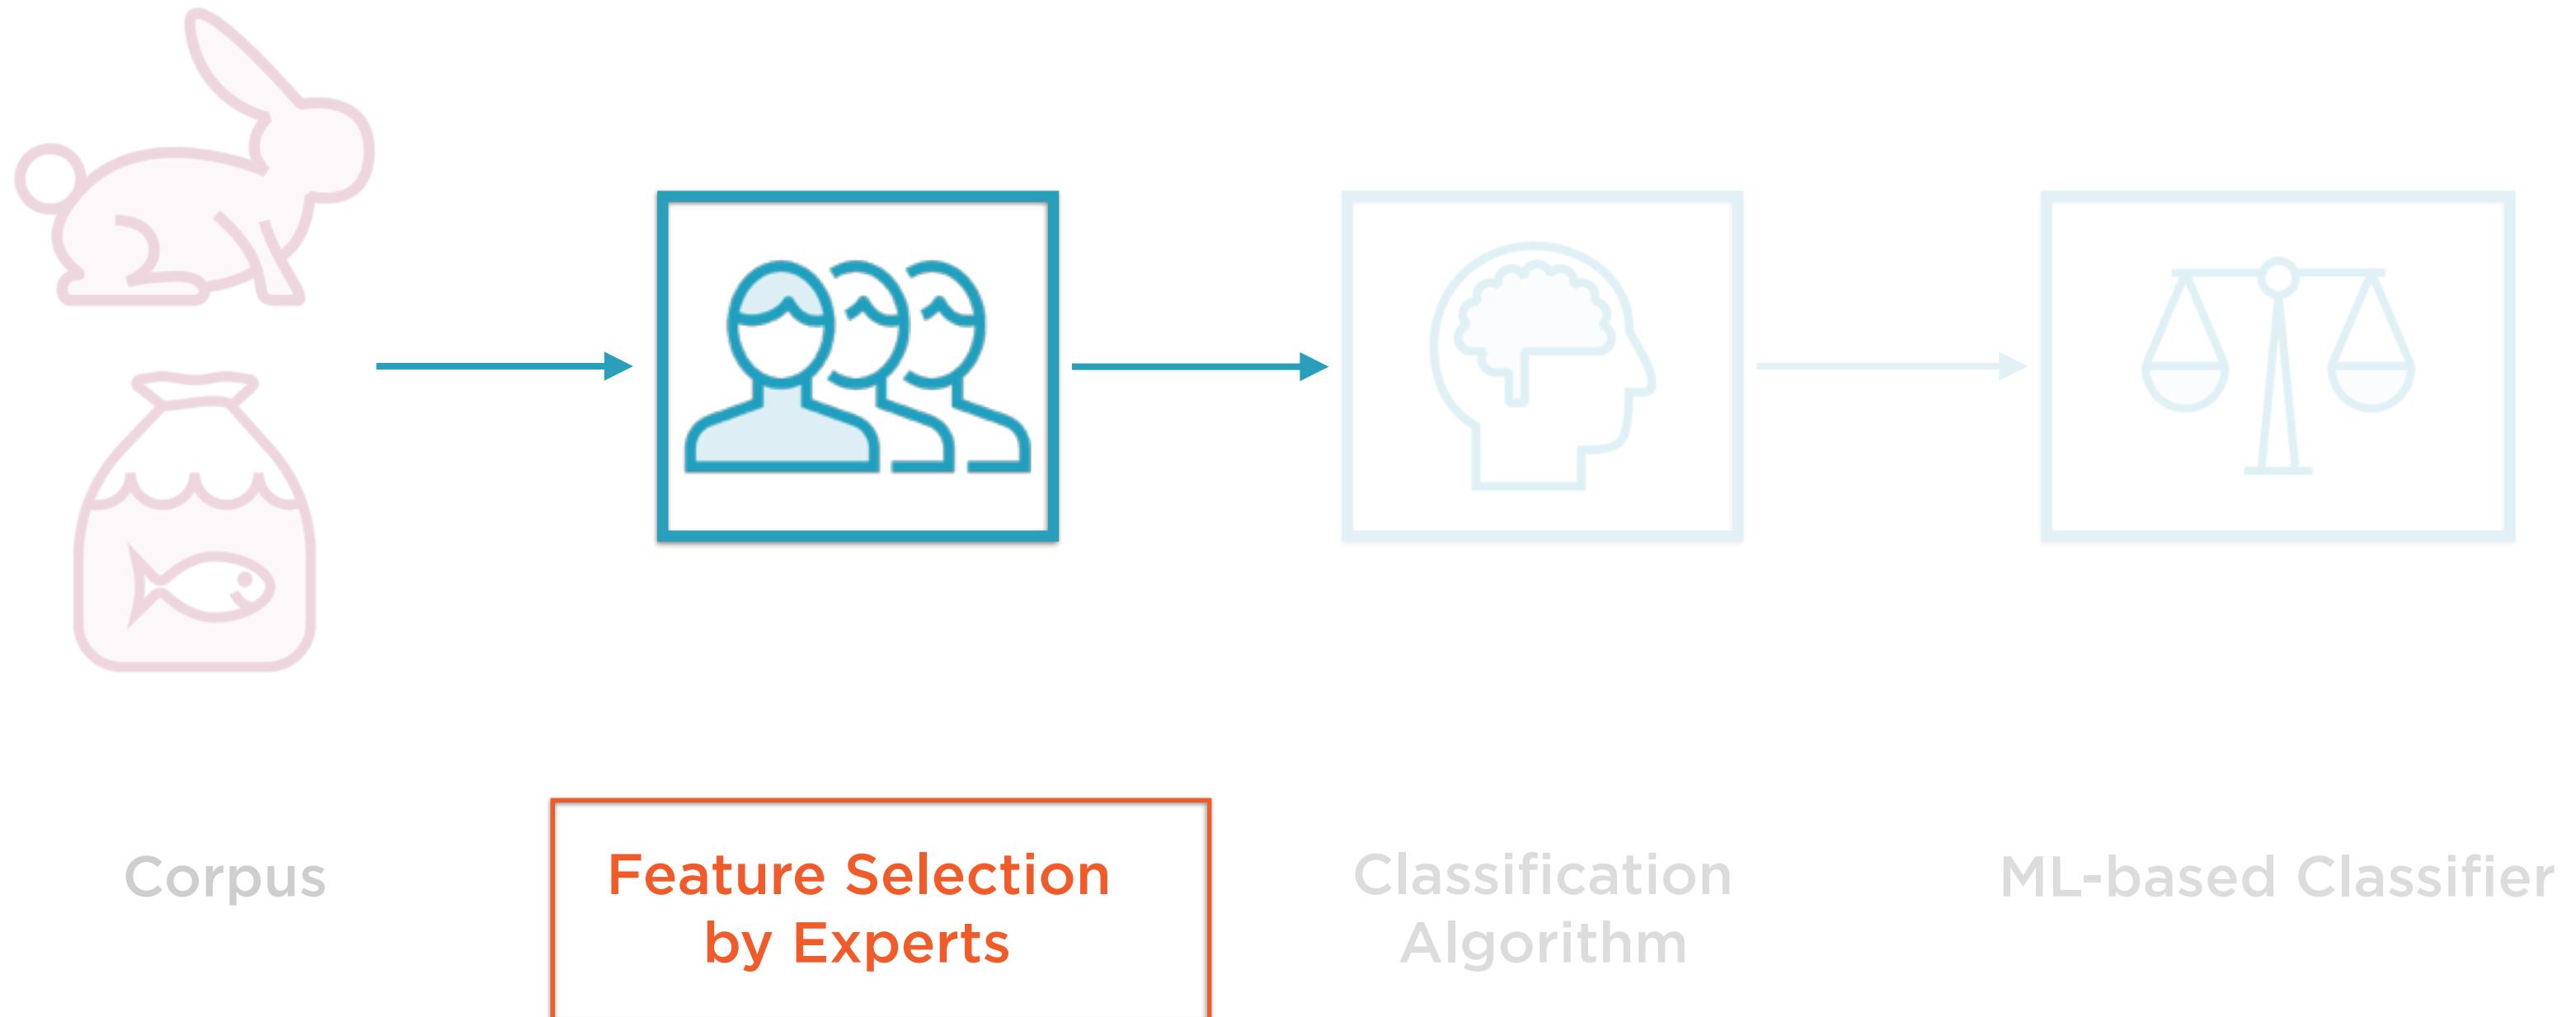
“Traditional” ML-based Binary Classifier



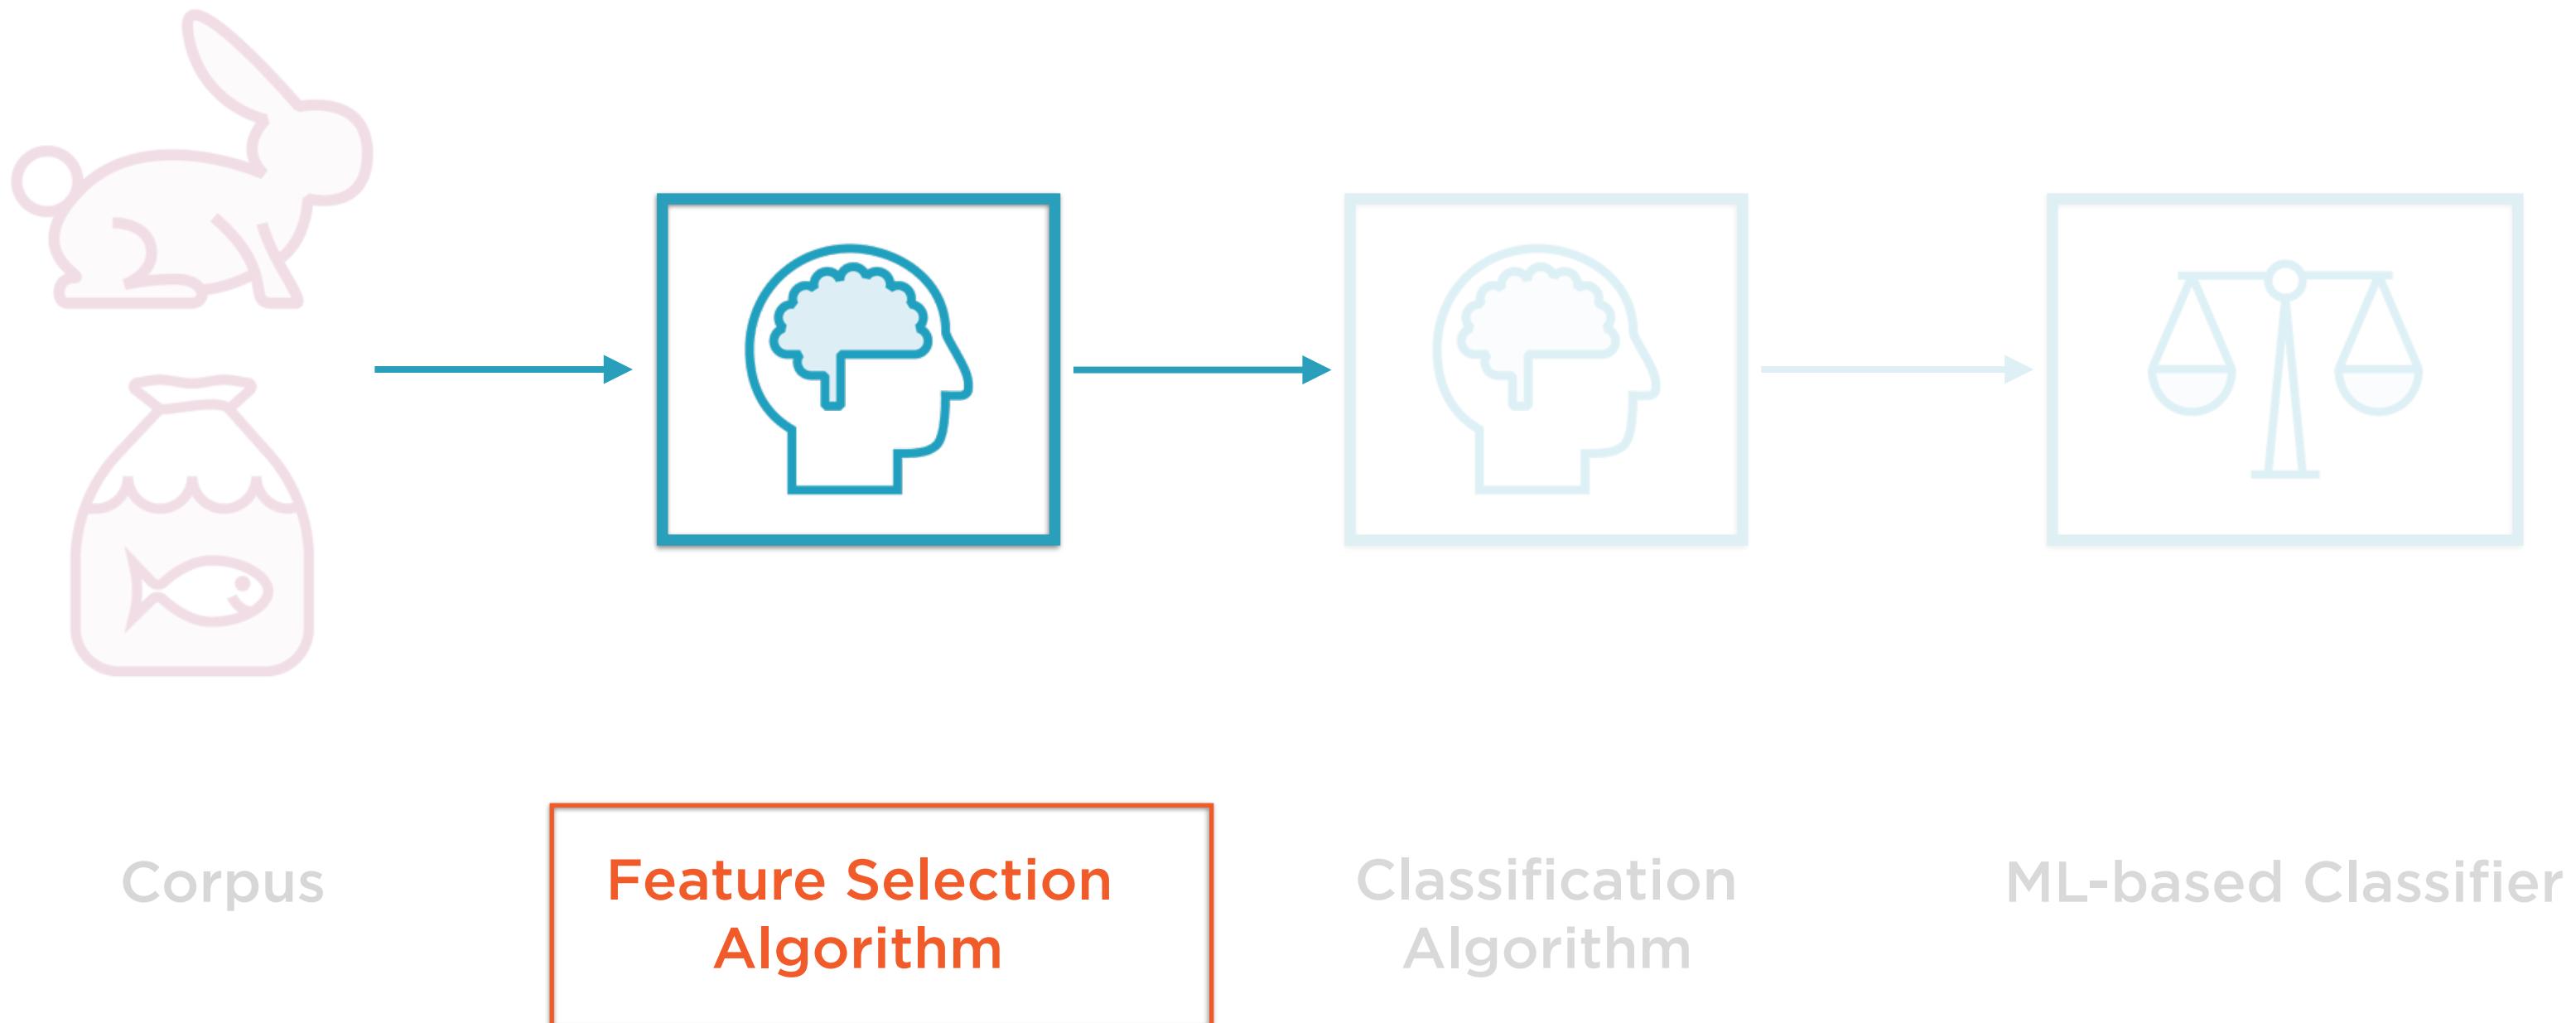
“Traditional” ML-based Binary Classifier



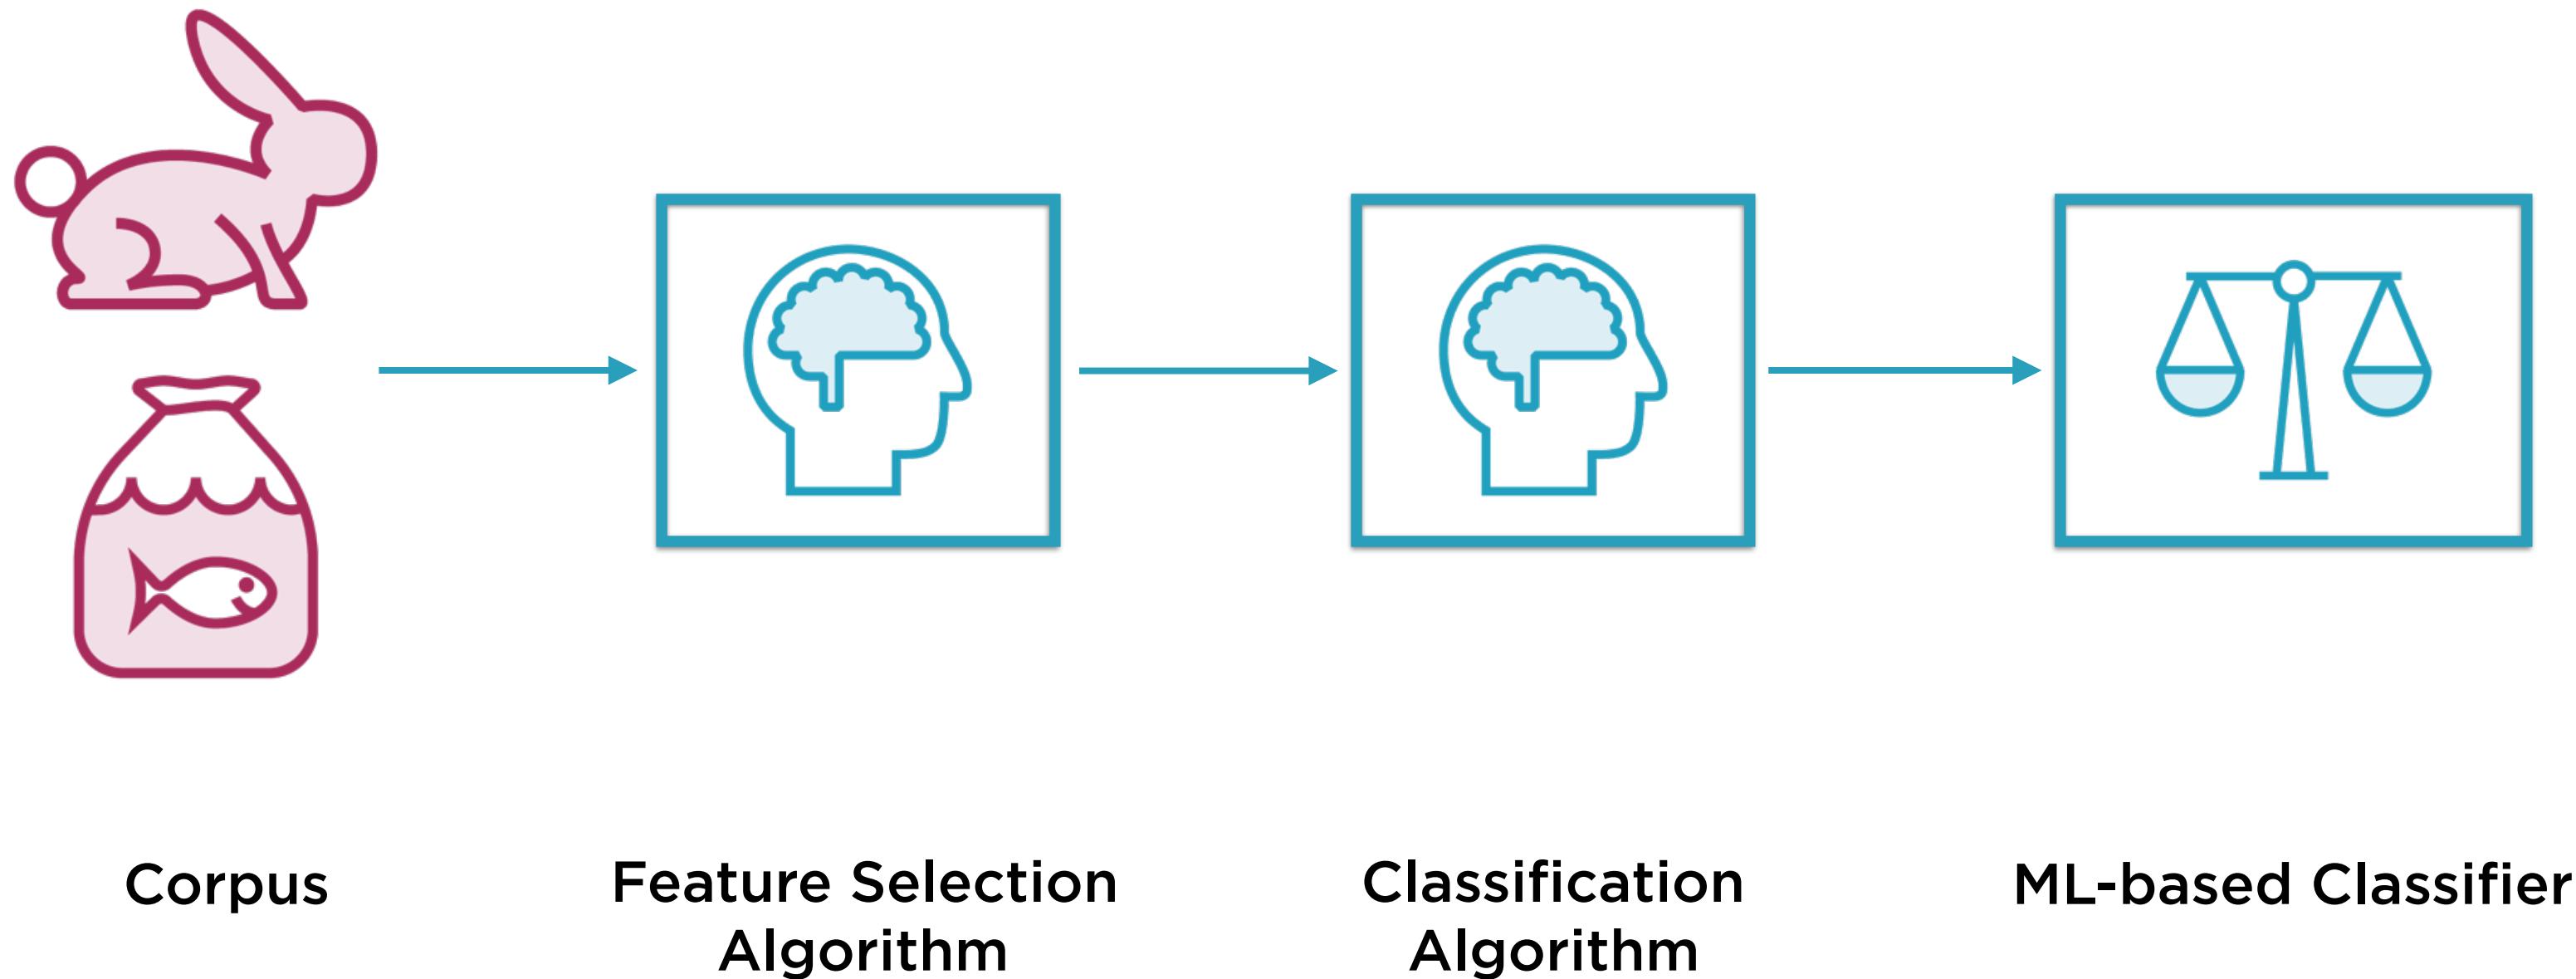
“Traditional” ML-based Binary Classifier



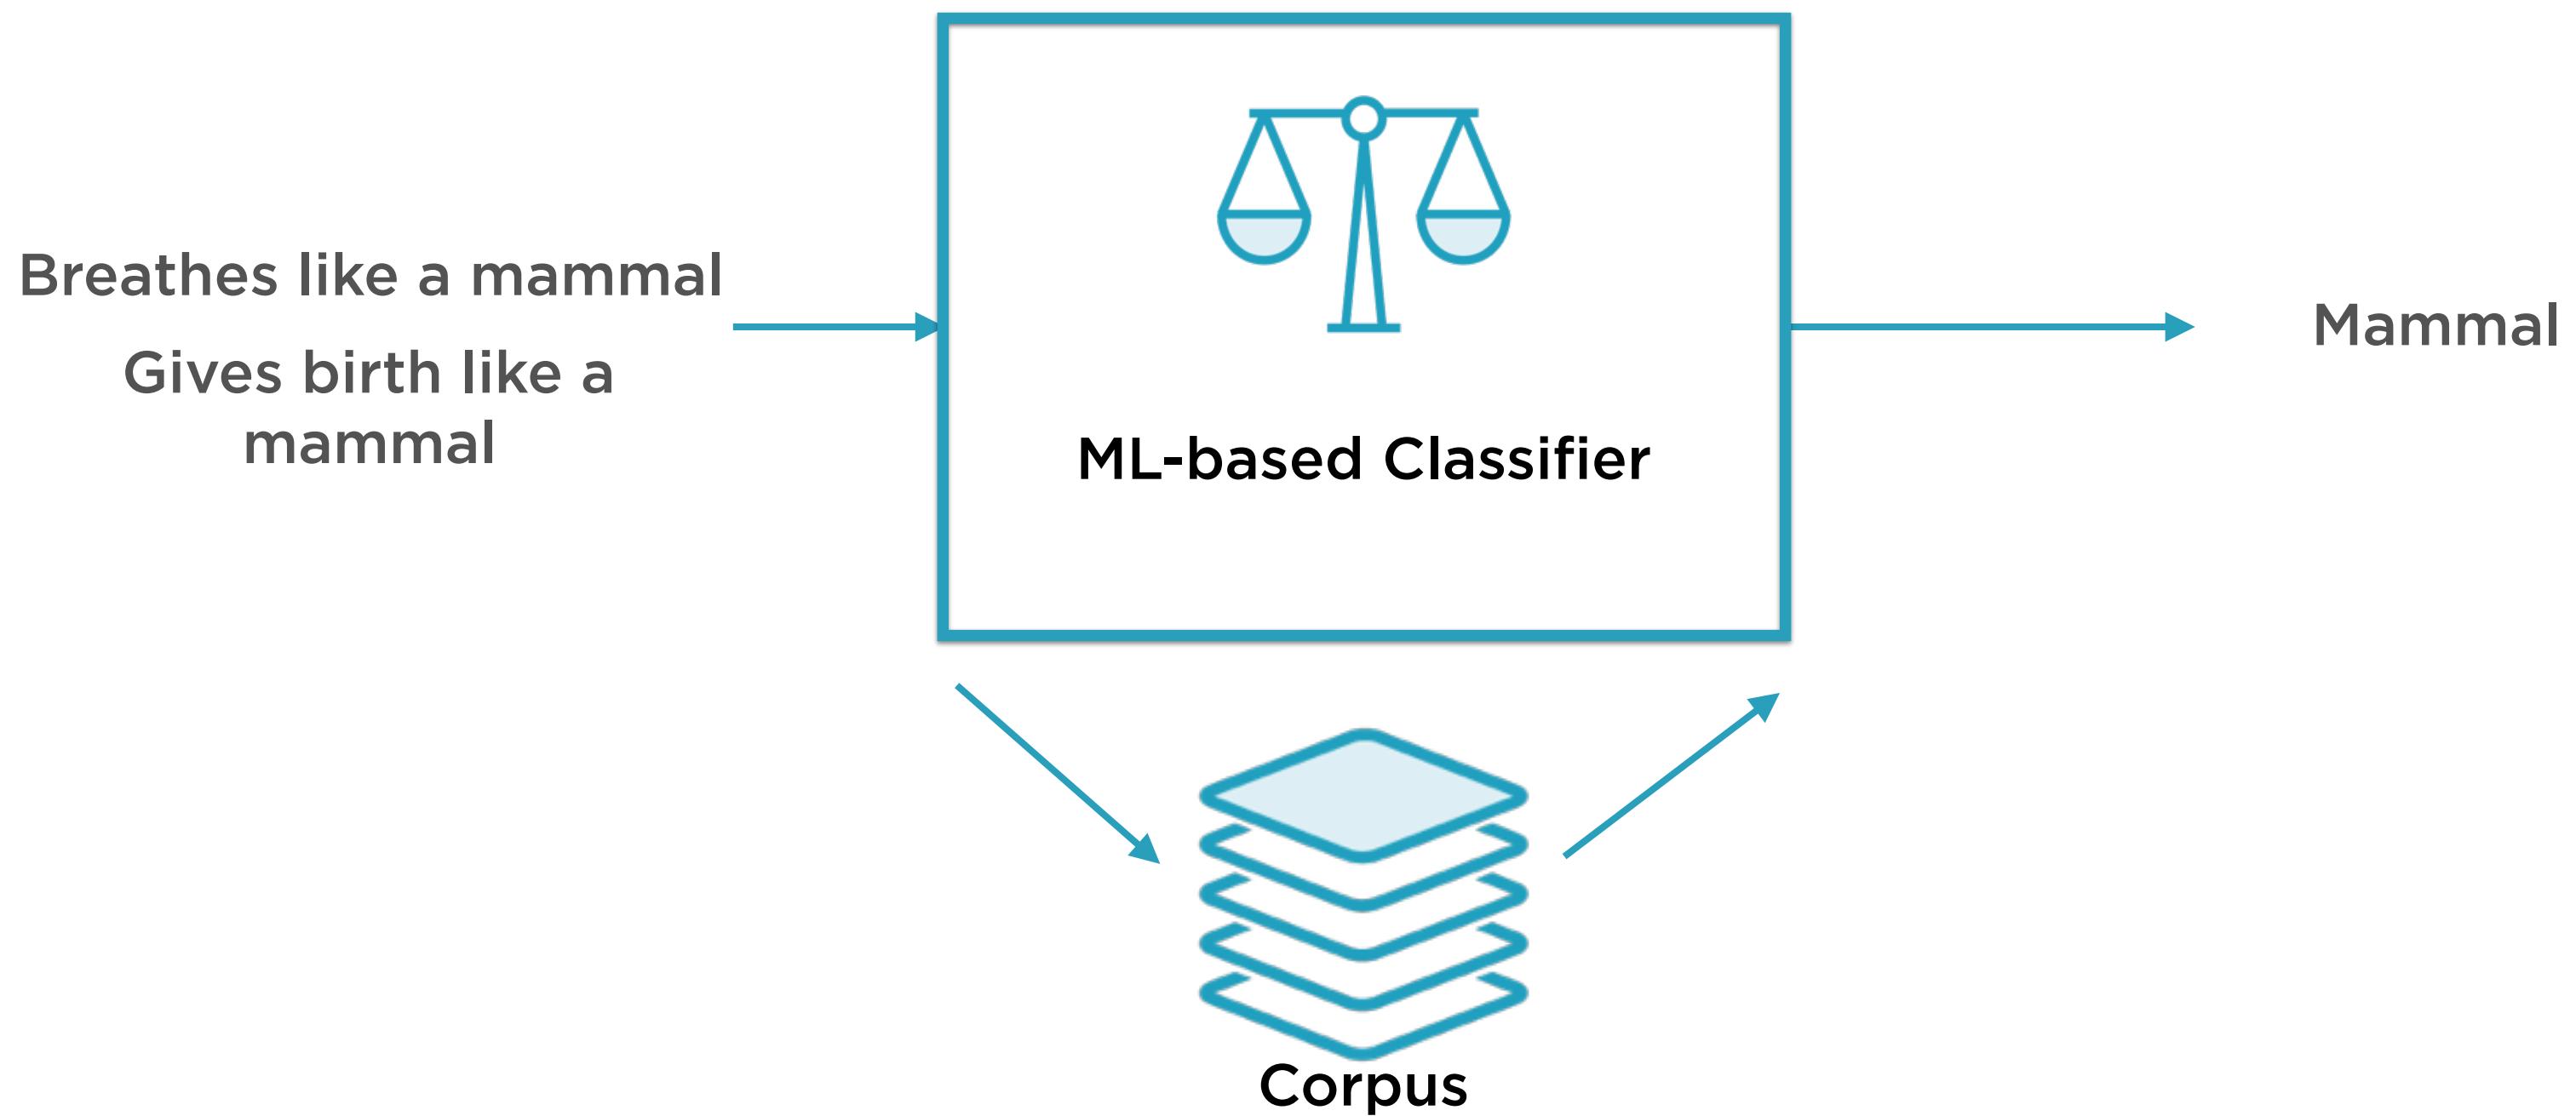
“Representation” ML-based Binary Classifier



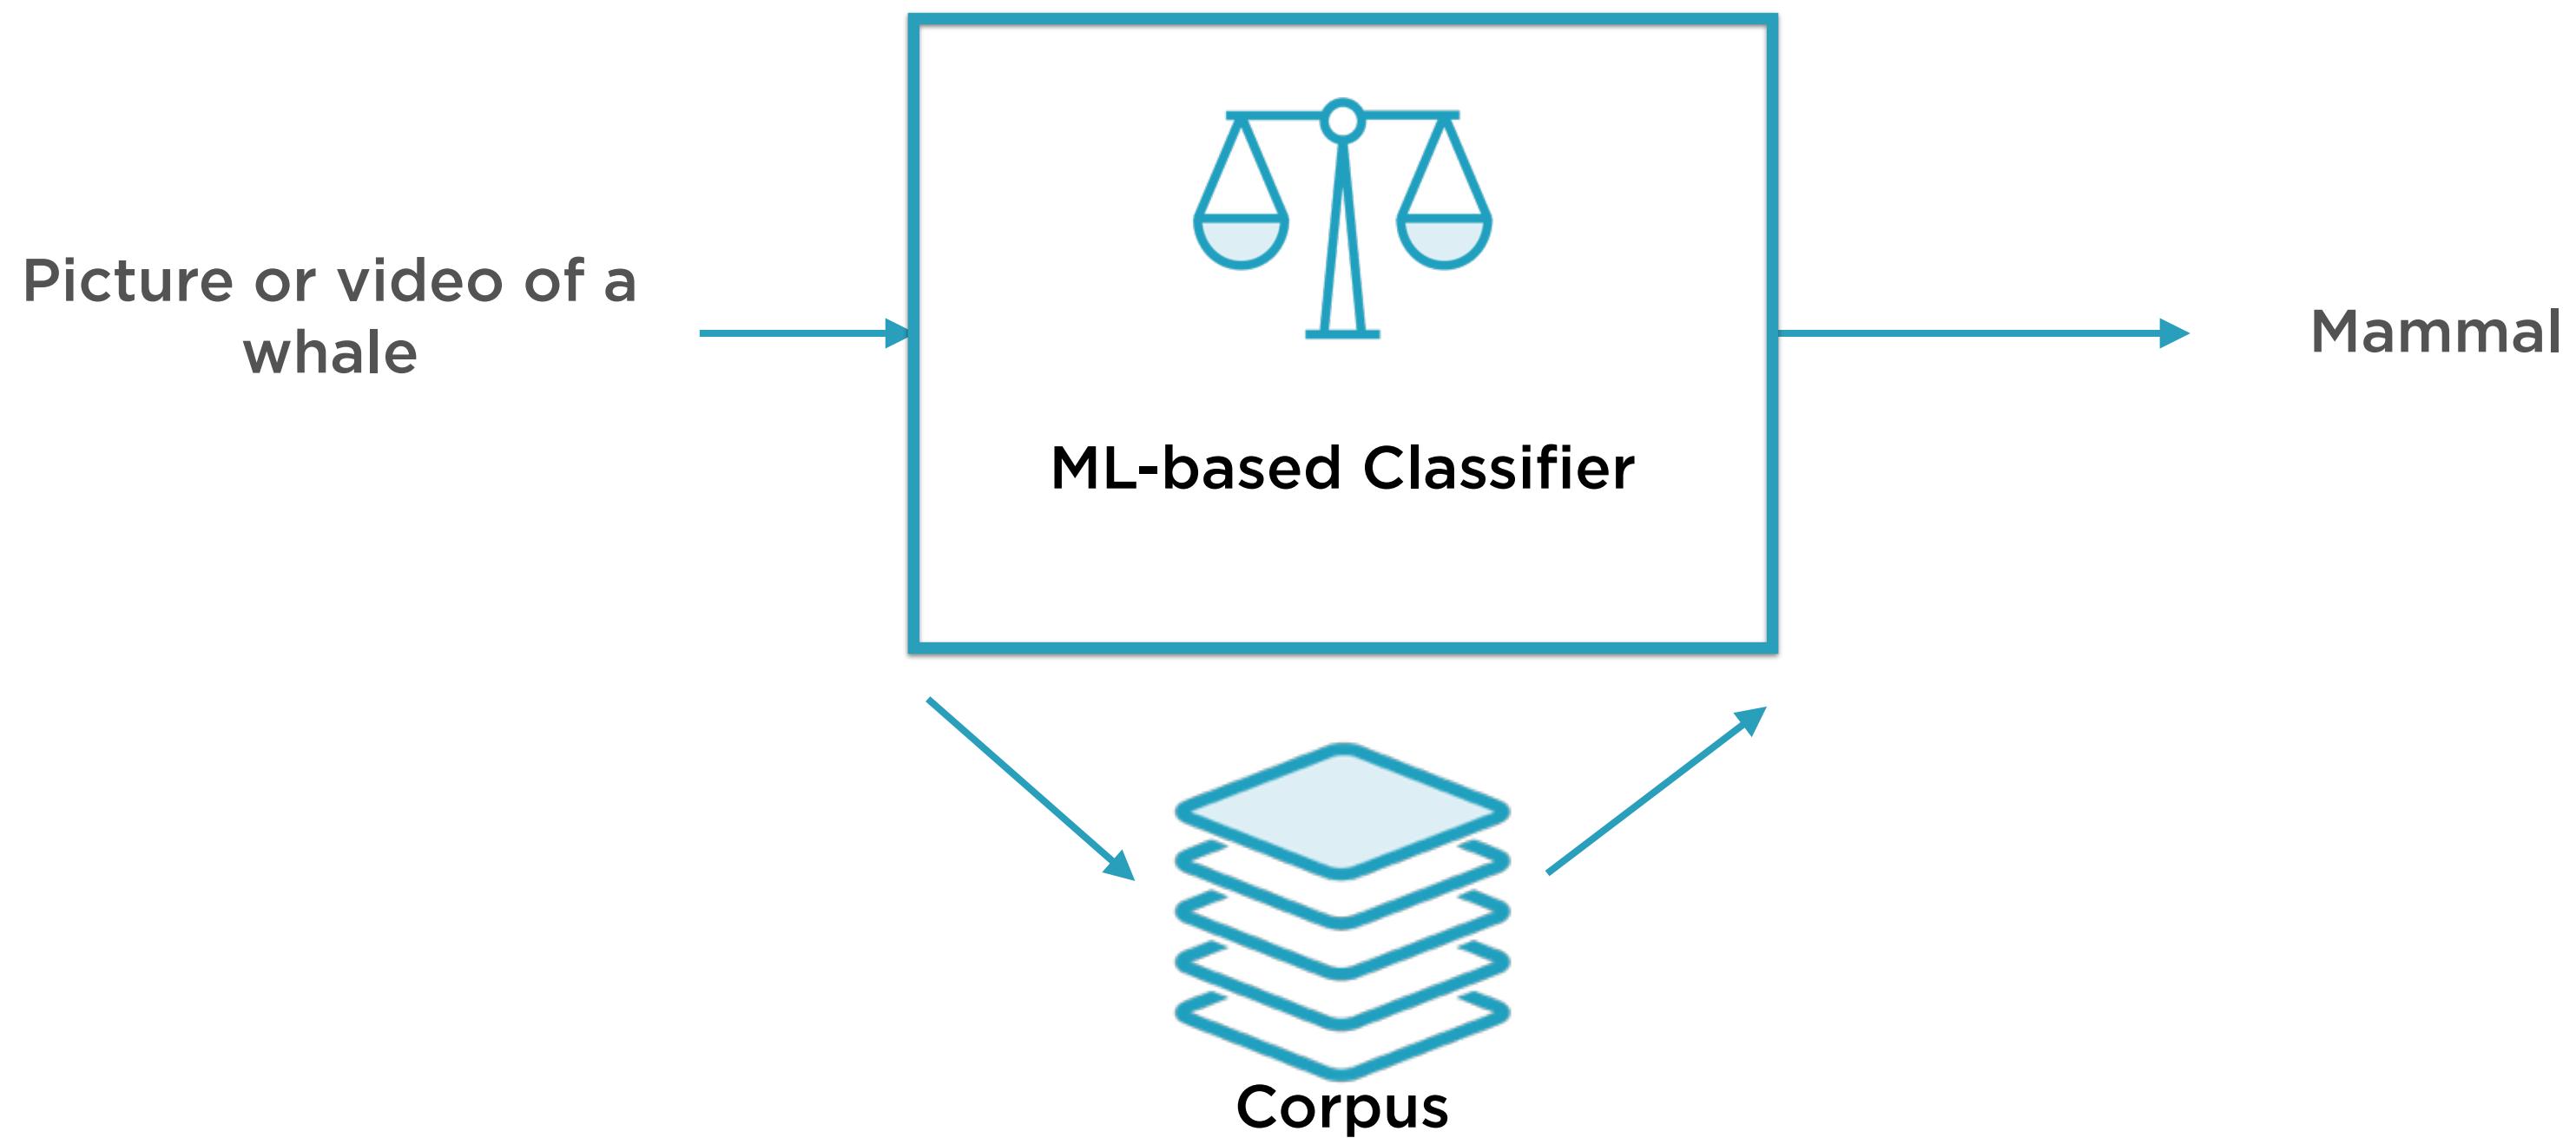
“Representation” ML-based Binary Classifier



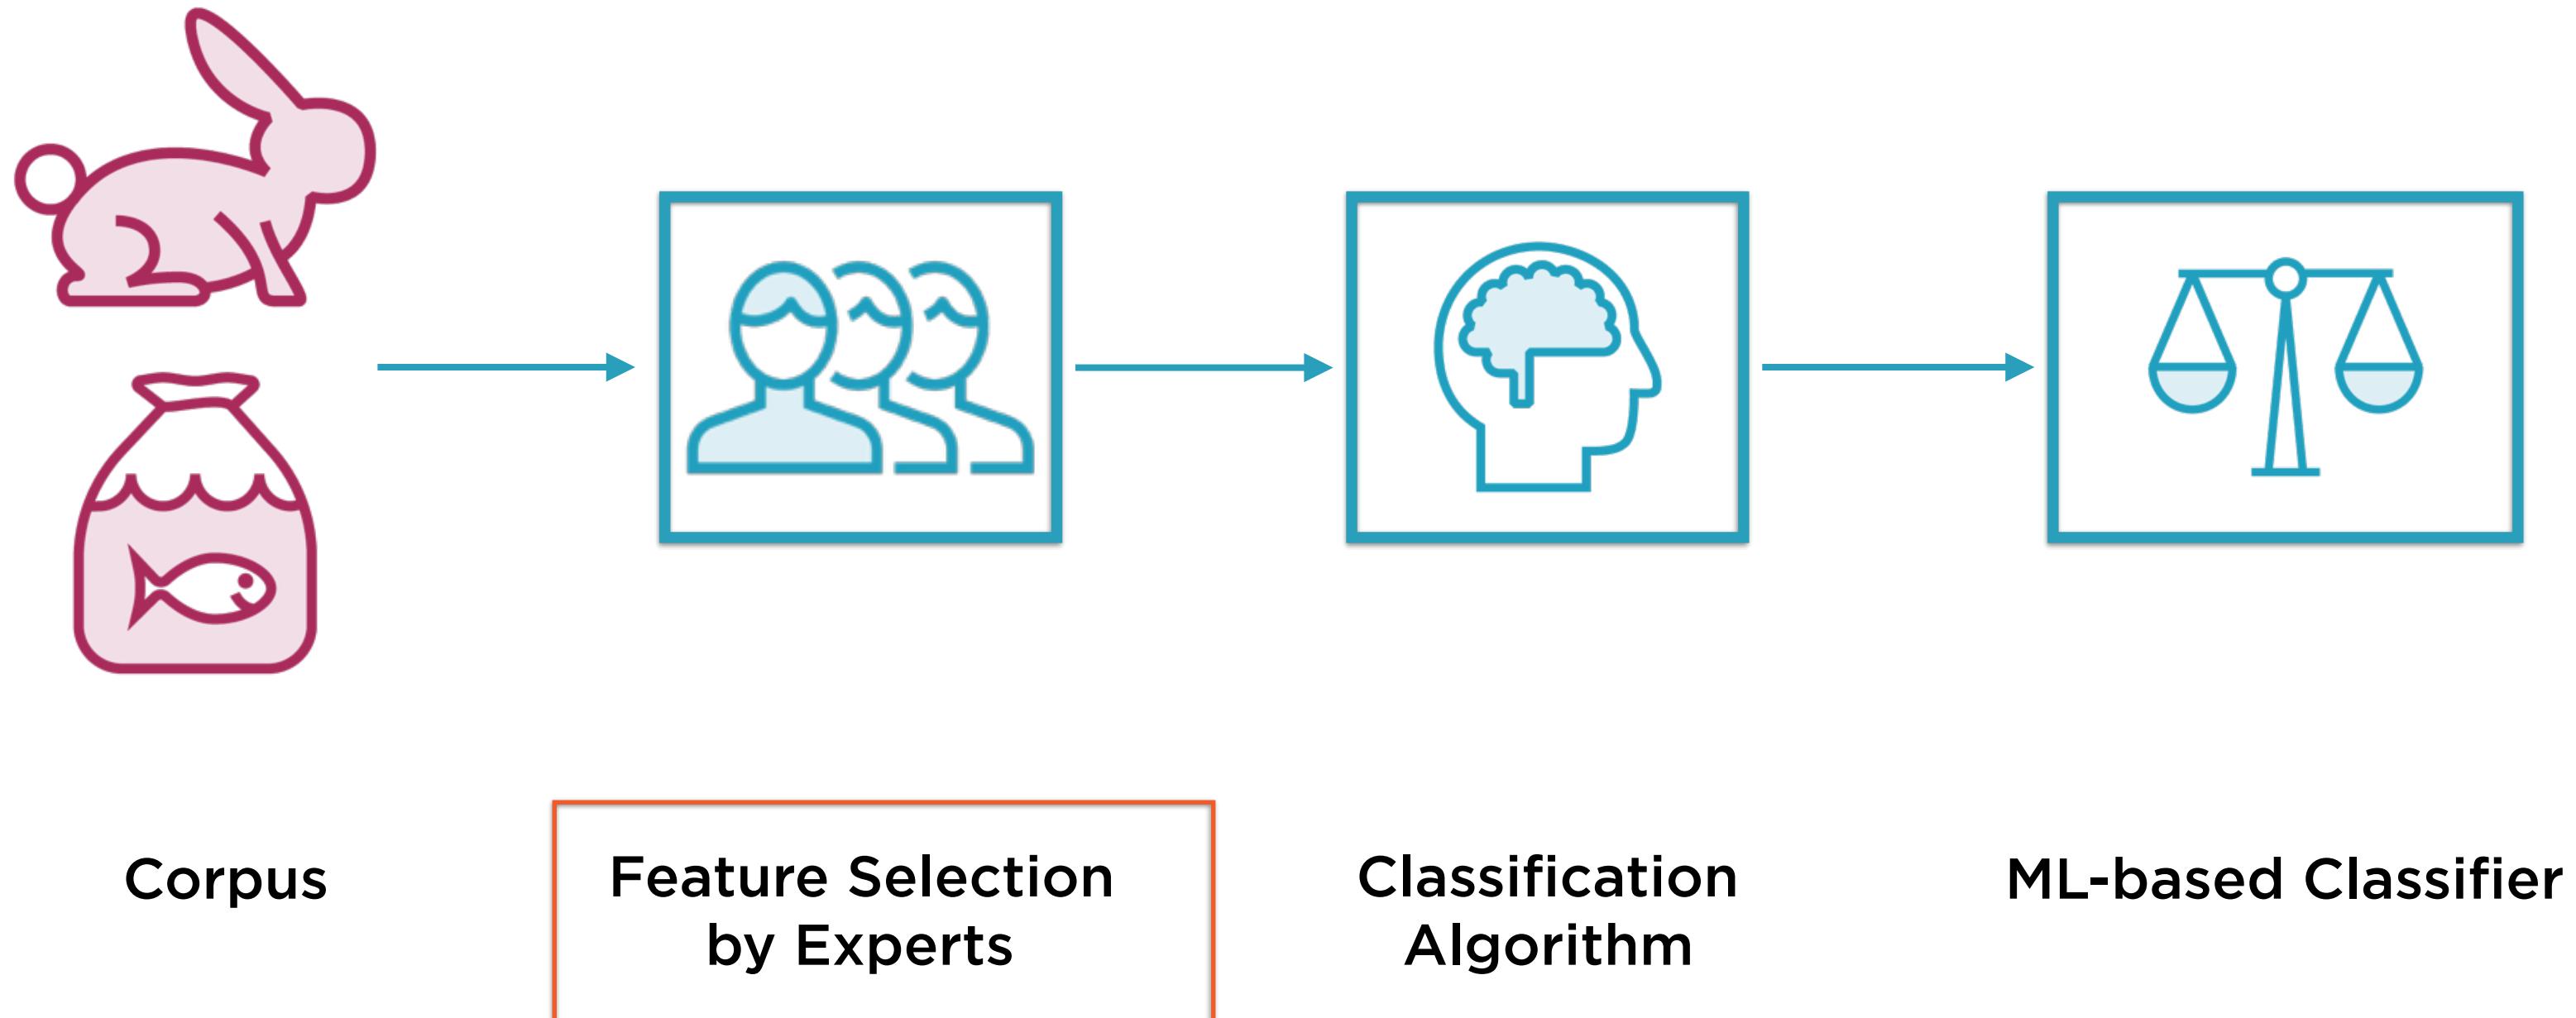
“Traditional” ML-based Binary Classifier



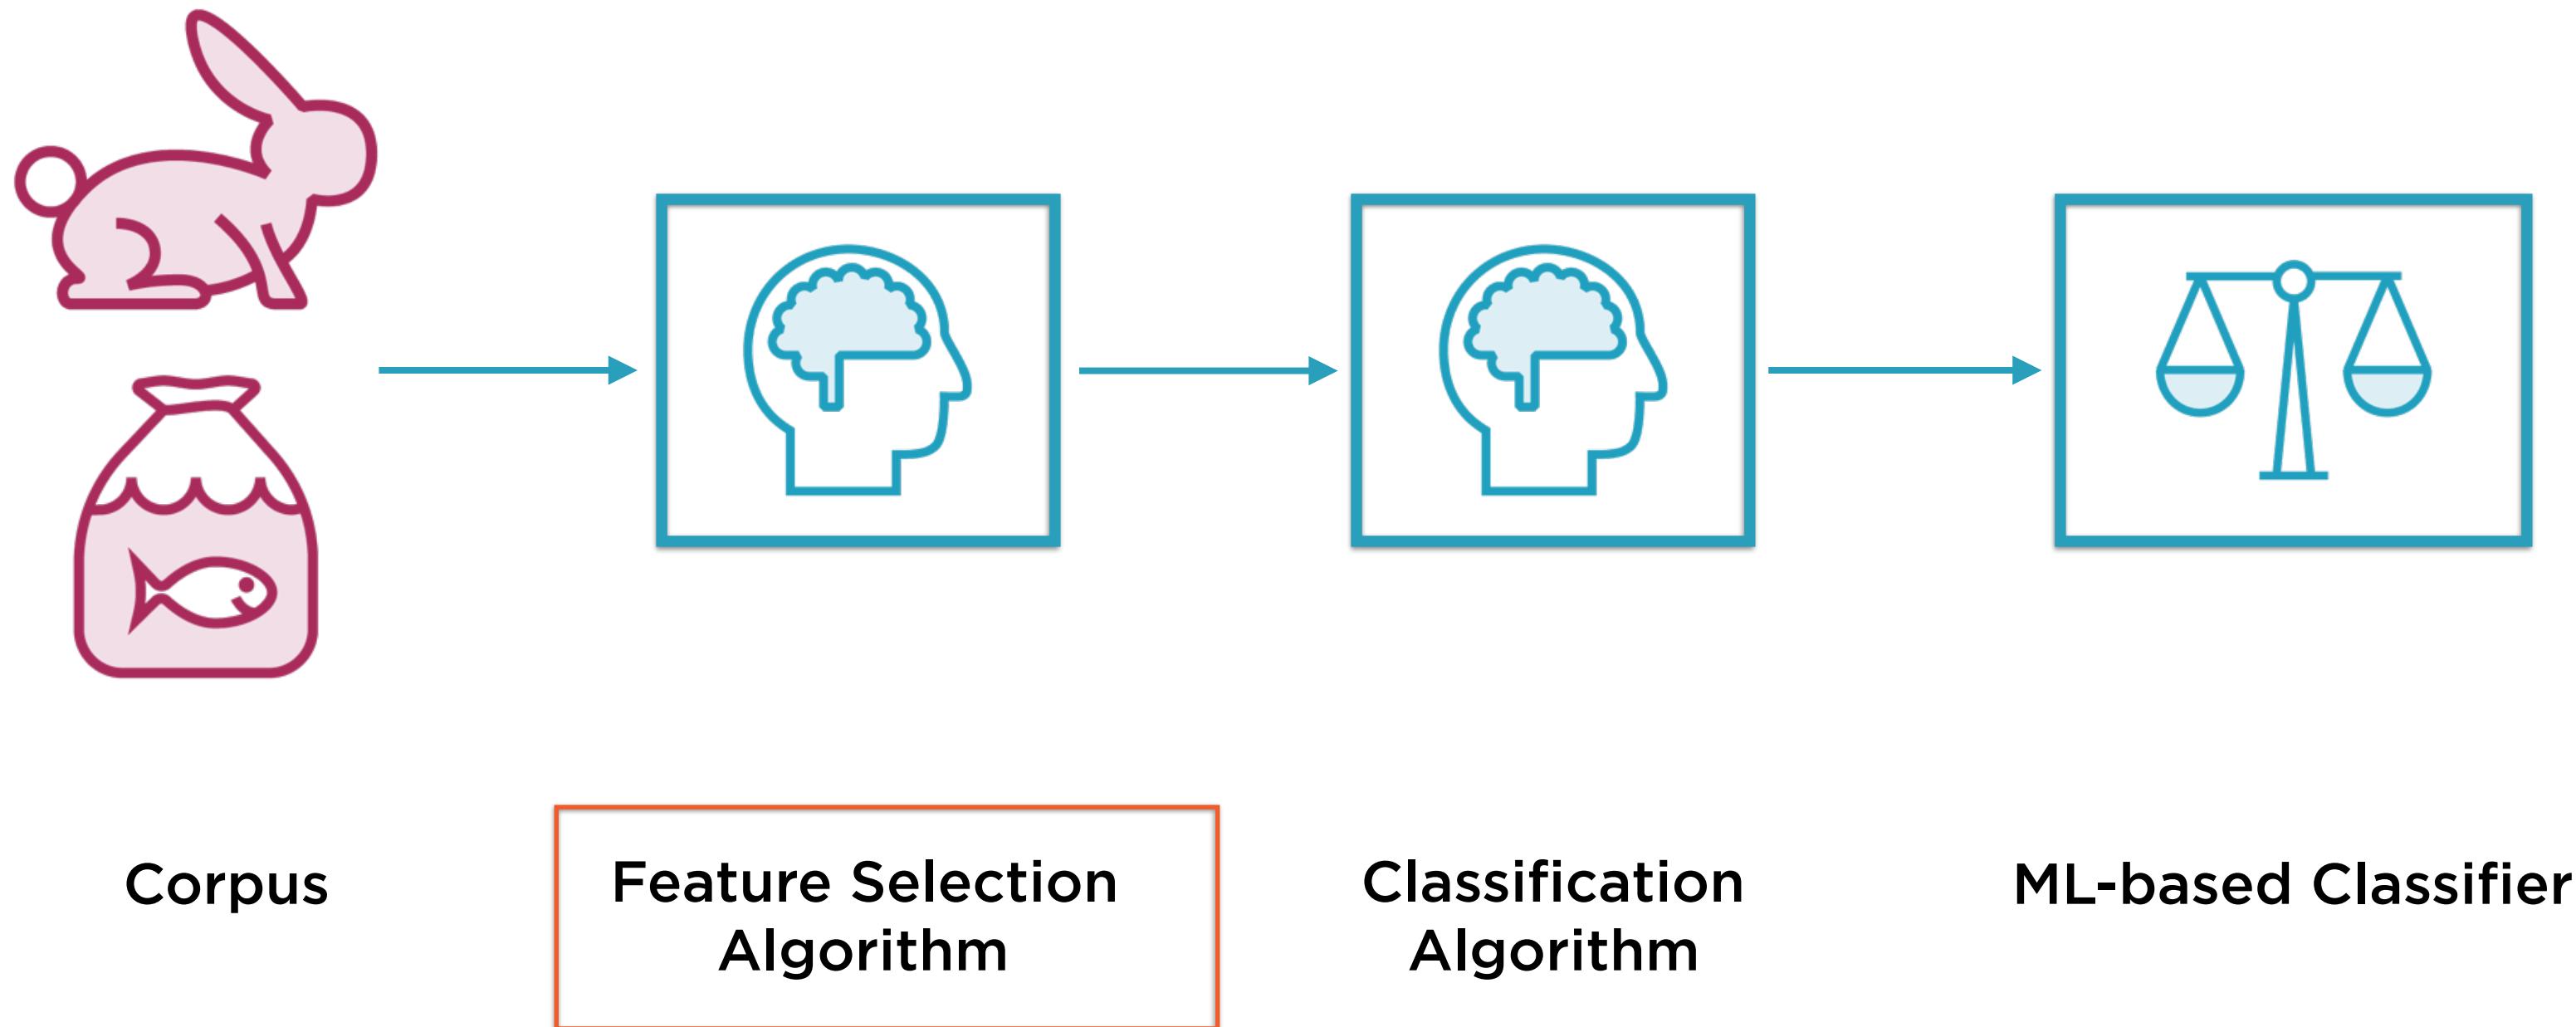
“Representation” ML-based Binary Classifier



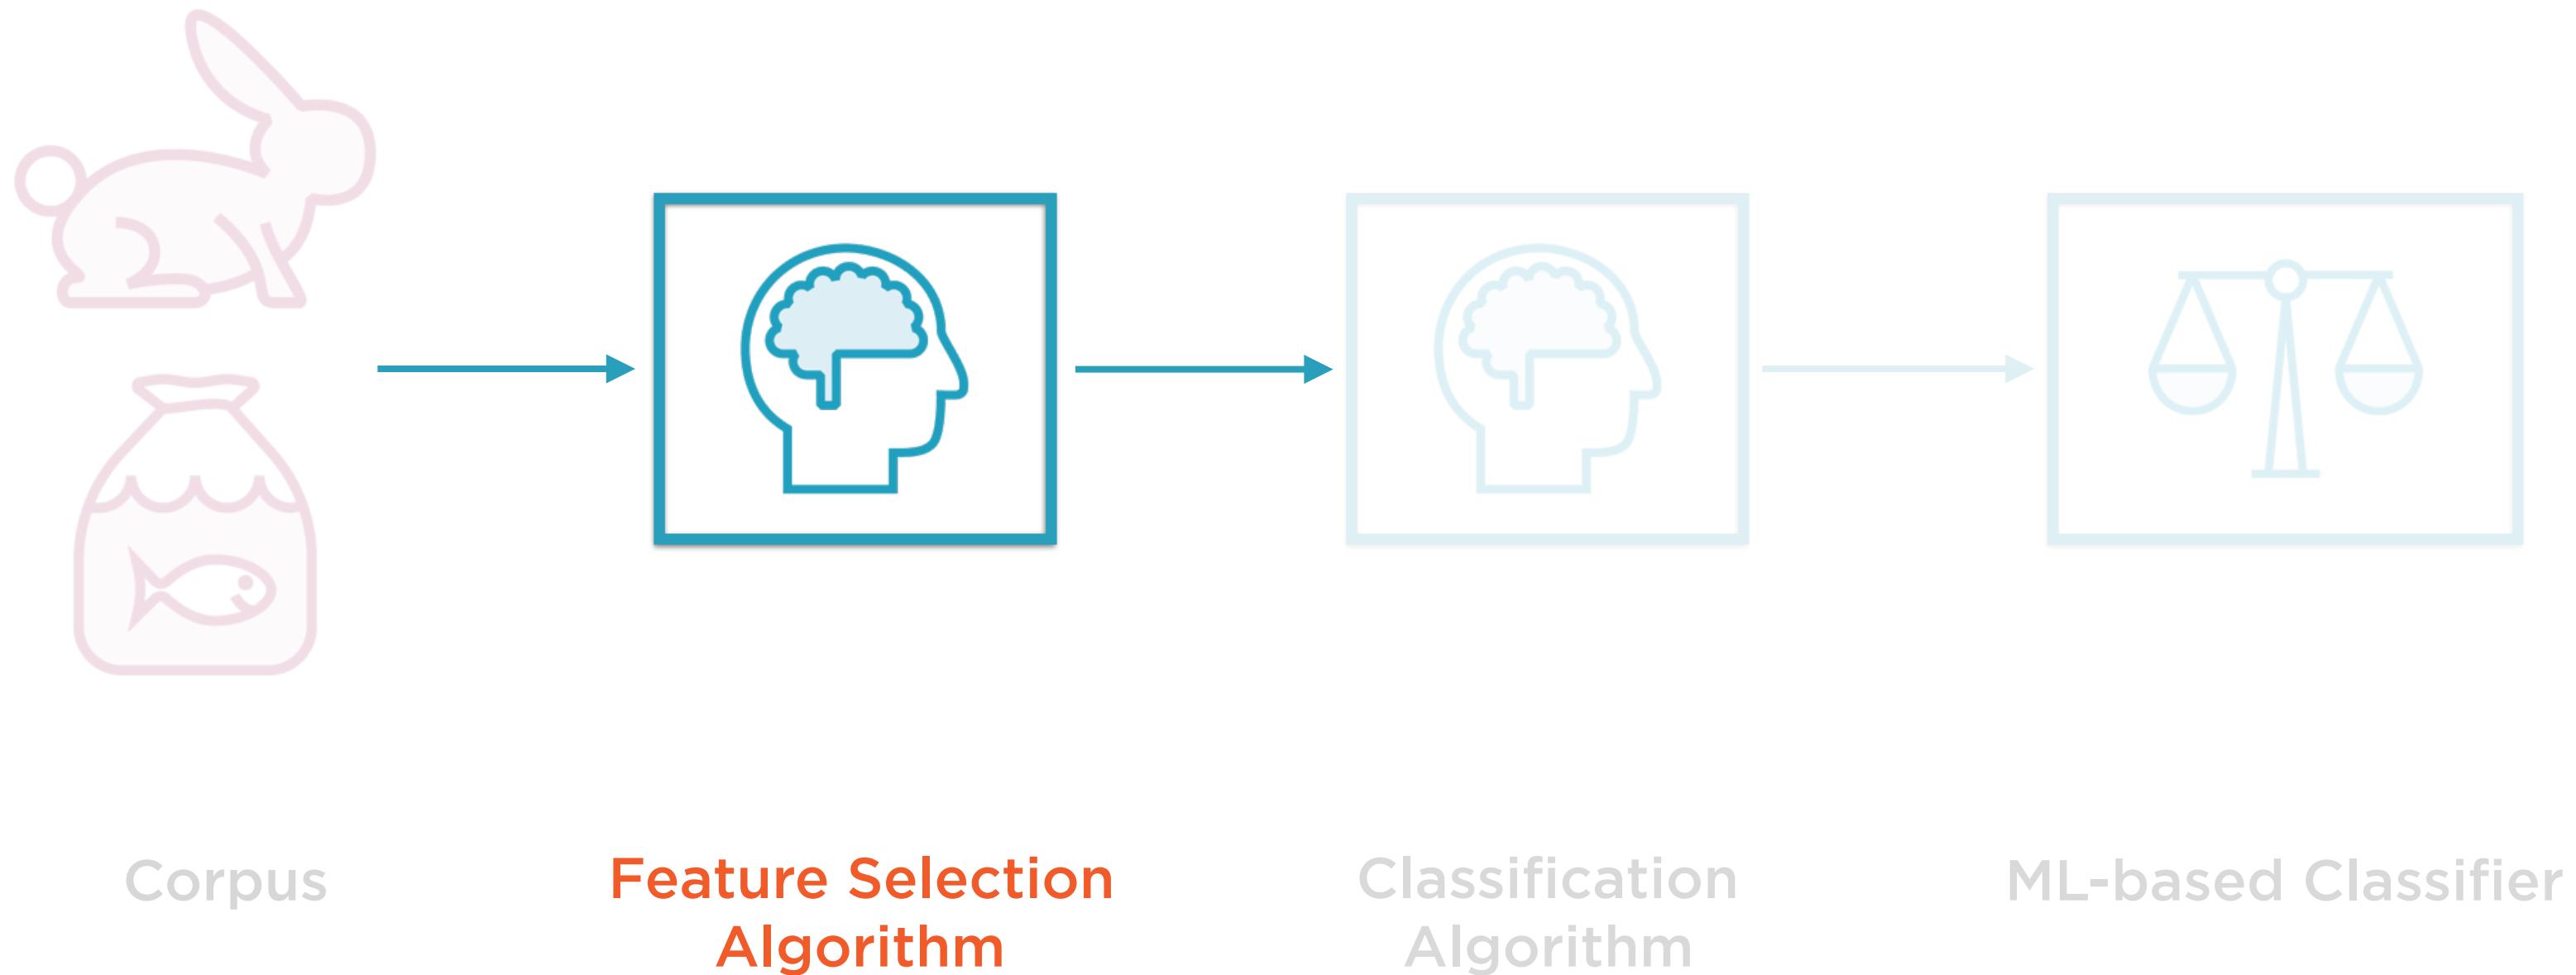
“Traditional” ML-based Binary Classifier



“Representation” ML-based Binary Classifier



“Representation” ML-based Binary Classifier



“Deep Learning” systems are one type of representation systems

Deep Learning and Neural Networks

Deep Learning

Algorithms that learn
what features matter

Neural Networks

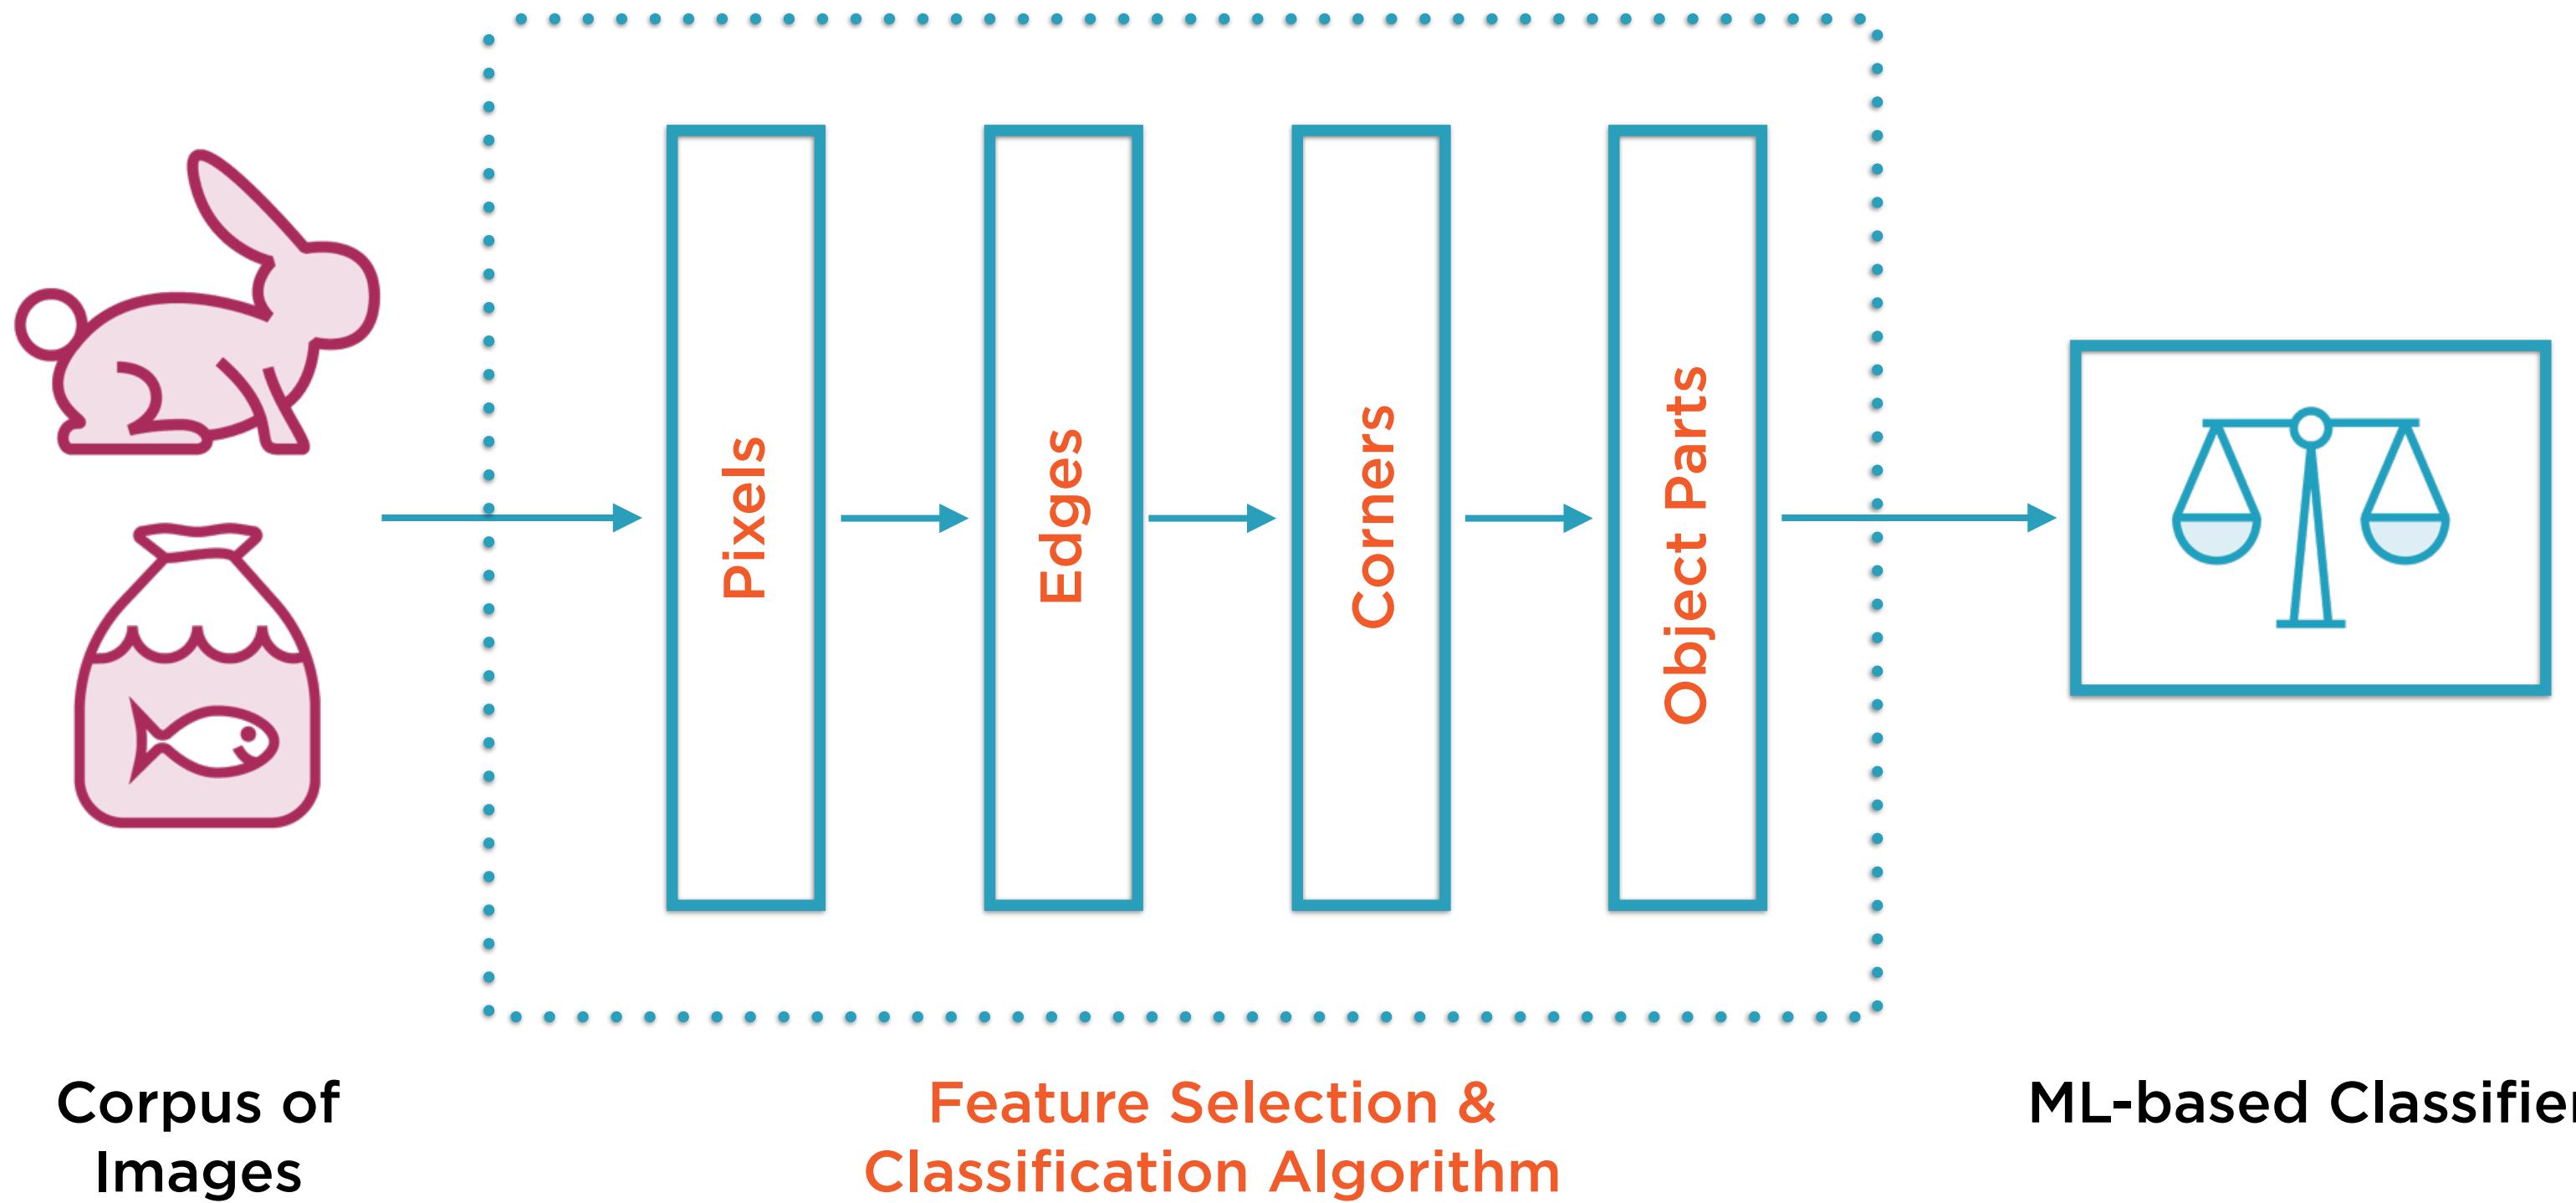
The most common class
of deep learning
algorithms

Neurons

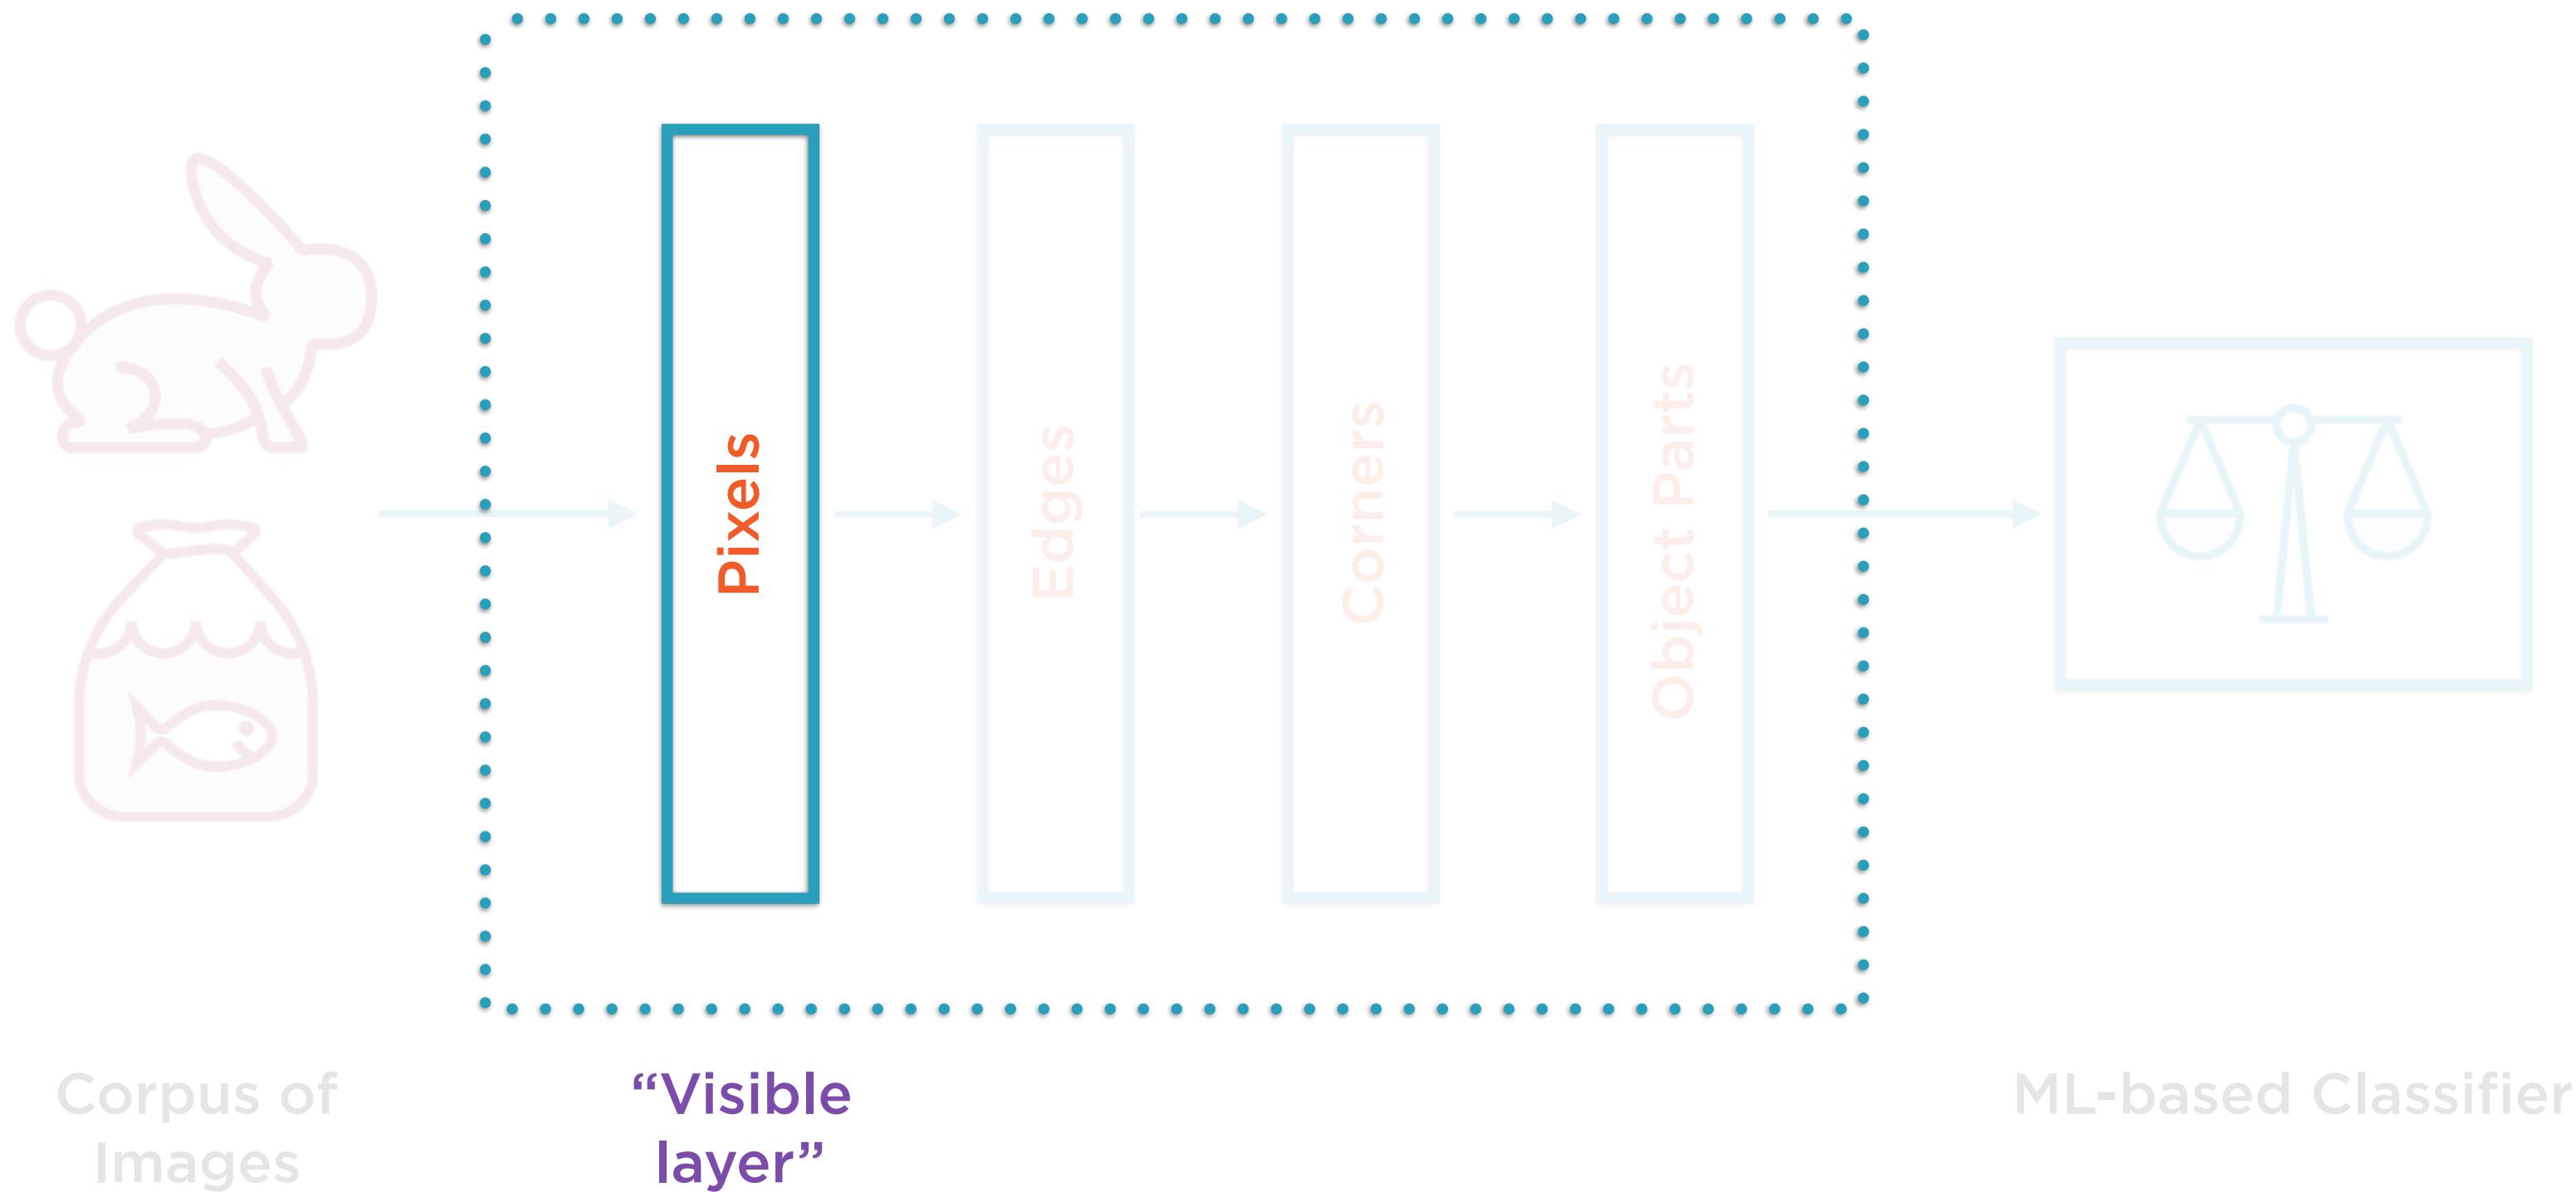
Simple building blocks
that actually “learn”

Deep Learning Book - Chapter 1 (intro), page 6

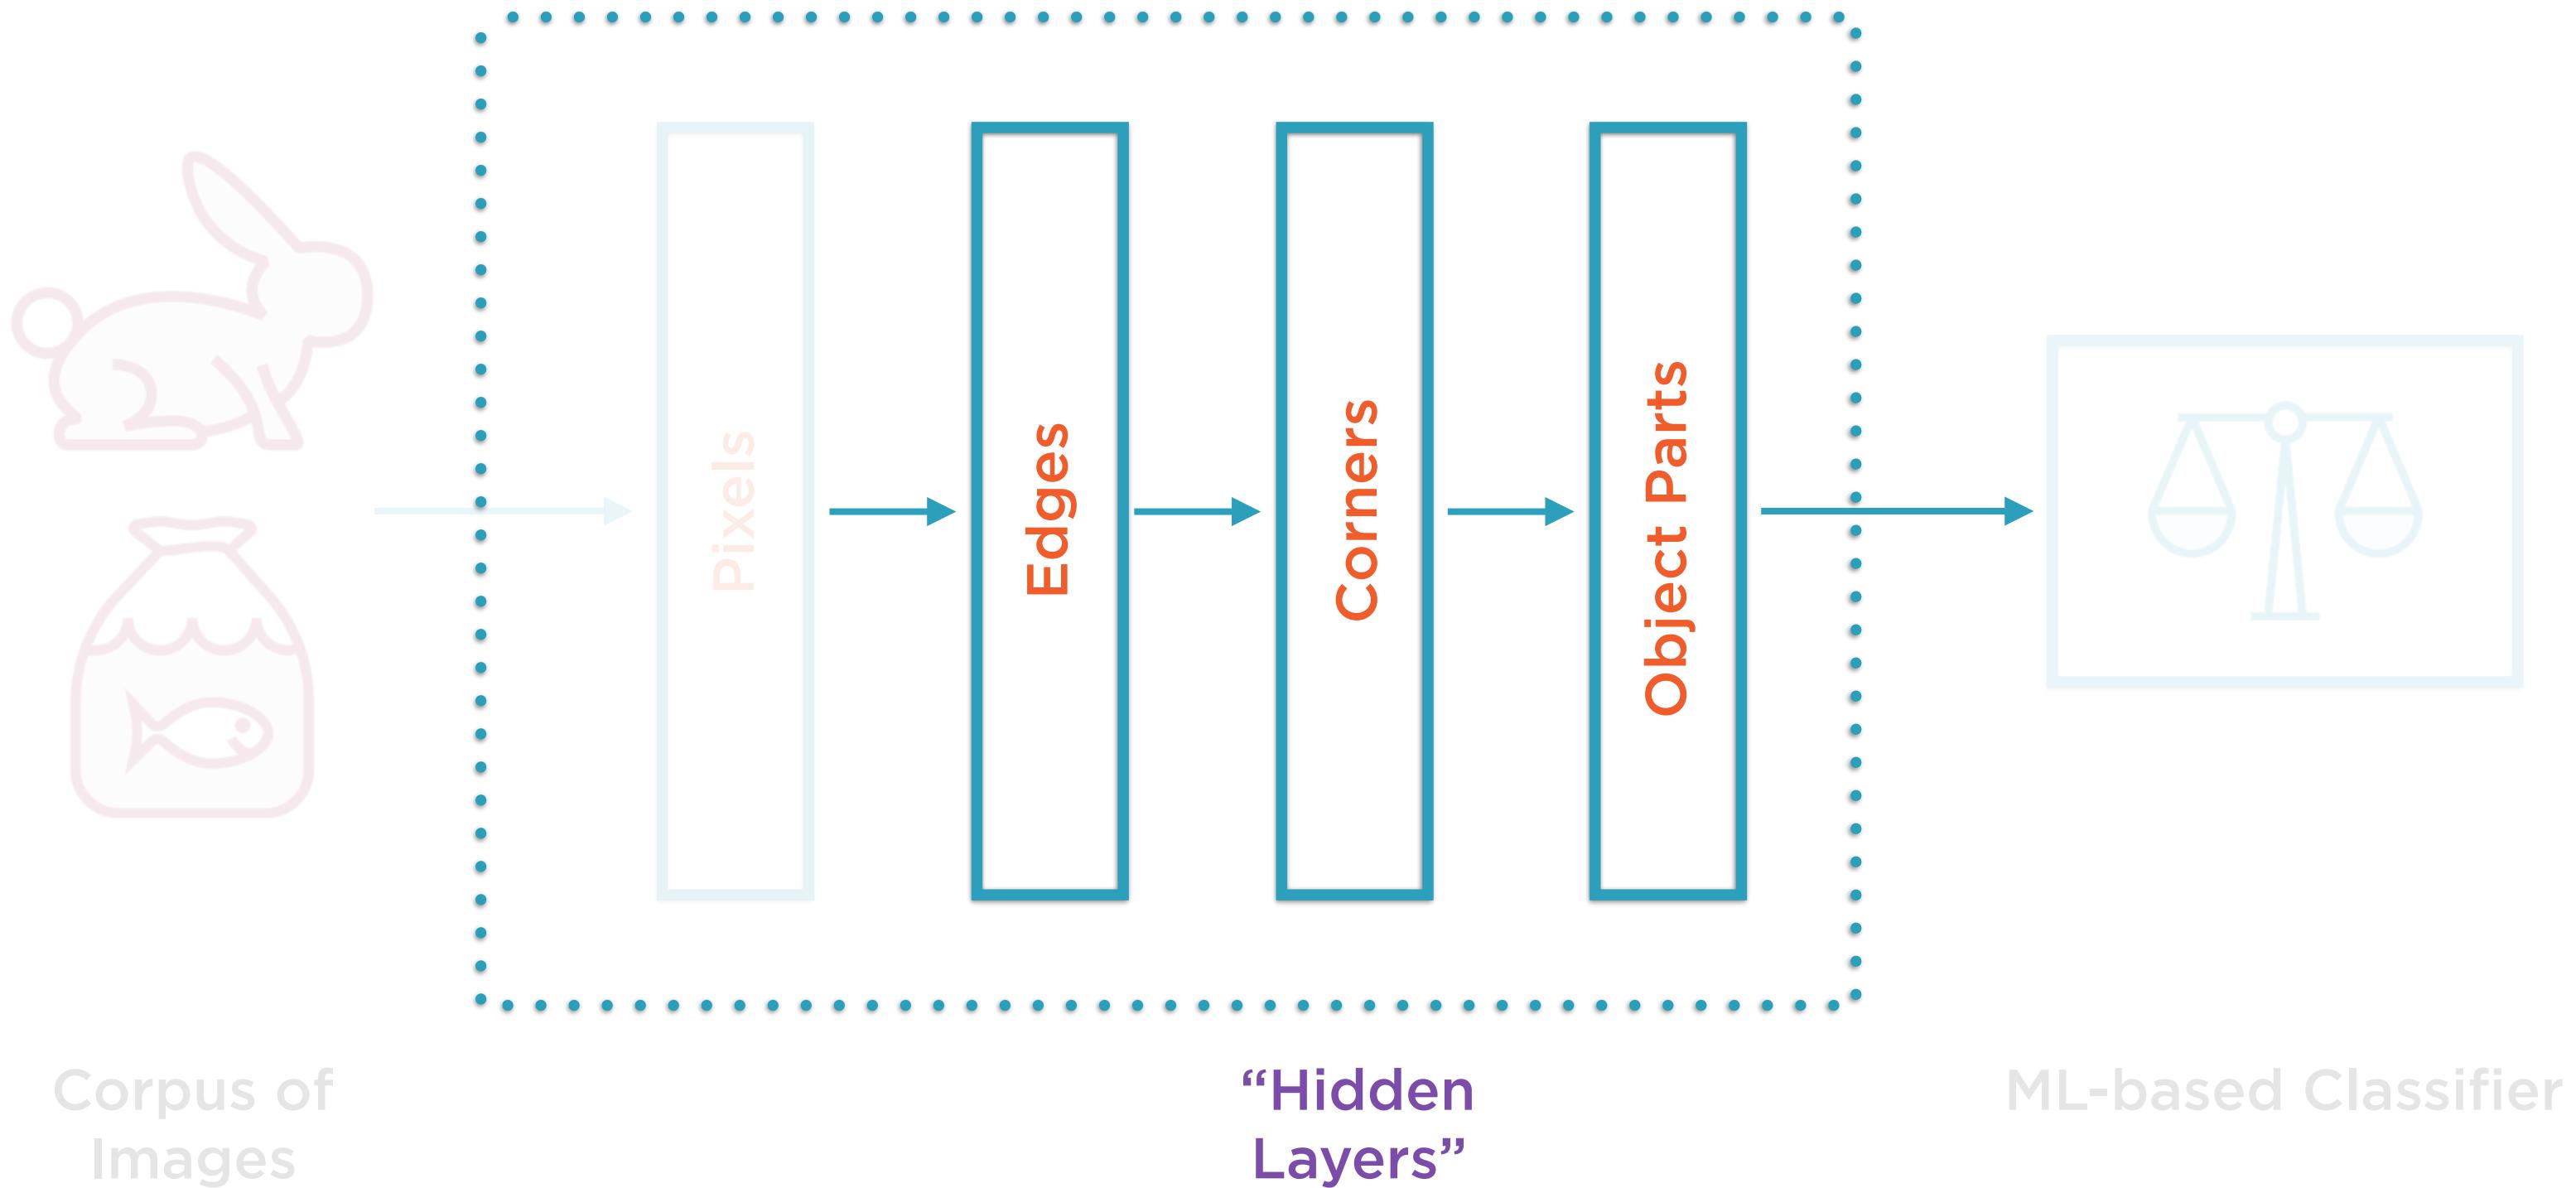
“Deep Learning”-based Binary Classifier



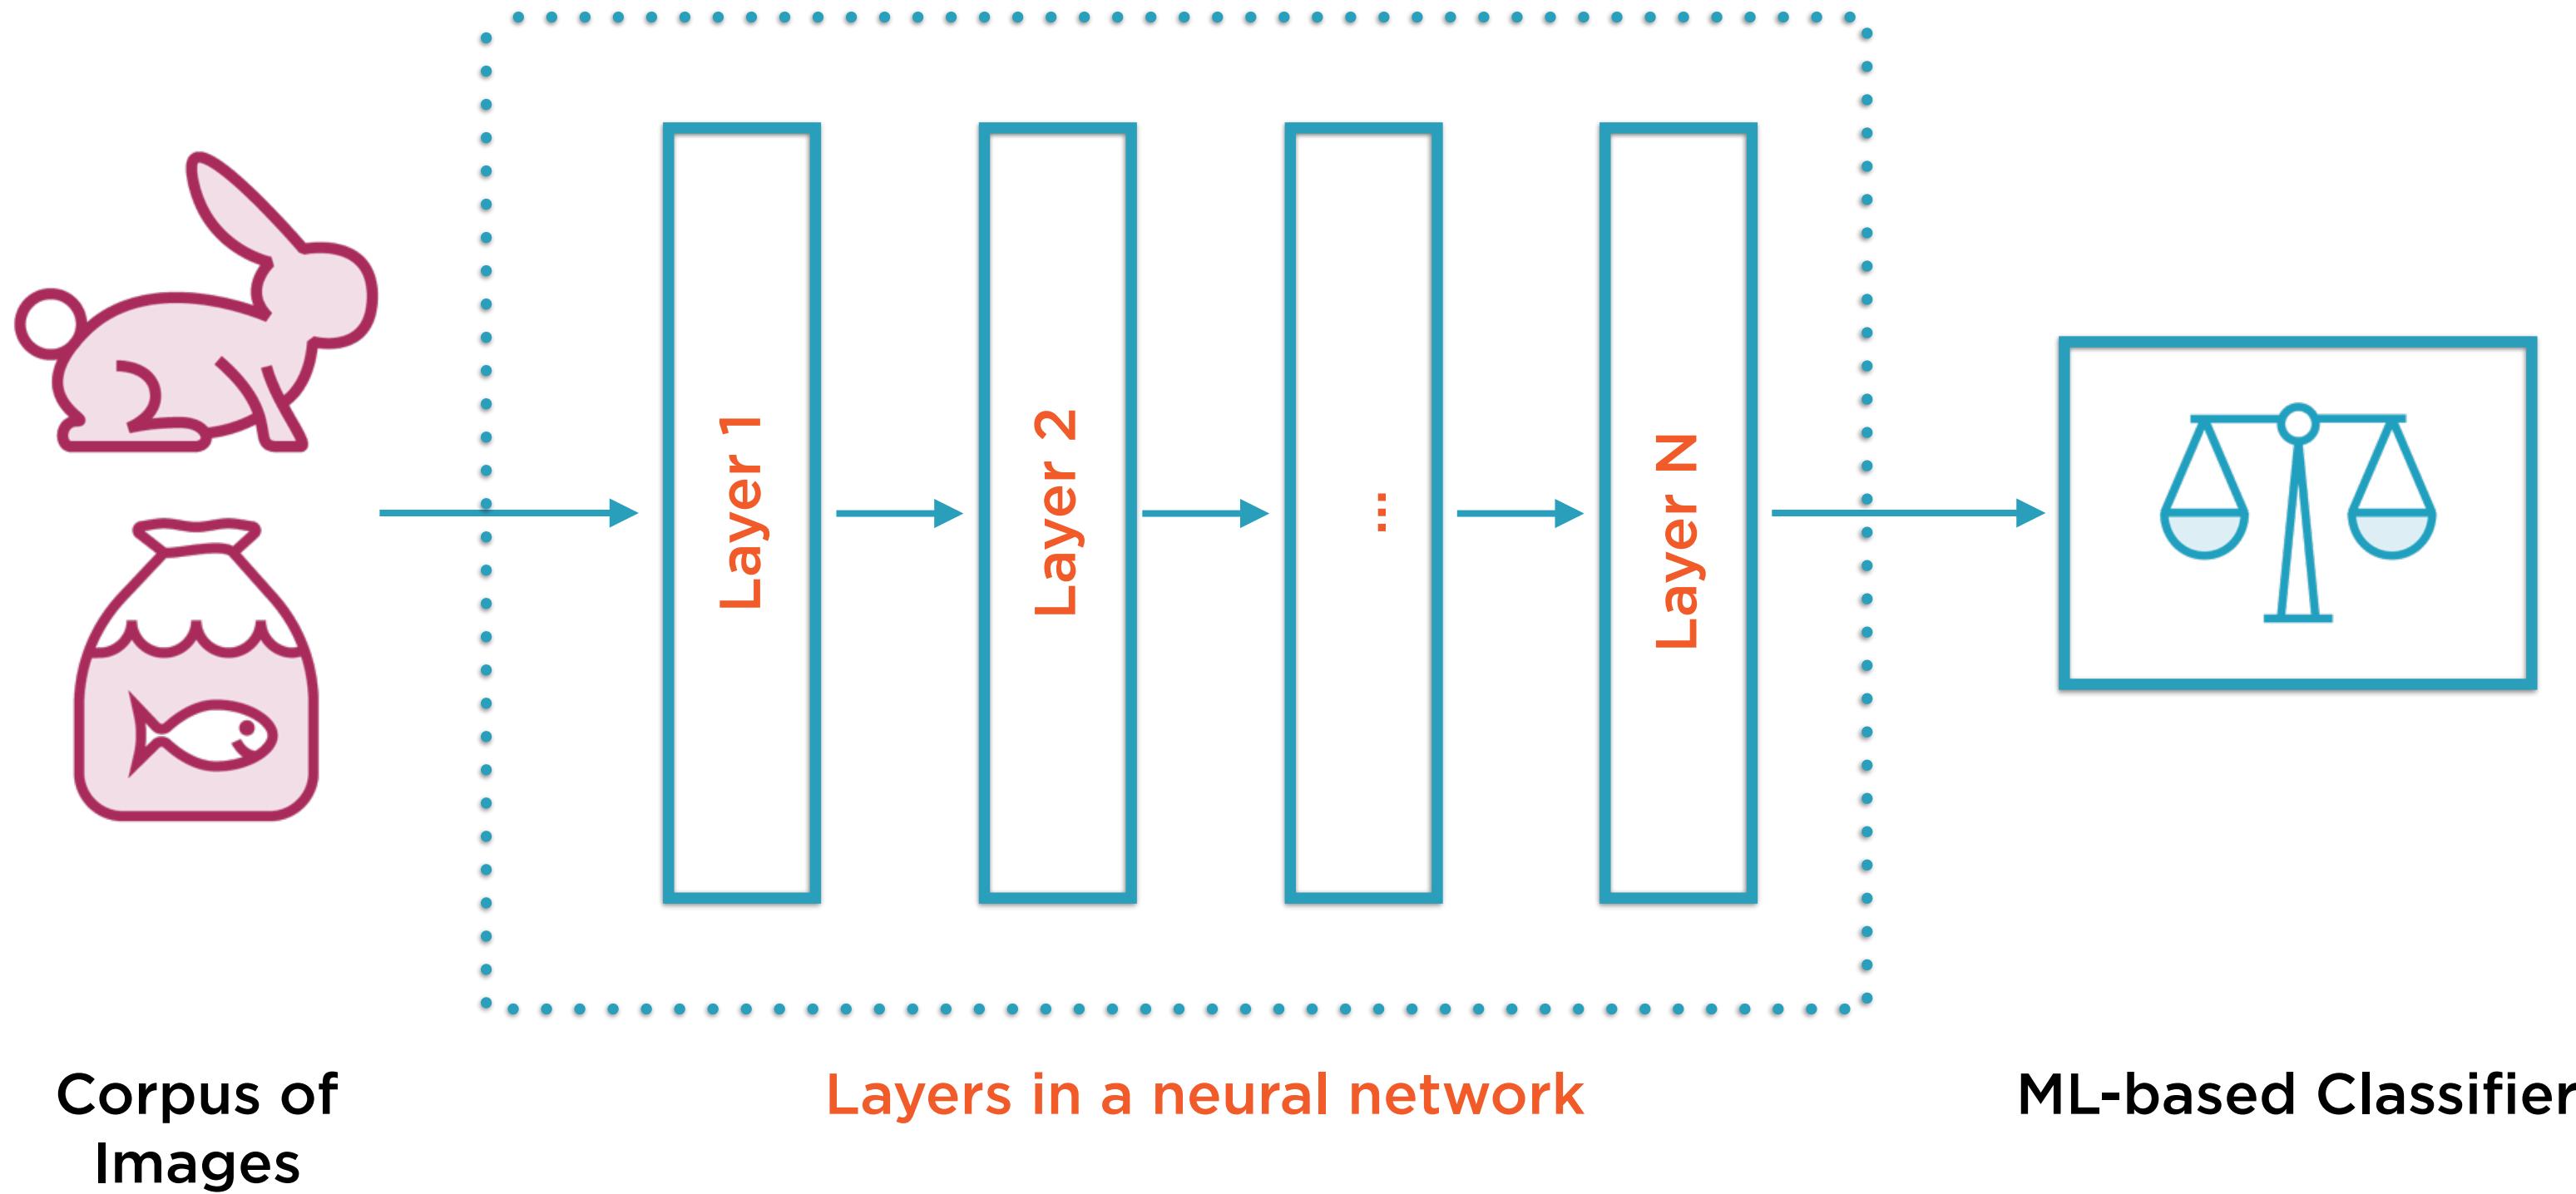
“Deep Learning”-based Binary Classifier



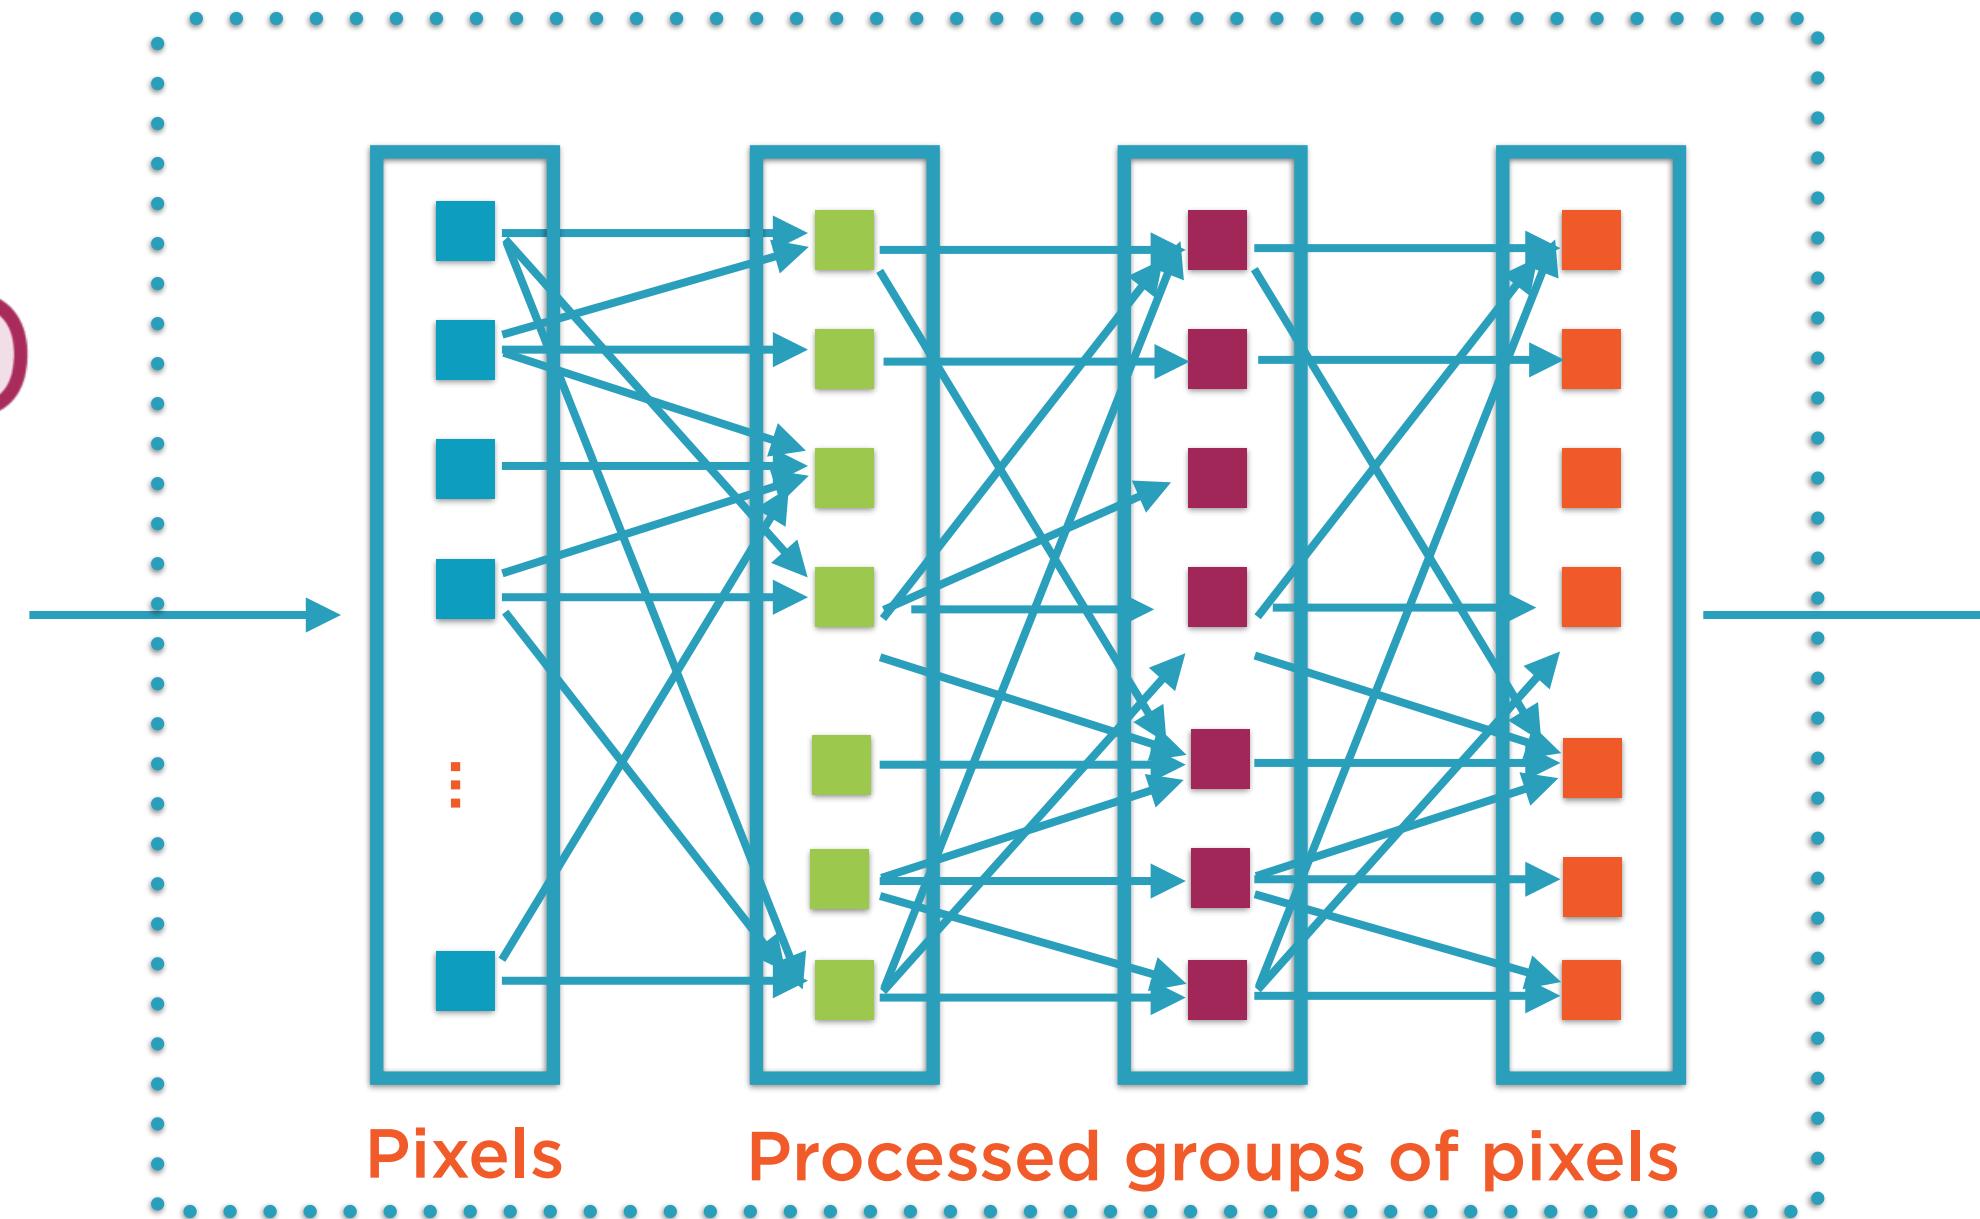
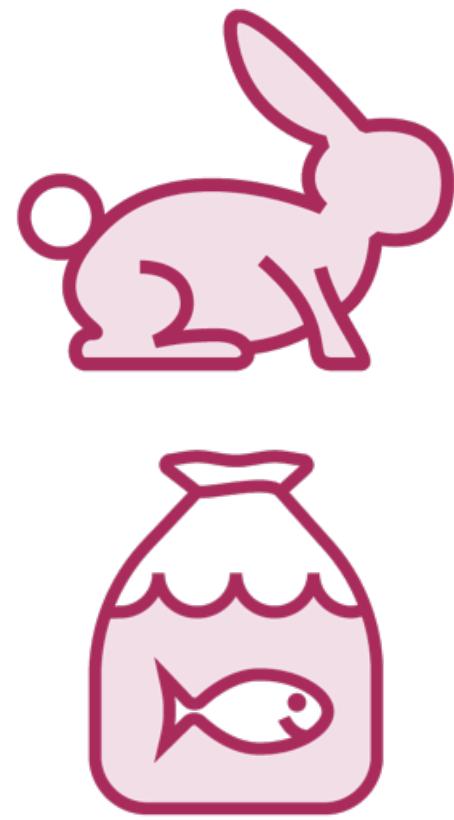
“Deep Learning”-based Binary Classifier



Neural Networks Introduced



Neural Networks Introduced



Corpus of
Images

Each layer consists of individual
interconnected neurons

ML-based Classifier

Neurons as Learning Units

A machine learning algorithm is an algorithm that is able to learn from data

Learning Algorithms

A computer program is said to learn from **experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in T, as measured by P, **improves** with experience E

Most common tasks in ML:
Classification, regression

Learning Algorithms

A computer program is said to learn from experience E with respect to some class of **tasks** T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E



Learning Algorithms

A computer program is said to learn from experience E with respect to some class of tasks T and **performance measure** P, if its performance at tasks in T, as measured by P, improves with experience E

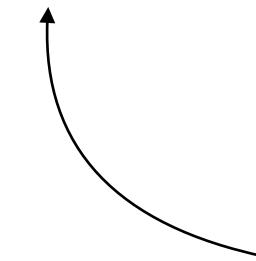
Accuracy in classification,
residual variance in regression



Learning Algorithms

A computer program is said to learn from **experience E** with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E

Training using a corpus
of labelled instances



Learning Algorithms

A computer program is said to **learn from experience E** with respect to some class of tasks T and performance measure P, if its **performance at tasks** in T, as measured by P, **improves with experience E**

Deep Learning and Neural Networks

Deep Learning

Algorithms that learn
what features matter

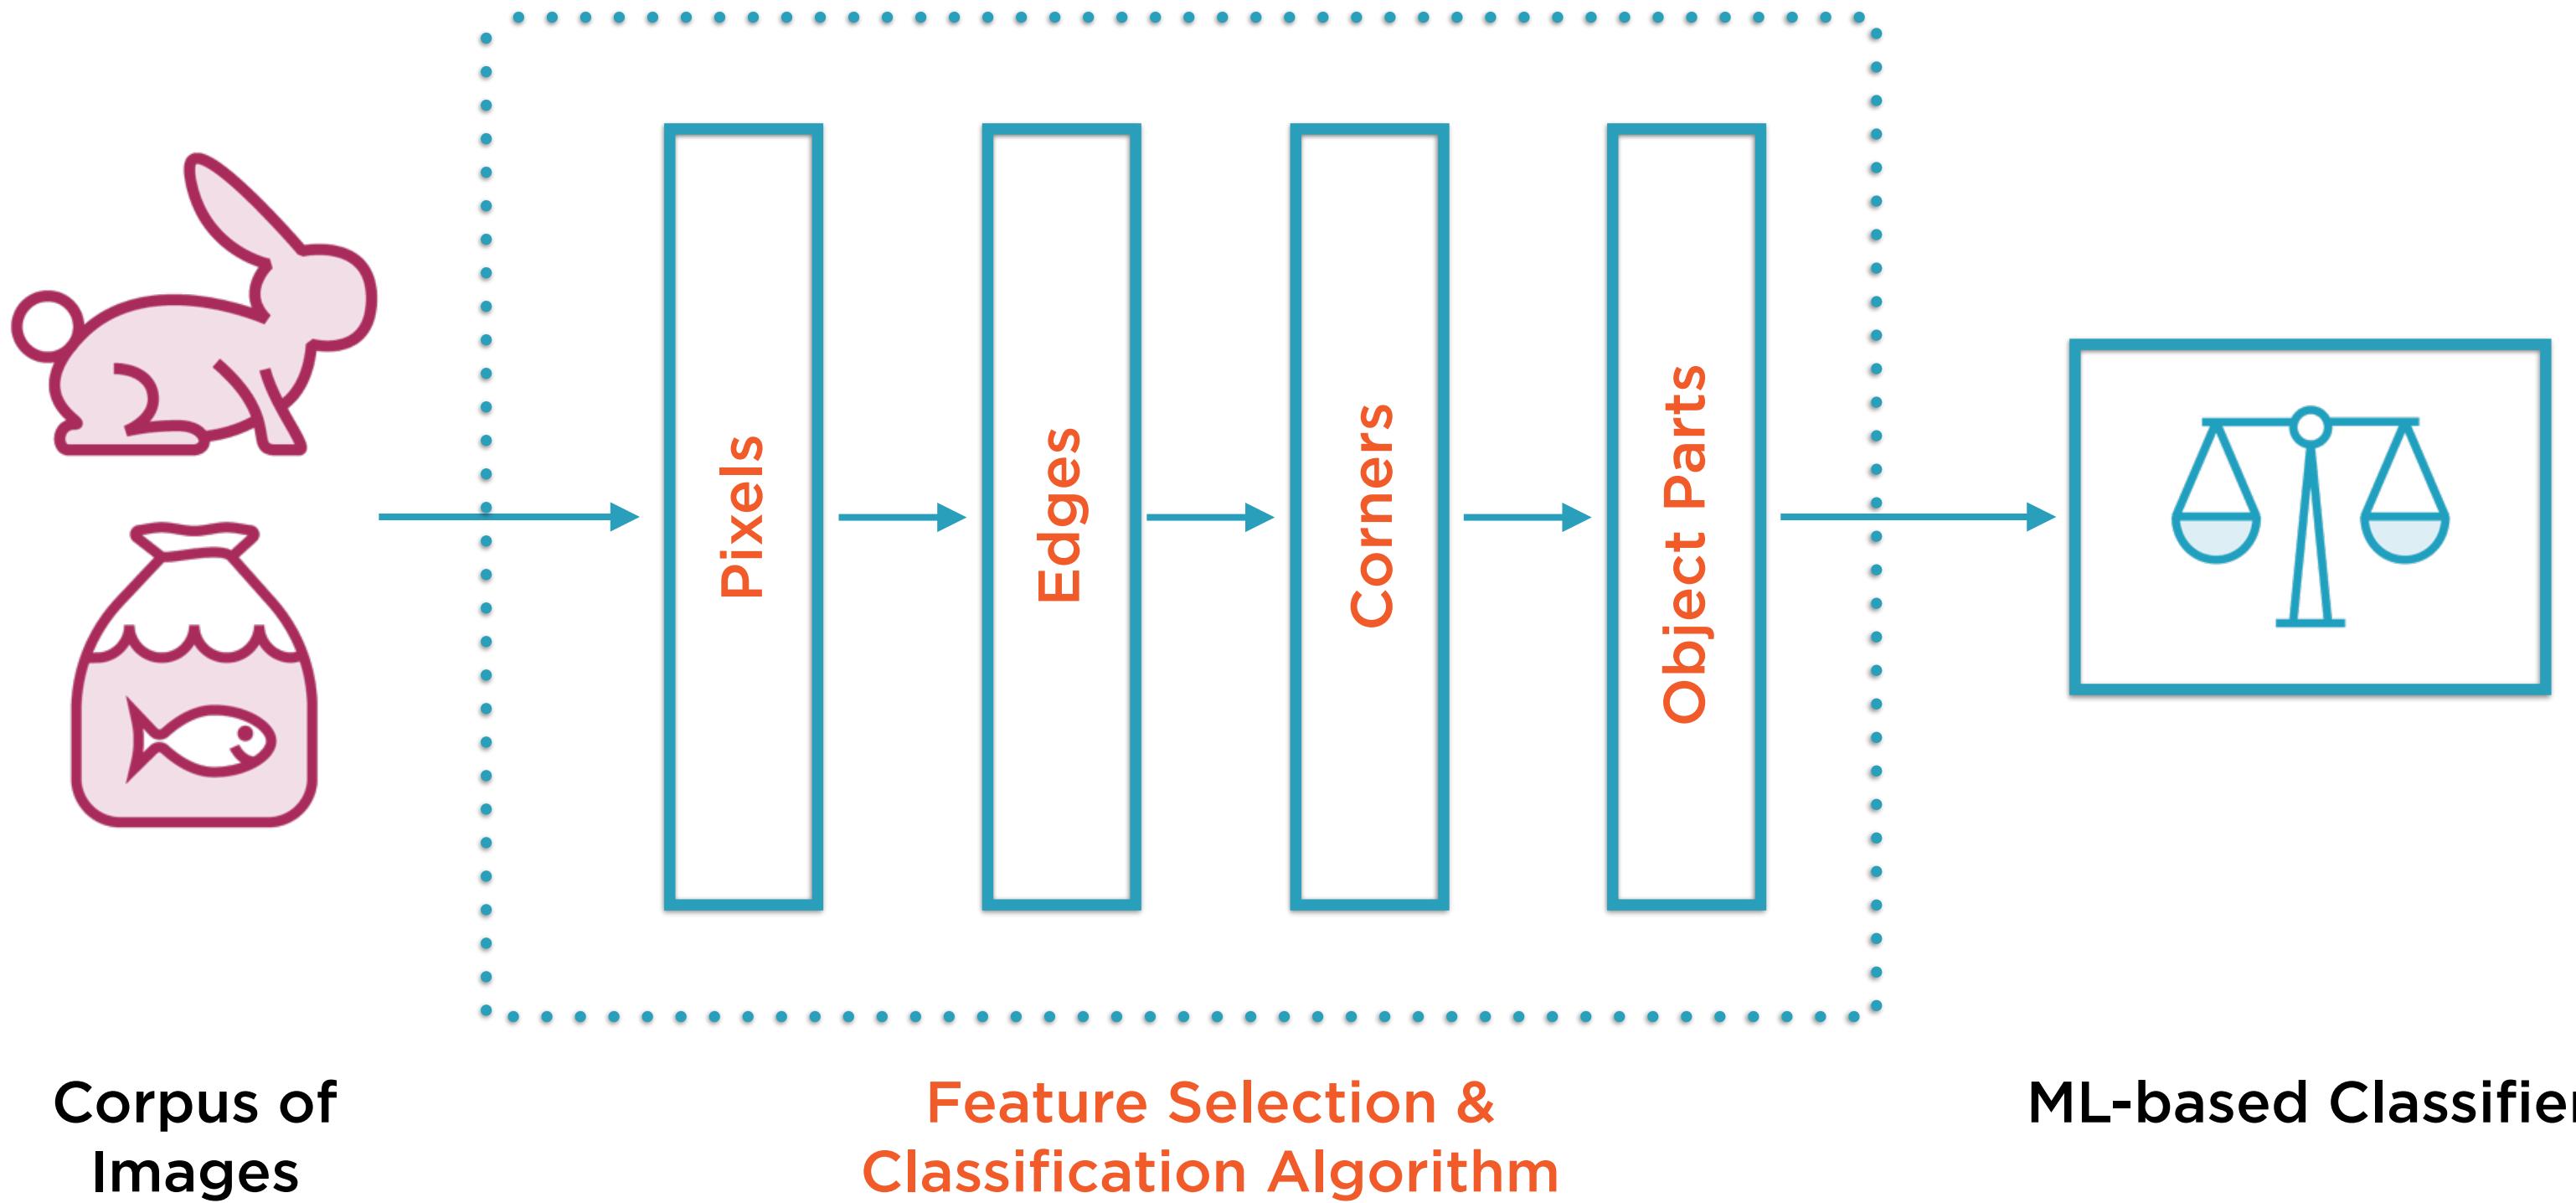
Neural Networks

The most common class
of deep learning
algorithms

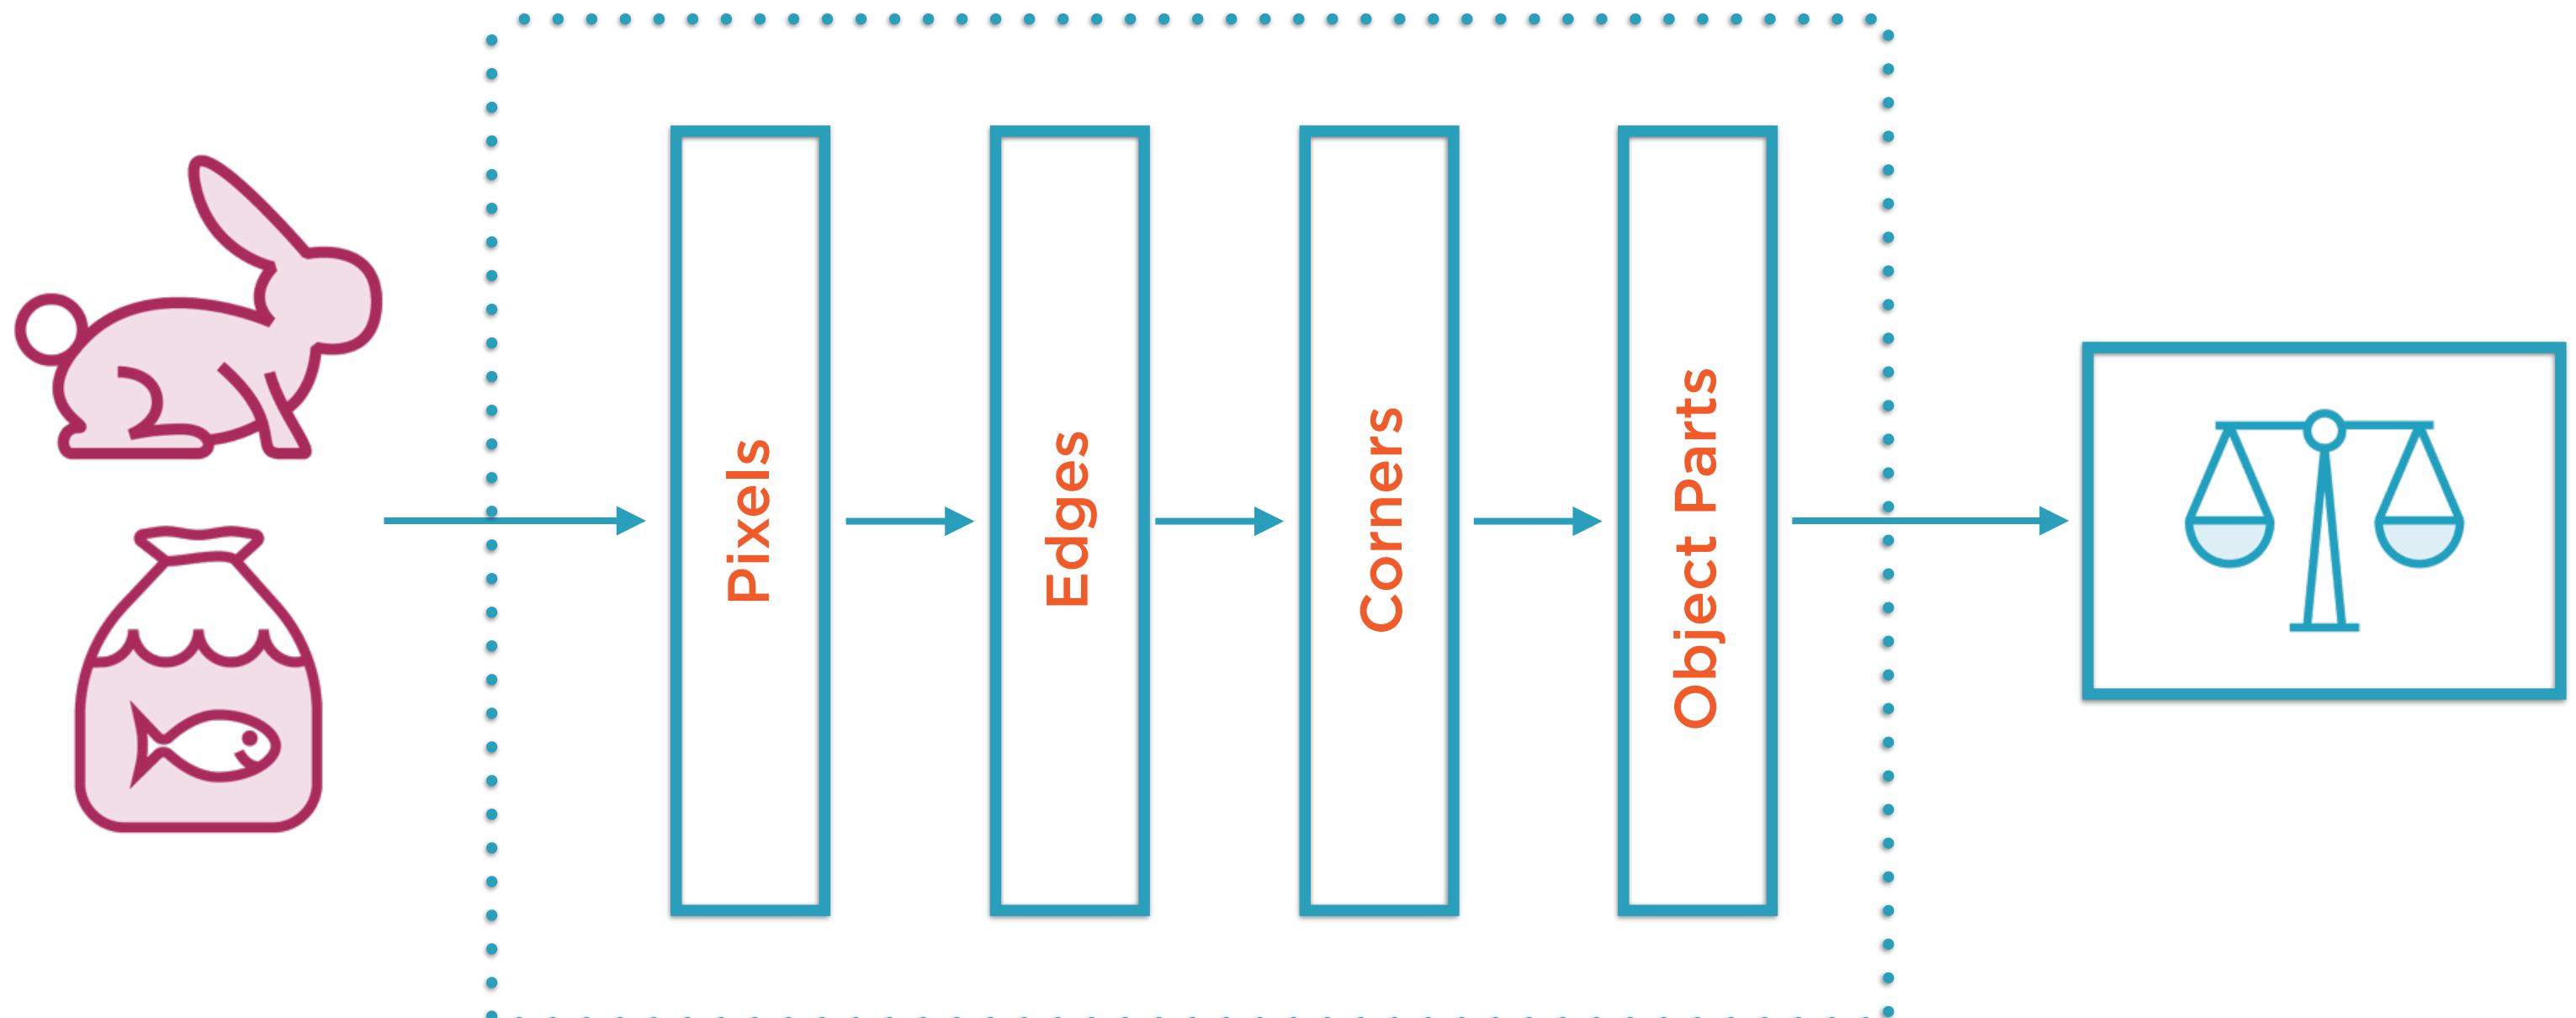
Neurons

Simple building blocks
that actually “learn”

“Deep Learning”-based Binary Classifier



Layers in the Computation Graph

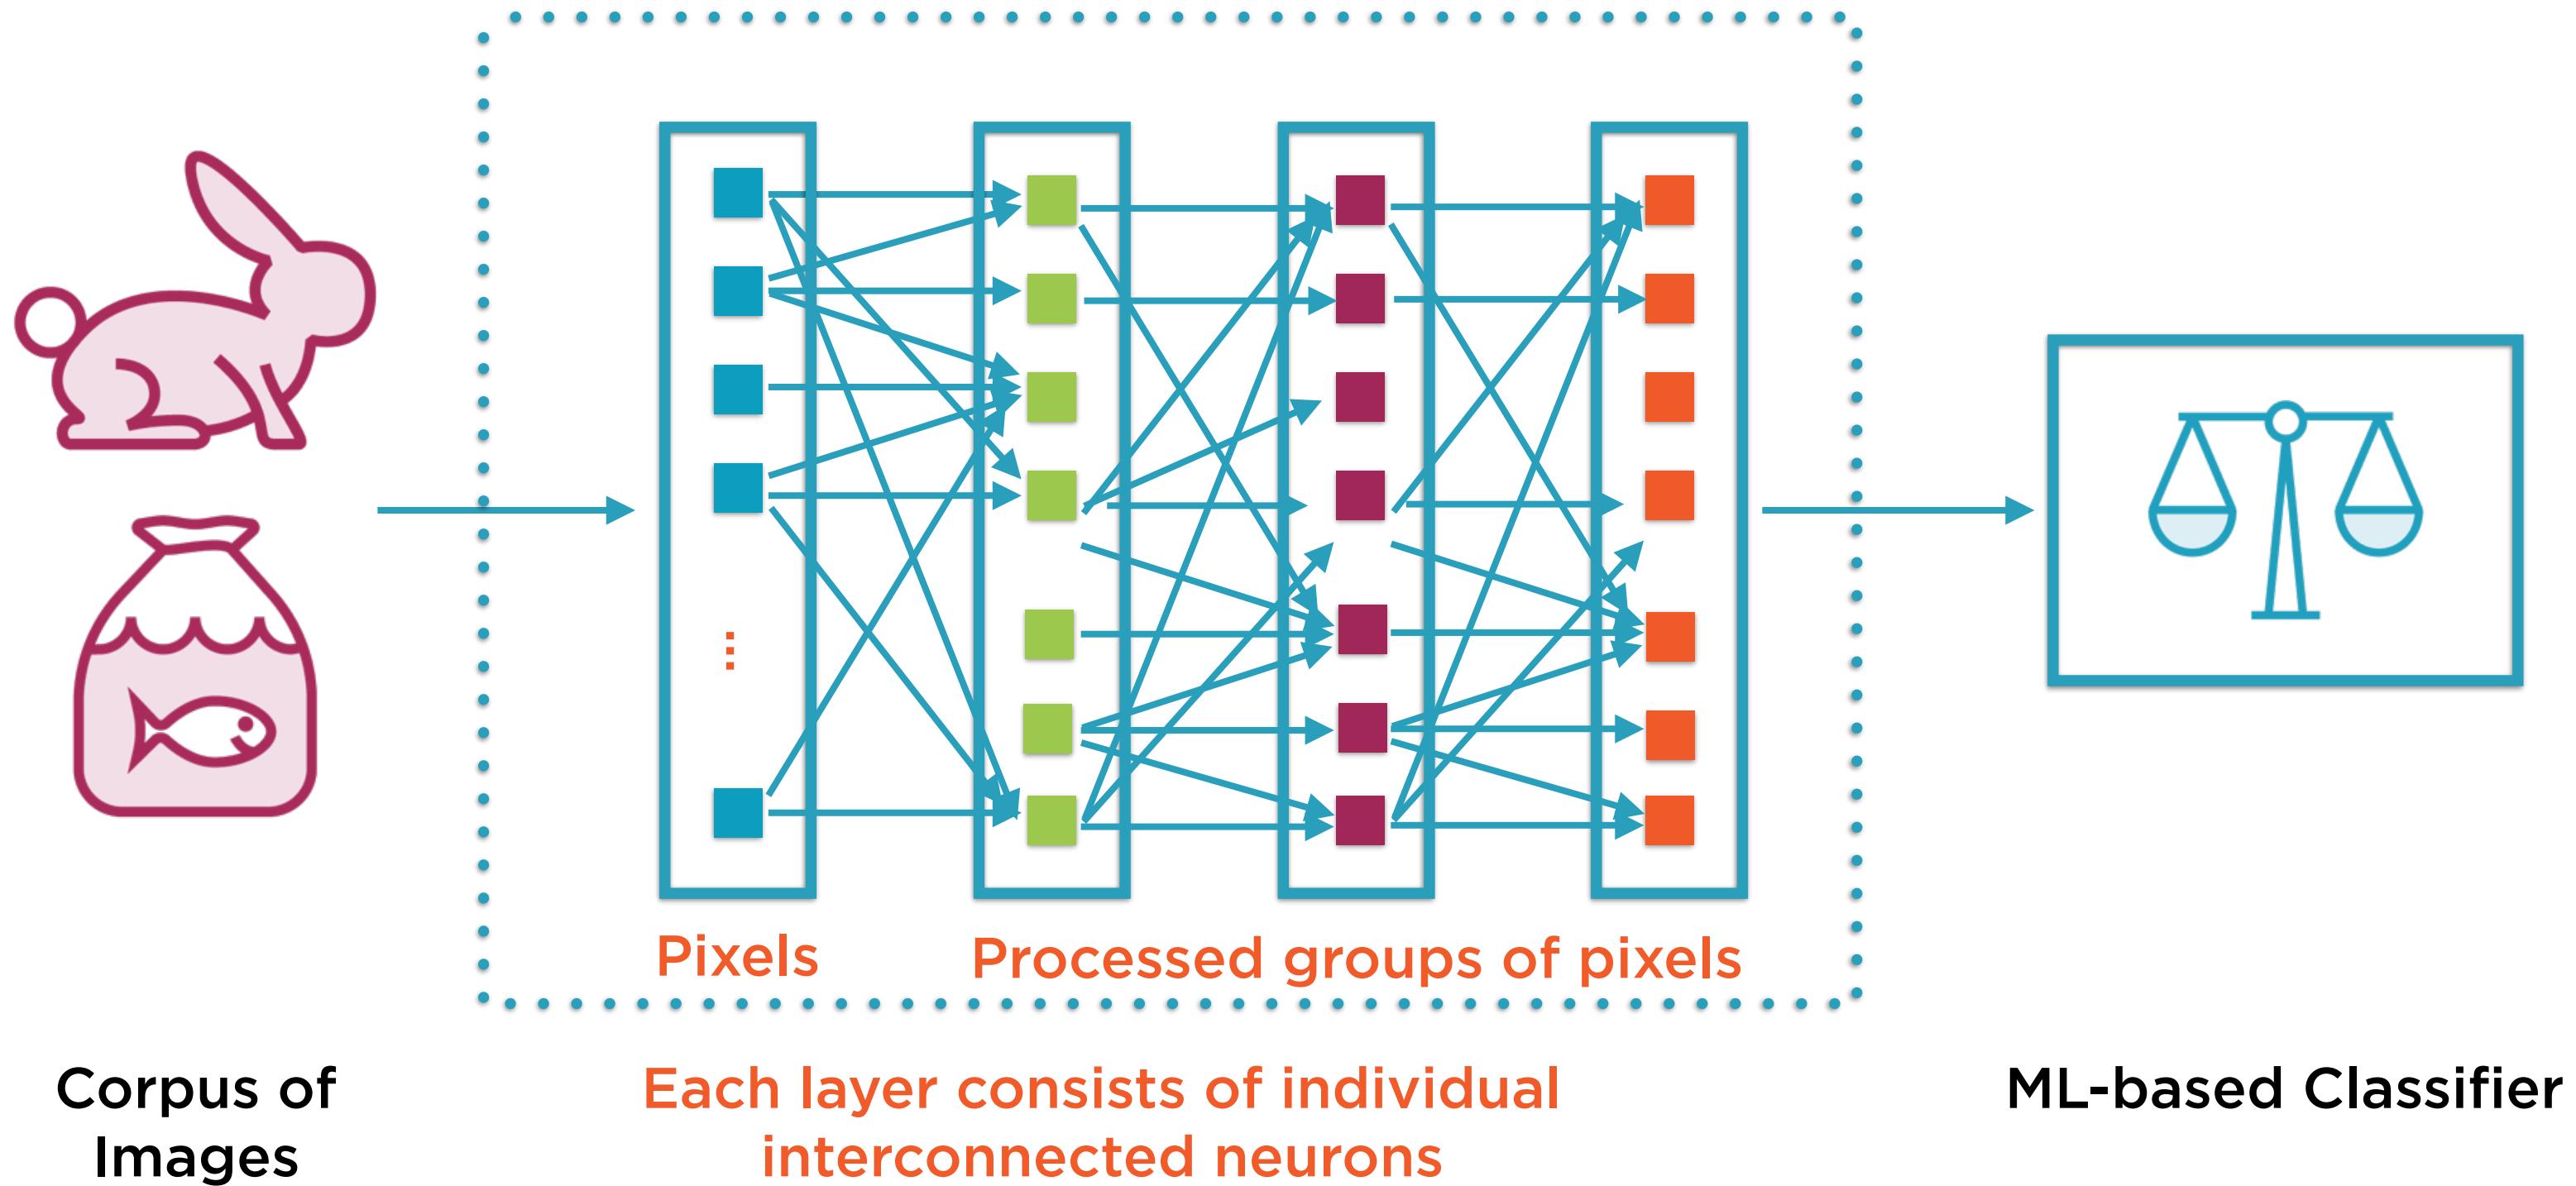


**Corpus of
Images**

**Groups of neurons that perform similar
functions are aggregated into layers**

ML-based Classifier

Neural Networks: Networks of Neurons



Deep Learning

Directed computation graphs “learn” relationships between data

The more complex the graph, the more the relationships it can “learn”

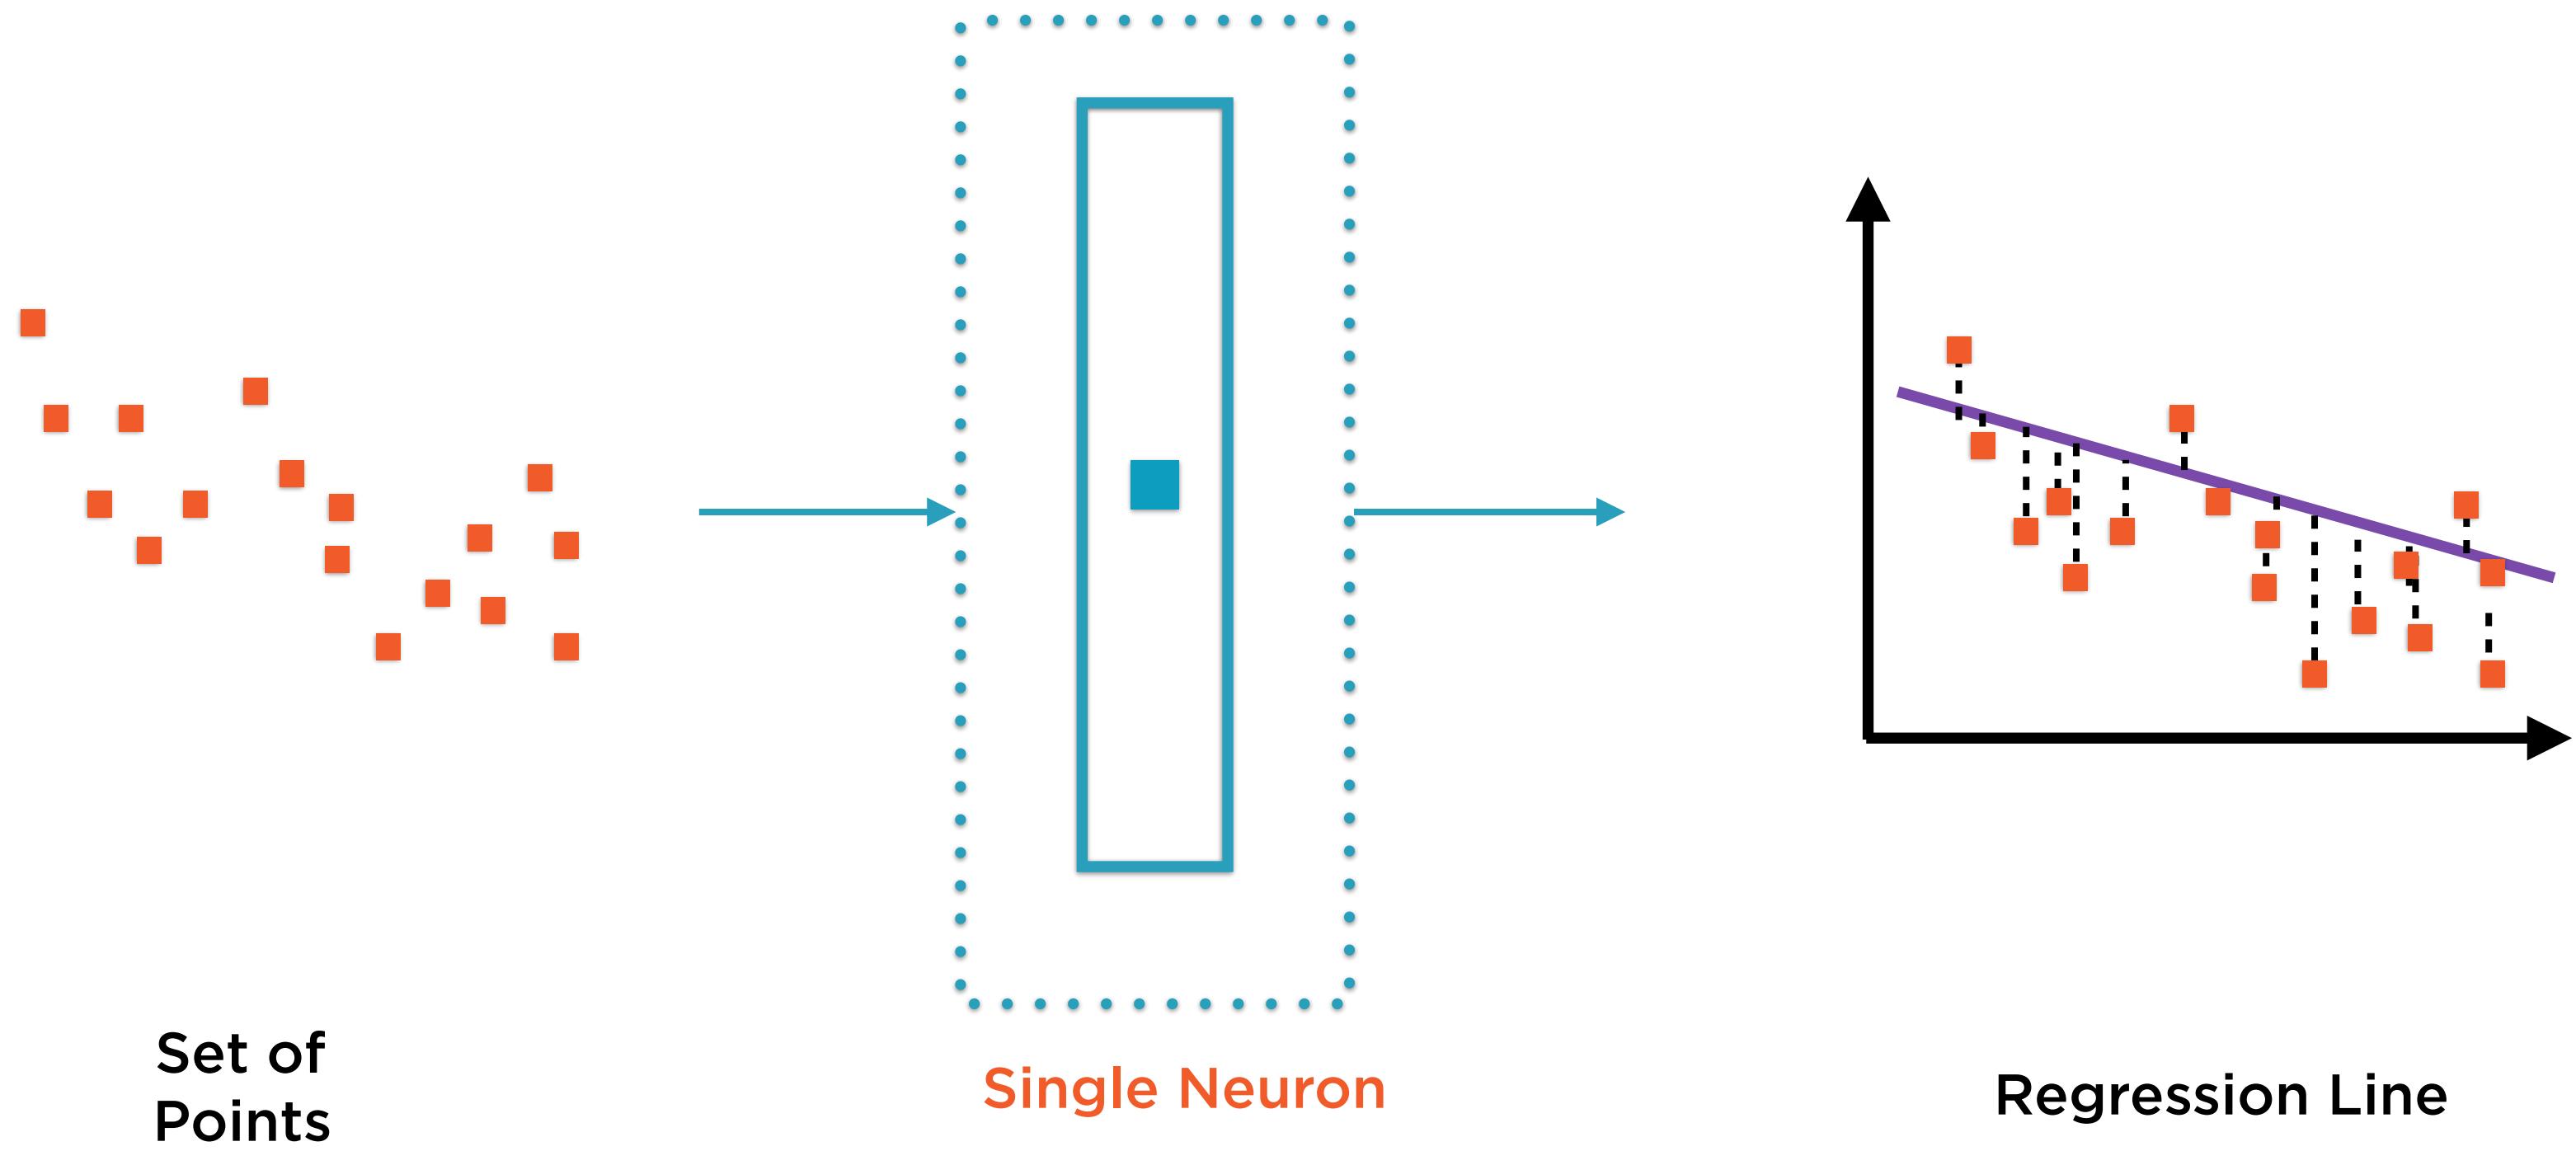
“Deep” Learning: Depth of the computation graph

$$y = Wx + b$$

“Learning” Regression

Regression can be reverse-engineered by a single neuron

Regression: The Simplest Neural Network

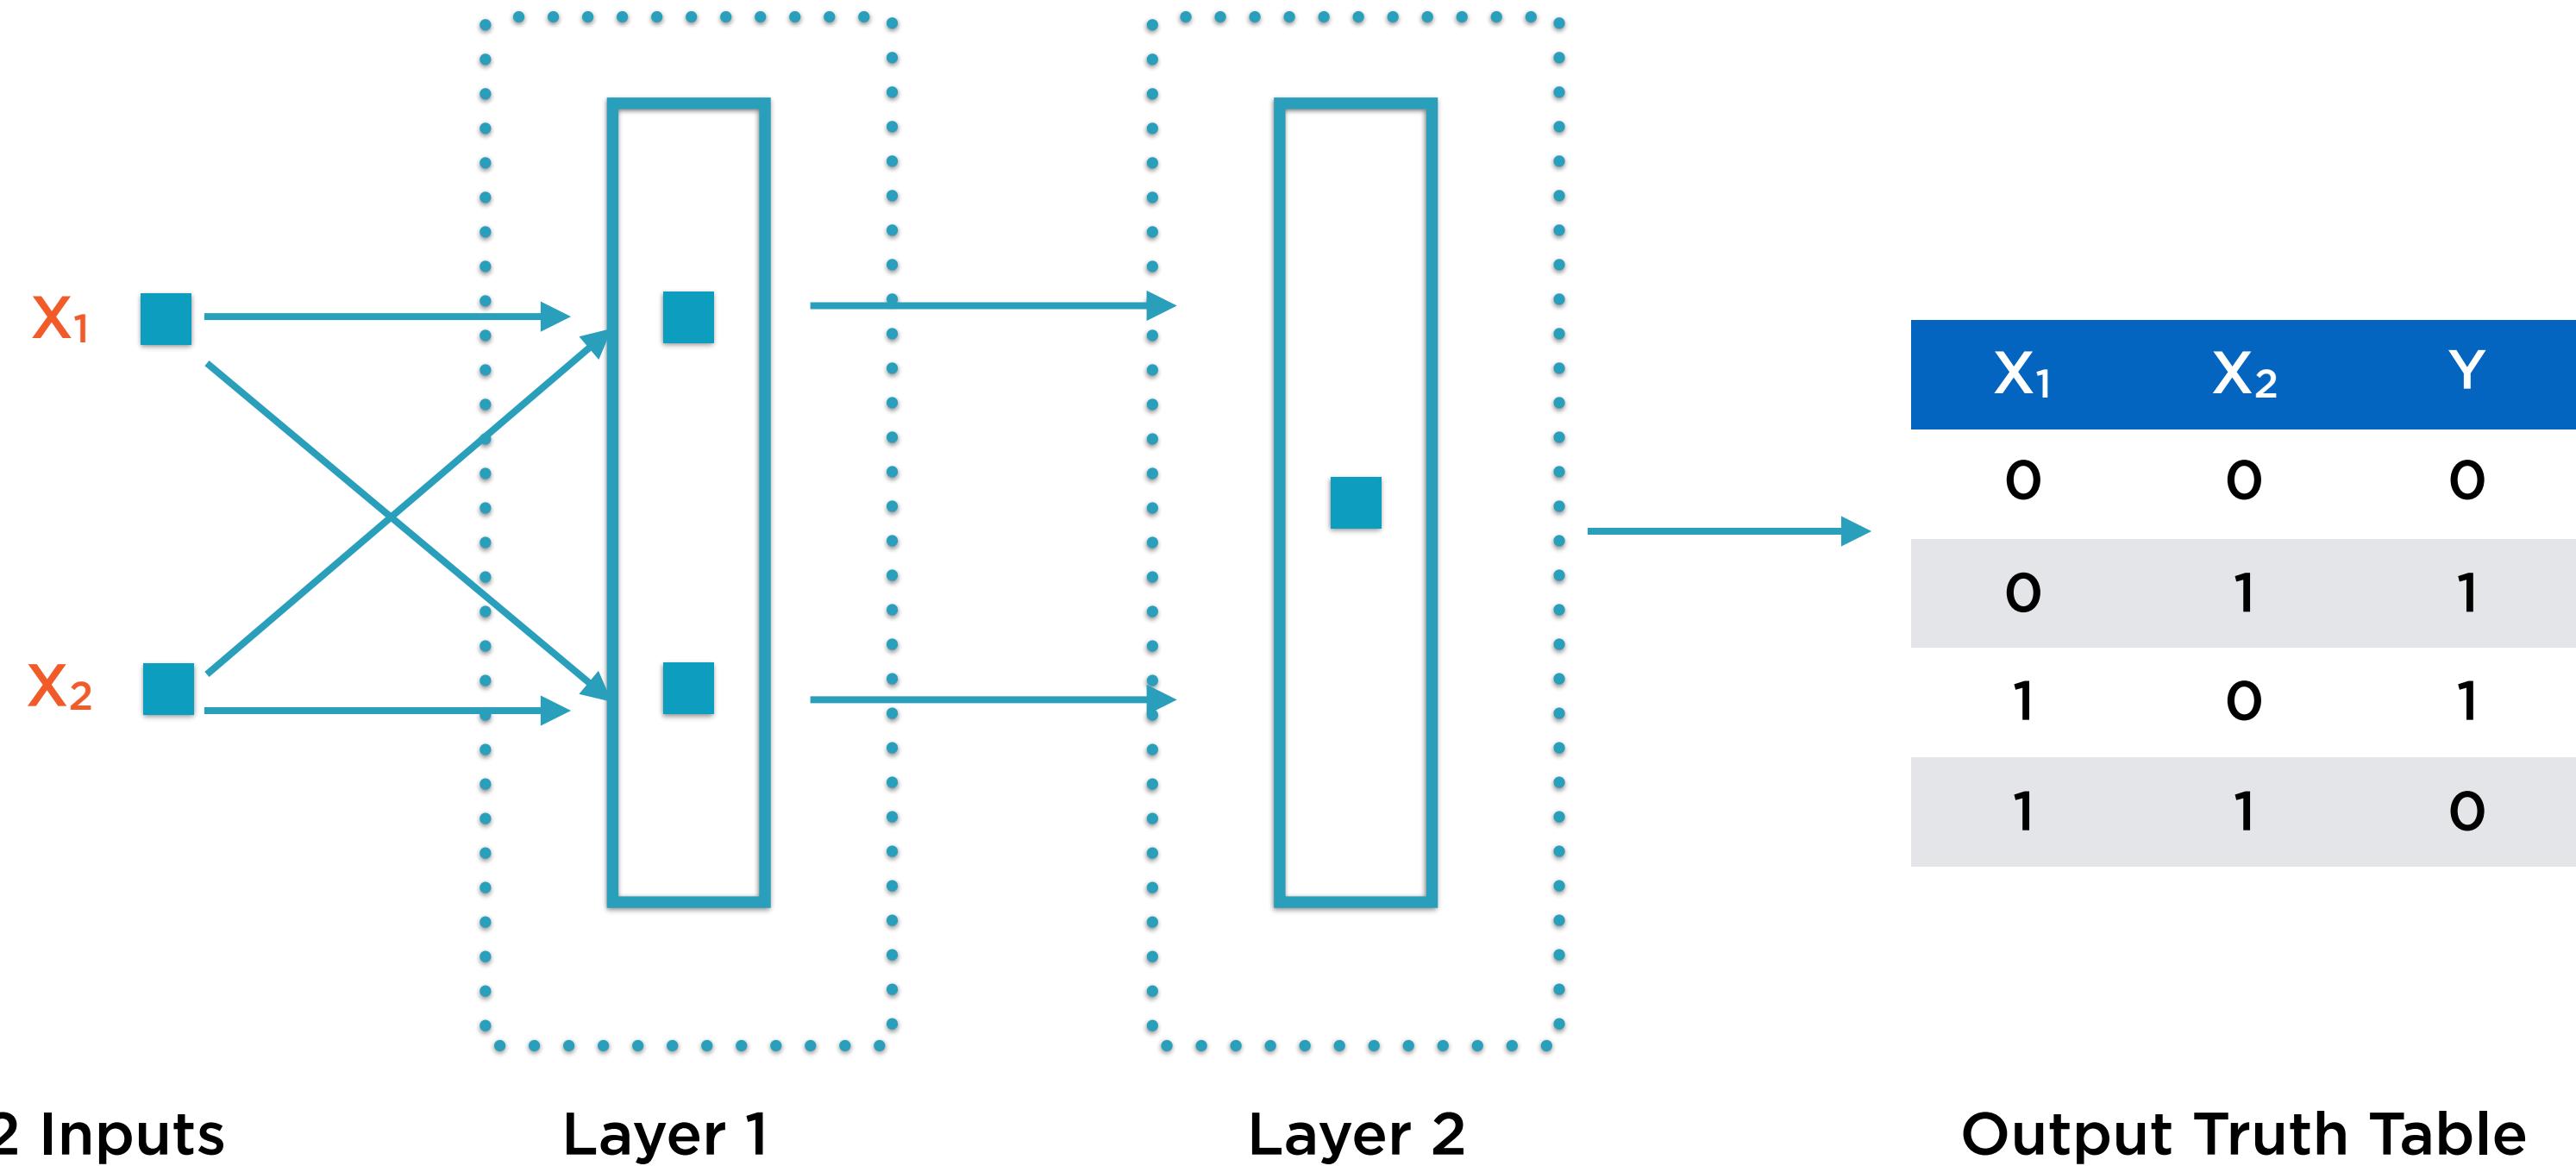


```
def XOR(x1, x2):  
    if (x1 == x2):  
        return 0  
    return 1
```

“Learning” XOR

The XOR function can be reverse-engineered using 3 neurons arranged in 2 layers

XOR: 3 Neurons, 2 Layers

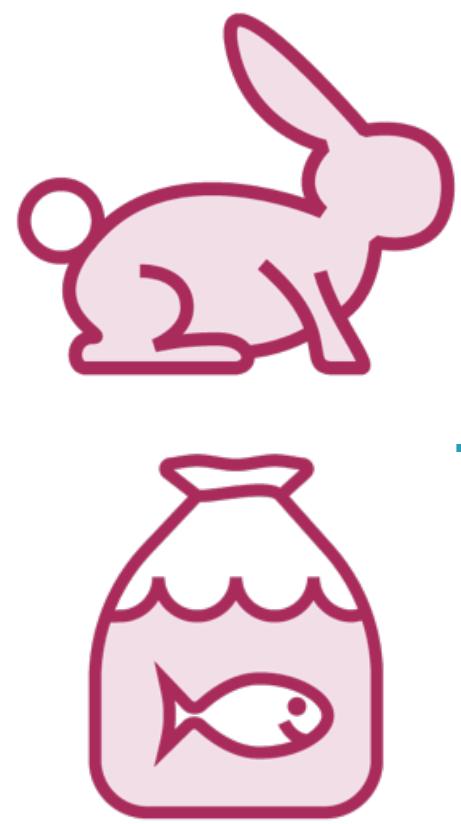


```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

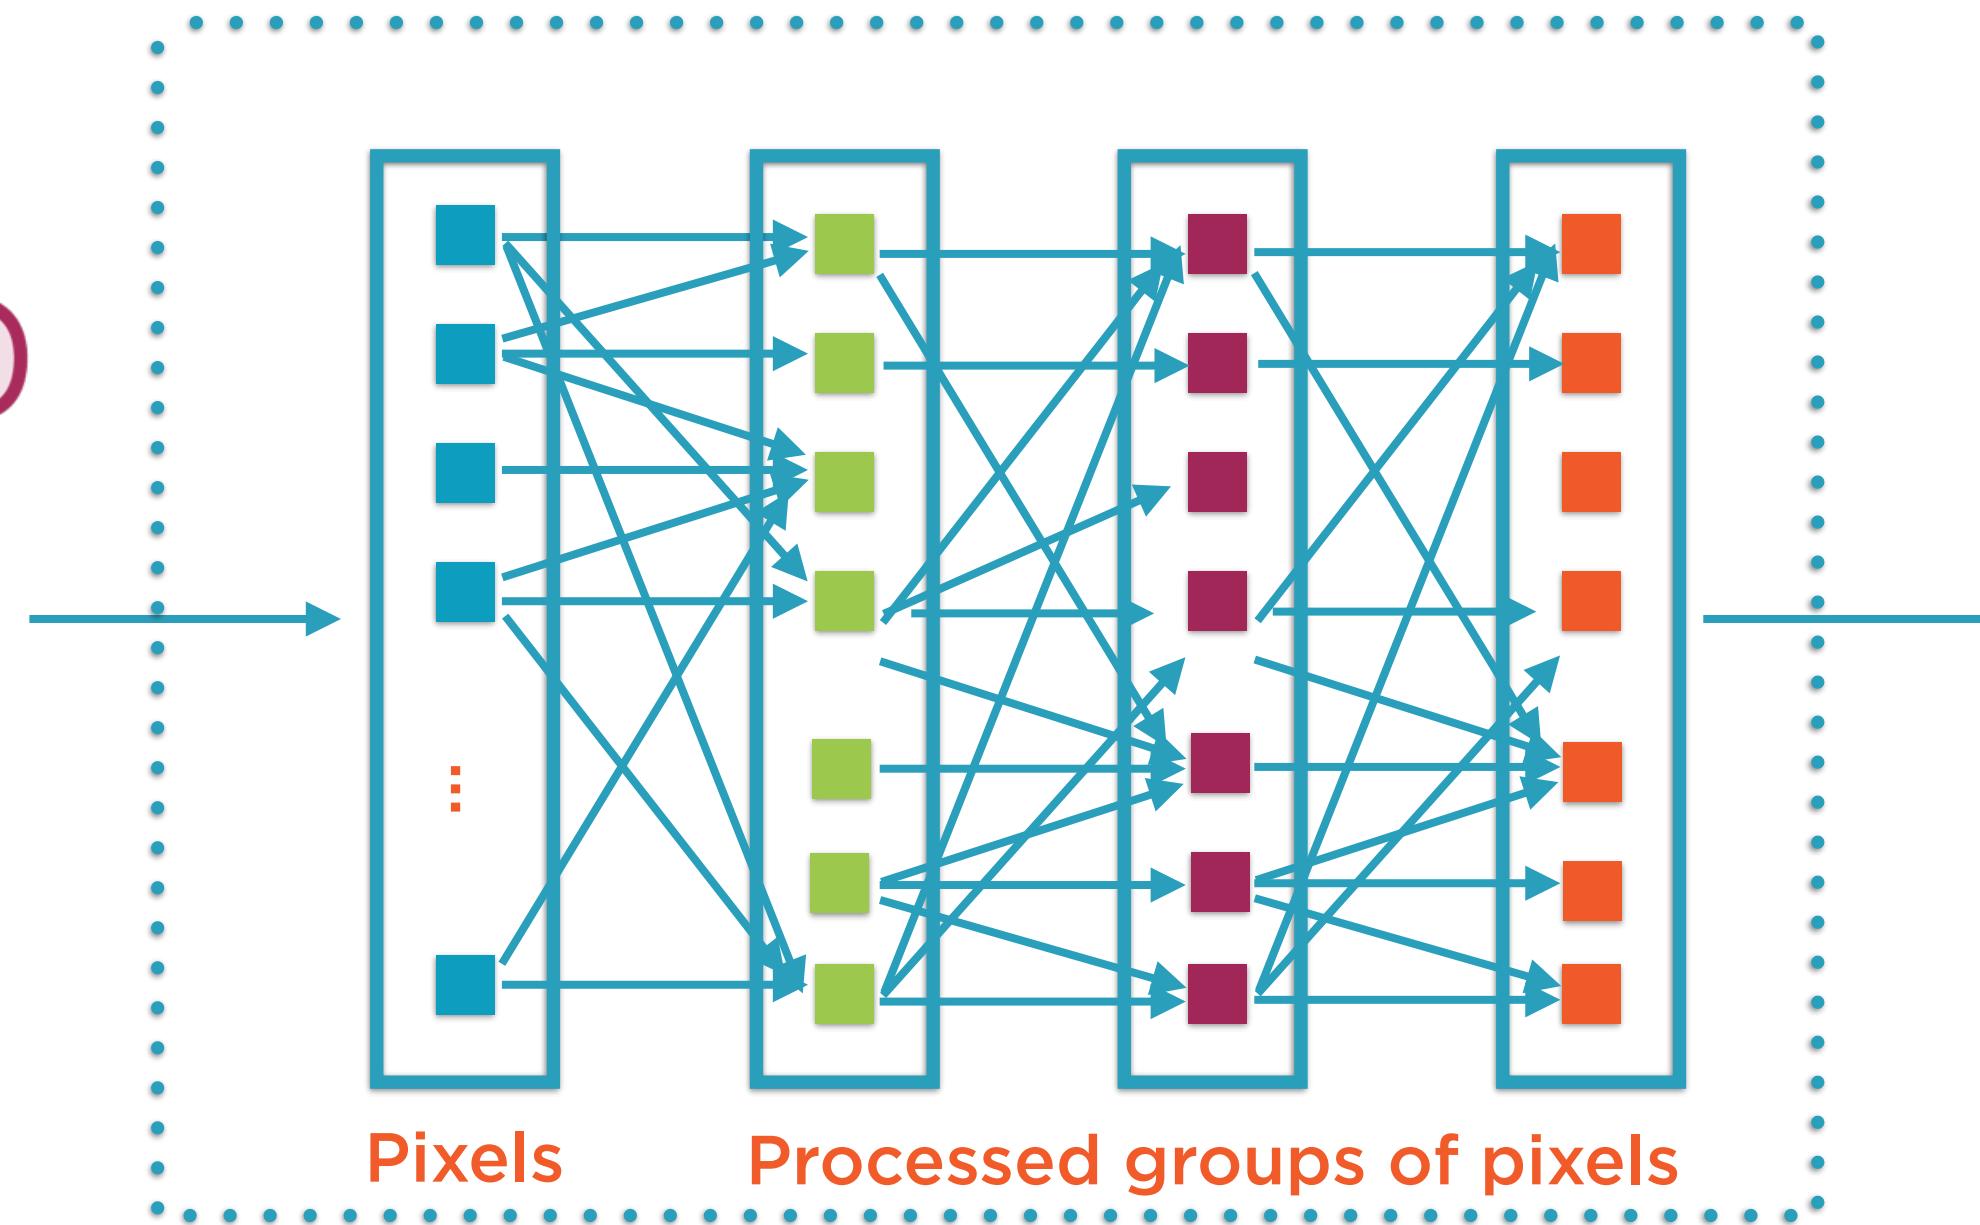
“Learning” Arbitrarily Complex Functions

Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything

Arbitrarily Complex Function



Corpus of
Images

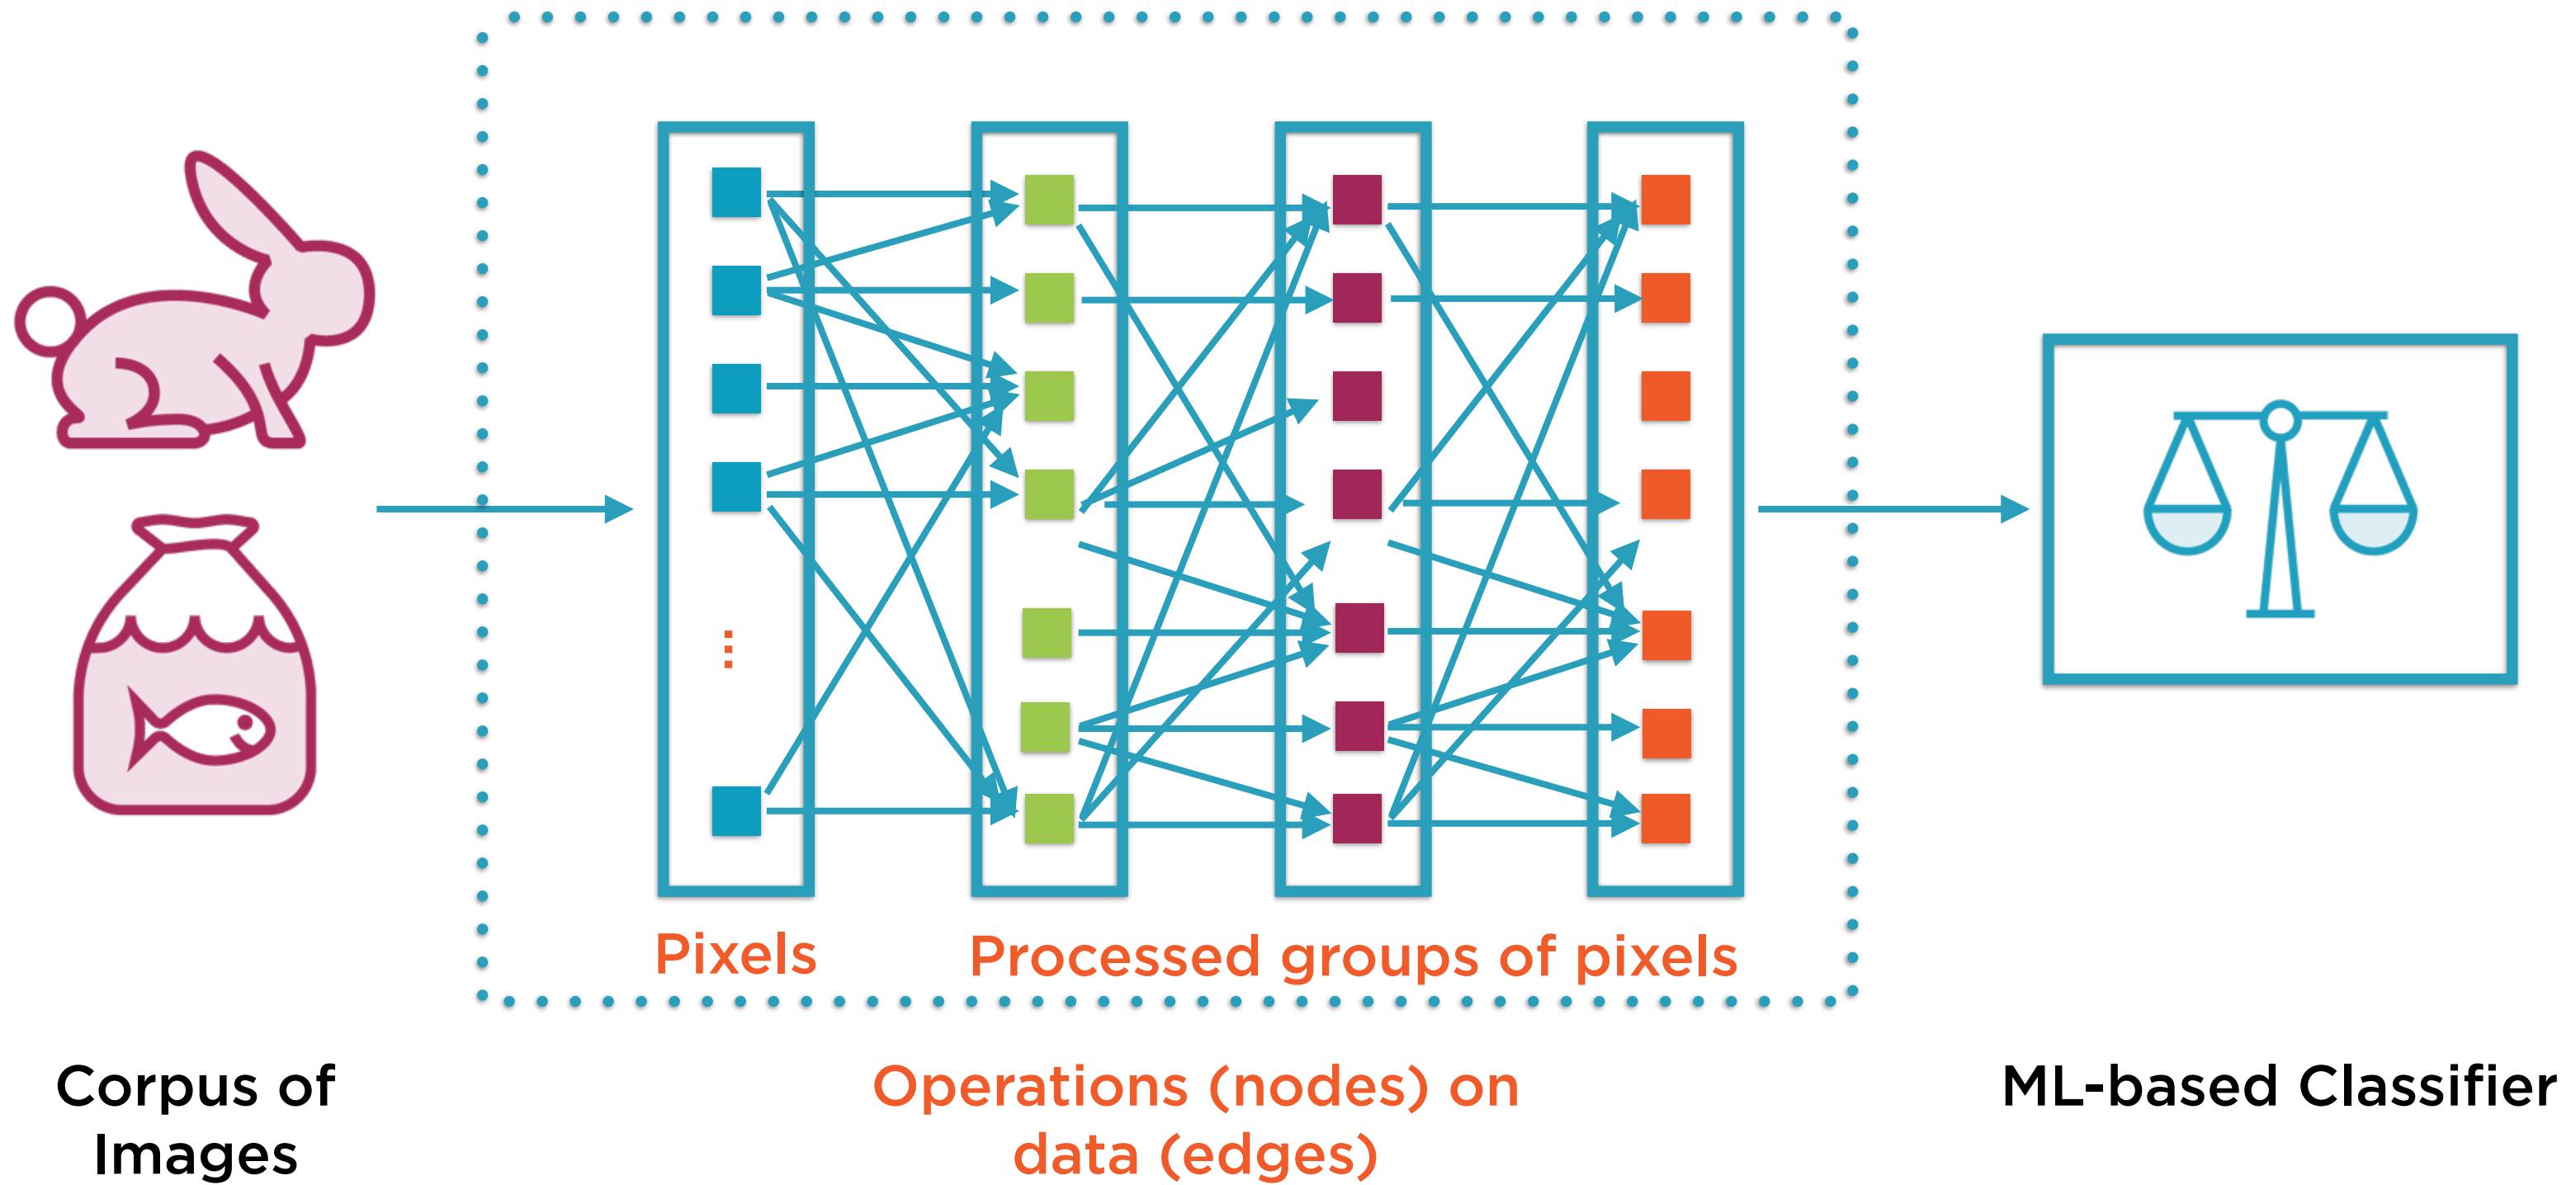


Operations (nodes) on
data (edges)

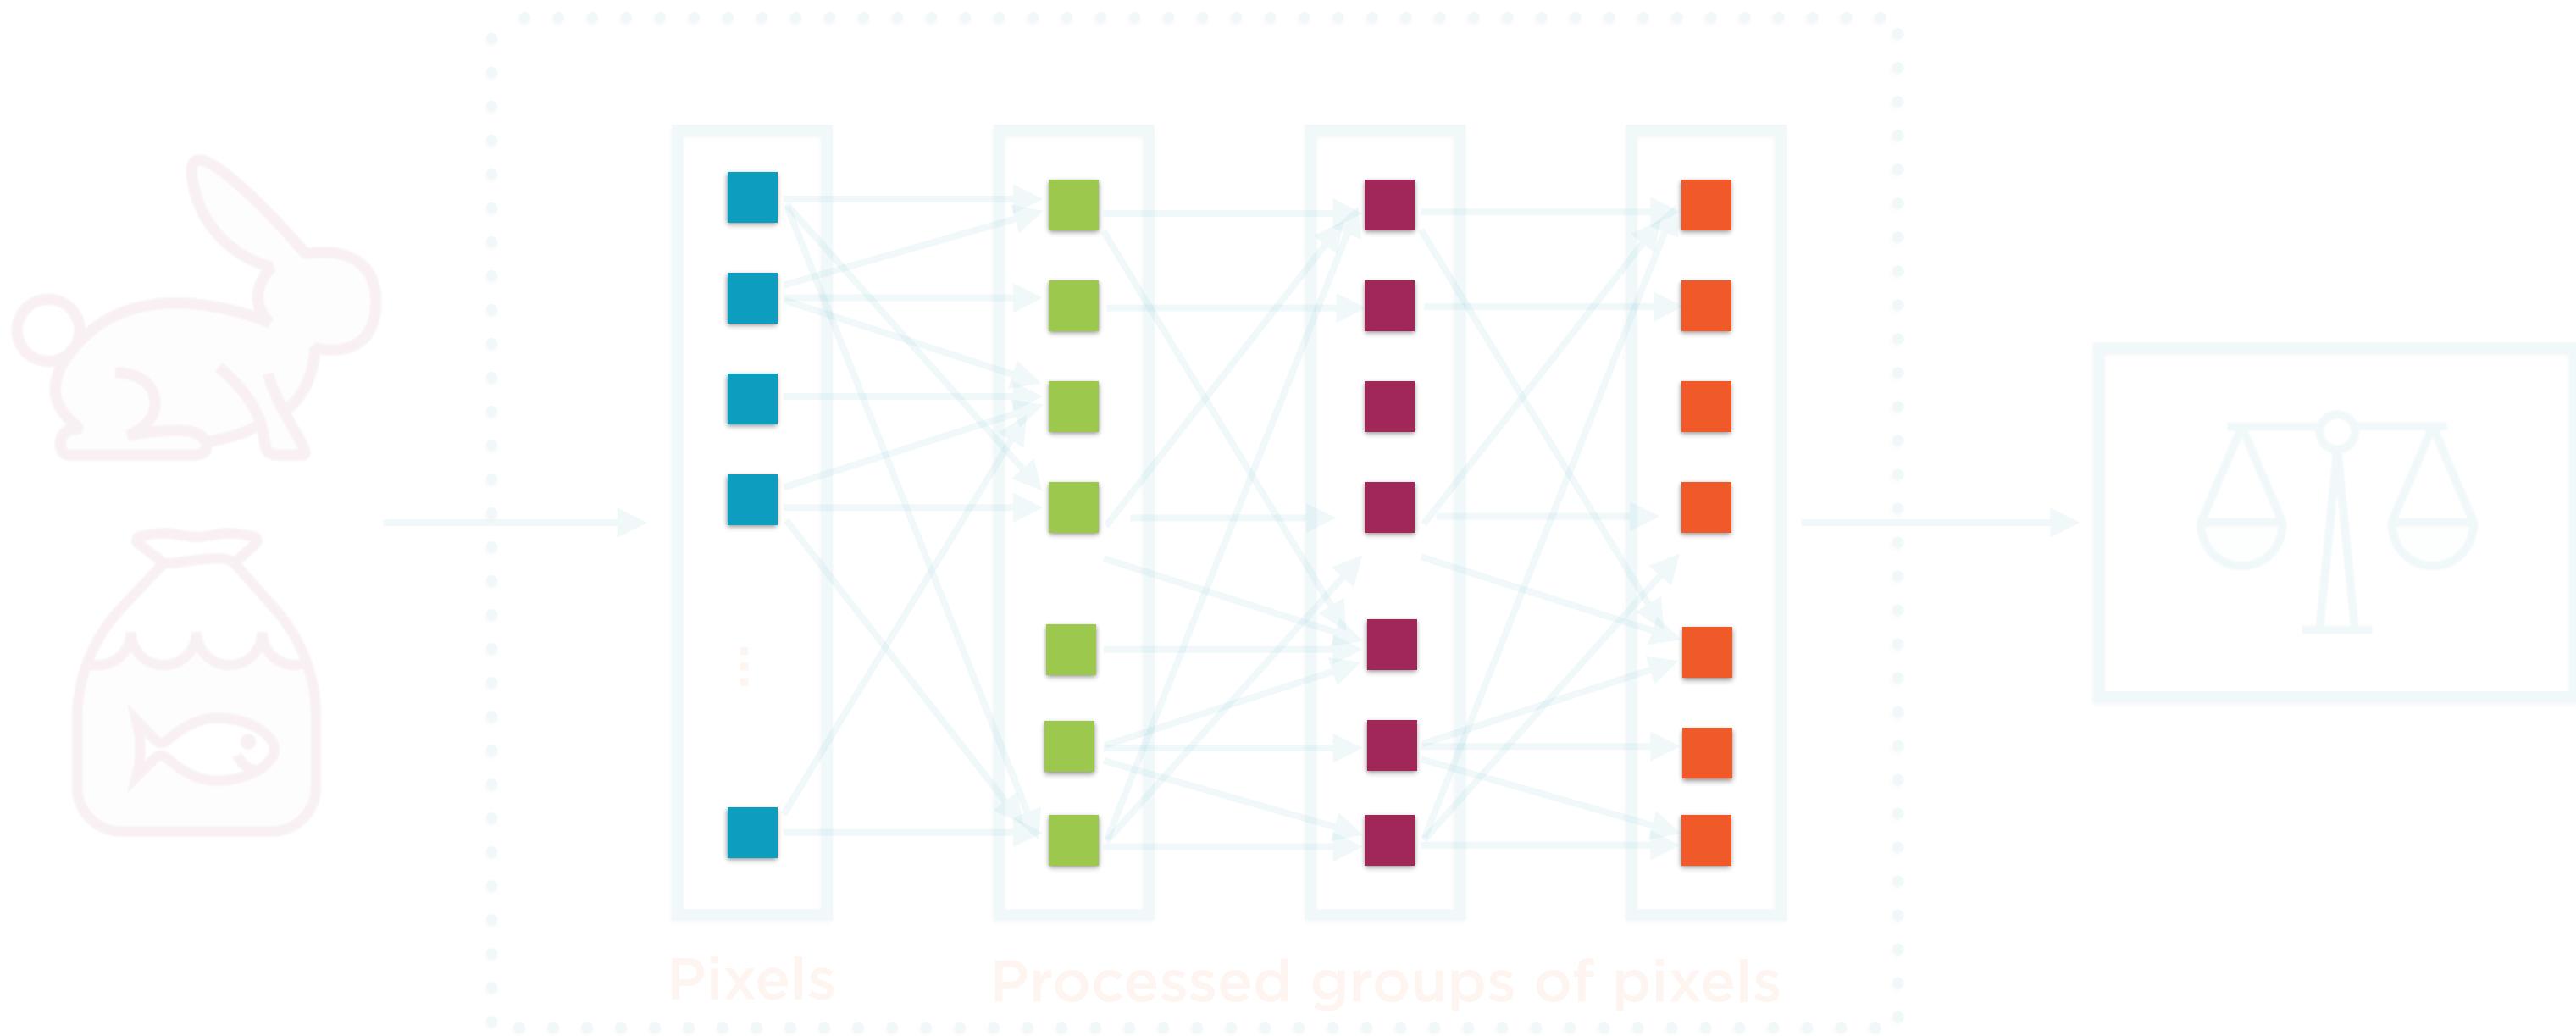


ML-based Classifier

The Computational Graph

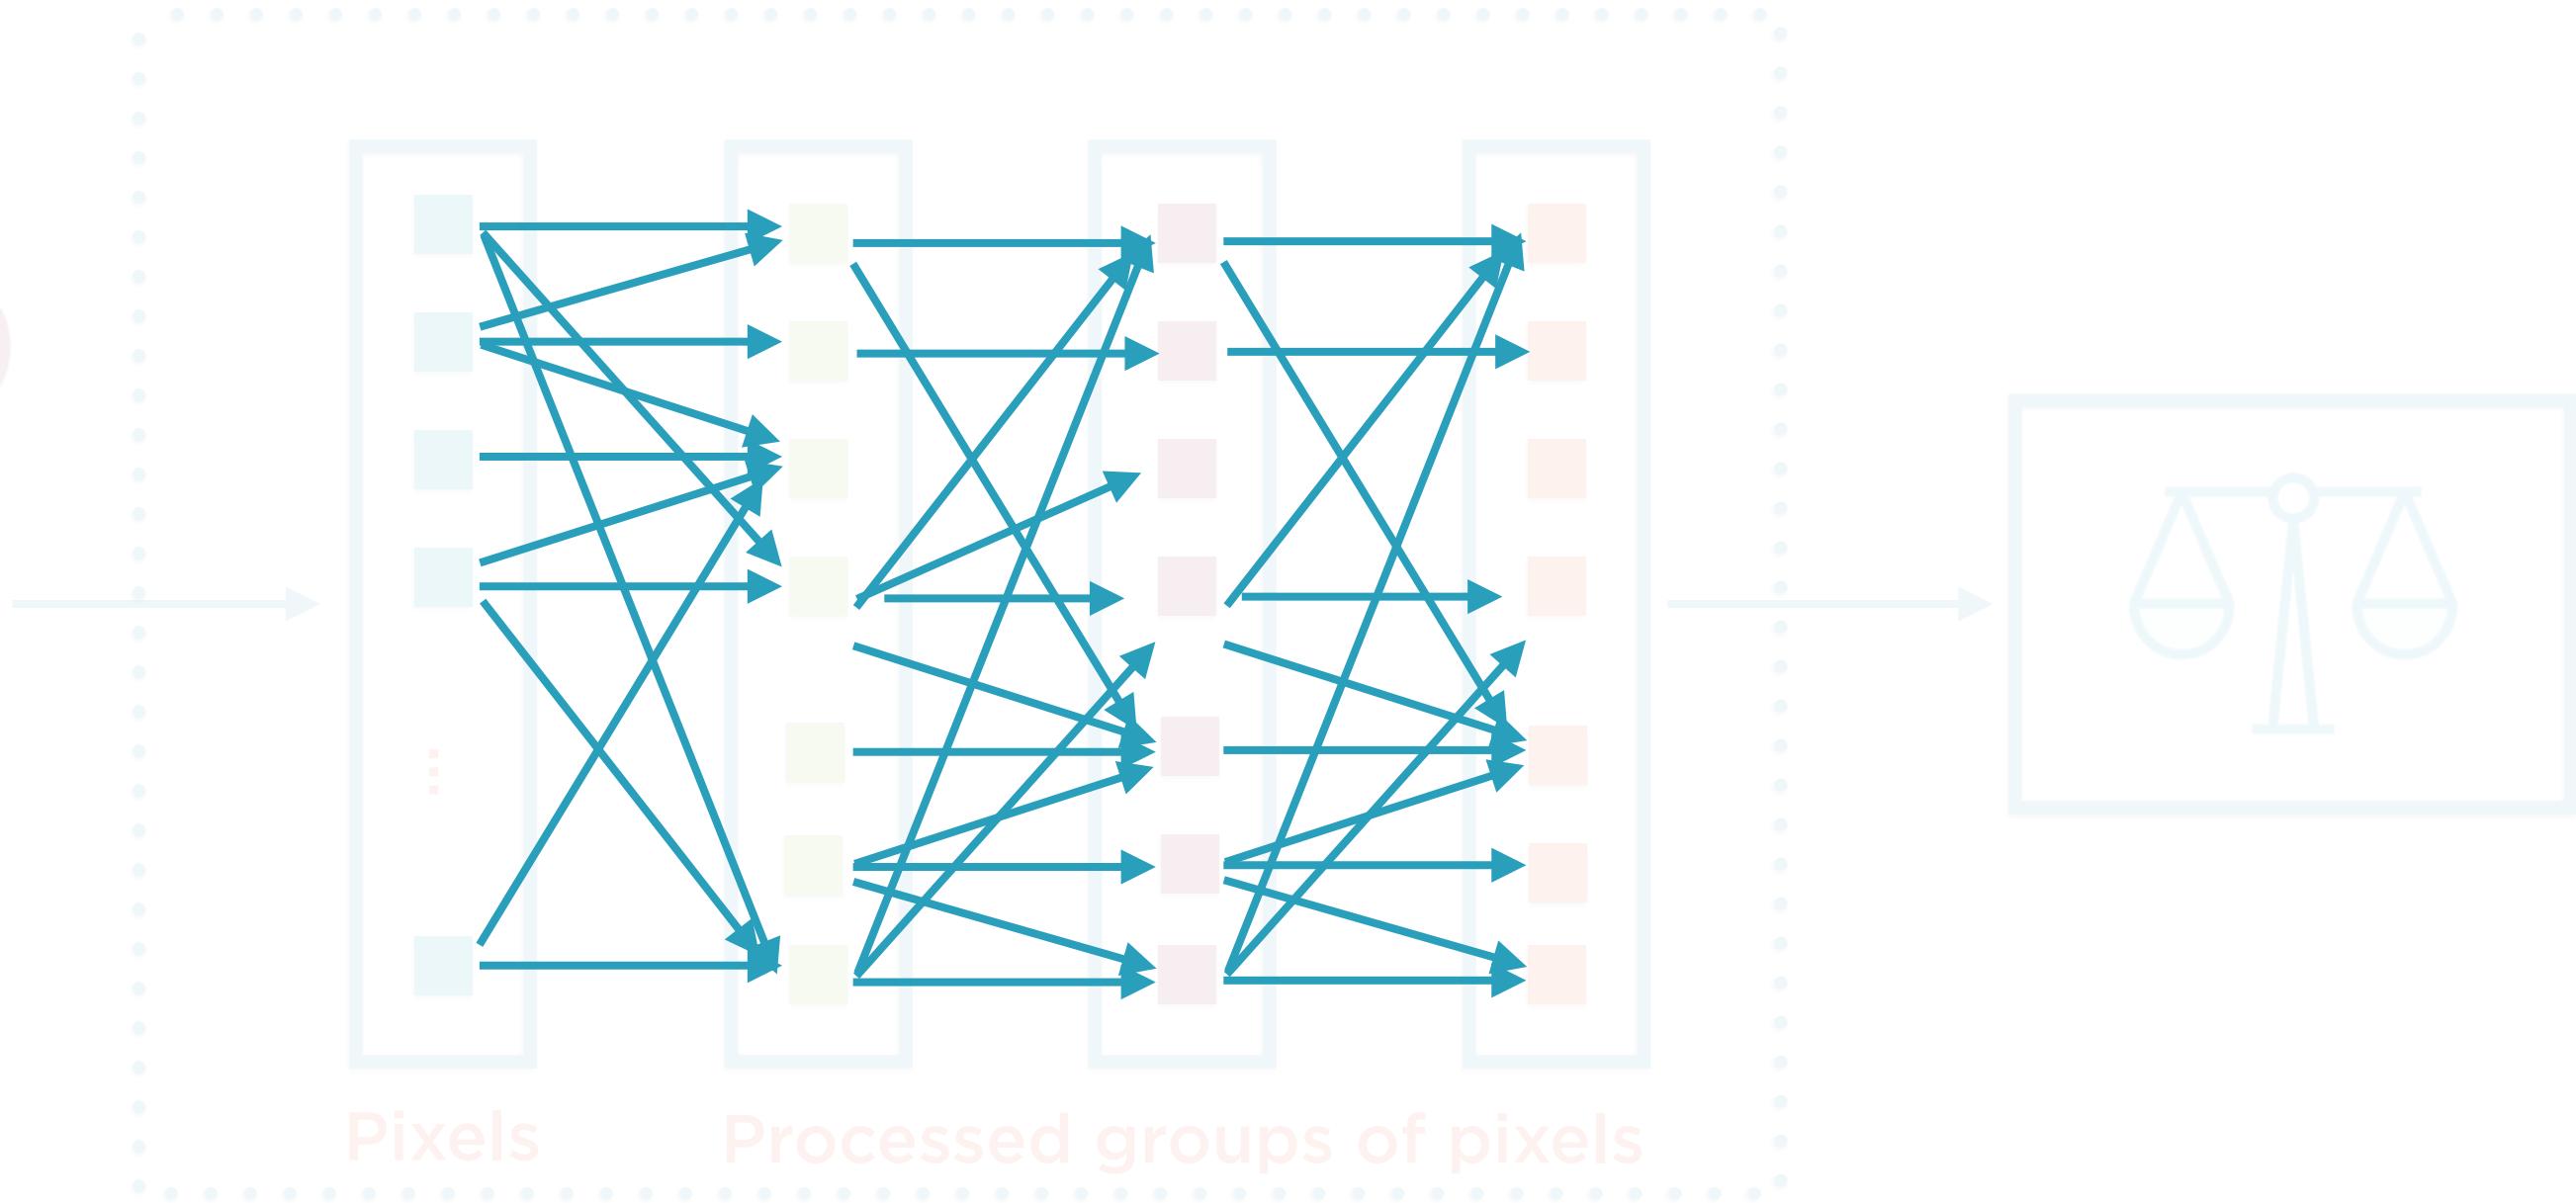
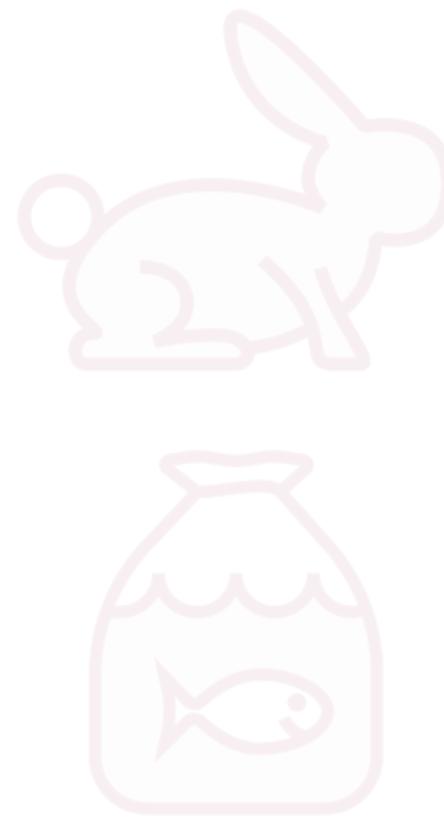


The Computational Graph



The nodes in the computation graph are neurons (simple building blocks)

The Computational Graph



Corpus of
Images

**The edges in the computation graph
are data items called tensors**

ML-based Classifier

Neurons

**The nodes in the computation graph
are simple entities called neurons**

**Each neuron performs very simple
operations on data**

**The neurons are connected in very
complex, sophisticated ways**

Neural Networks

The complex interconnections between simple neurons

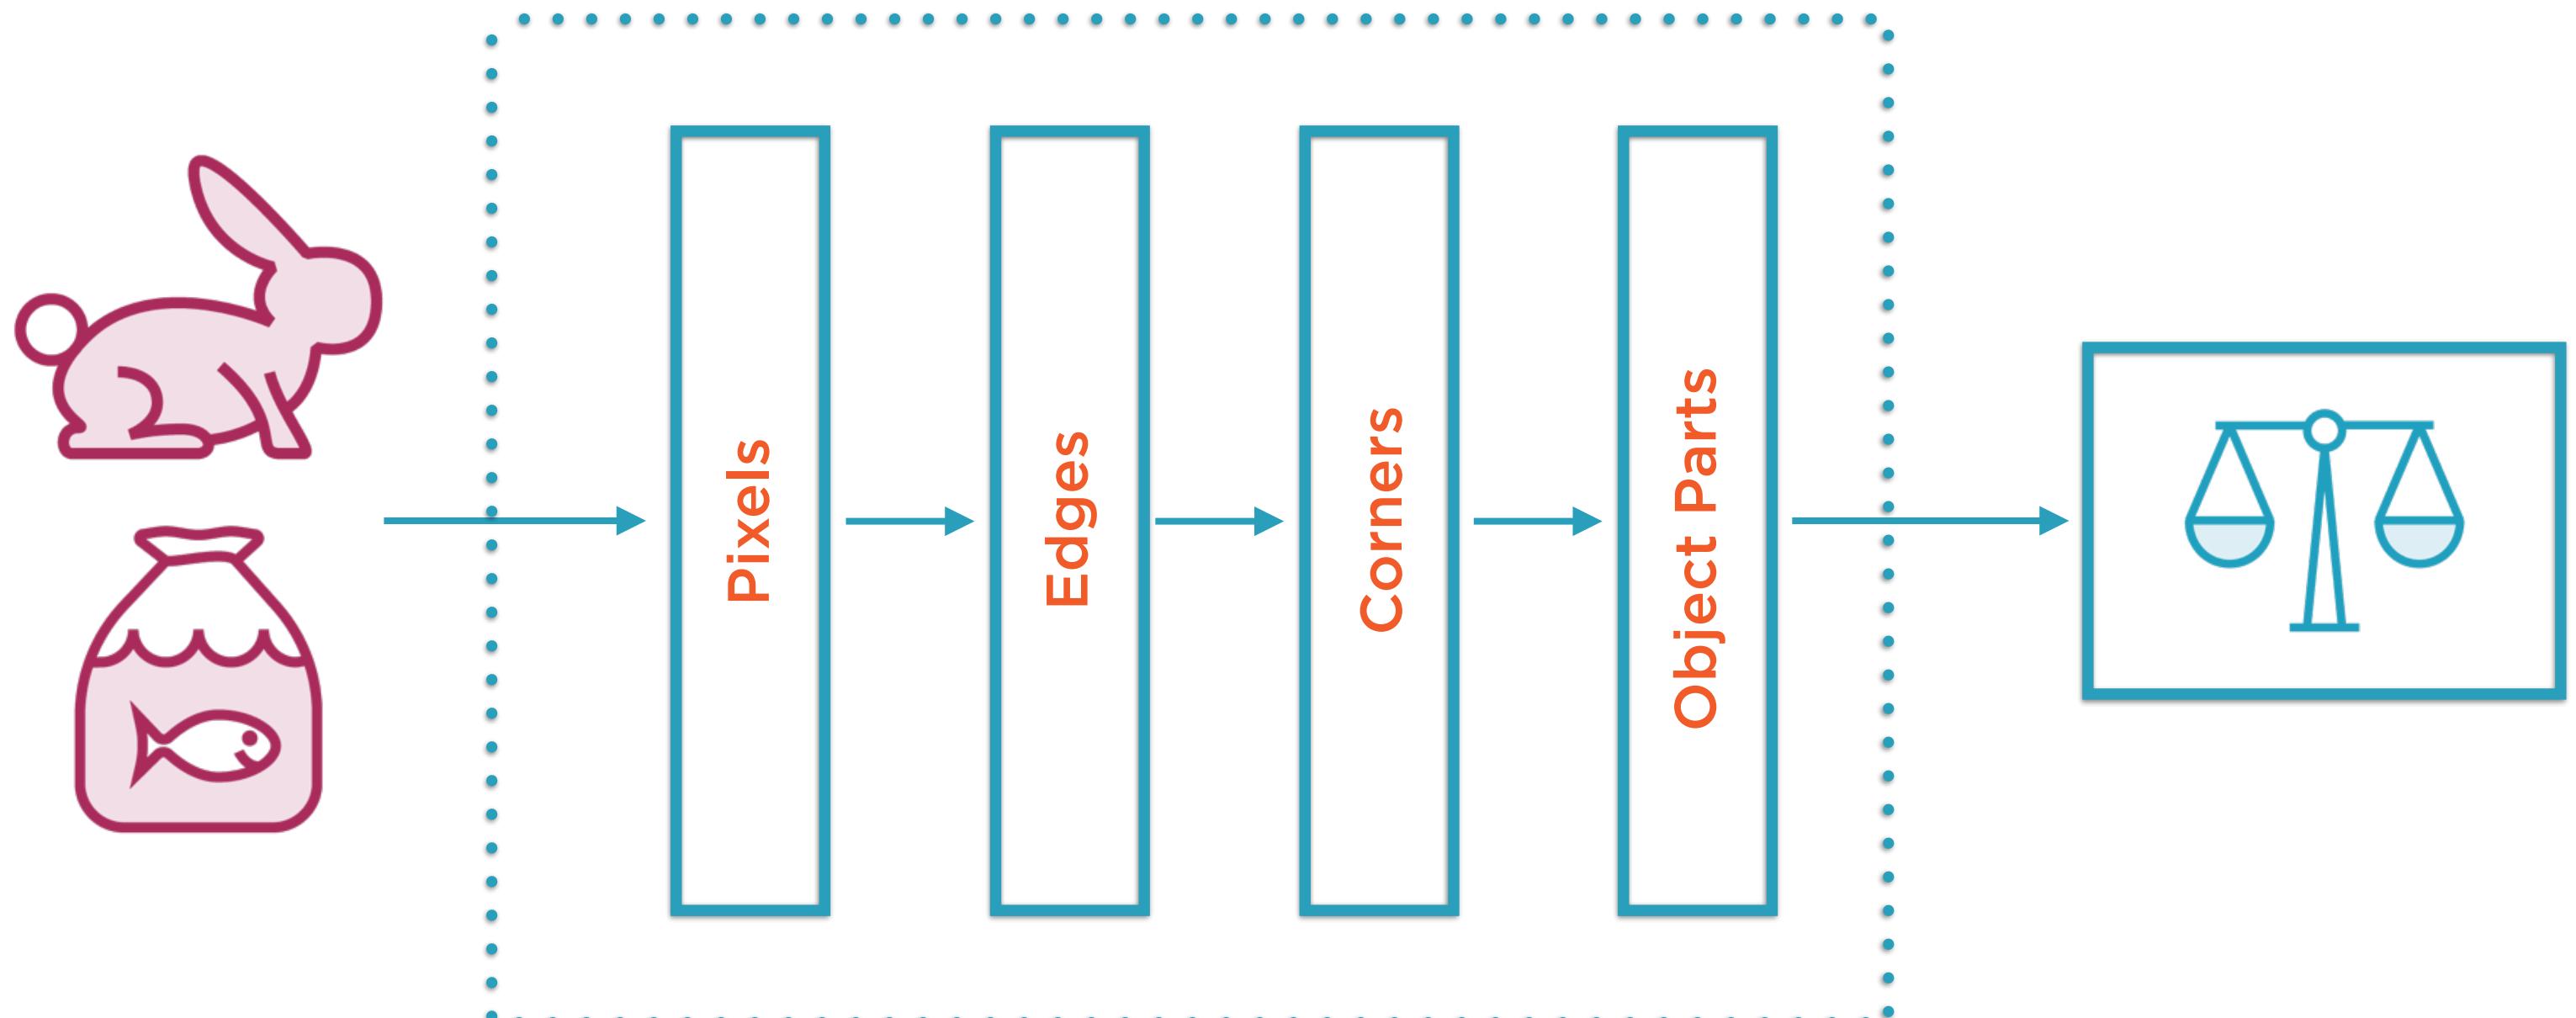
Different network configurations => different types of neural networks

- Convolutional
- Recurrent

Neural Networks

Groups of neurons that perform similar functions are aggregated into layers

Layers in the Computation Graph

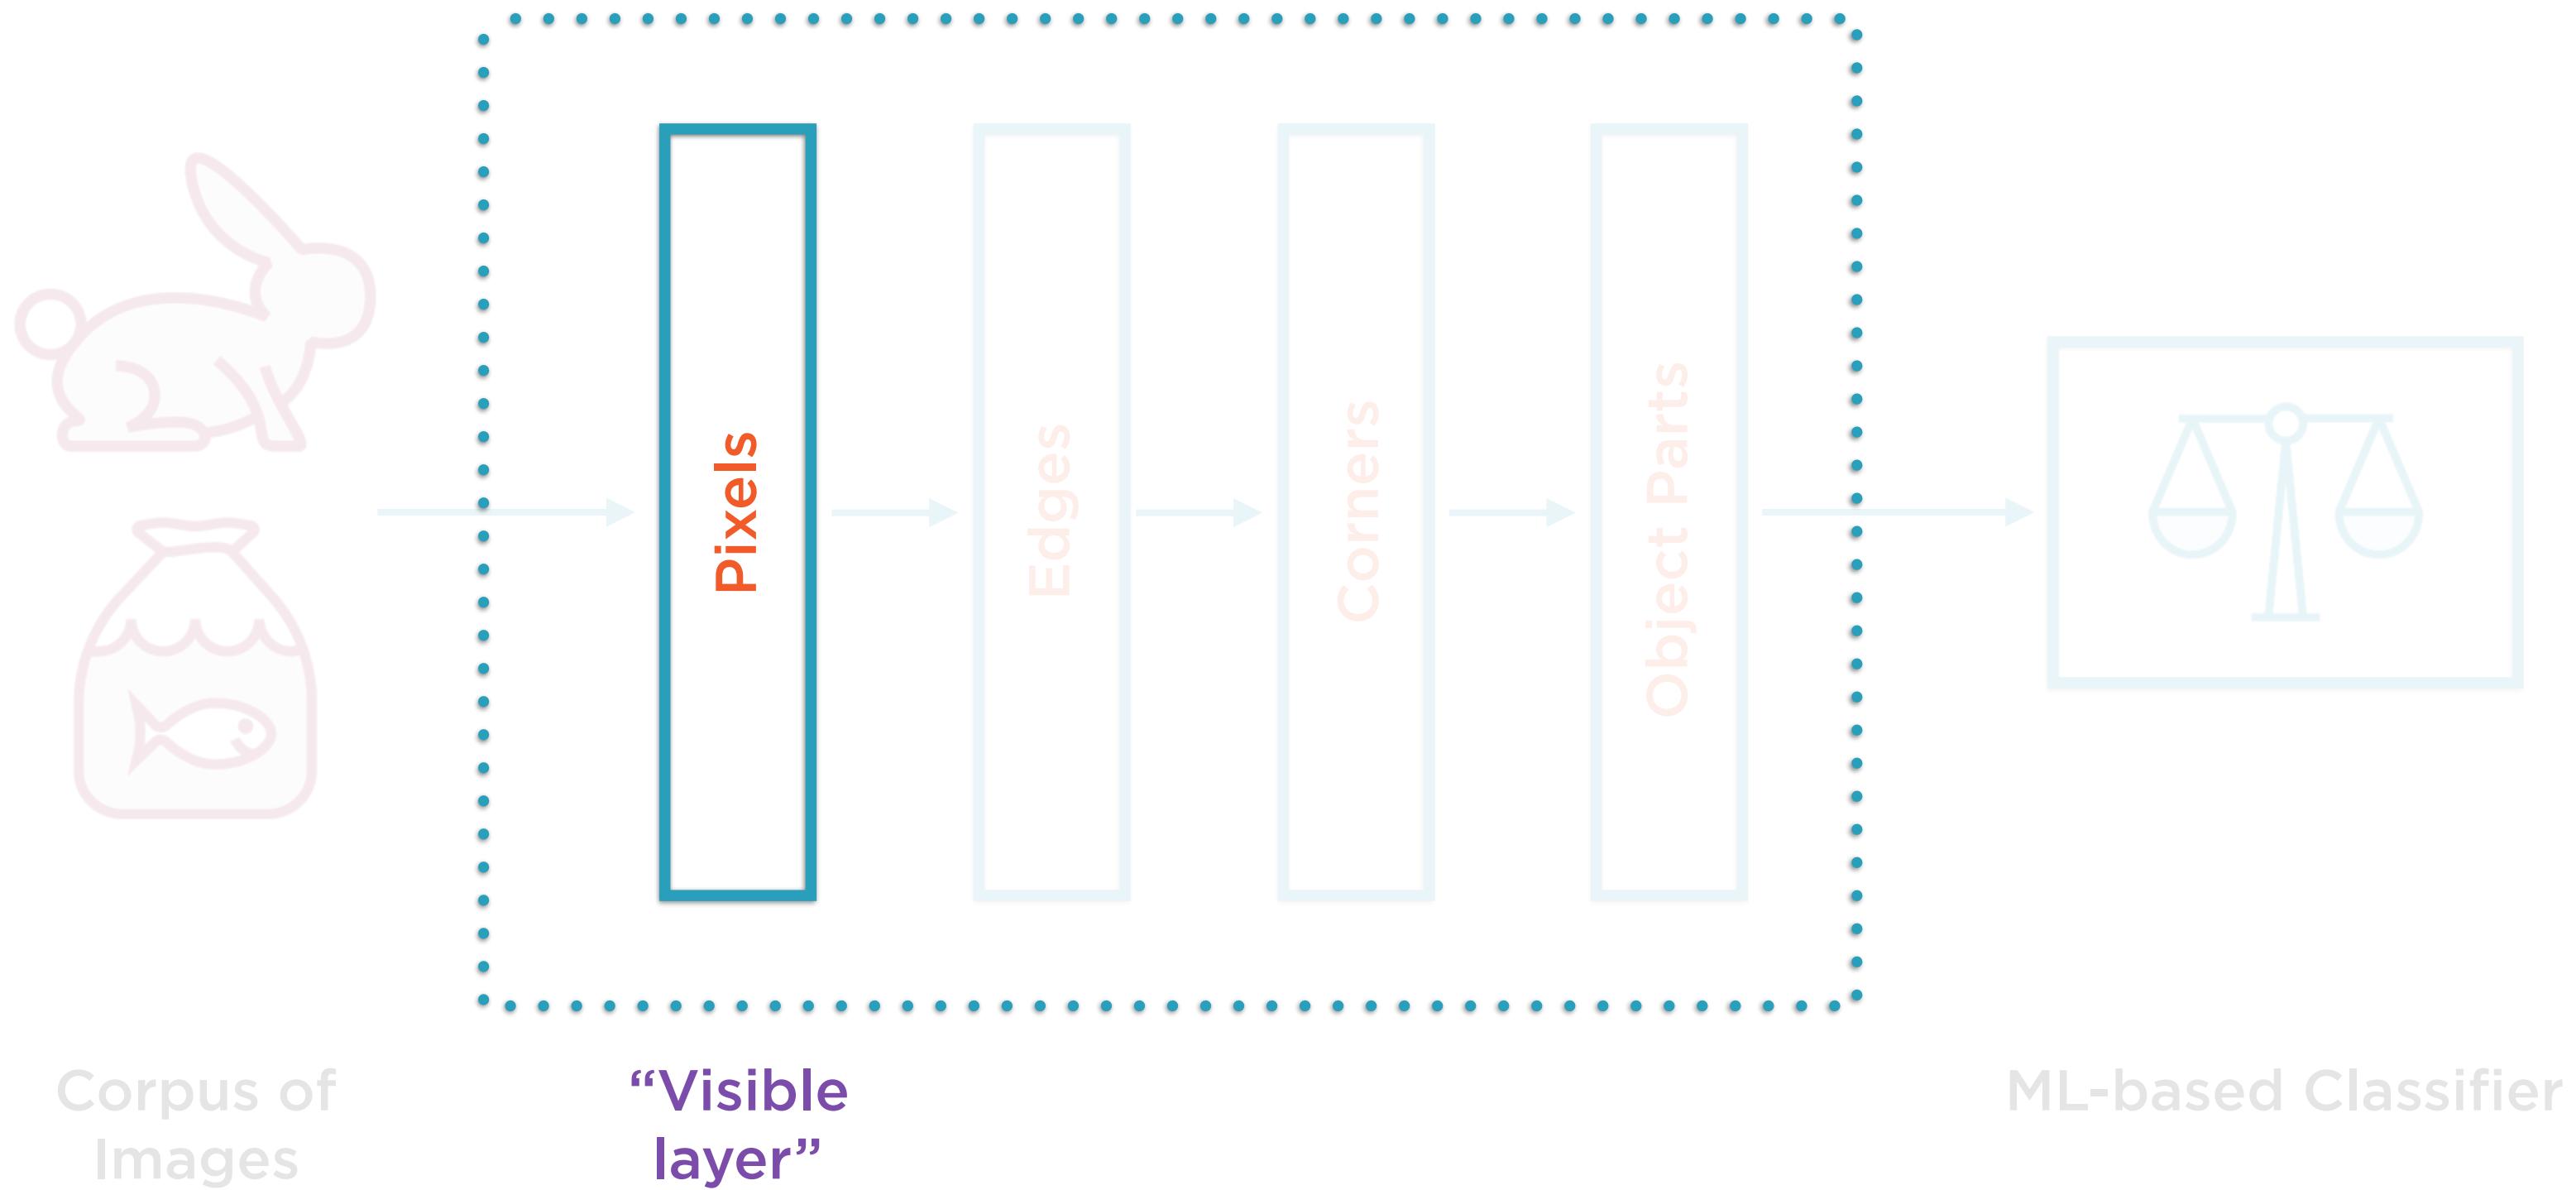


**Corpus of
Images**

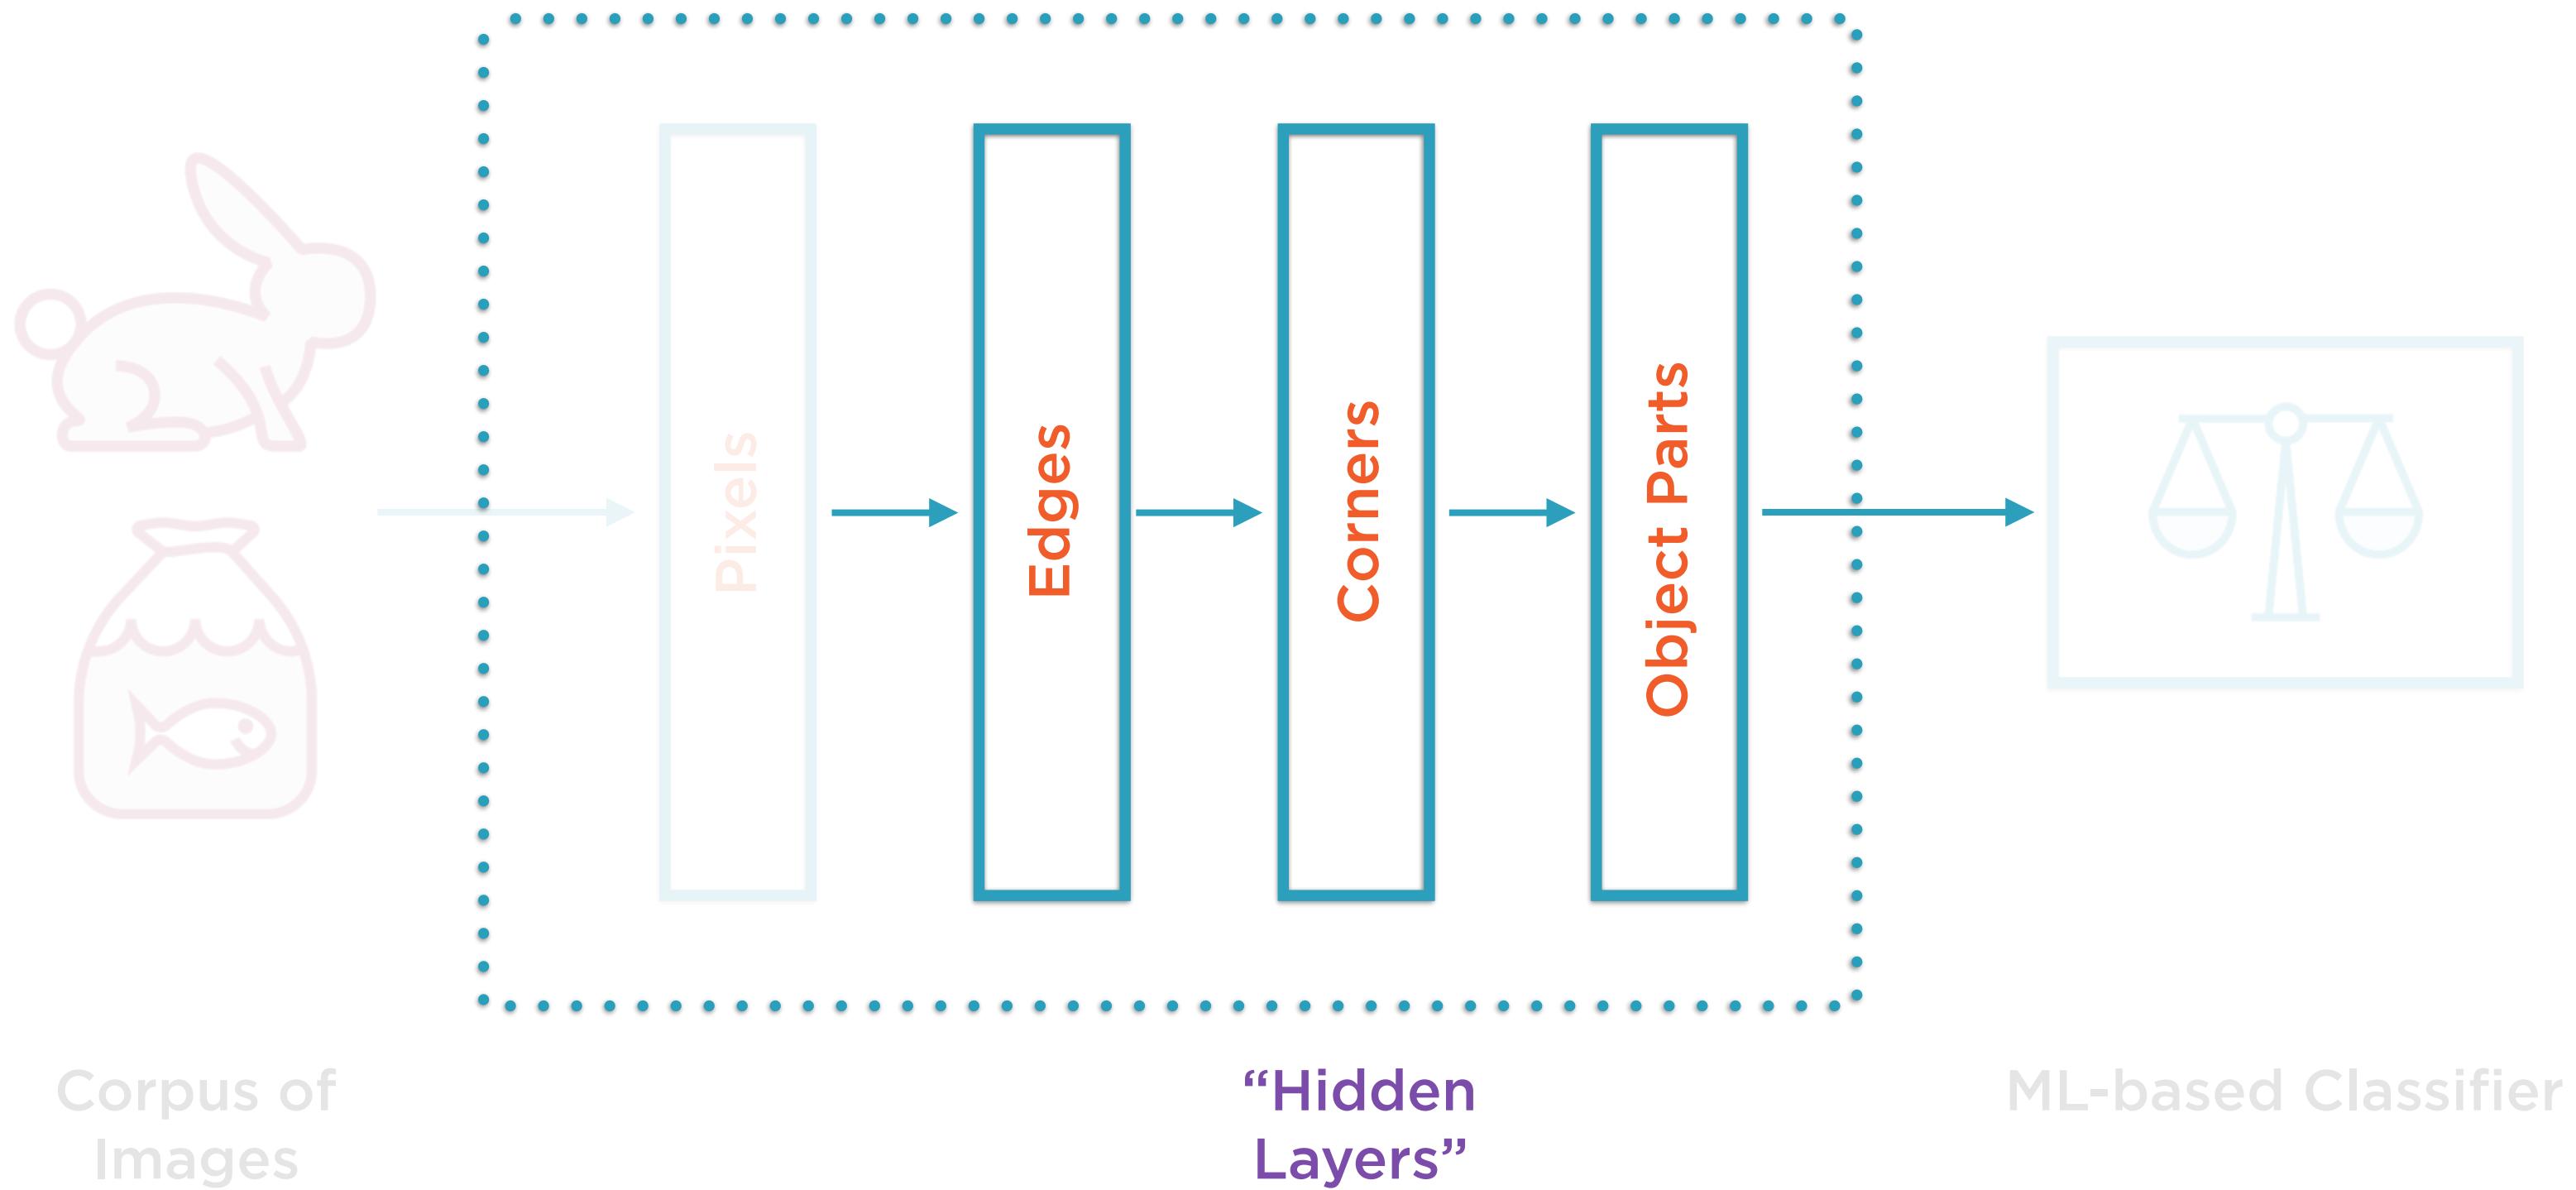
**Groups of neurons that perform similar
functions are aggregated into layers**

ML-based Classifier

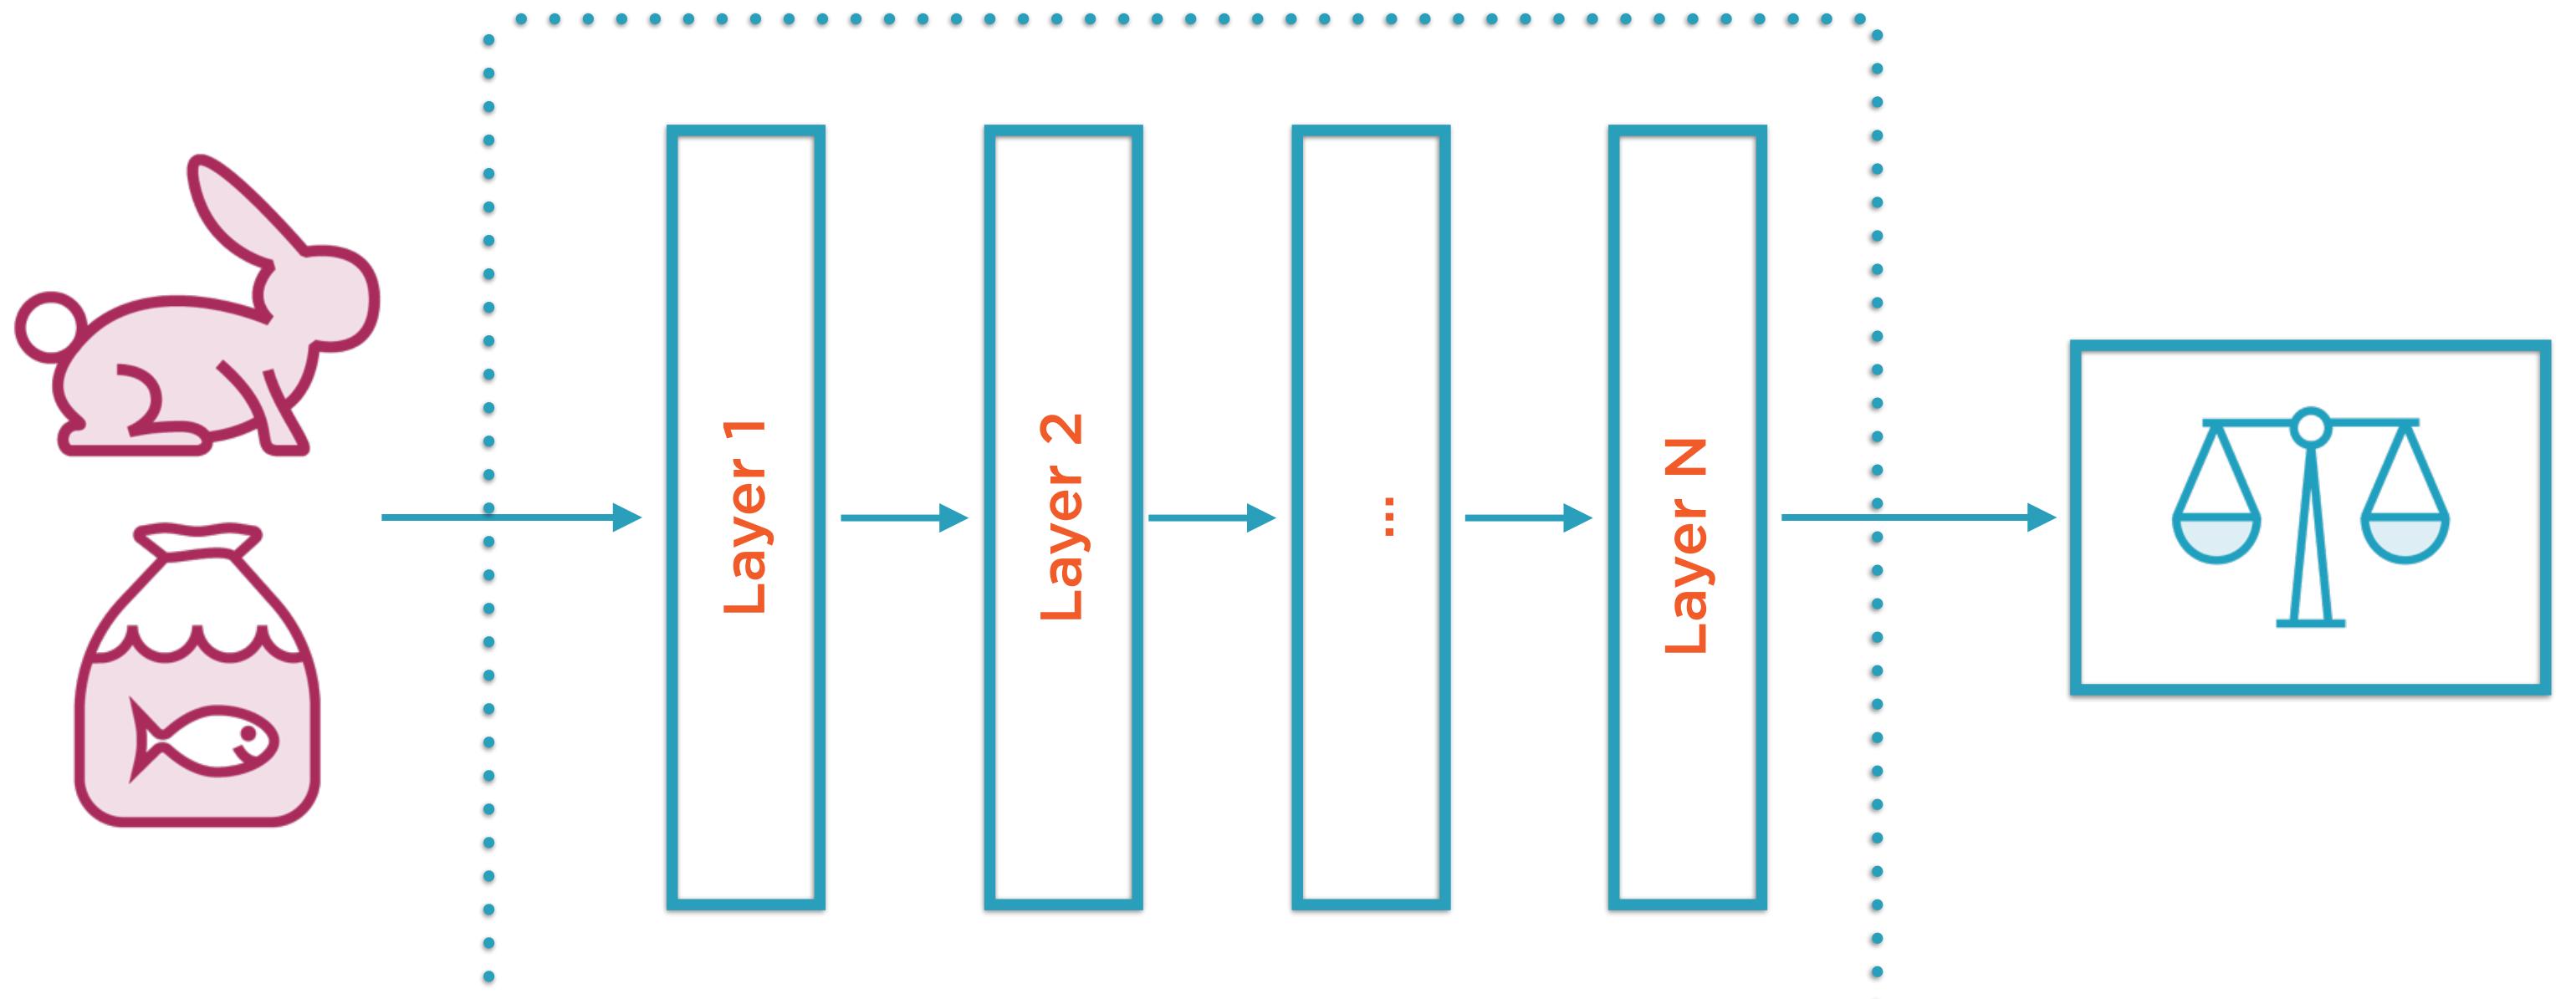
Layers in the Computation Graph



Layers in the Computation Graph



Layers in the Computation Graph



**Corpus of
Images**

Layers in a neural network

ML-based Classifier

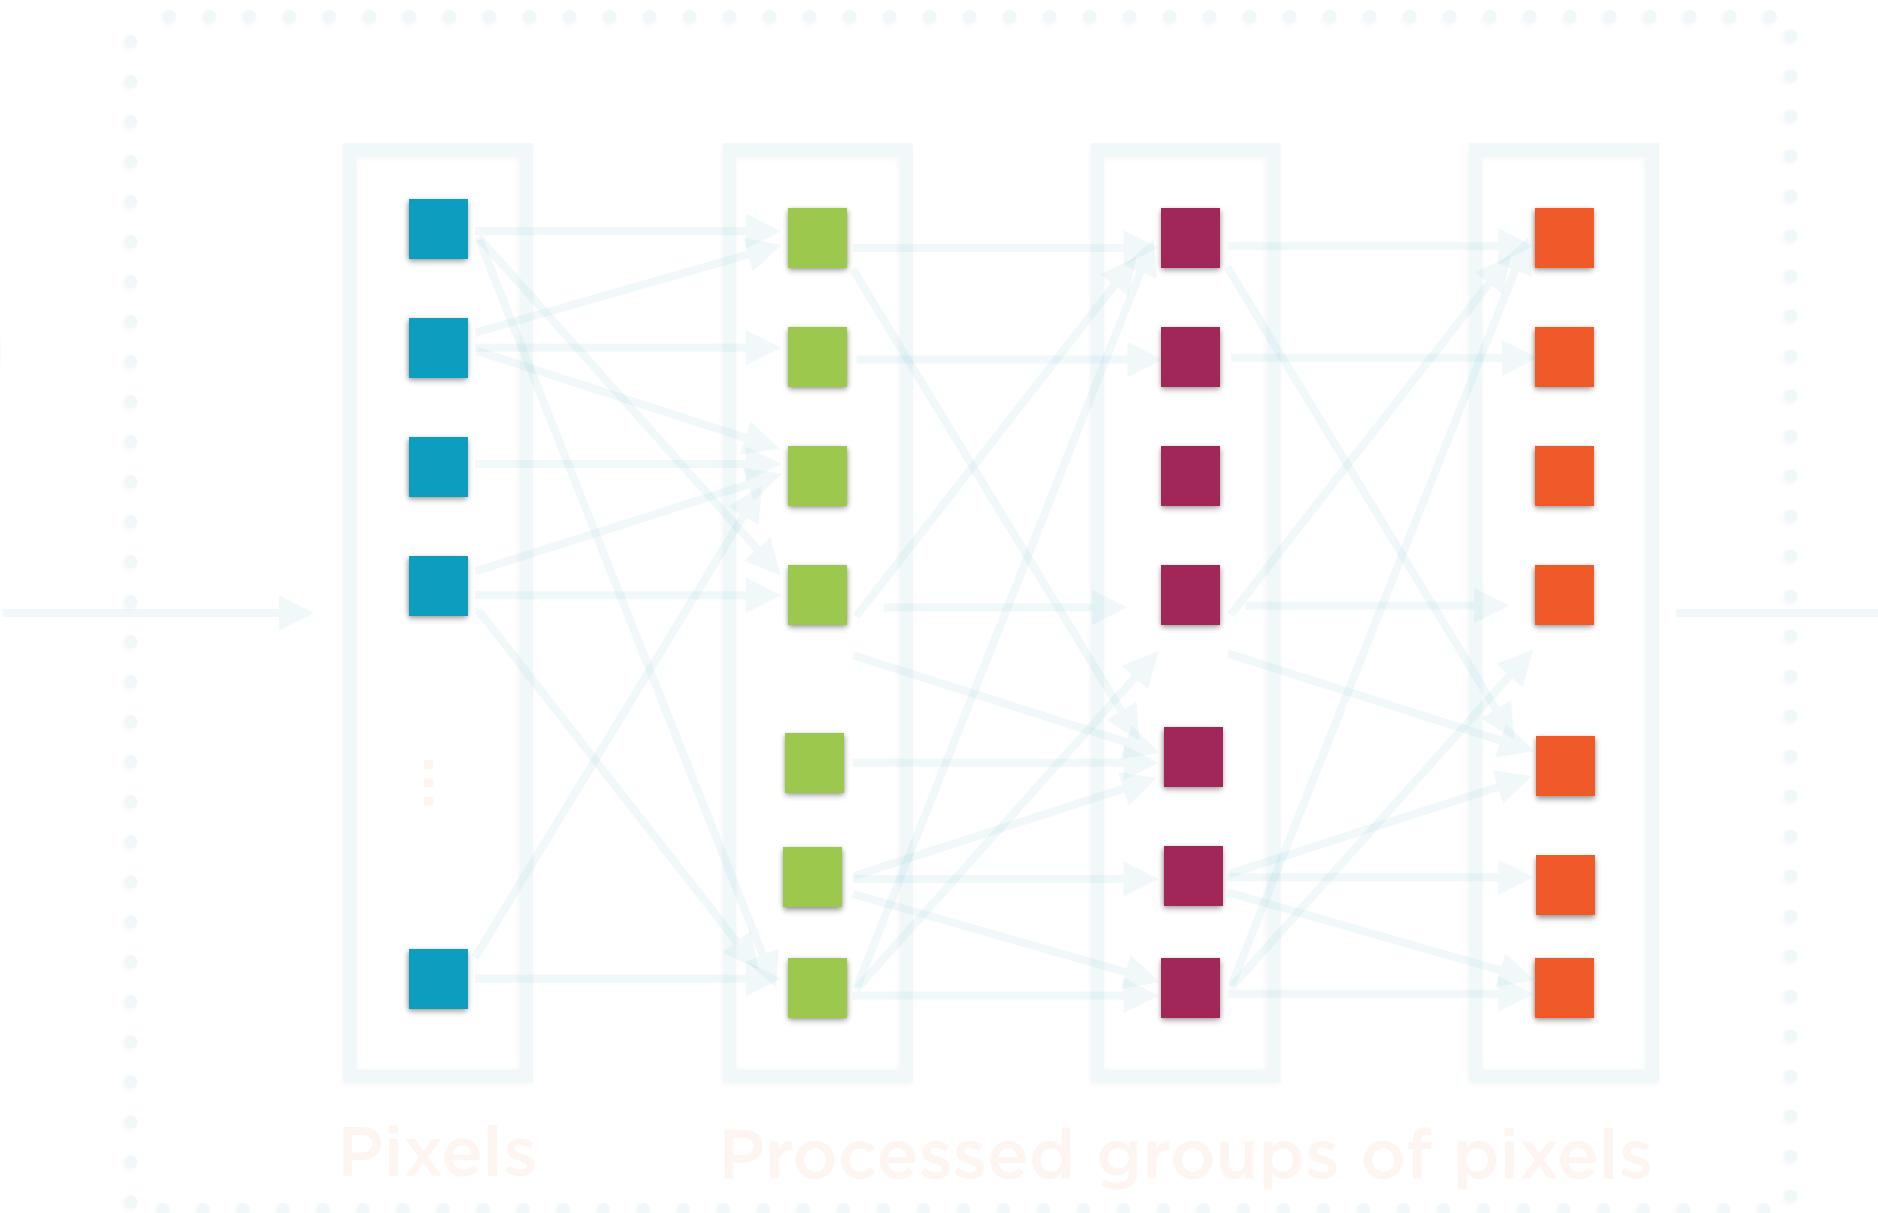
Neurons

Each layer consists of units called neurons

Neurons



Corpus of
Images



Neural Network

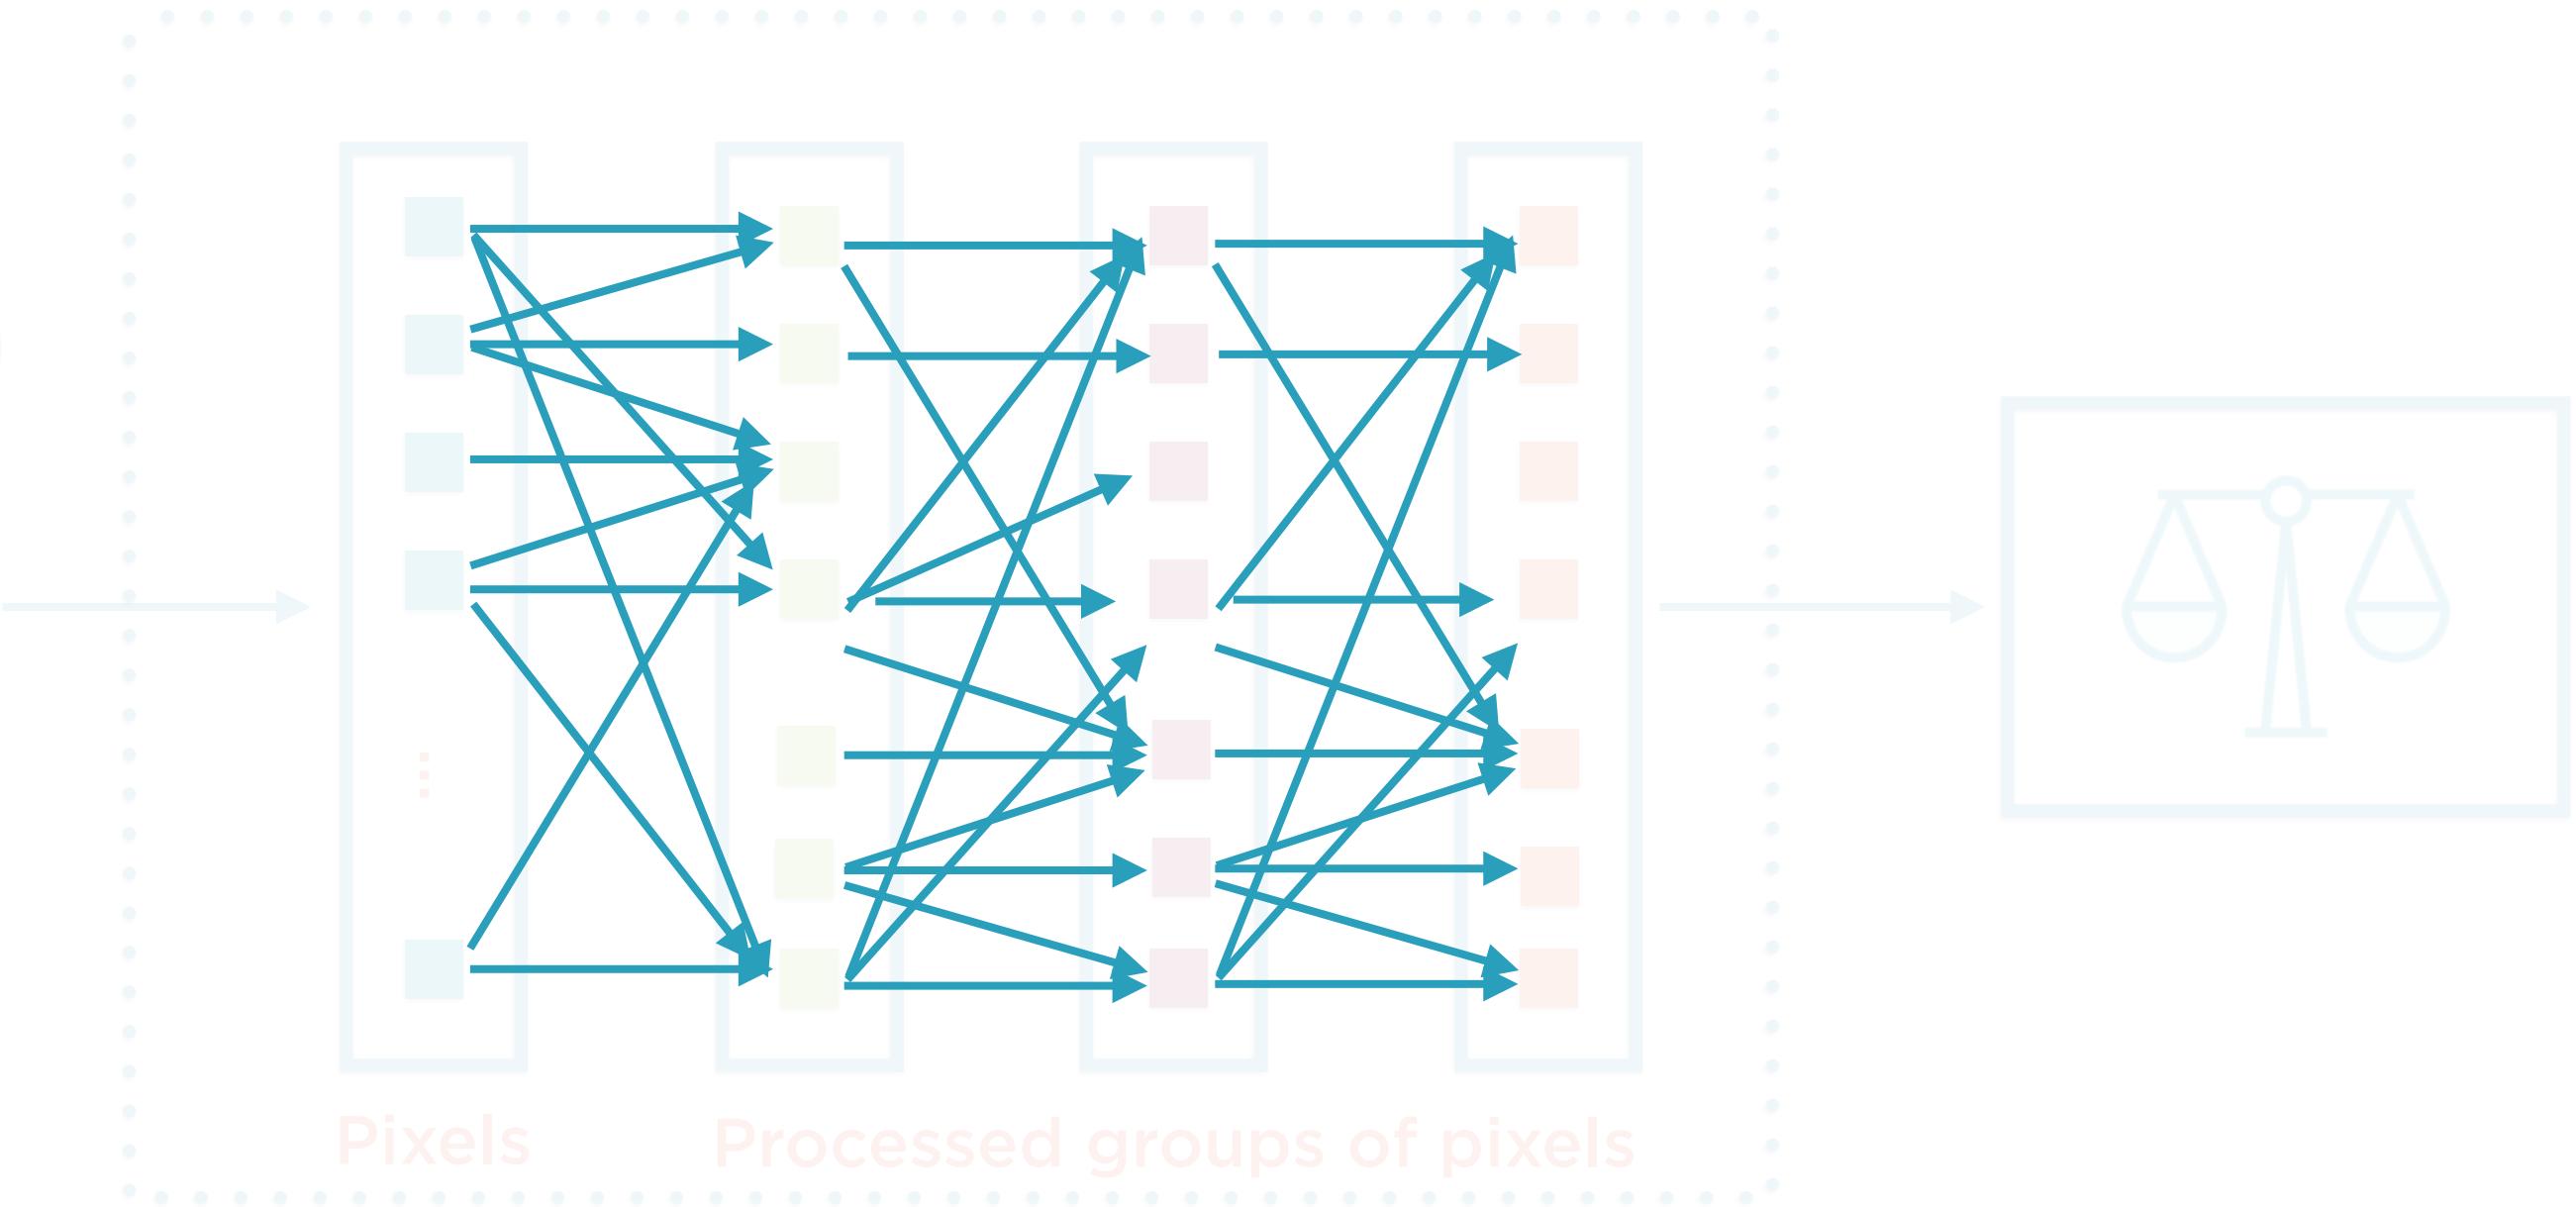
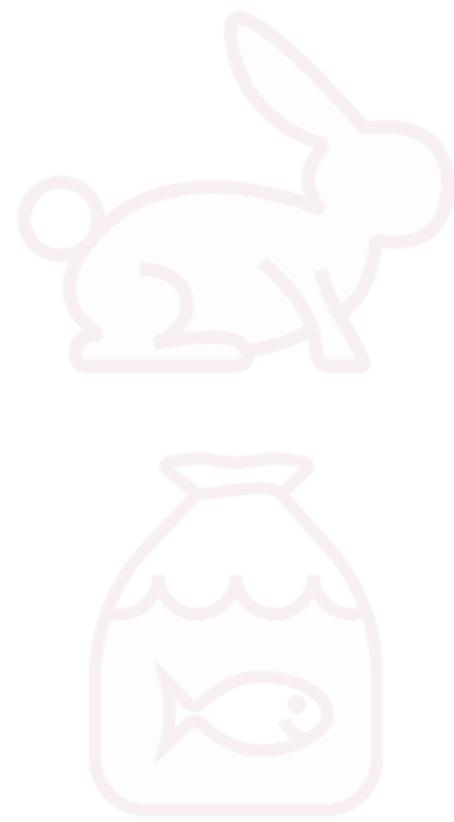
ML-based Classifier



Neural Networks

Neurons in a neural network can be connected in very complex ways...

Neural Networks Introduced



Corpus of
Images

Neurons in a neural network can be
connected in very complex ways...

ML-based Classifier

Neural Networks

Neurons in a neural network can be connected in very complex ways...

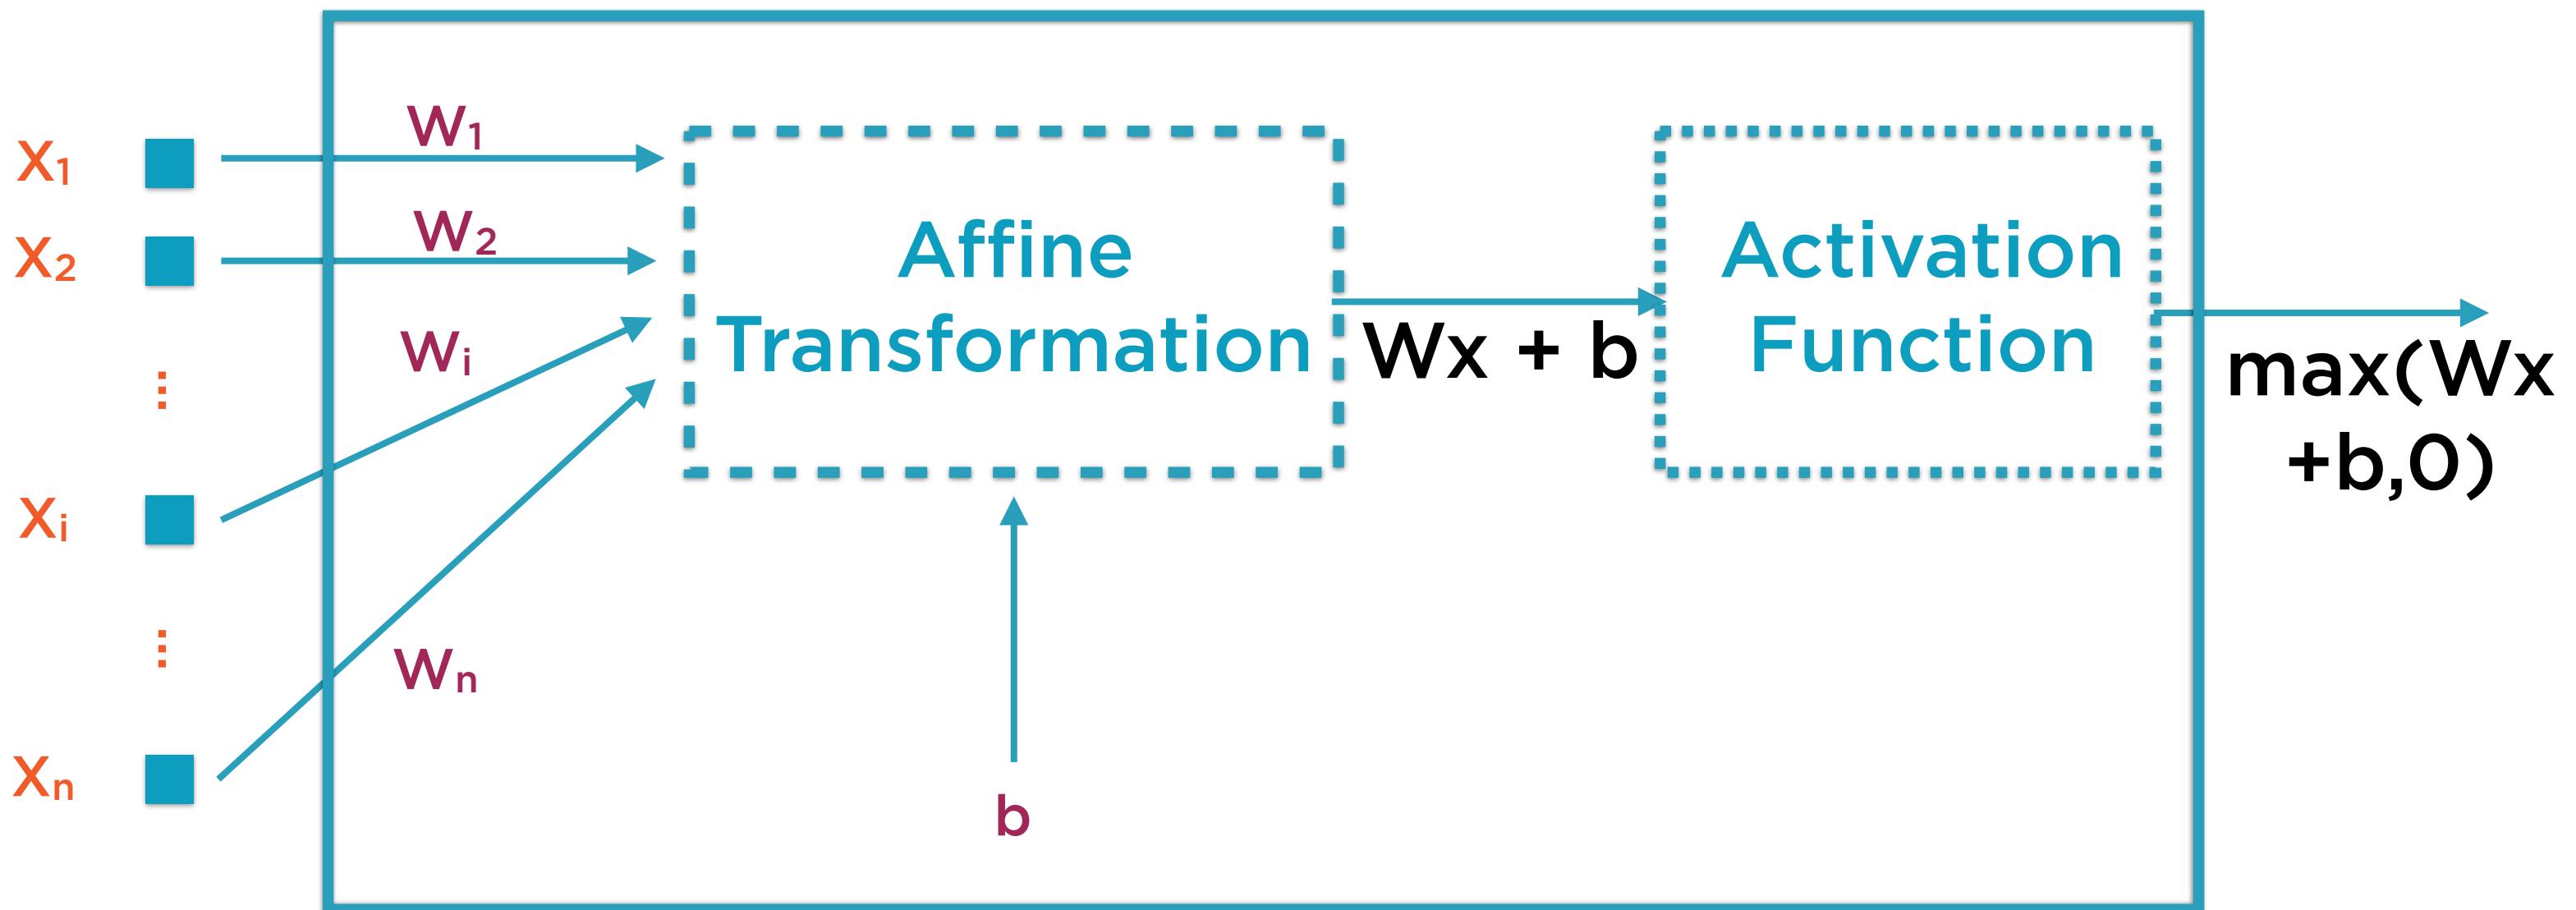
...But each neuron only applies **two simple functions to its inputs**

Neural Networks

**Neurons in a neural network can be connected in very complex ways...
...But each neuron only applies two simple functions to its inputs**

- A linear (affine) transformation
- An activation function

Operation of a Single Neuron



Each neuron only applies two simple functions to its inputs

Operation of a Single Neuron

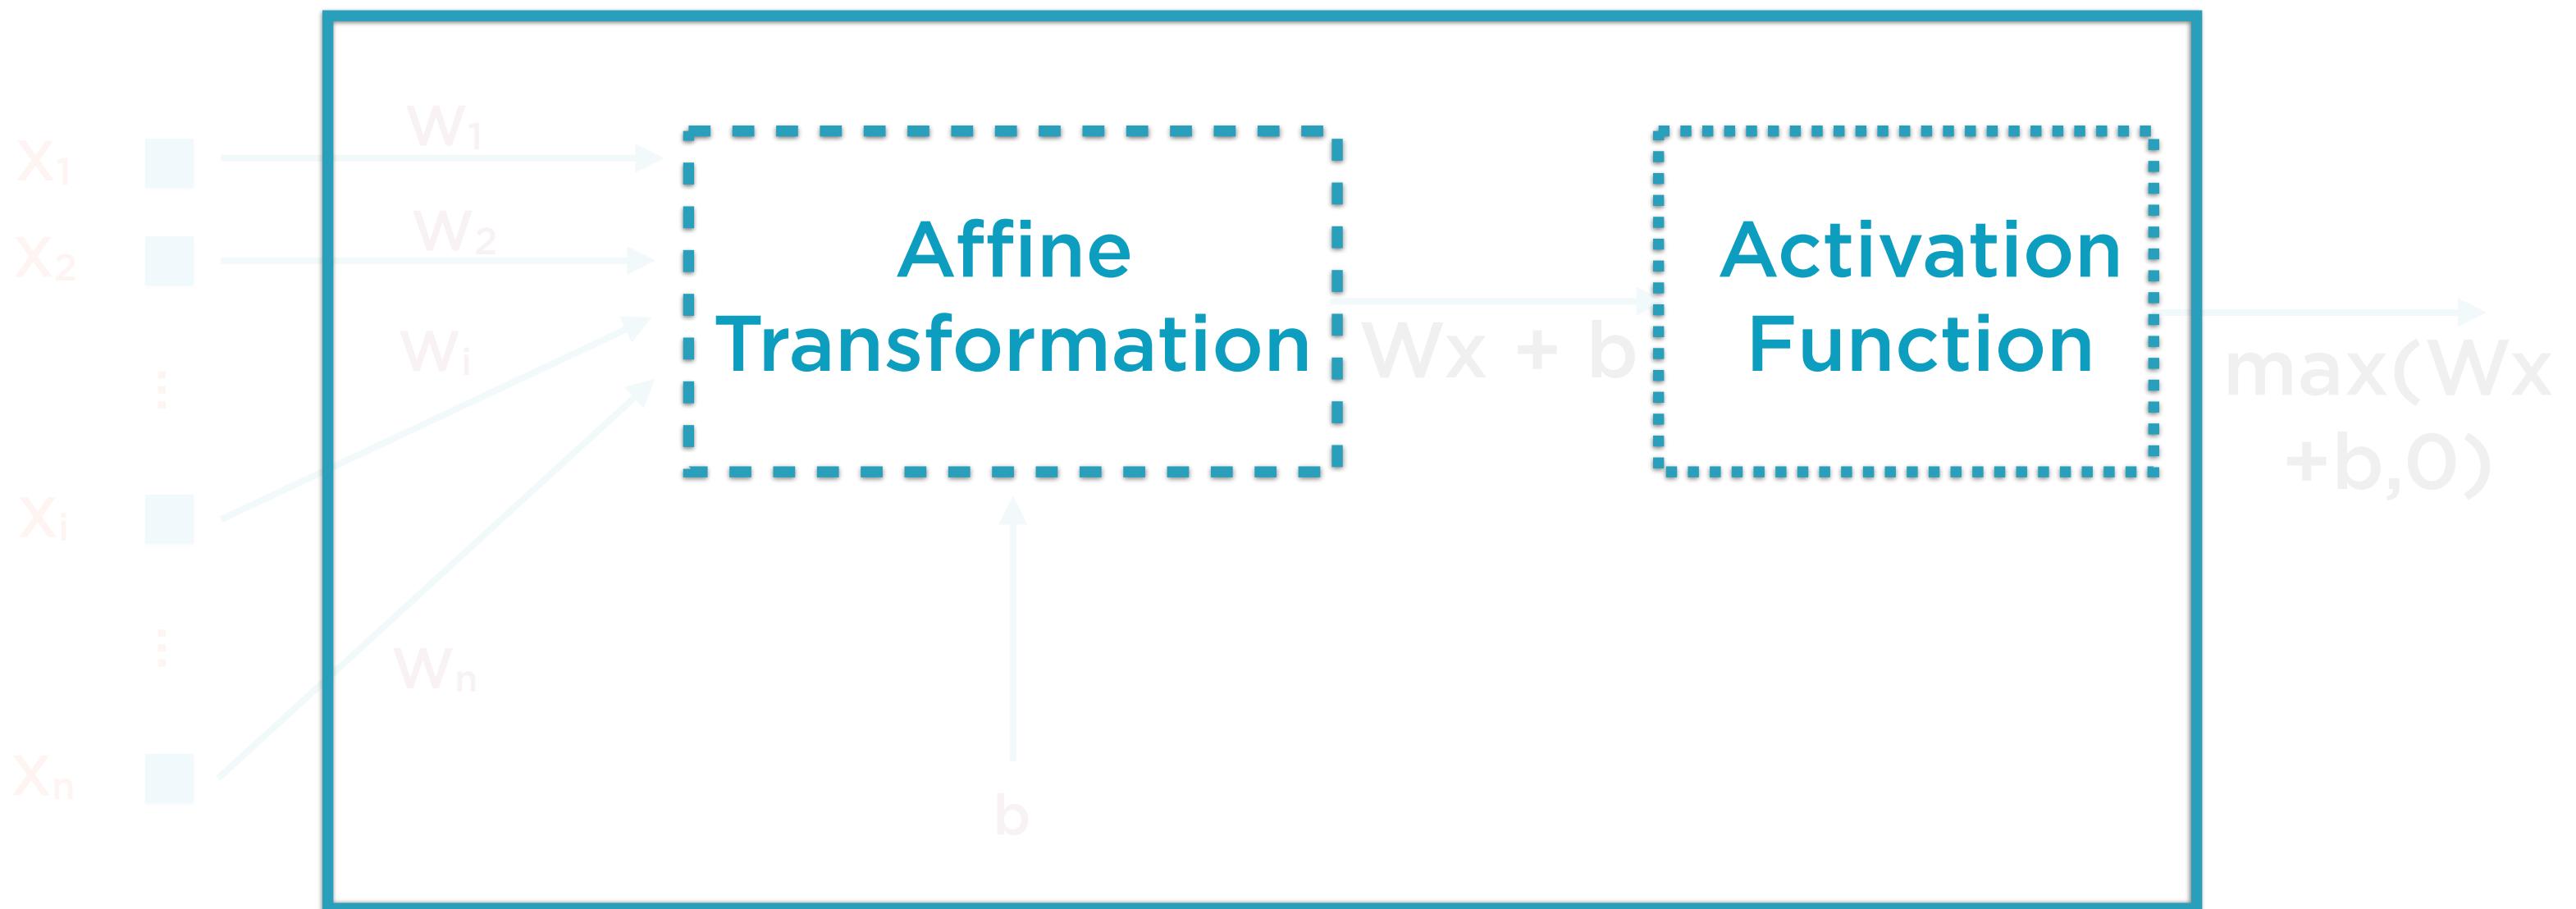


Inputs into
the neuron

Operation of a Single Neuron

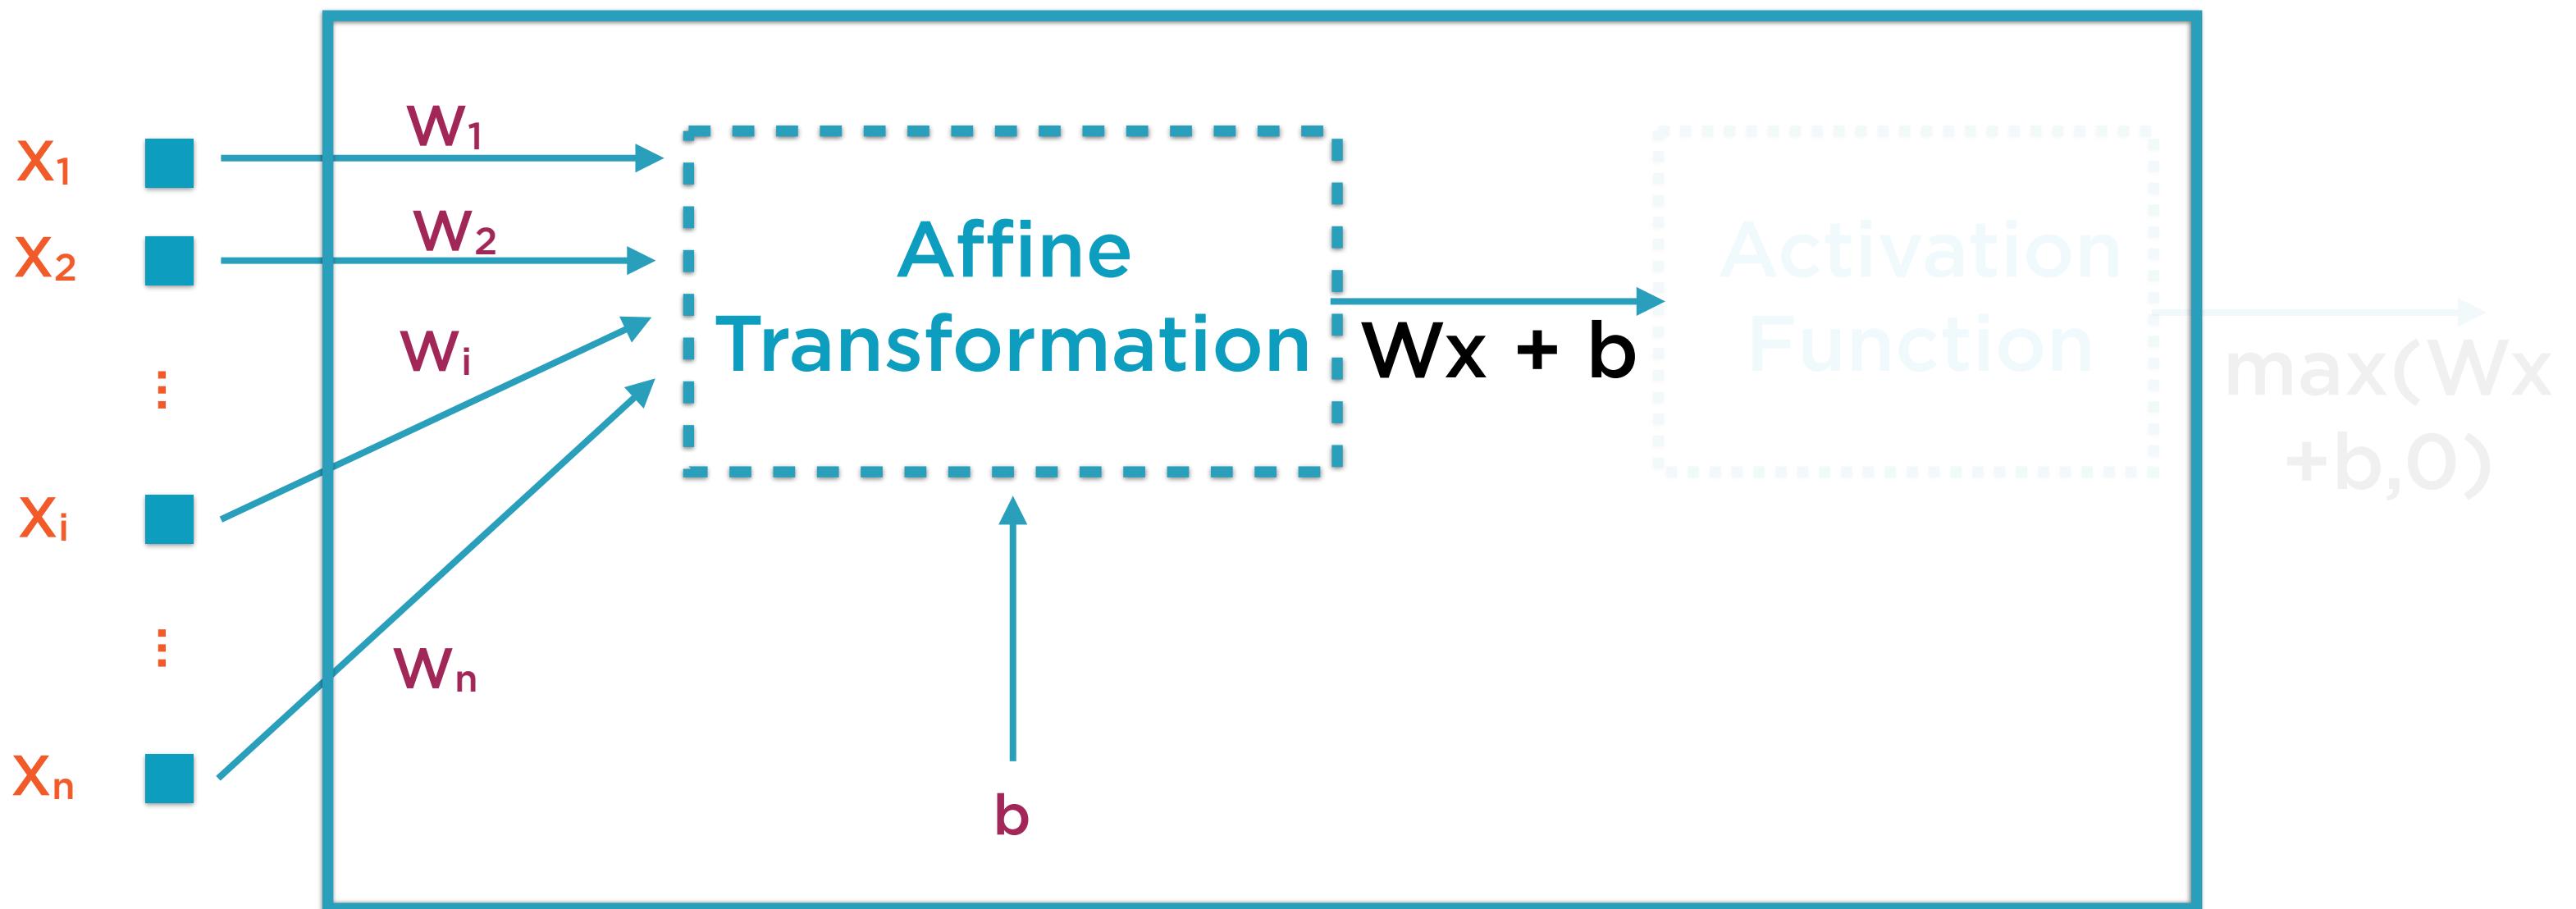


Operation of a Single Neuron



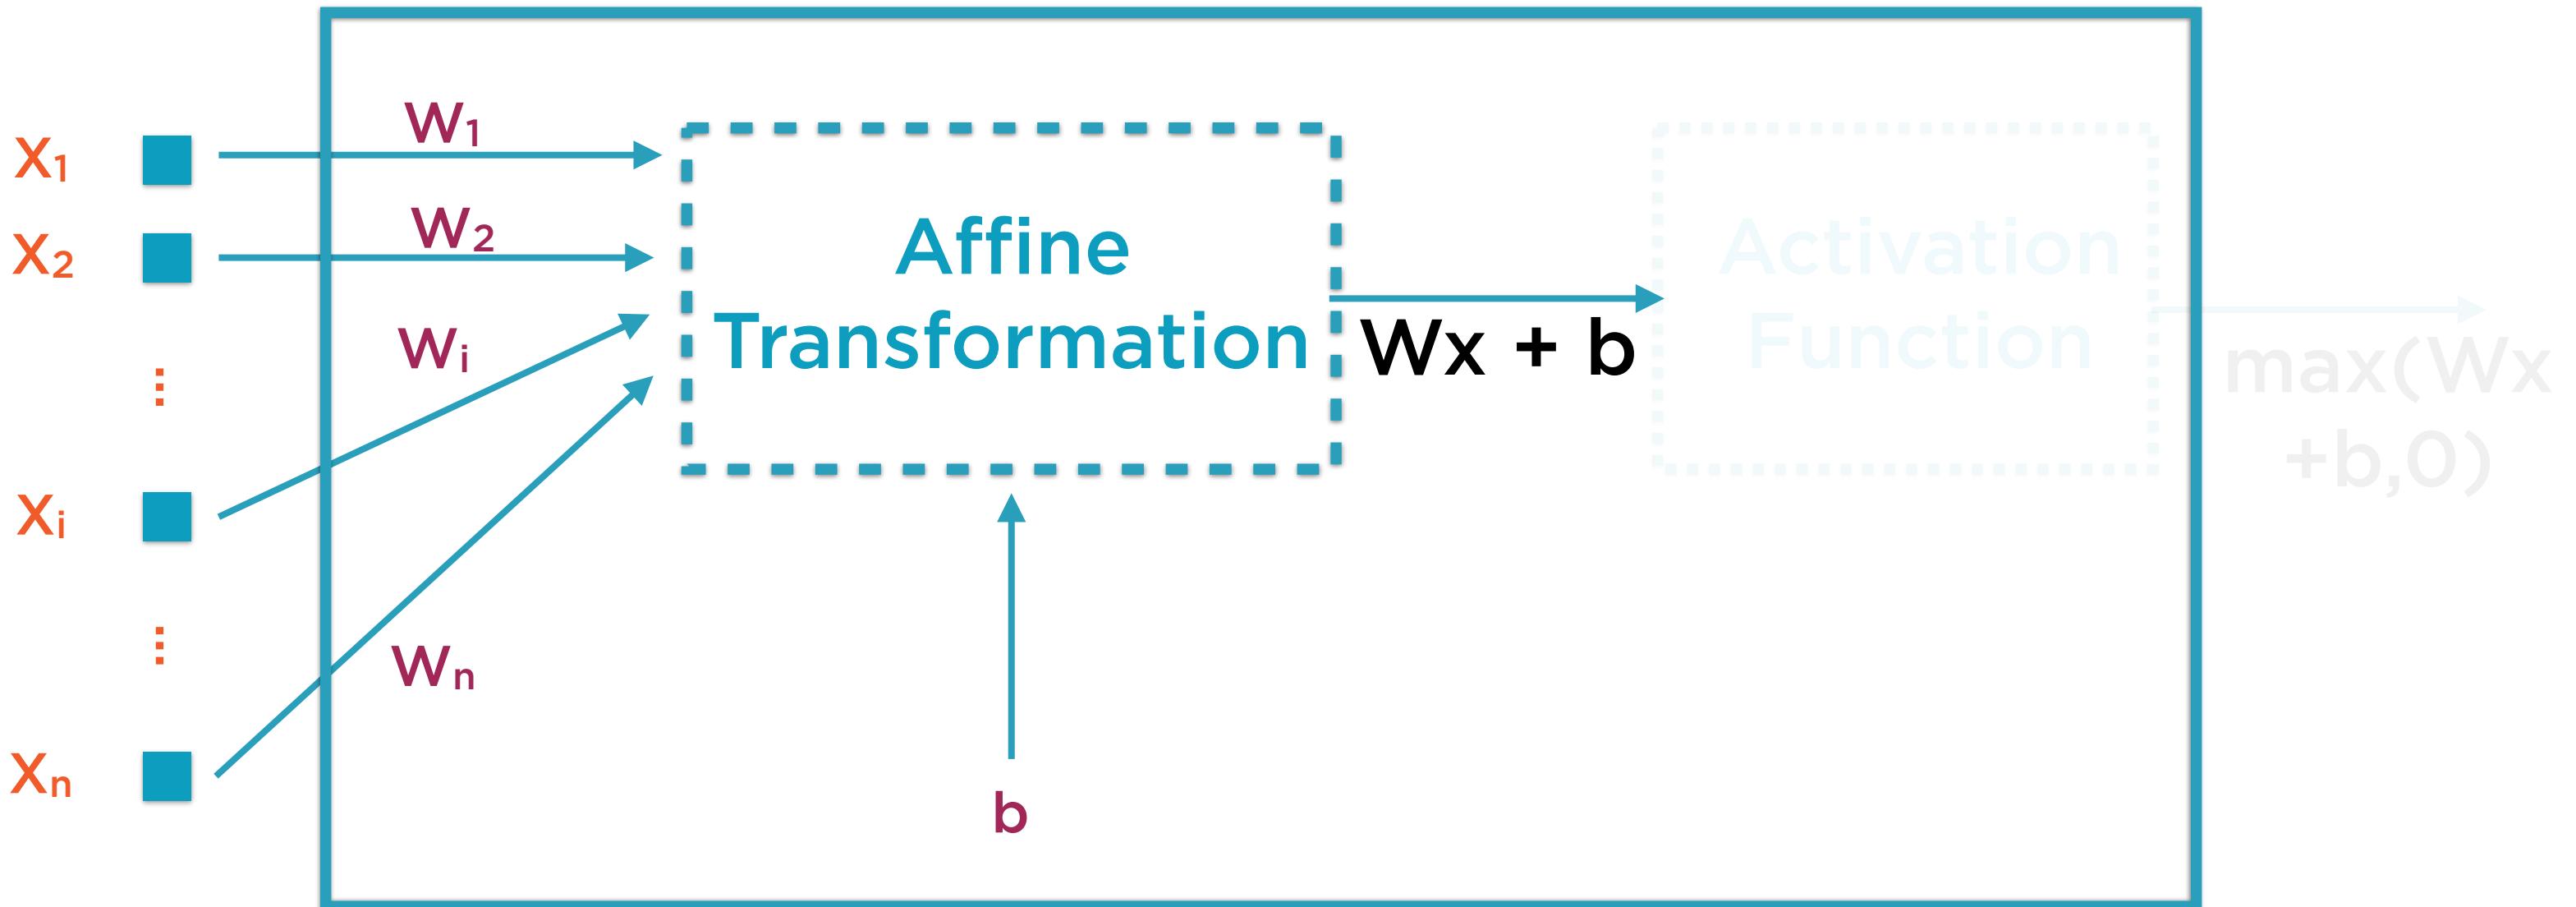
Each neuron only applies two simple functions to its inputs

Operation of a Single Neuron



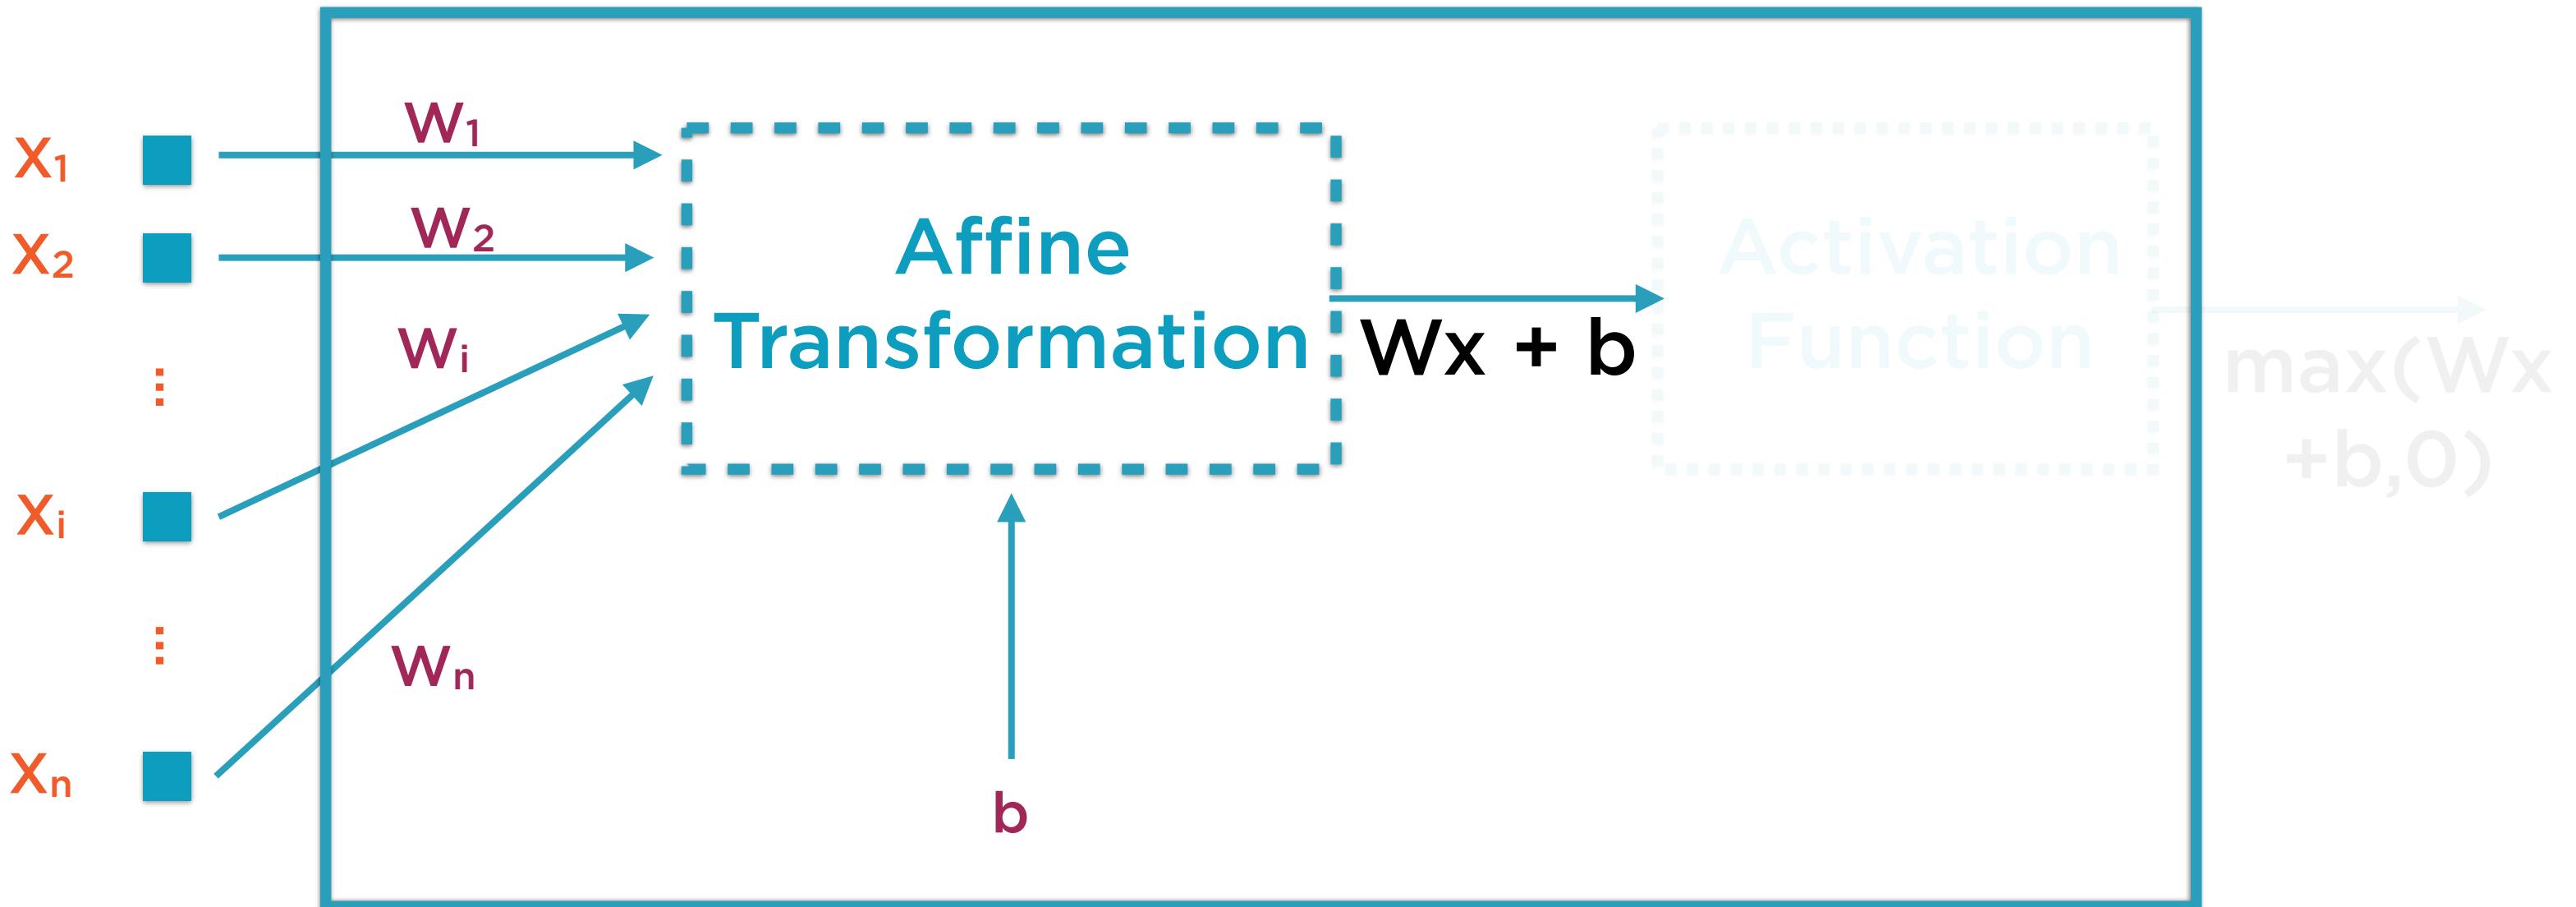
The affine transformation is just a weighted sum with a bias added

Operation of a Single Neuron



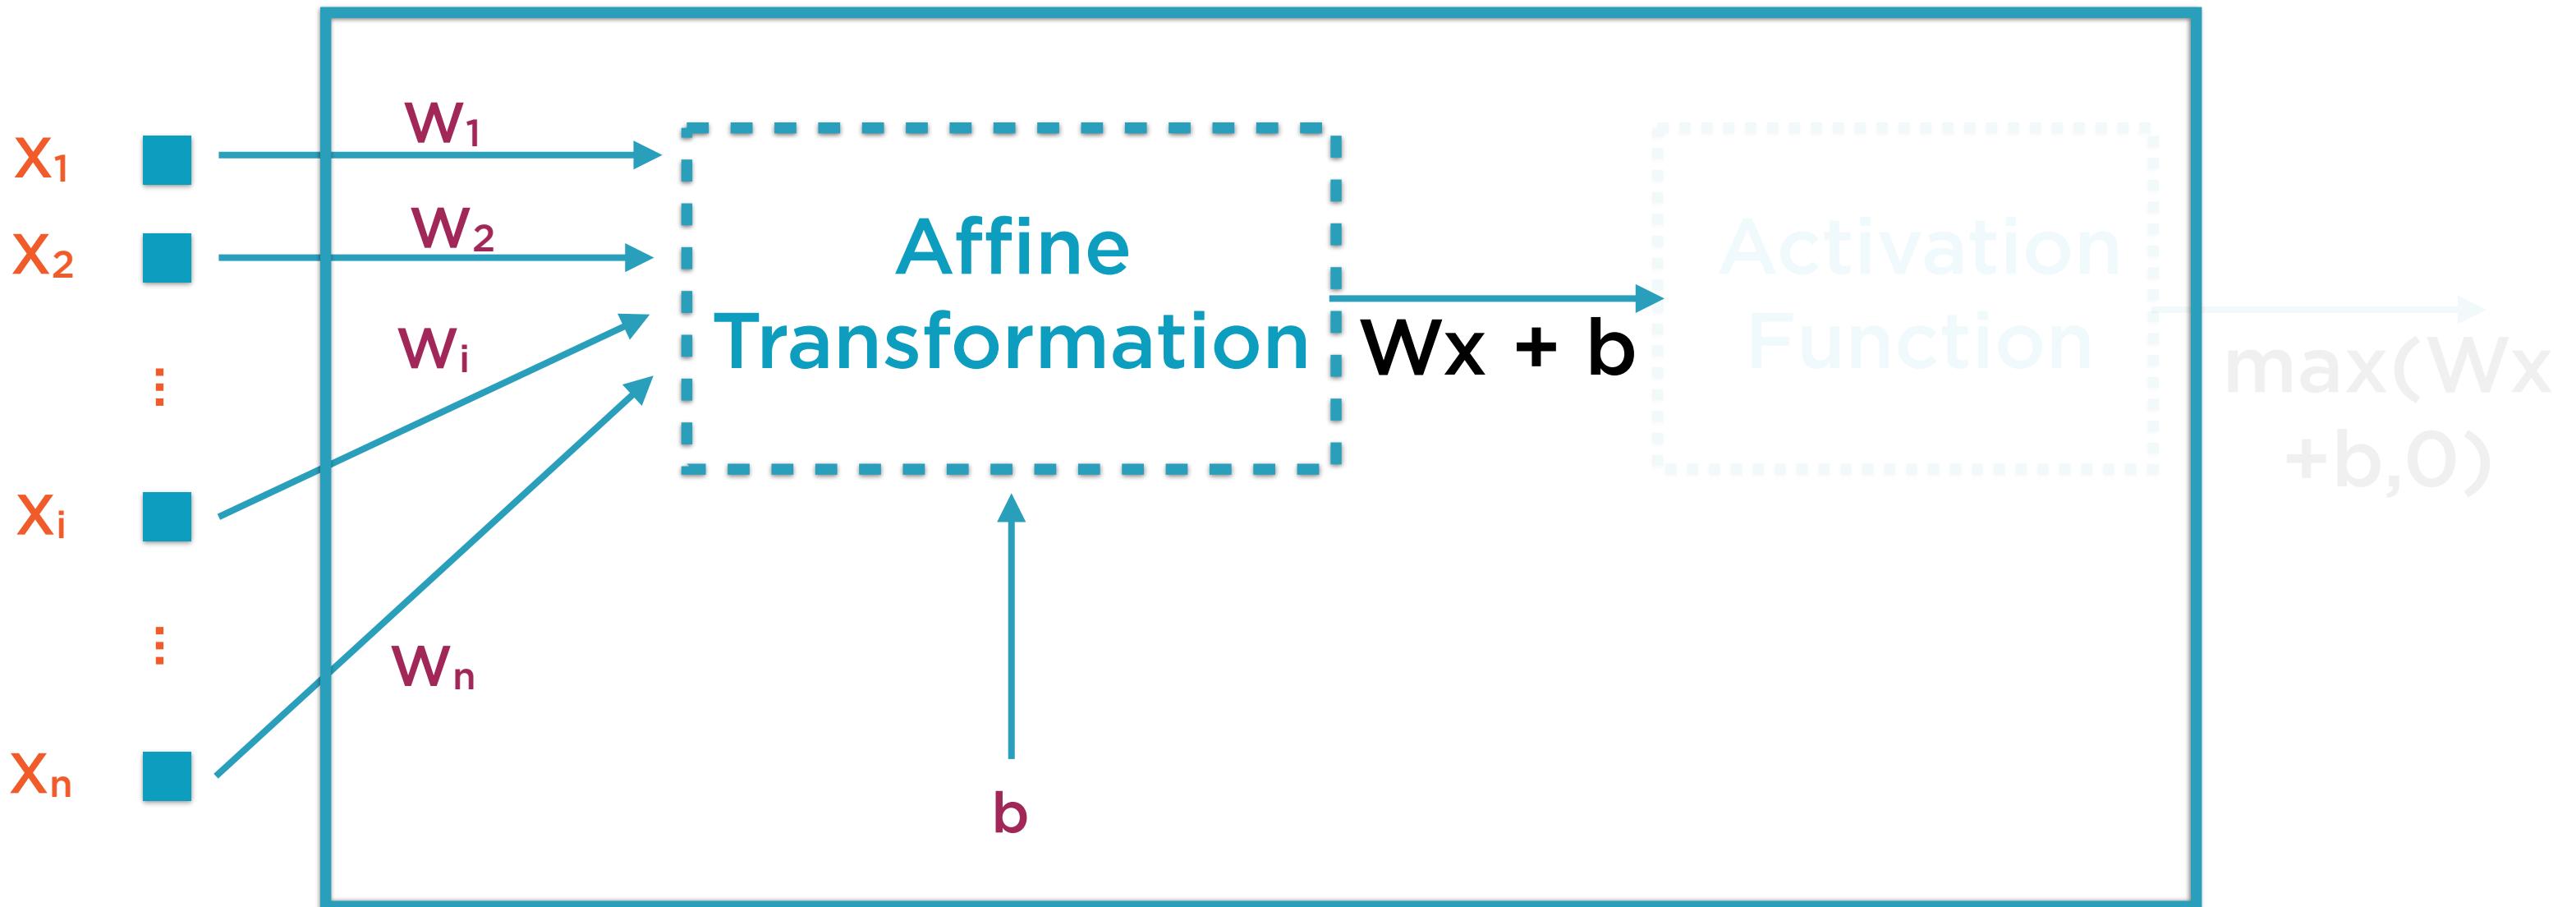
The values $w_1, w_2 \dots w_n$ are called the weights

Operation of a Single Neuron



The value b is called the bias

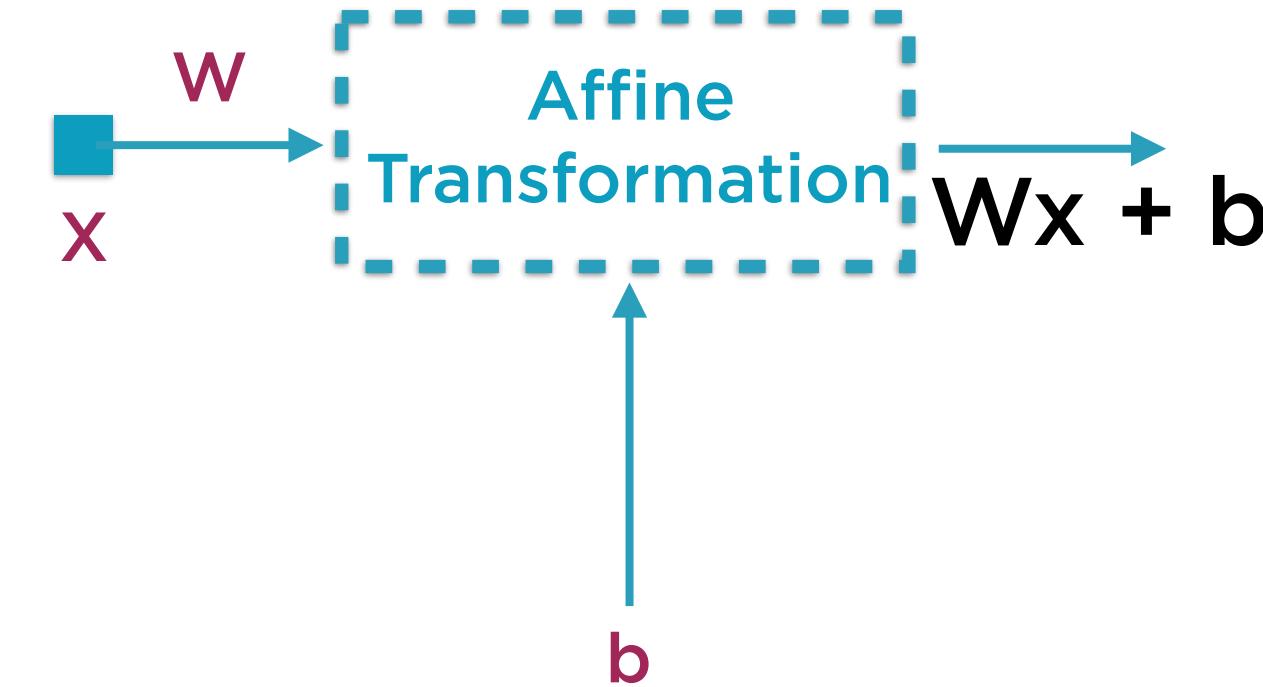
Operation of a Single Neuron



Where do the values of W and b come from?

The weights and biases of individual neurons are determined during the training process

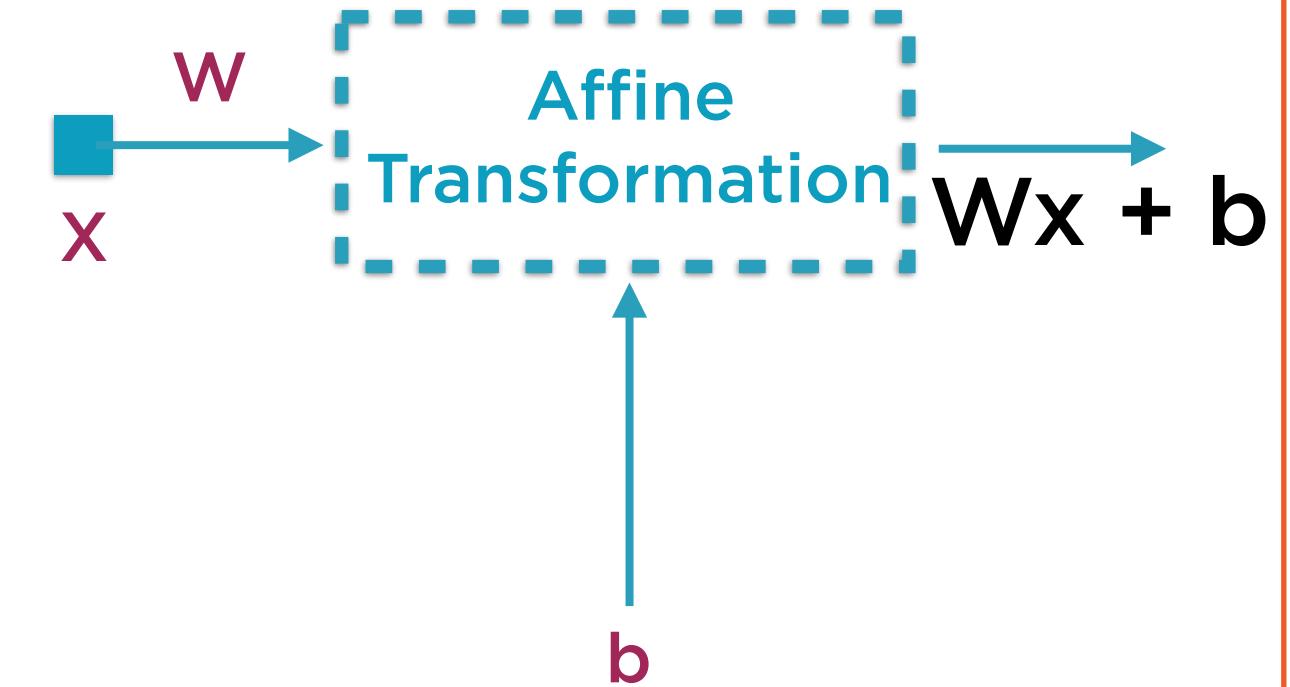
The actual training of a neural network is managed by TensorFlow



Finding the “best” values of W and b for each neuron is crucial

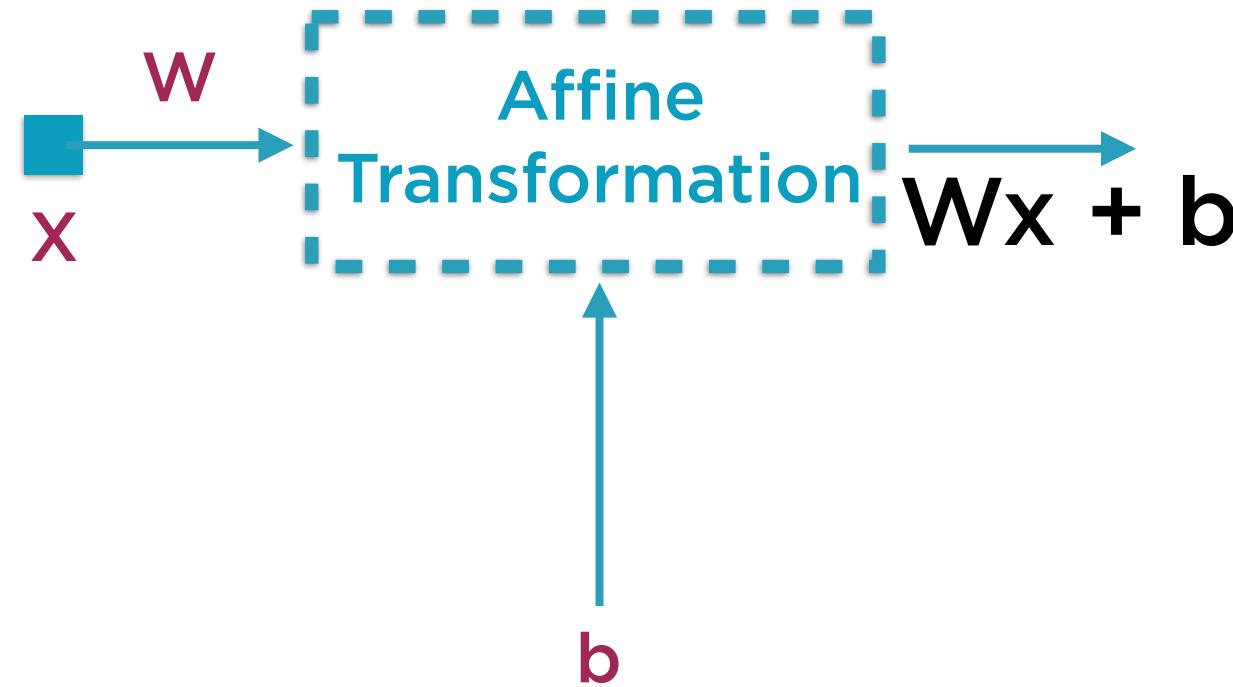
The “best” values are found using the cost function, optimiser and corpus...

...and the process of finding them is called the **training process**



Different types of neural networks wire up neurons in different ways

These interconnections can get very sophisticated...

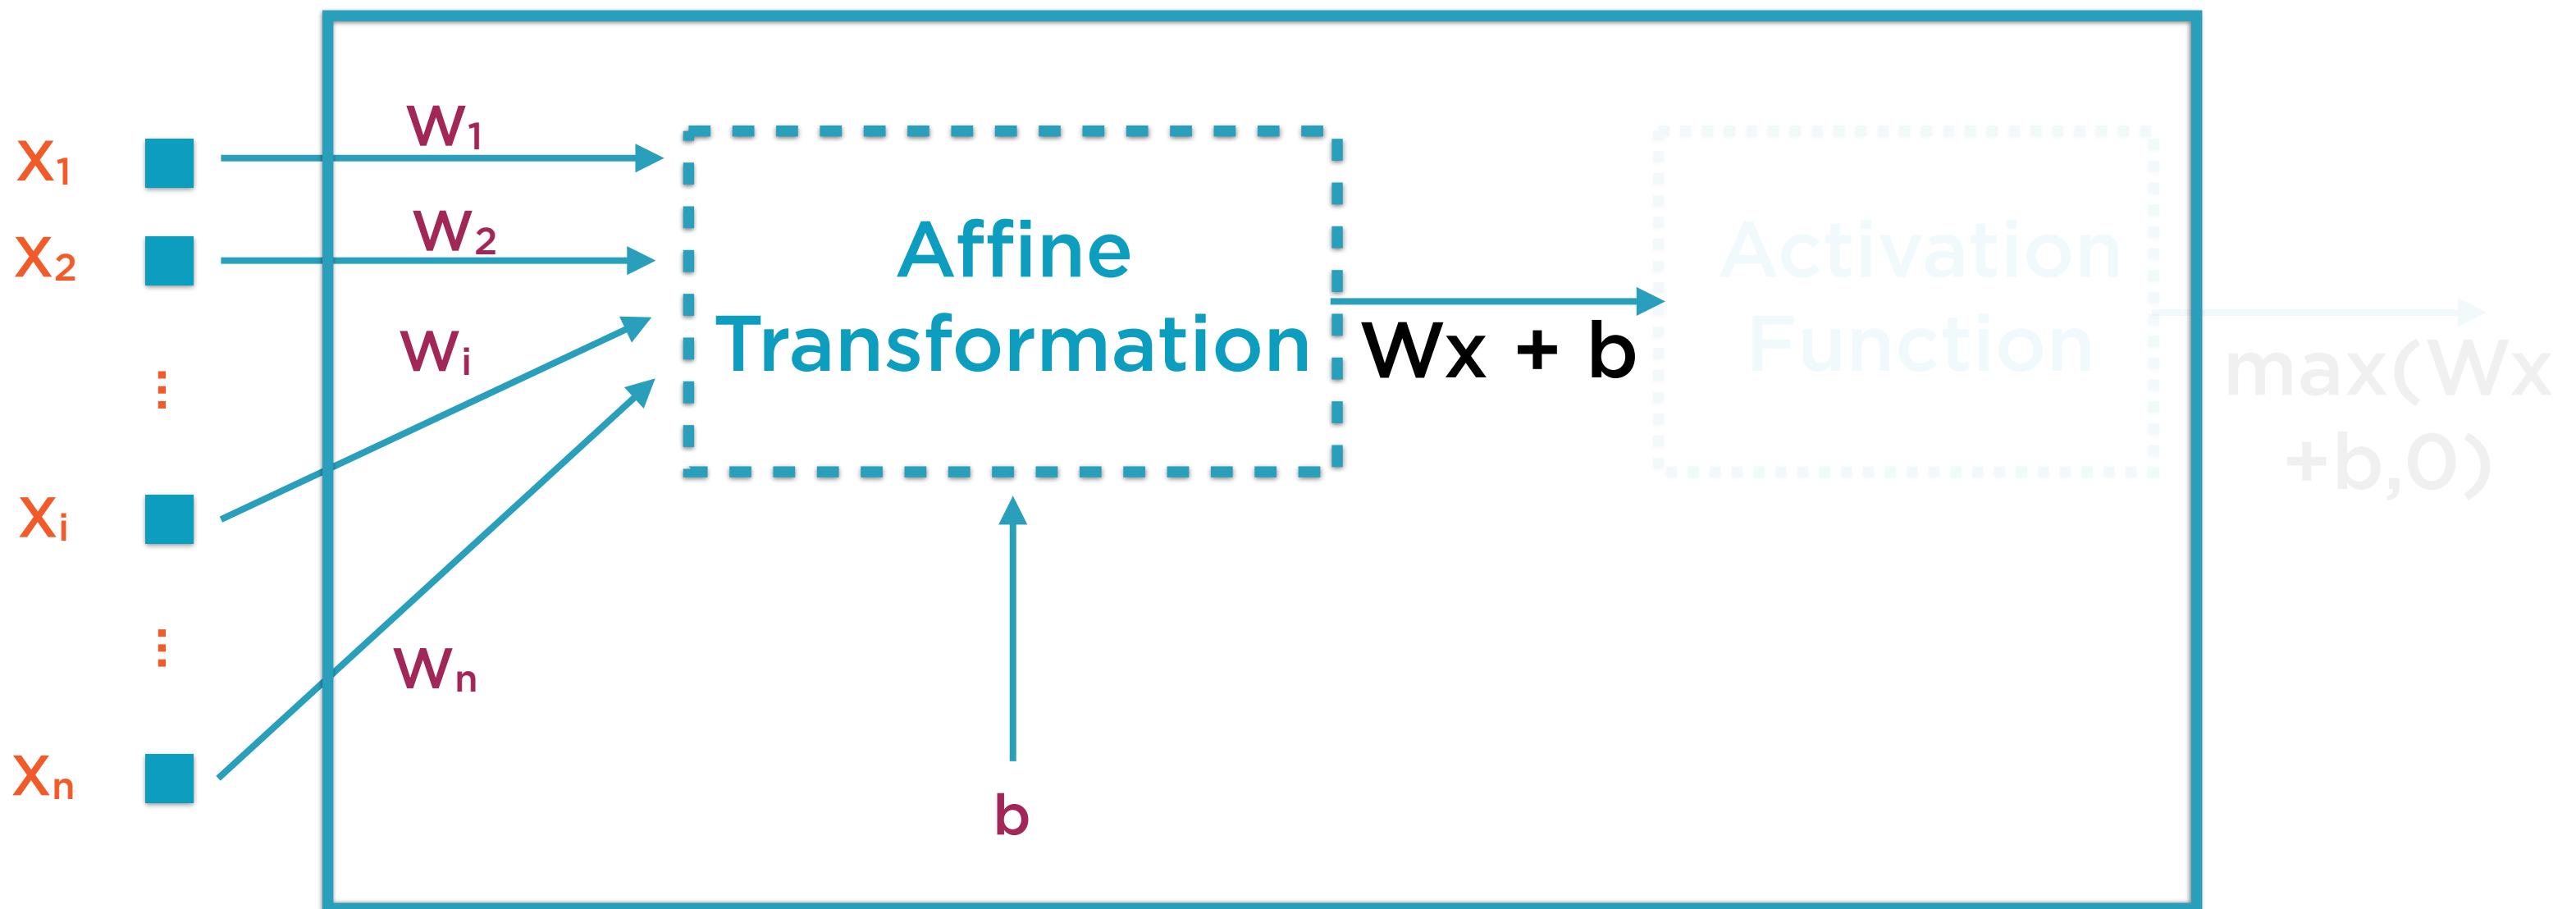


During training, the output of deeper layers may be “fed back” to find the best W, b

This is called back propagation

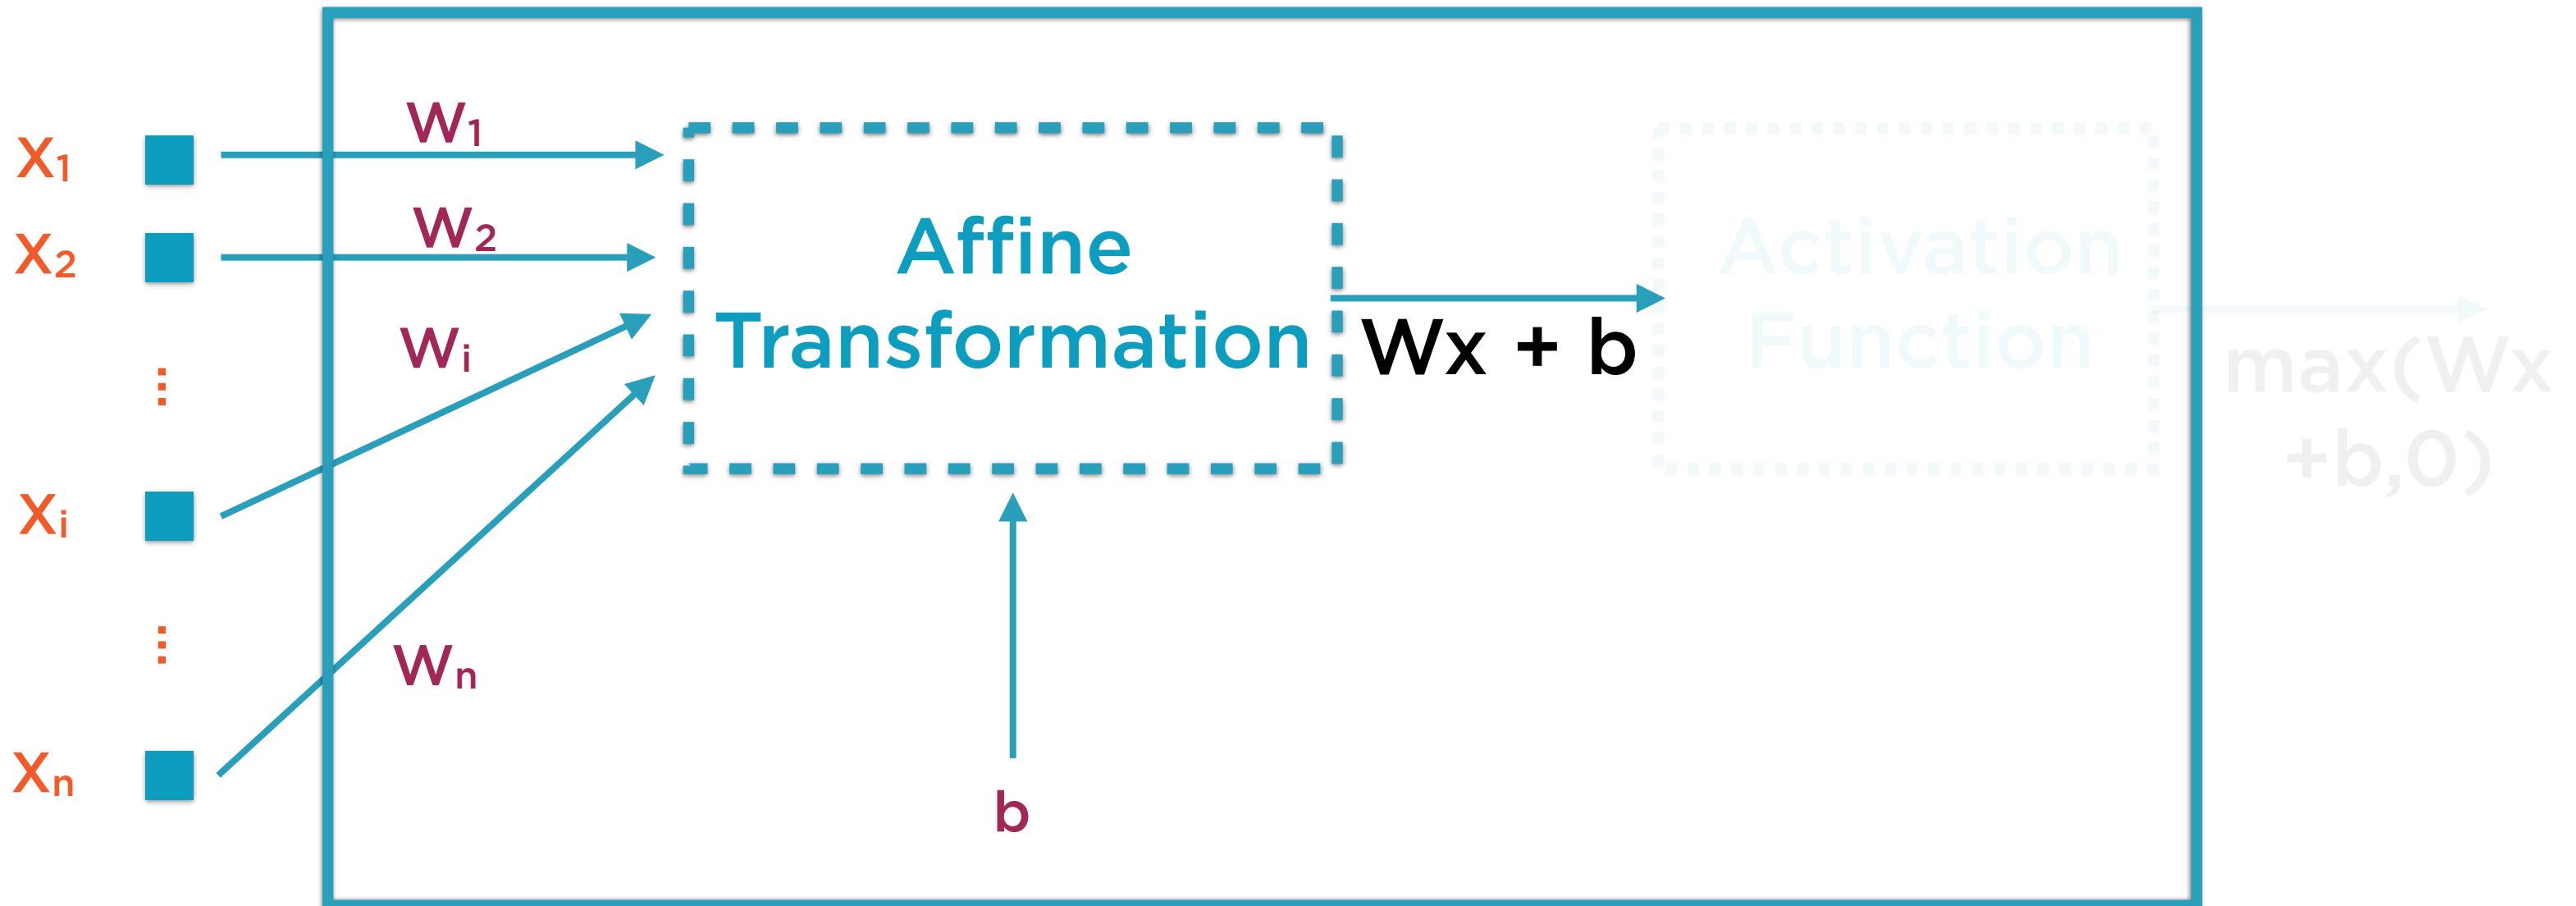
Back propagation is the standard algorithm for training neural networks

Operation of a Single Neuron



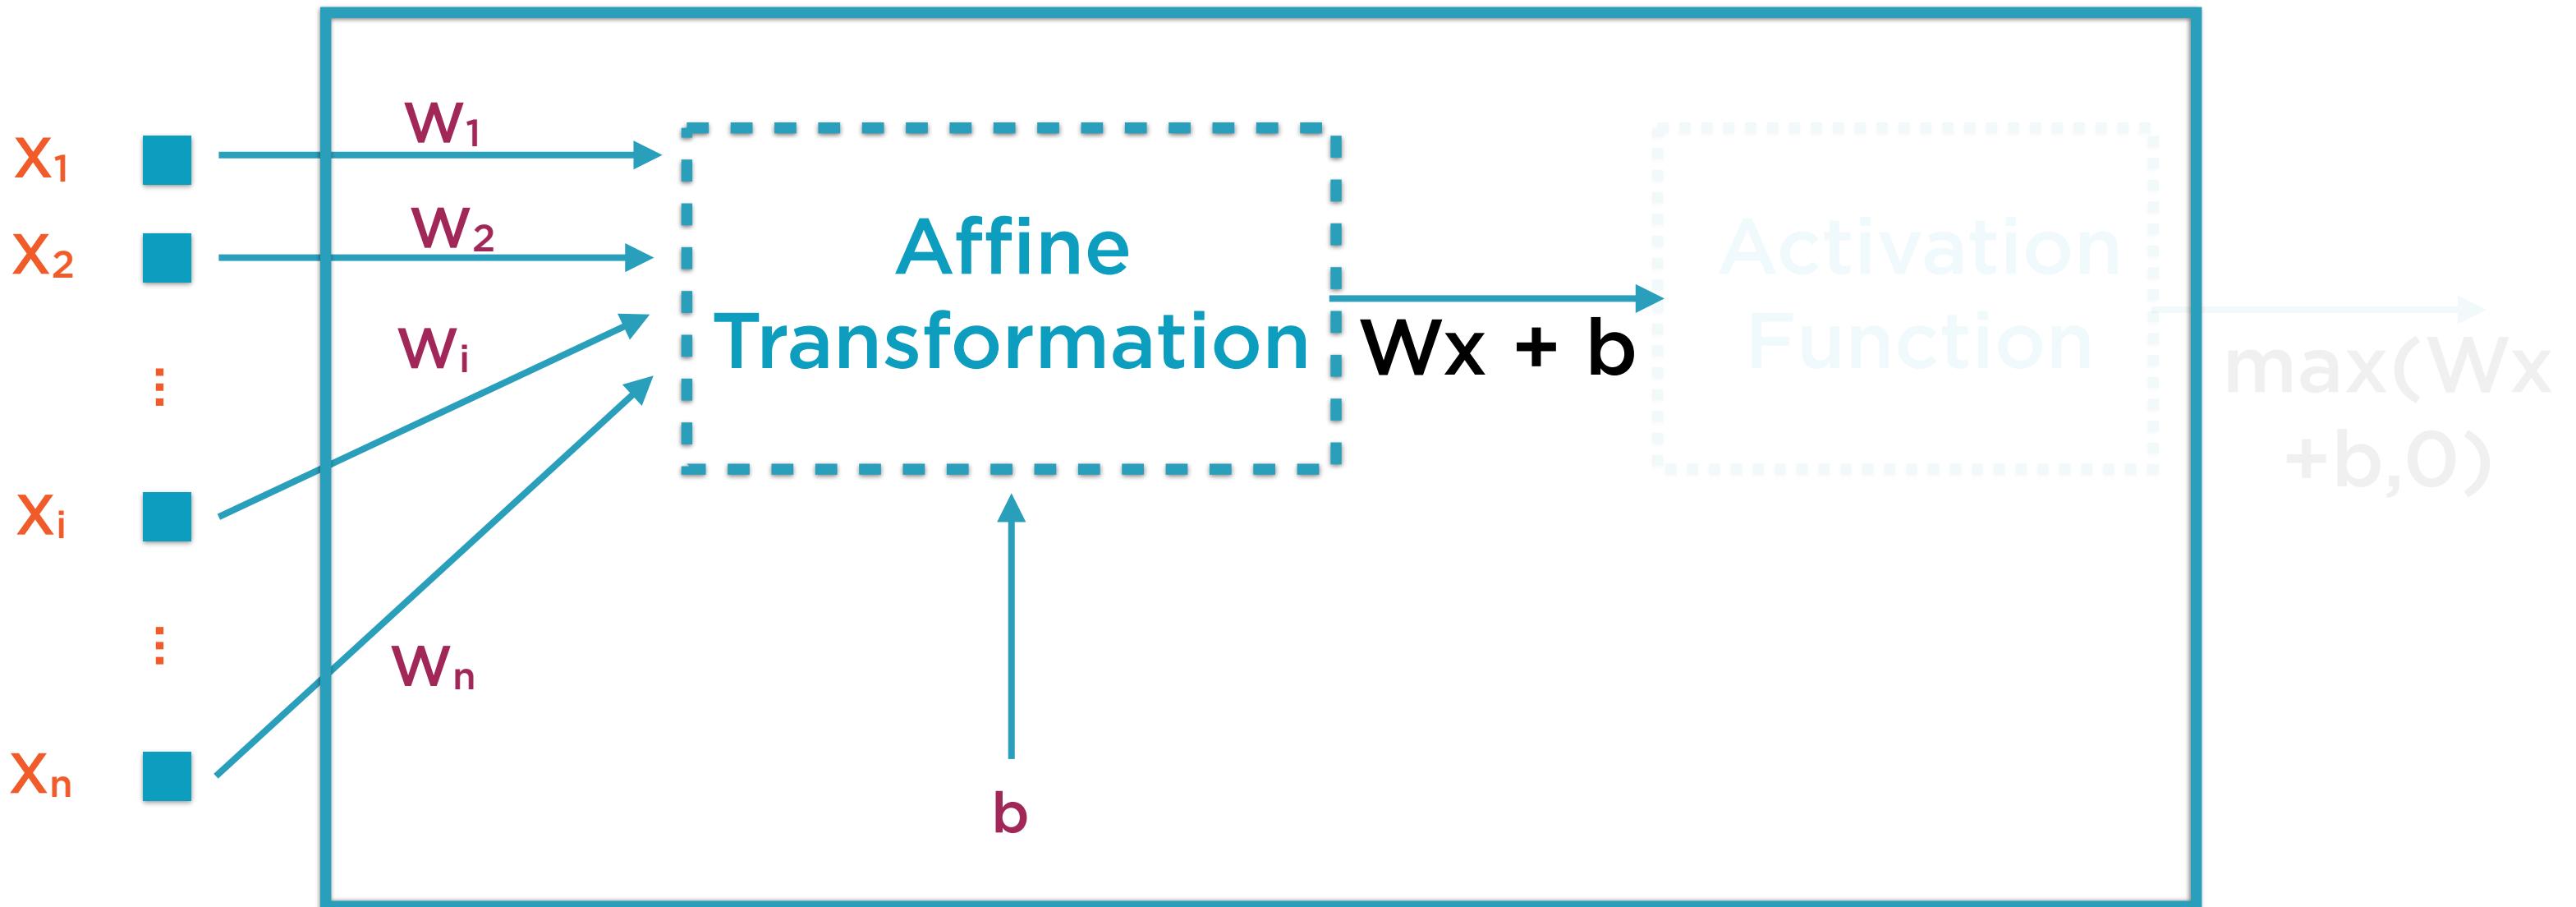
The training algorithm will use the weights to tell the neuron which inputs matter, and which do not...

Operation of a Single Neuron



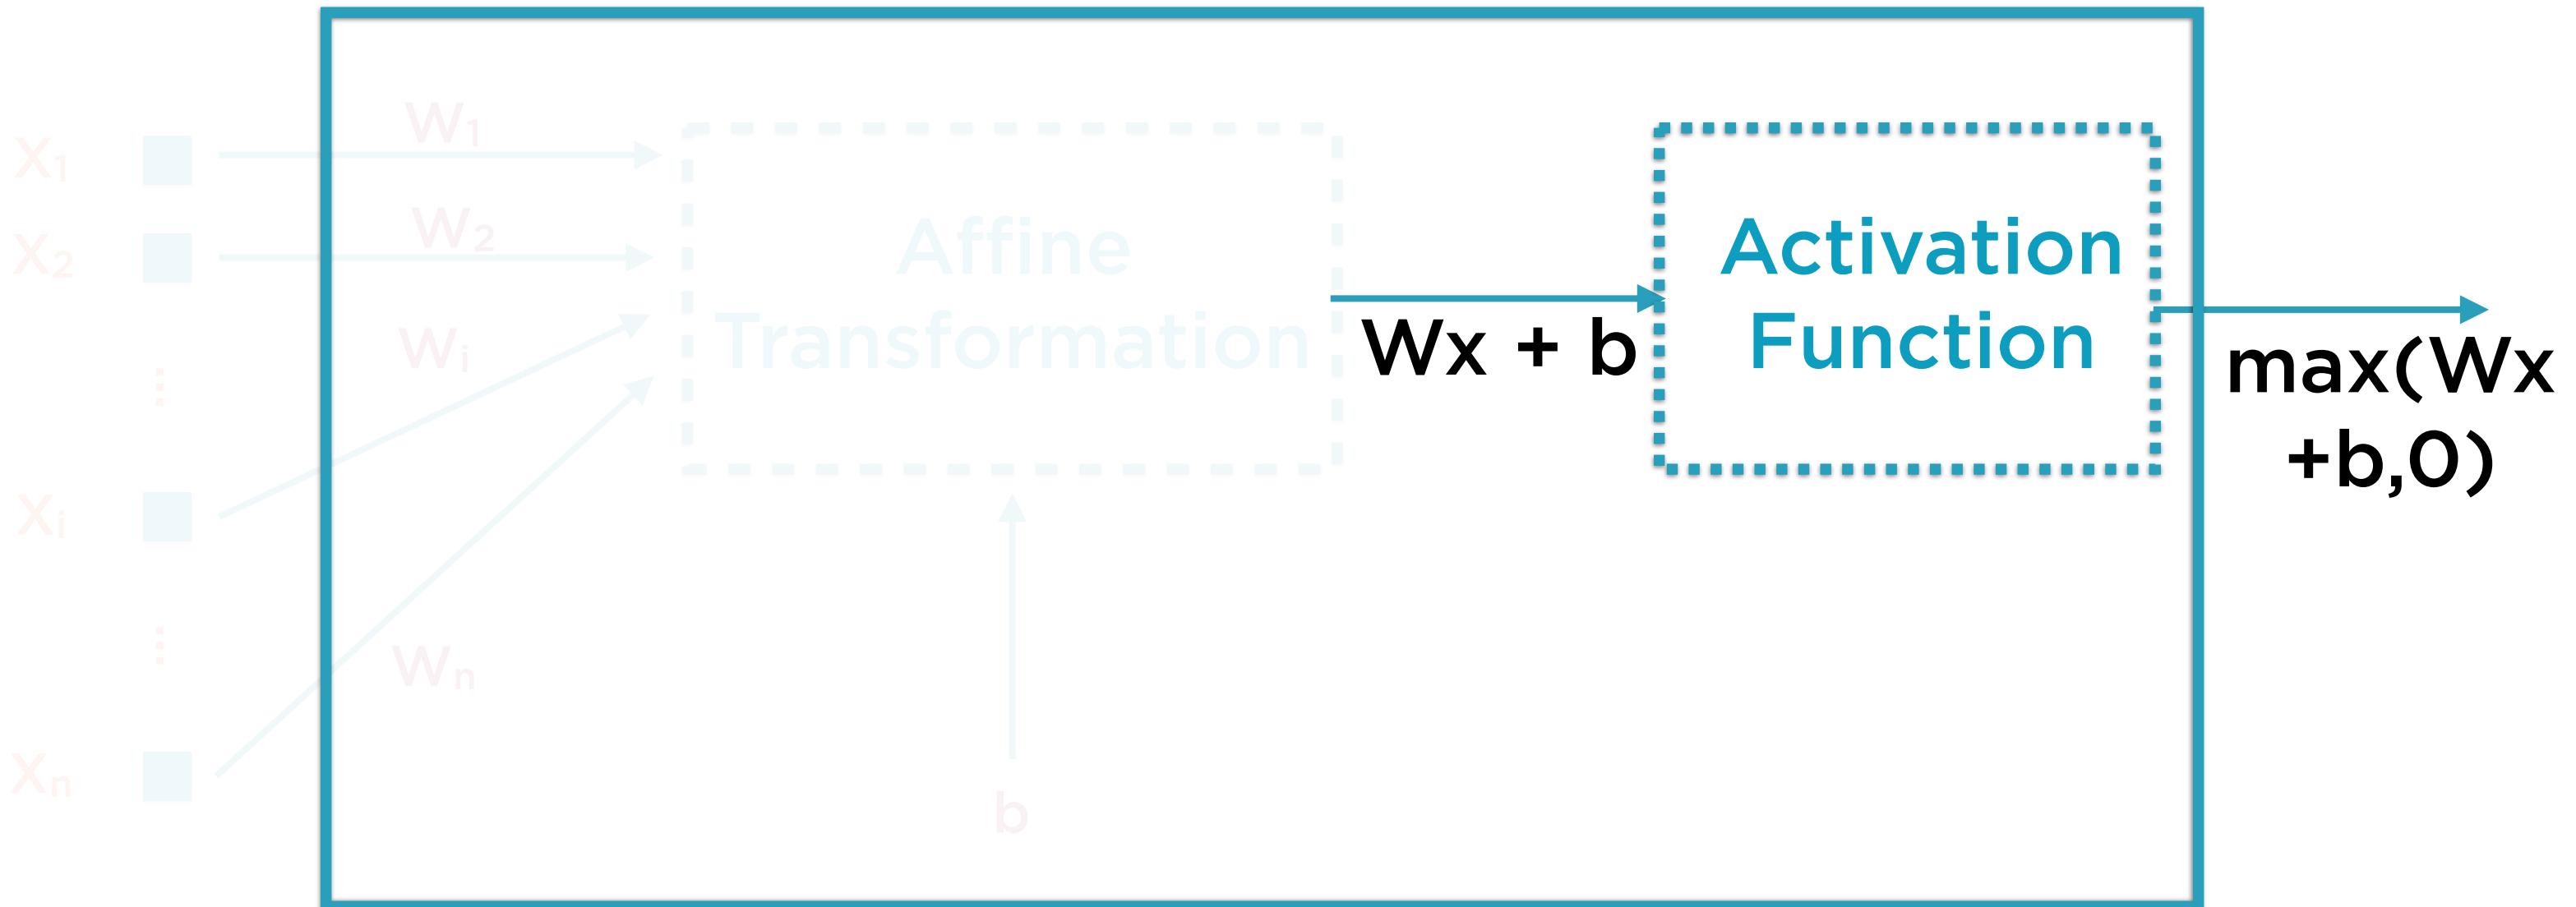
...and apply a corrective bias if needed

Operation of a Single Neuron



The linear output can only be used to learn linear functions, but we can easily generalize this...

Operation of a Single Neuron



The Activation function is a non-linear function, very often simply the $\max(0, \dots)$ function



The output of the affine transformation
is chained into an activation function



The activation function is needed for the neural network to predict non-linear functions



The most common form of the activation function is the ReLU

ReLU : Rectified Linear Unit

$$\text{ReLU}(x) = \max(0, x)$$

Regression: The Simplest Neural Network

```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

“Learning” Arbitrarily Complex Functions

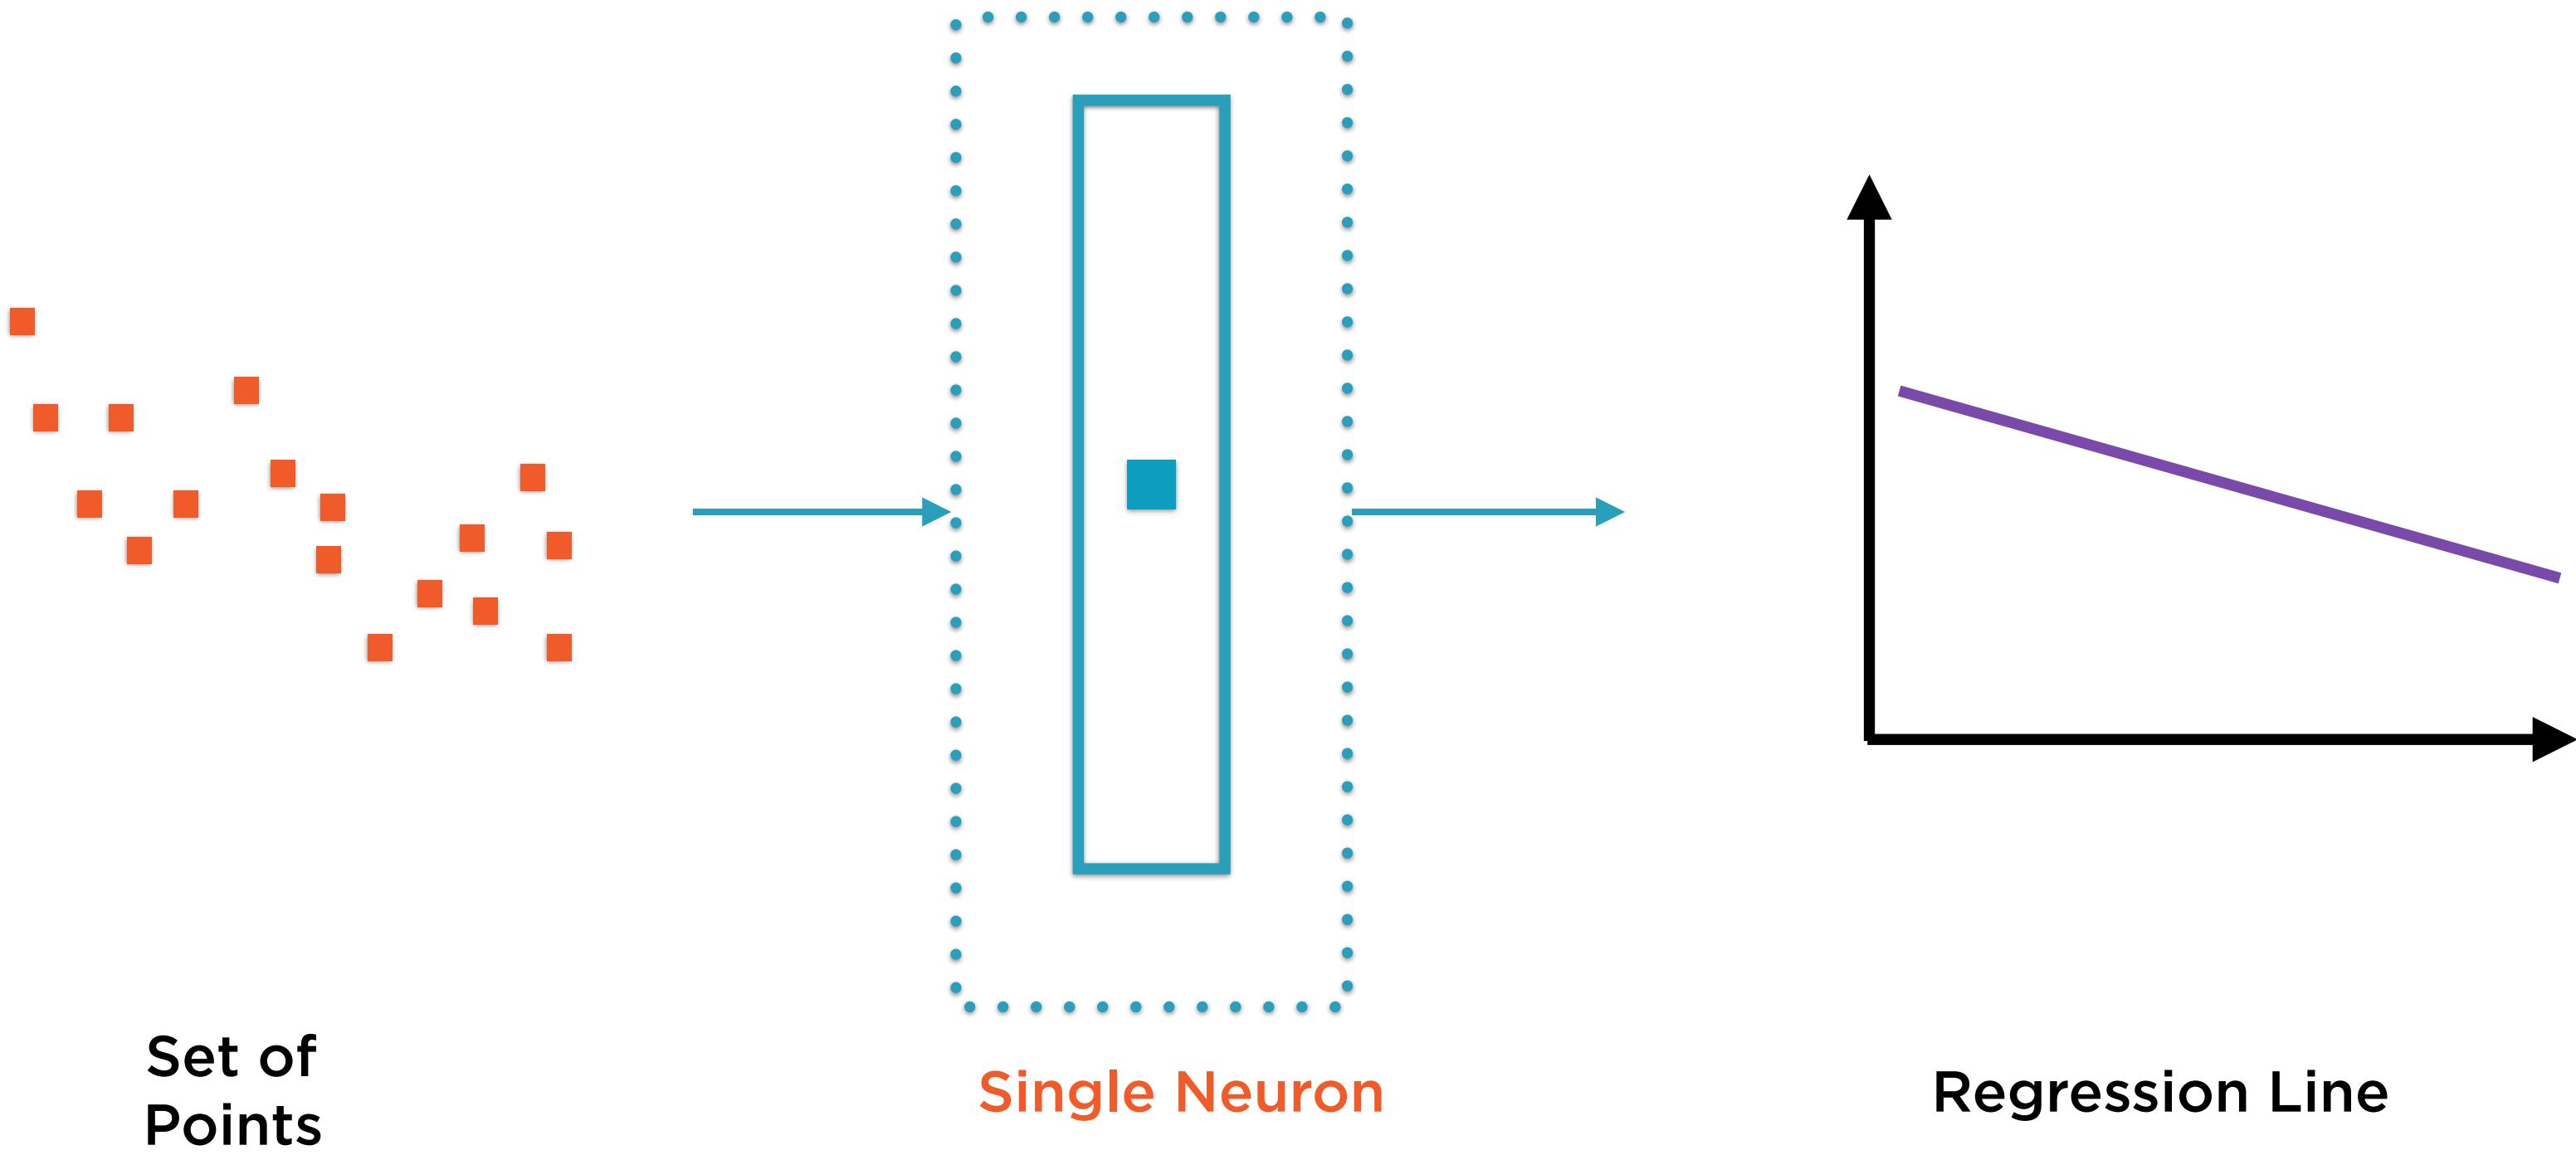
Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything

$$y = Wx + b$$

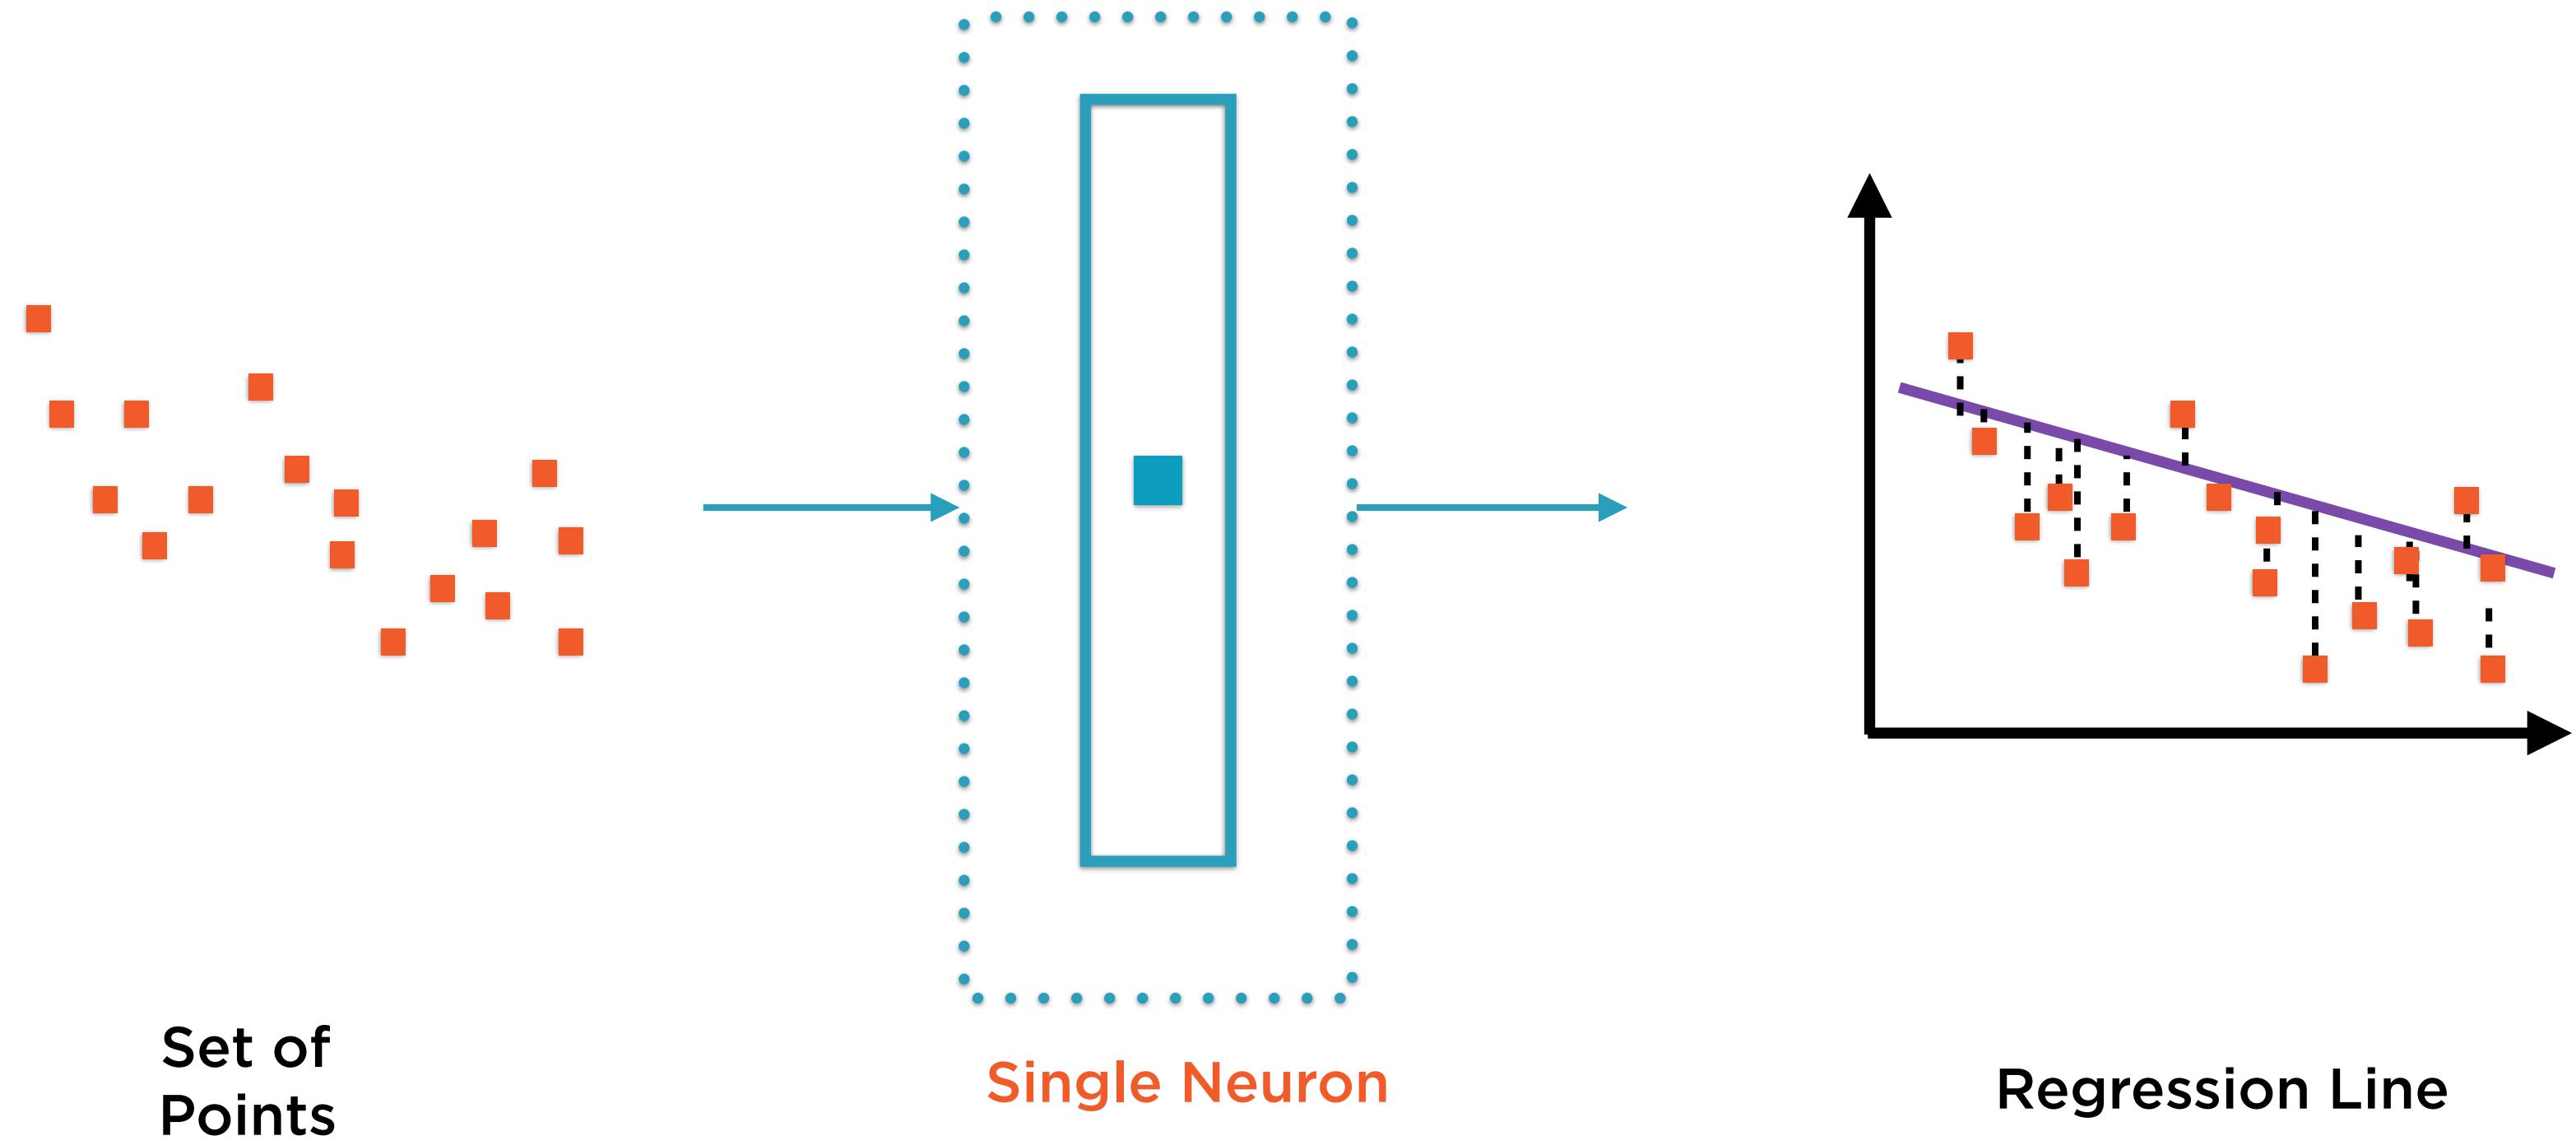
“Learning” Regression

Regression can be reverse-engineered by a single neuron

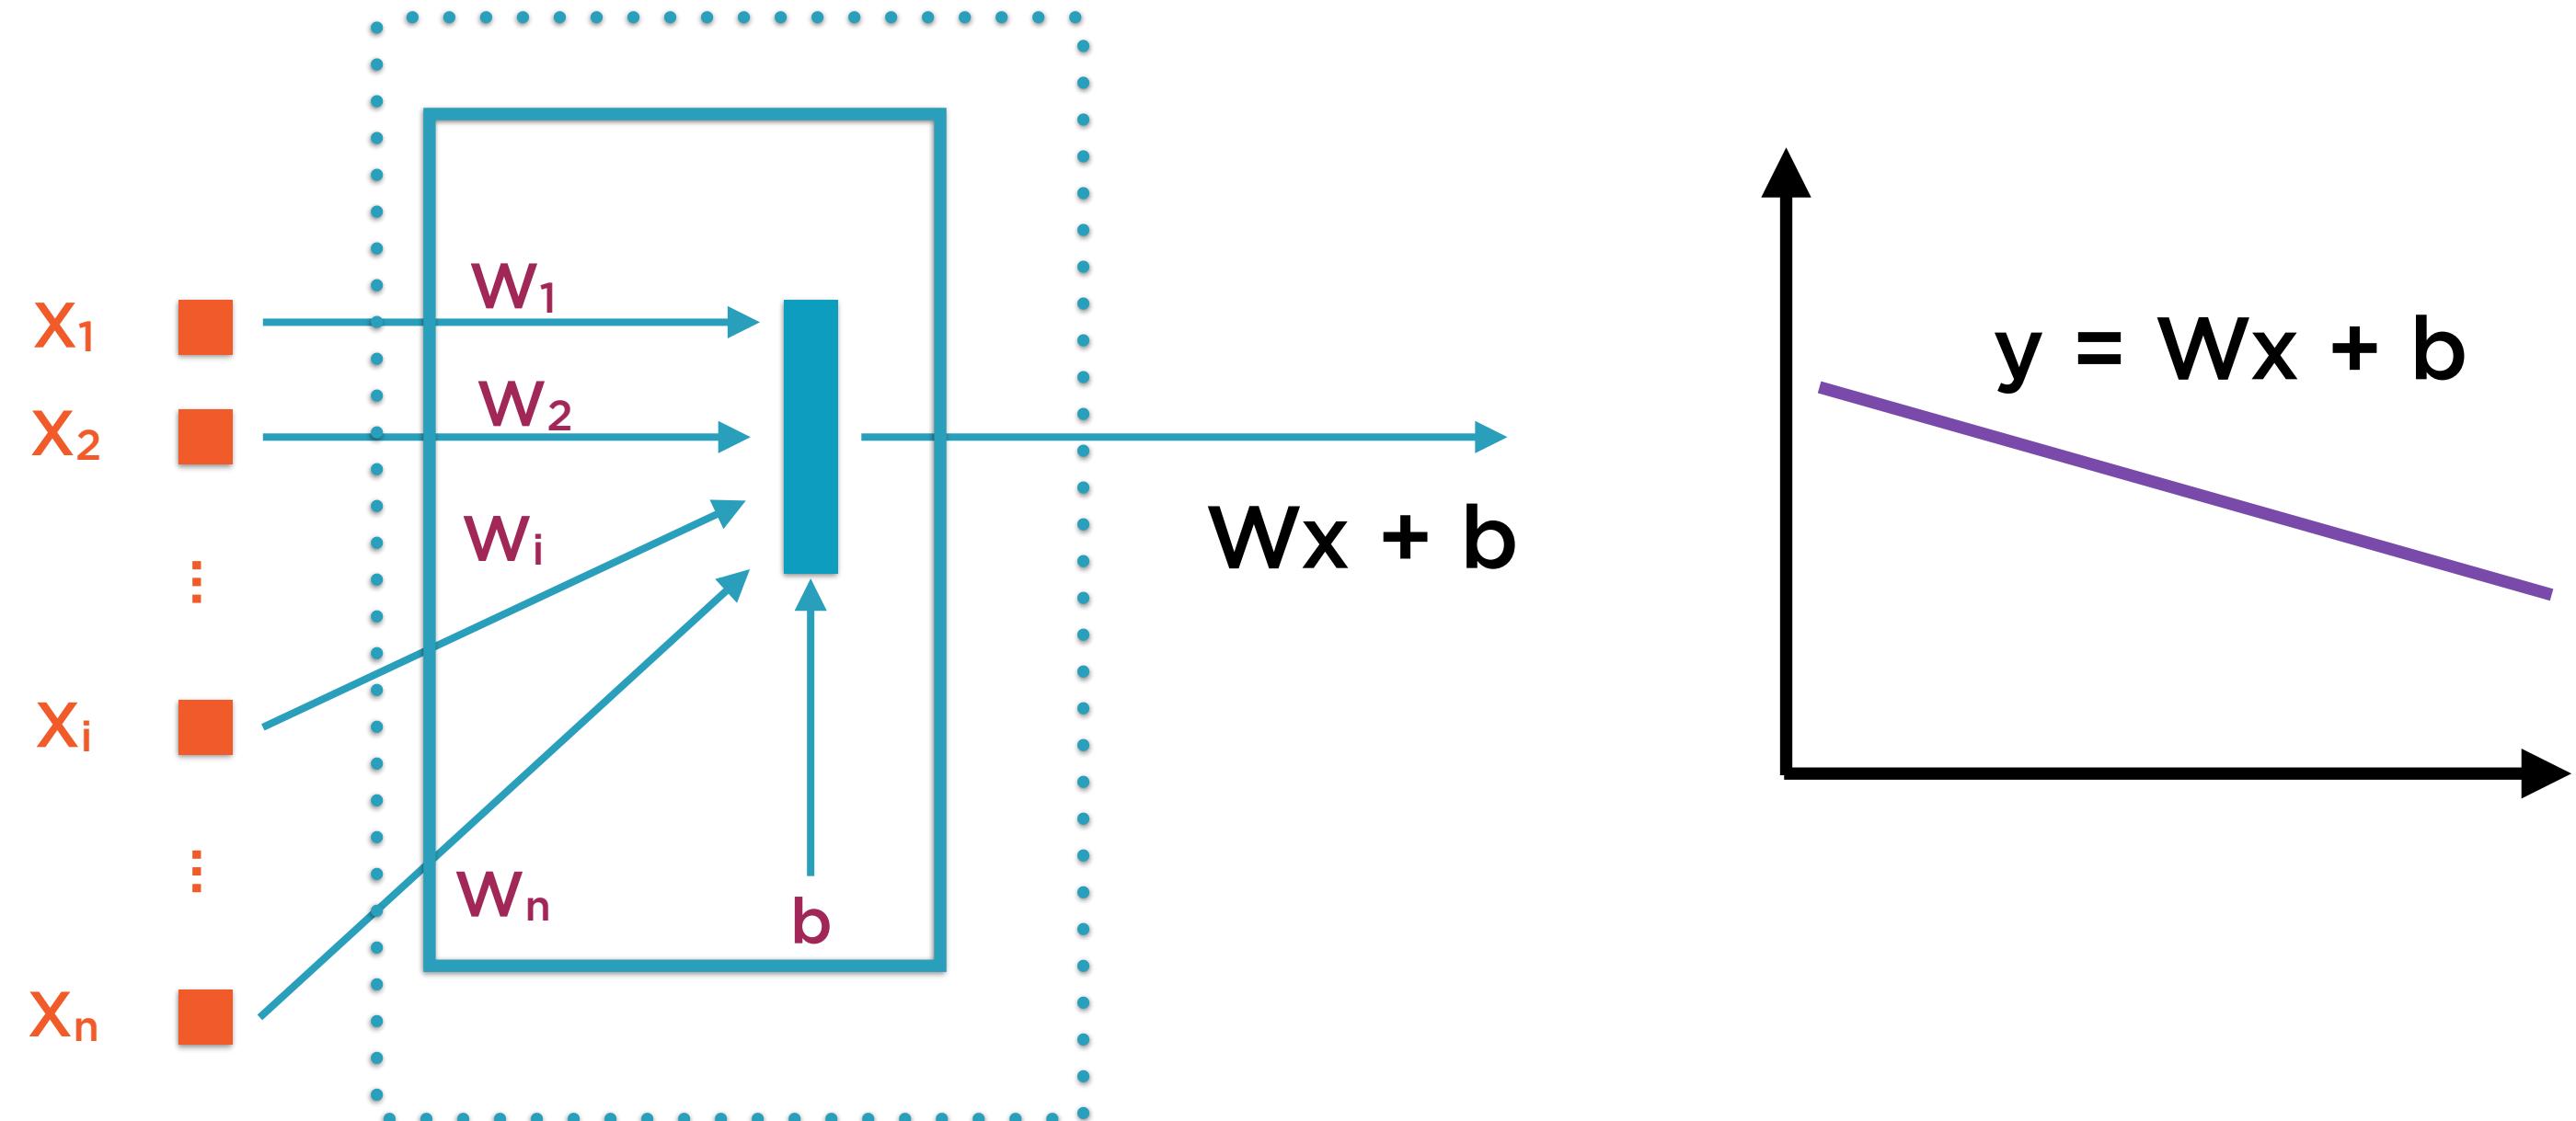
Regression: The Simplest Neural Network



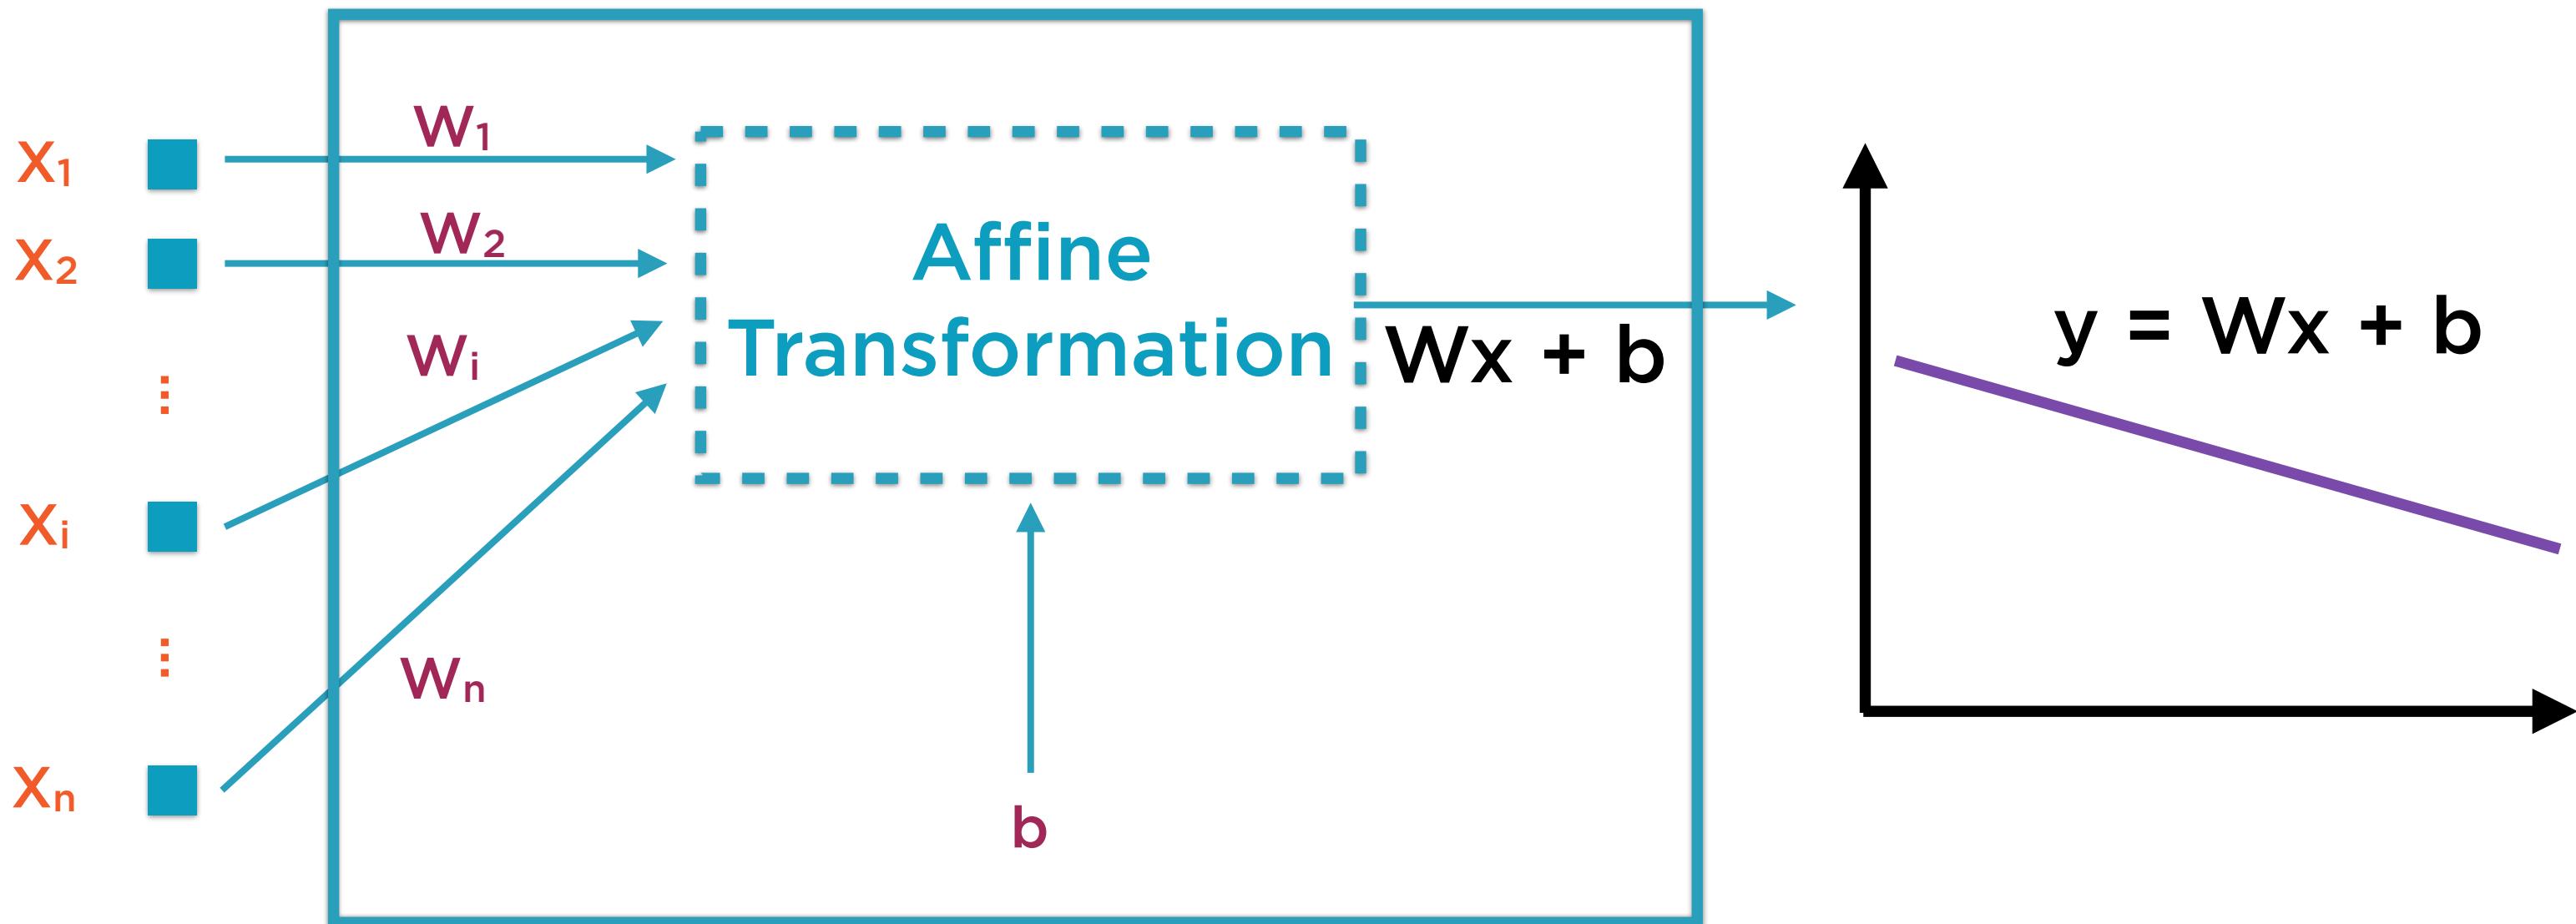
Regression: The Simplest Neural Network



Operation of a Single Neuron

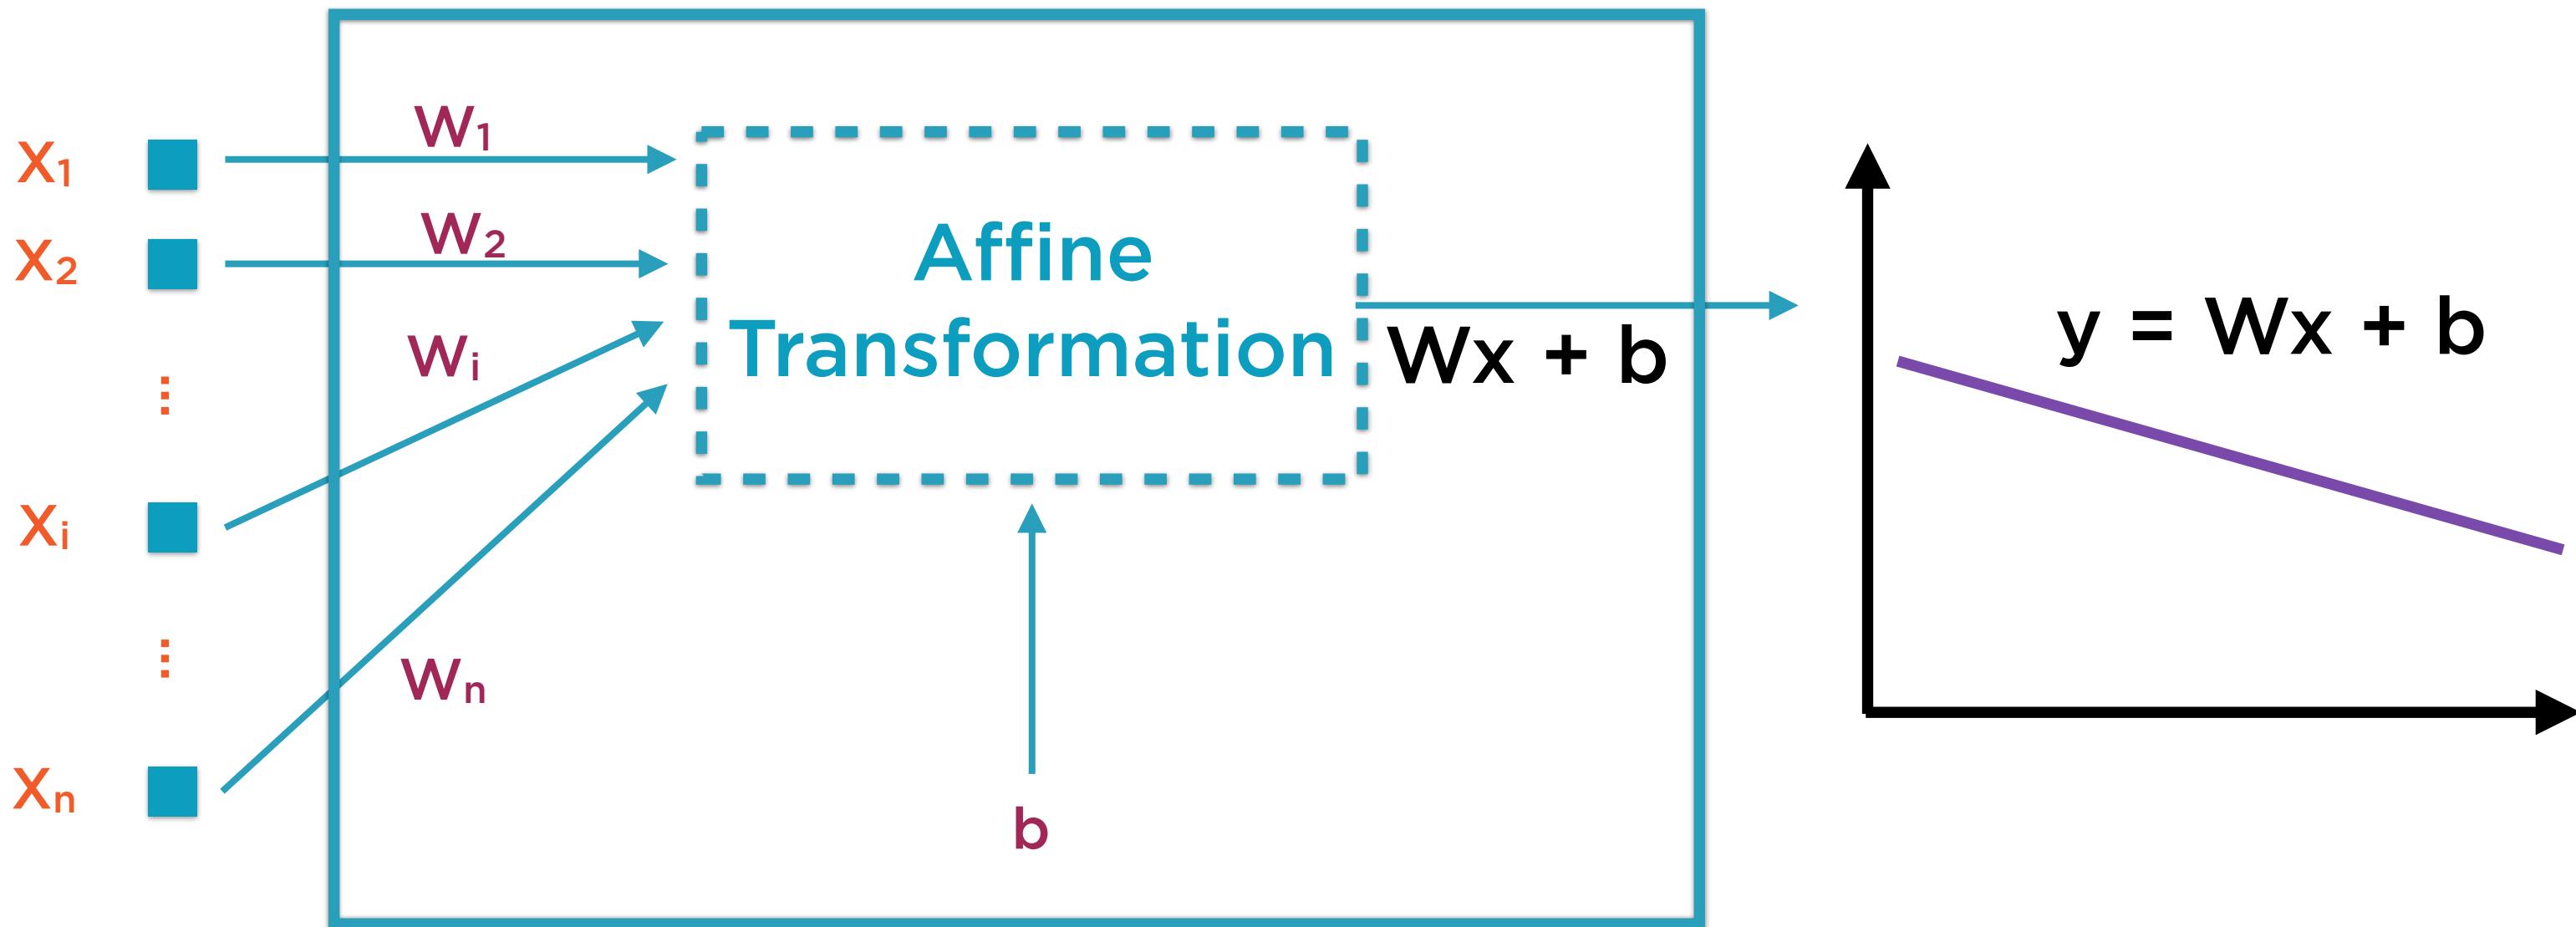


Operation of a Single Neuron



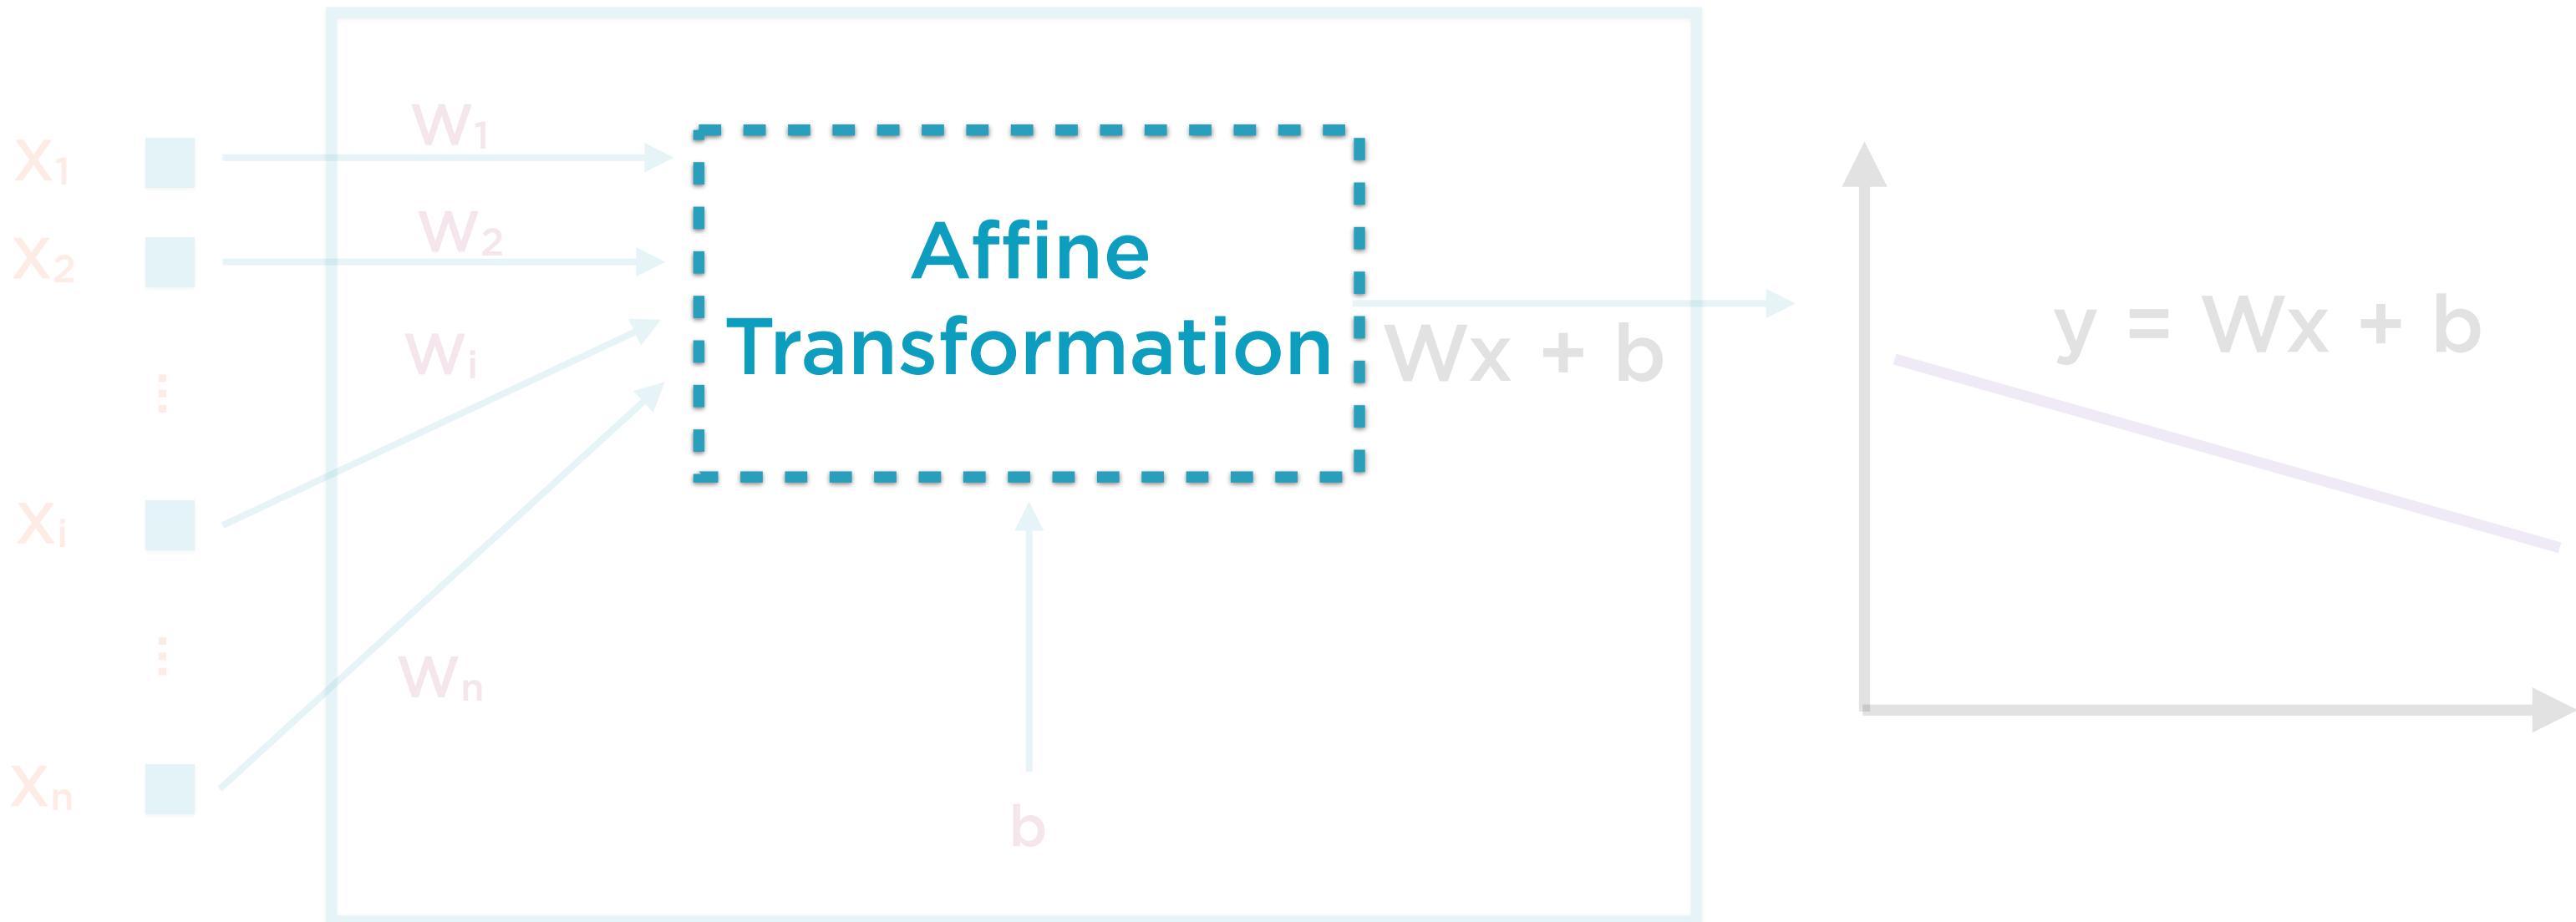
Here the neuron is an entity that finds the “best fit” line through a set of points

Operation of a Single Neuron



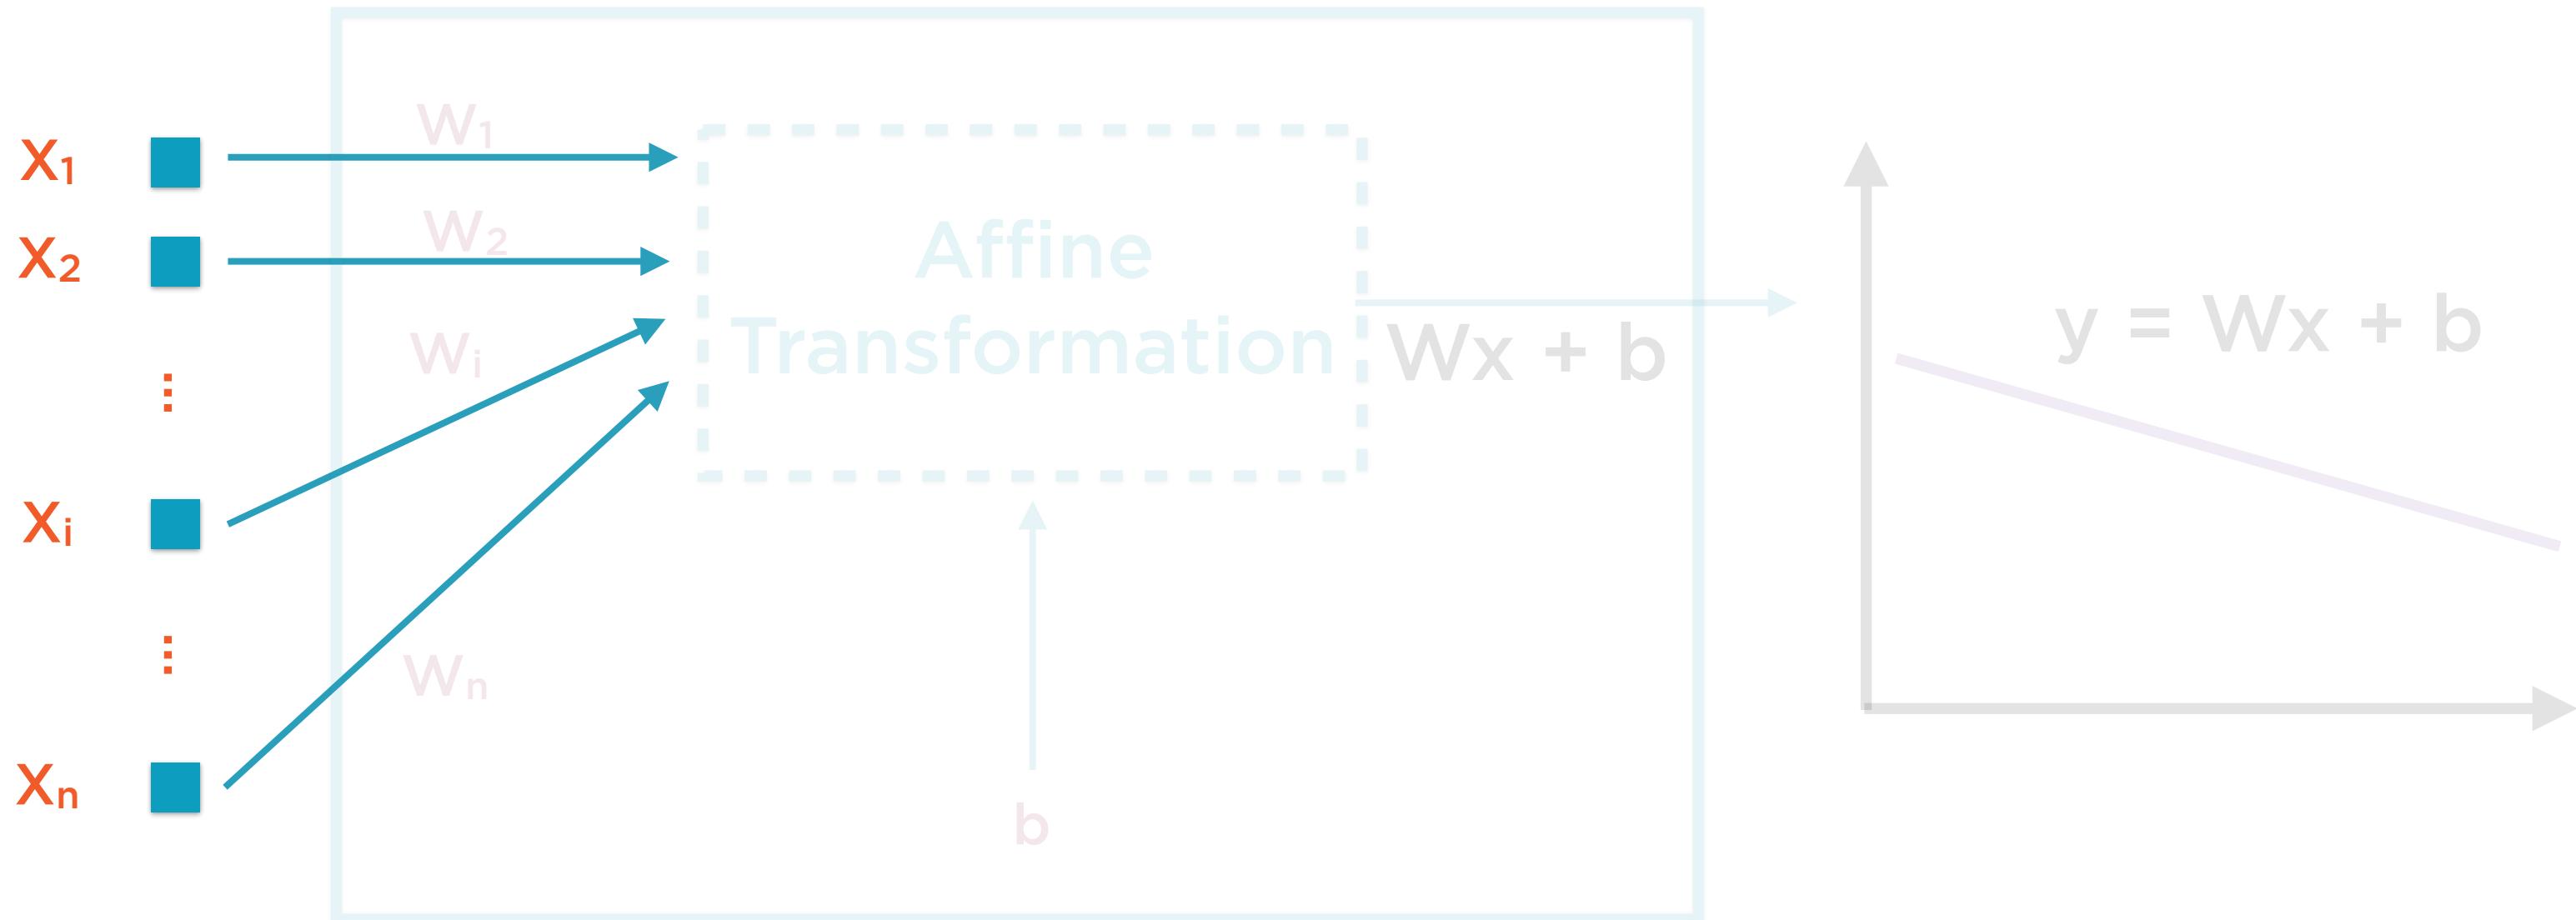
We are instructing the neuron to learn a linear function
- so no activation function is required at all

Operation of a Single Neuron



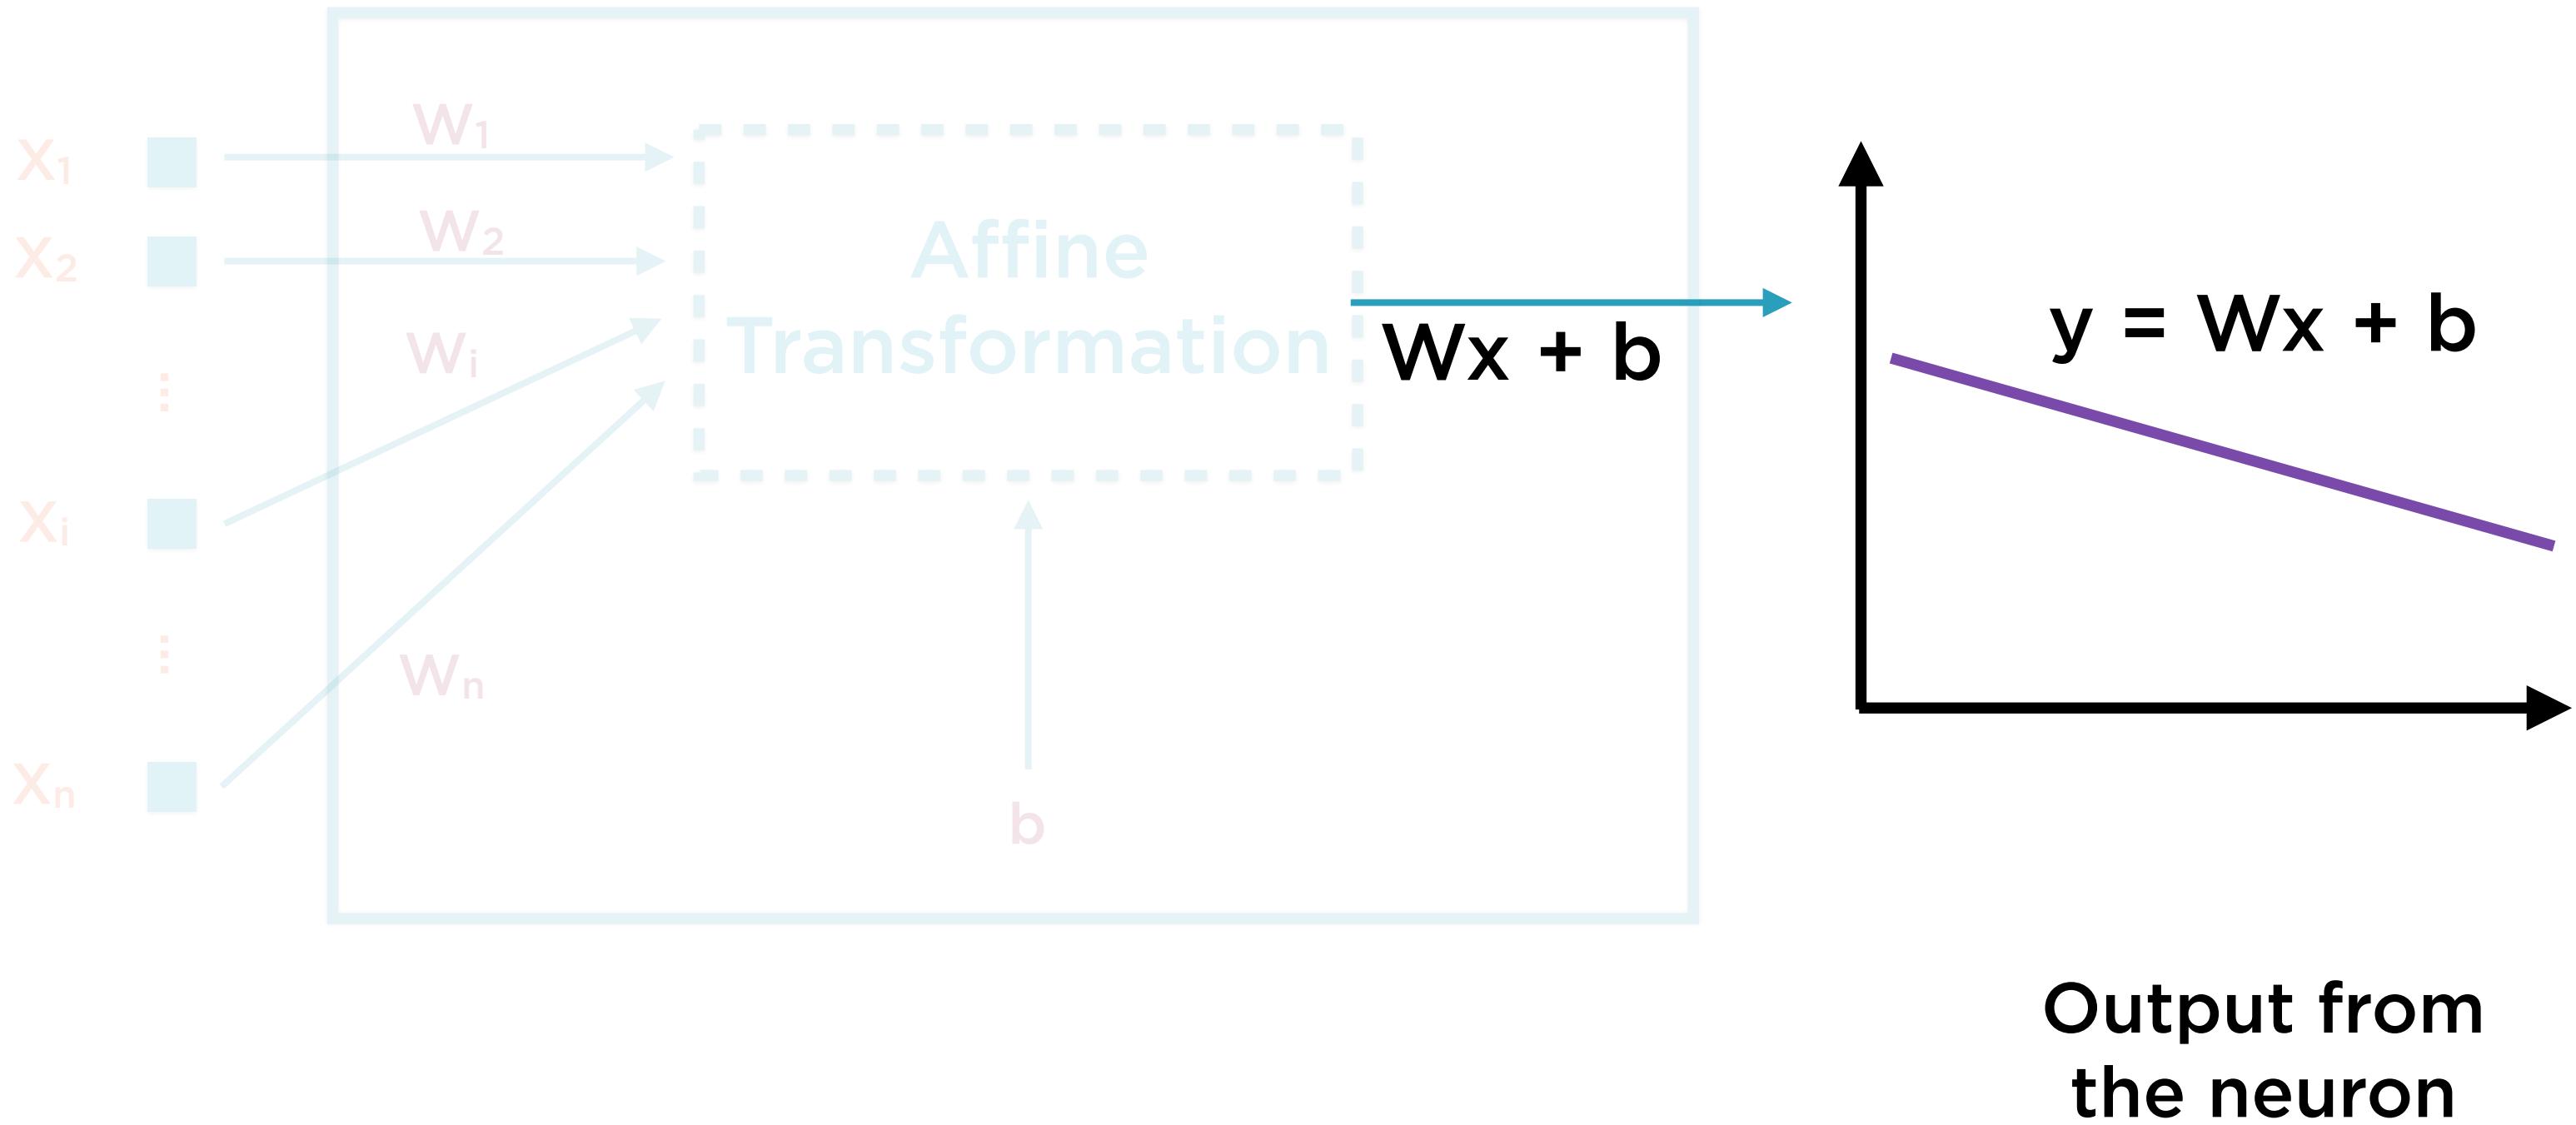
The affine transformation is just a weighted sum of the inputs with a bias added

Operation of a Single Neuron

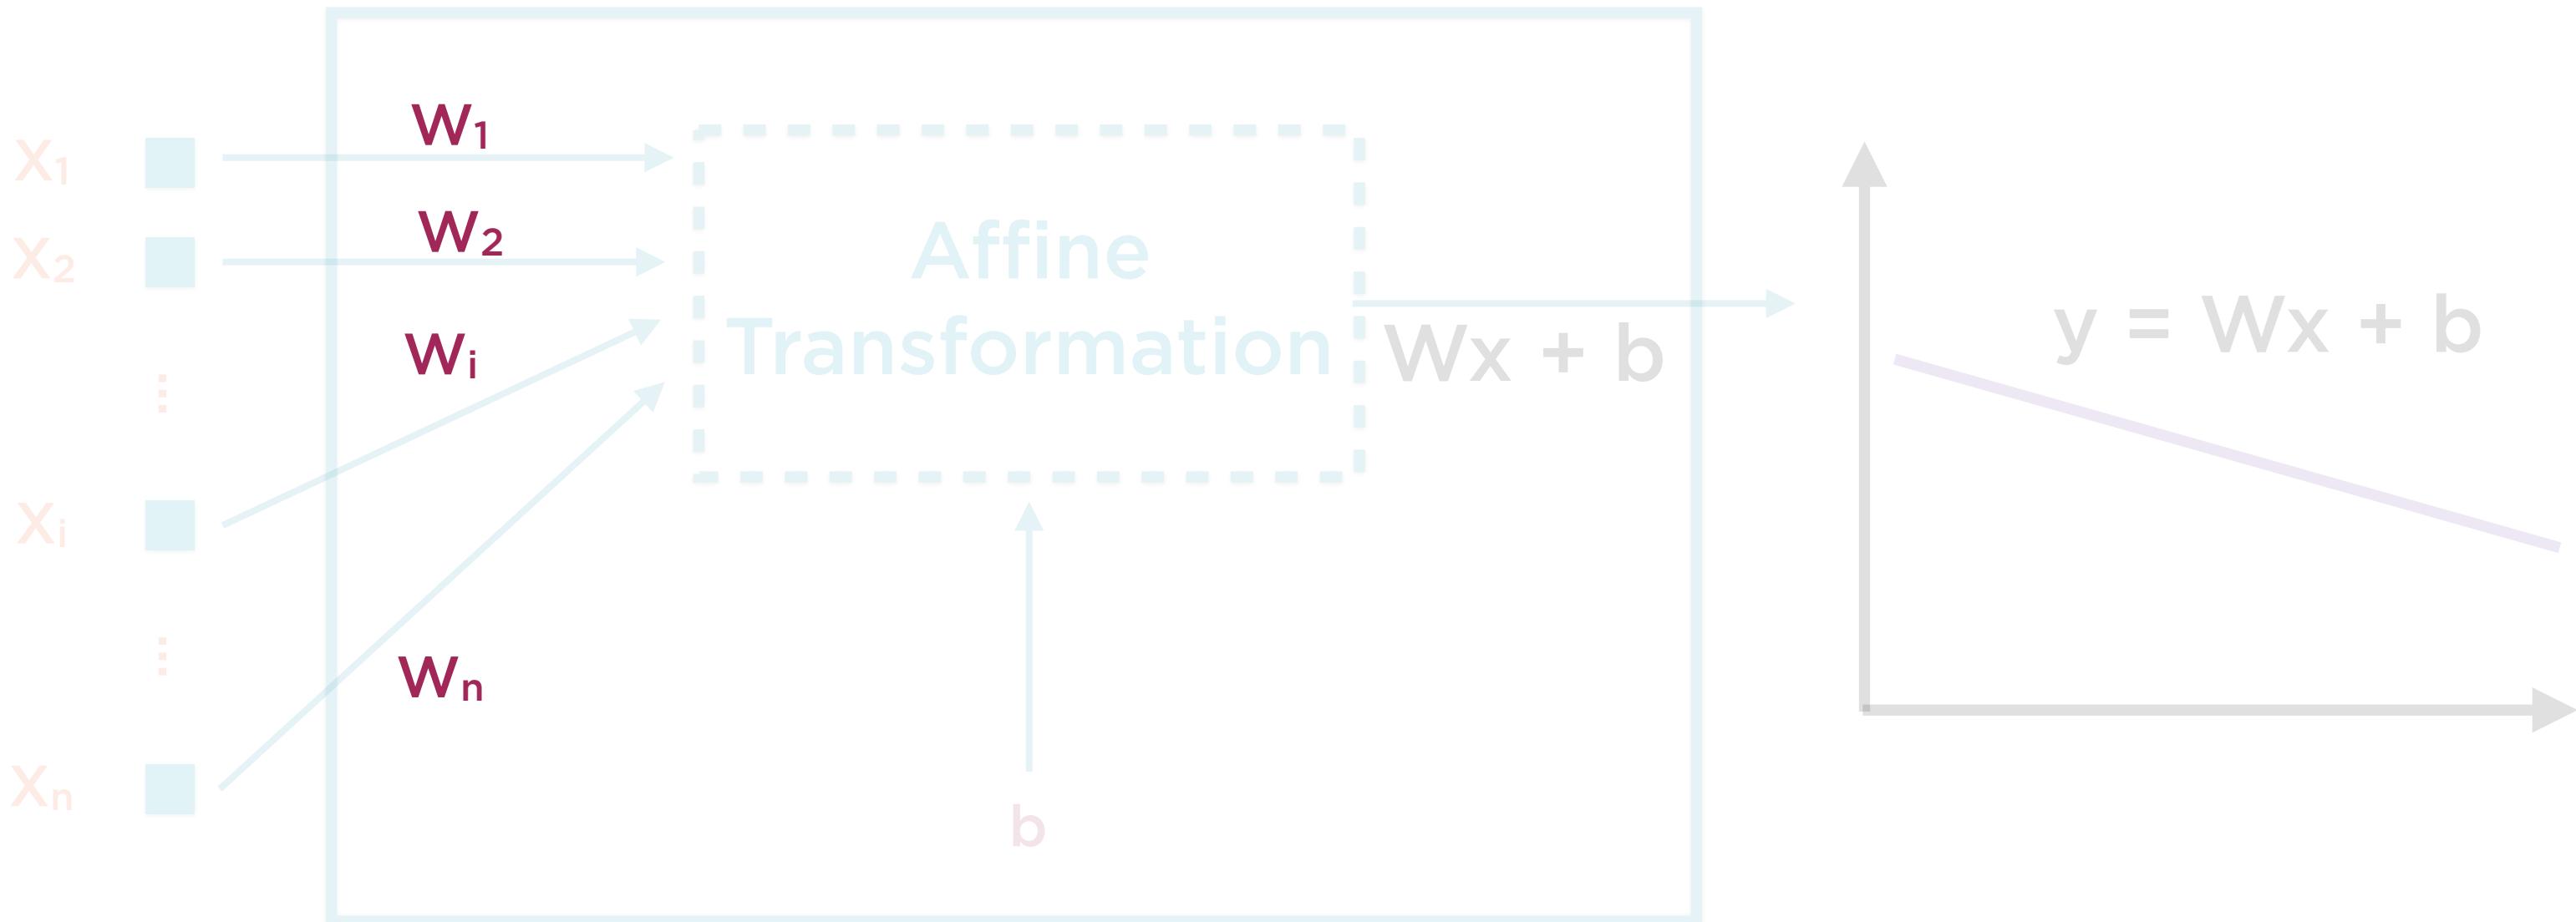


**Inputs into
the neuron**

Operation of a Single Neuron

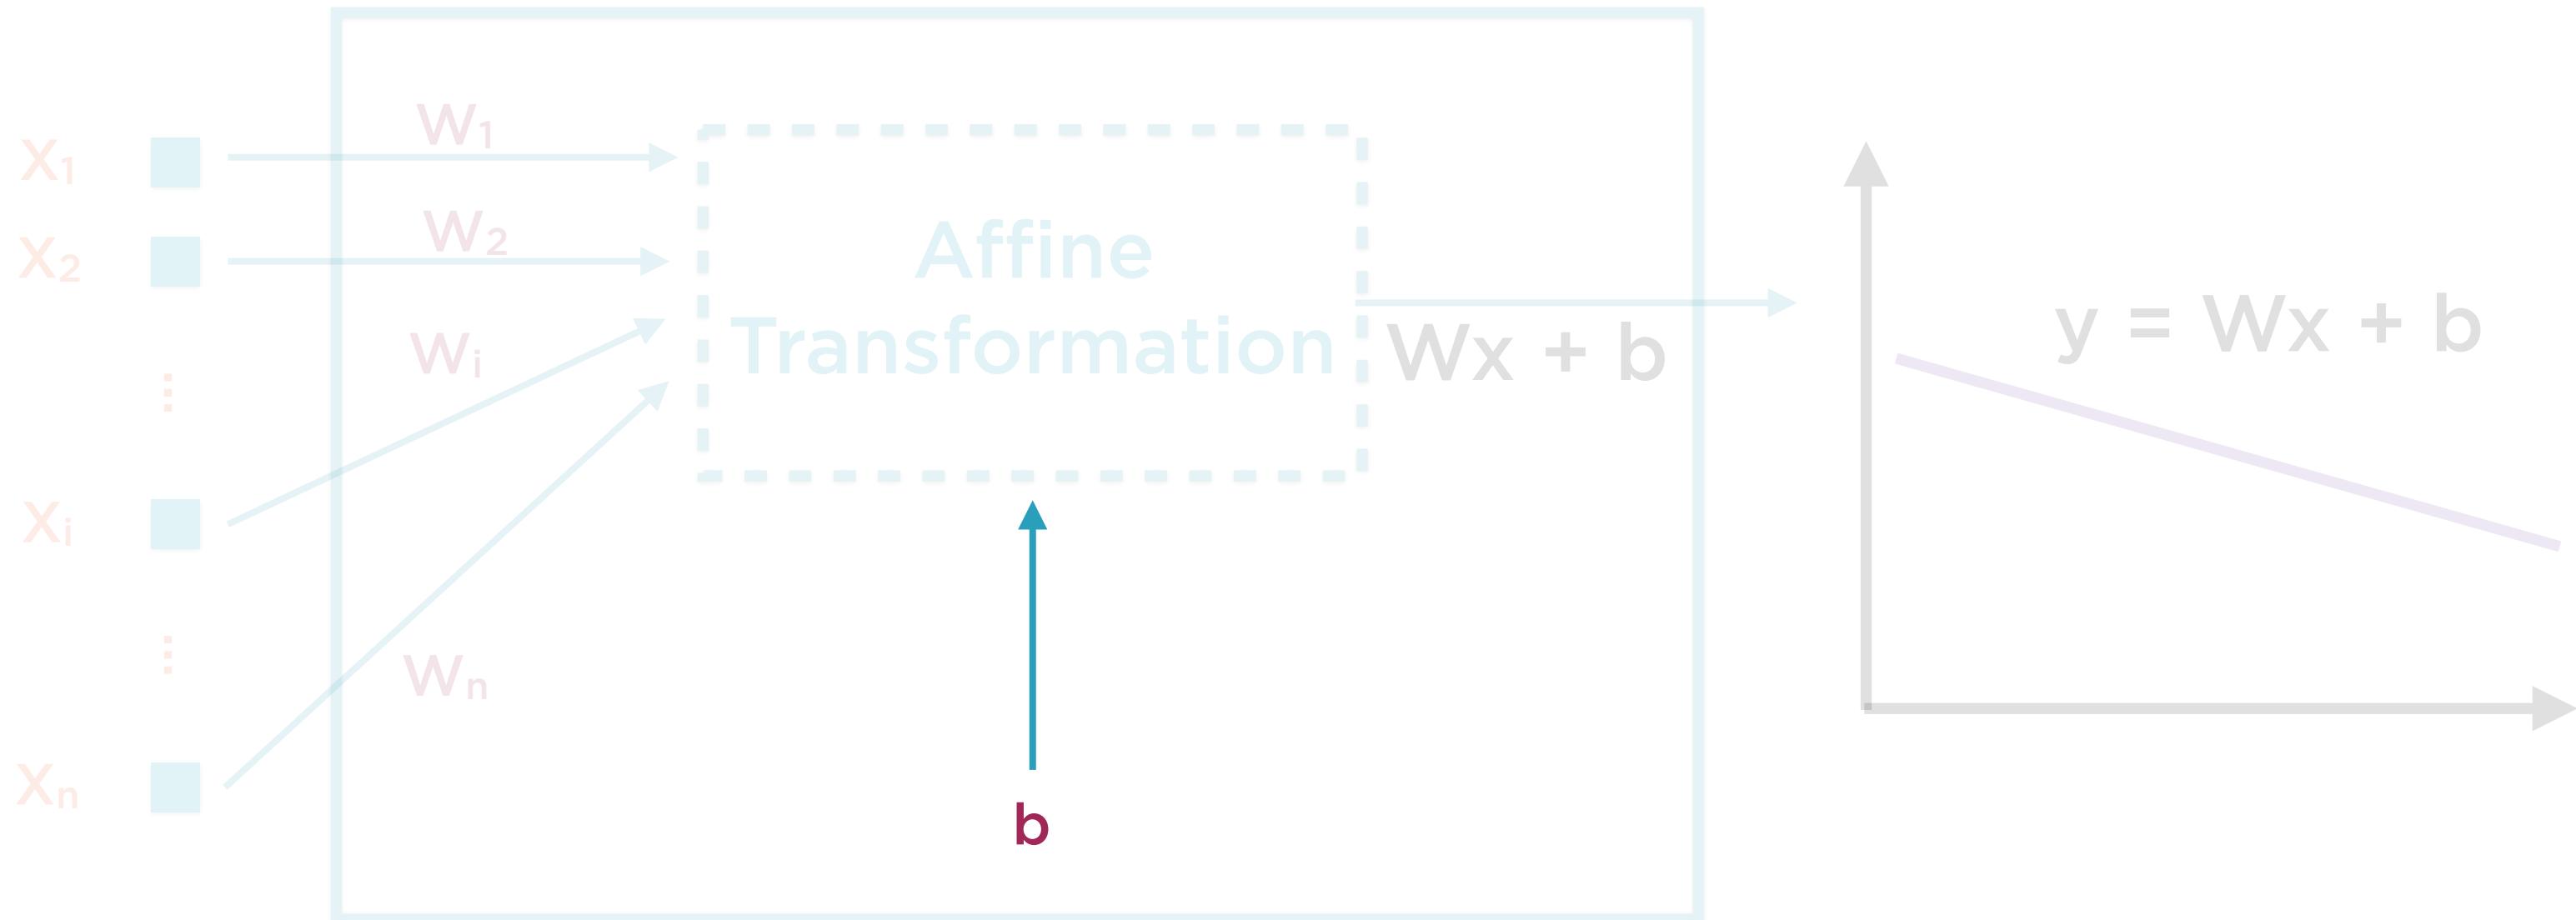


Operation of a Single Neuron



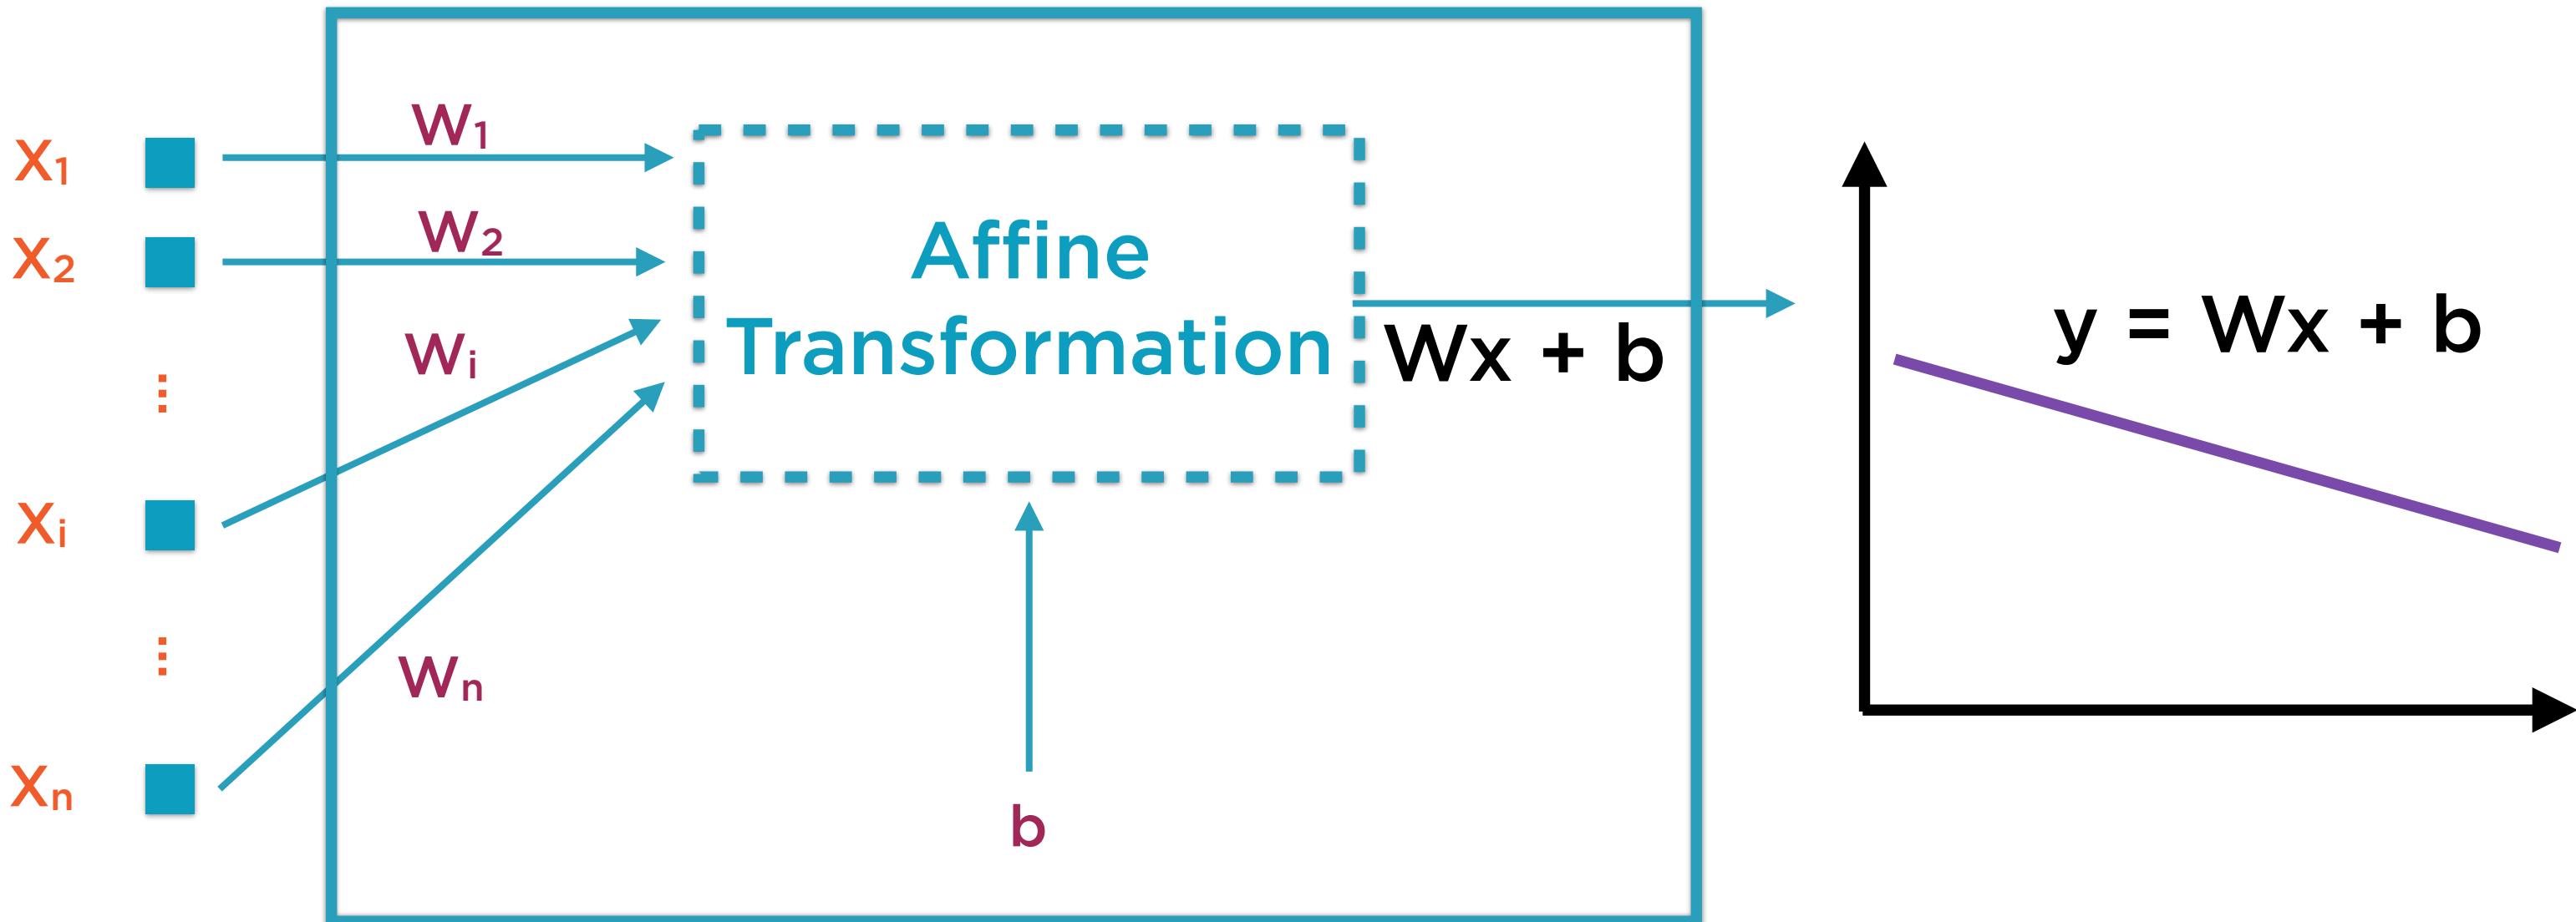
The values $W_1, W_2\dots W_n$ are called the weights

Operation of a Single Neuron



The value b is called the bias

Operation of a Single Neuron



Where do the weights W and the bias b come from?

The weights and biases of individual neurons are determined during the training process

The actual training of a neural network is managed by TensorFlow

Simple Regression

Regression Equation:

$$y = A + Bx$$

$$y_1 = A + Bx_1$$

$$y_2 = A + Bx_2$$

$$y_3 = A + Bx_3$$

...

...

$$y_n = A + Bx_n$$

Simple Regression

Regression Equation:

$$y = A + Bx$$

$$y_1 = A + Bx_1 + e_1$$

$$y_2 = A + Bx_2 + e_2$$

$$y_3 = A + Bx_3 + e_3$$

...

...

$$y_n = A + Bx_n + e_n$$

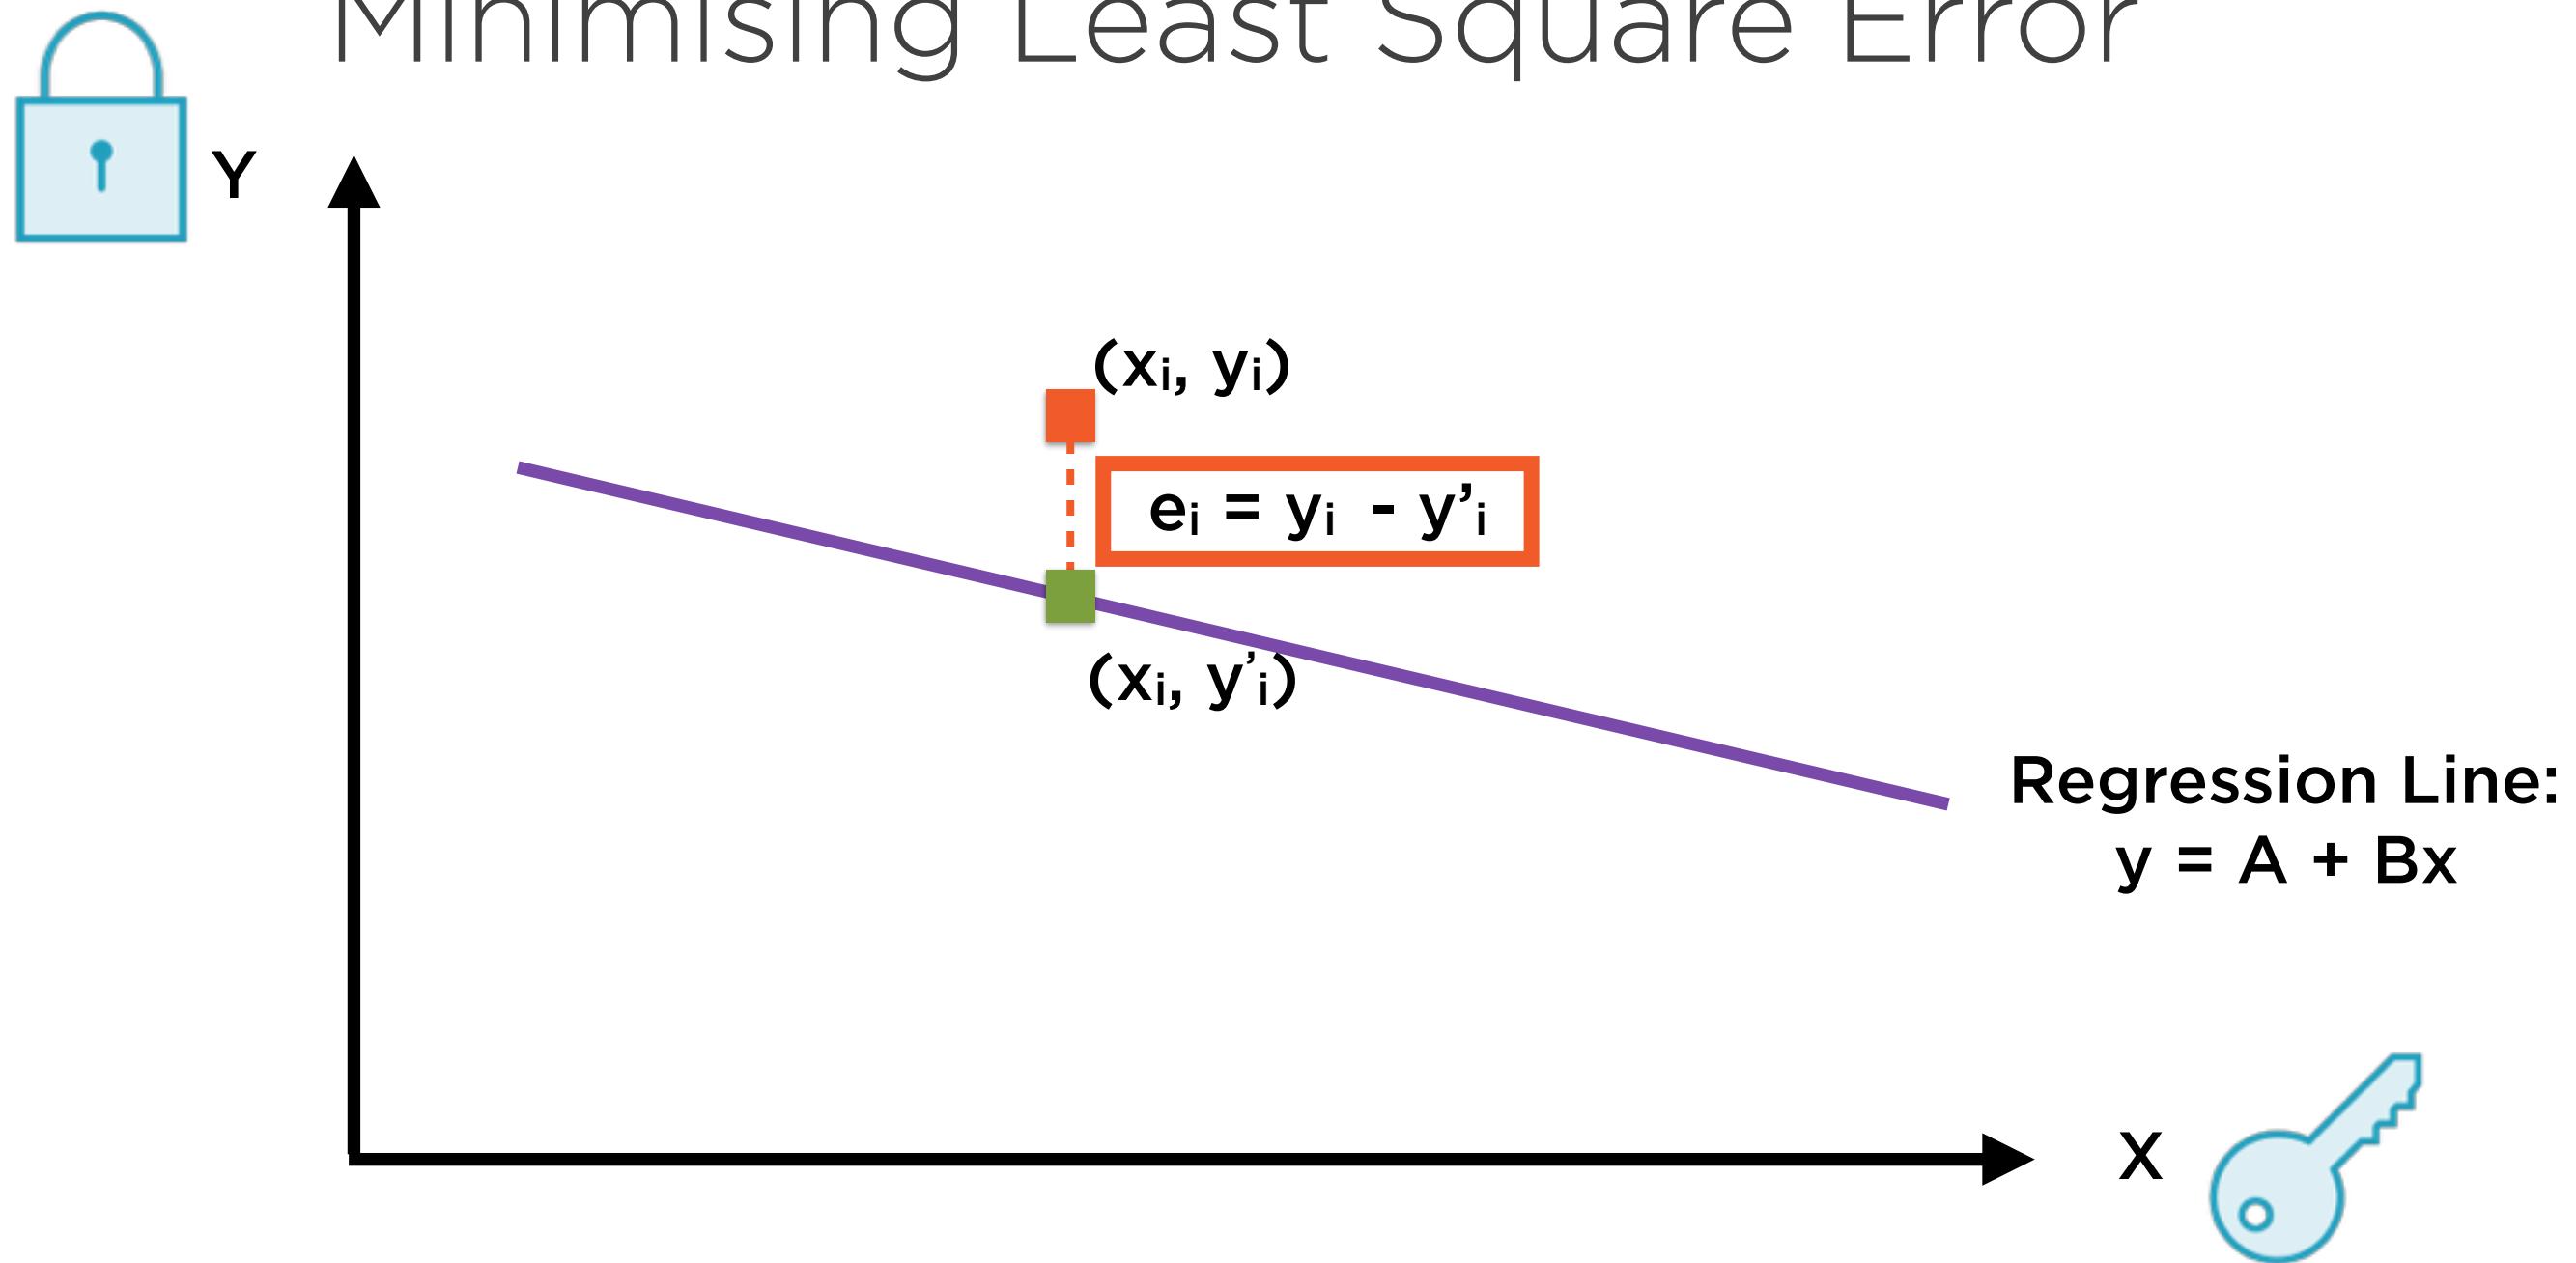
Simple Regression

Regression Equation:

$$y = A + Bx$$

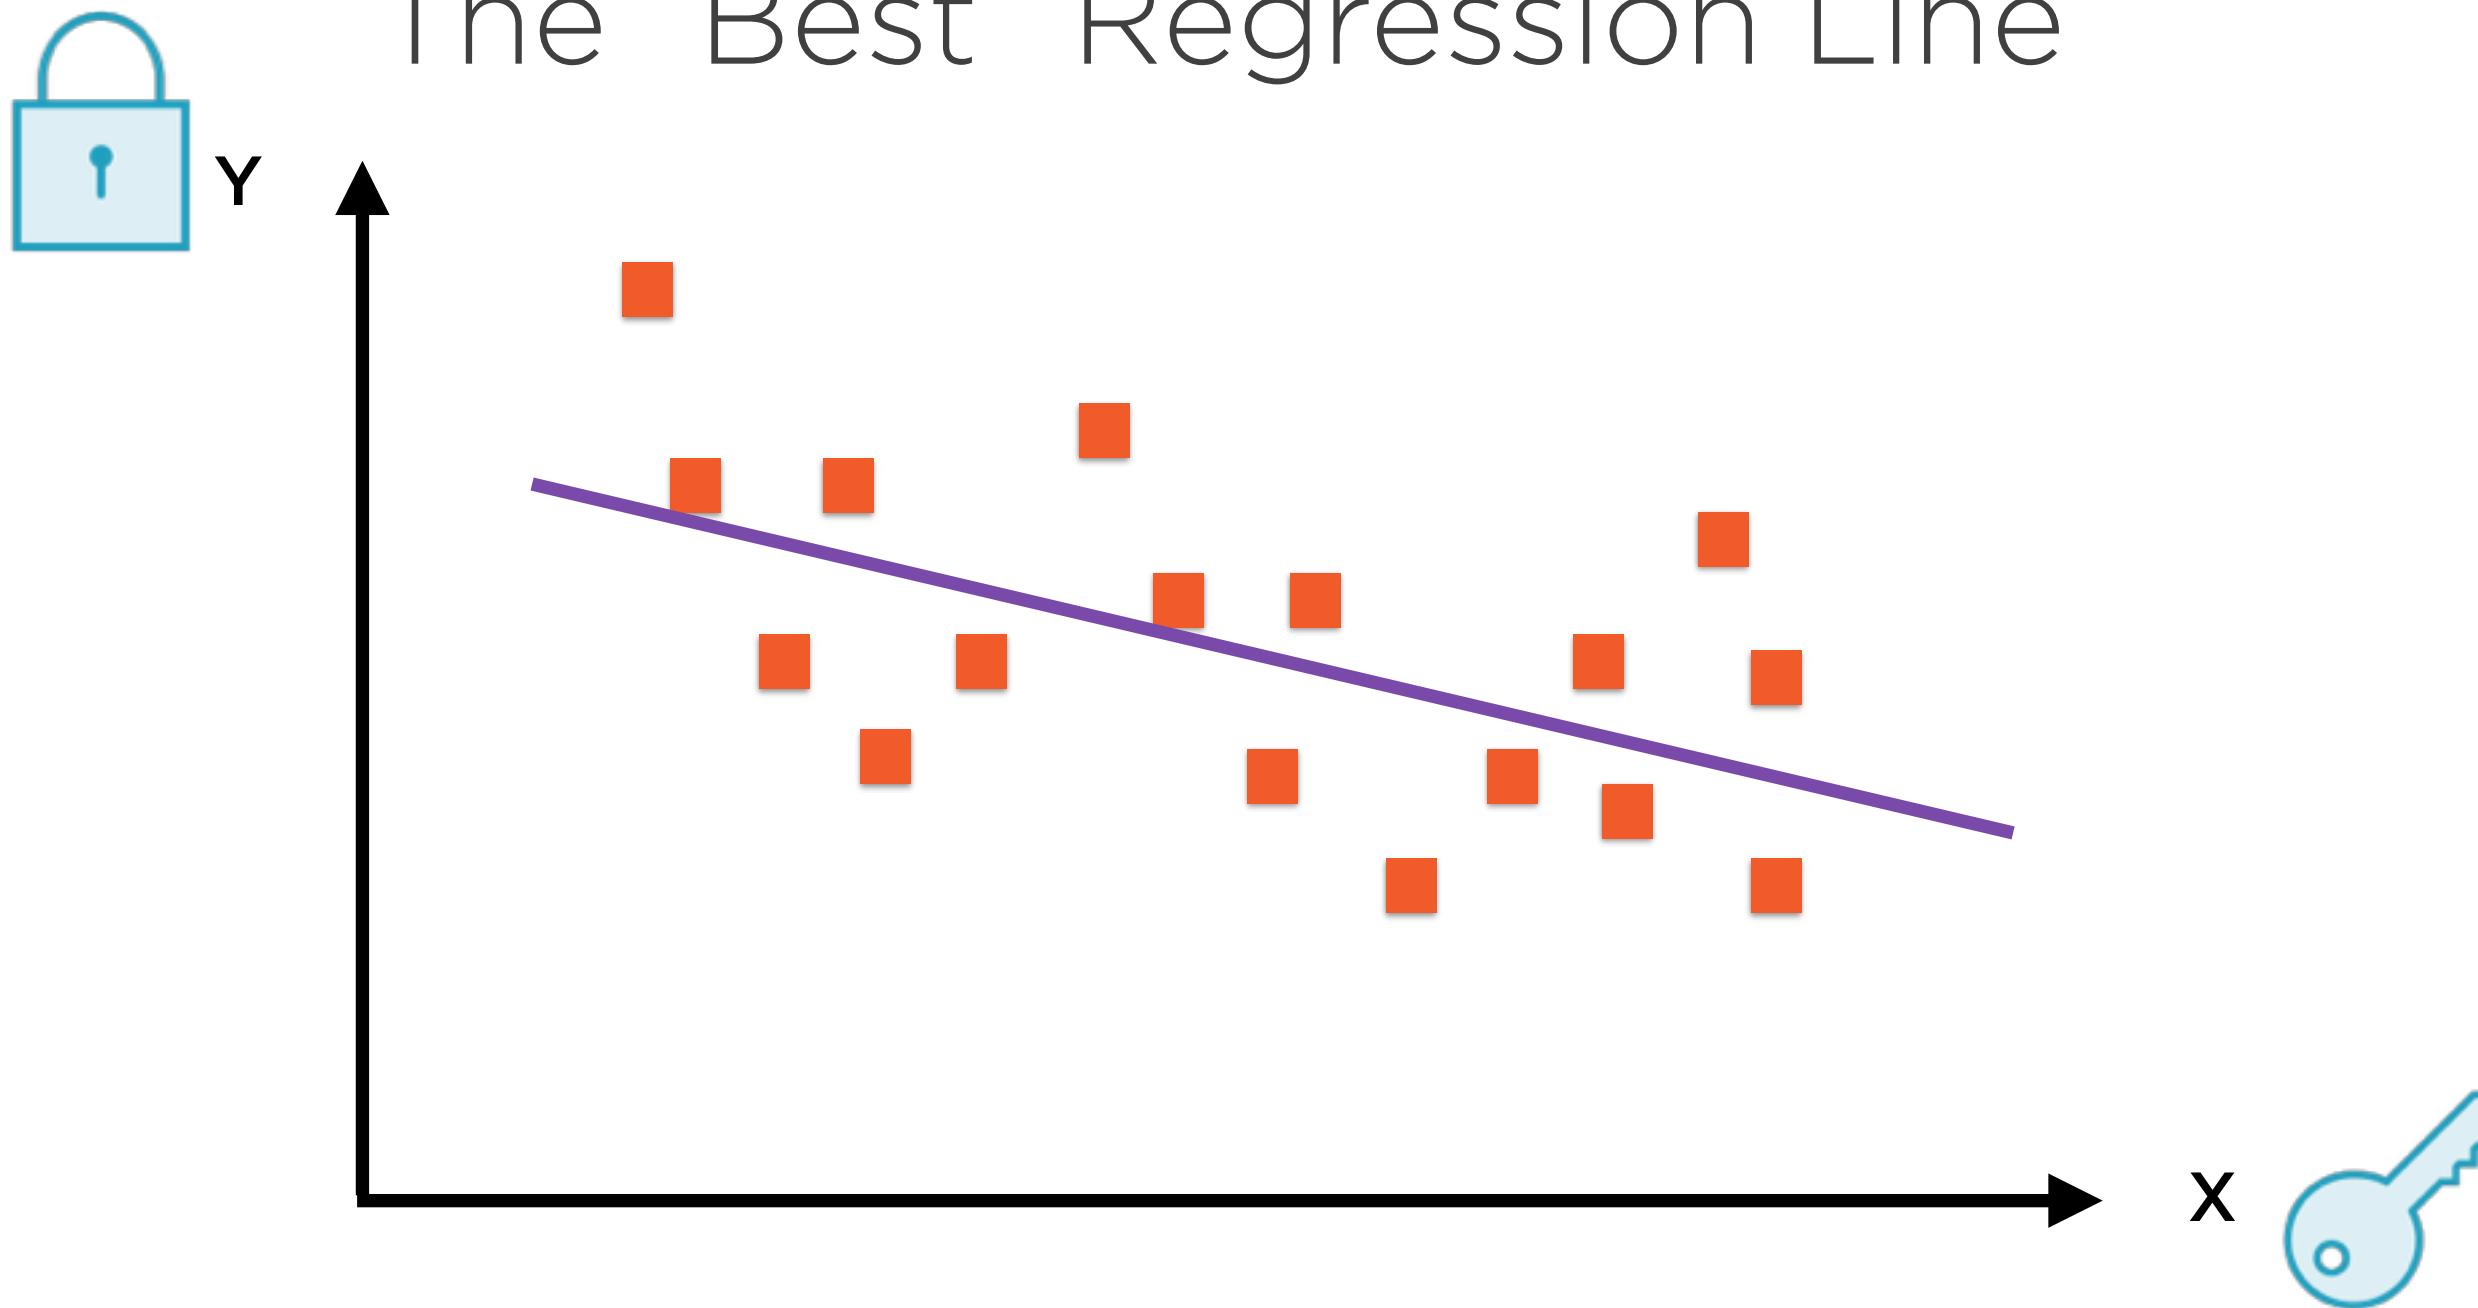
$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} = A \begin{bmatrix} 1 \\ 1 \\ 1 \\ \dots \\ 1 \end{bmatrix} + B \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \dots \\ x_n \end{bmatrix} + \begin{bmatrix} e_1 \\ e_2 \\ e_3 \\ \dots \\ e_n \end{bmatrix}$$

Minimising Least Square Error



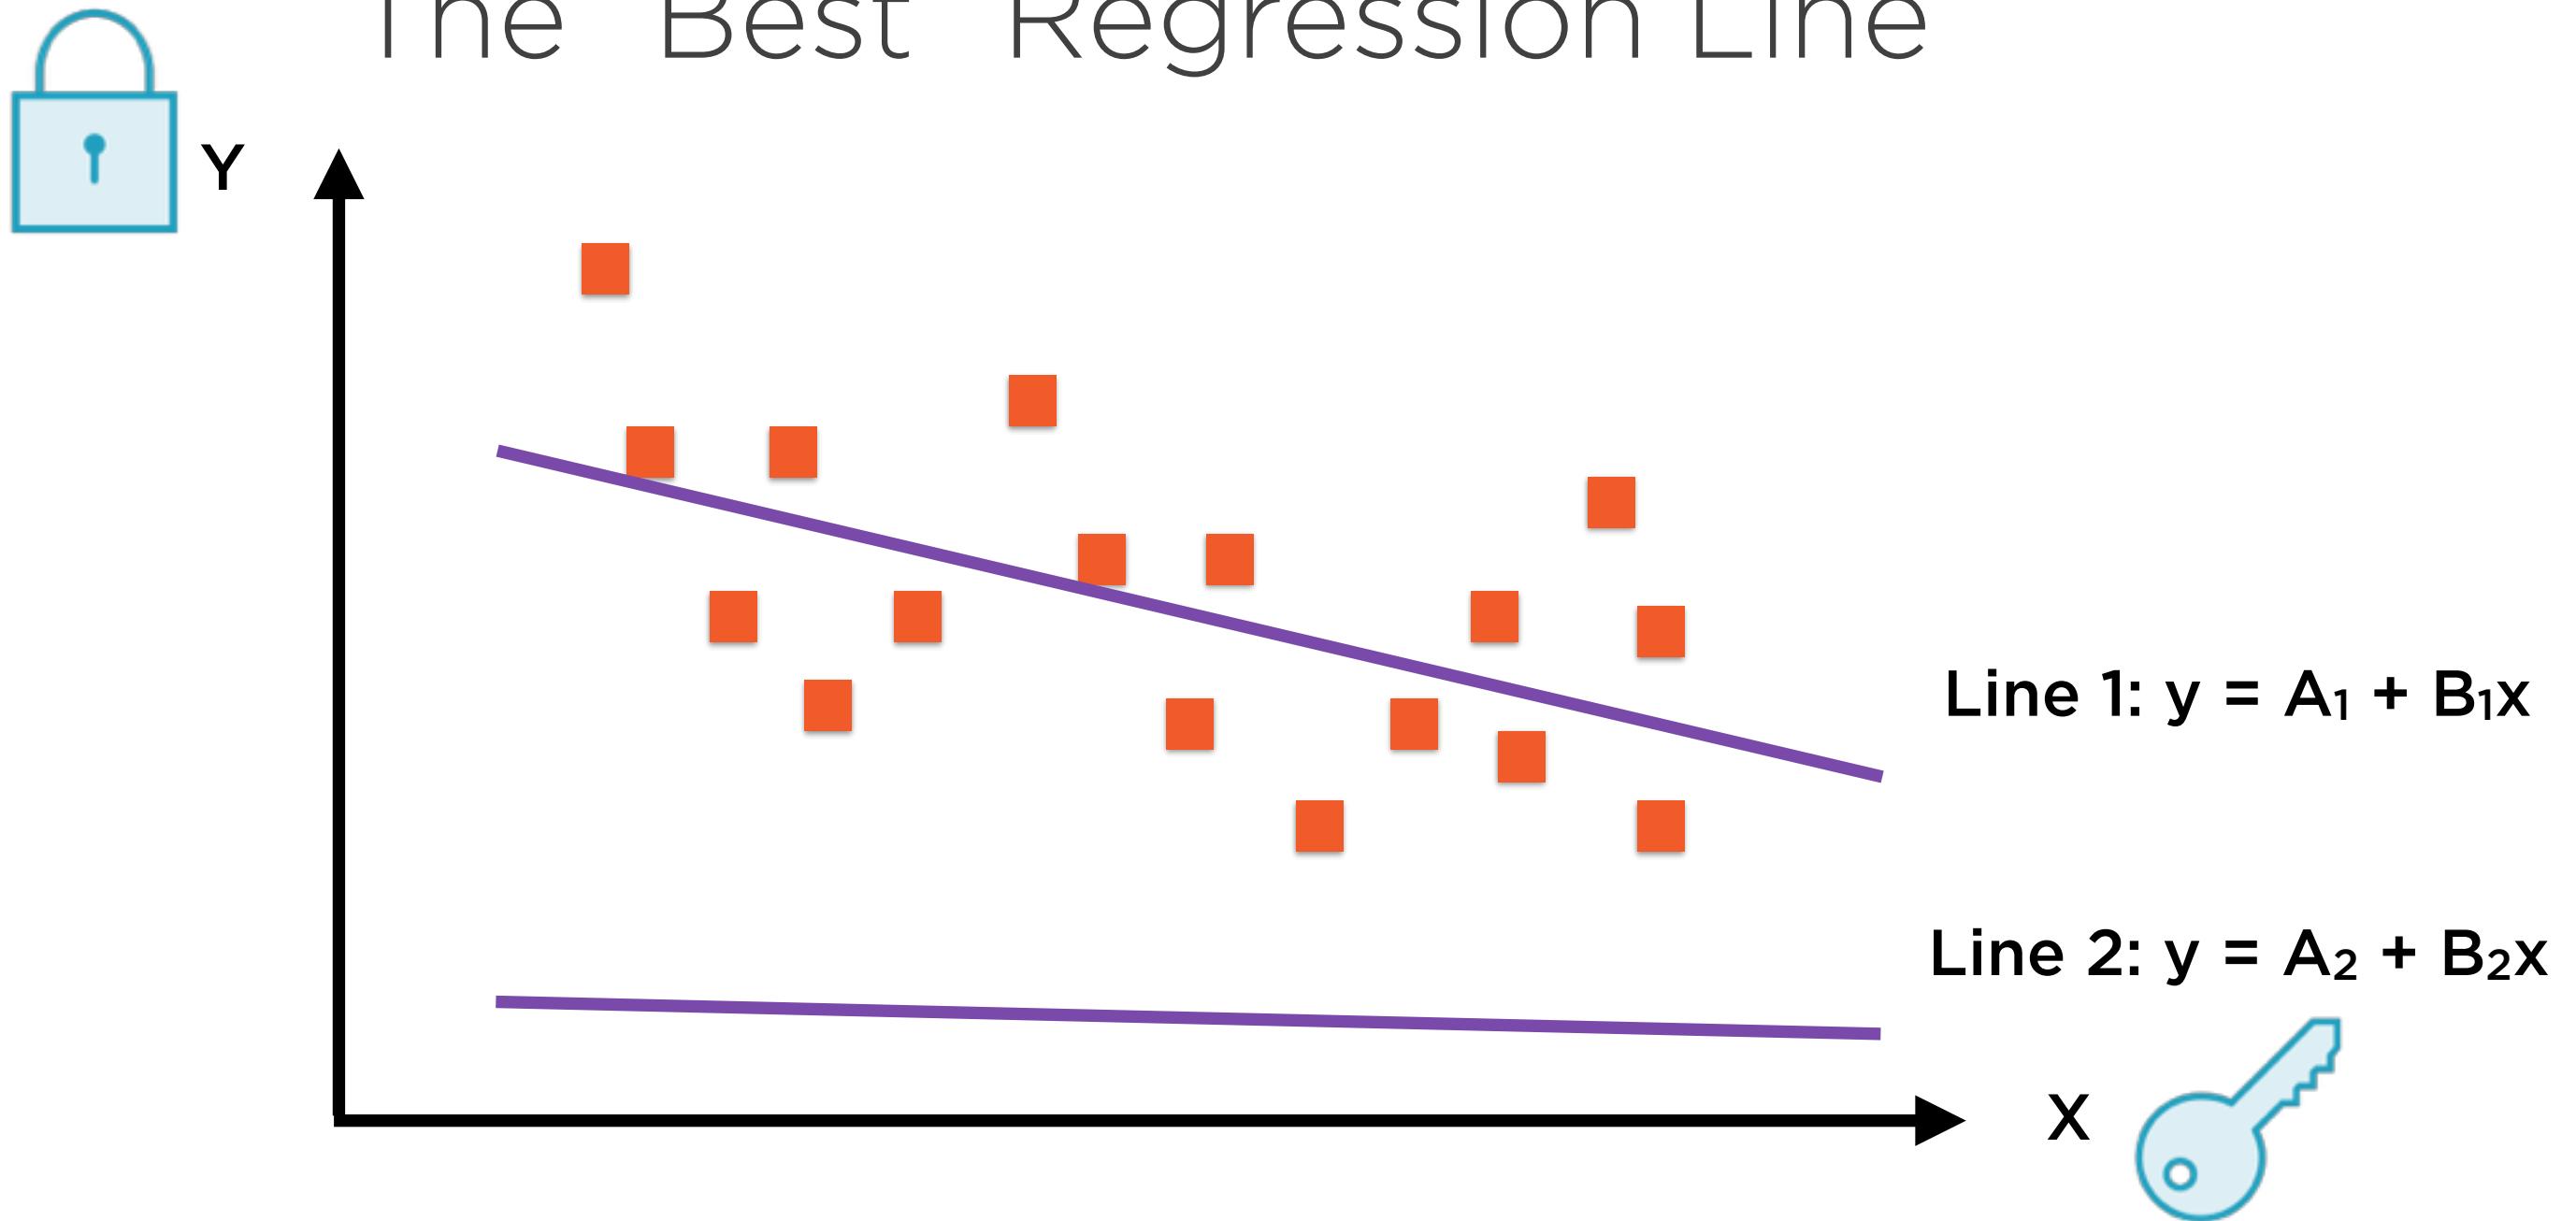
Residuals of a regression are the difference between actual and fitted values of the dependent variable

The “Best” Regression Line



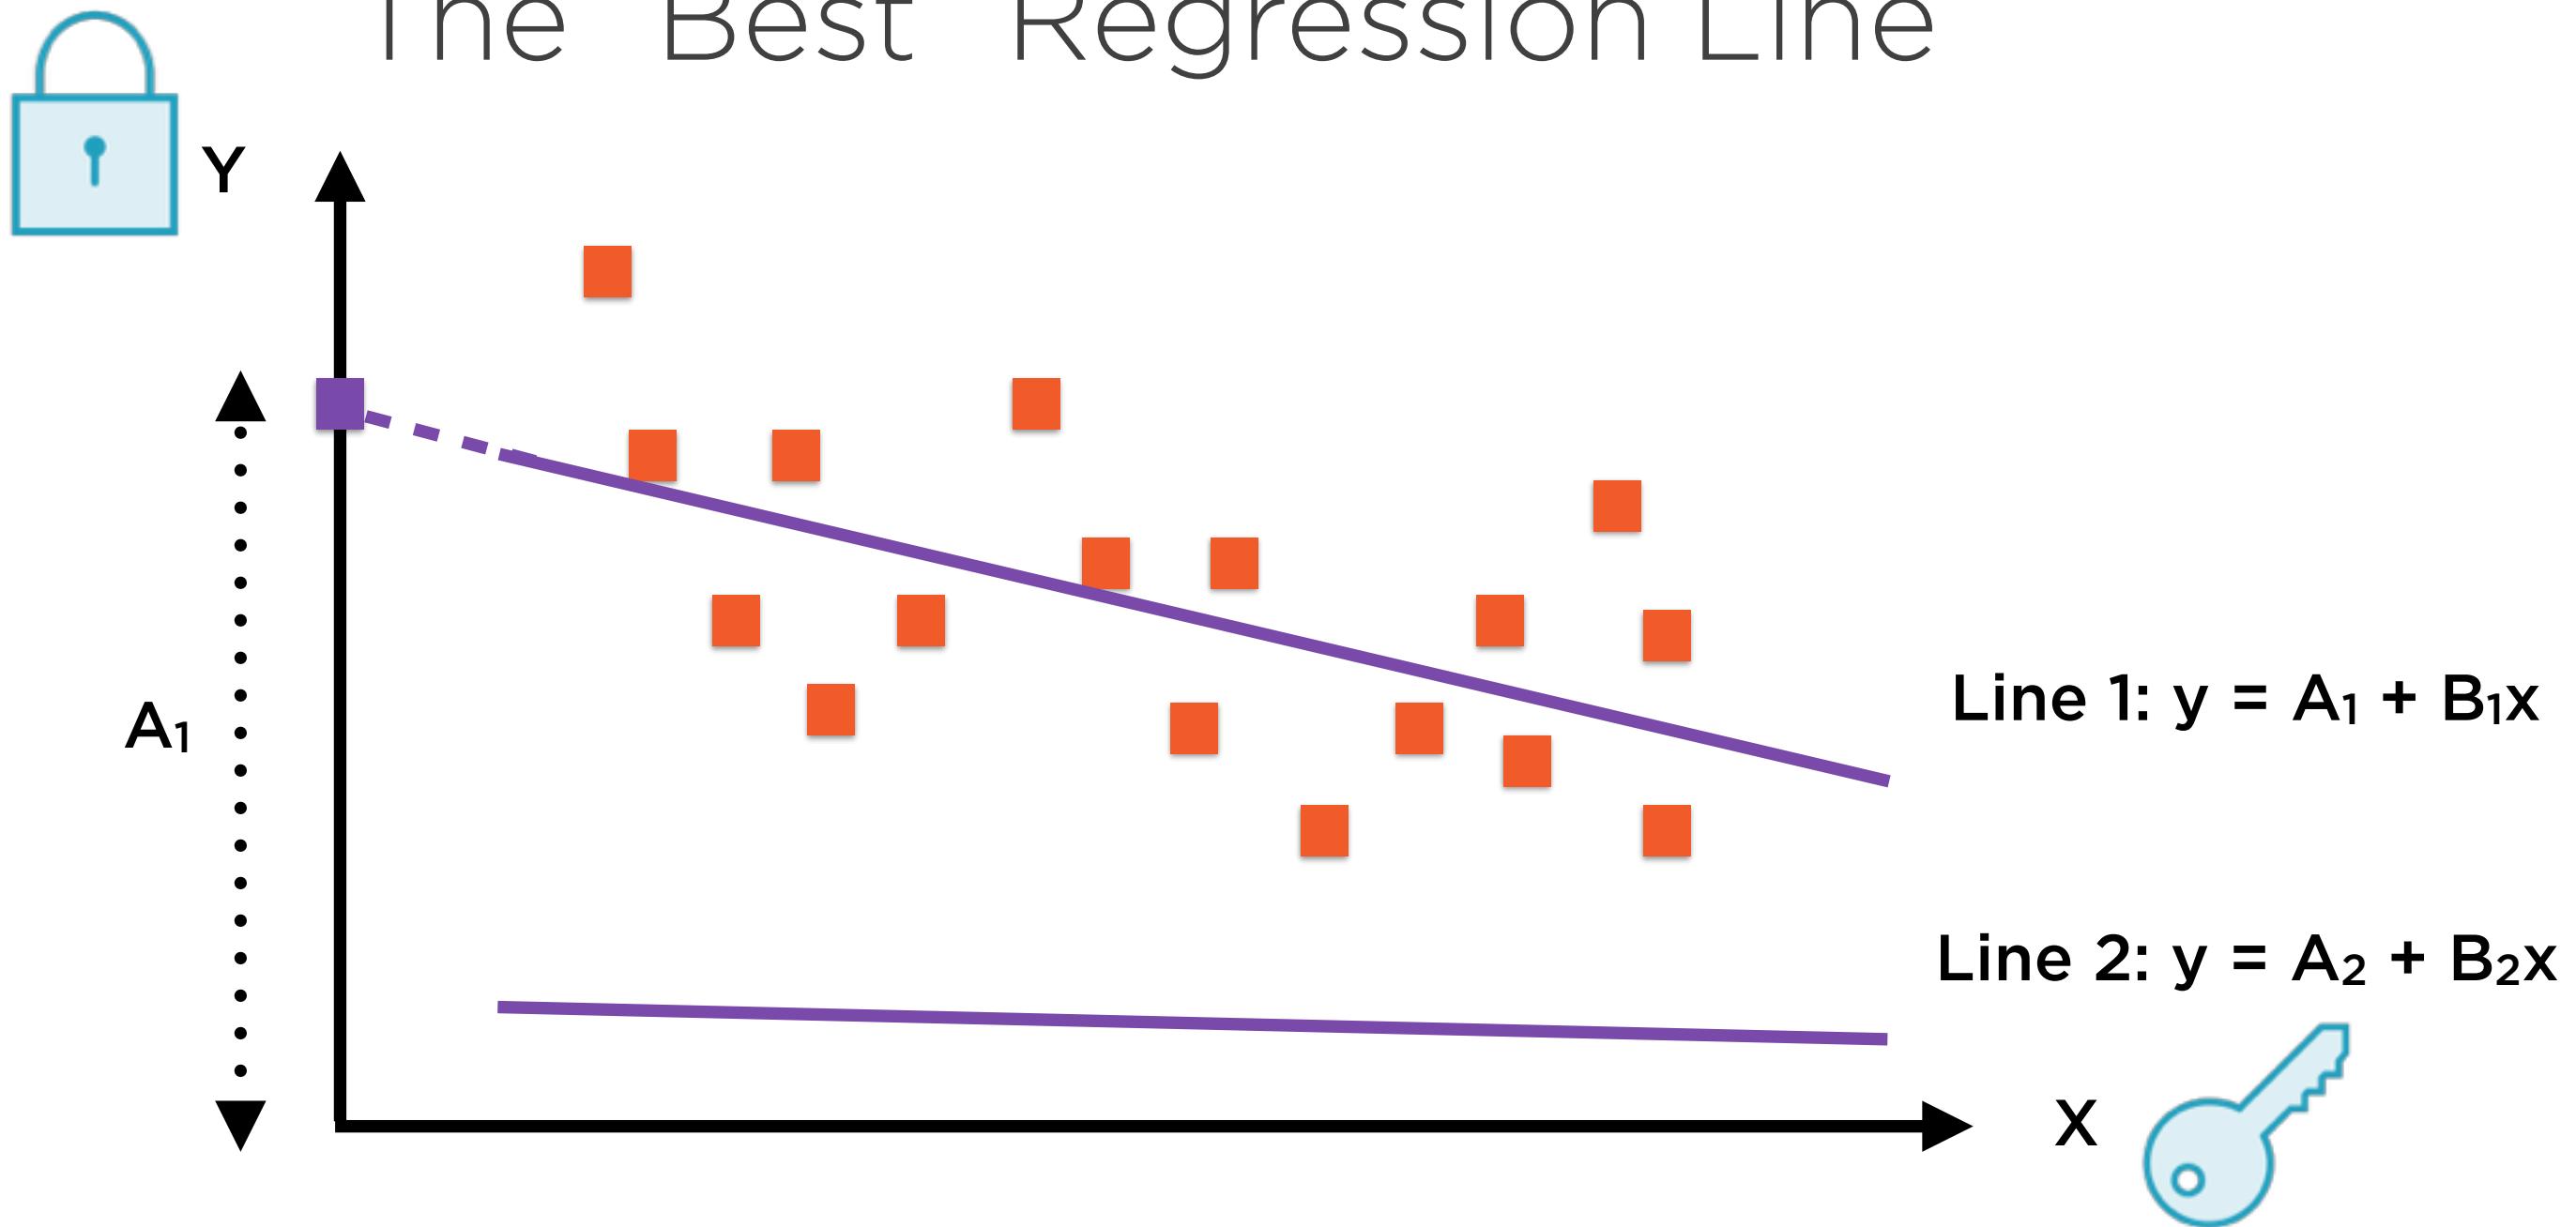
Linear Regression involves finding the “best fit” line

The “Best” Regression Line

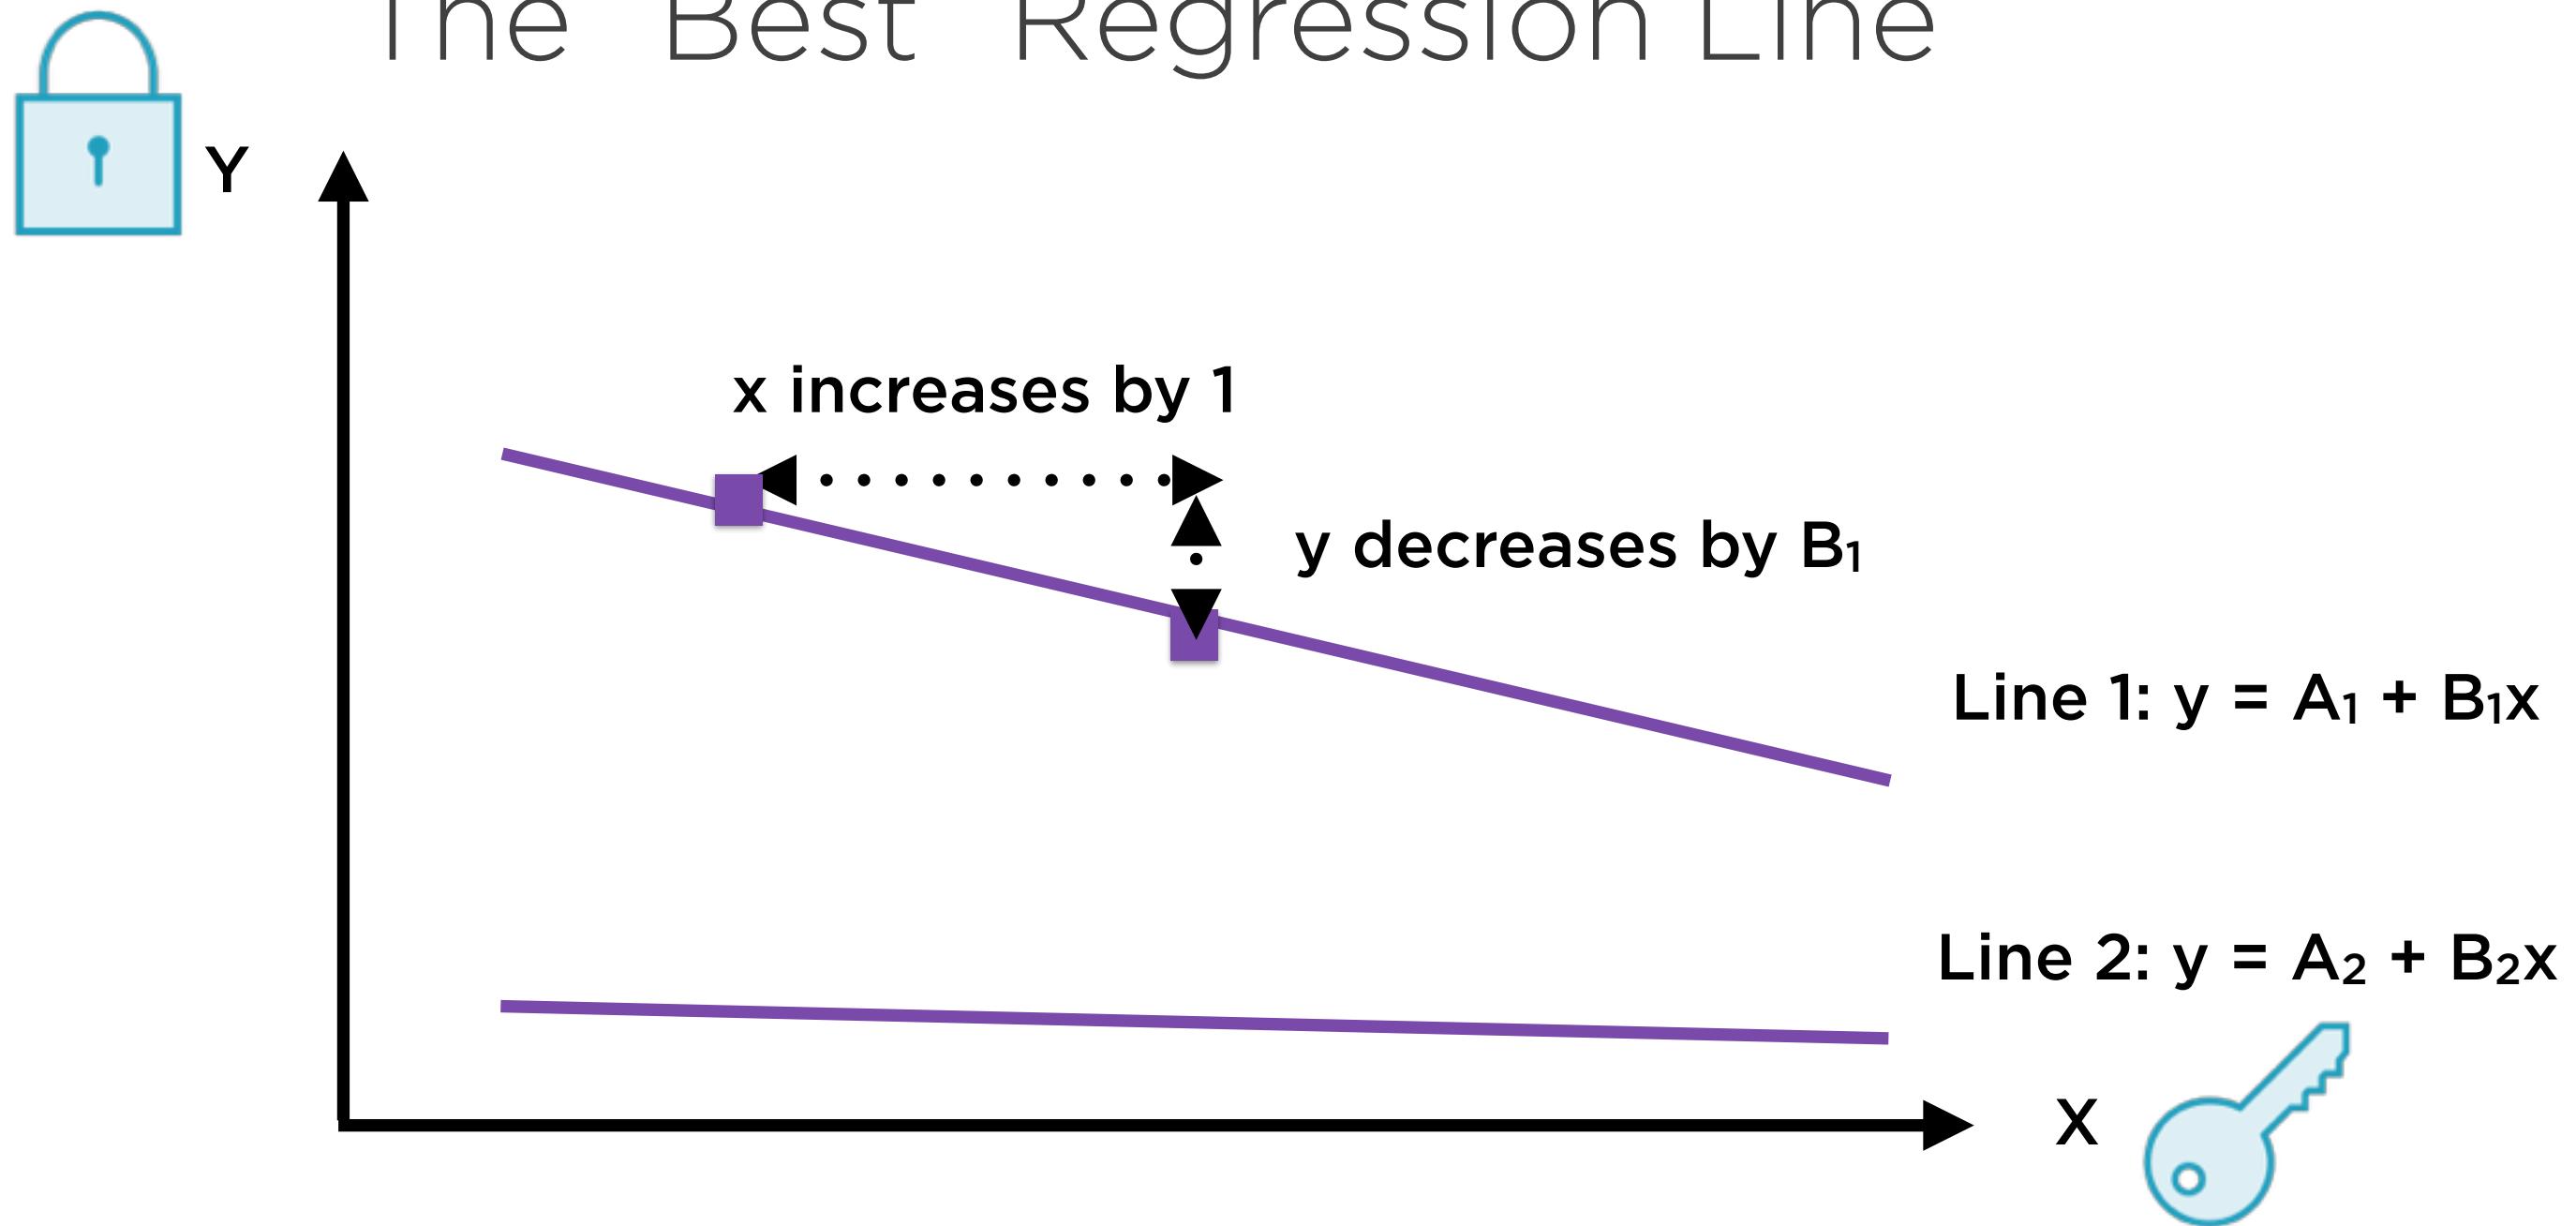


Let's compare two lines, Line 1 and Line 2

The “Best” Regression Line

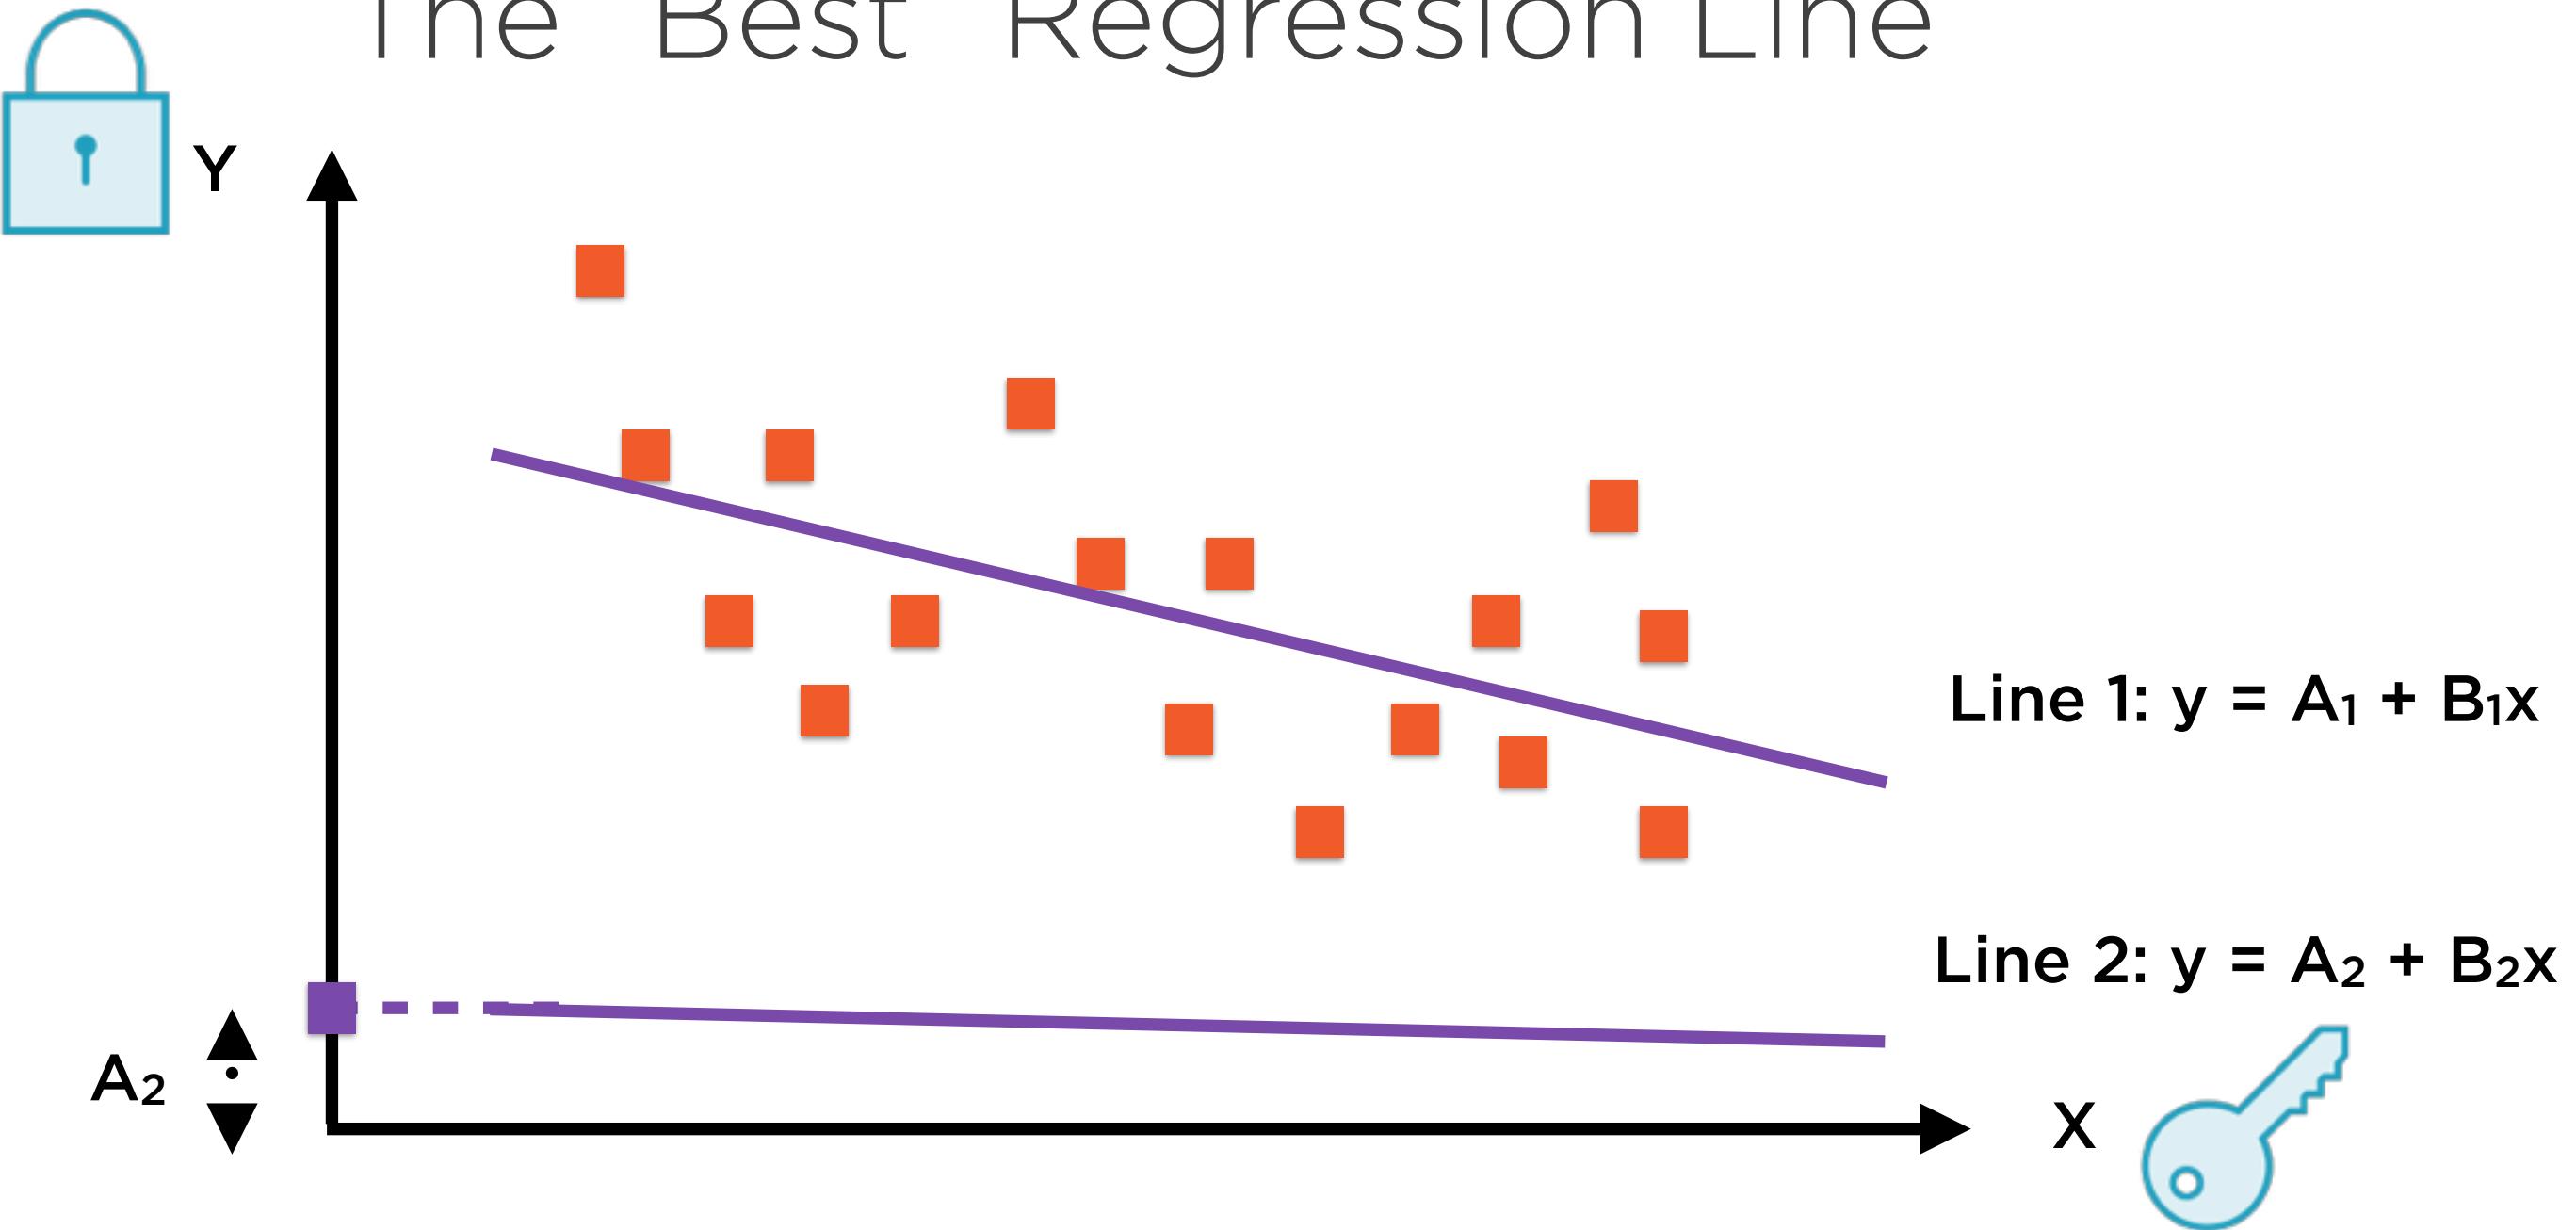


The “Best” Regression Line



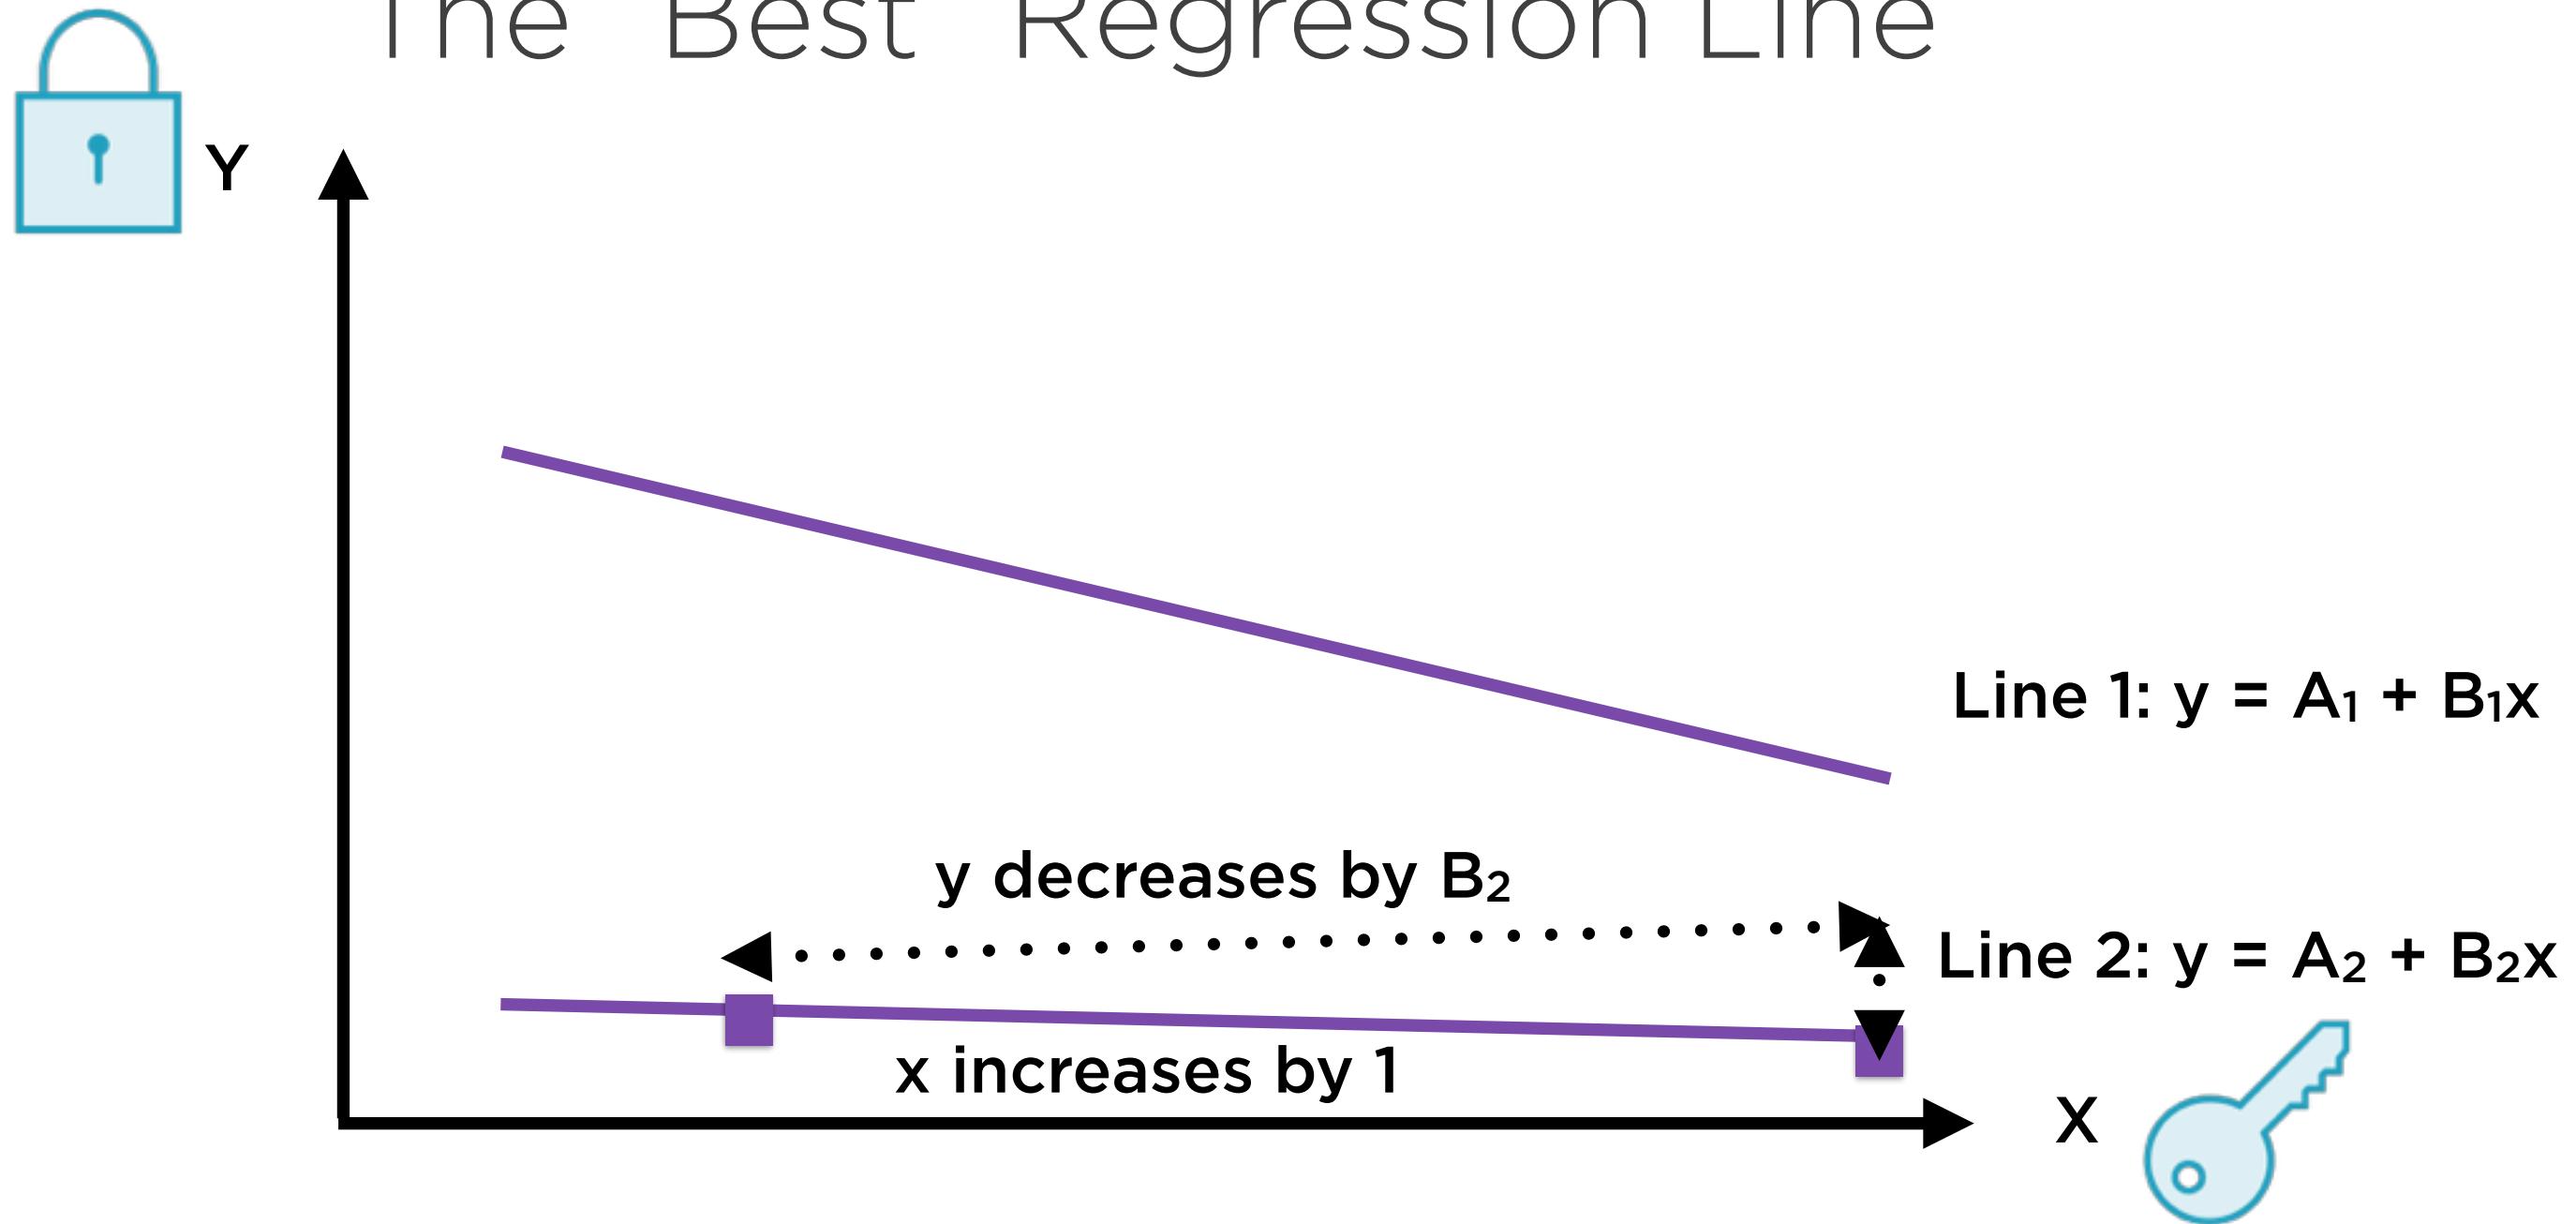
In the first line, if x increases by 1 unit, y decreases by B_1 units

The “Best” Regression Line



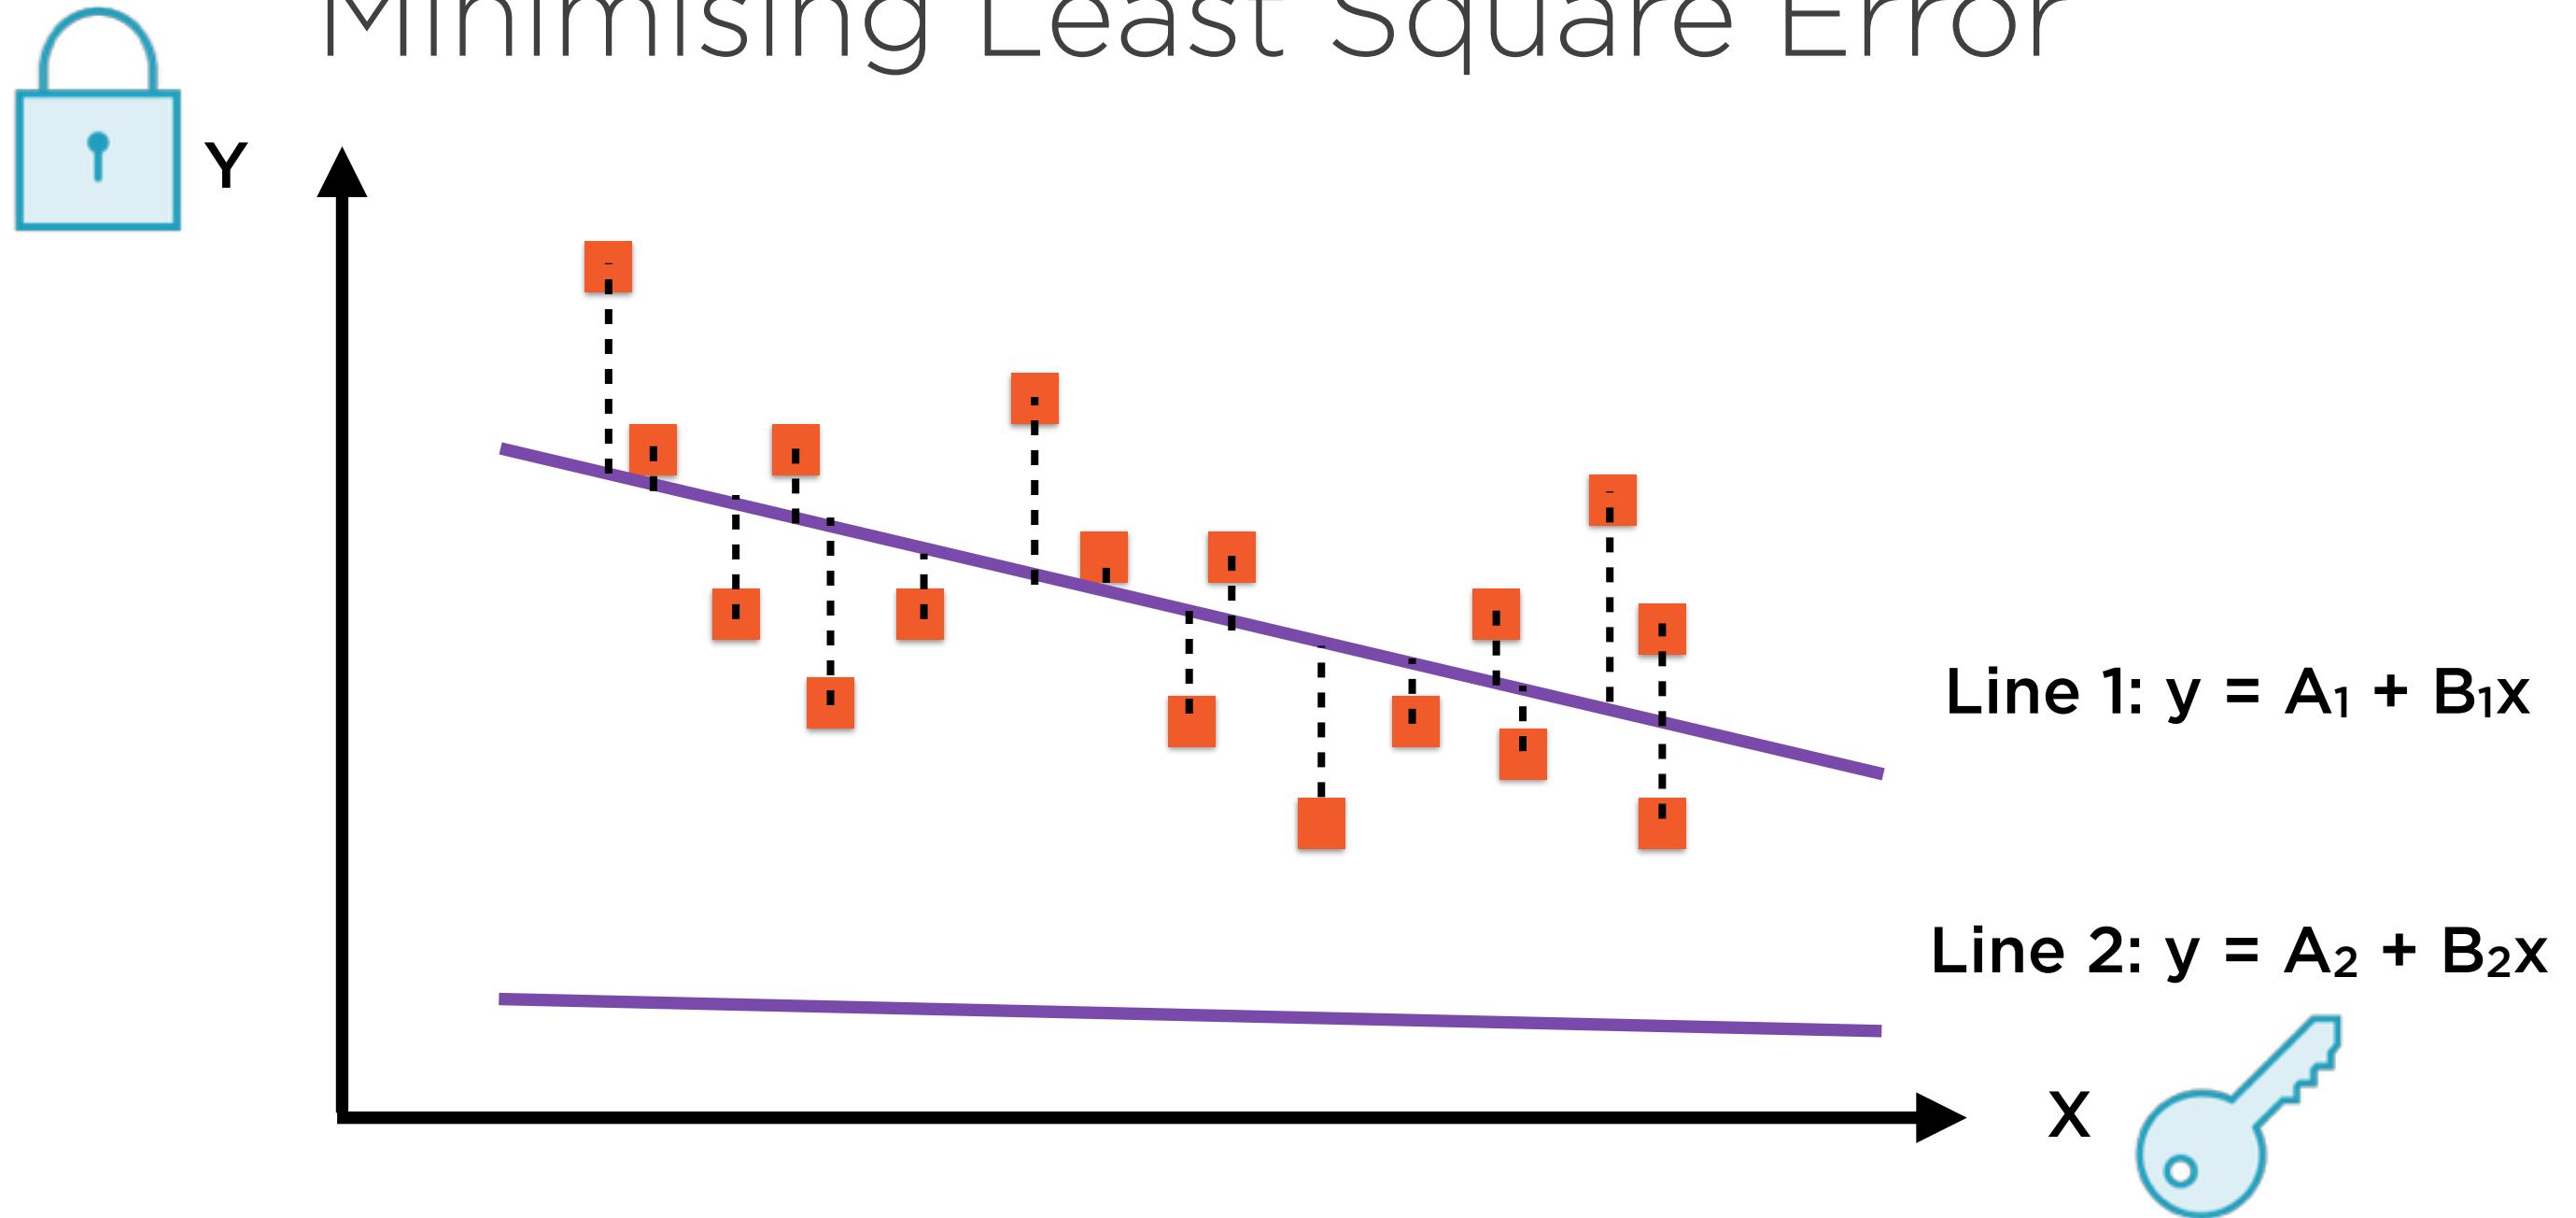
The second line has y-intercept A_2

The “Best” Regression Line



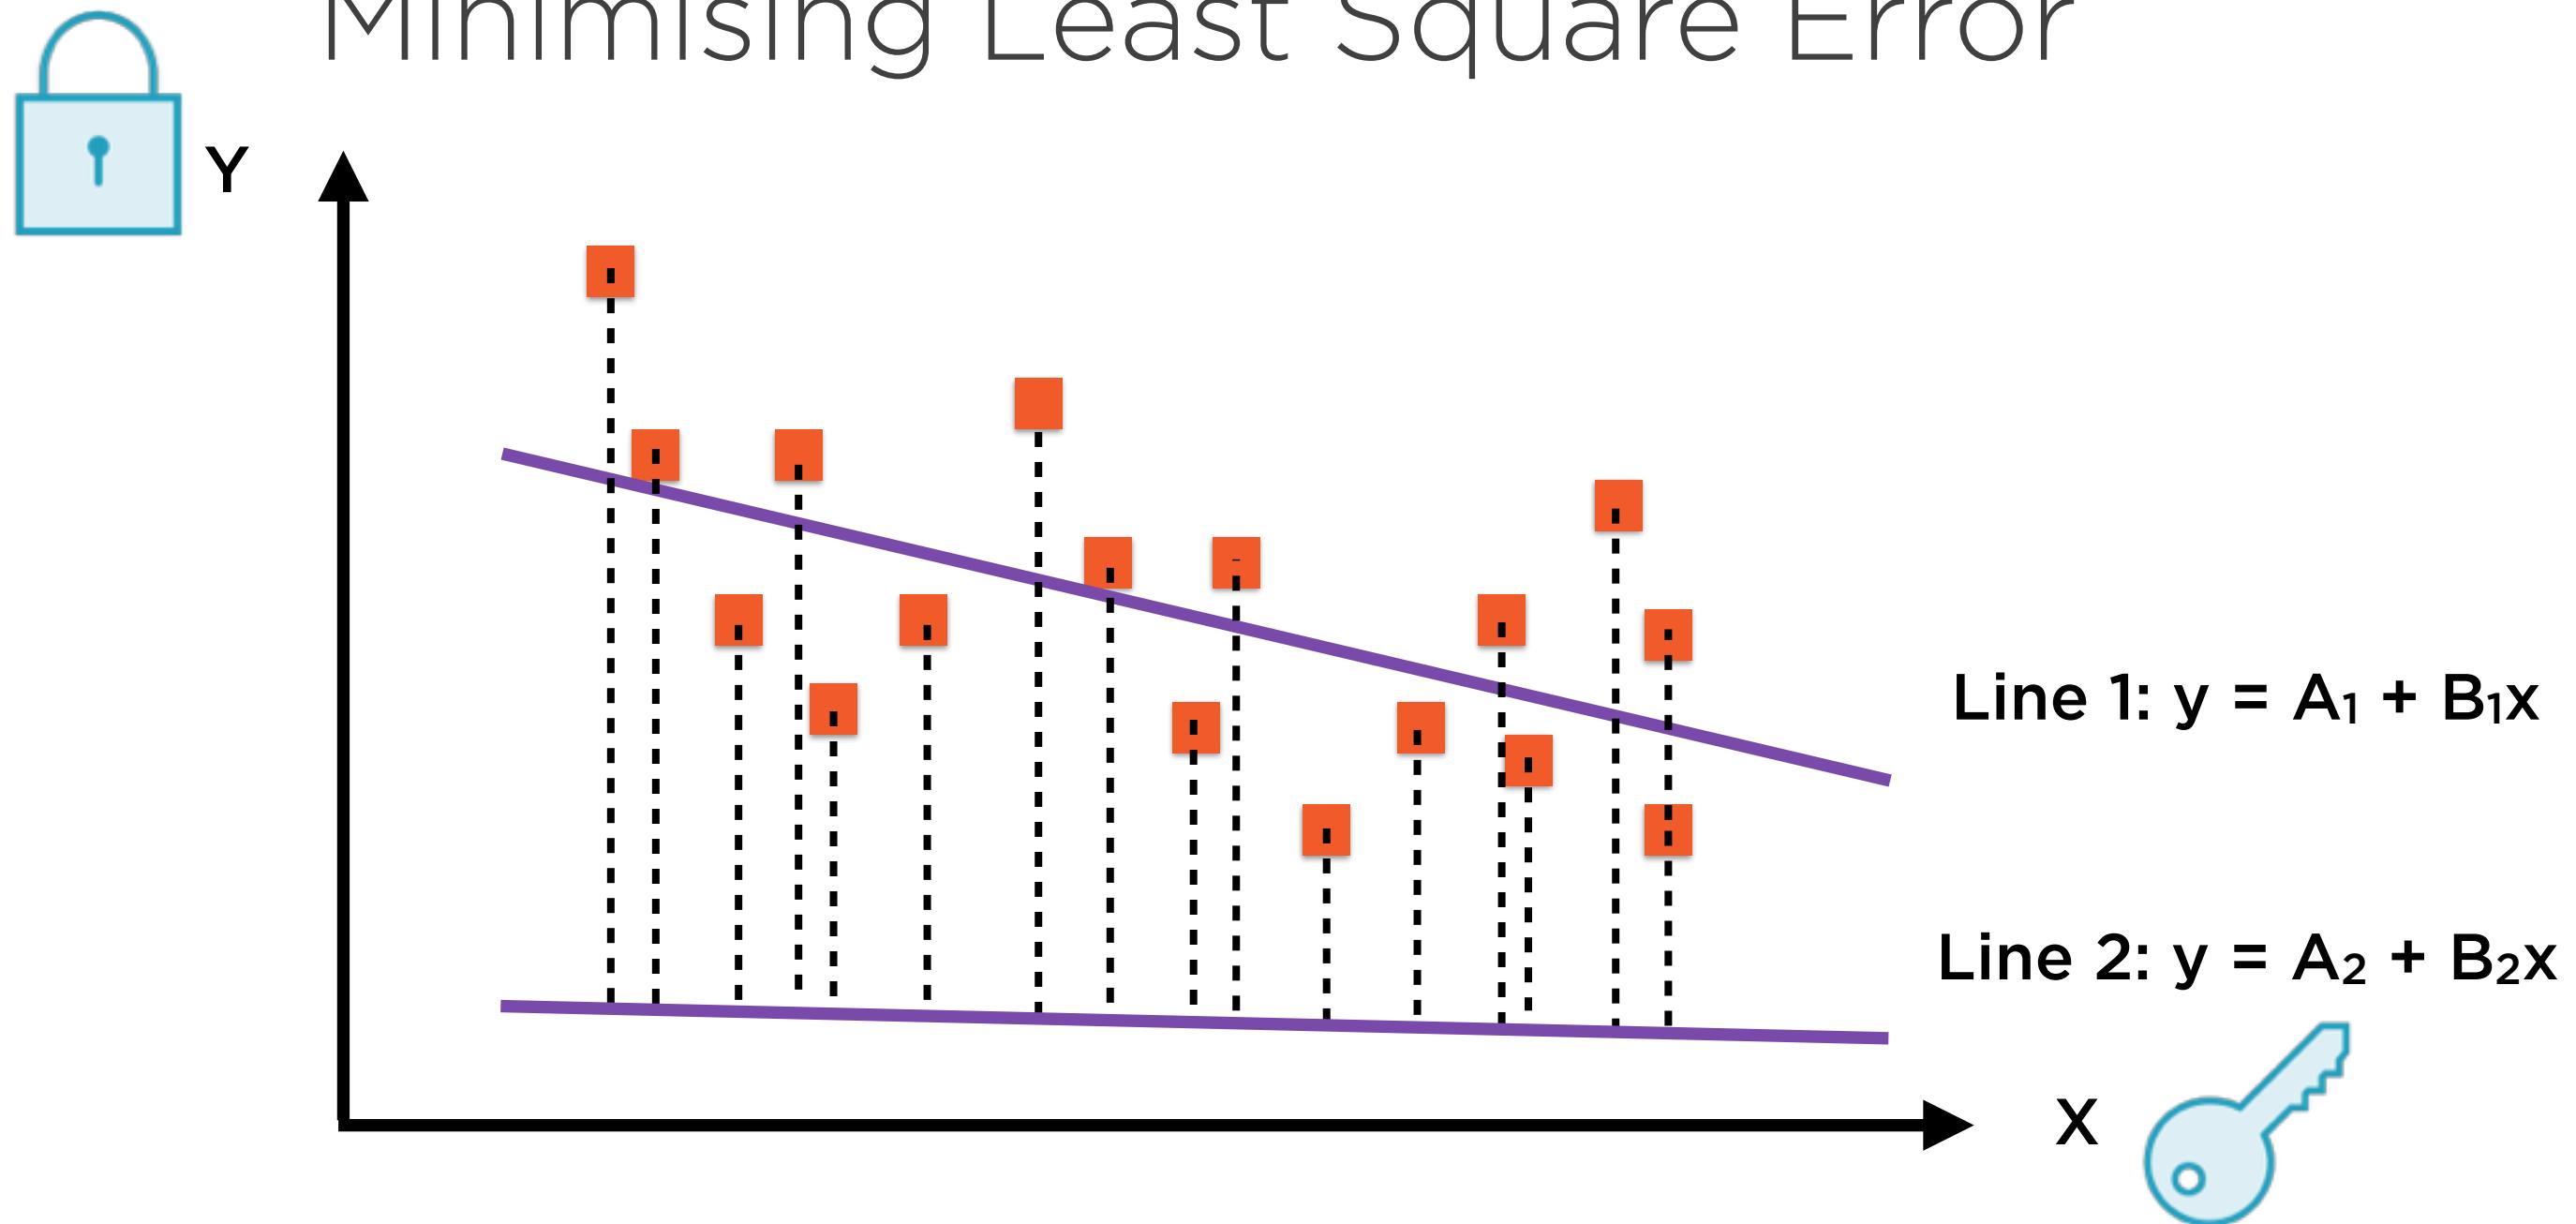
In the second line, if x increases by 1 unit, y decreases by B_2 units

Minimising Least Square Error



Drop vertical lines from each point to
the lines A and B

Minimising Least Square Error



Drop vertical lines from each point to
the lines A and B

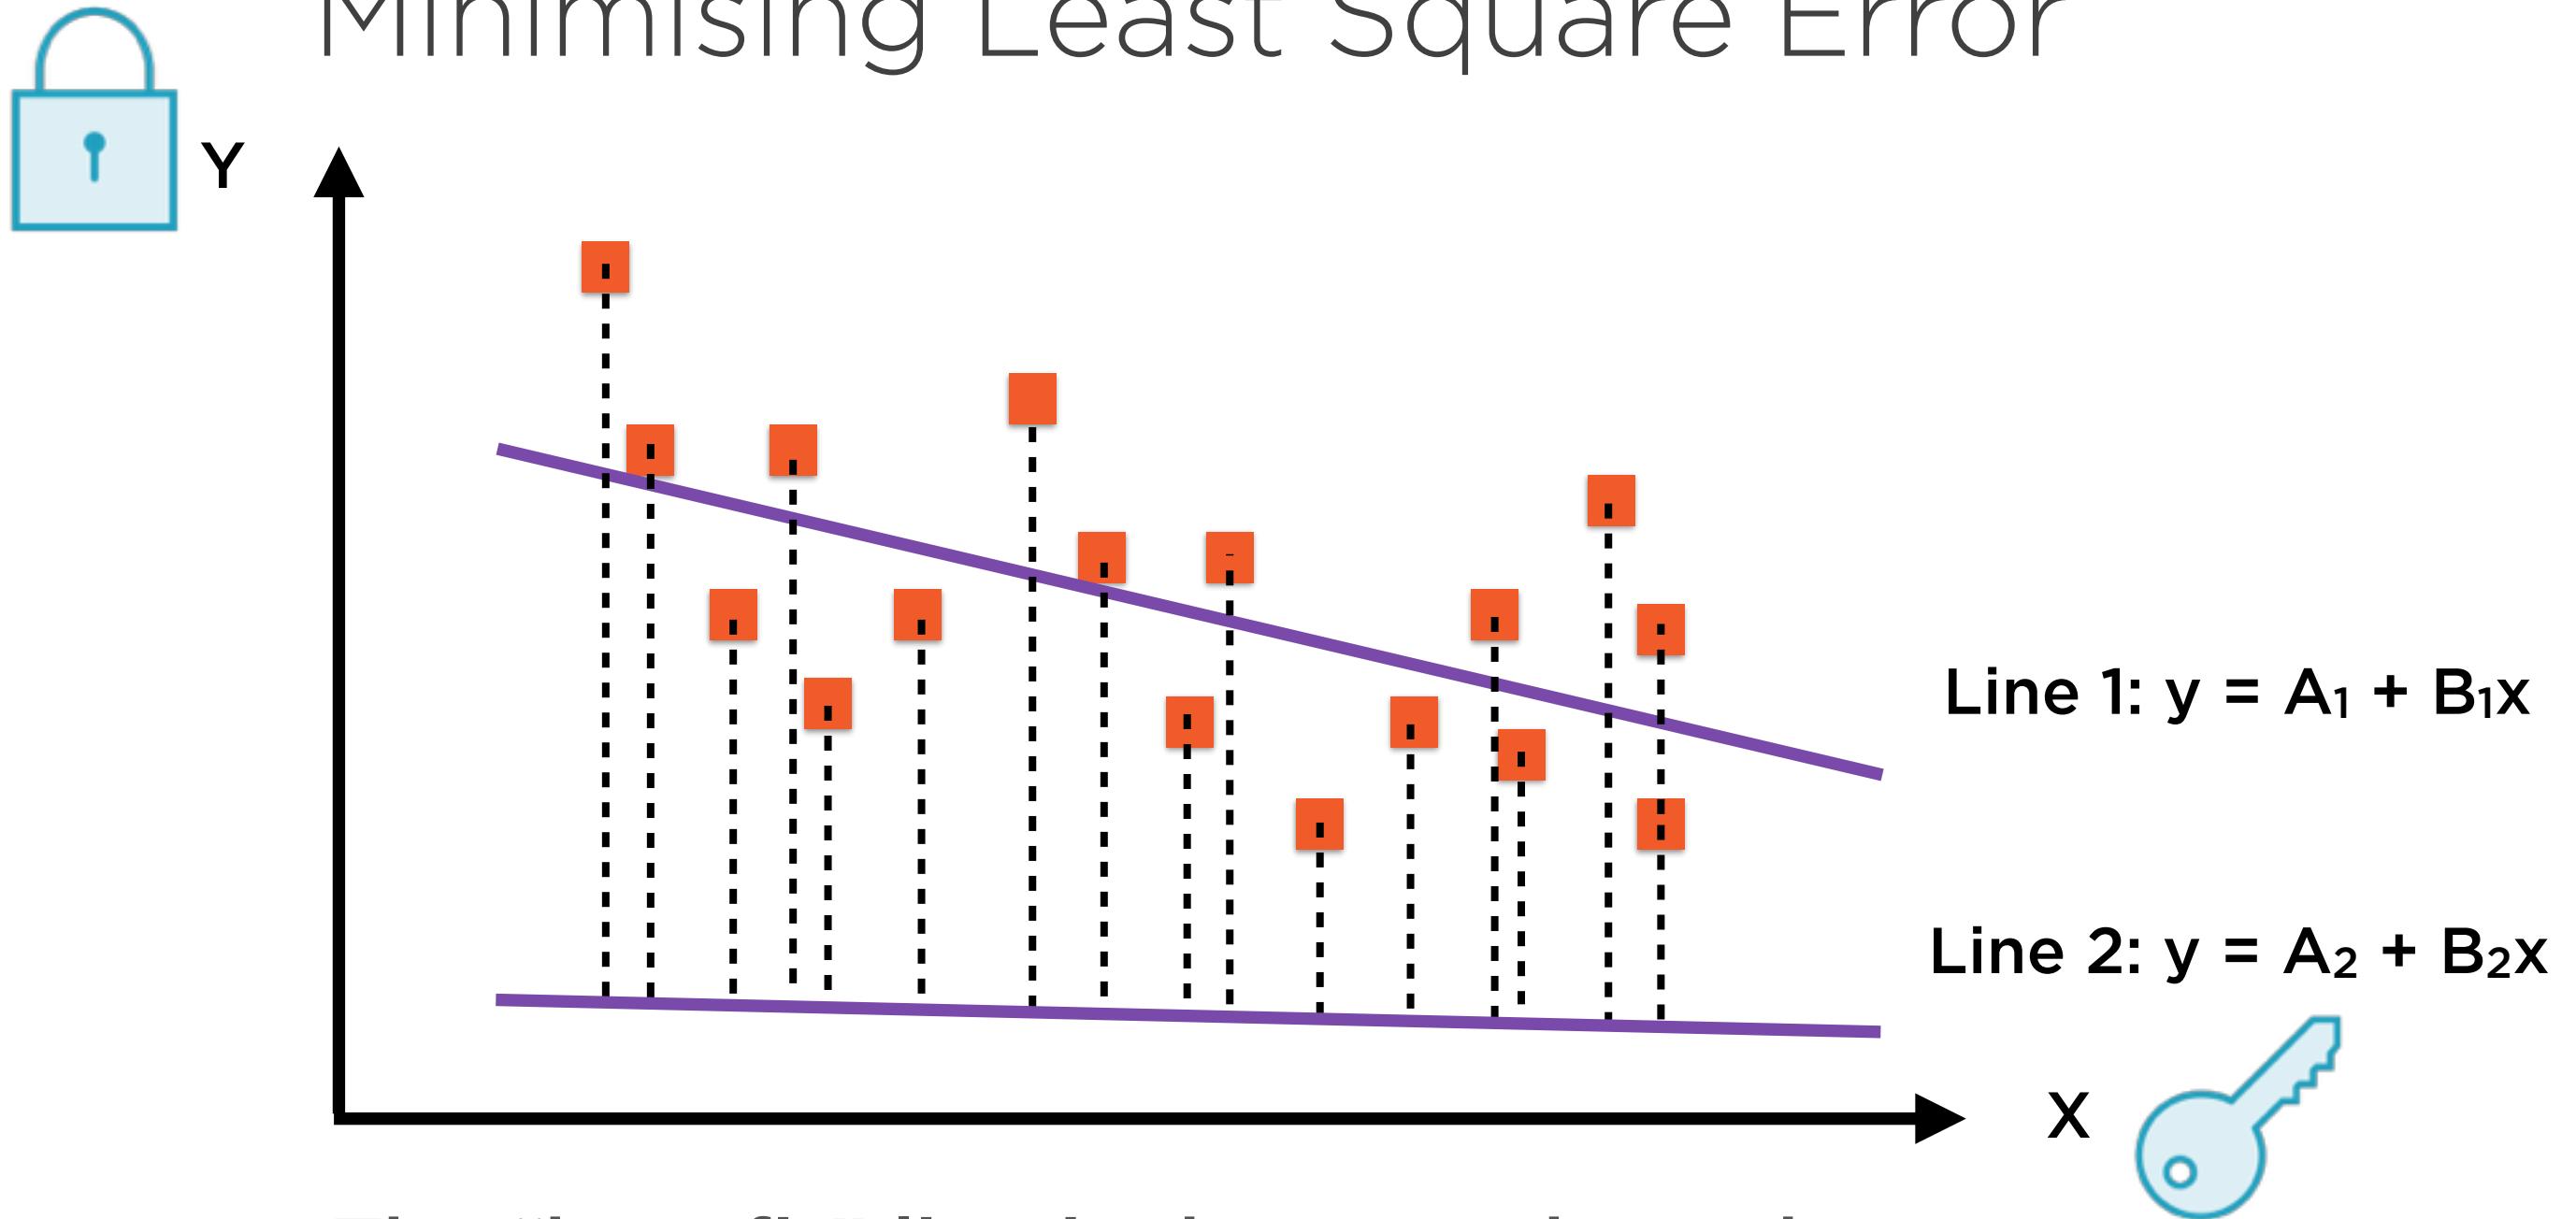
Simple Regression

Regression Equation:

$$y = A + Bx$$

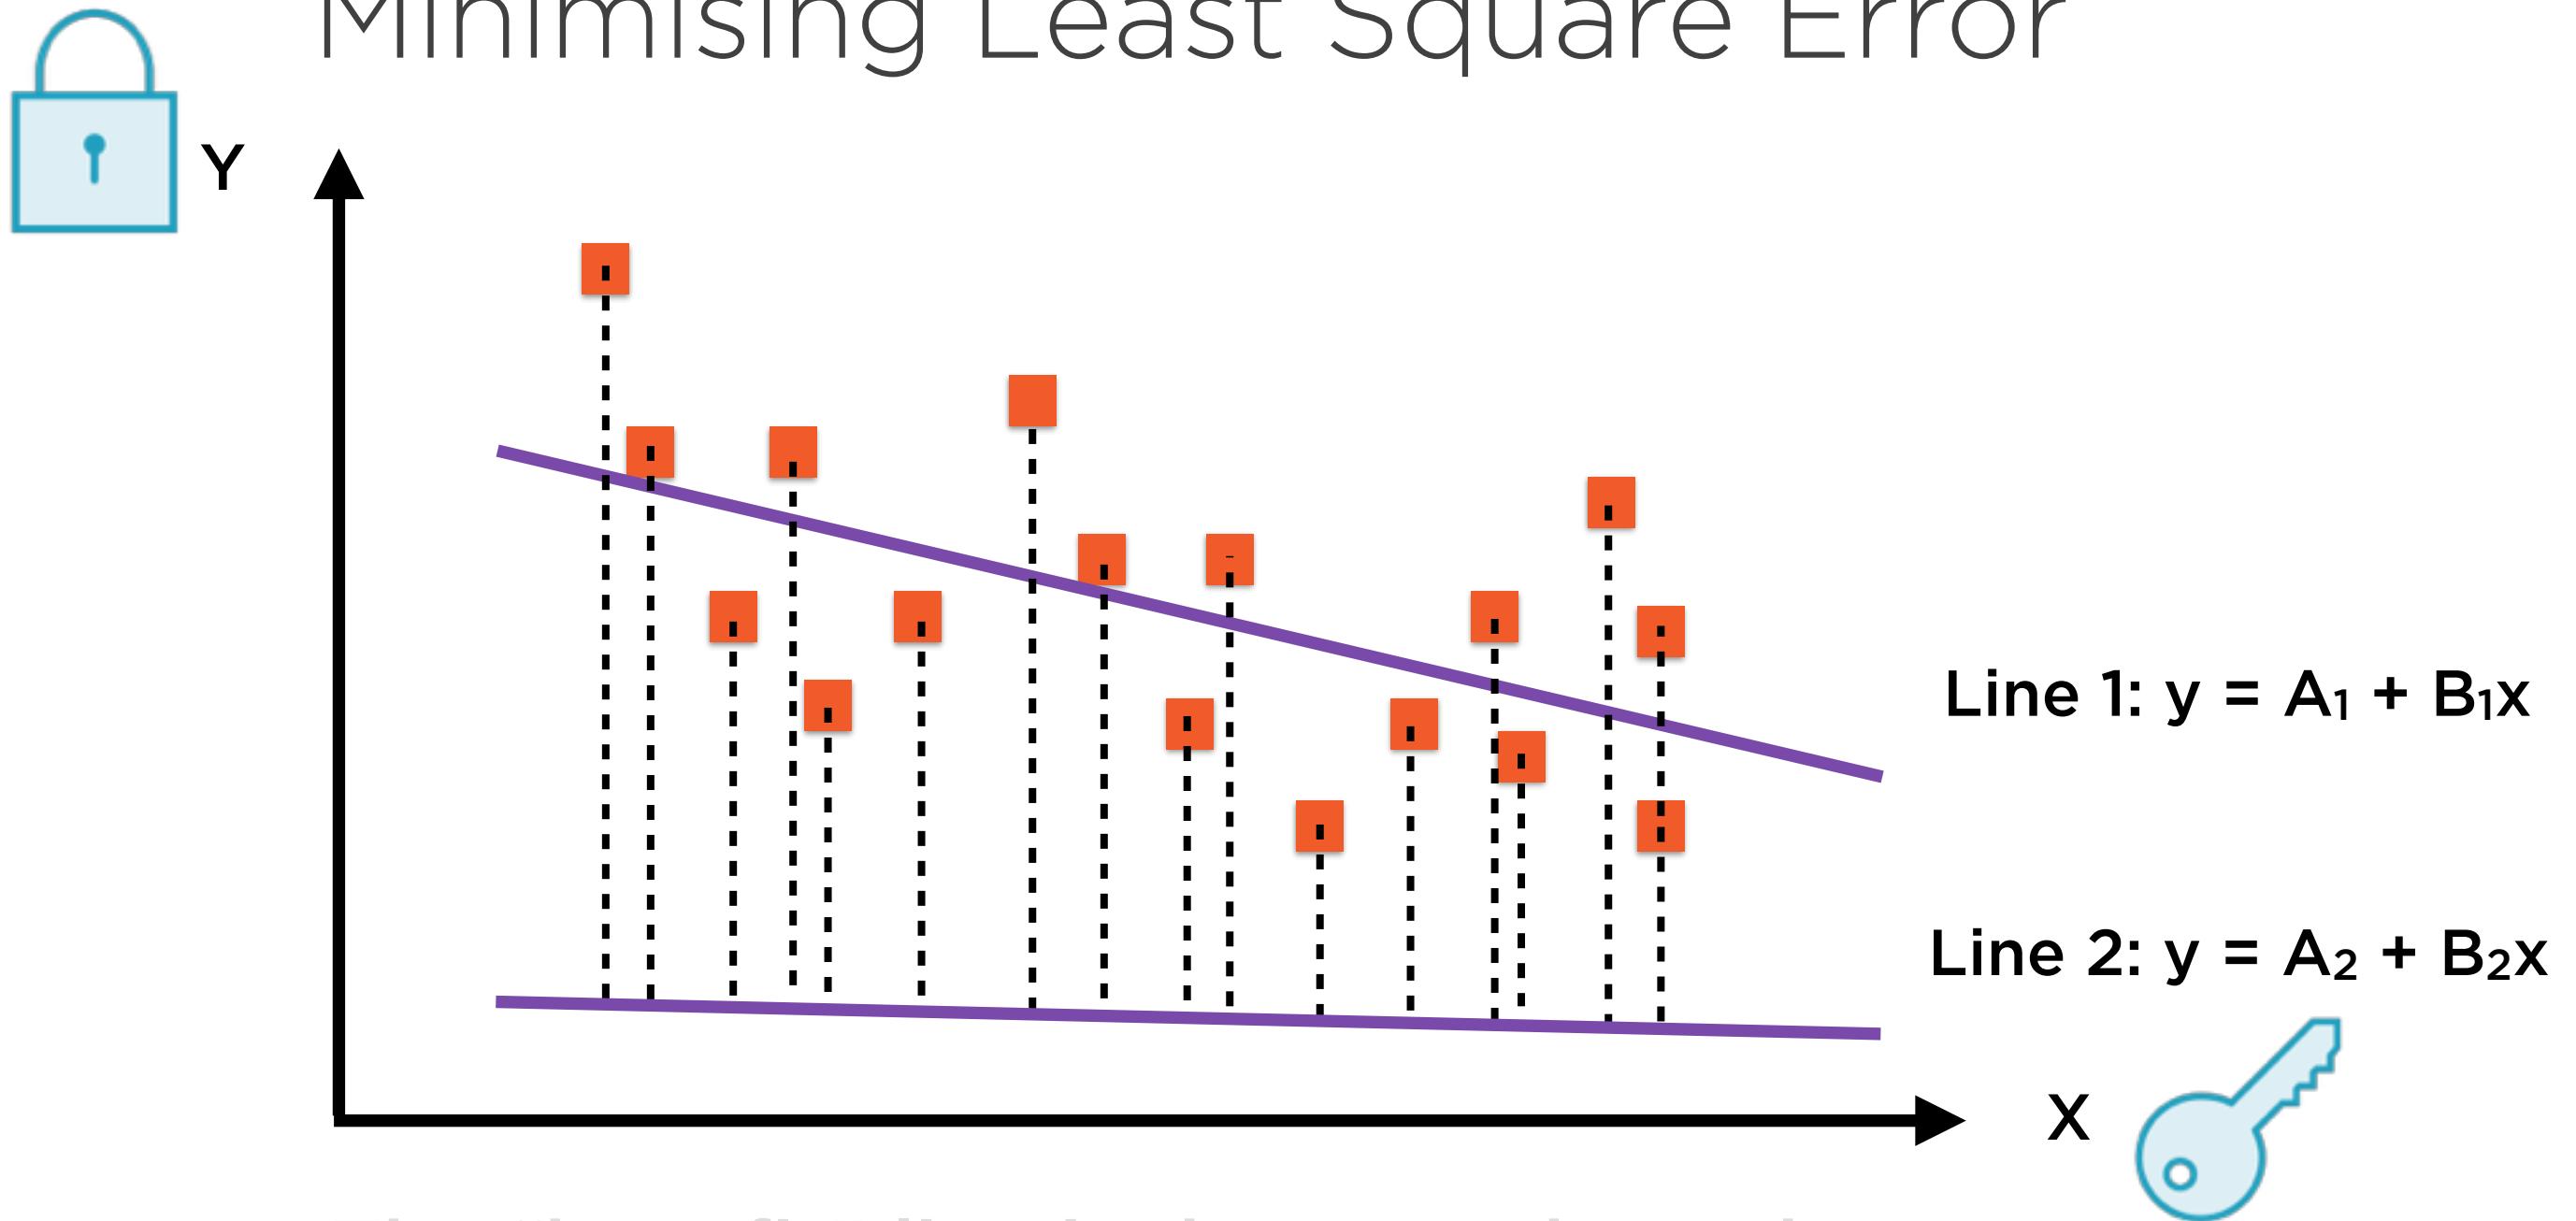
$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} = A \begin{bmatrix} 1 \\ 1 \\ 1 \\ \dots \\ 1 \end{bmatrix} + B \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \dots \\ x_n \end{bmatrix} + \begin{bmatrix} e_1 \\ e_2 \\ e_3 \\ \dots \\ e_n \end{bmatrix}$$

Minimising Least Square Error



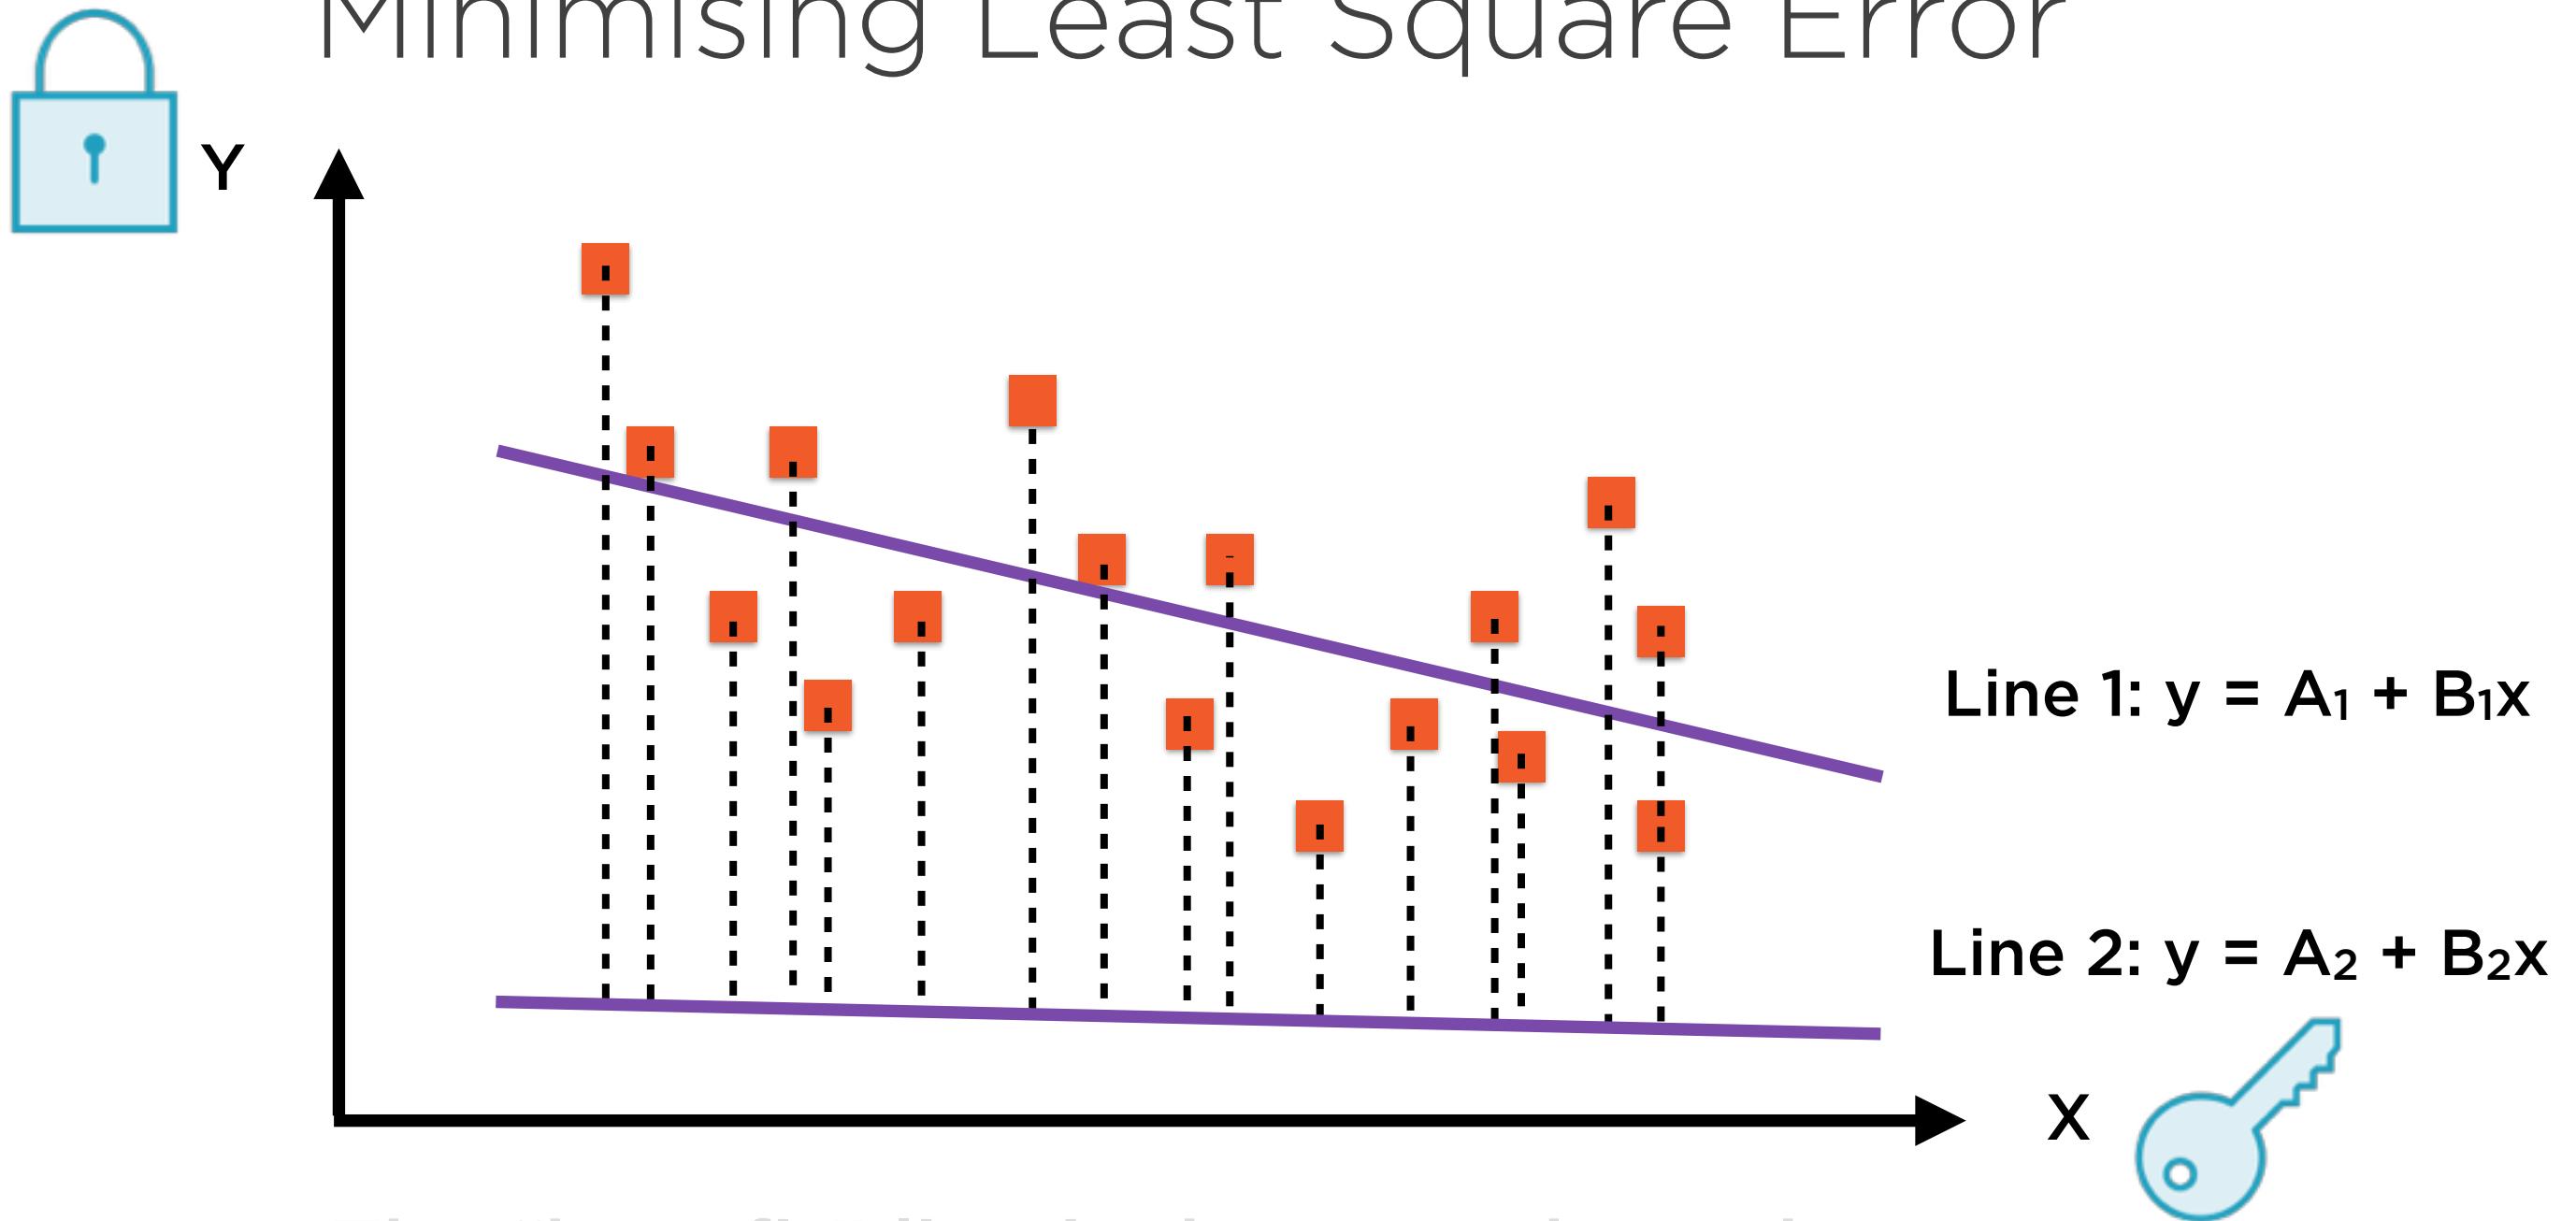
The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum

Minimising Least Square Error



The “best fit” line is the one where the sum of the squares of the lengths of these dotted lines is minimum

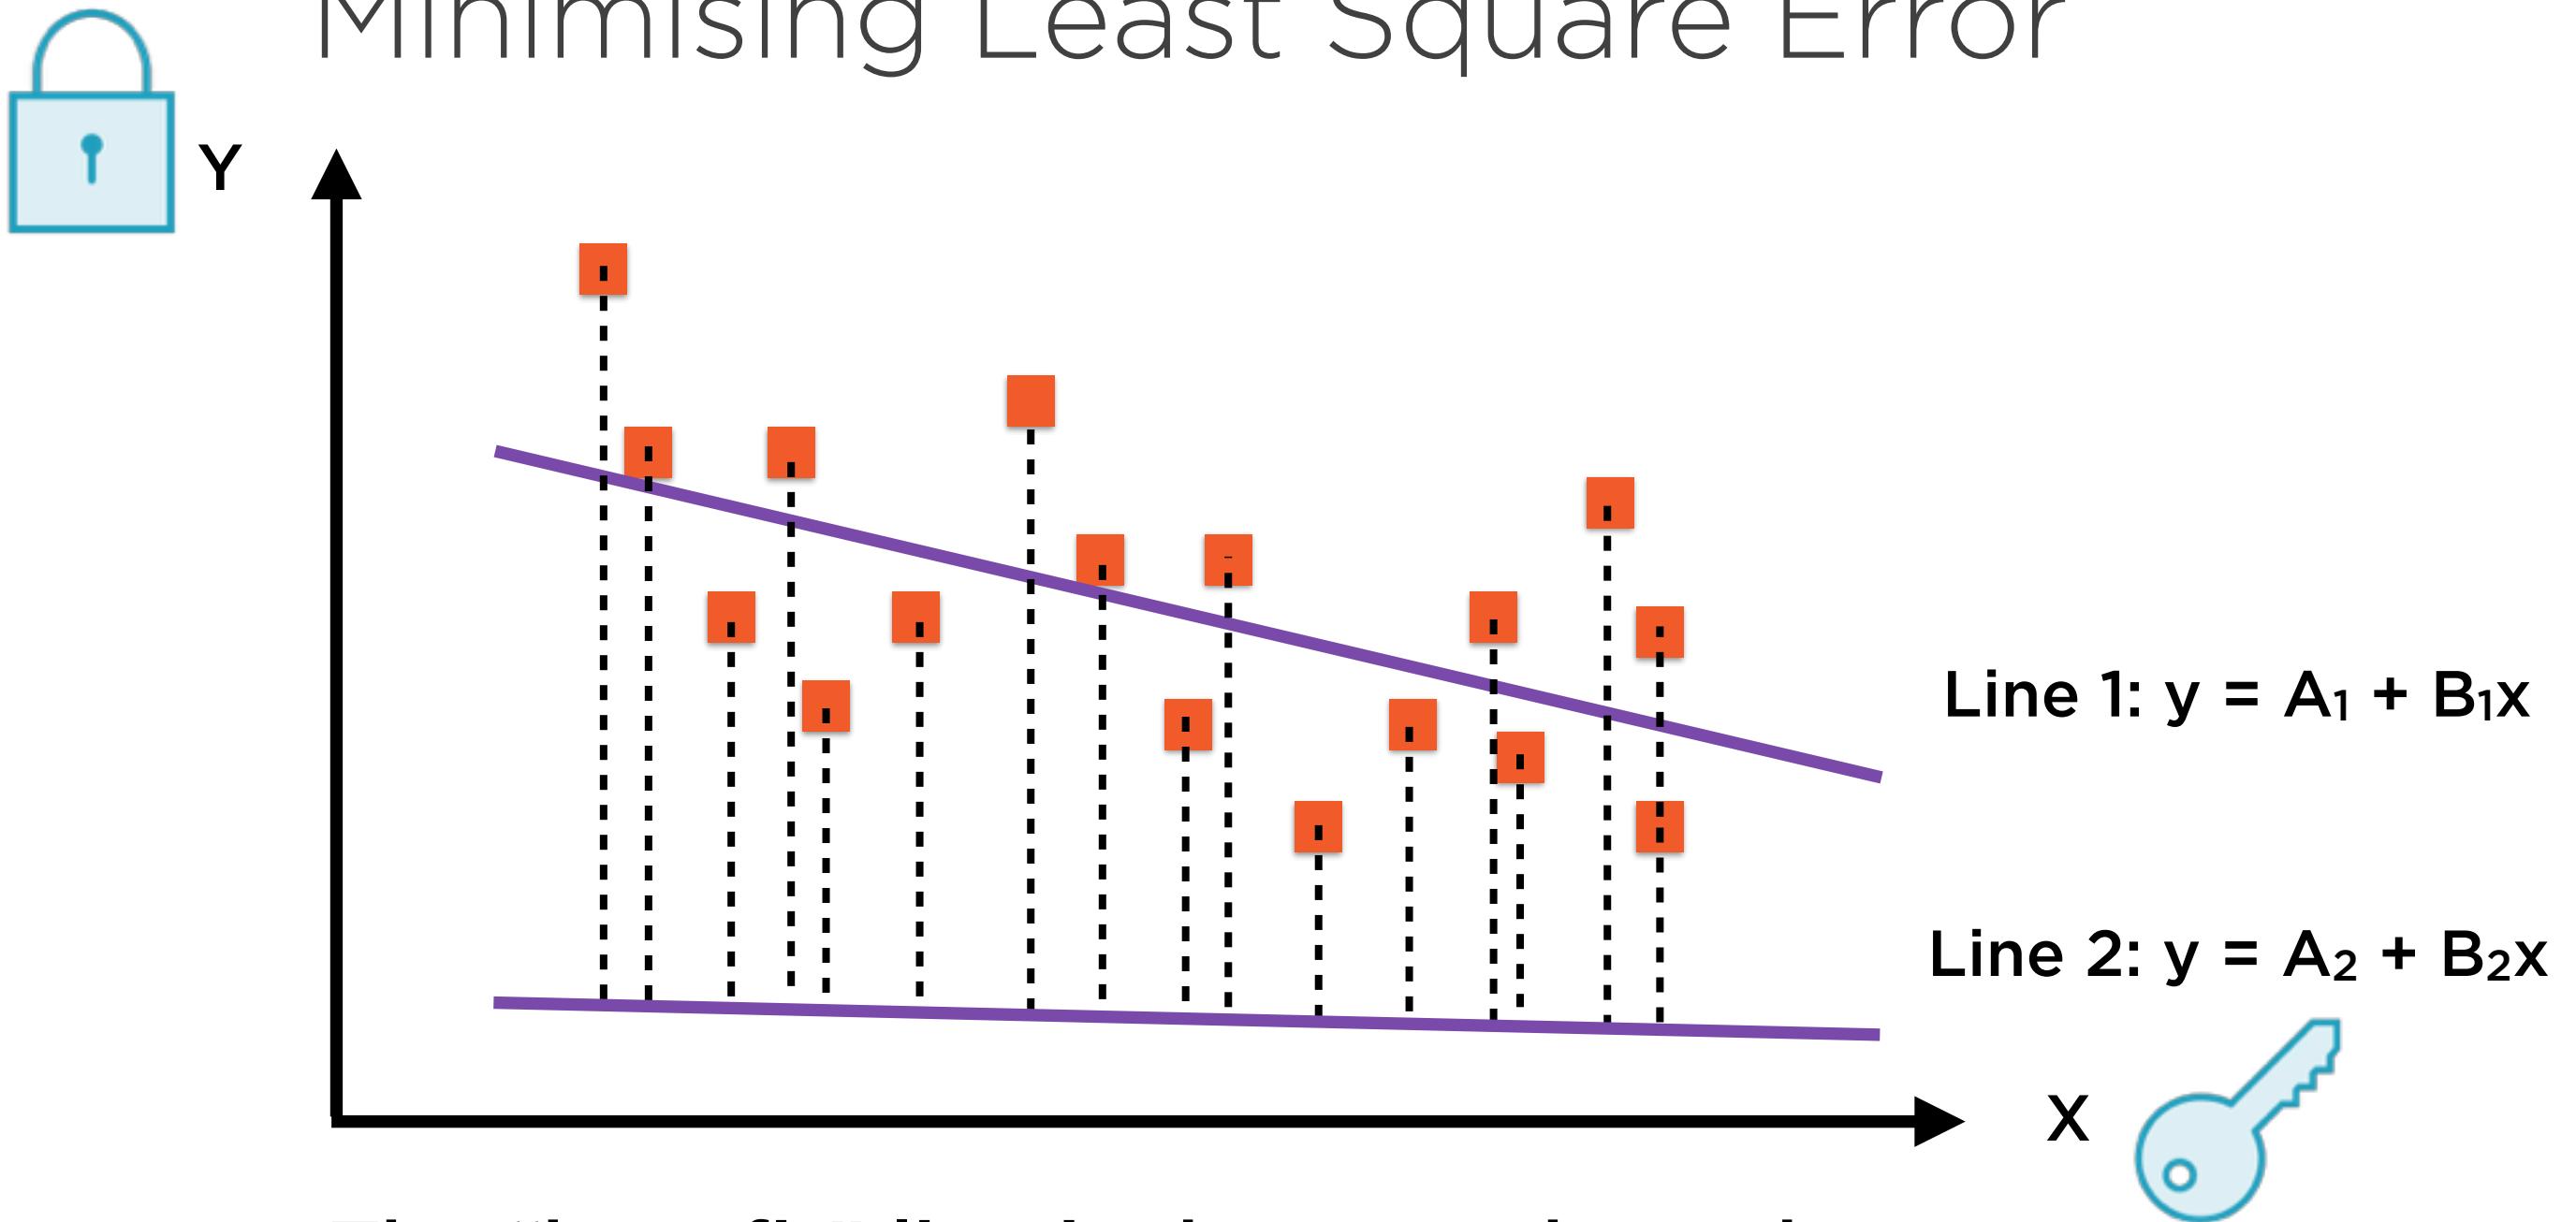
Minimising Least Square Error



The “best fit” line is the one where the sum of the squares of the lengths of the **errors** is minimum



Minimising Least Square Error



The “best fit” line is the one where the sum of the squares of the lengths of the errors is minimum

Simple Regression

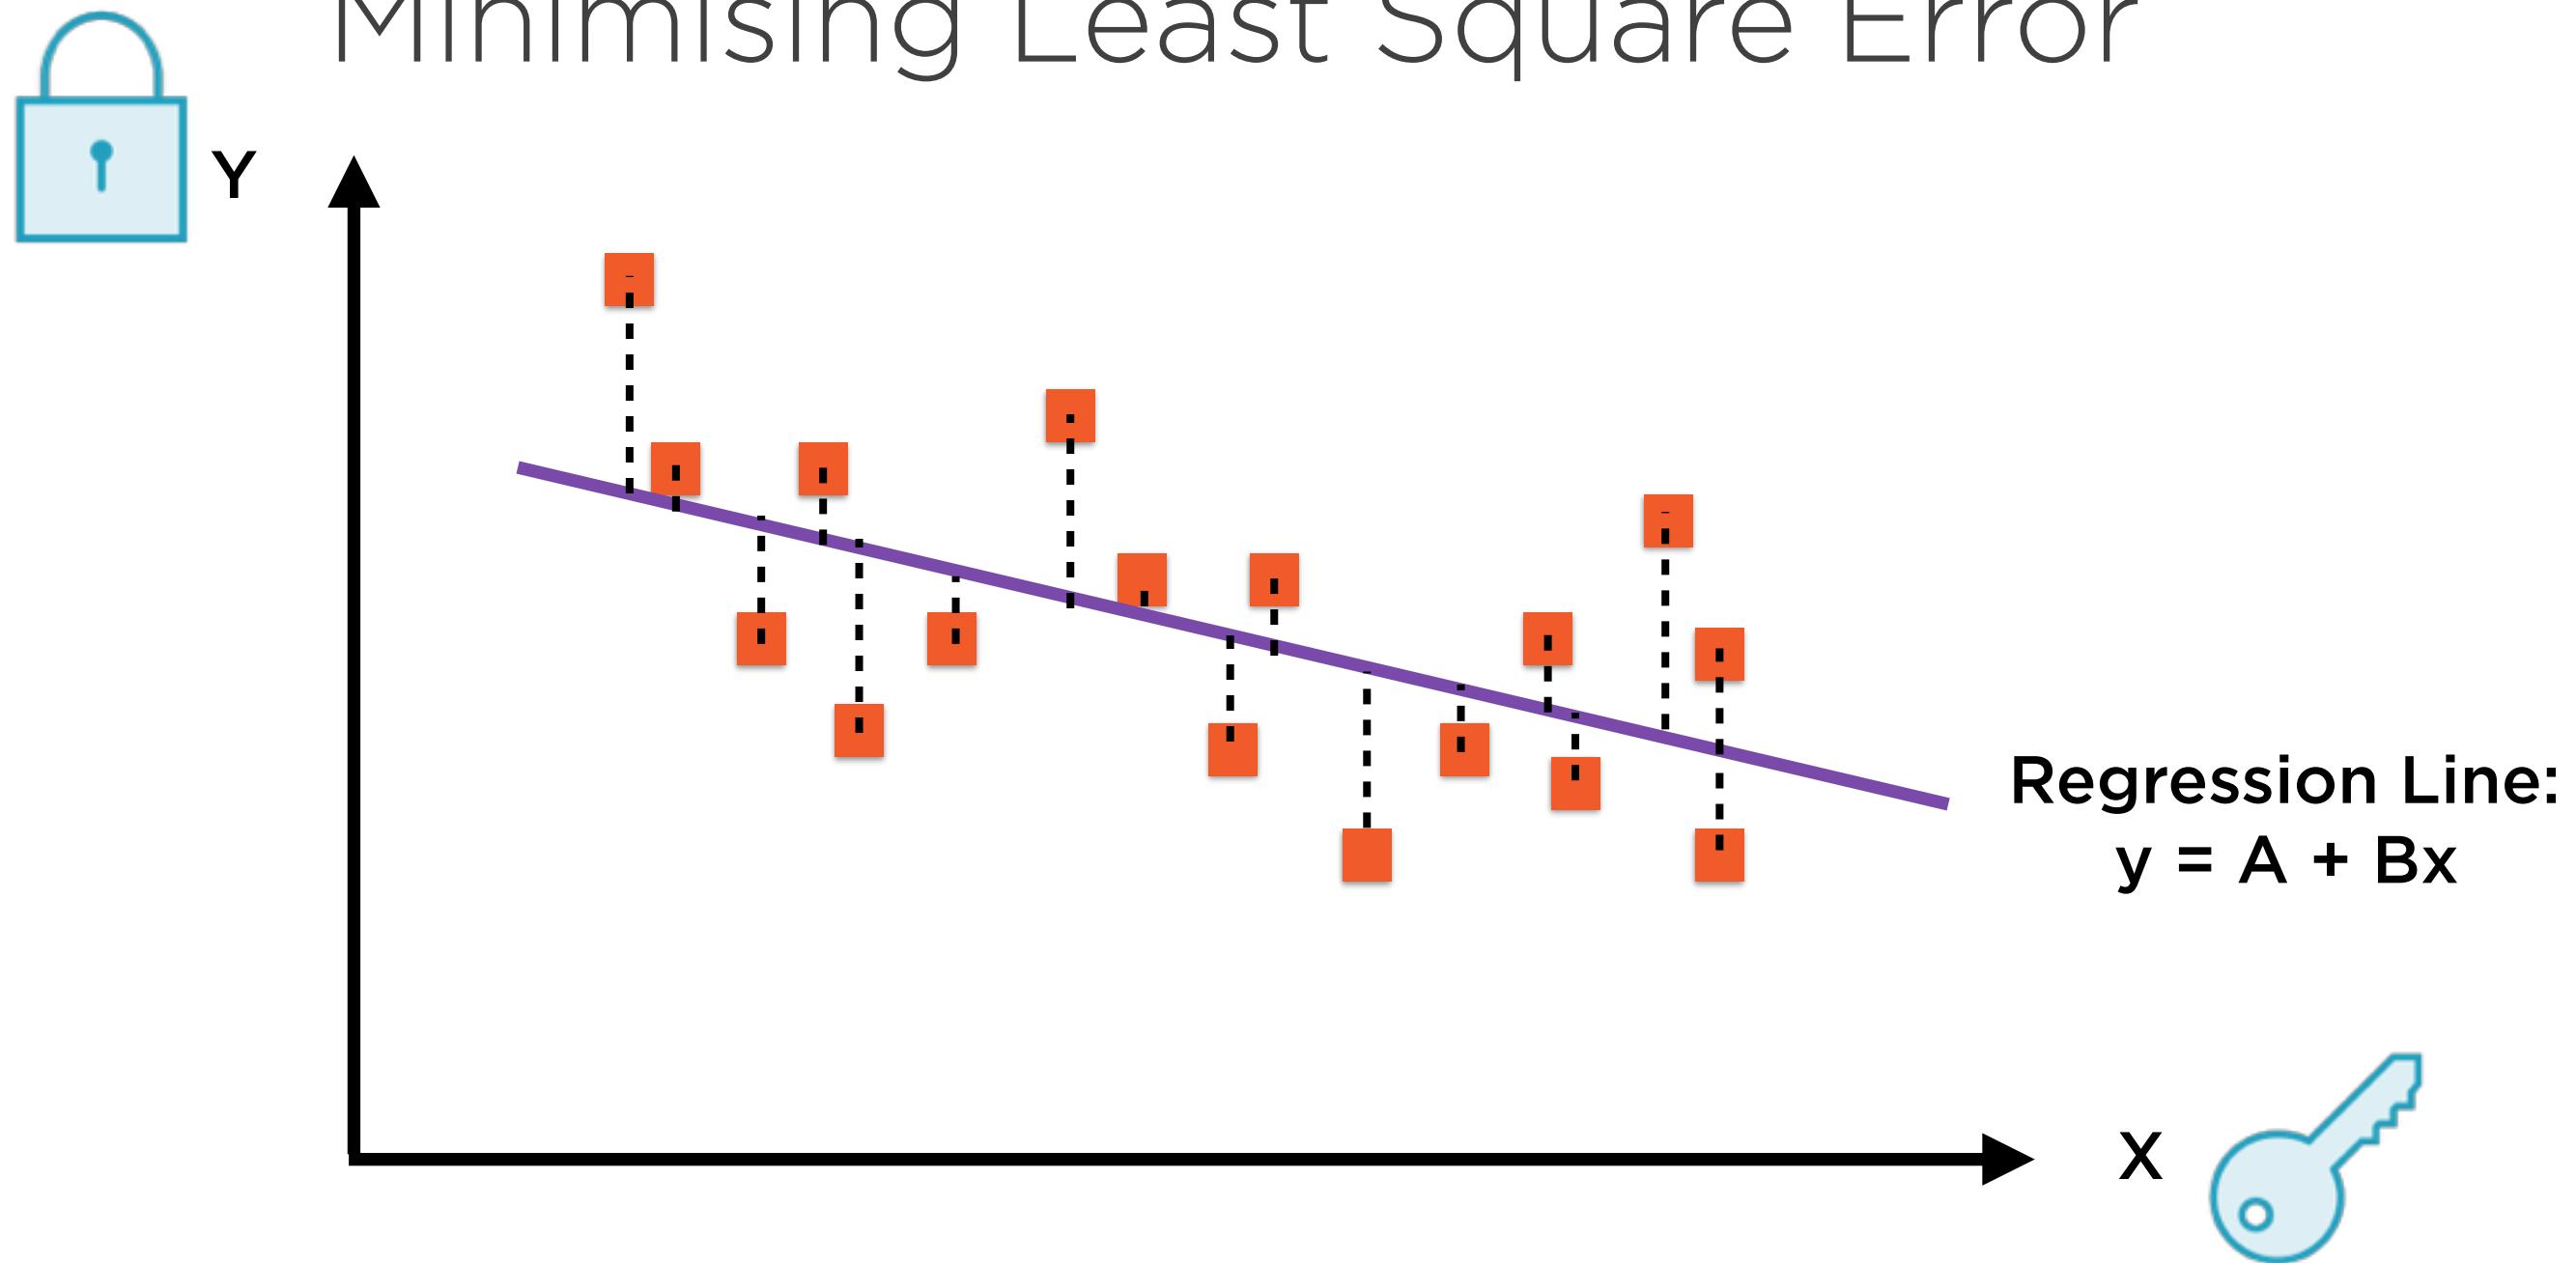
Regression Equation:

$$y = A + Bx$$

$$\begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ \dots \\ y_n \end{bmatrix} = A \begin{bmatrix} 1 \\ 1 \\ 1 \\ \dots \\ 1 \end{bmatrix} + B \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ \dots \\ x_n \end{bmatrix} + \begin{bmatrix} e_1 \\ e_2 \\ e_3 \\ \dots \\ e_n \end{bmatrix}$$

The “best fit” line is the one where the sum of the squares of the lengths of the errors is minimum

Minimising Least Square Error



The “best fit” line is called the
regression line

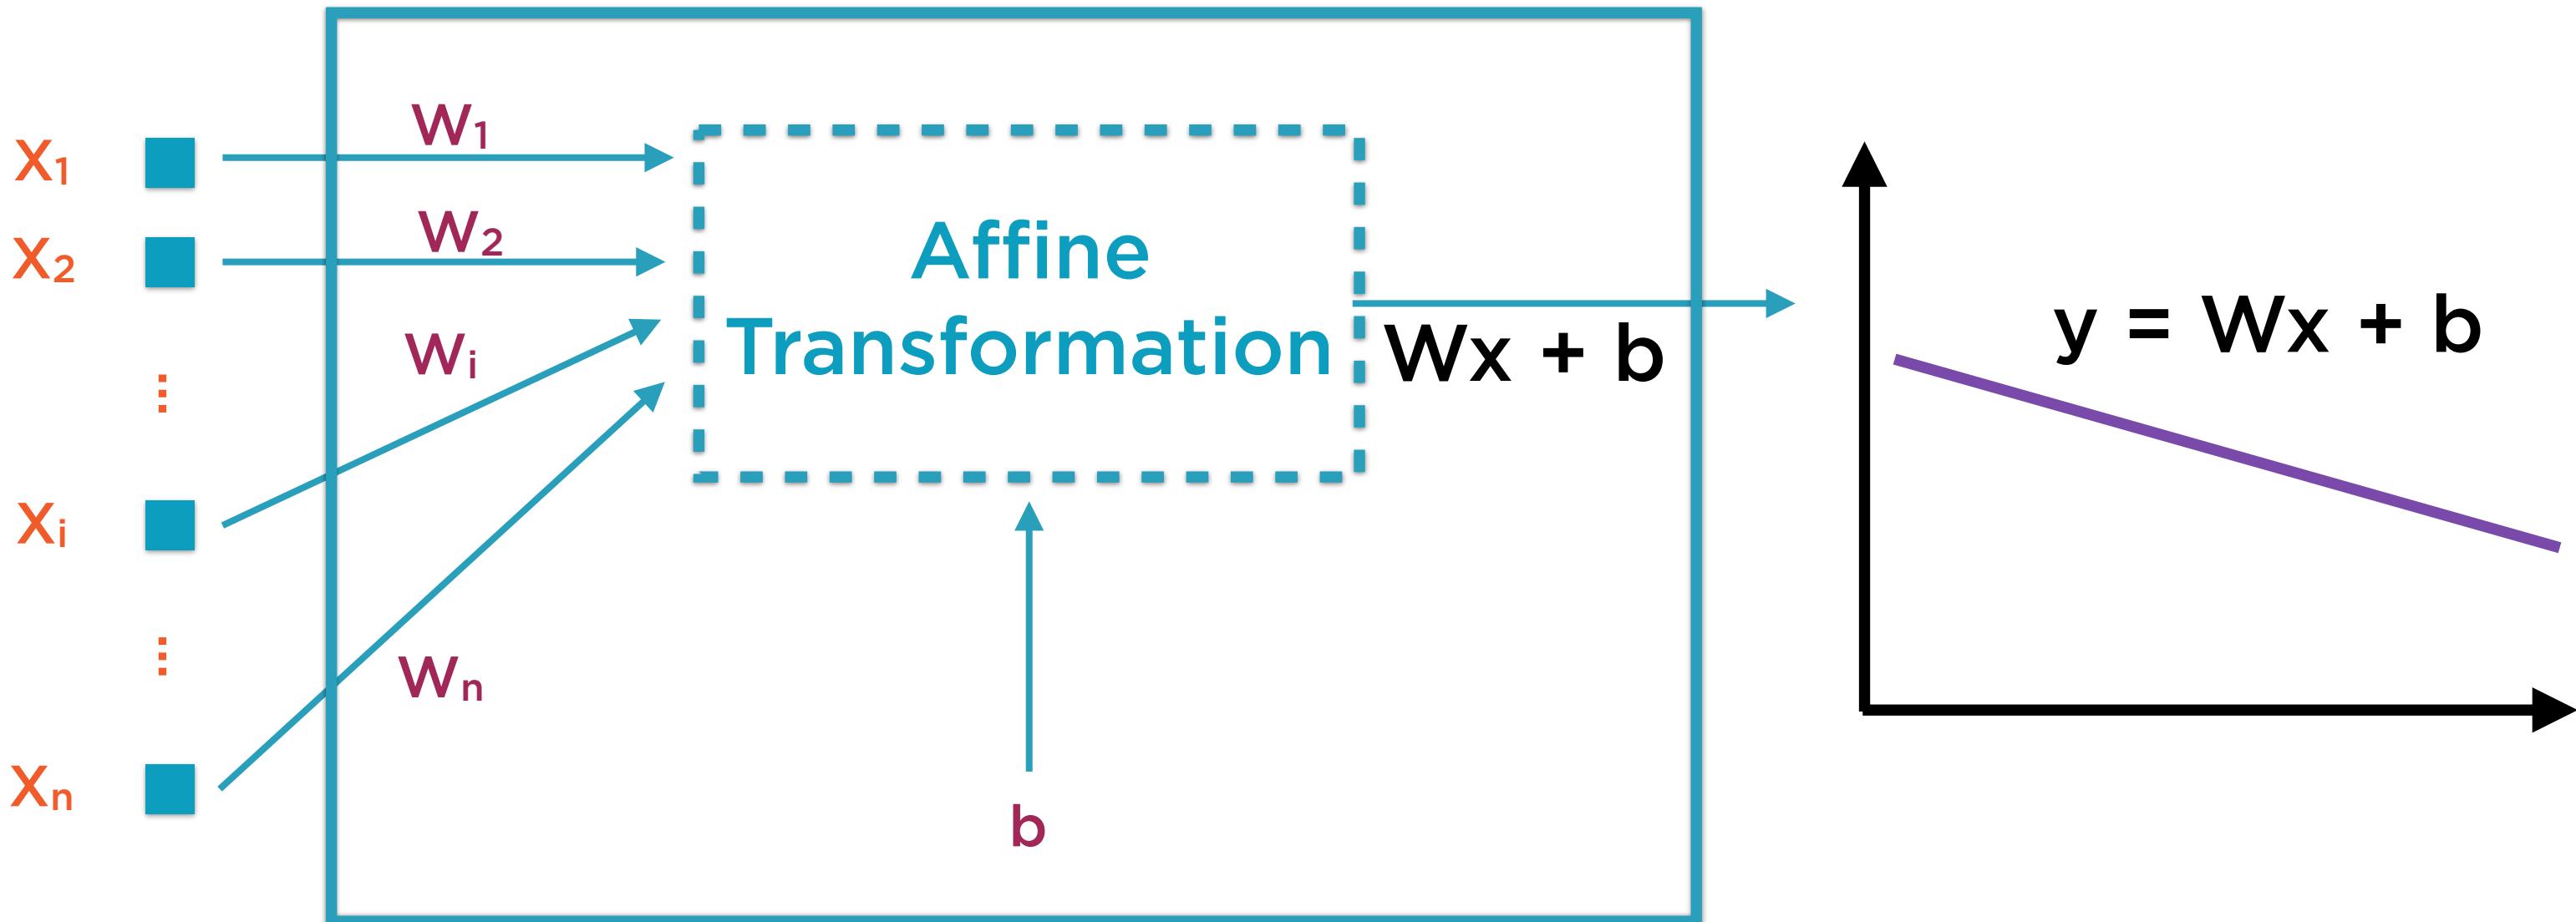
Optimizers for the “Best-fit”

Method of
moments

Method of least
squares

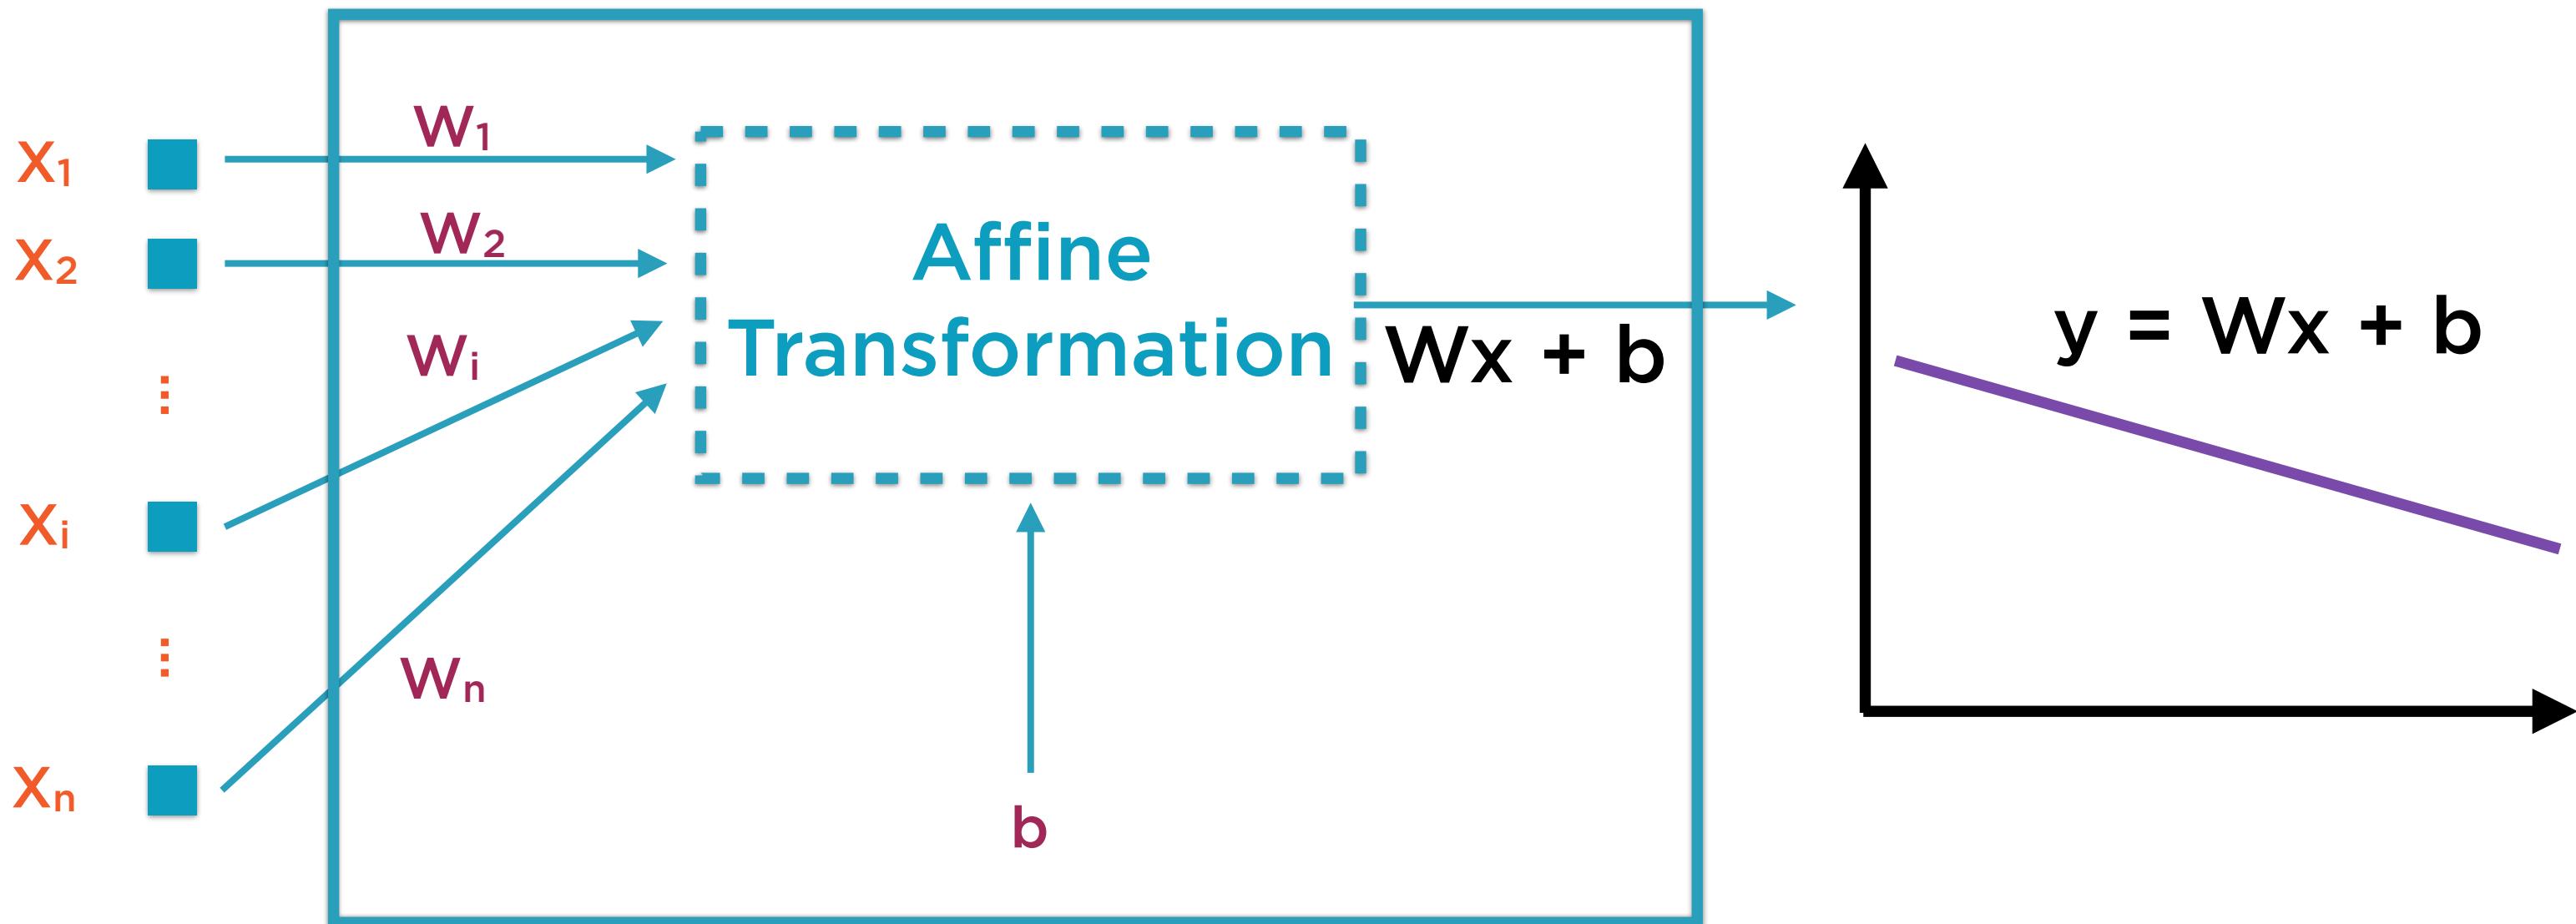
Maximum
likelihood
estimation

Operation of a Single Neuron



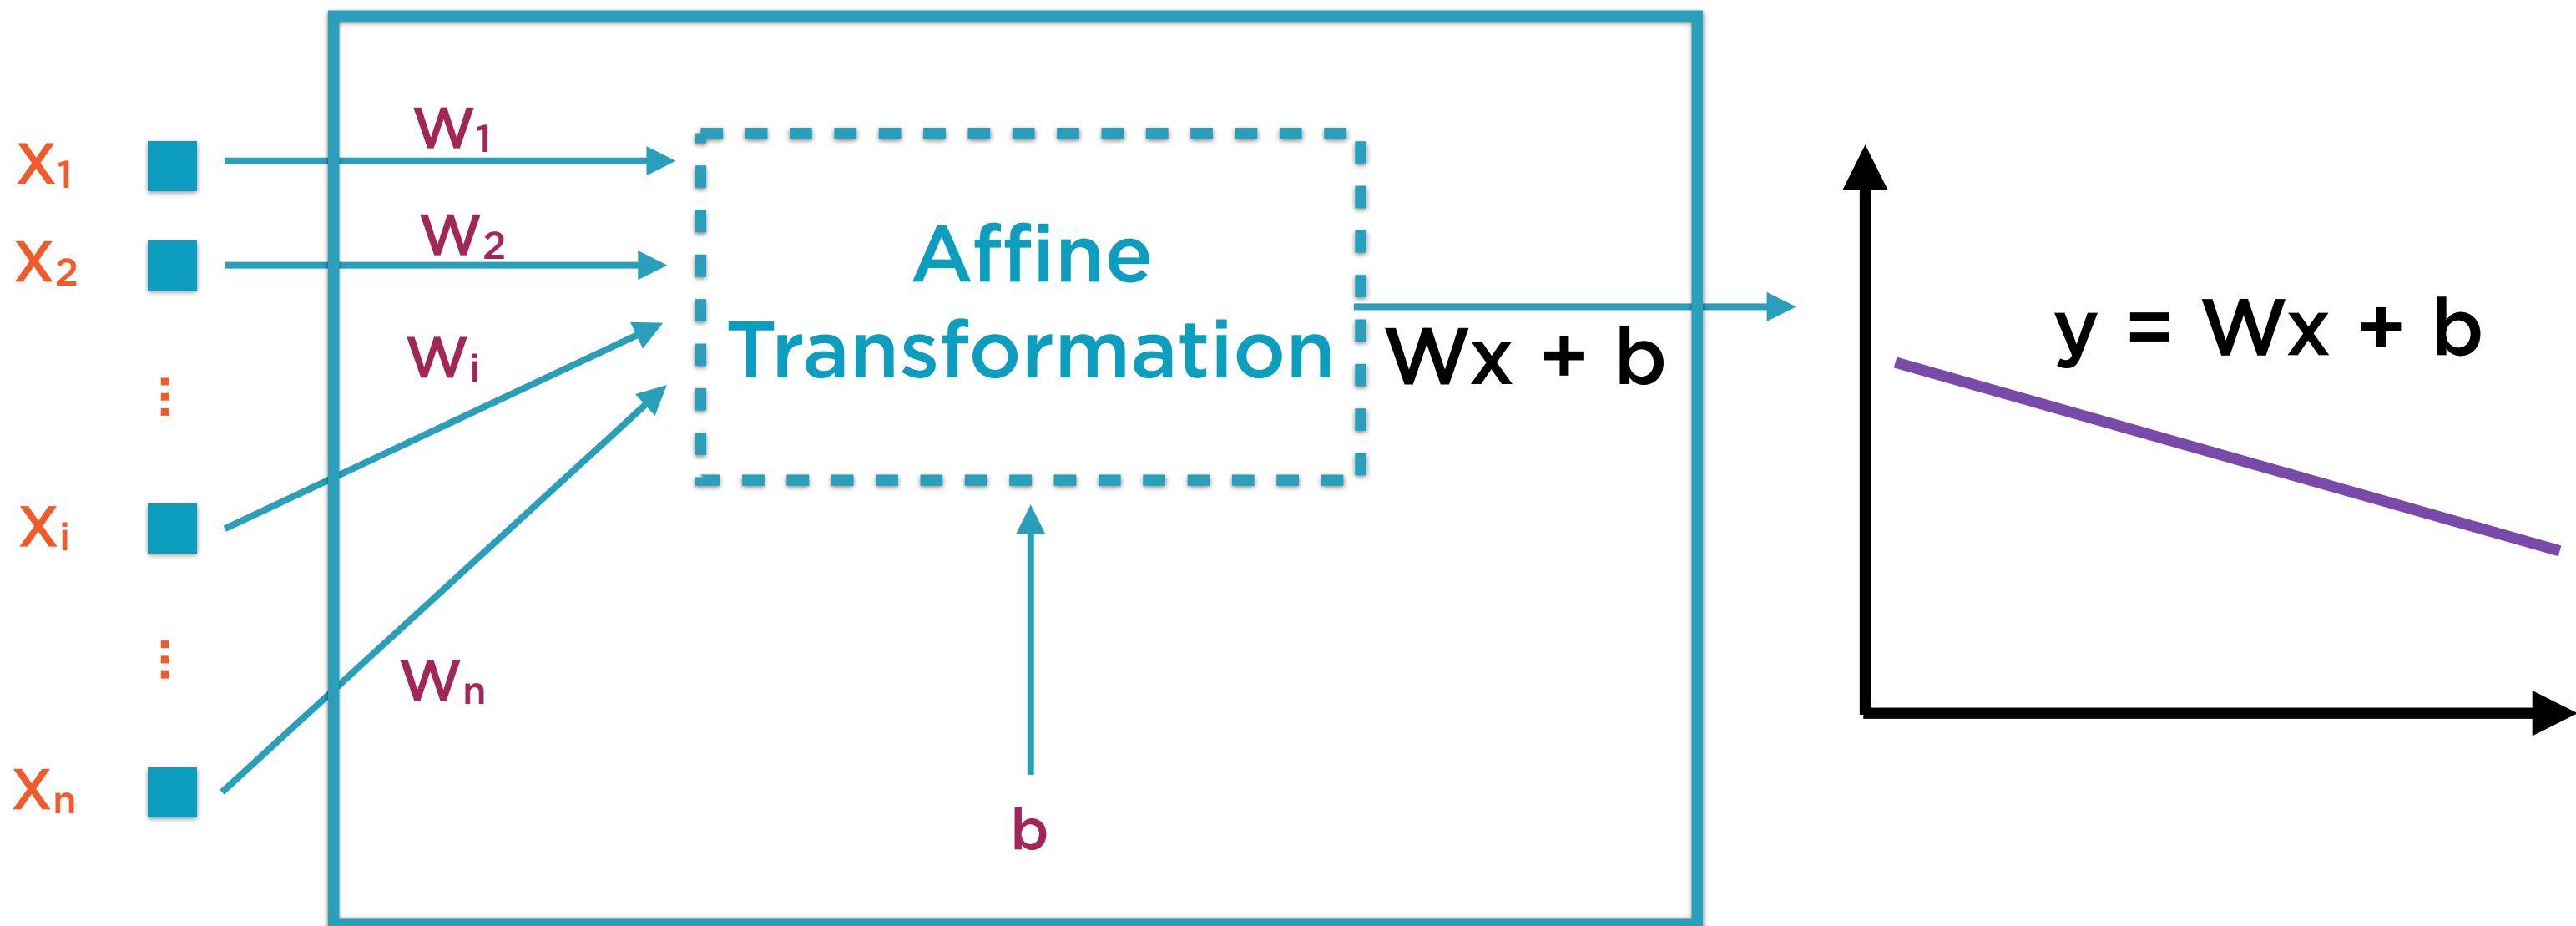
Where do the weights W and the bias b come from?

Operation of a Single Neuron



Where do the weights W and the bias b come from?
They are determined during the training process

Operation of a Single Neuron



This optimization is **not** carried out by the individual neuron, rather a training algorithm will take care of this

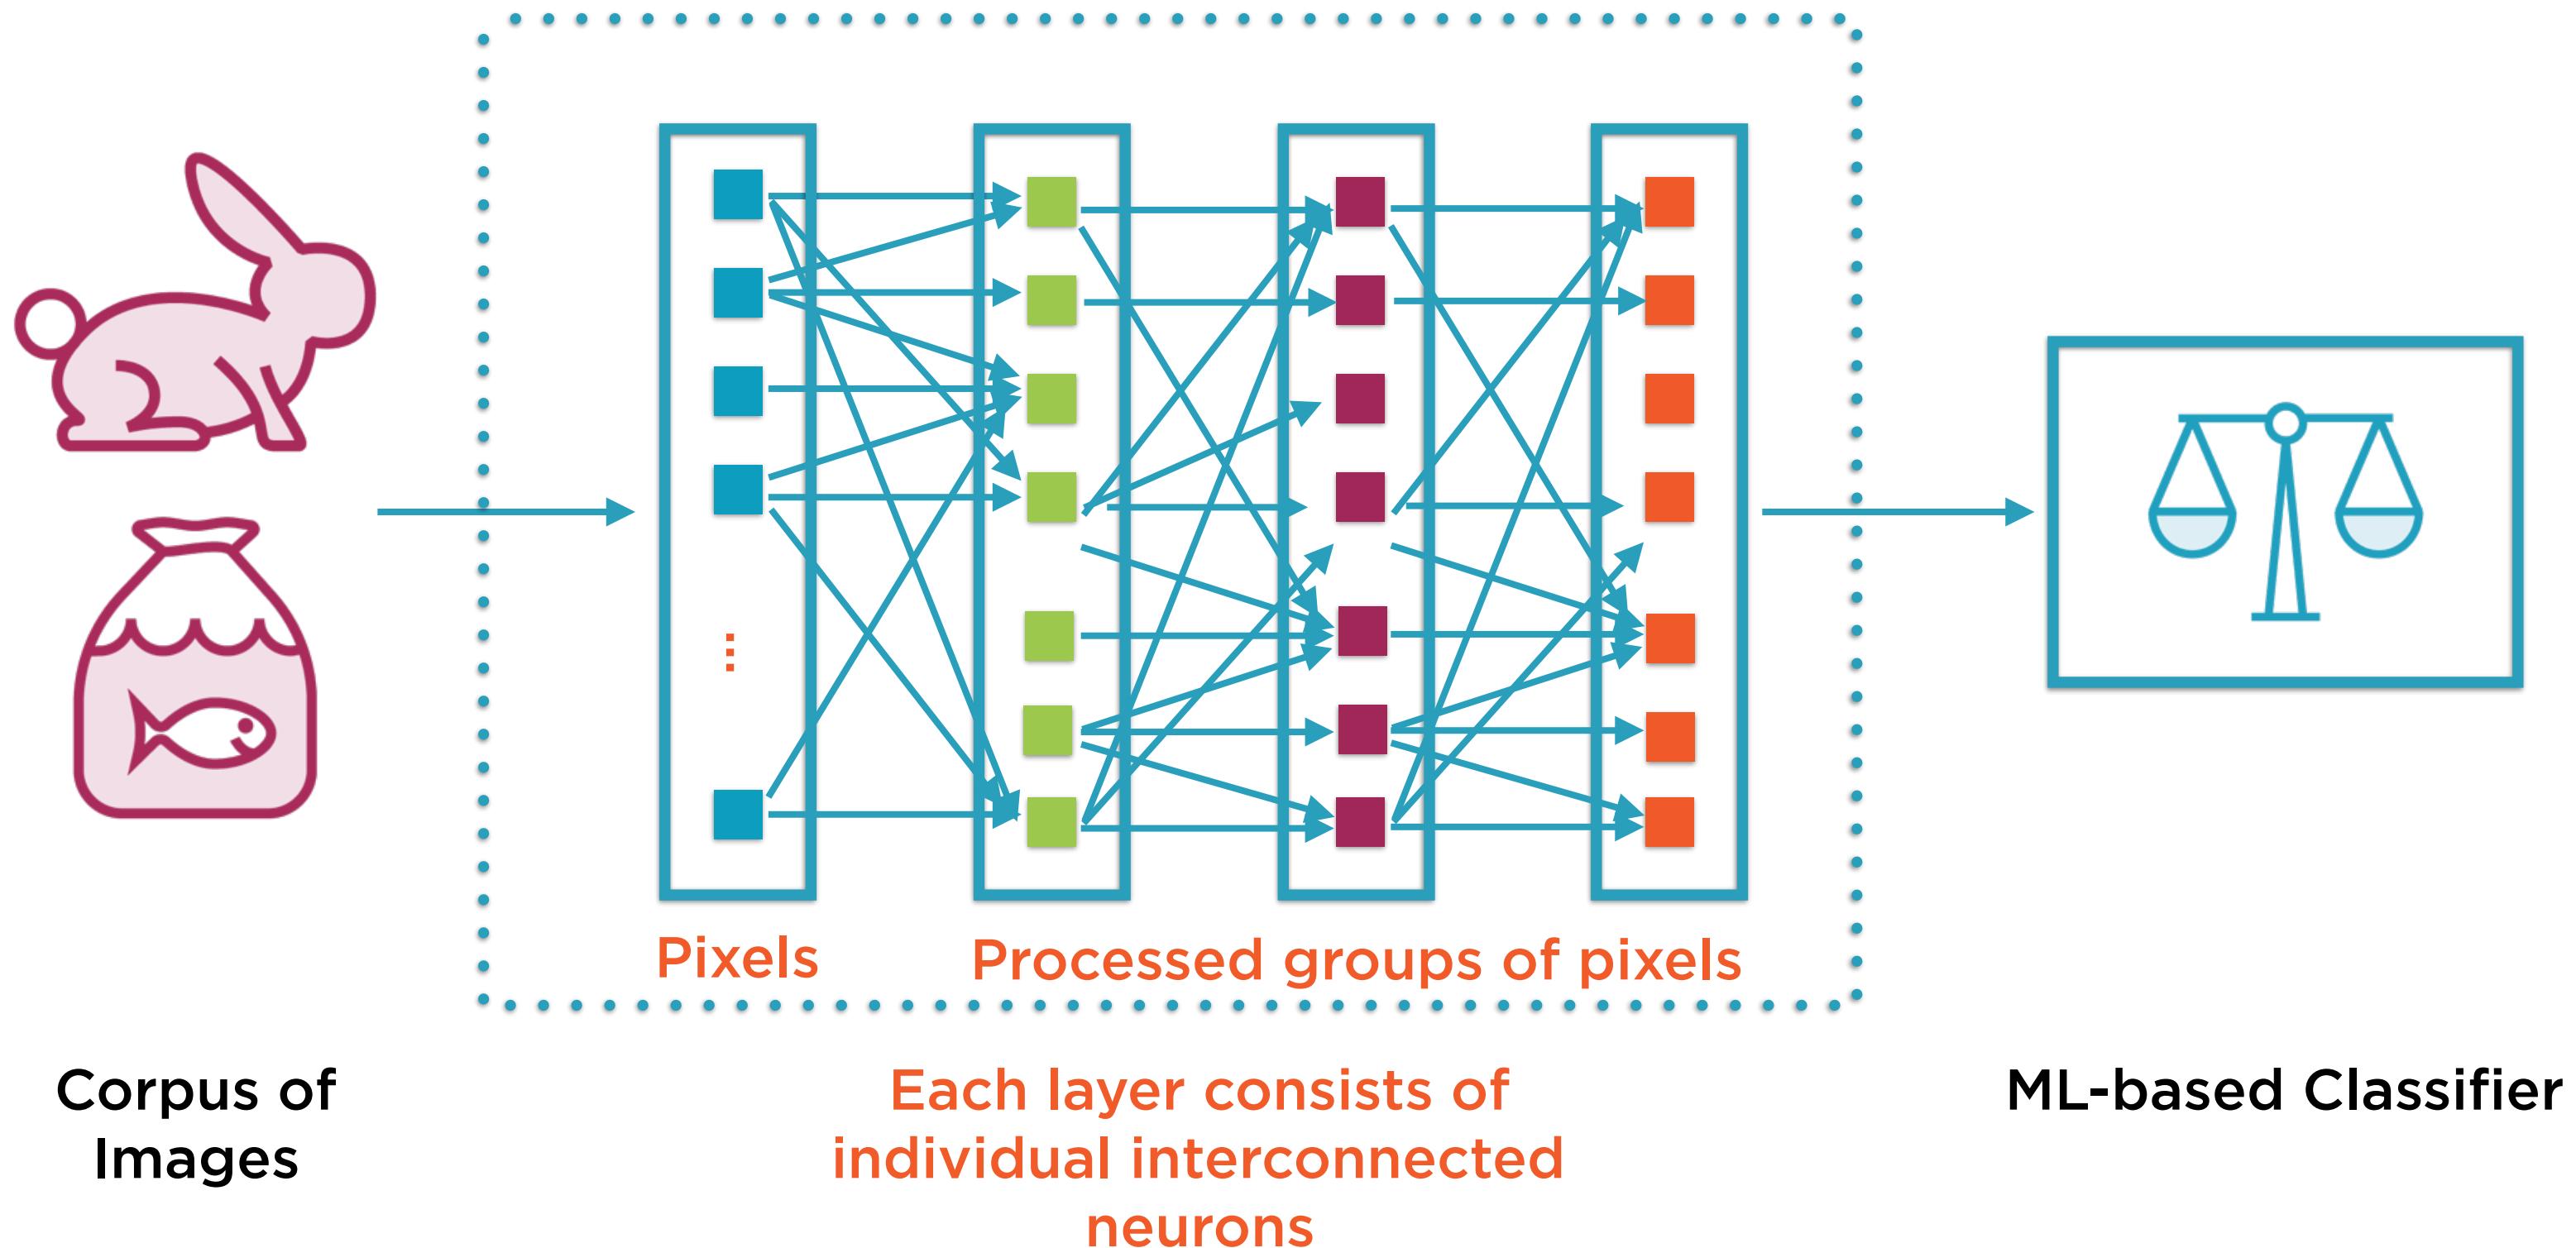
A Slightly More Complex Neural Network

```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

“Learning” Arbitrarily Complex Functions

Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything

An Arbitrarily Complex Function

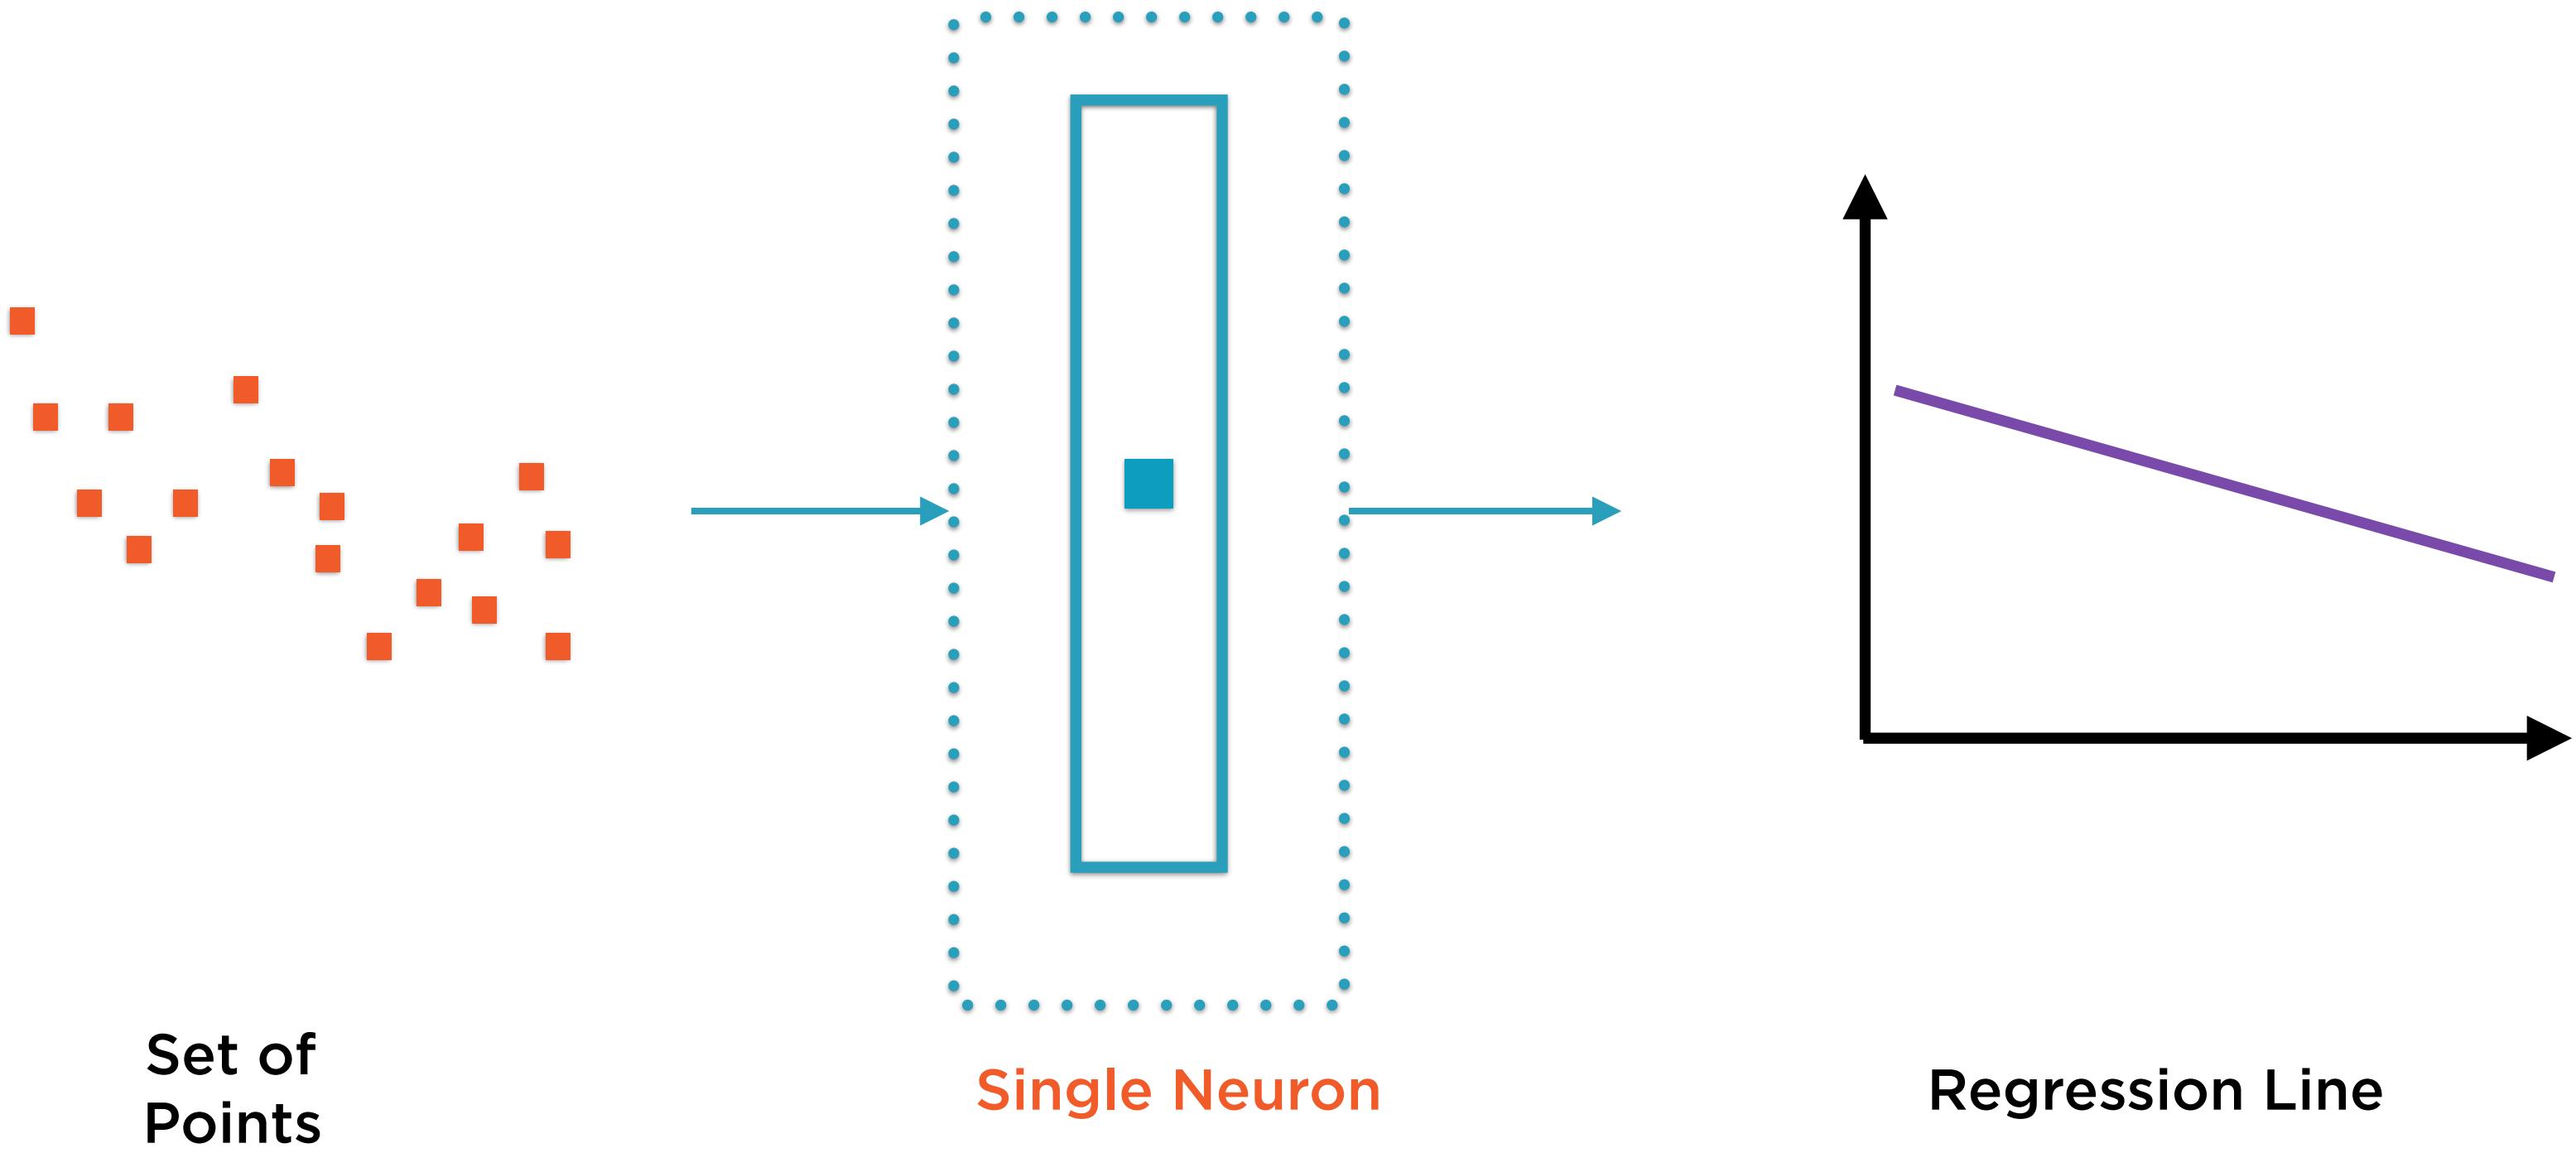


$$y = Wx + b$$

“Learning” Regression

Regression can be learnt by a single neuron using an affine transformation alone

Regression: The Simplest Neural Network

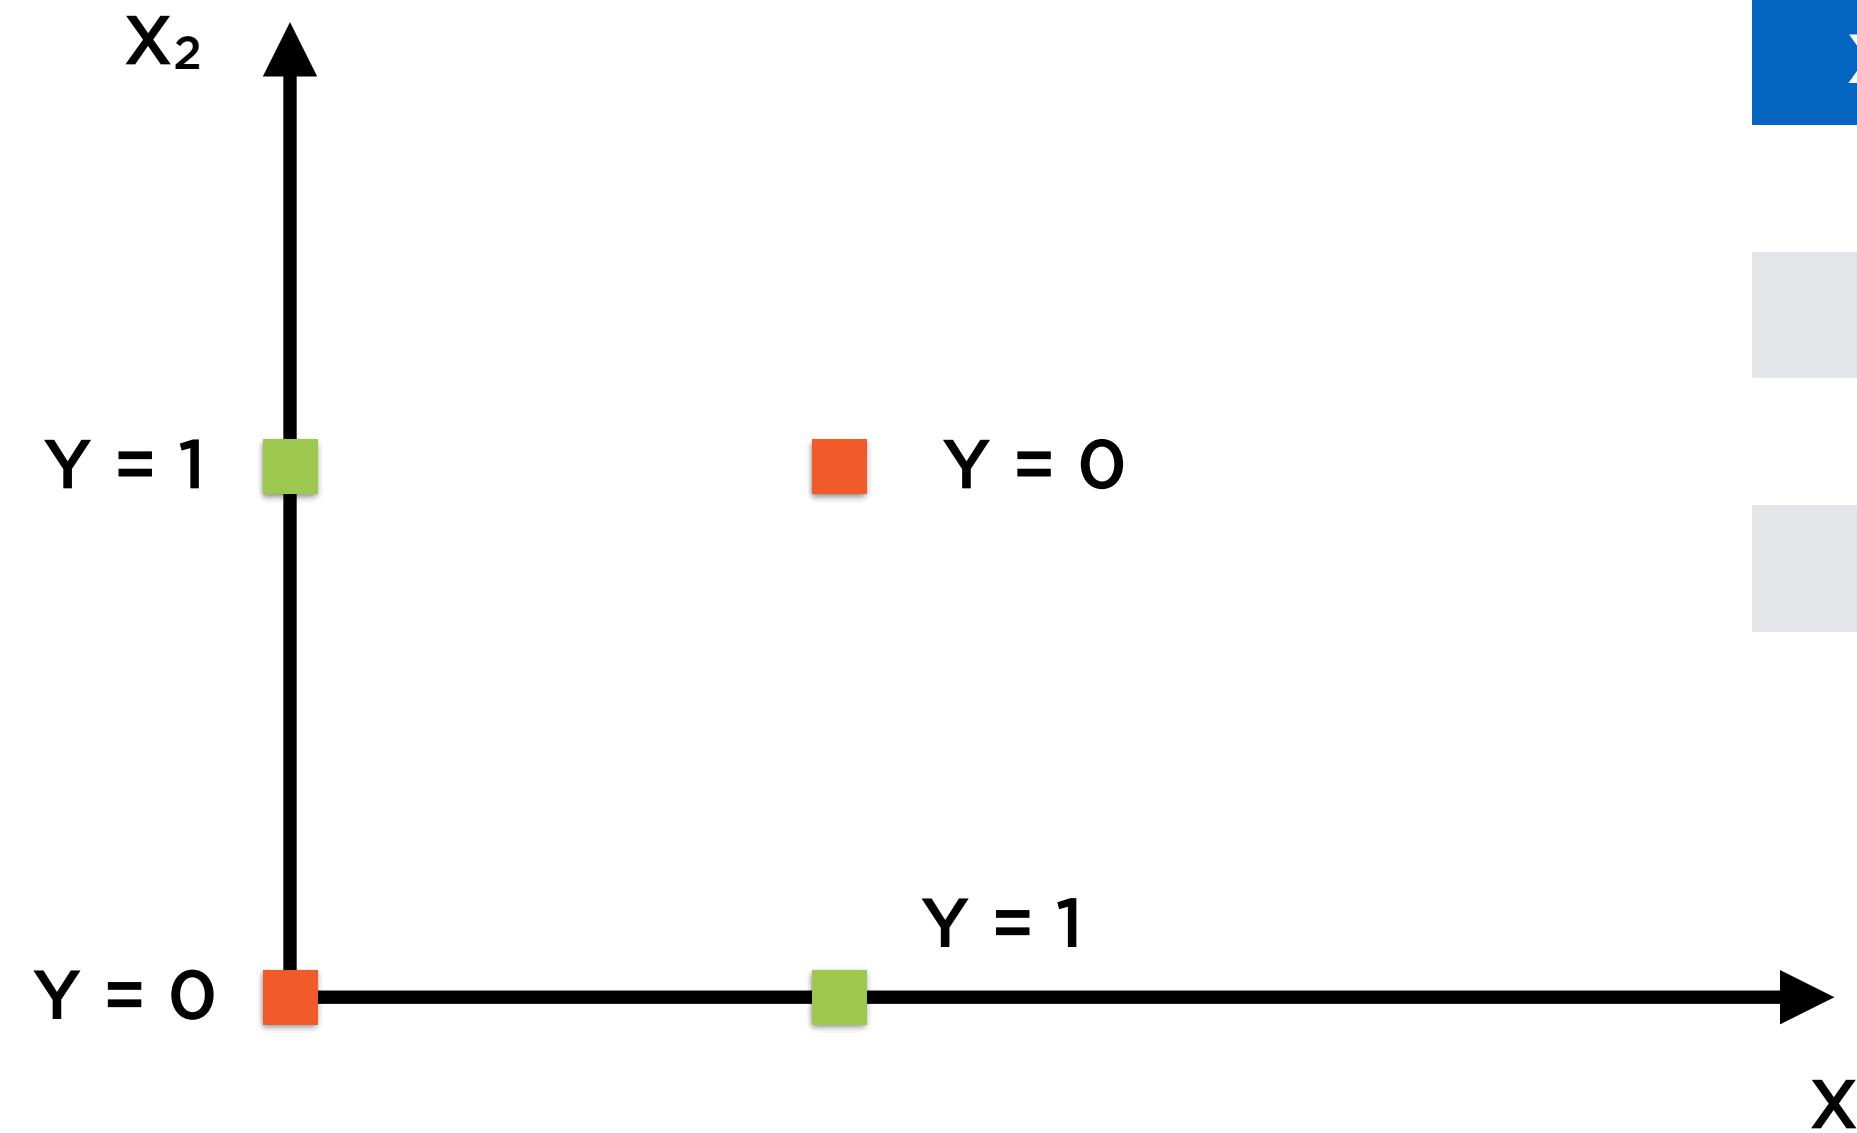


```
def XOR(x1, x2):  
    if (x1 == x2):  
        return 0  
    return 1
```

“Learning” XOR

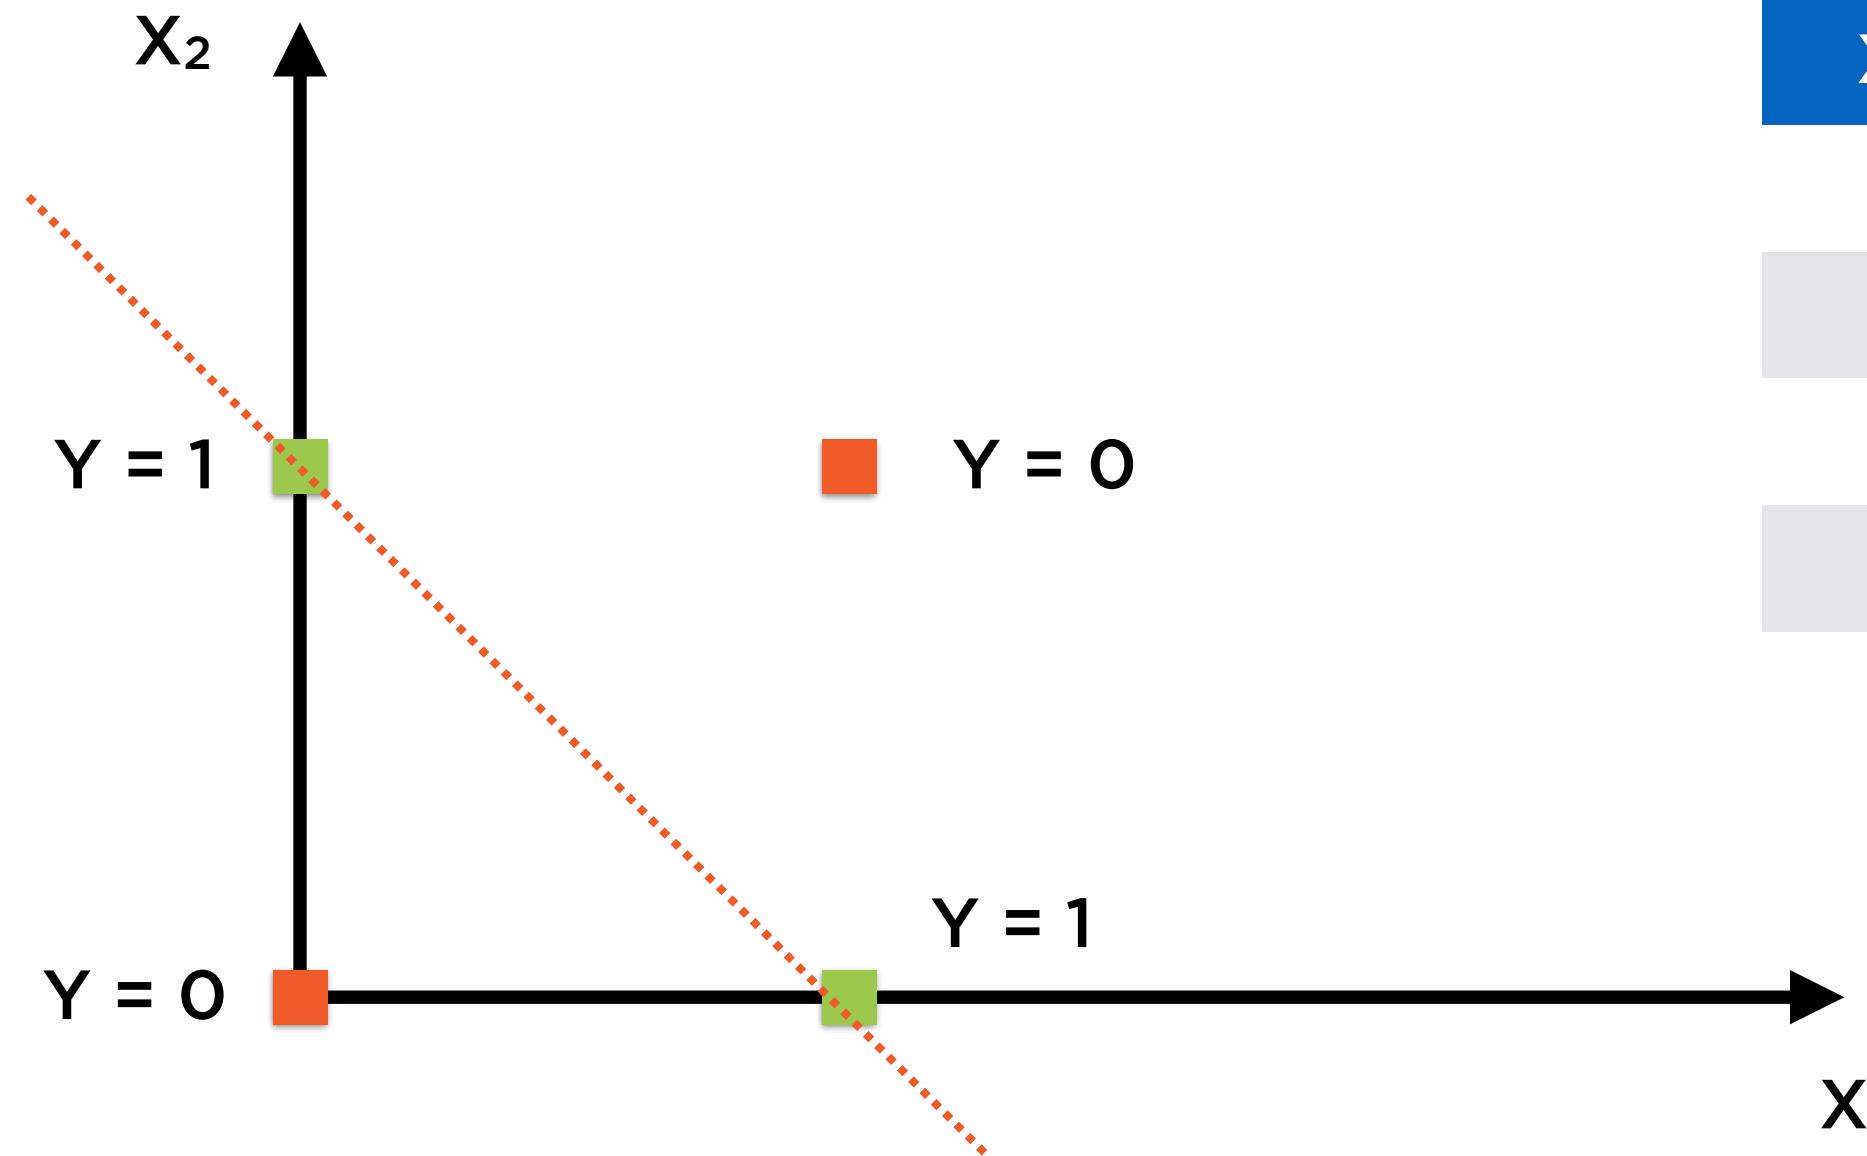
Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

XOR: Not Linearly Separable



X_1	X_2	Y
0	0	0
0	1	1
1	0	1
1	1	0

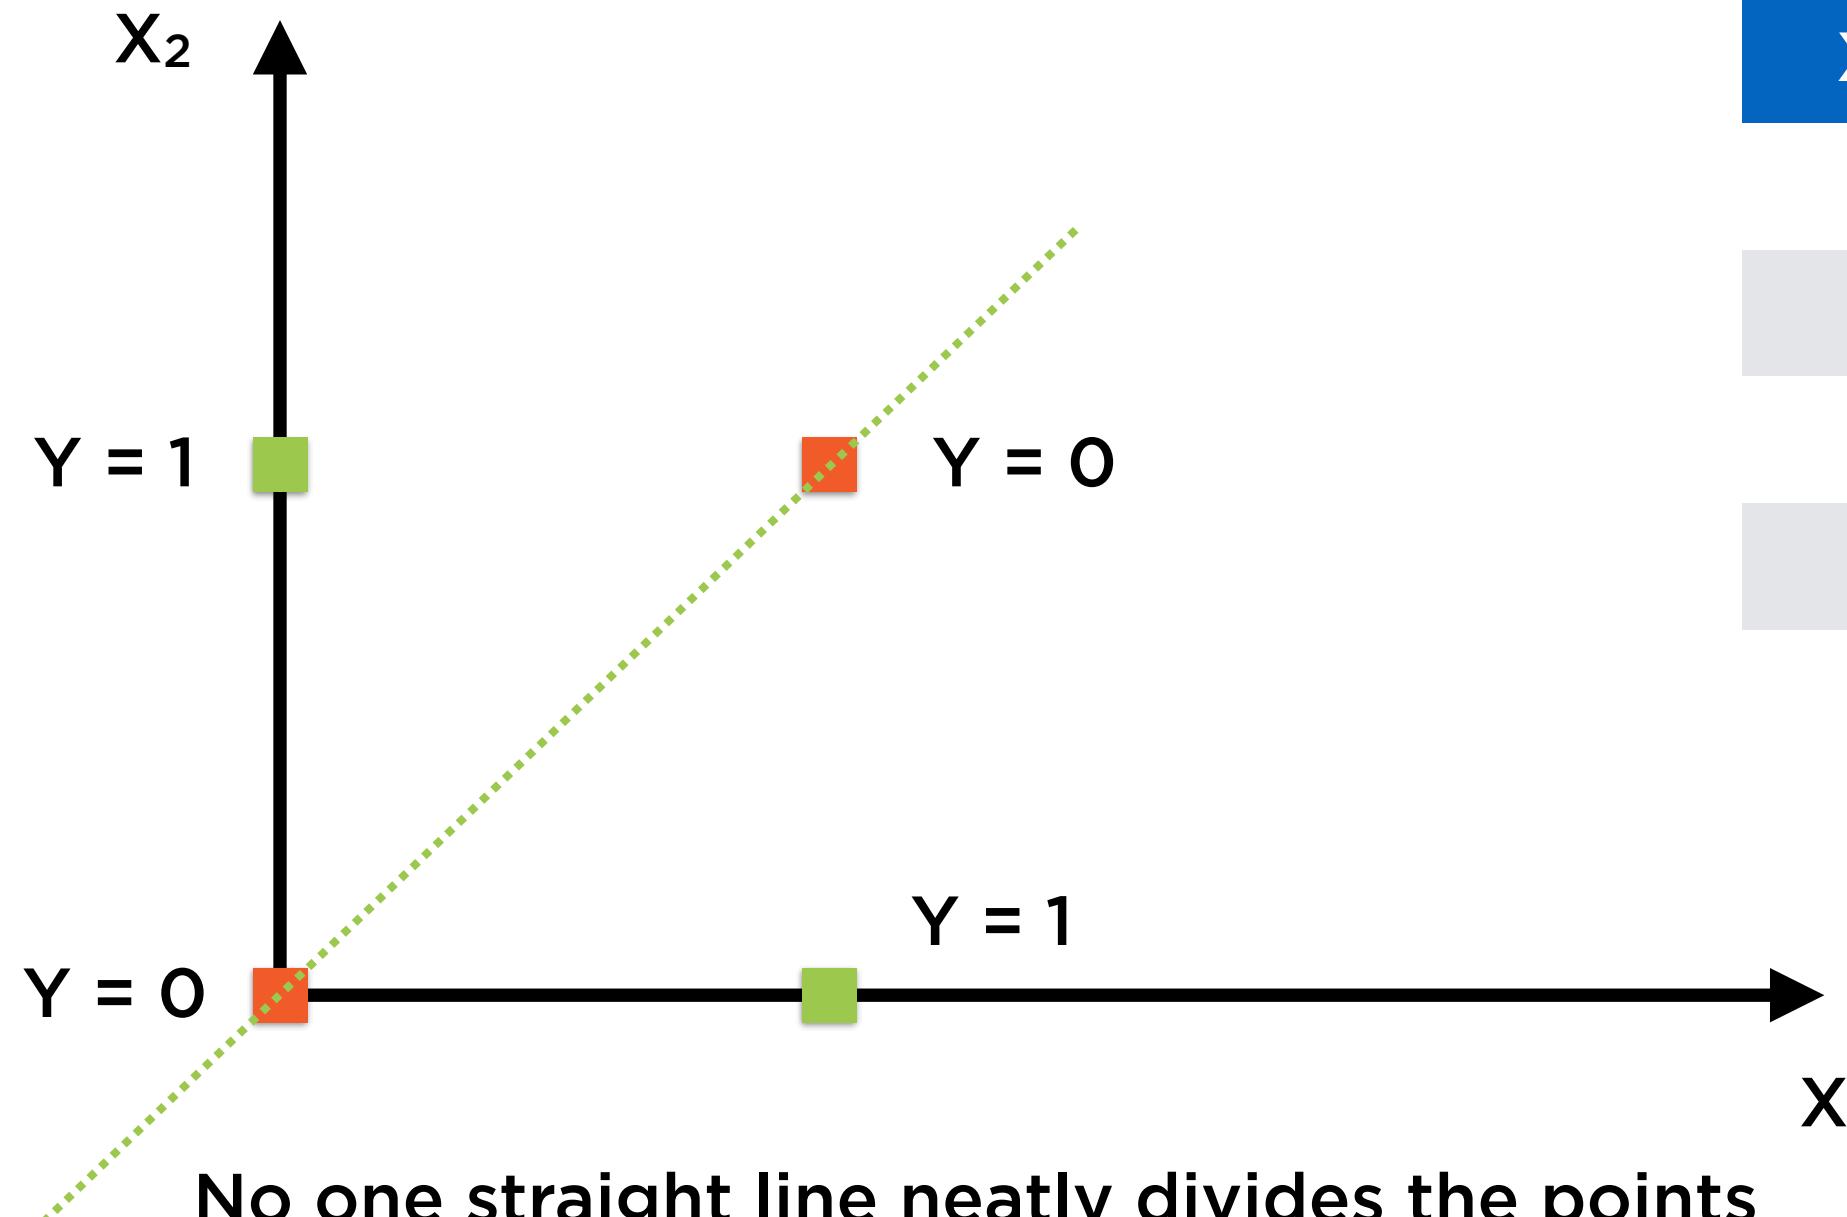
XOR: Not Linearly Separable



X_1	X_2	Y
0	0	0
0	1	1
1	0	1
1	1	0

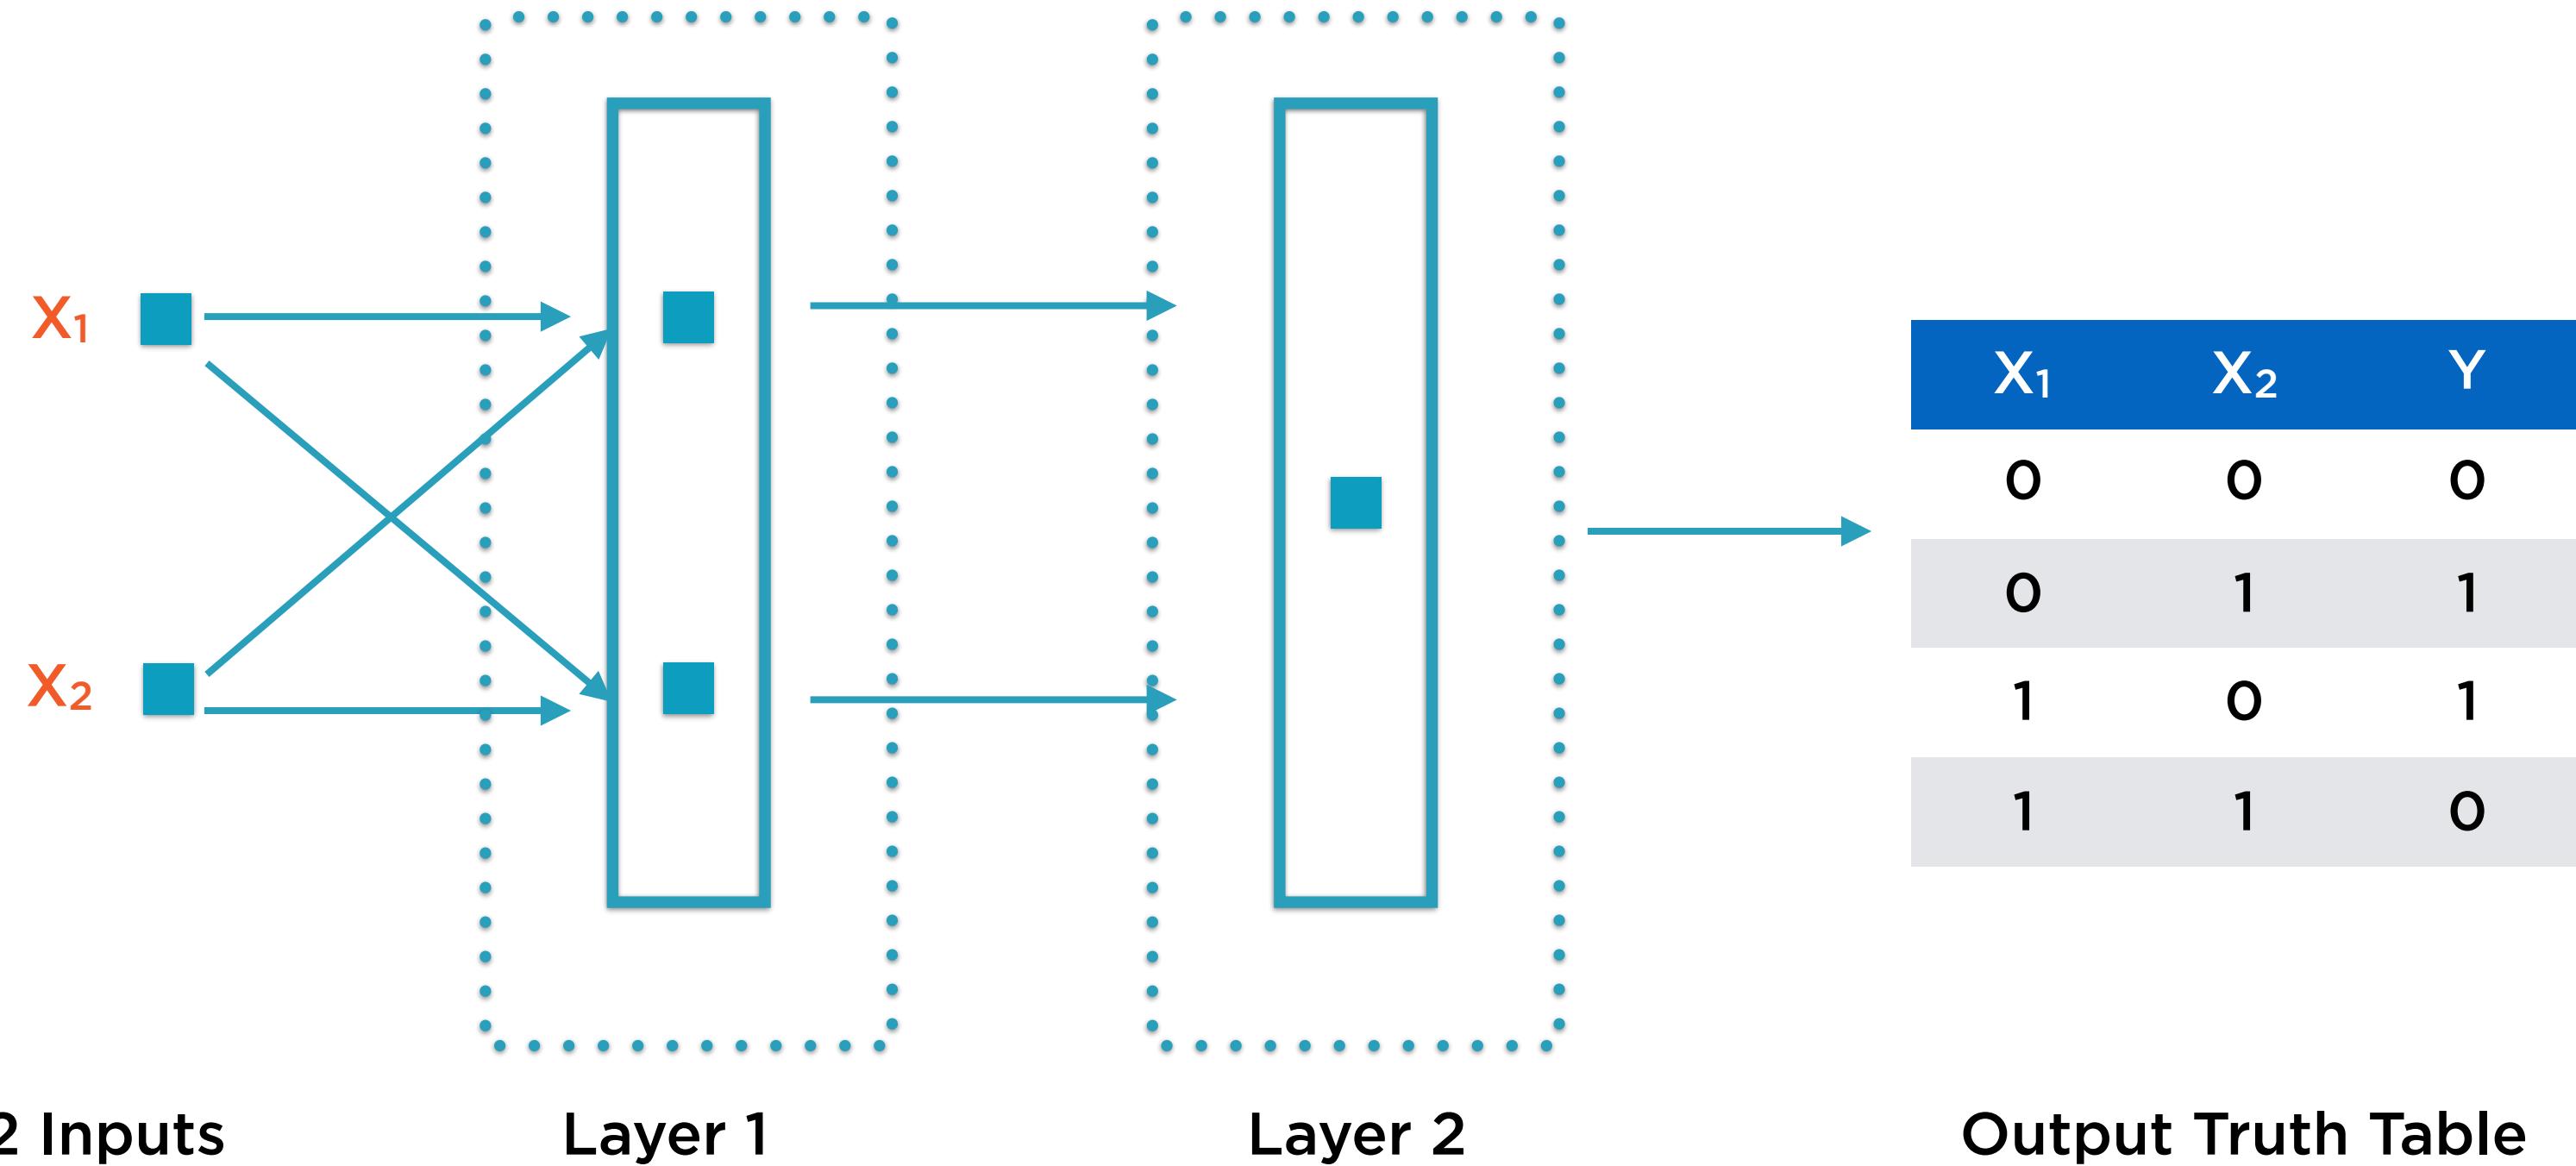
No one straight line neatly divides the points into disjoint regions where $Y = 0$ and $Y = 1$

XOR: Not Linearly Separable



X_1	X_2	Y
0	0	0
0	1	1
1	0	1
1	1	0

XOR: 3 Neurons, 2 Layers

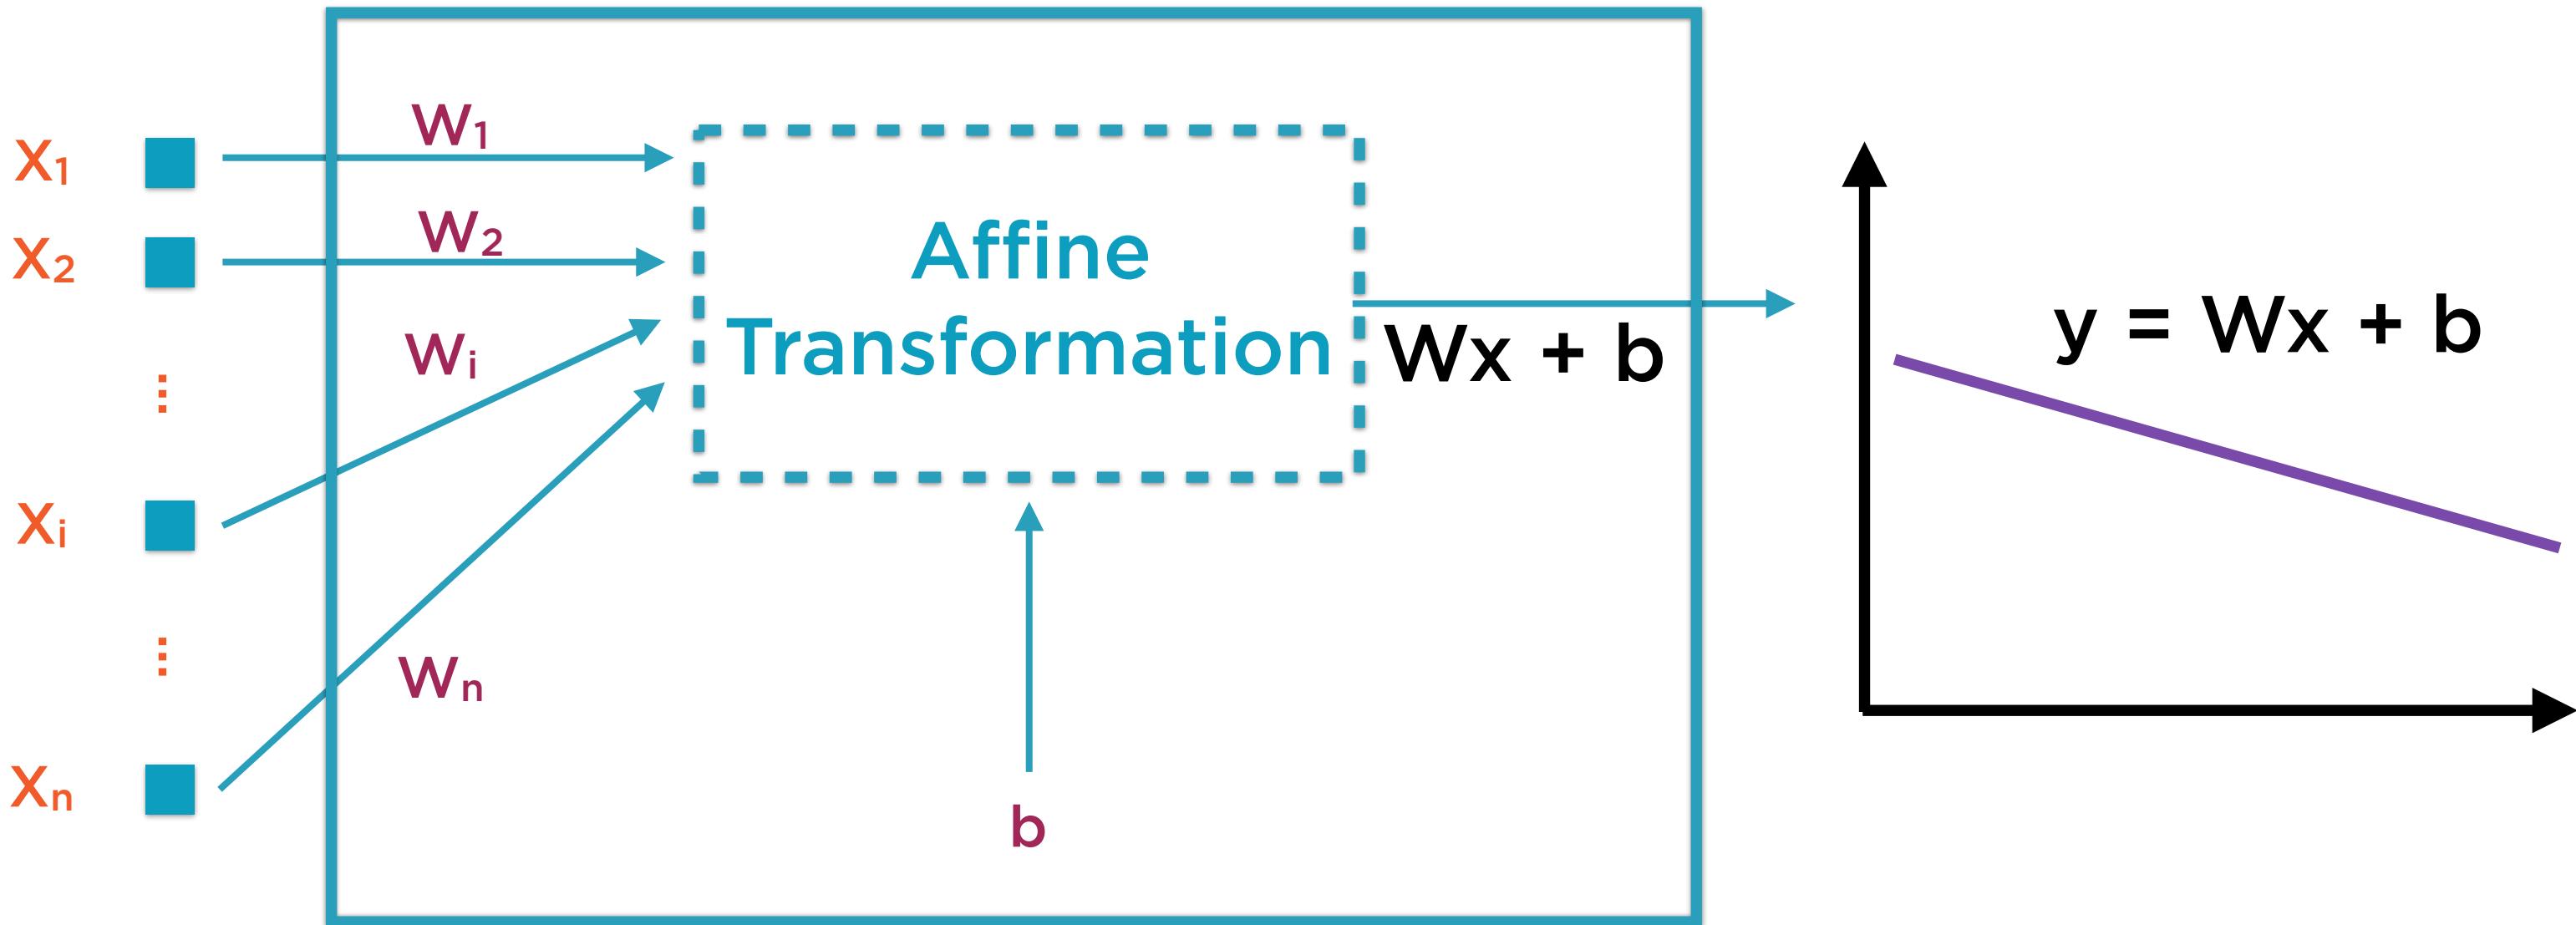


X ₁	X ₂	Y
0	0	0
0	1	1
1	0	1
1	1	0

“Learning” XOR

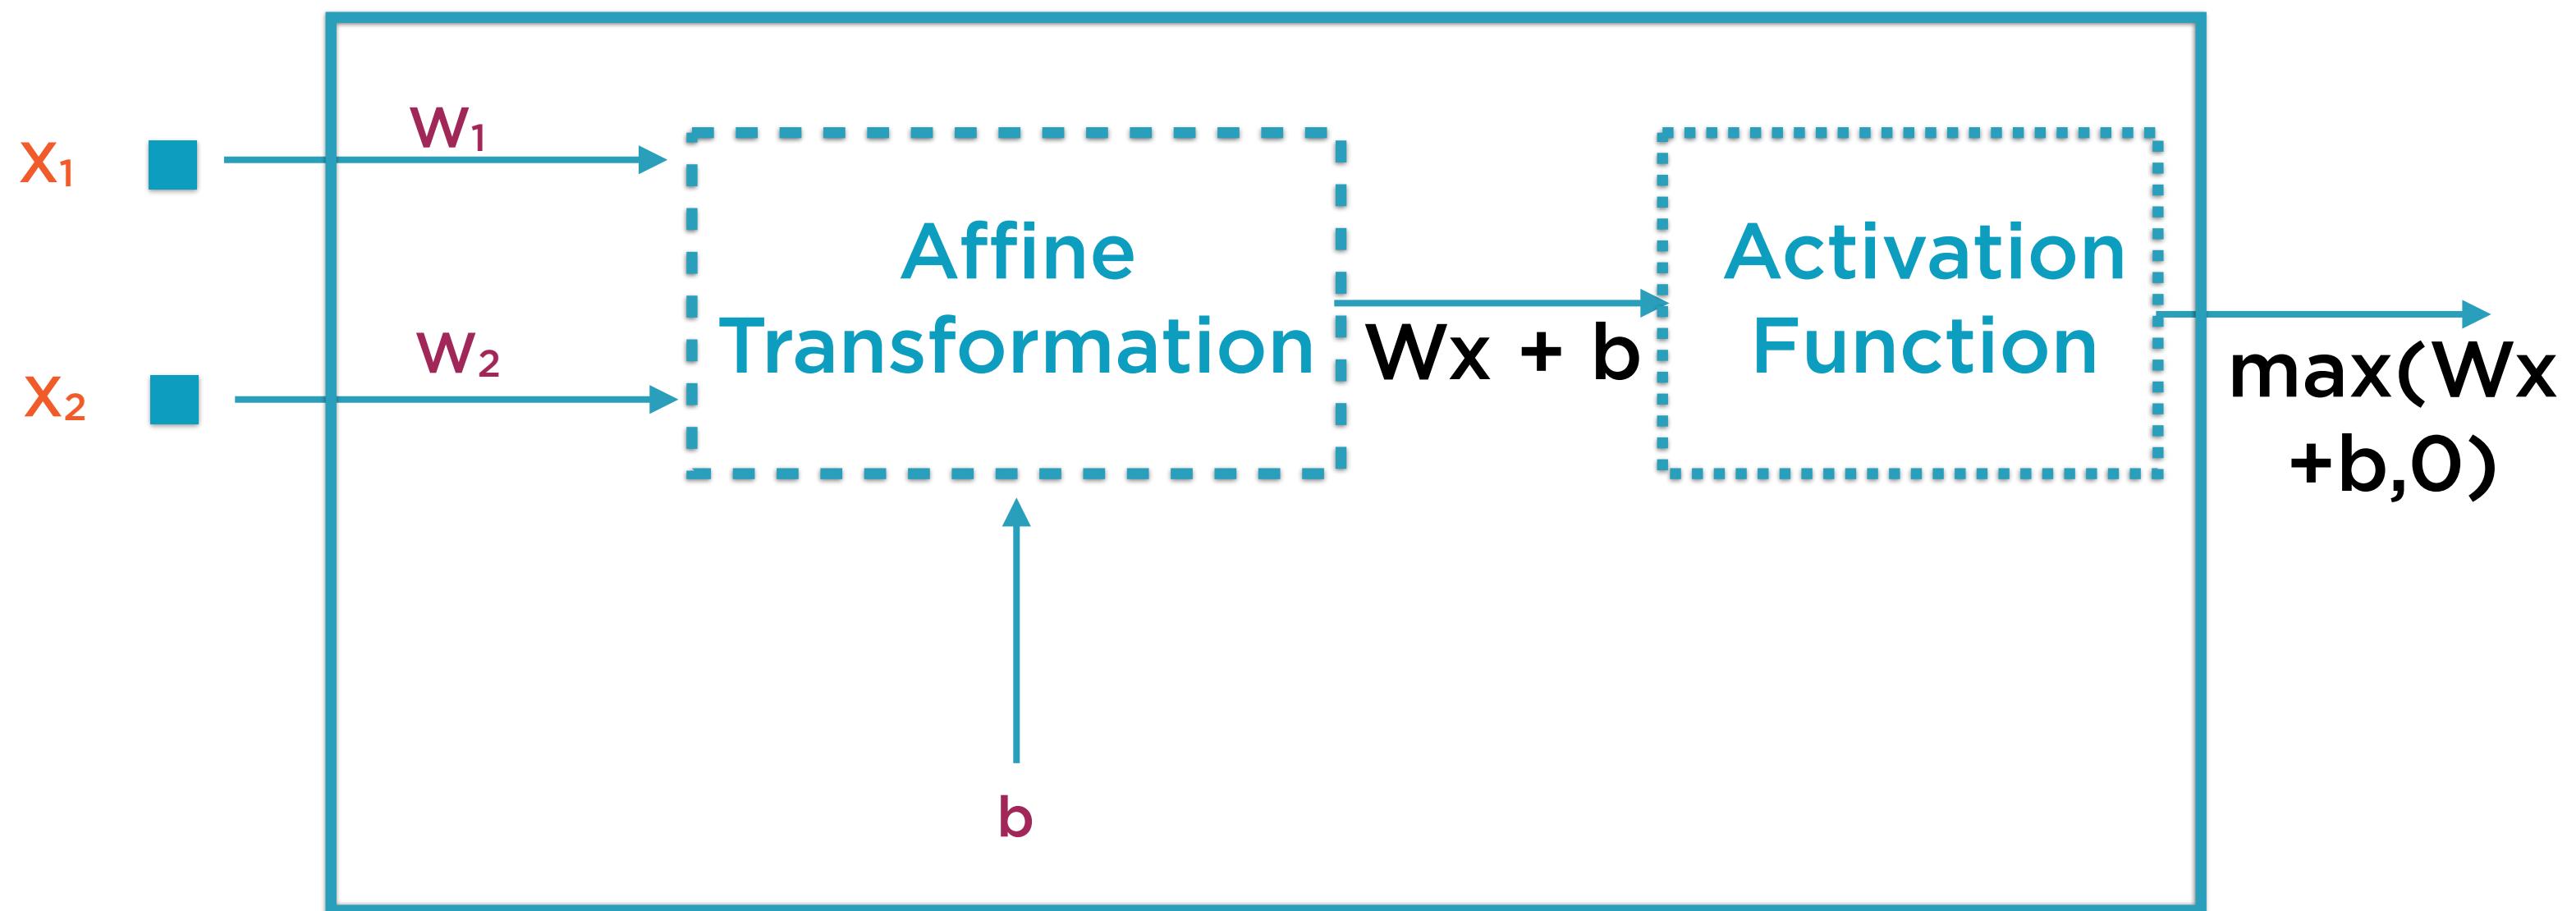
Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

1-Neuron Regression



Regression could be learnt by a single neuron using a
single, linear operation

Adding an Activation Function



XOR, a simple non-linear function, can be learnt by 3 neurons if we add an appropriate activation function



The activation function is needed for the neural network to predict non-linear functions

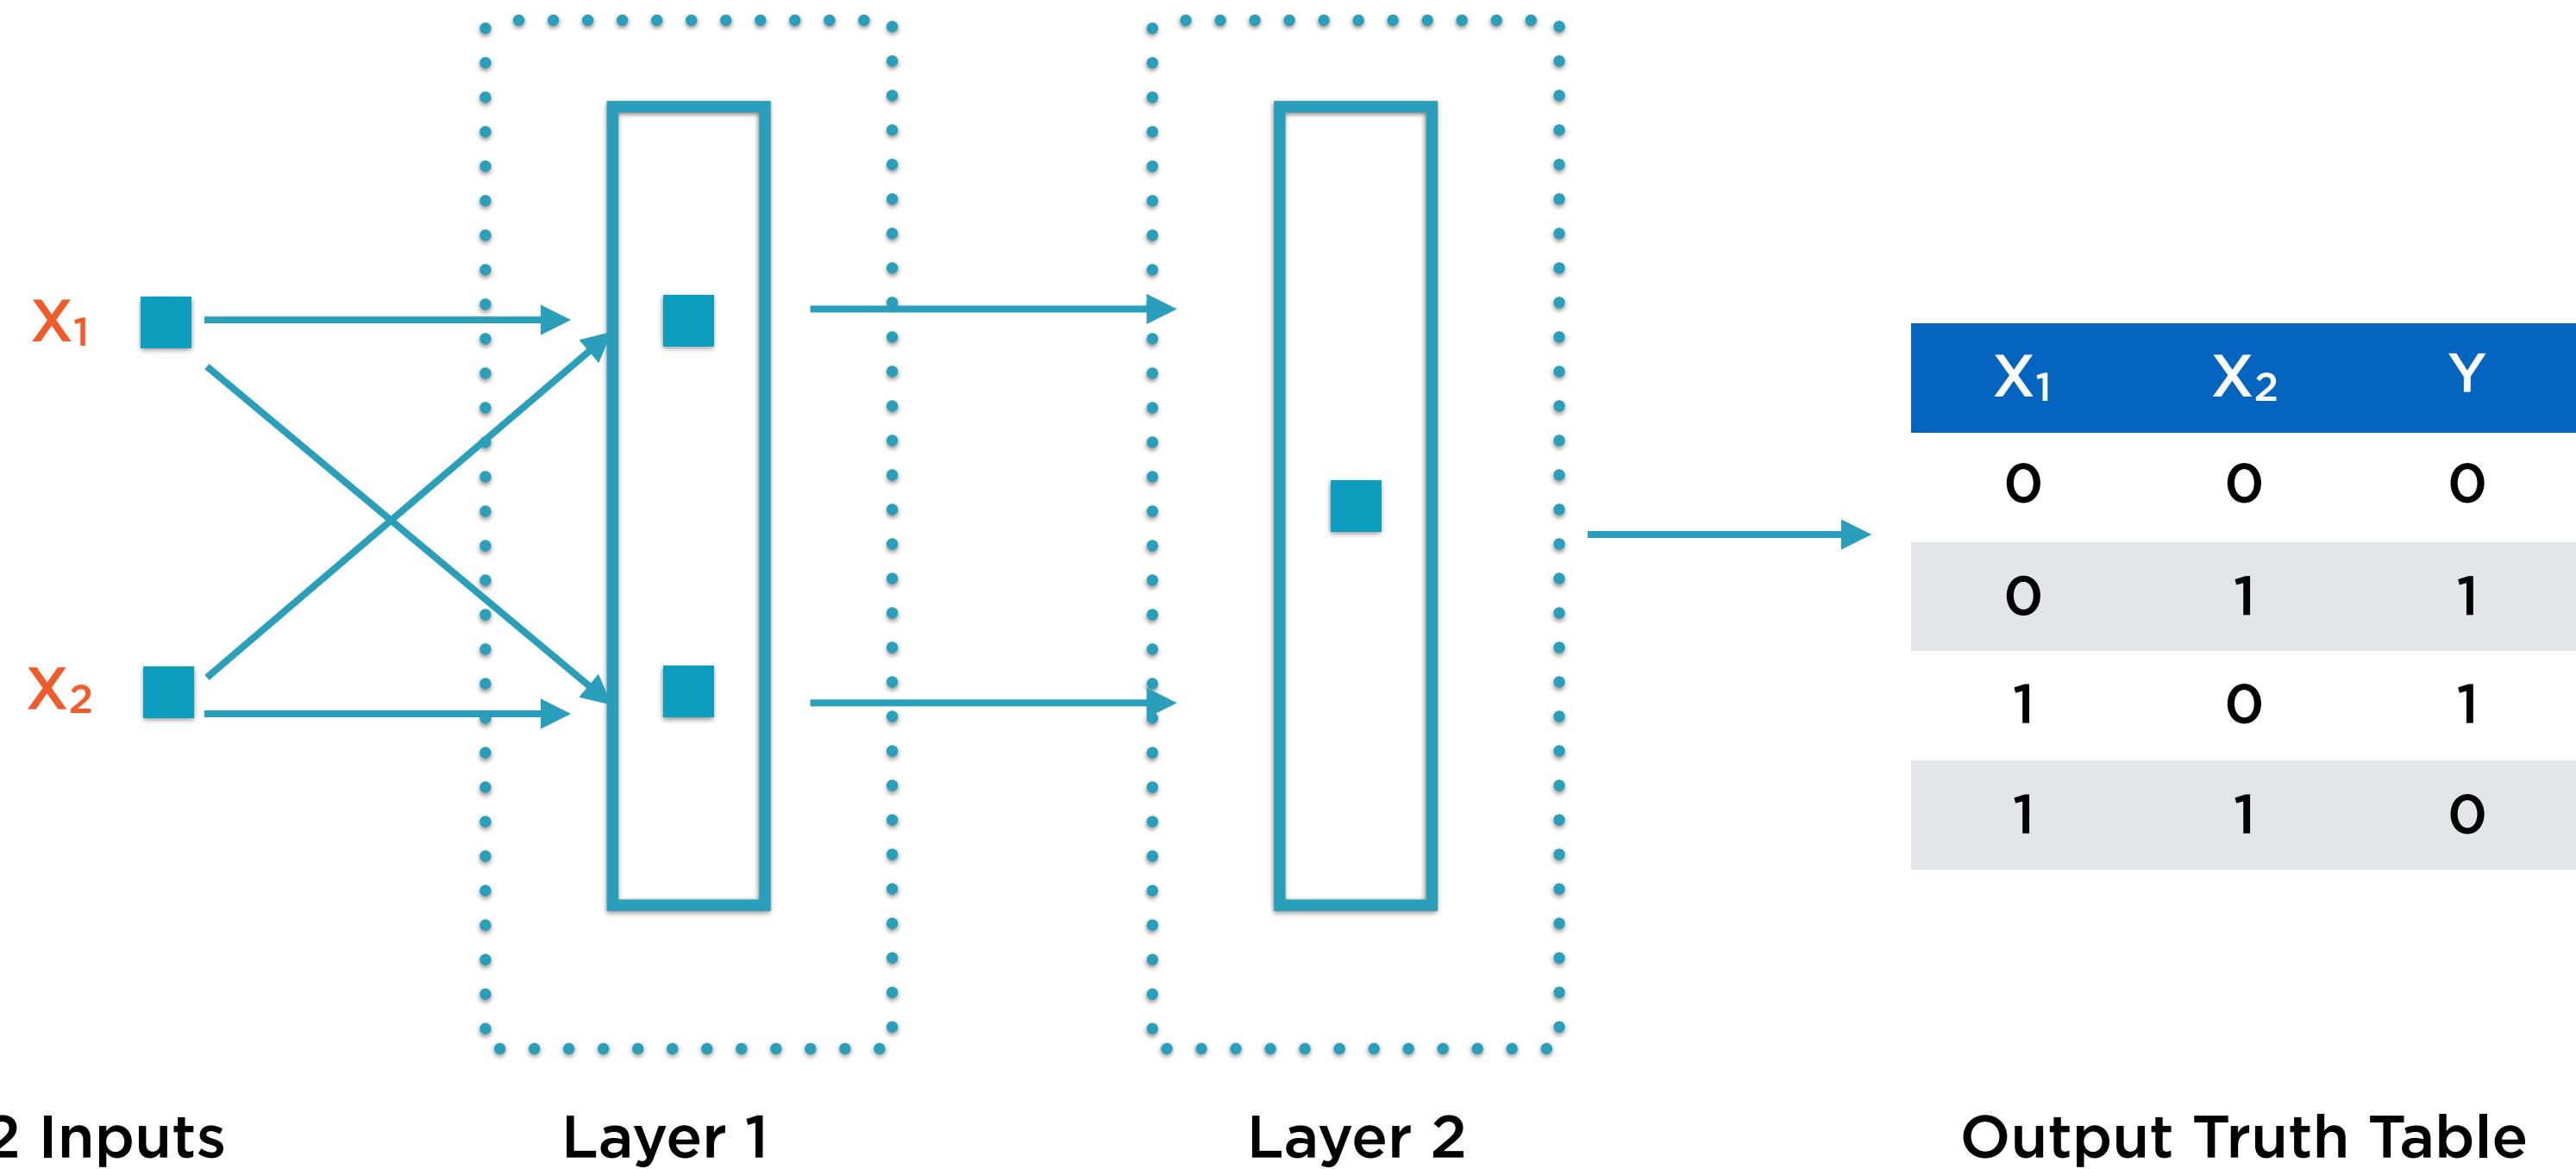


The most common form of the activation function is the ReLU

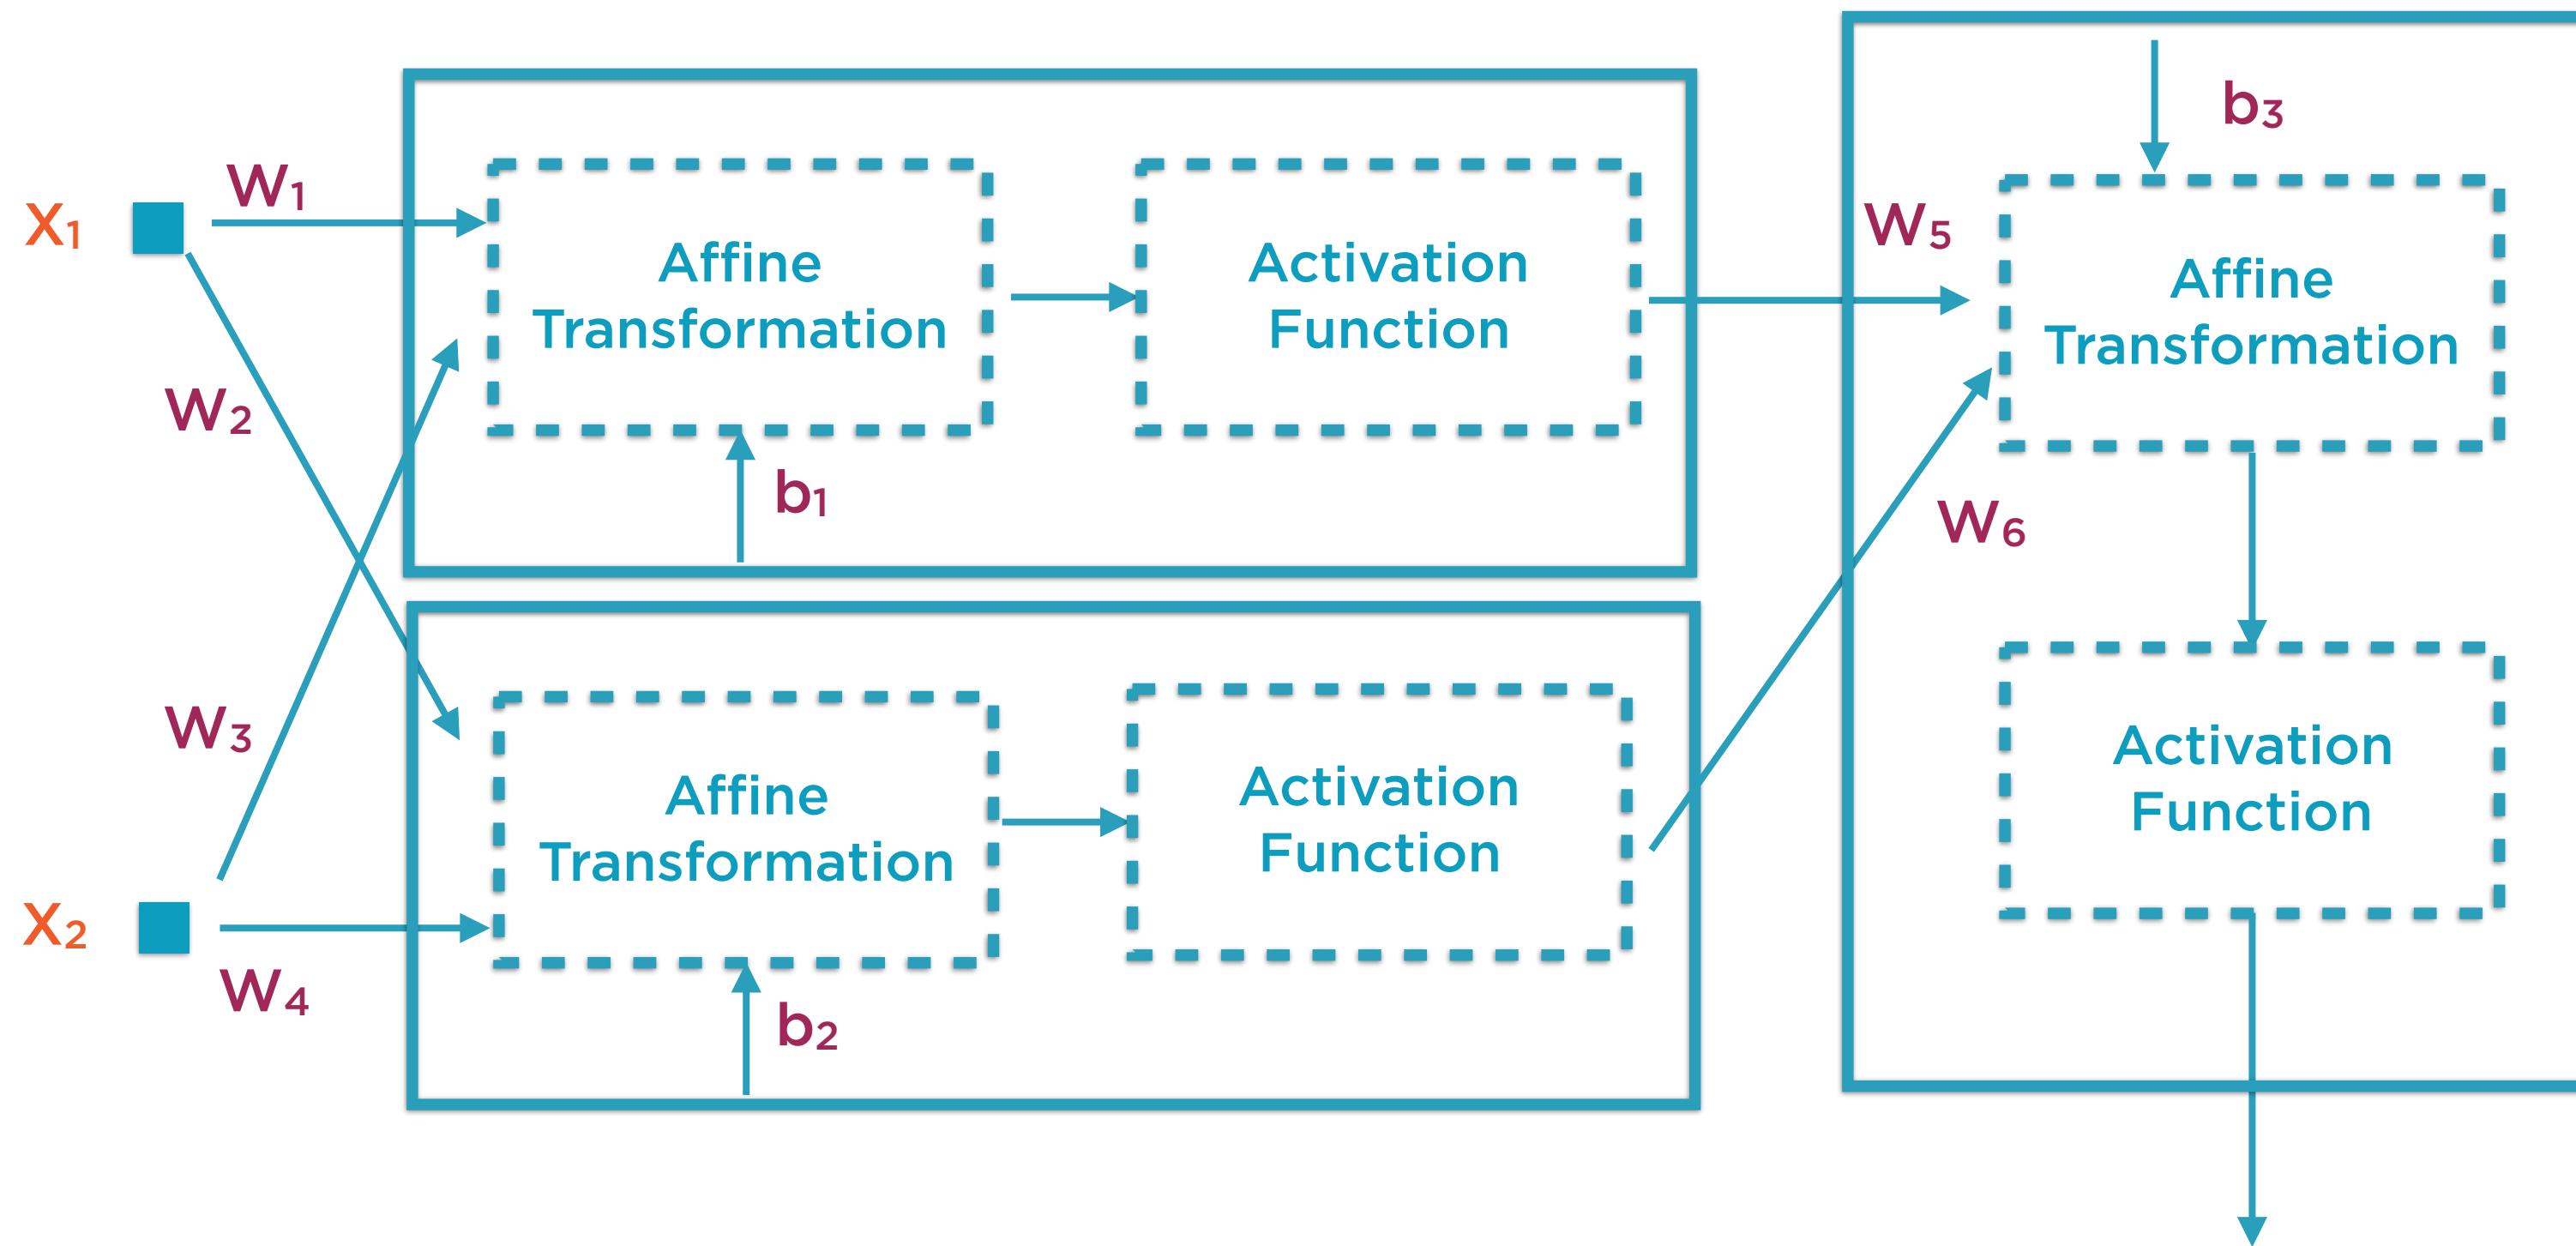
ReLU : Rectified Linear Unit

$$\text{ReLU}(x) = \max(0, x)$$

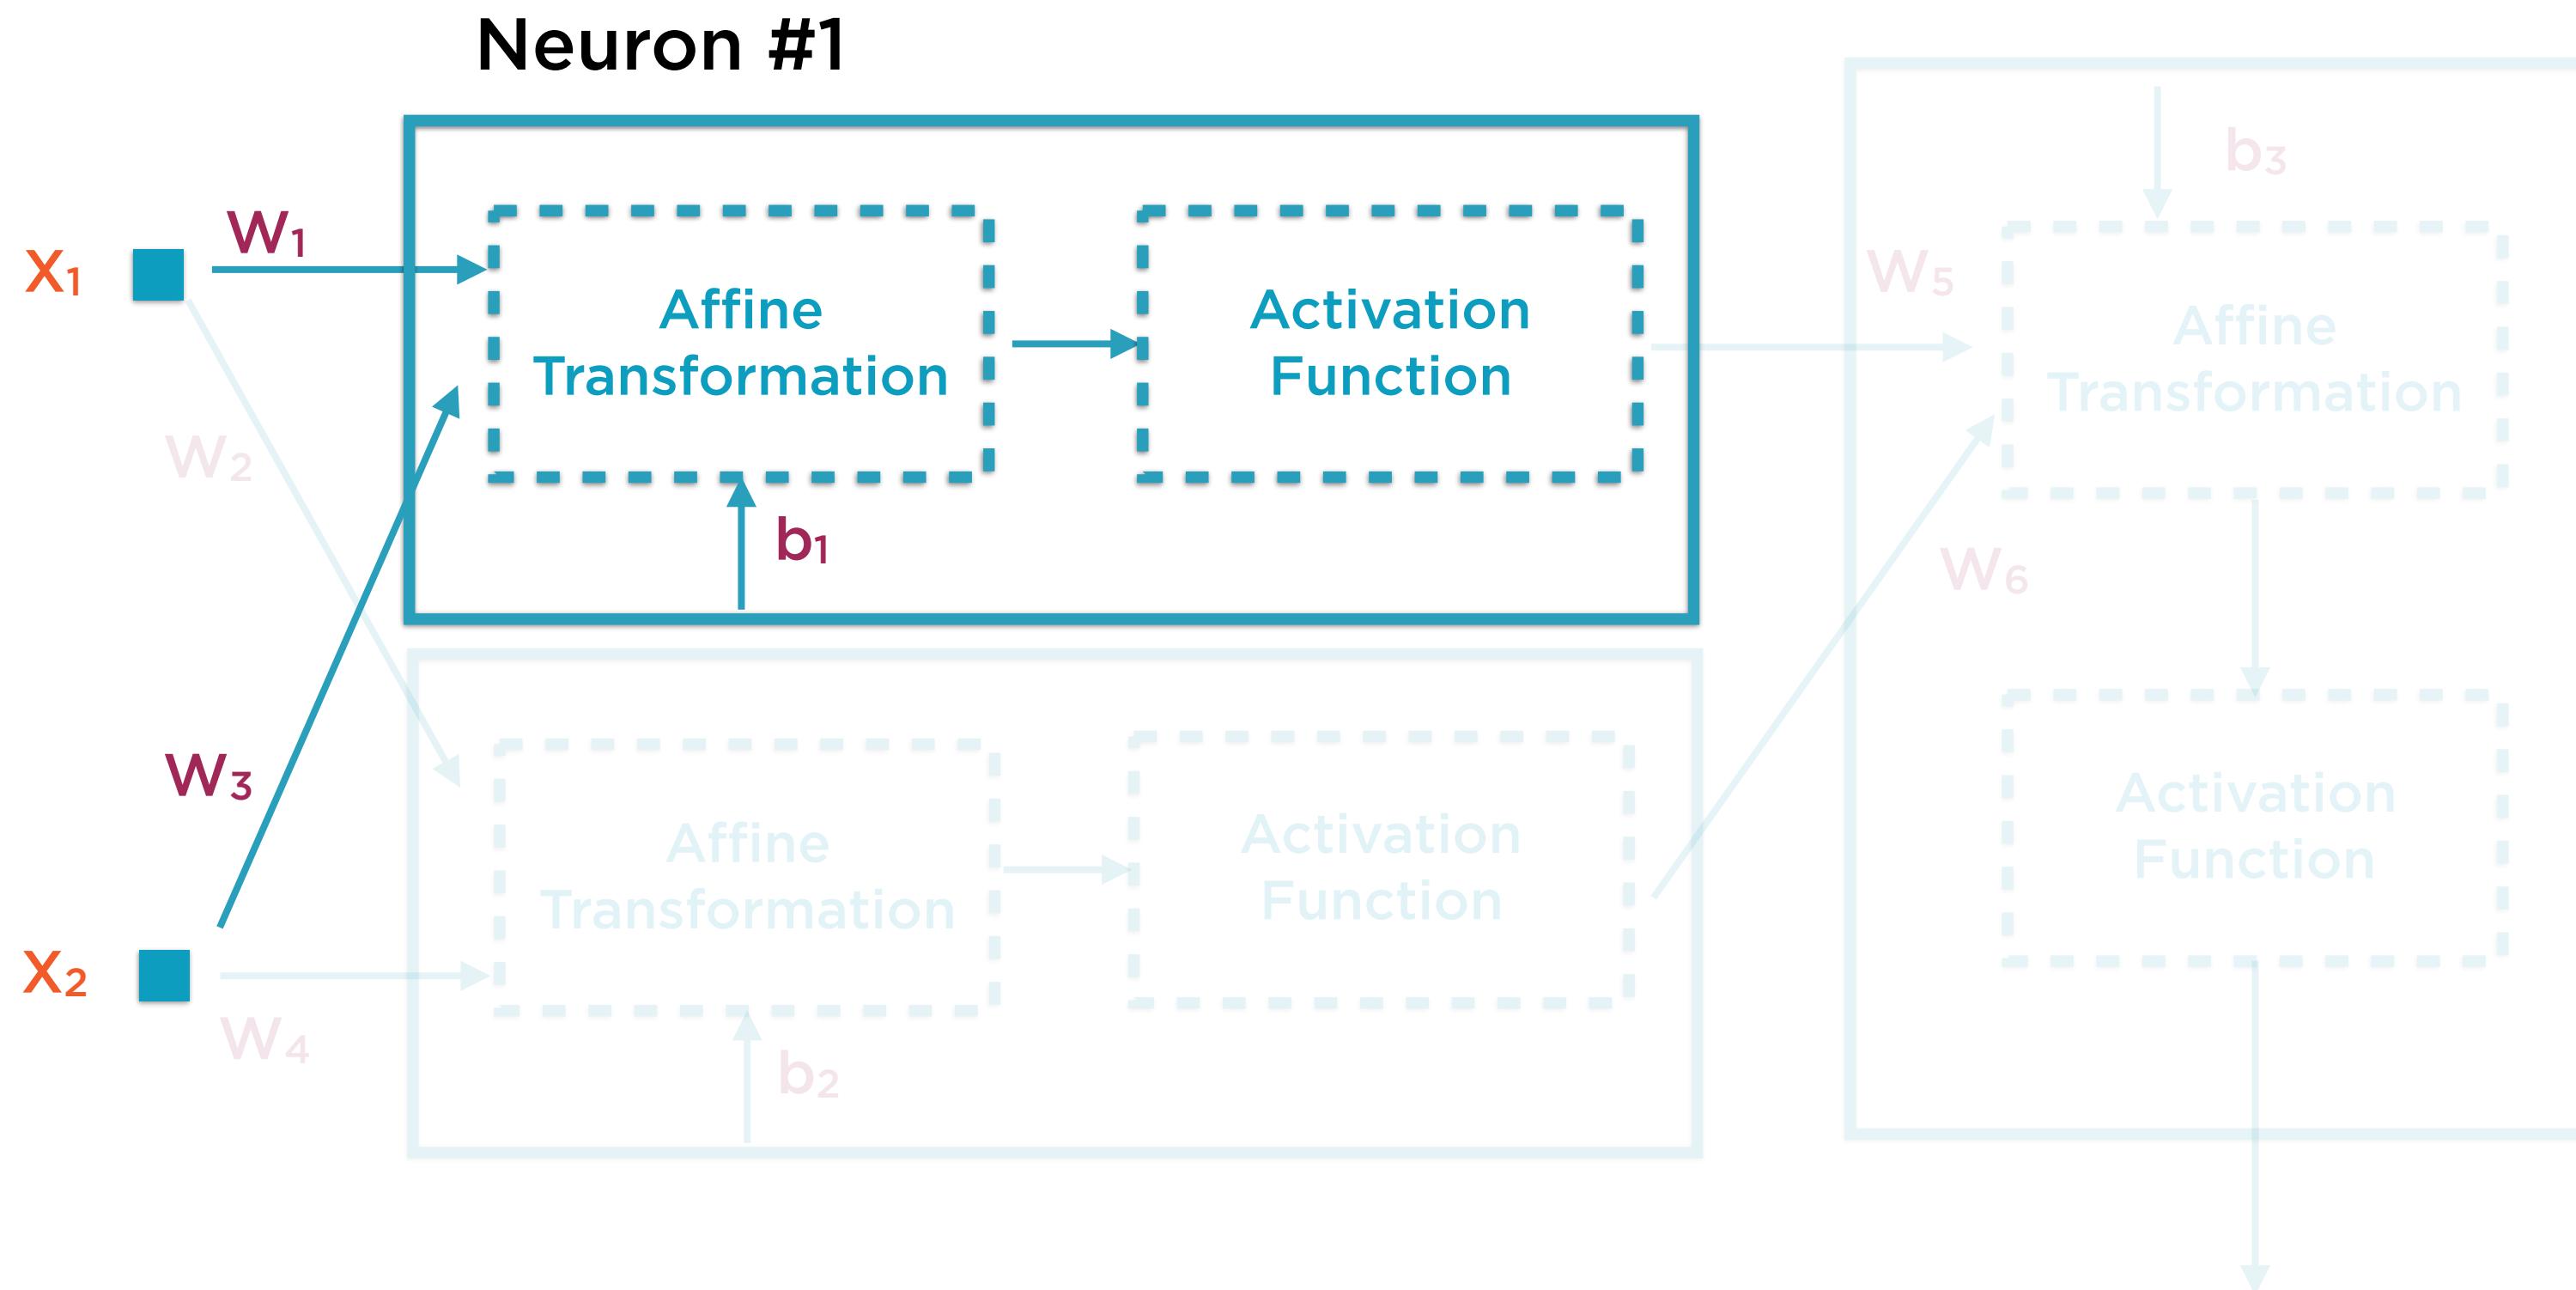
XOR: 3 Neurons, 2 Layers



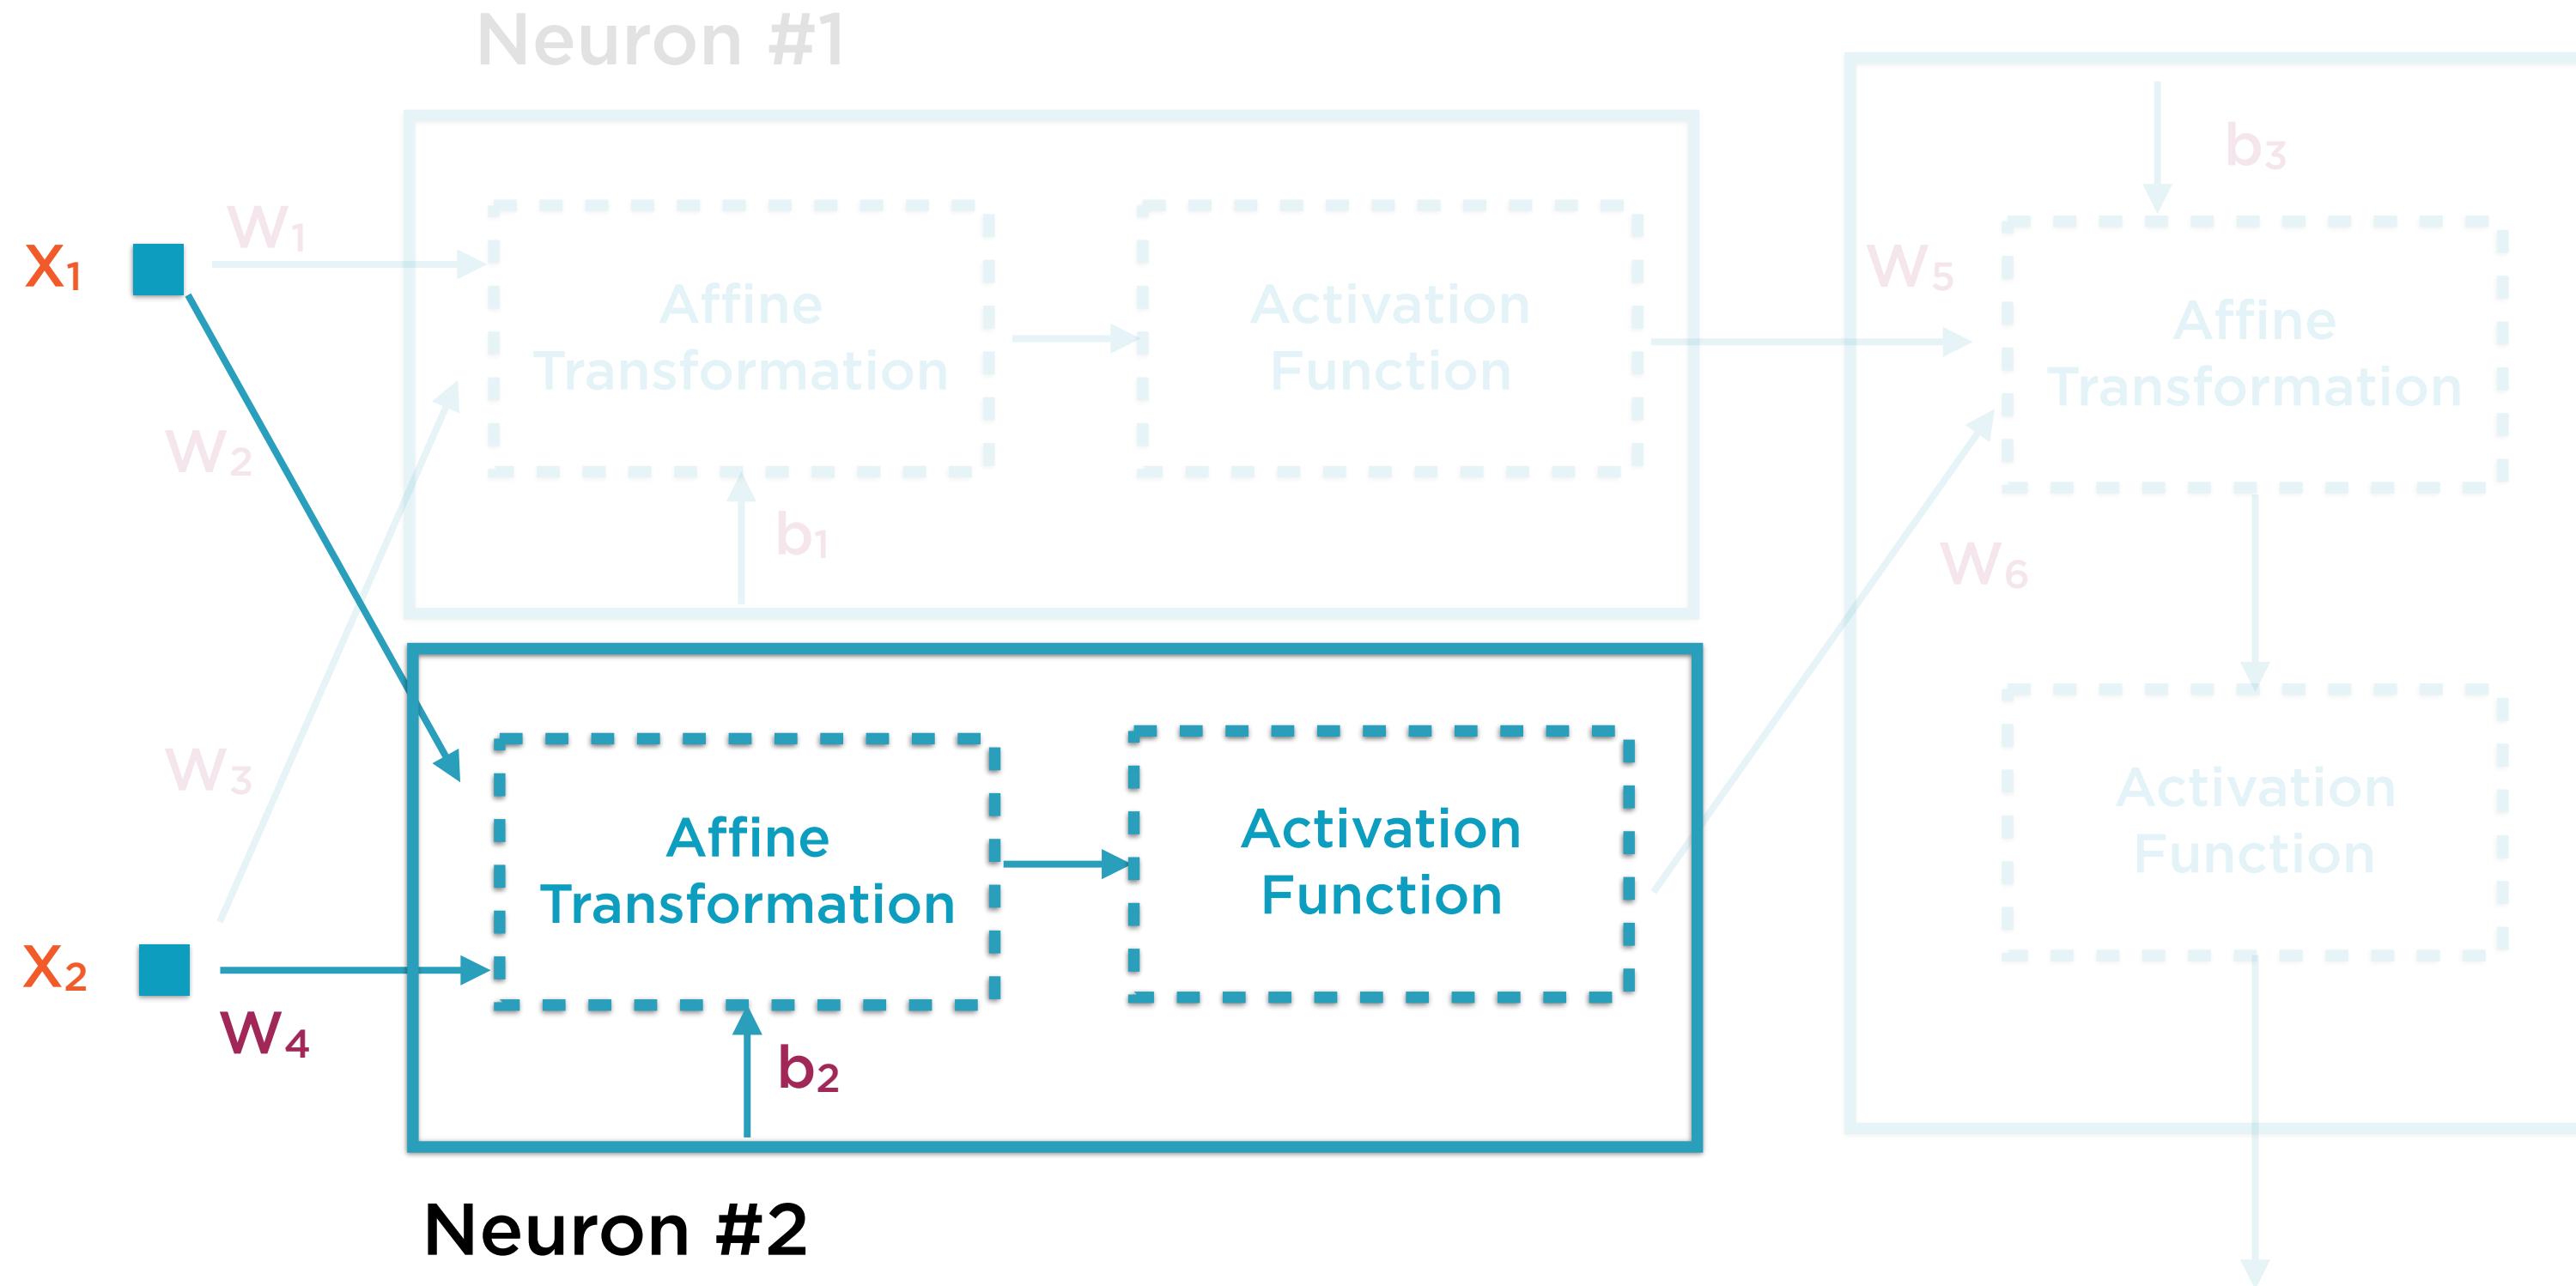
3-Neuron XOR



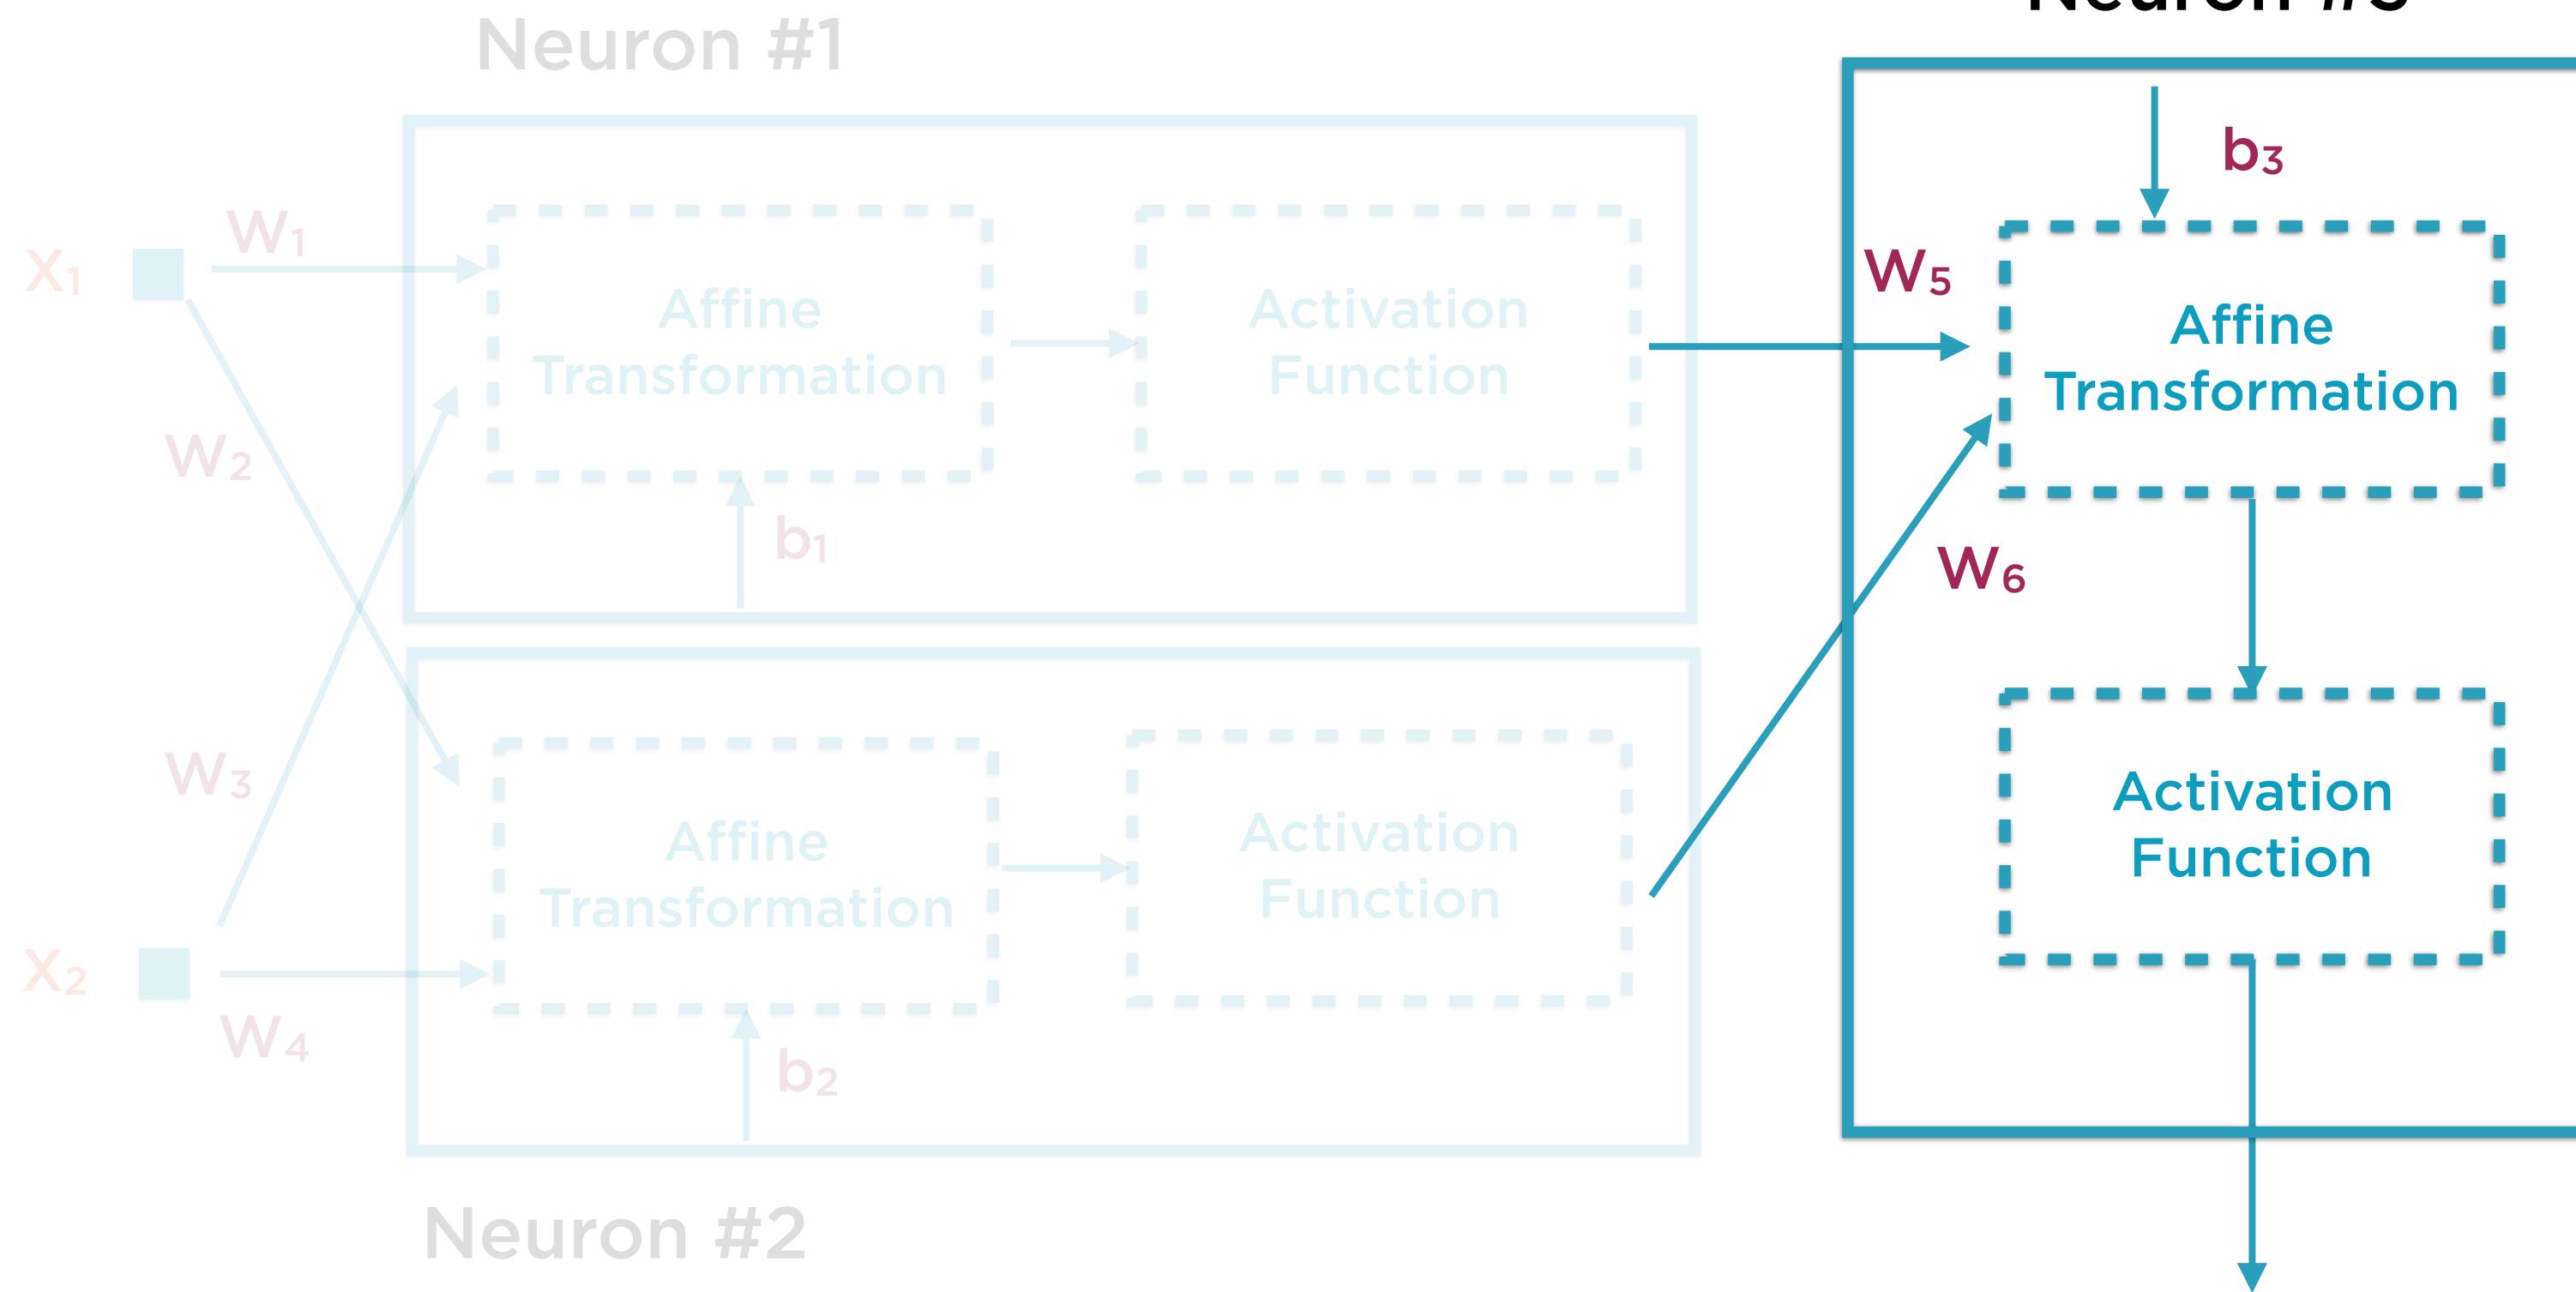
3-Neuron XOR



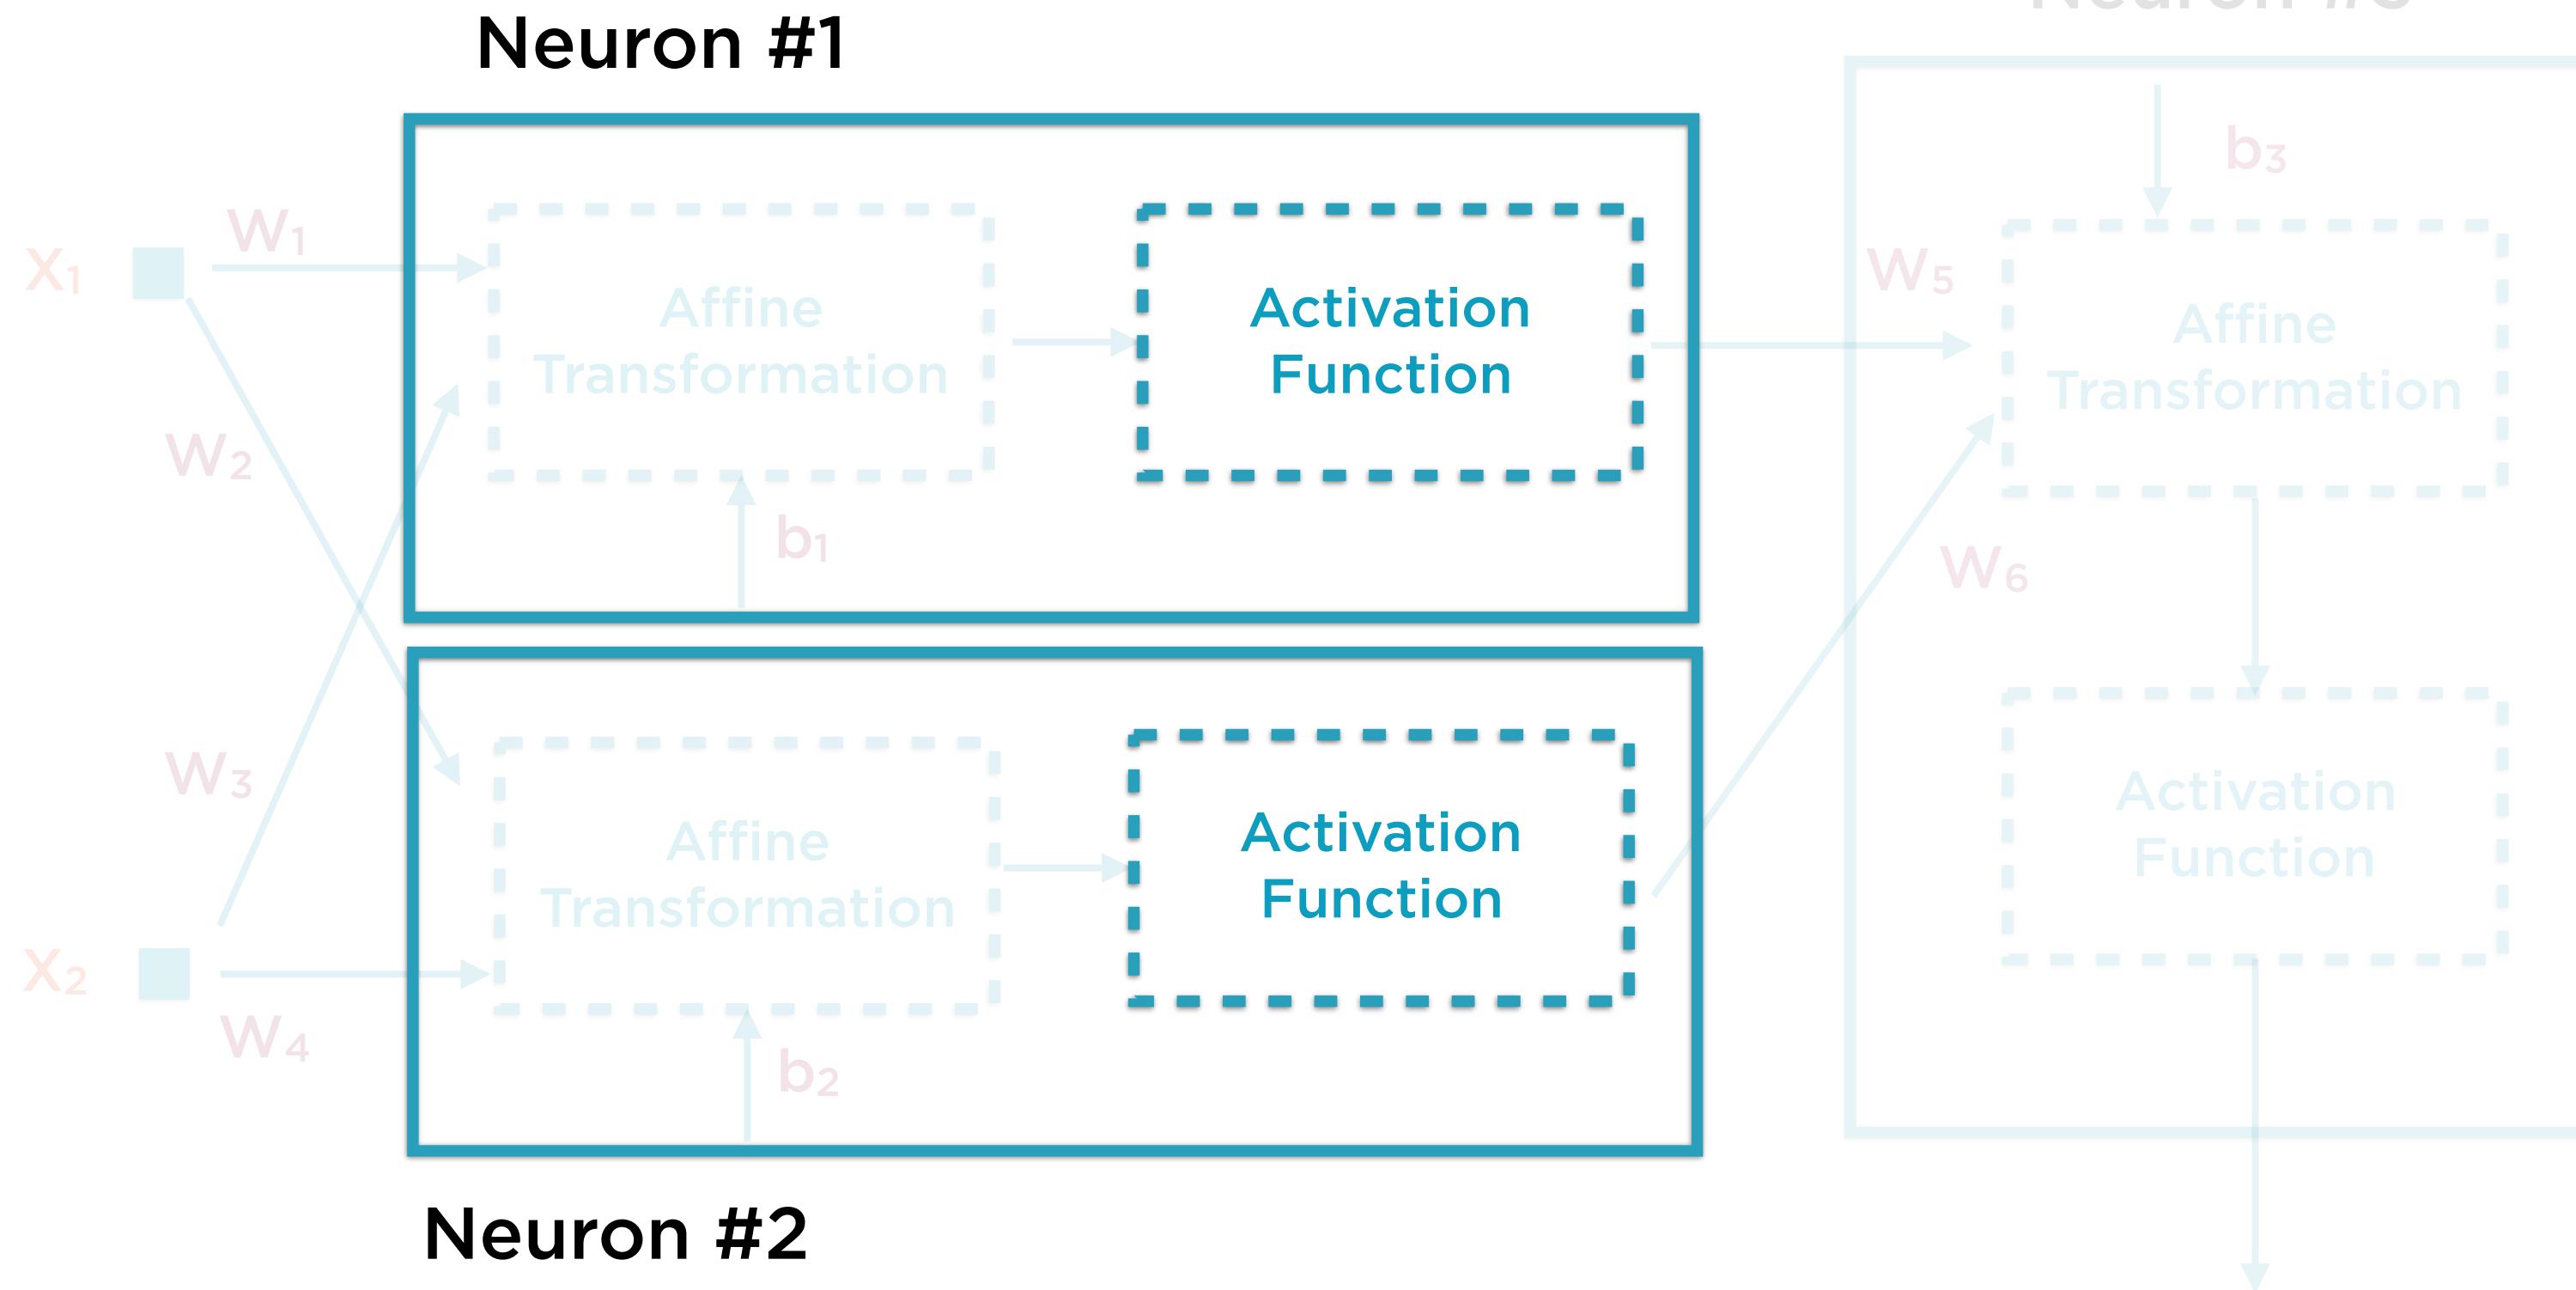
3-Neuron XOR



3-Neuron XOR



3-Neuron XOR



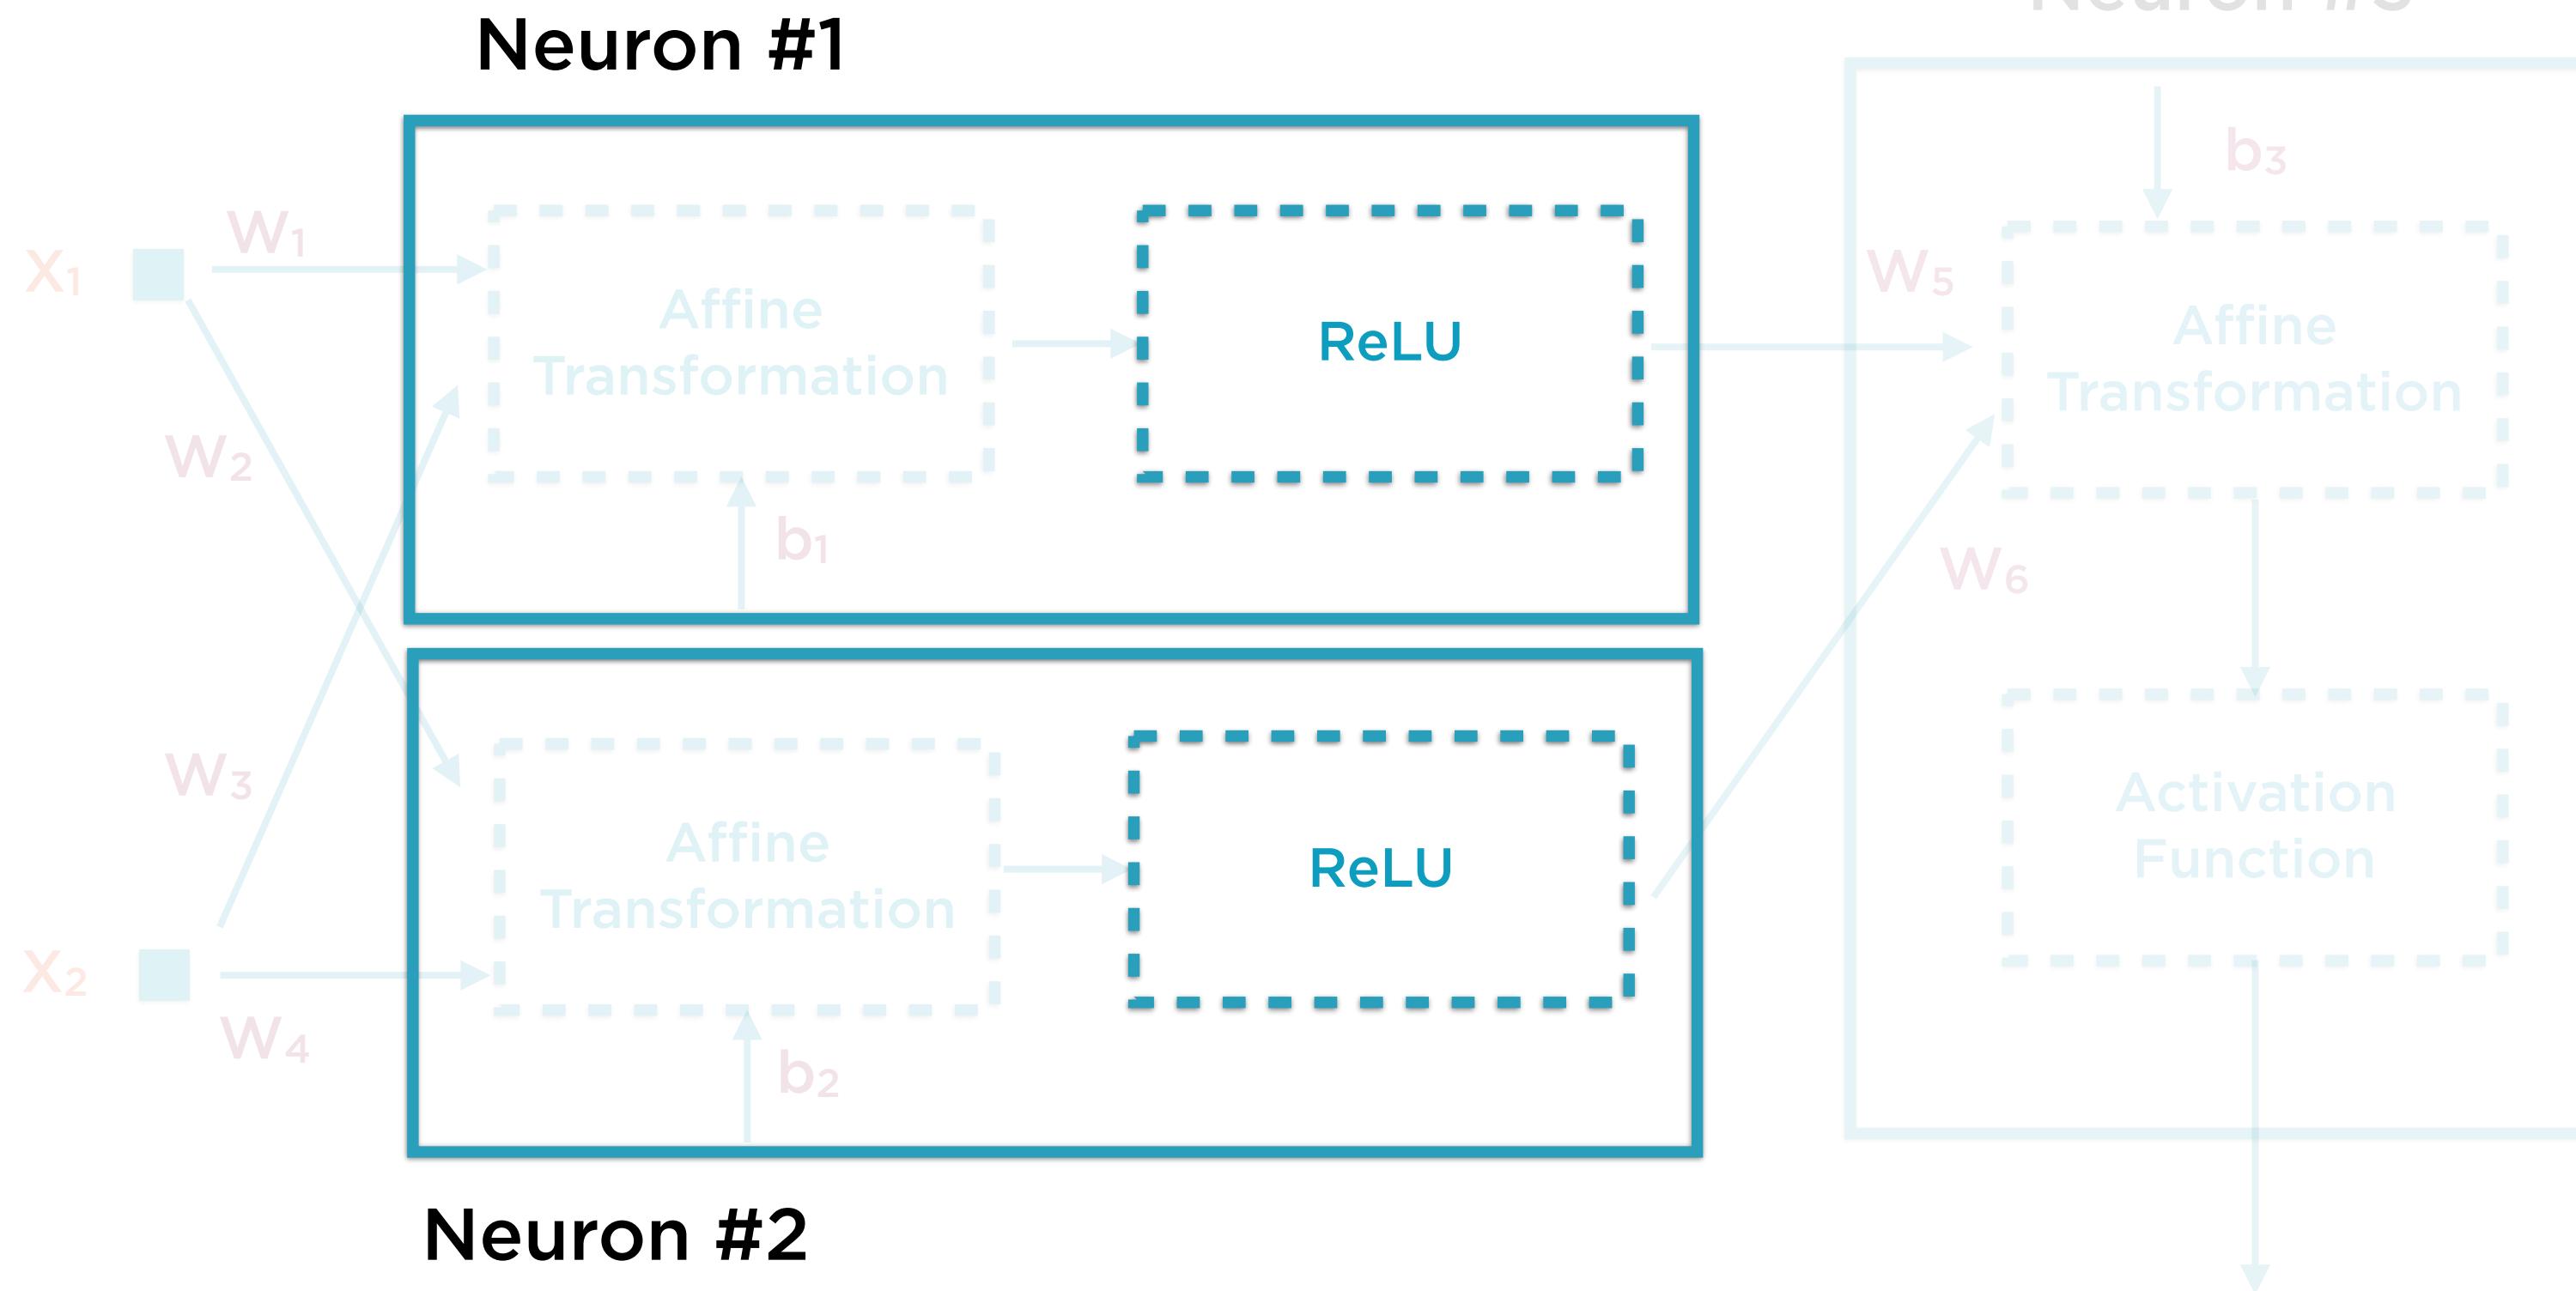


The most common form of the activation function is the ReLU

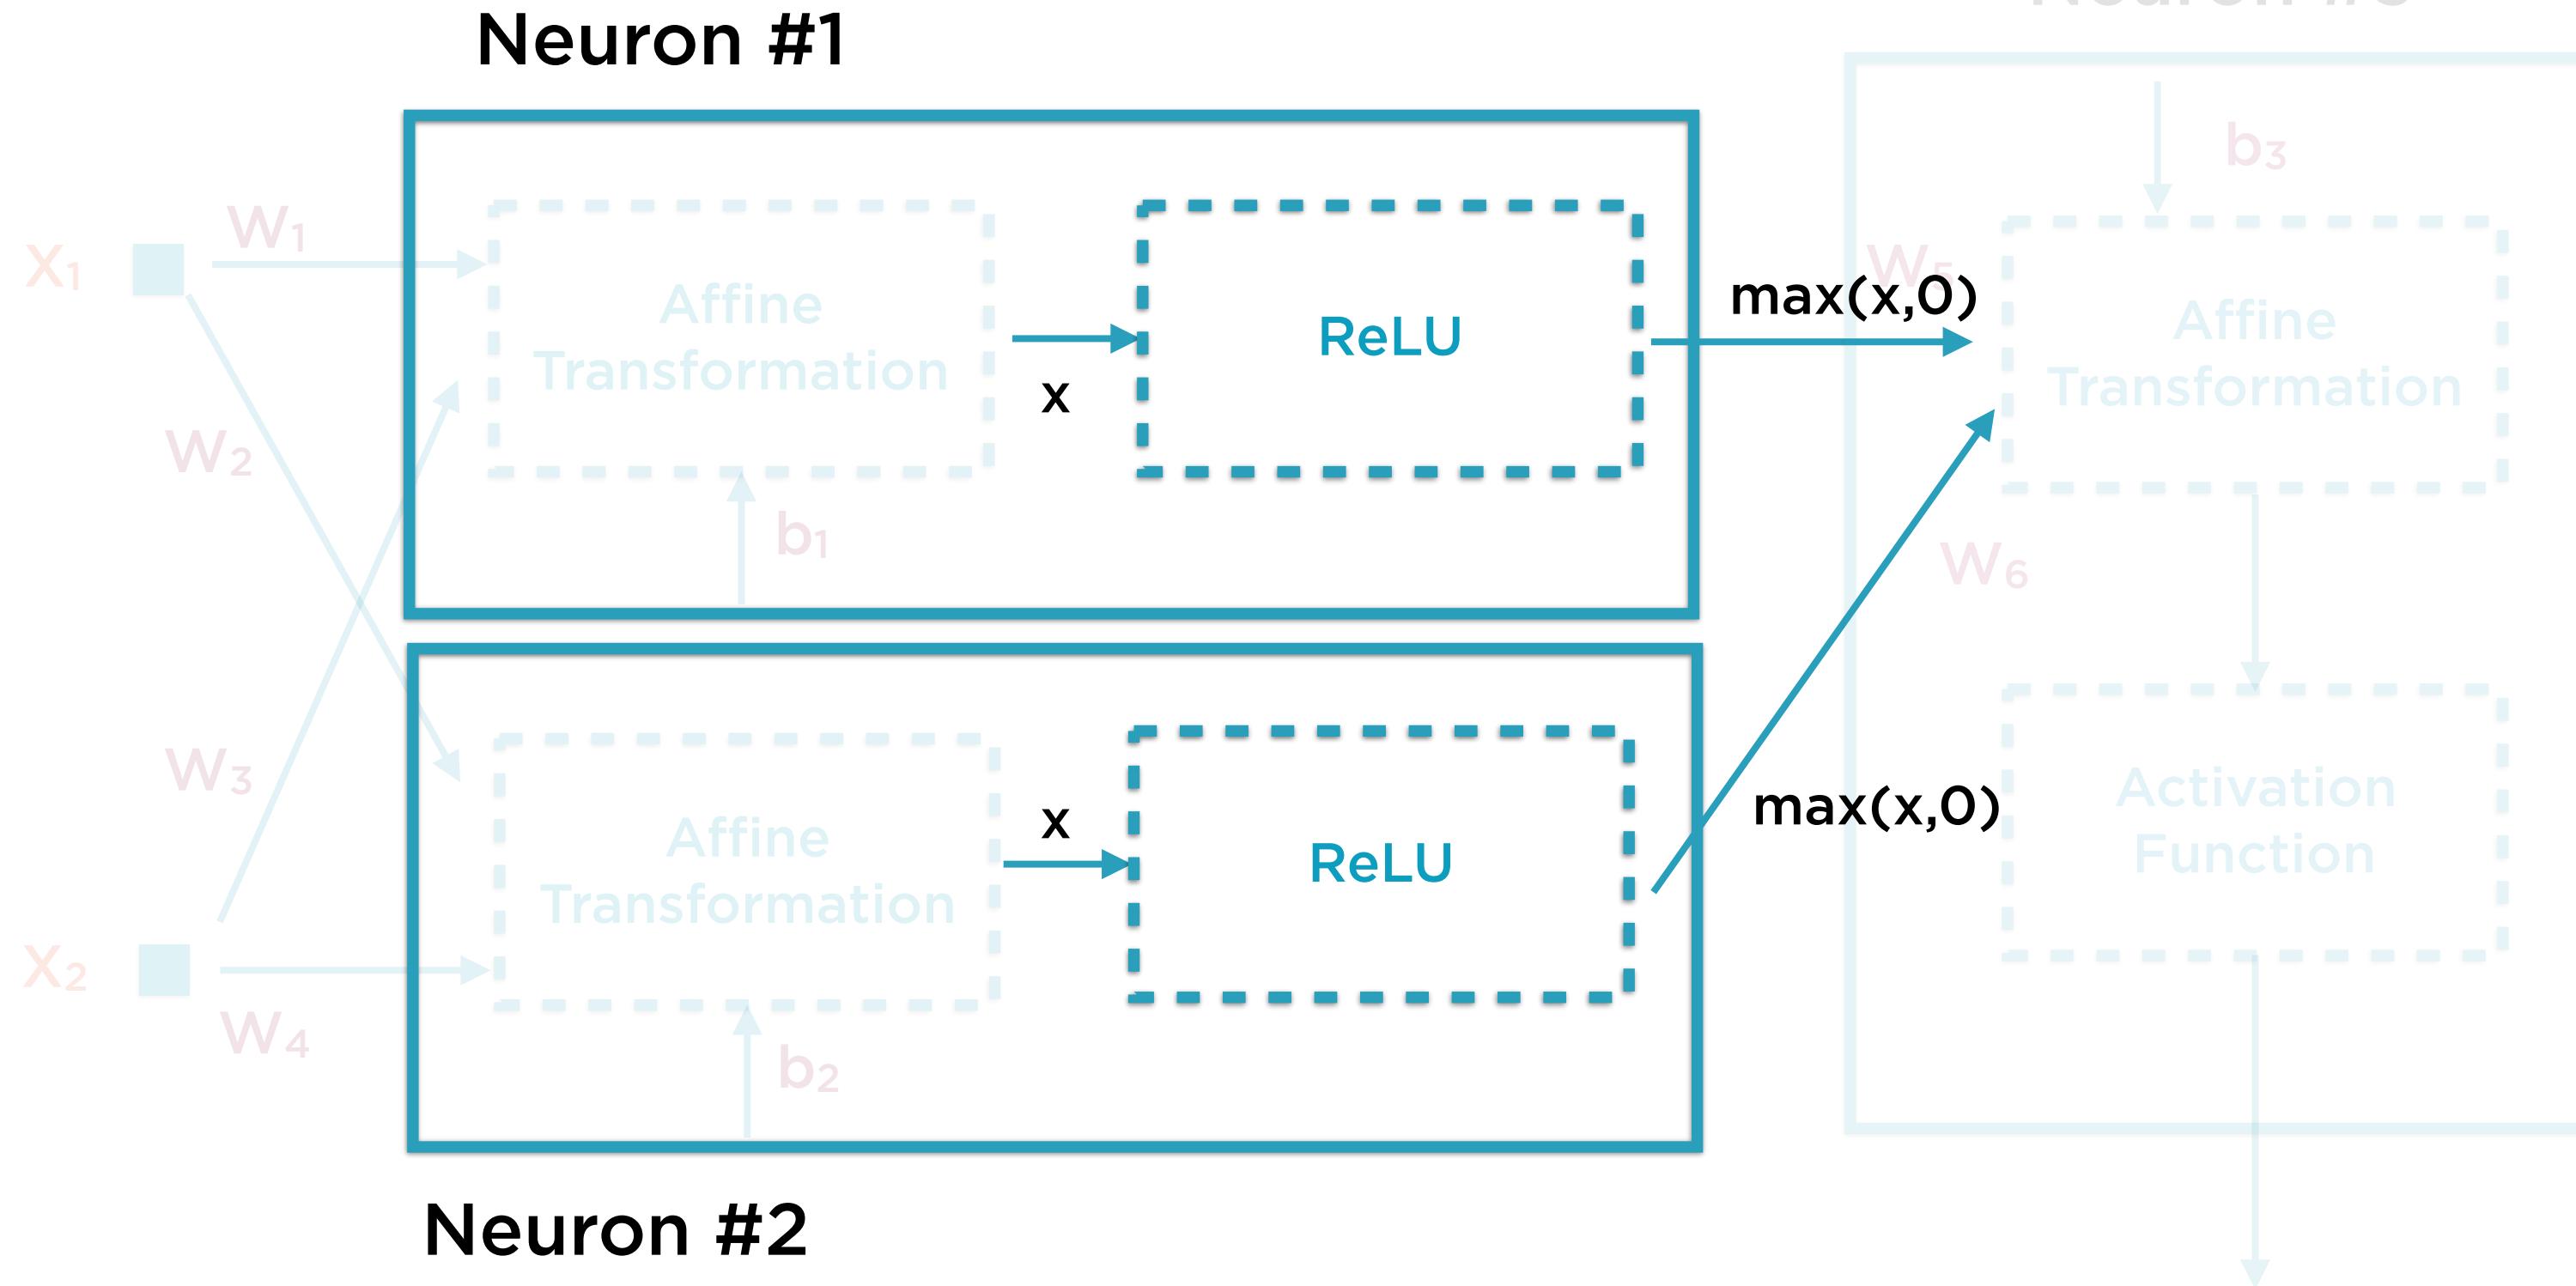
ReLU : Rectified Linear Unit

$$\text{ReLU}(x) = \max(0, x)$$

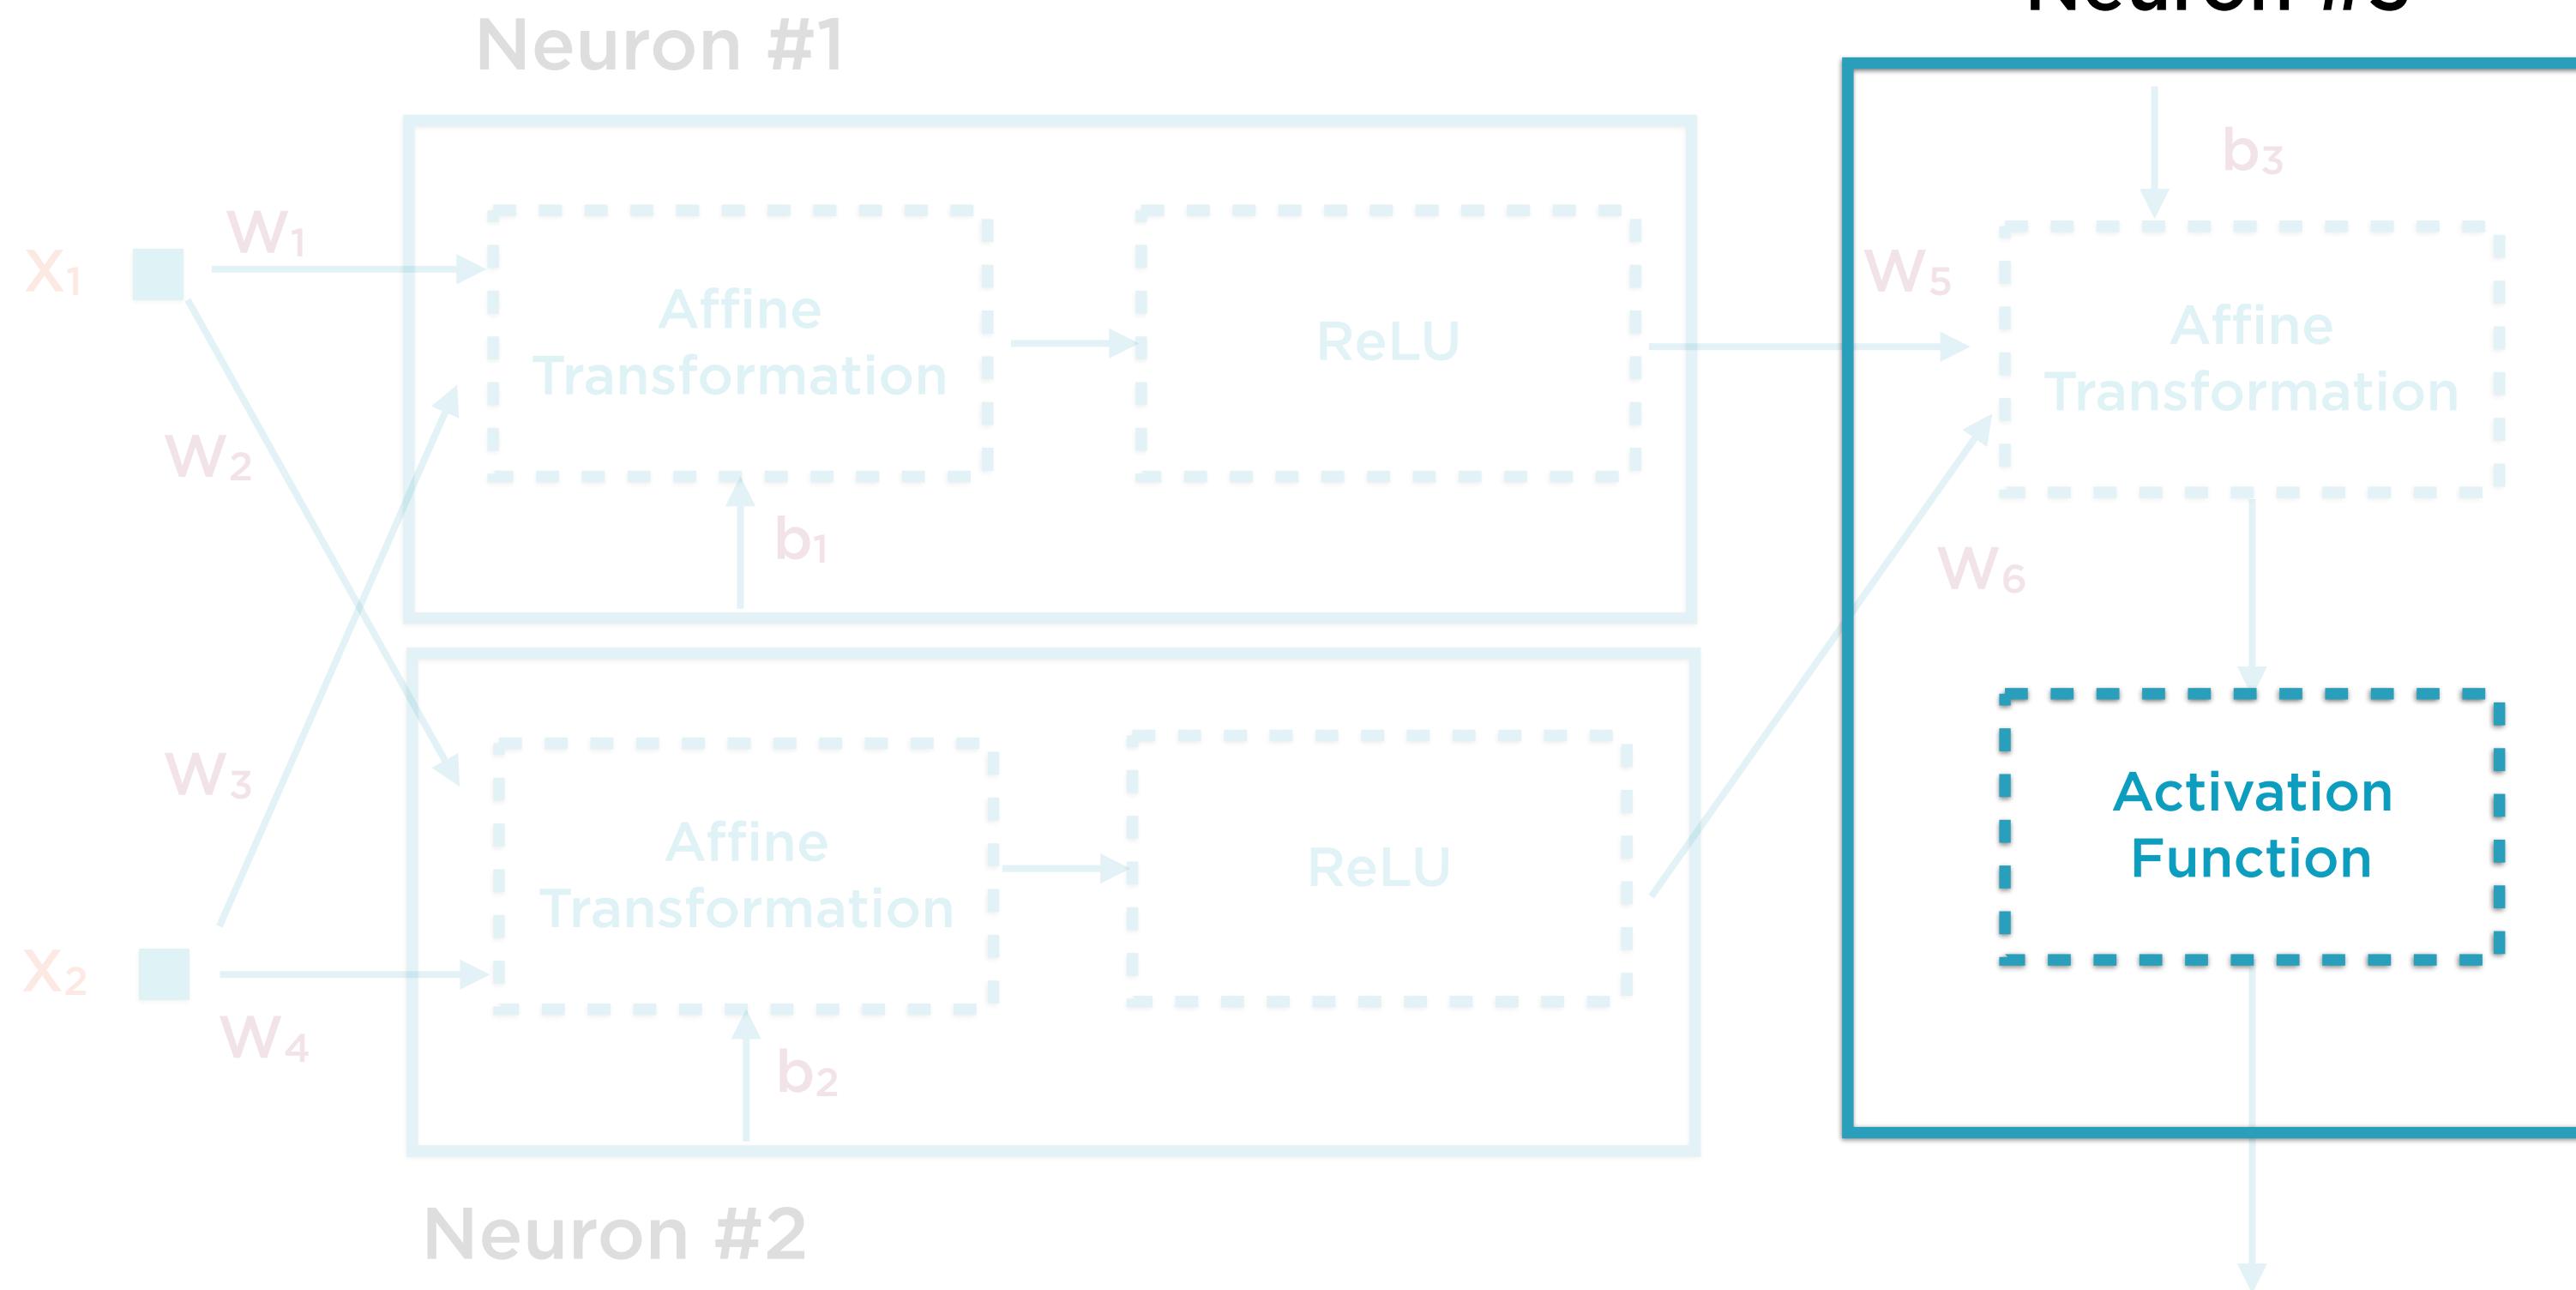
3-Neuron XOR



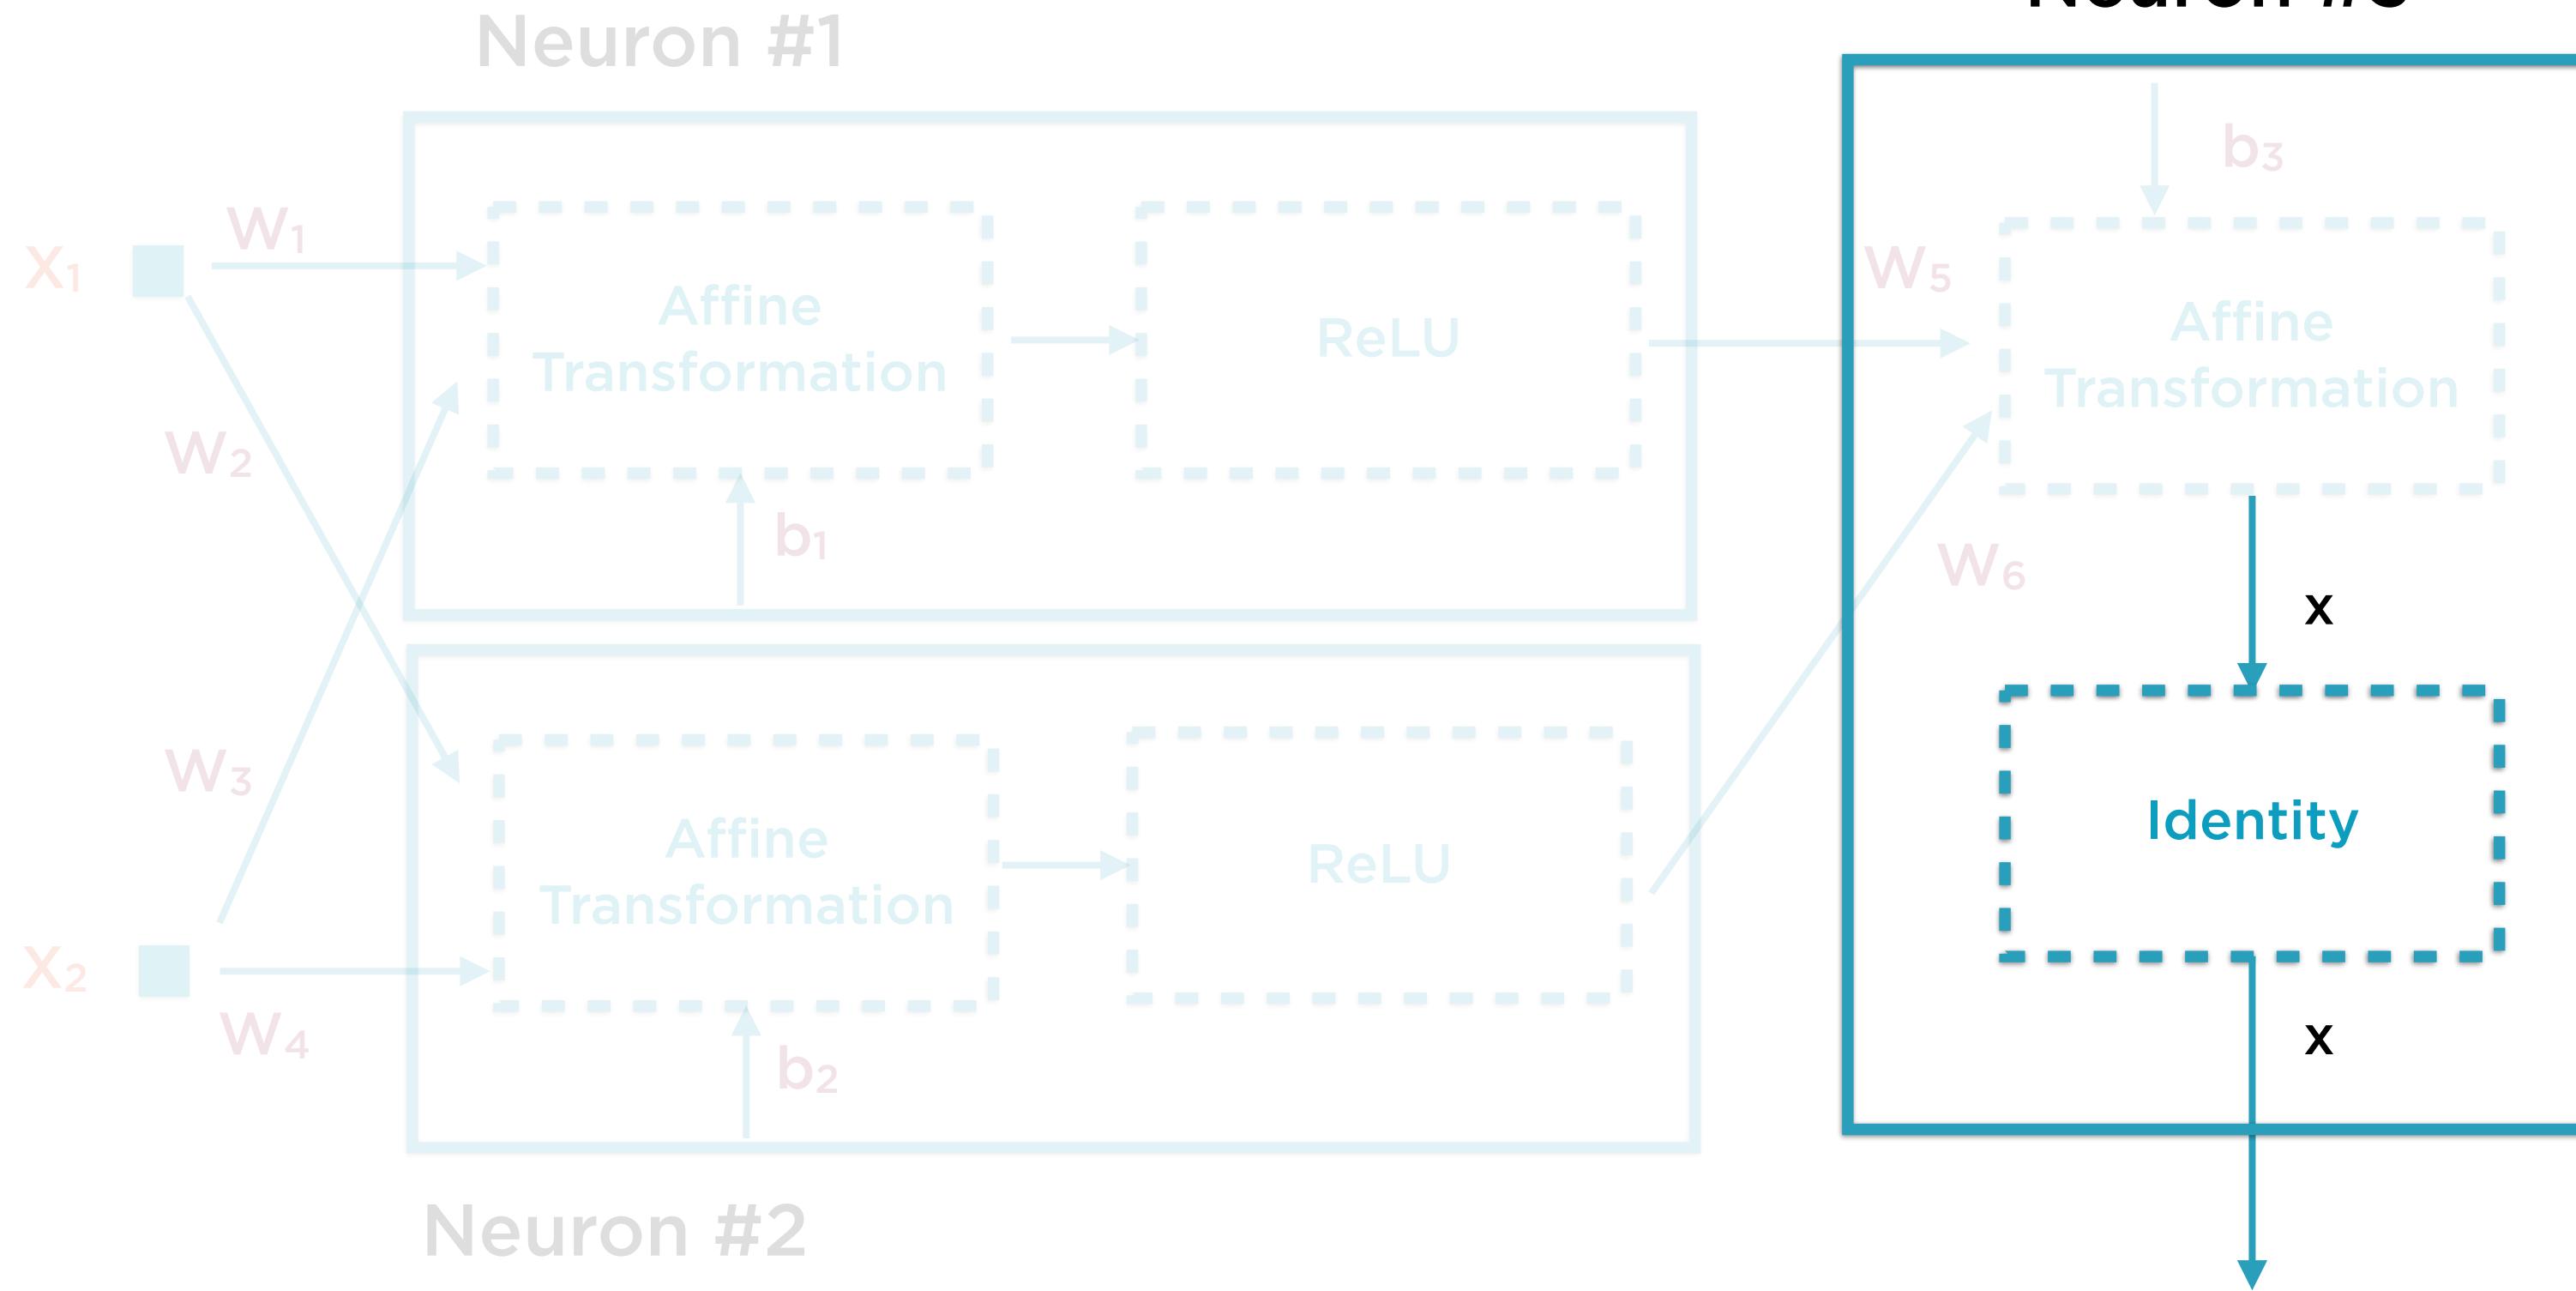
3-Neuron XOR



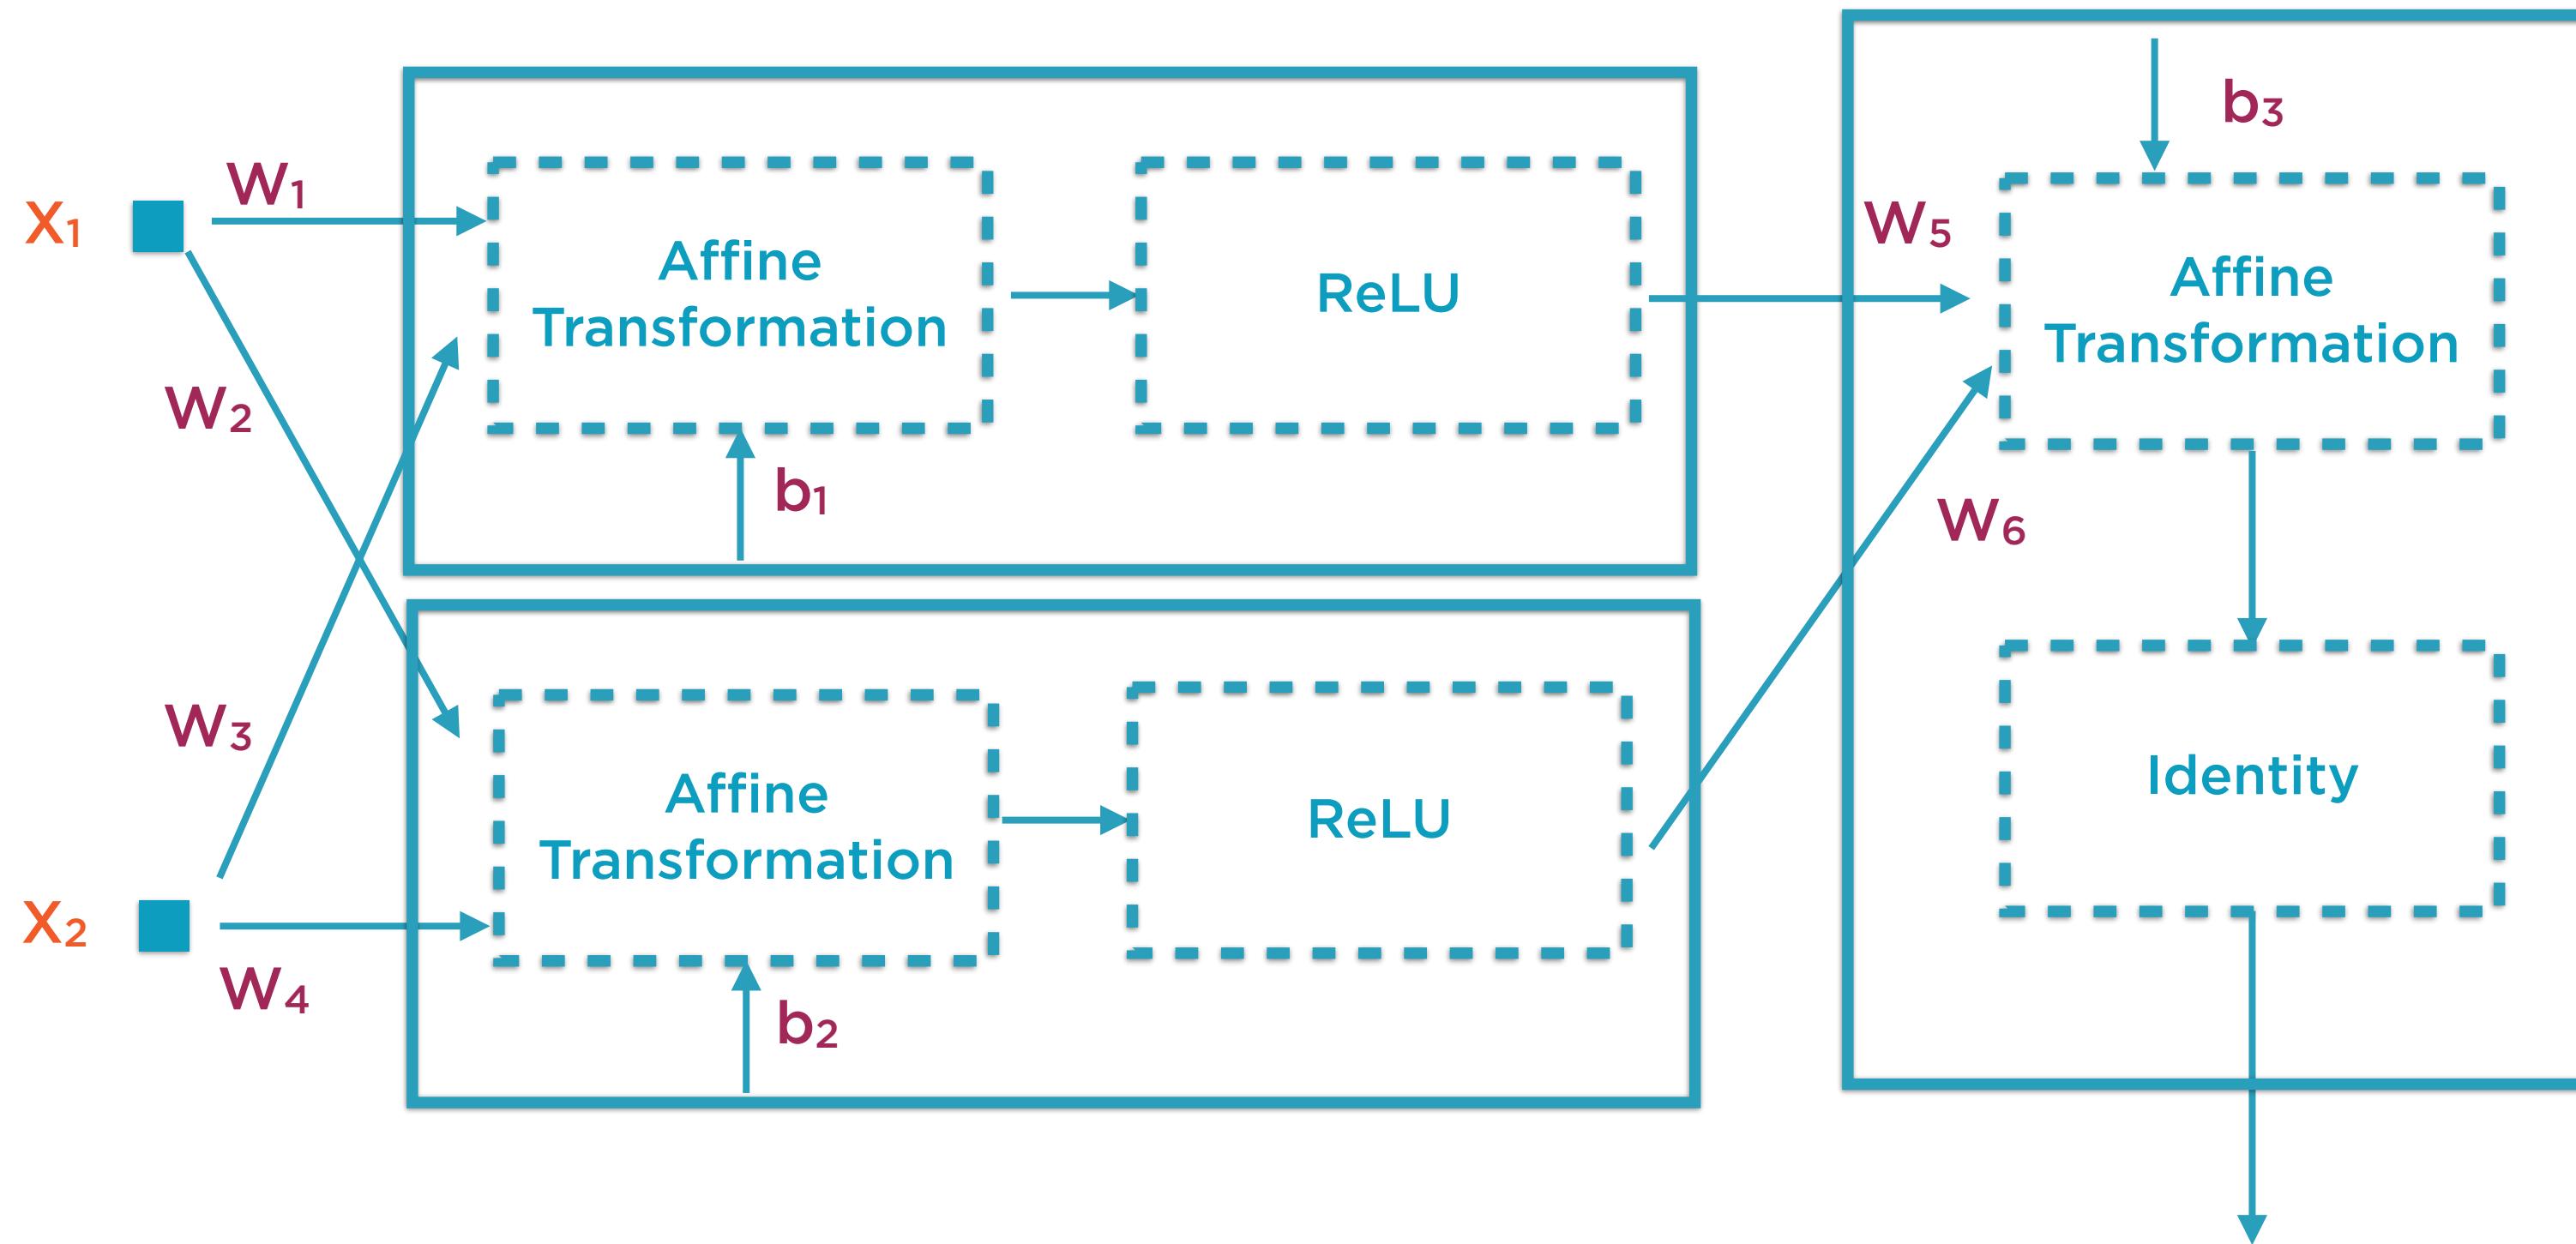
3-Neuron XOR



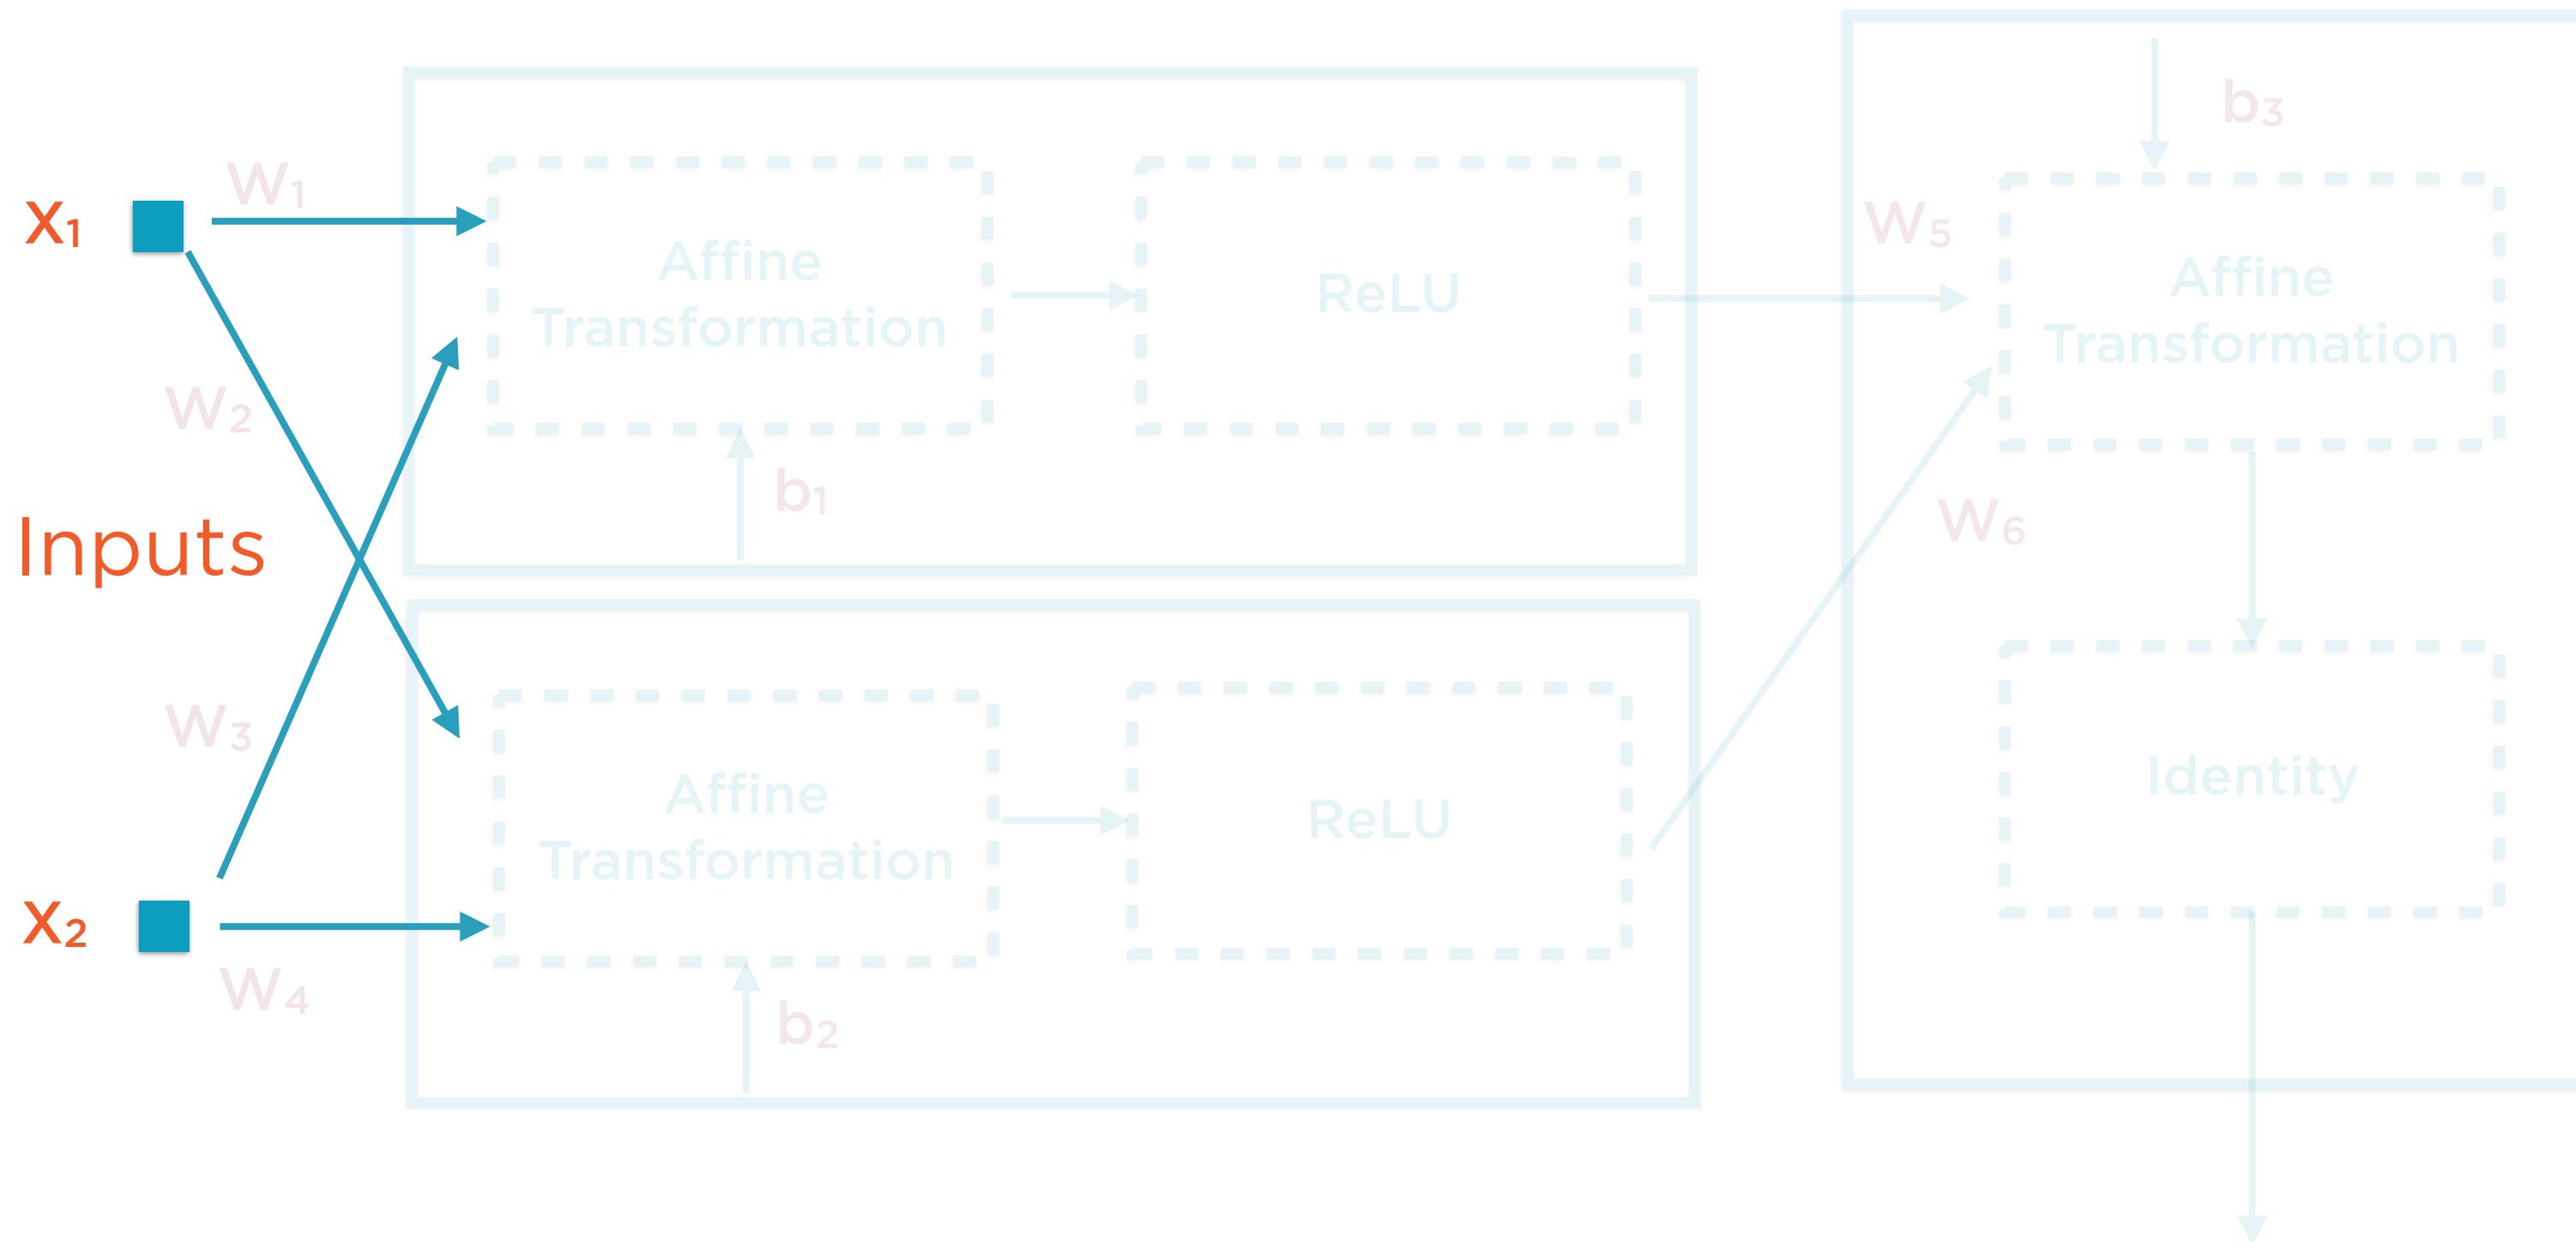
3-Neuron XOR



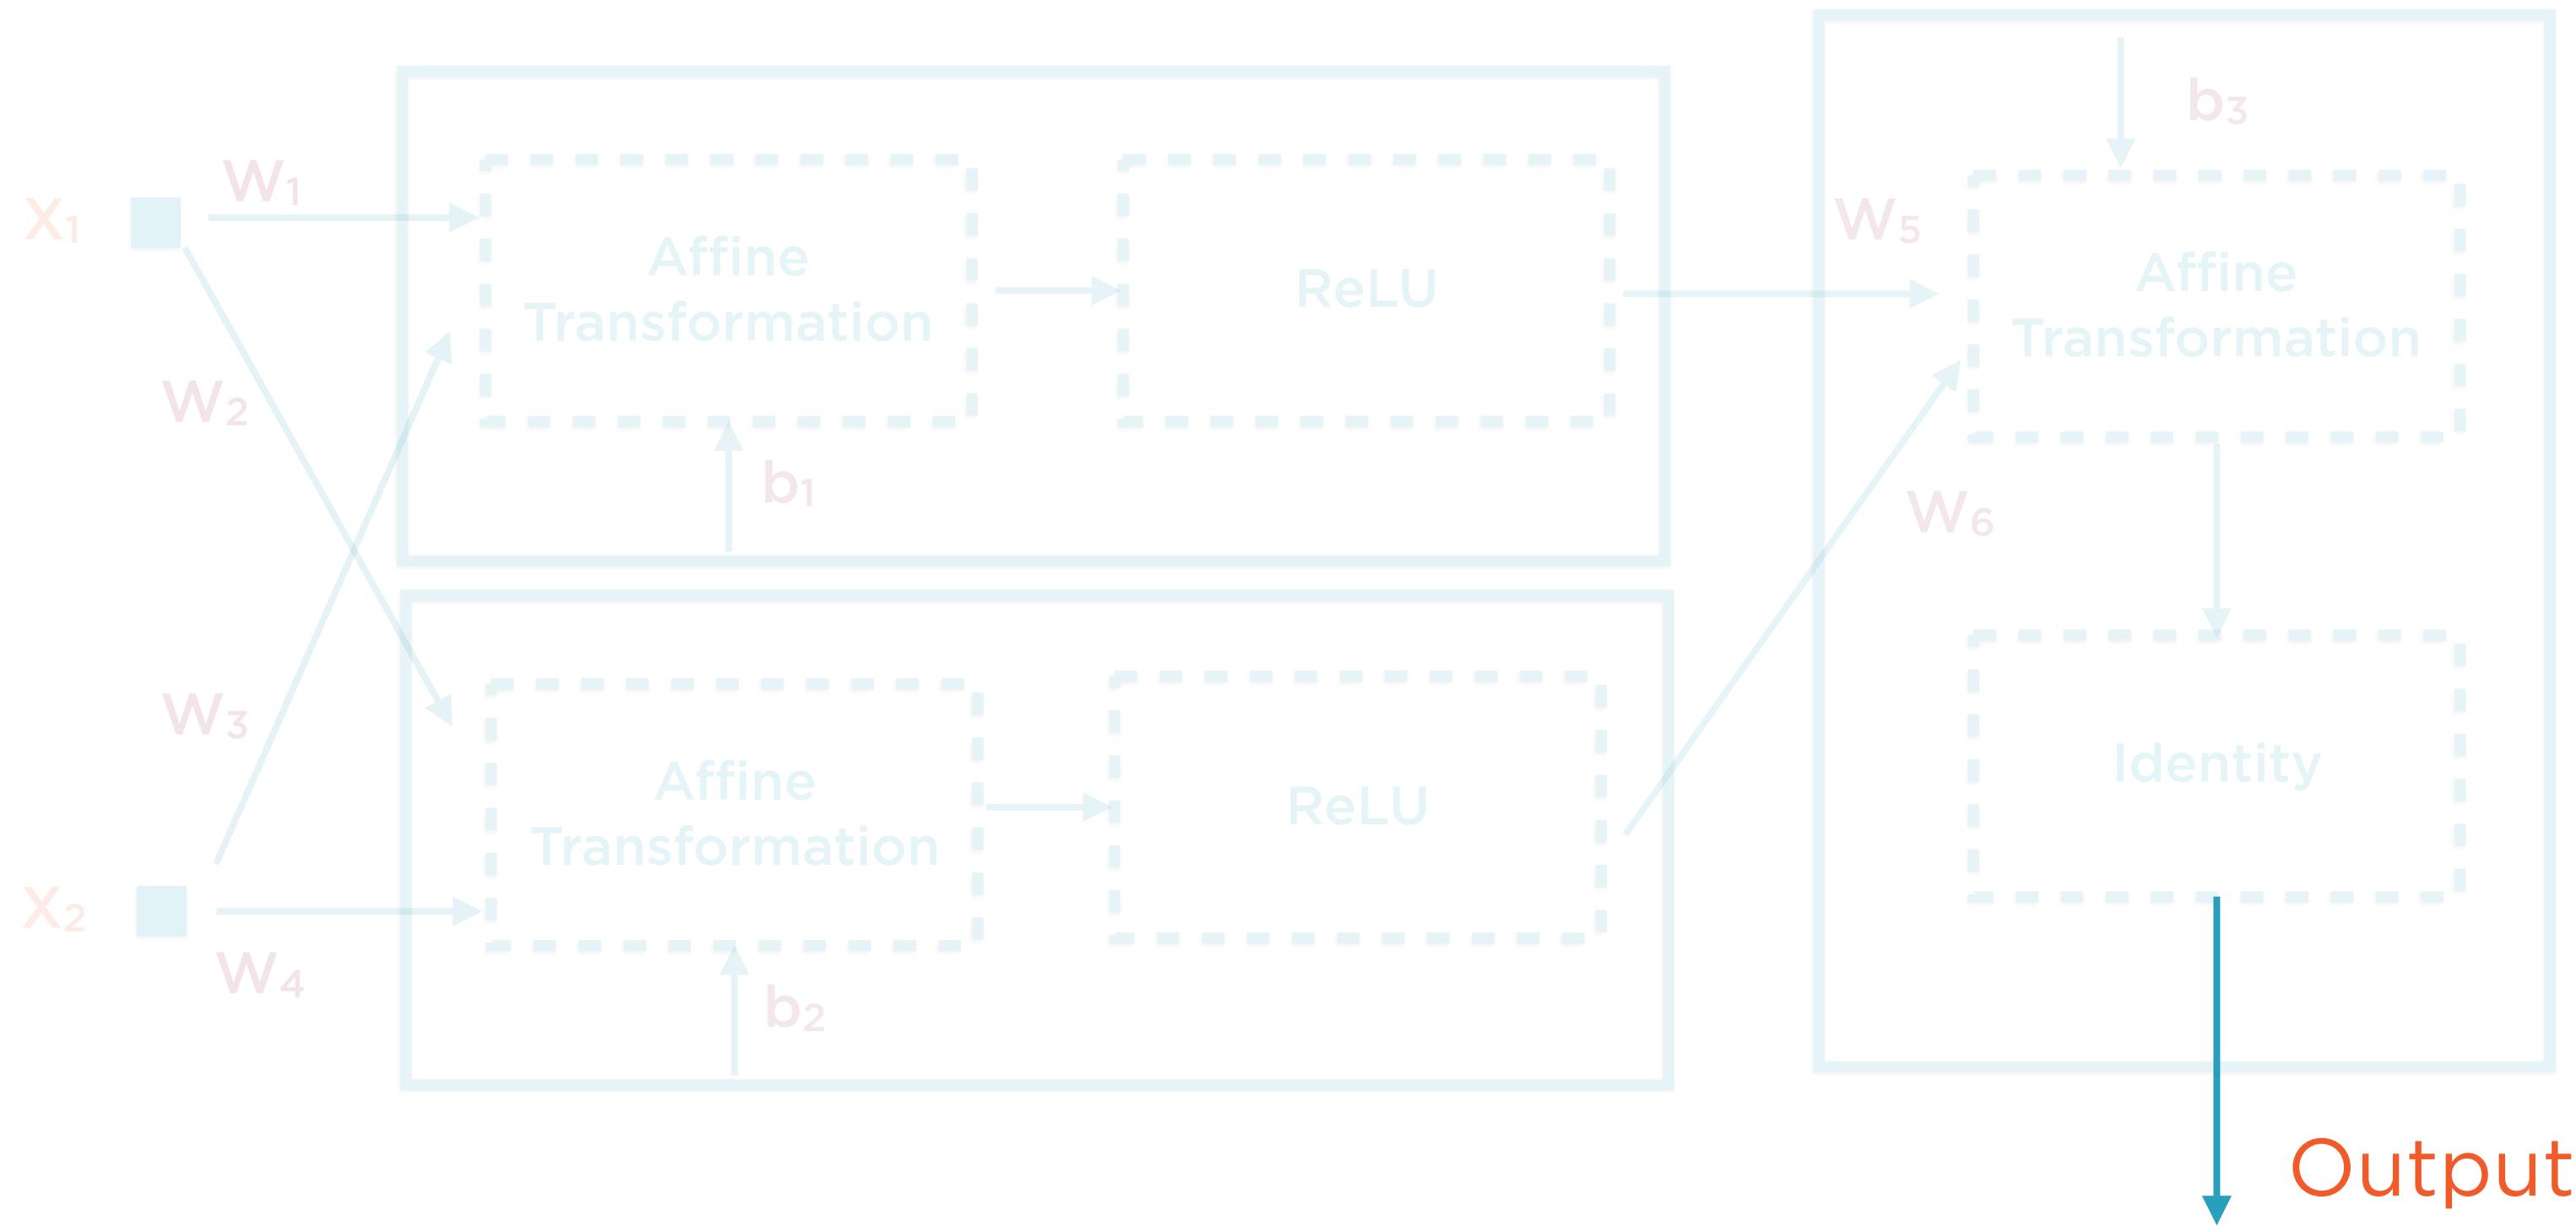
3-Neuron XOR



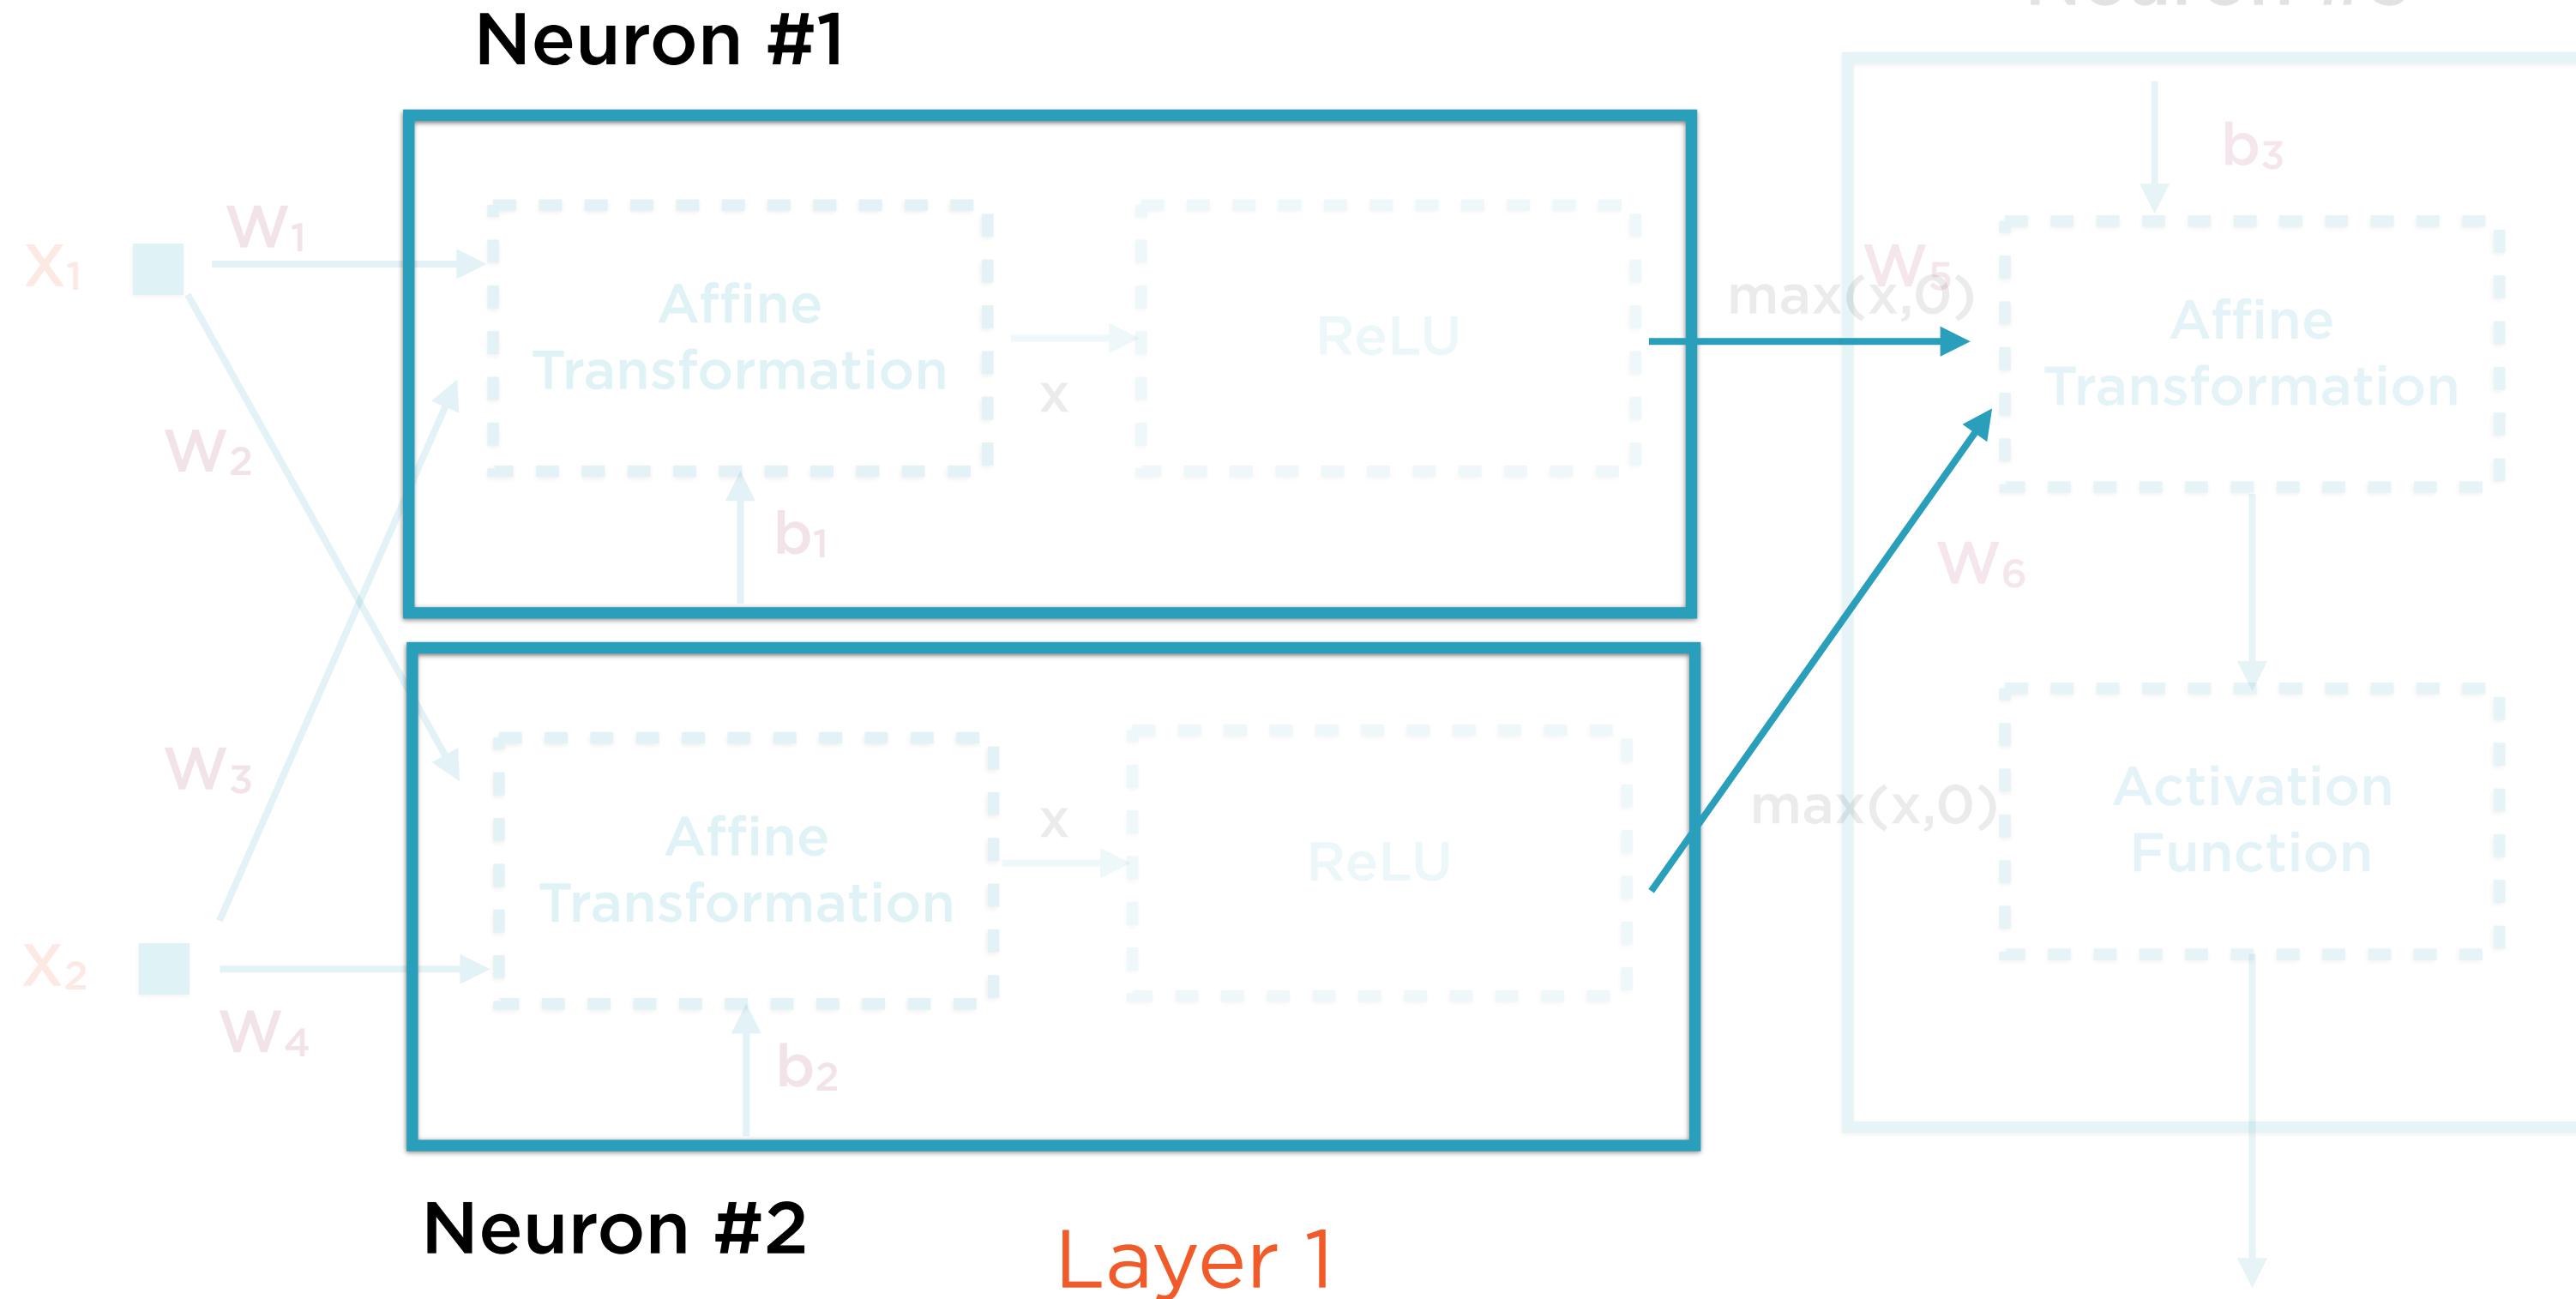
3-Neuron XOR



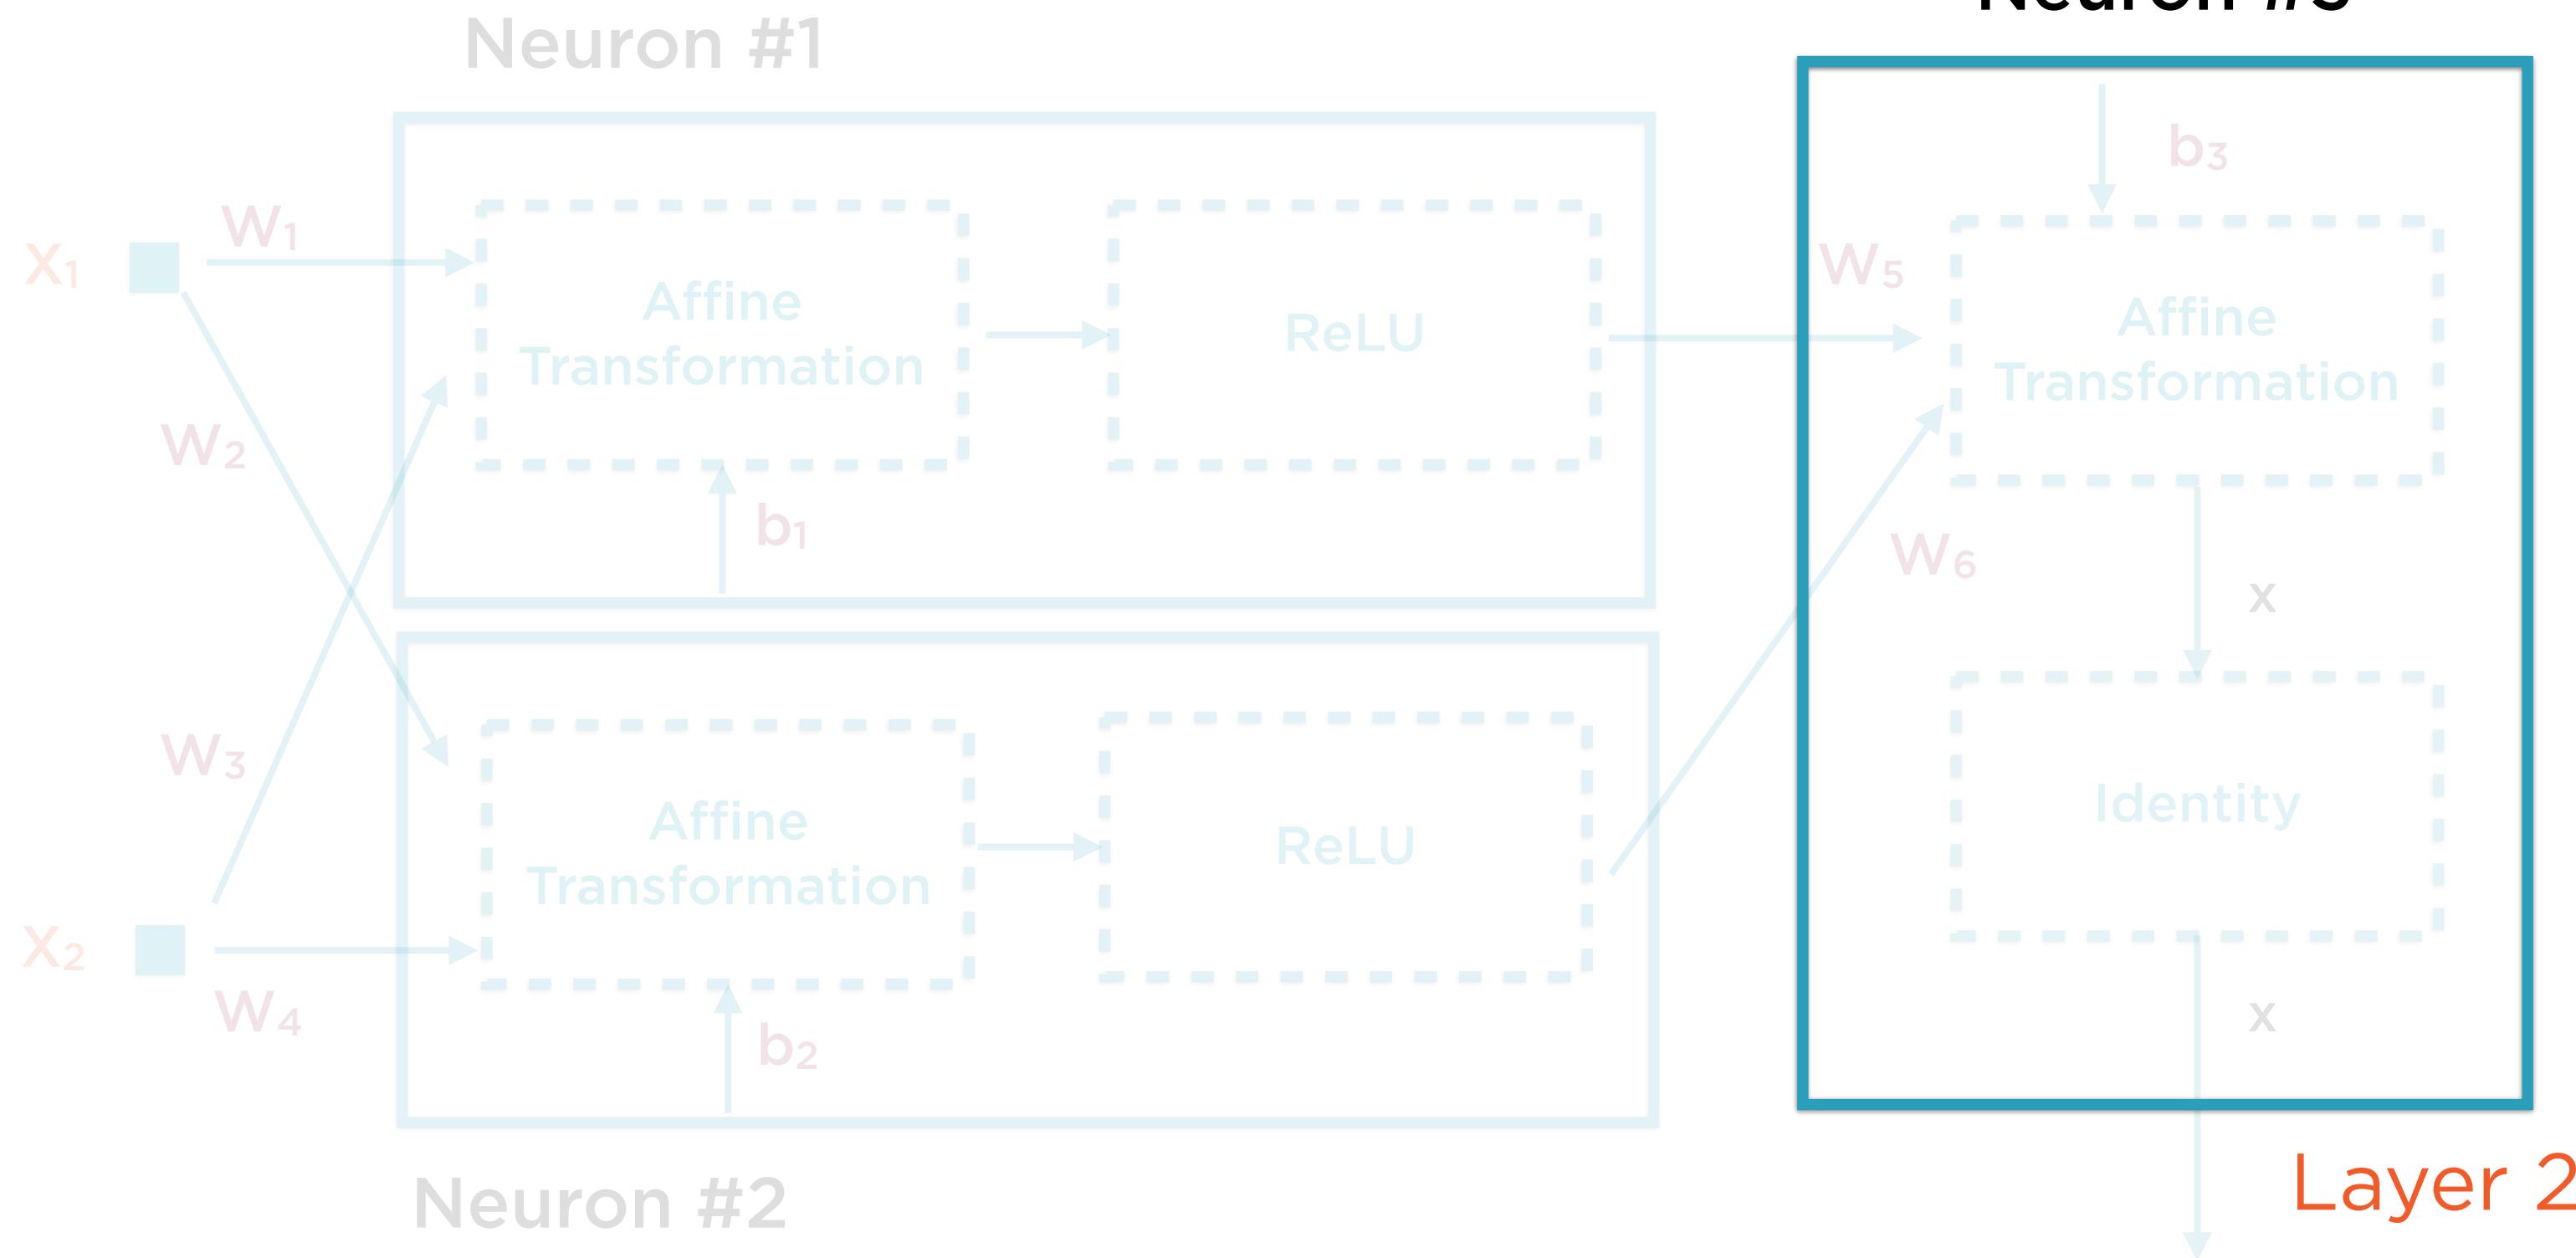
3-Neuron XOR



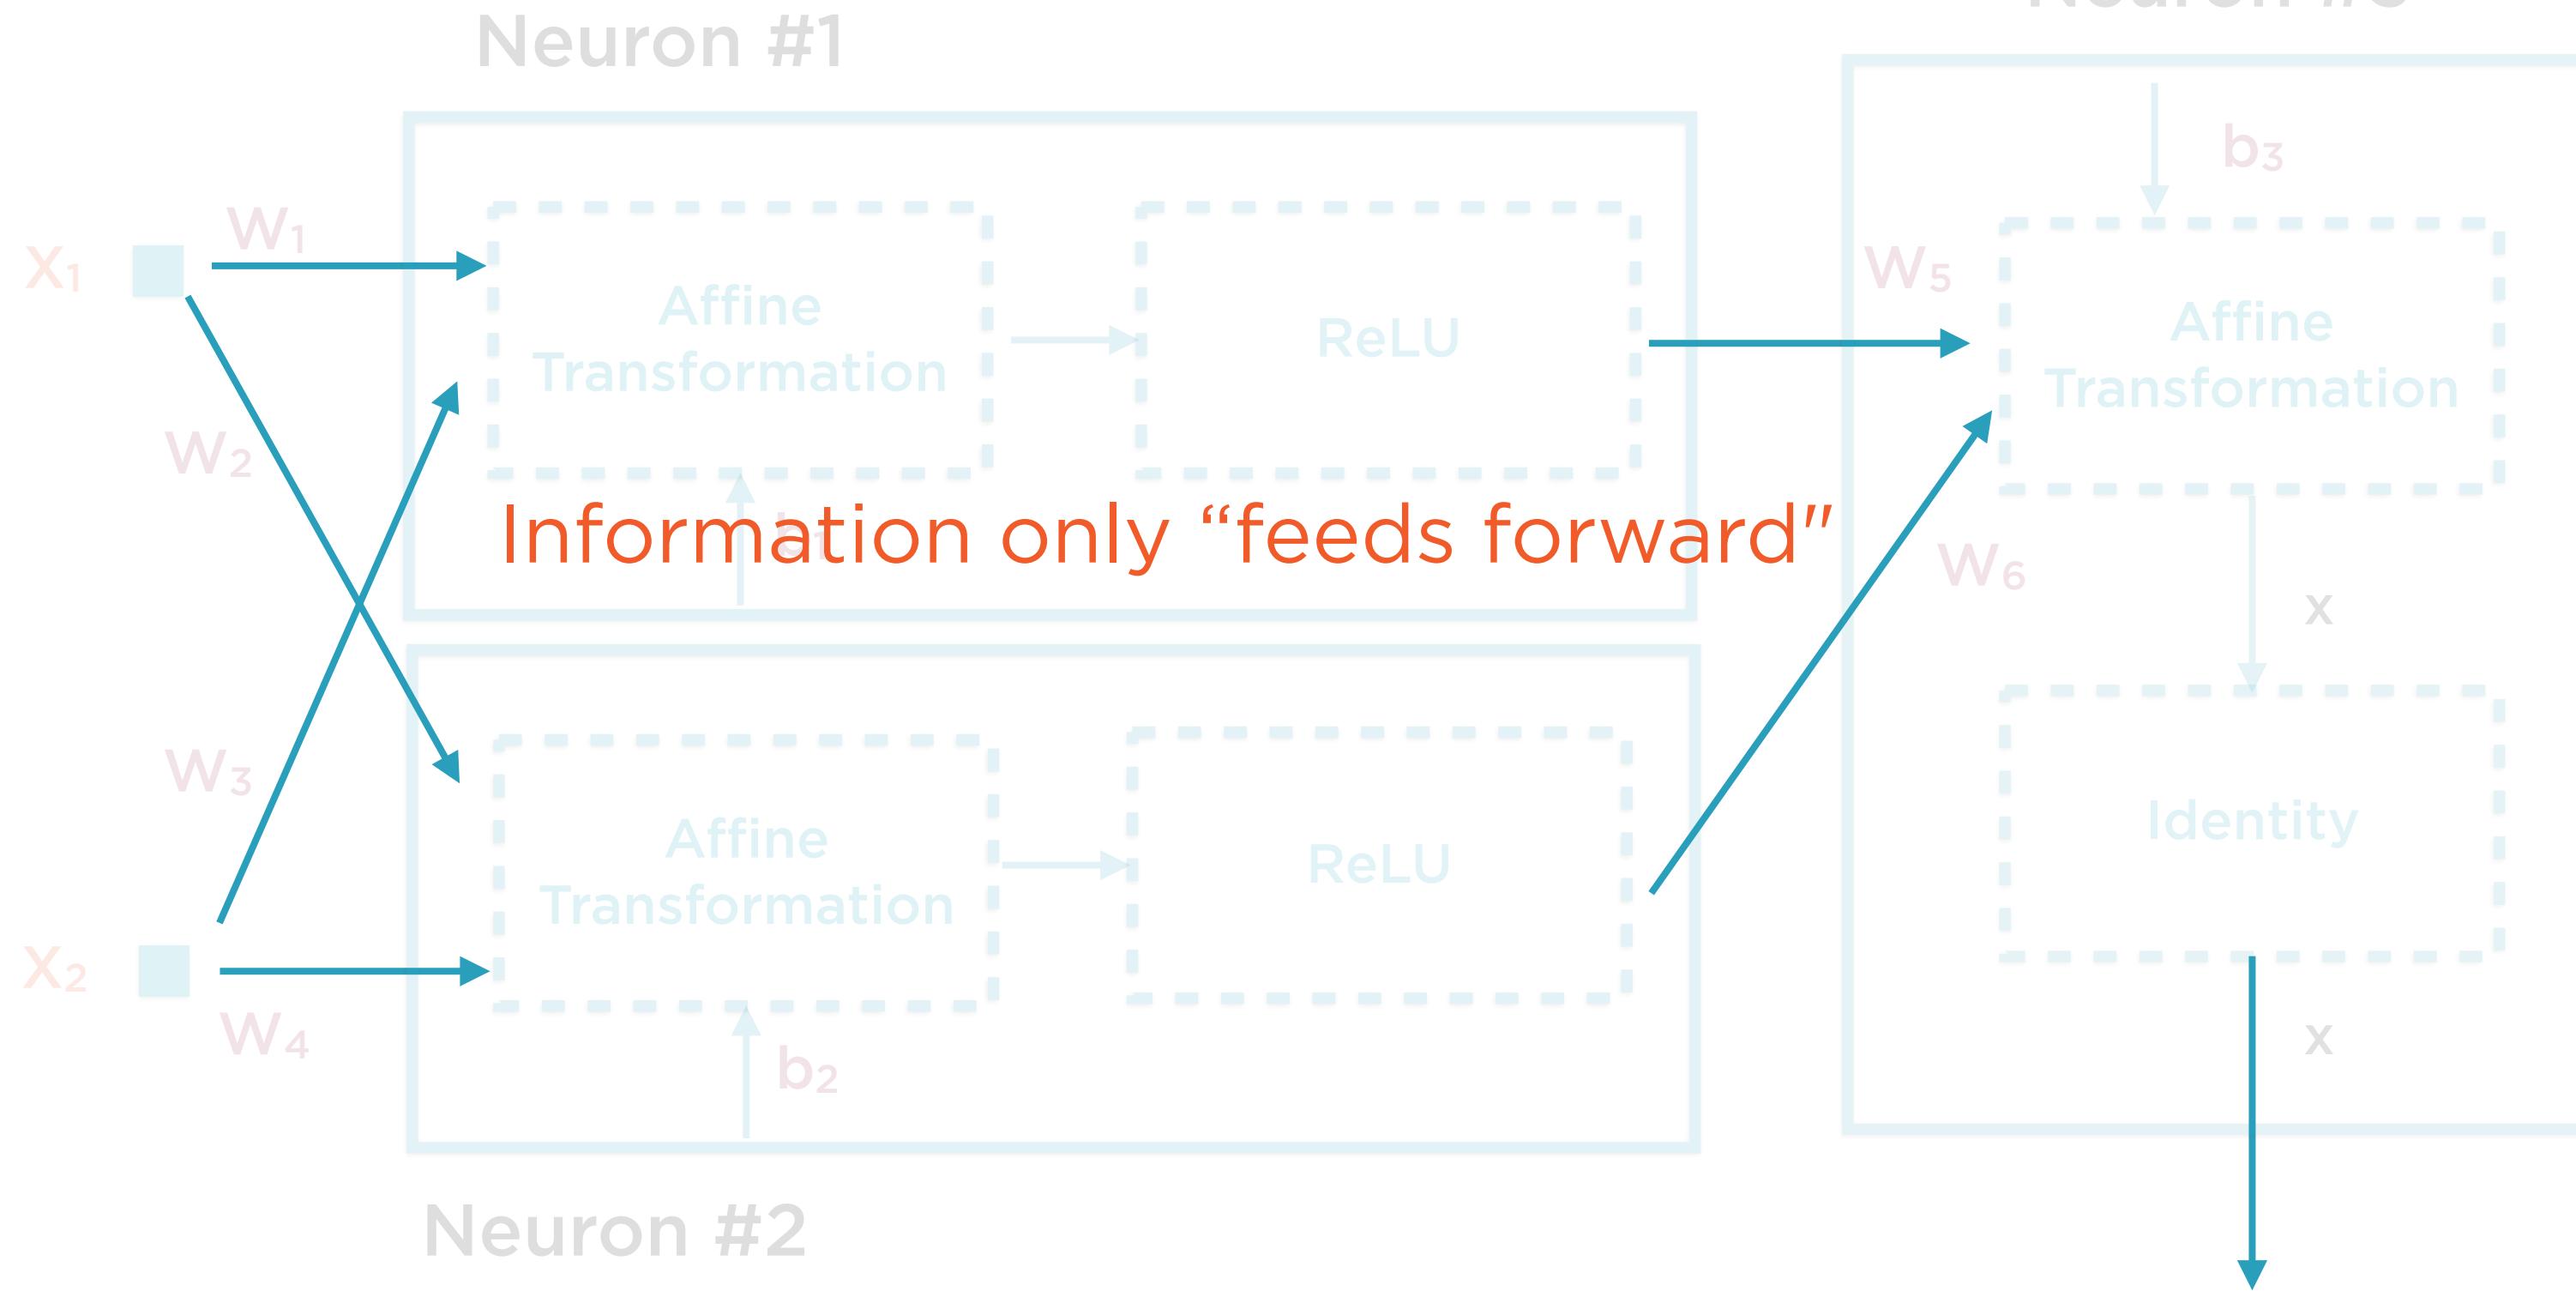
3-Neuron XOR



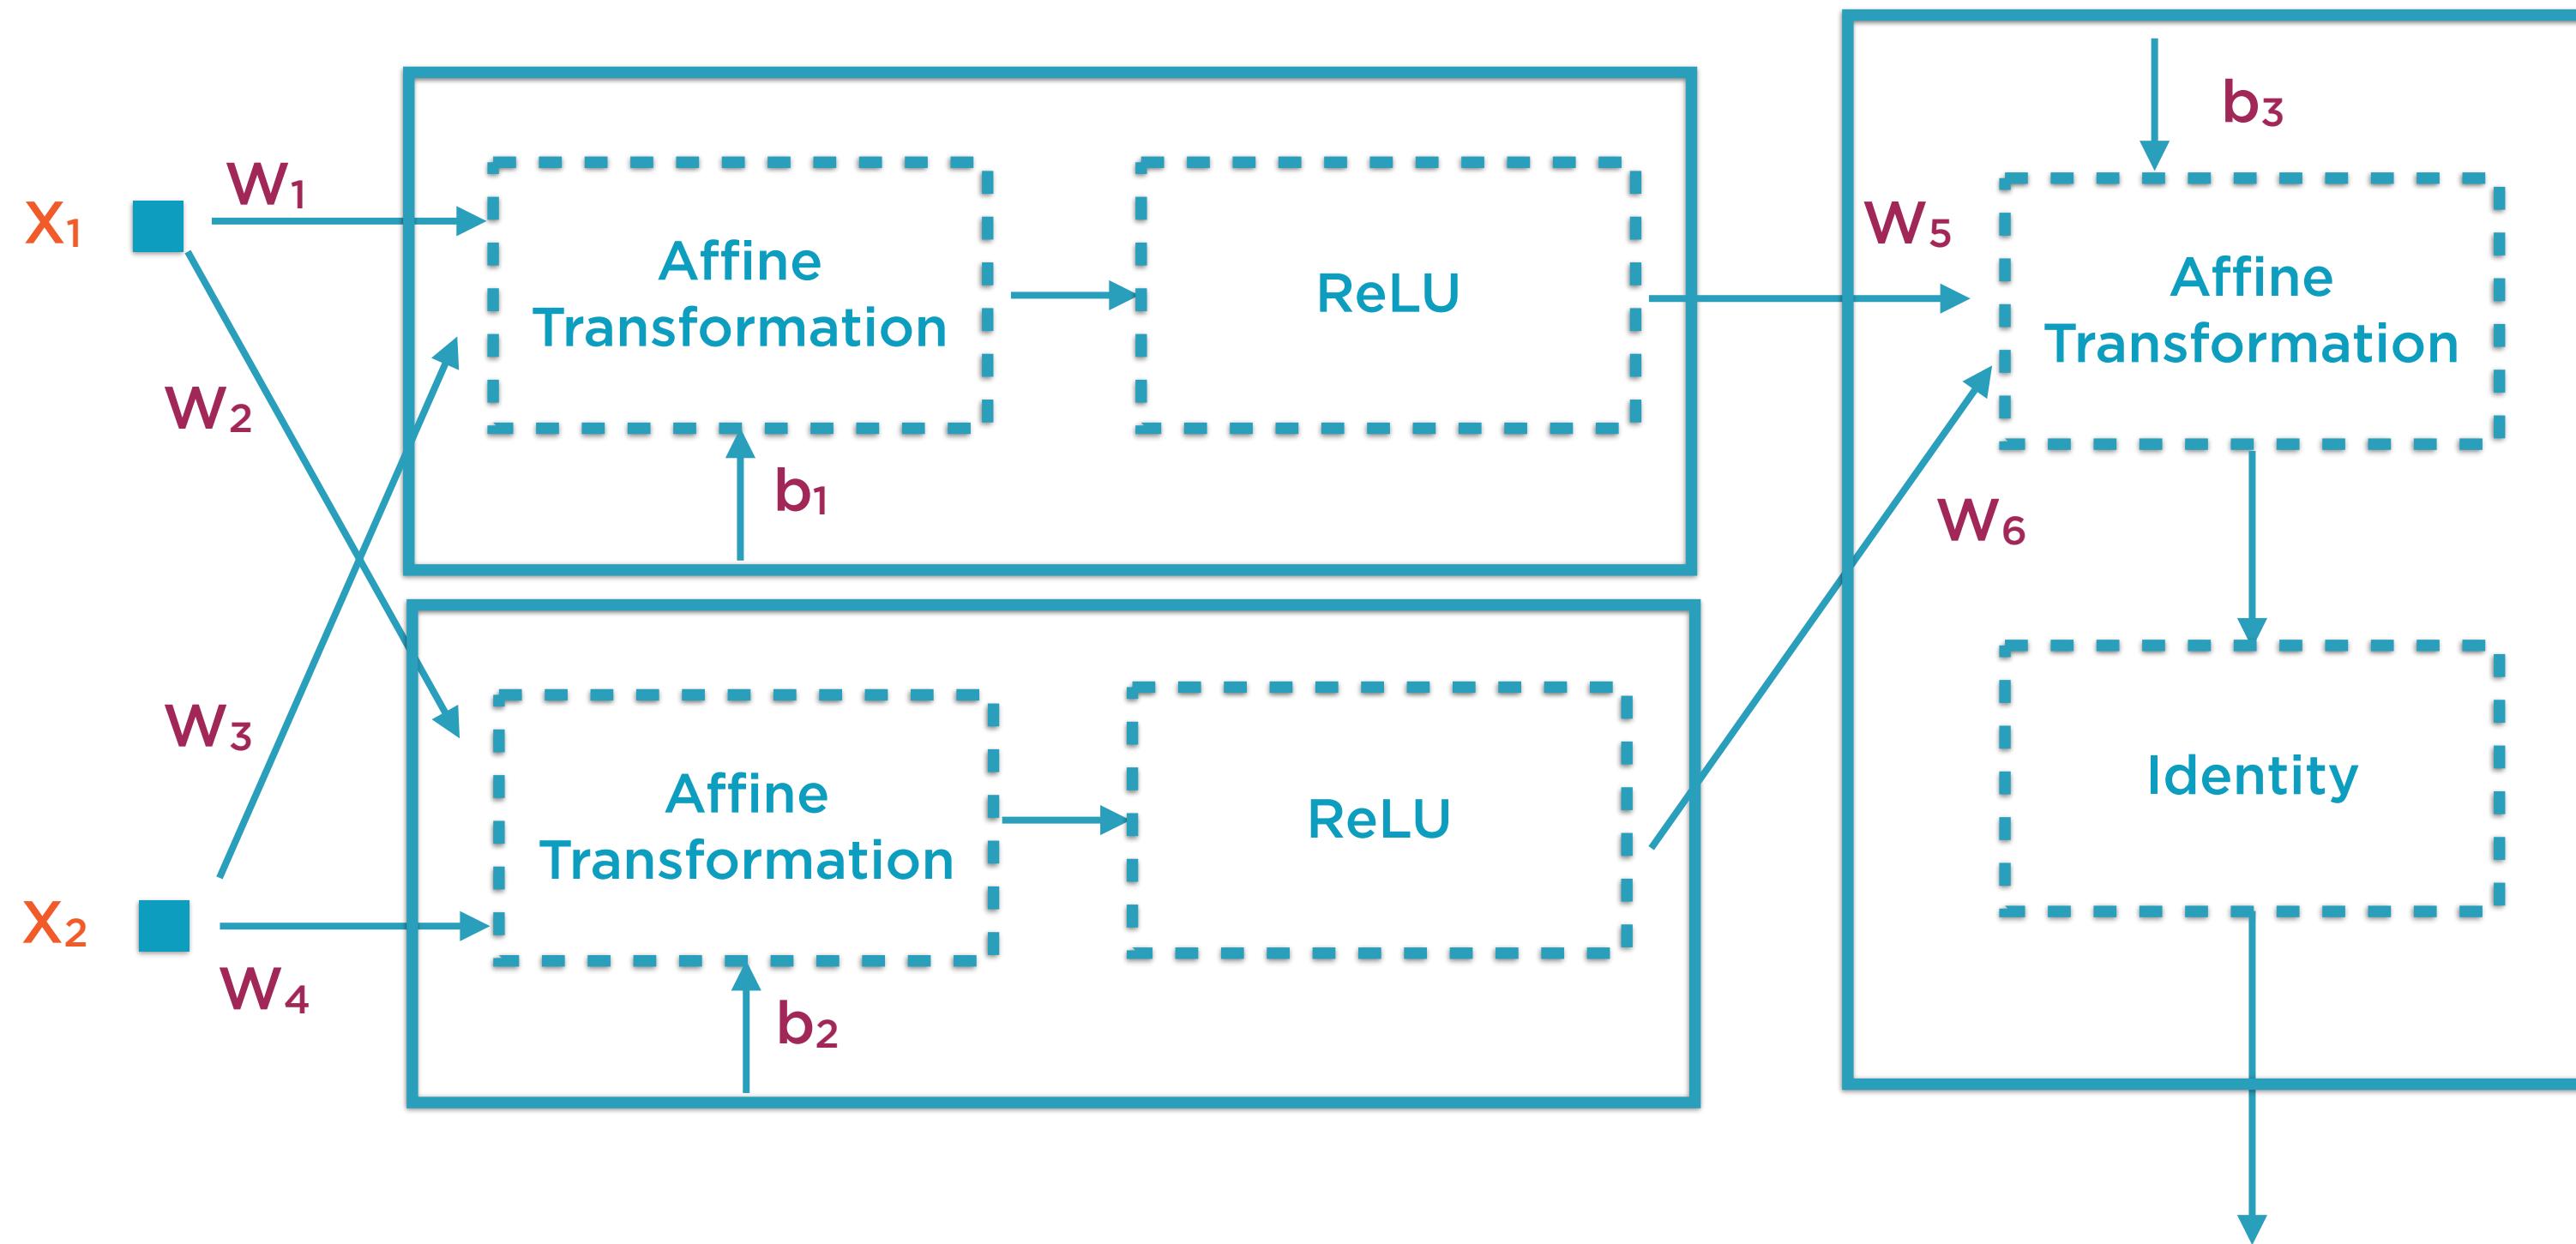
3-Neuron XOR



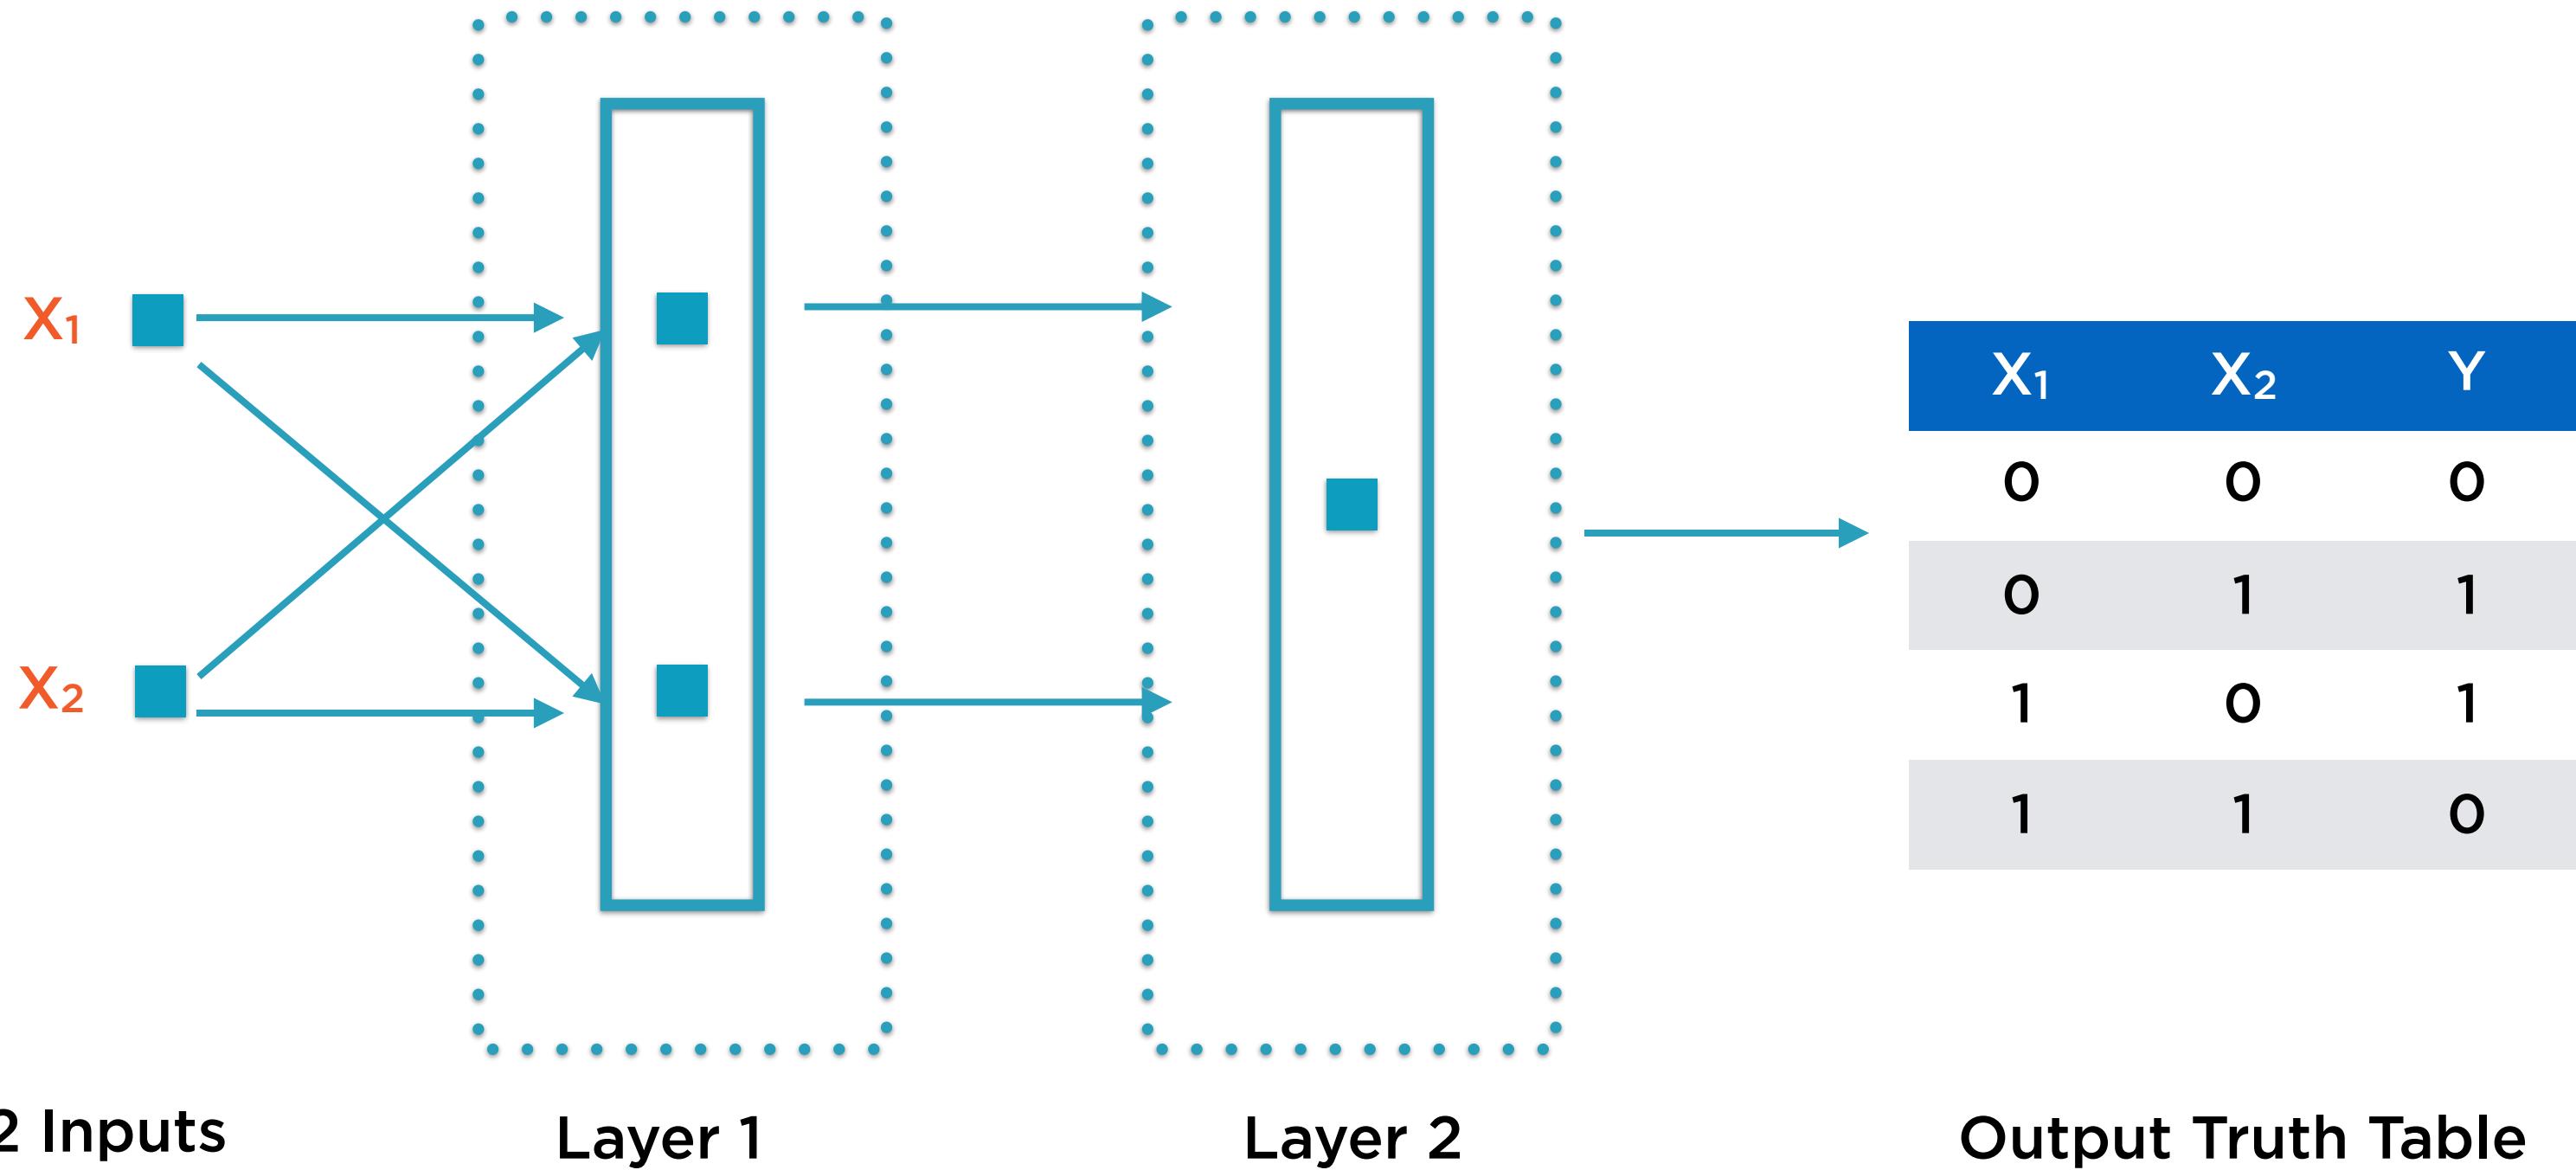
3-Neuron XOR



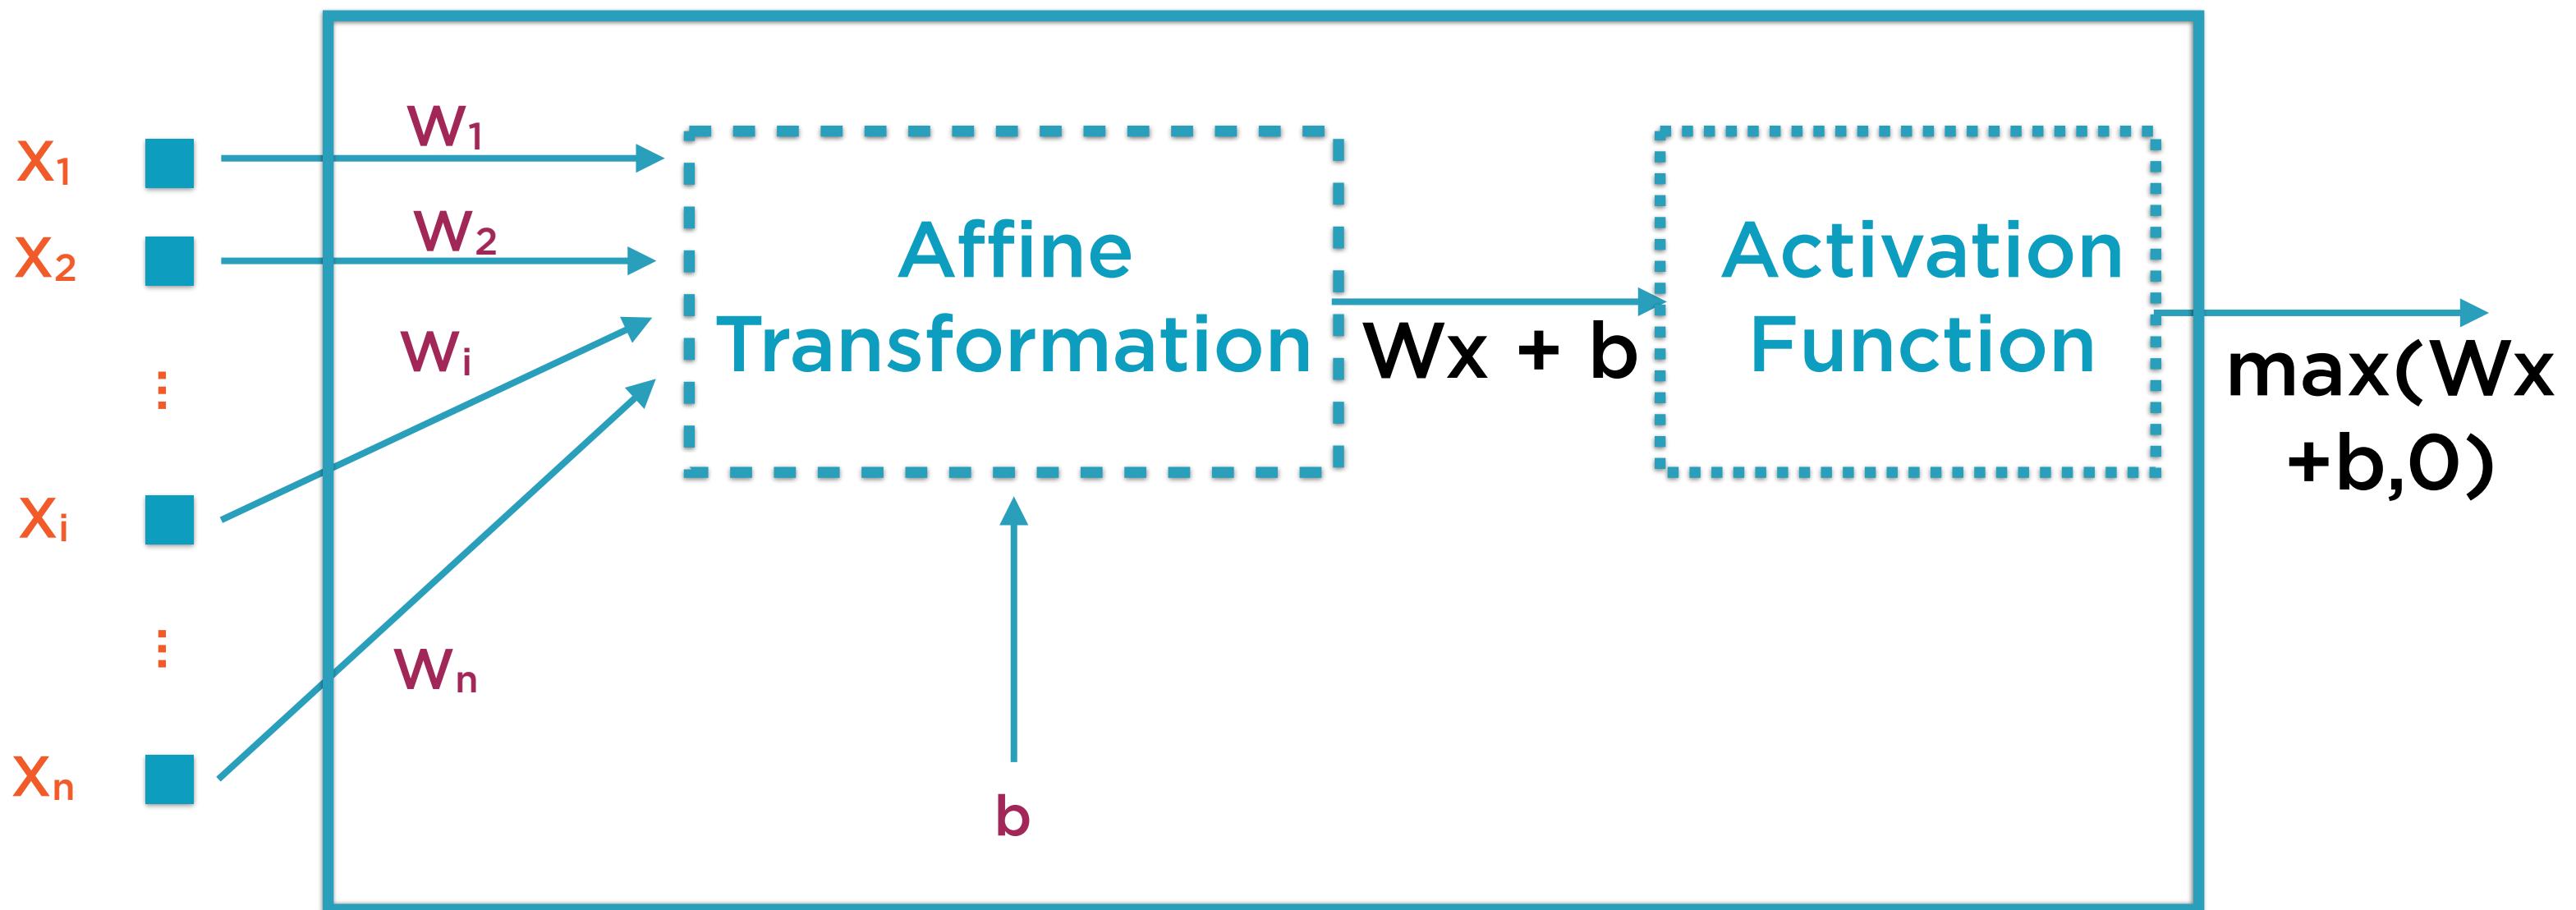
“2-Layer Feed-forward Neural Network”



XOR: 3 Neurons, 2 Layers

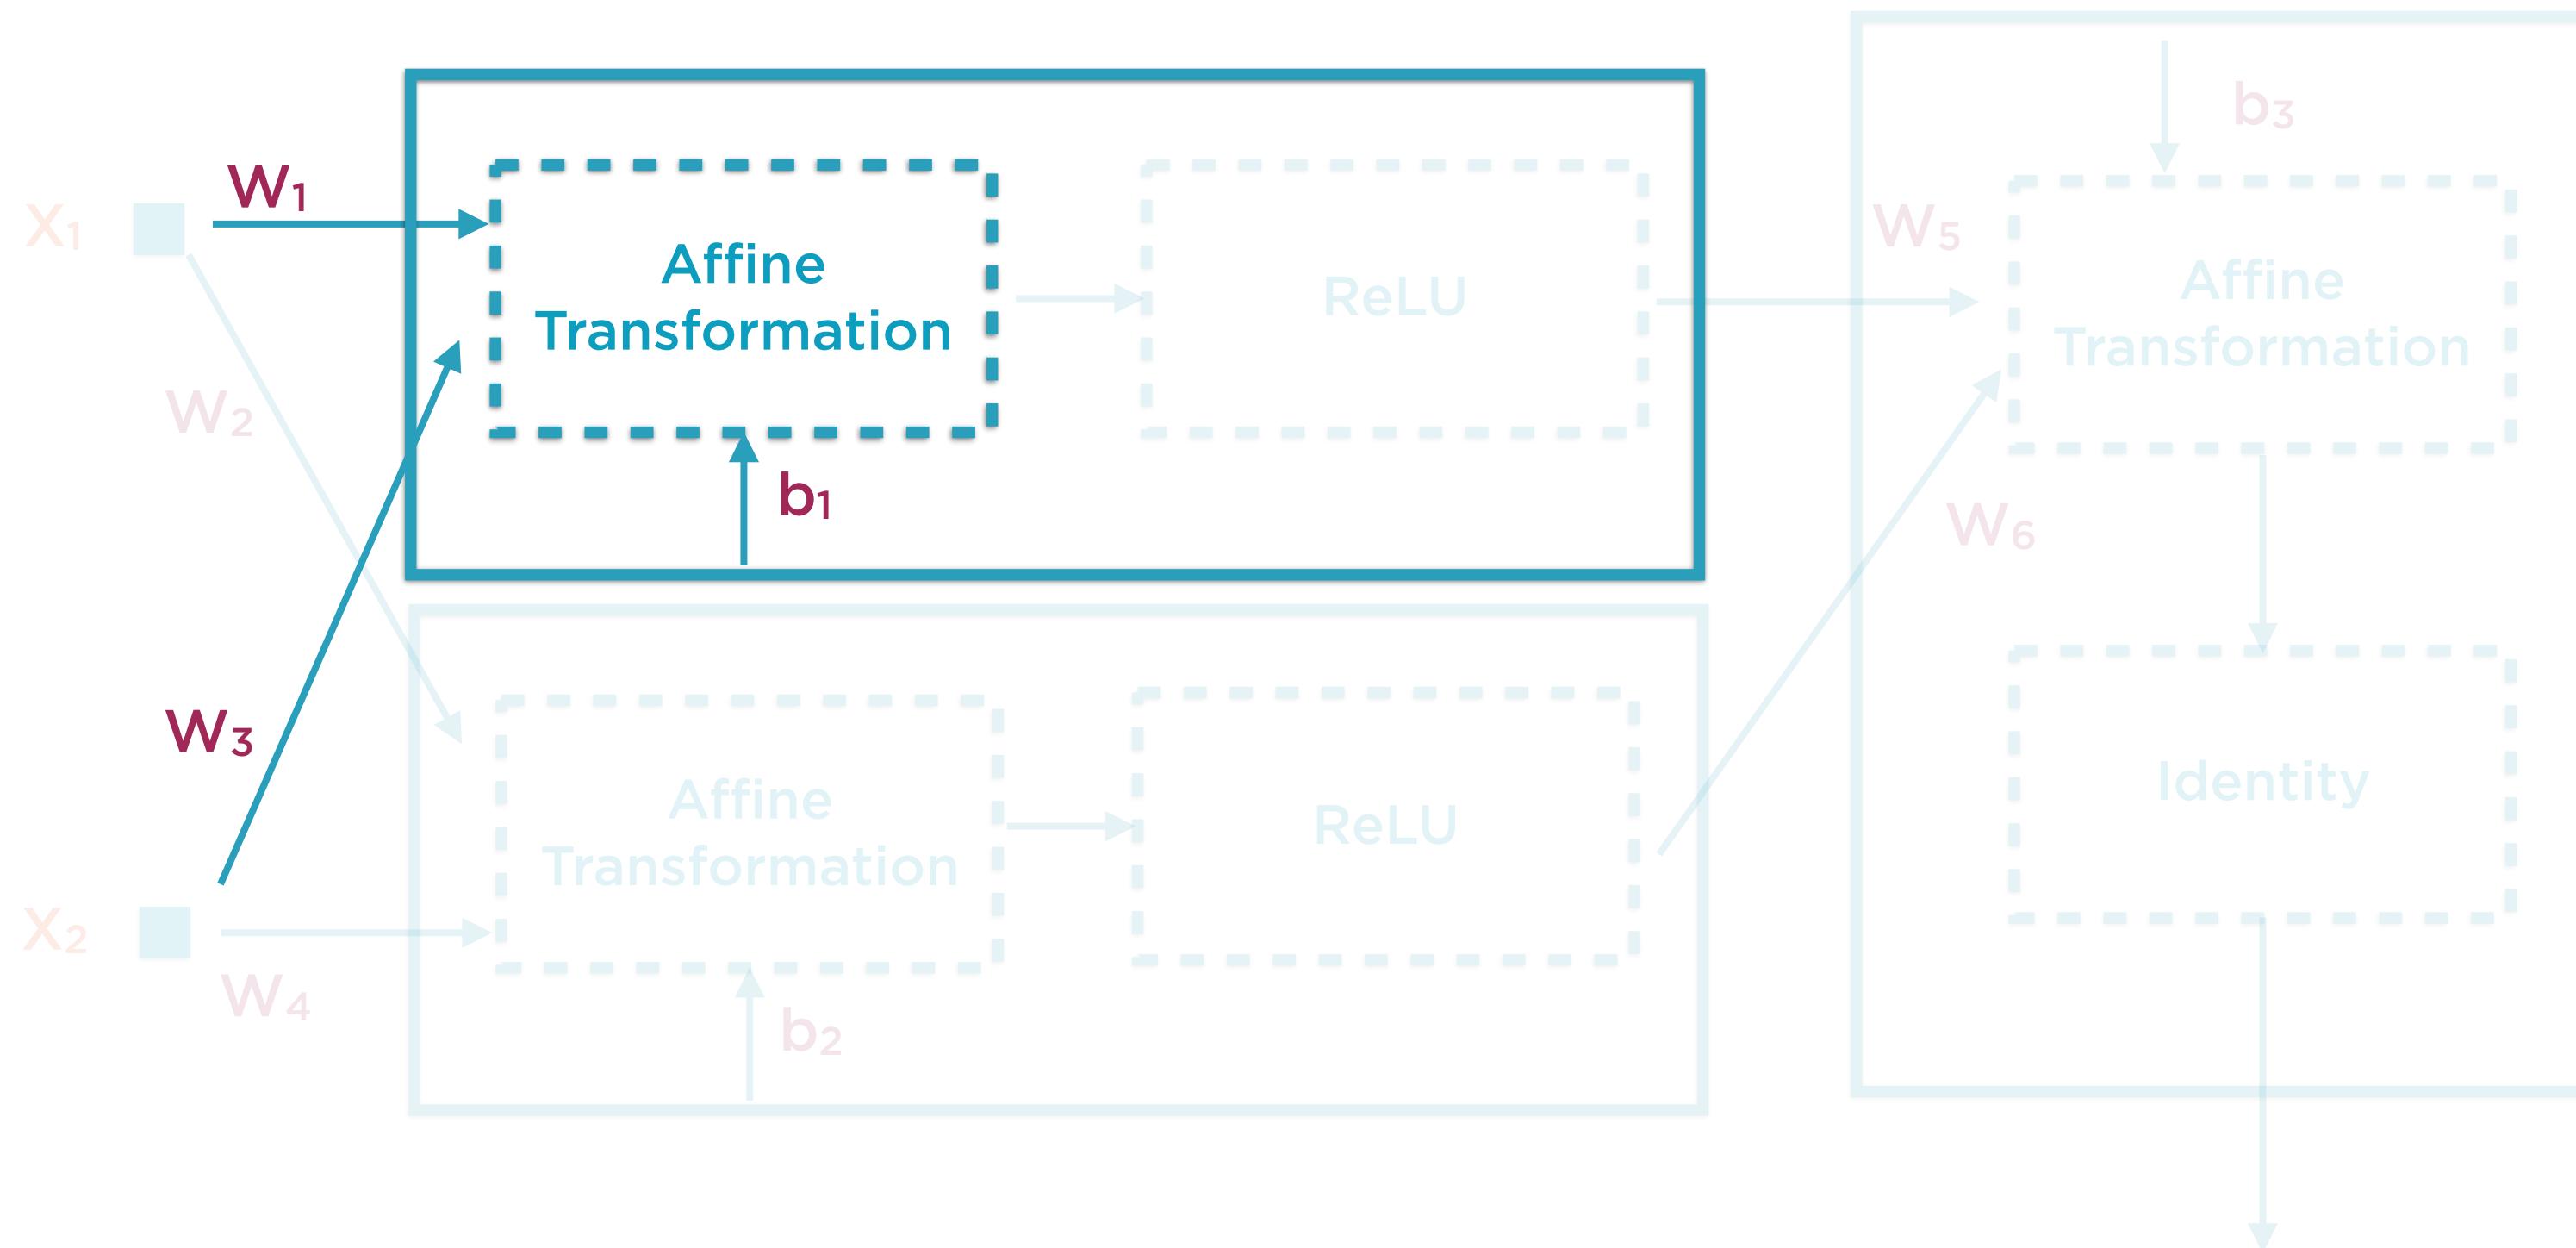


Operation of a Single Neuron

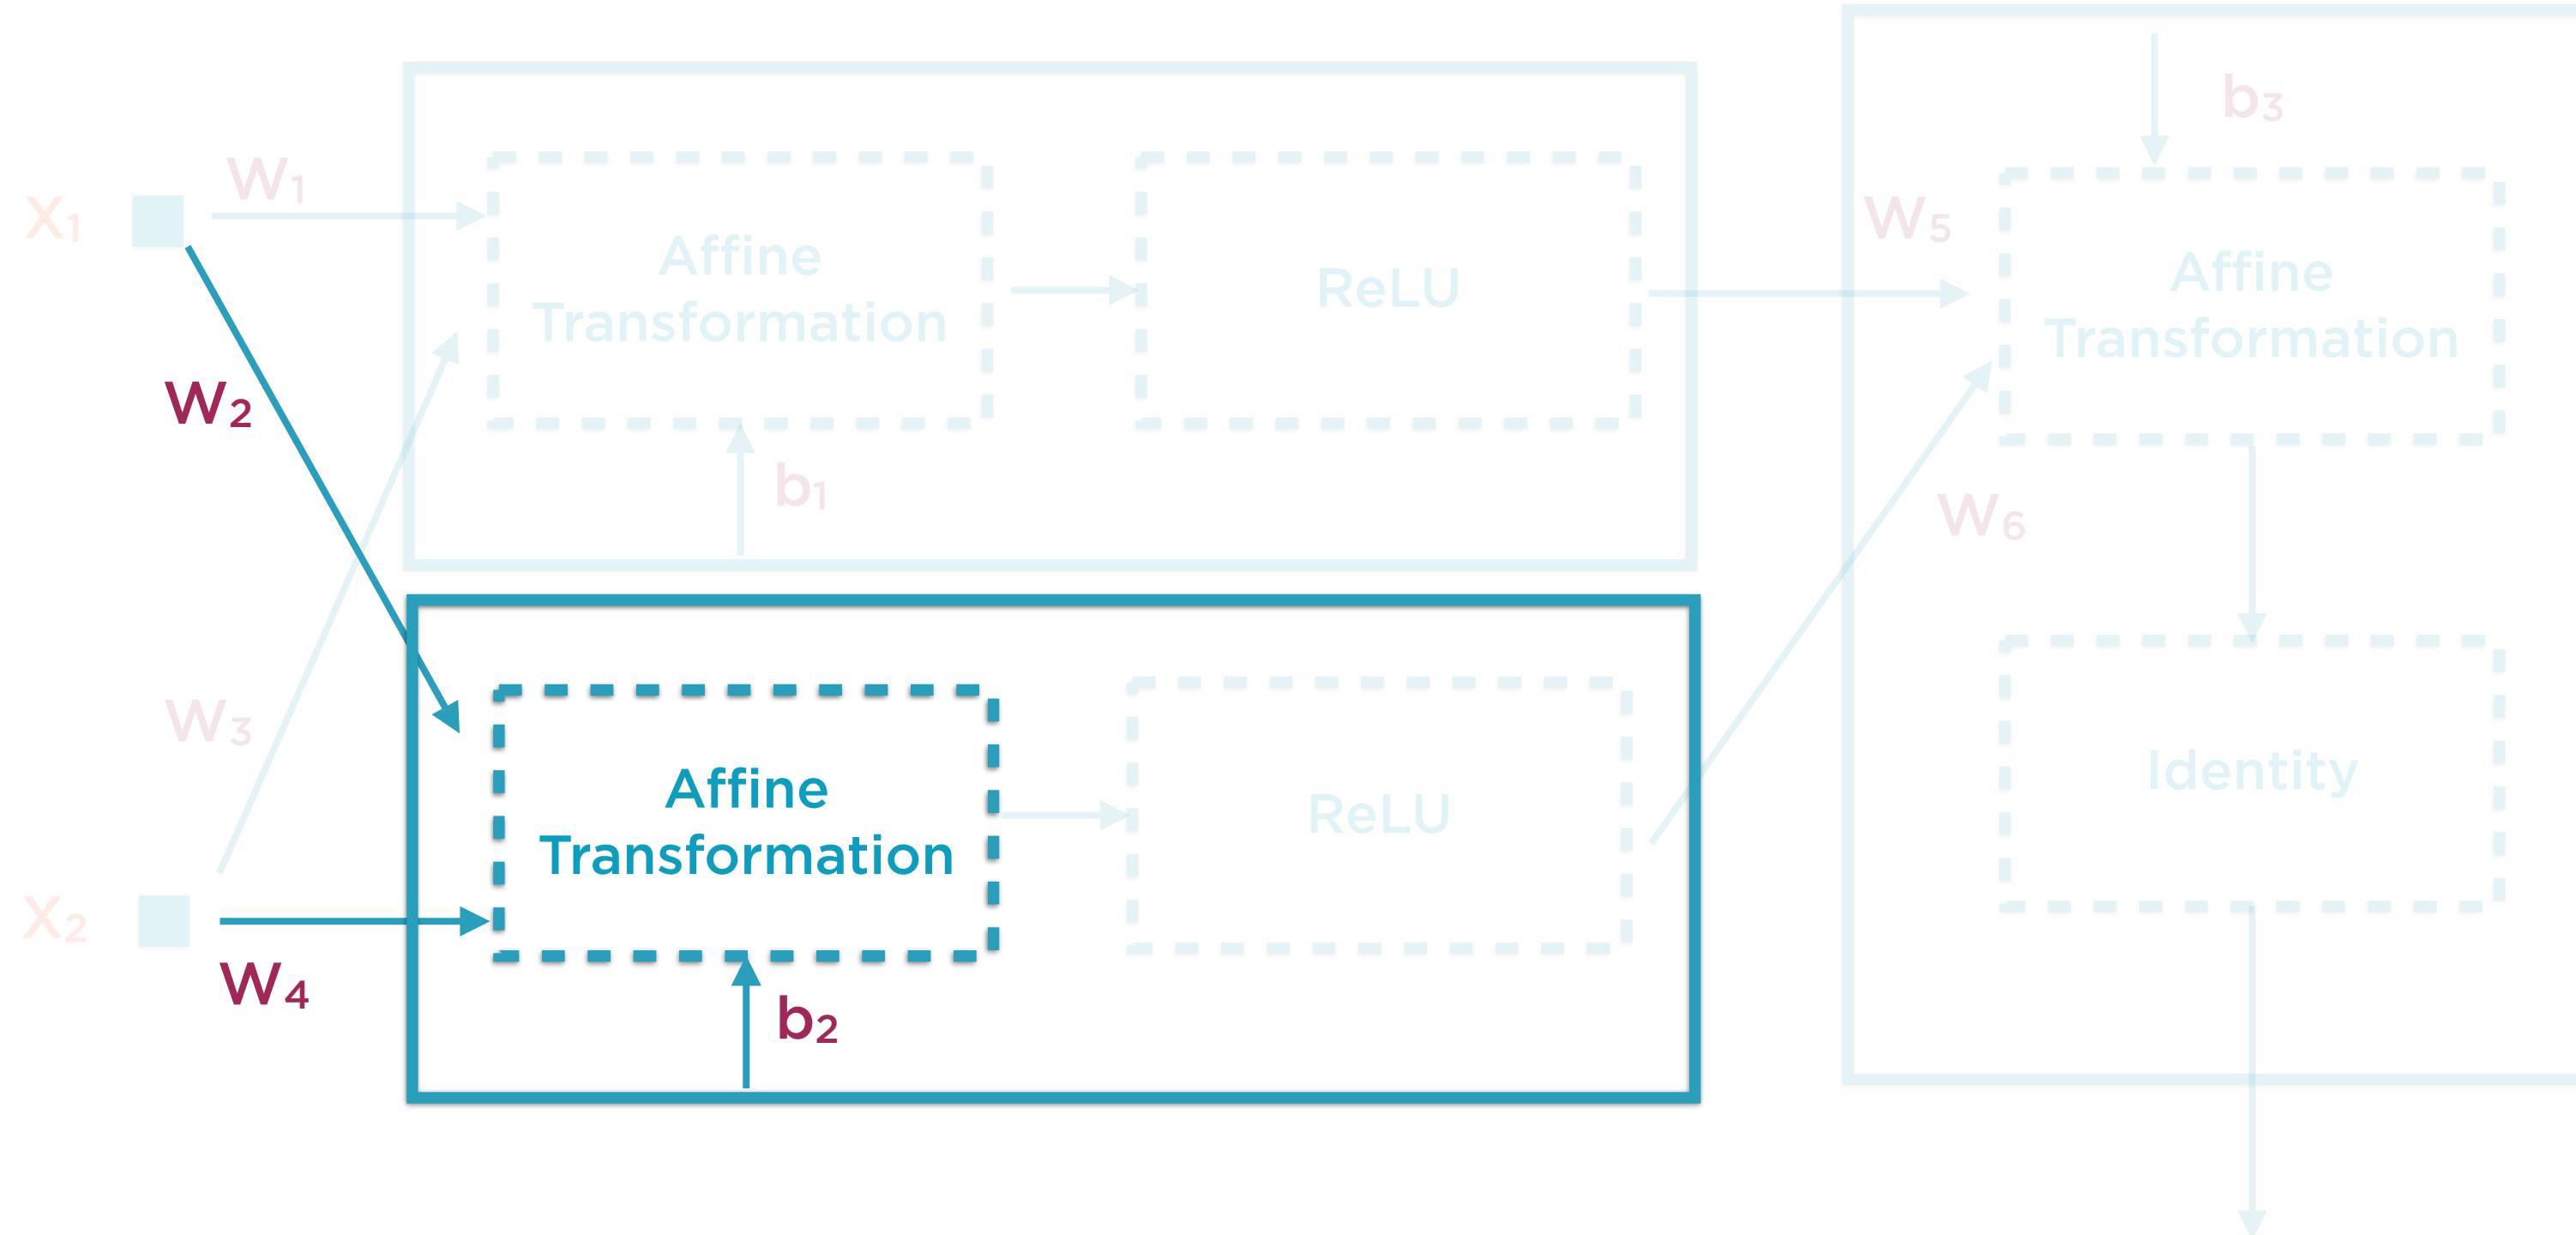


Each neuron has weights and a bias that must be calculated by the training algorithm (done for us by TensorFlow)

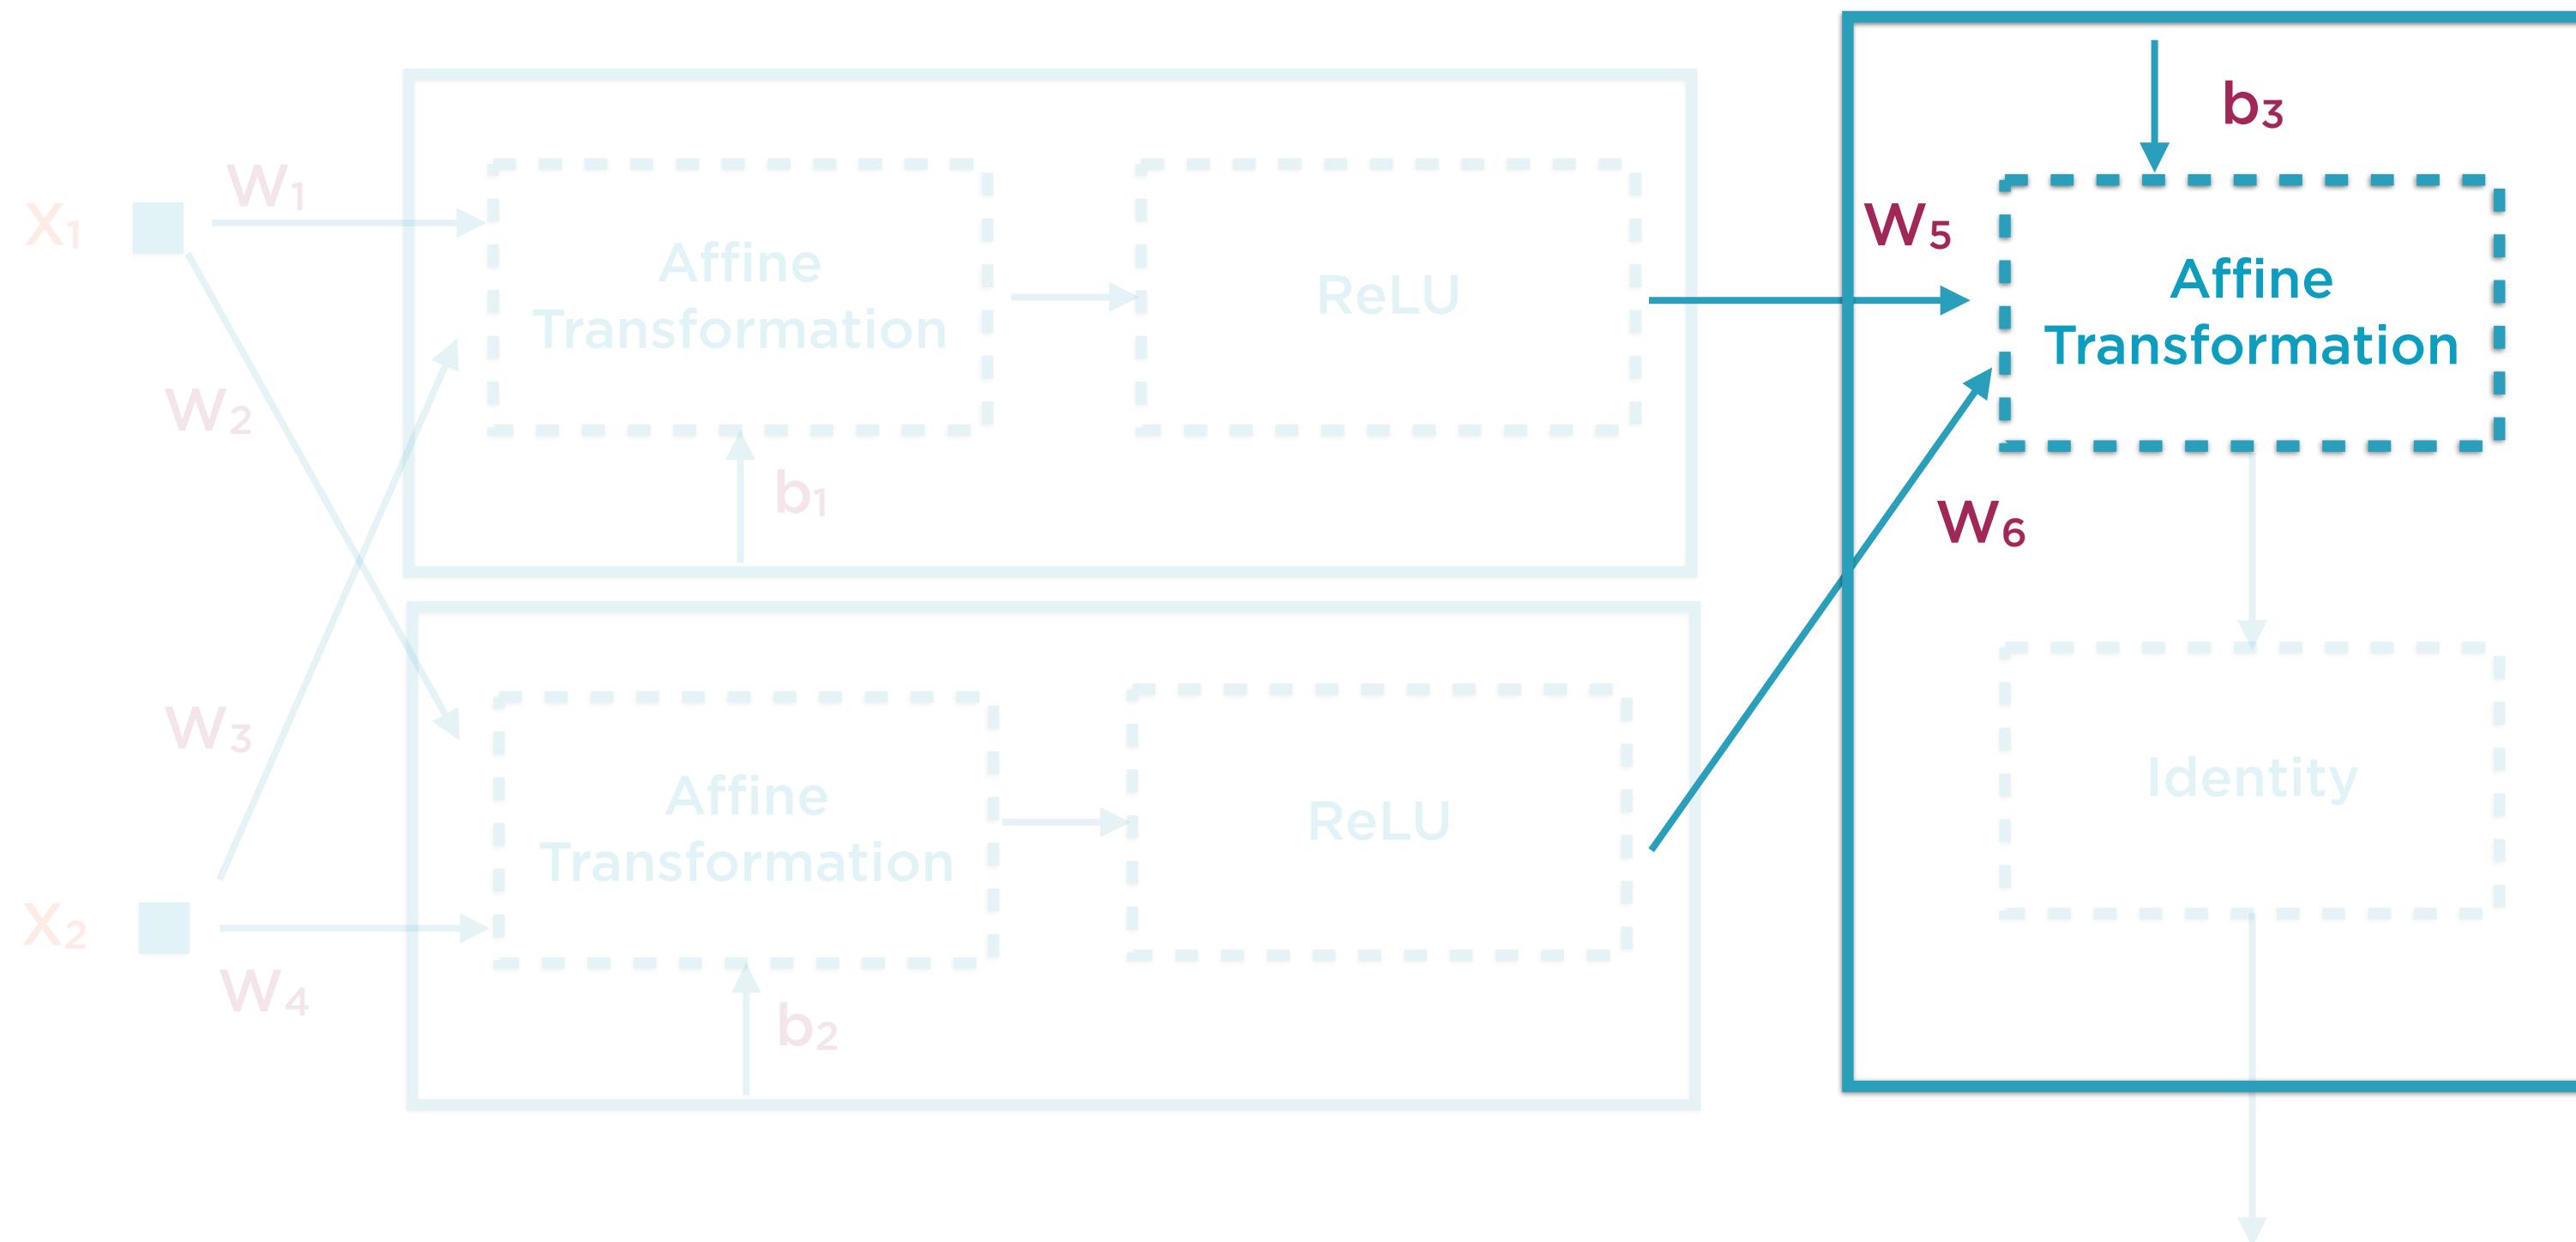
Weights and Bias of Neuron #1



Weights and Bias of Neuron #2

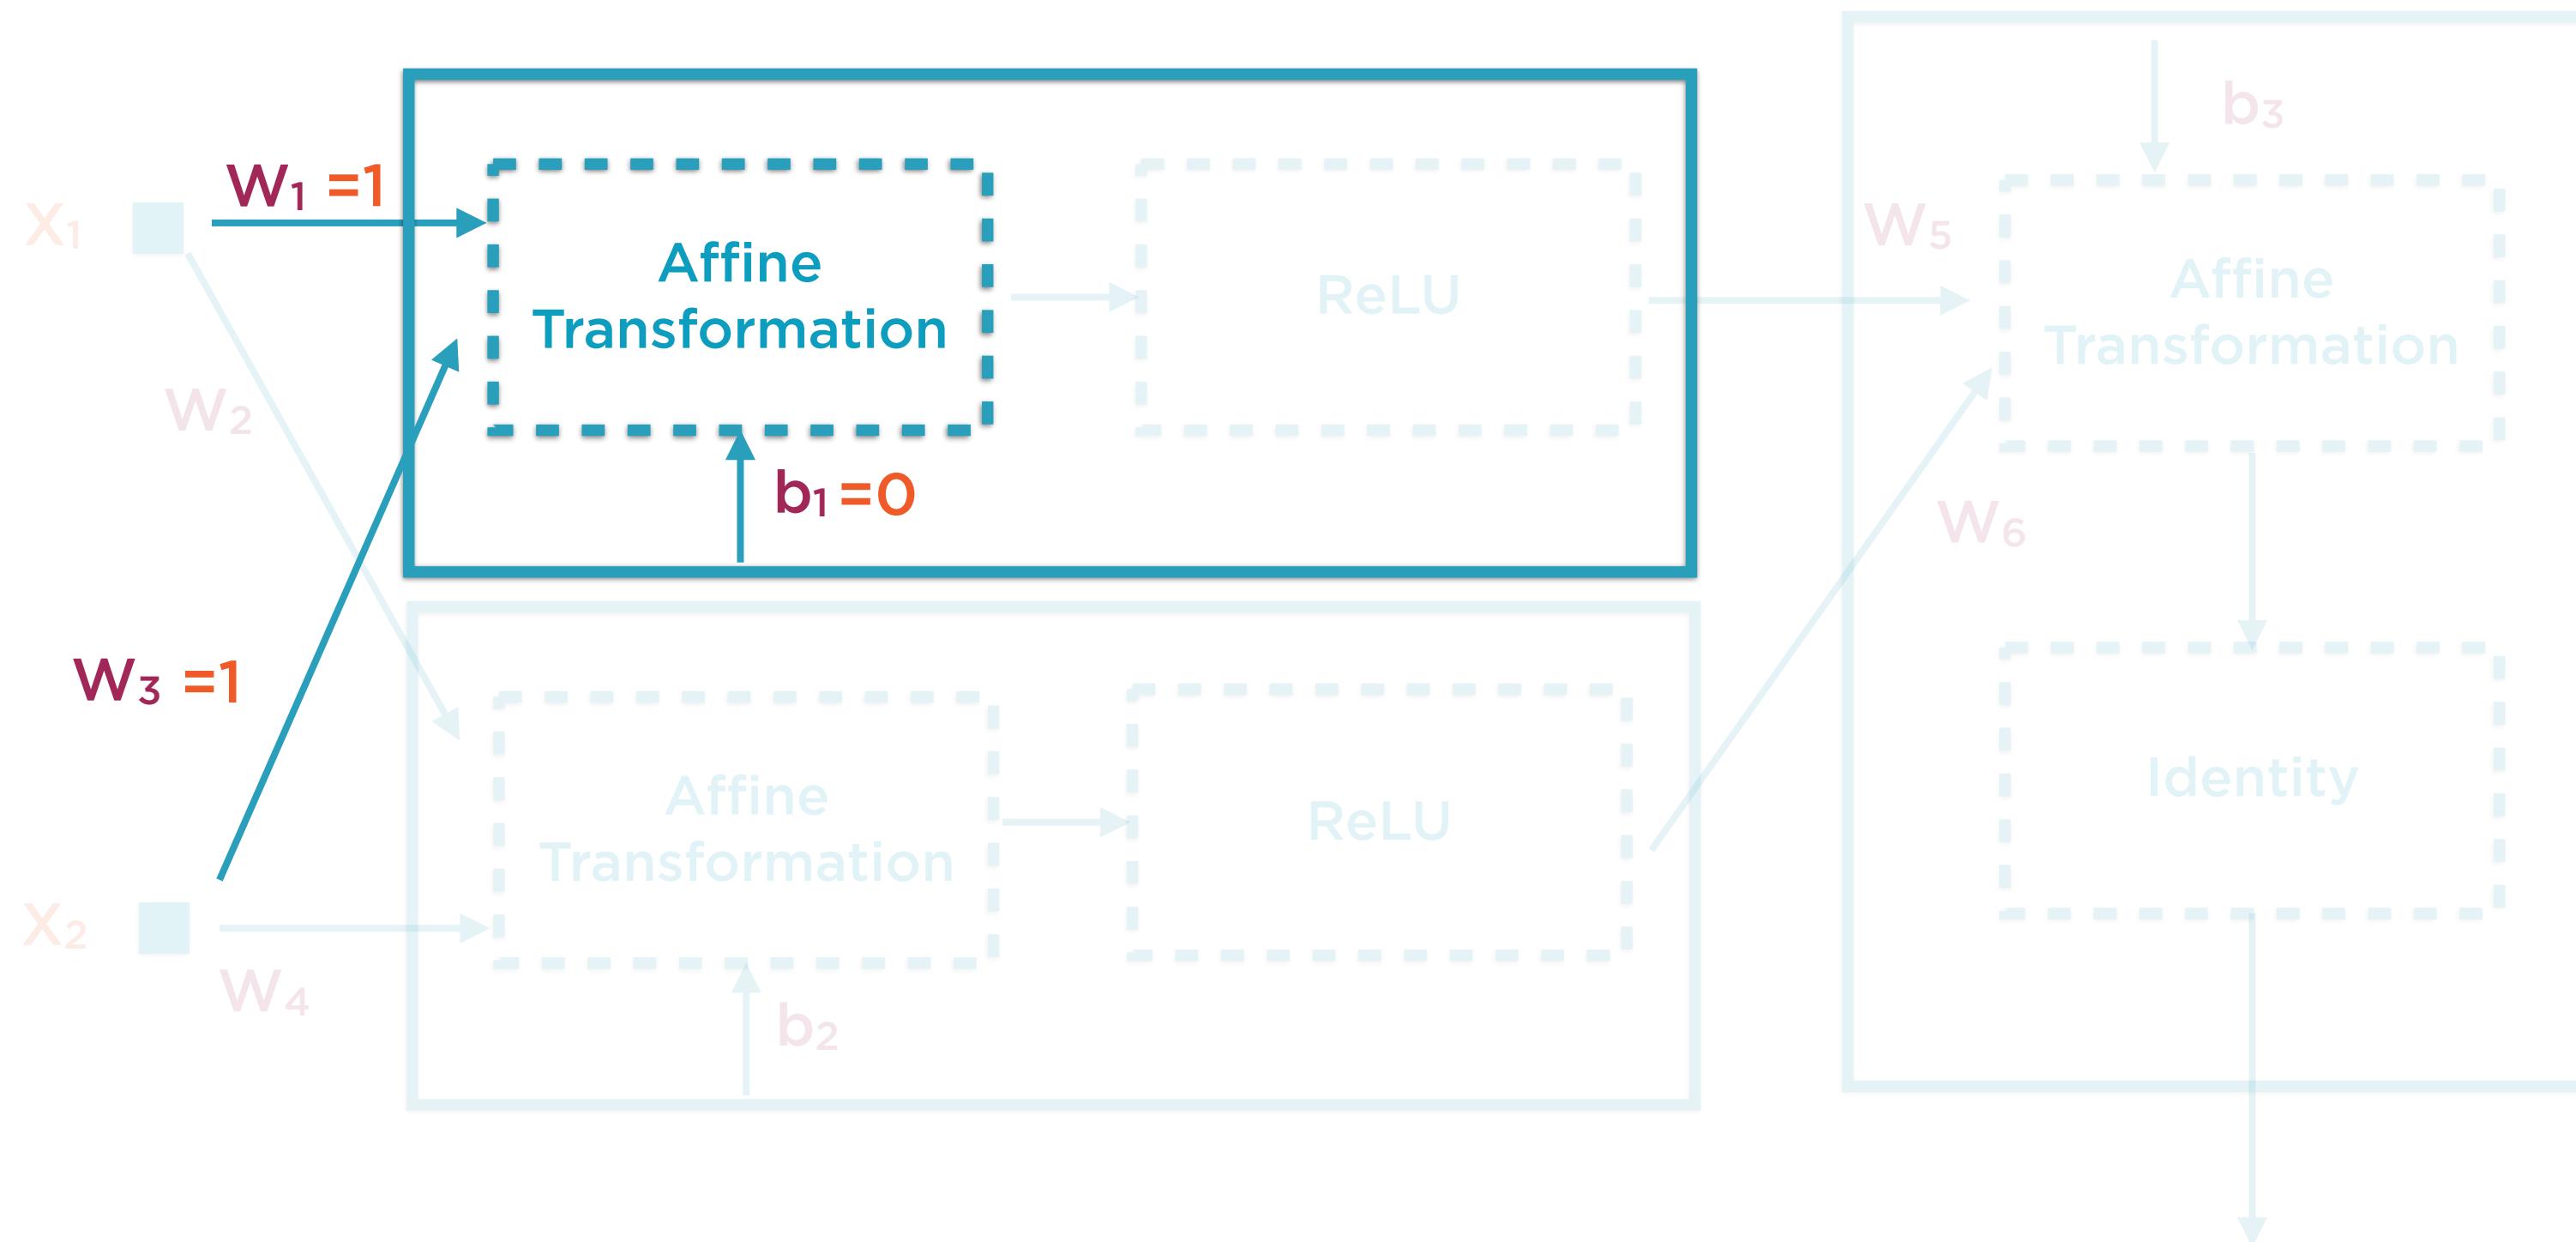


Weights and Bias of Neuron #3

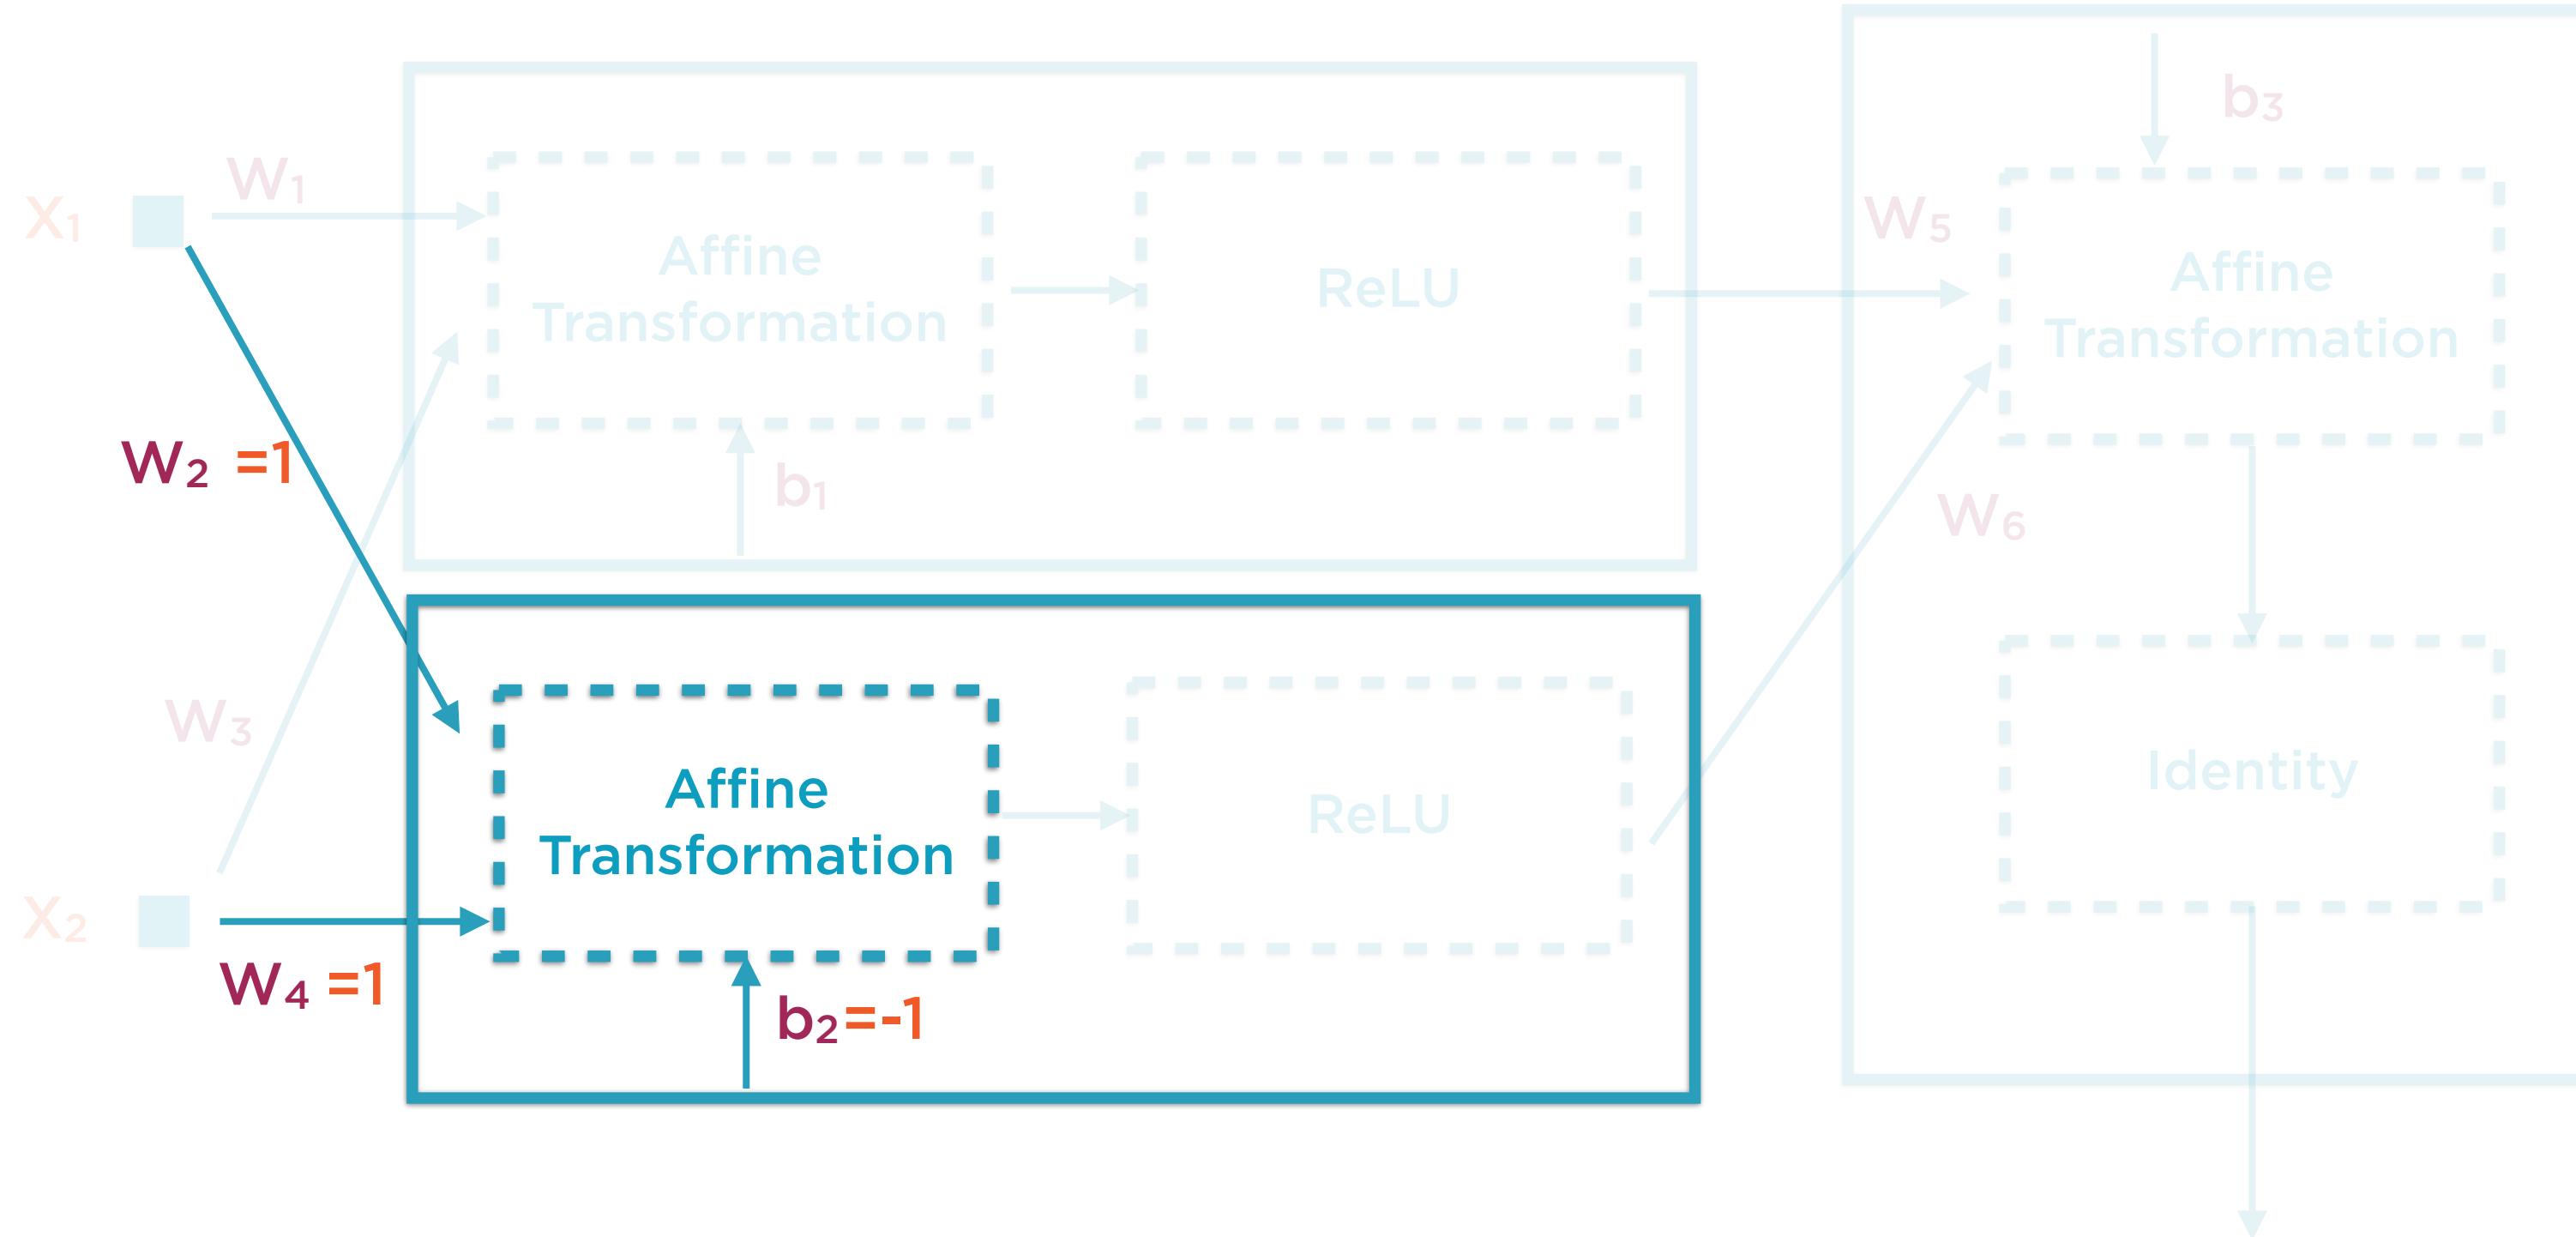


The weights and biases of individual neurons are determined during the training process

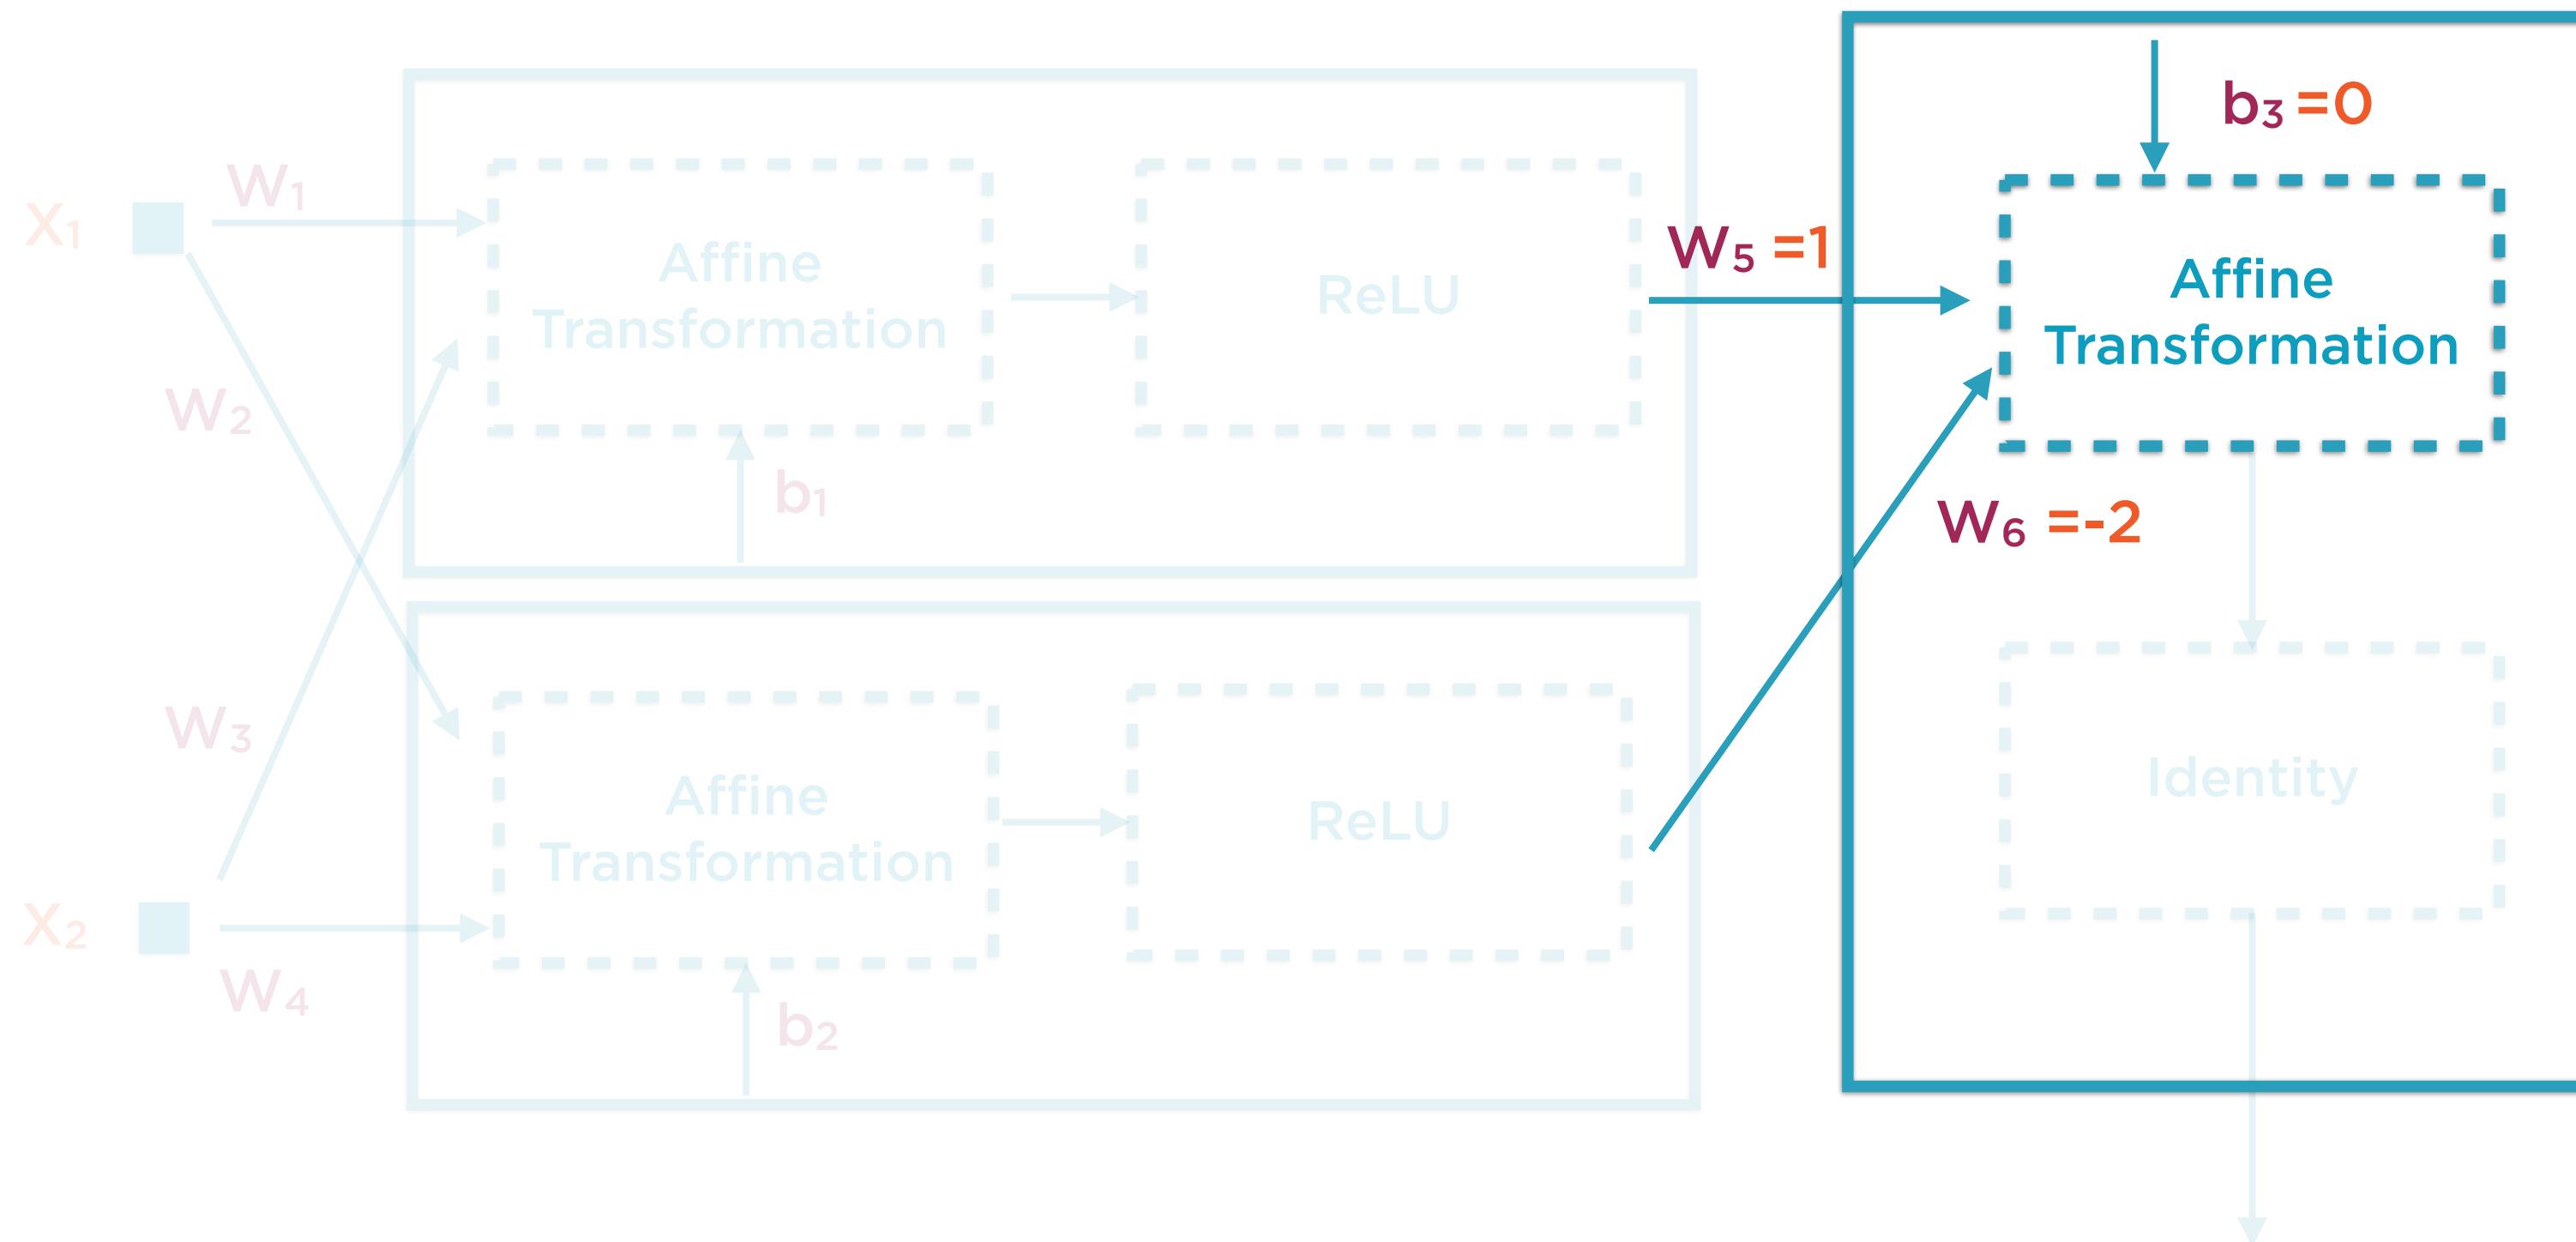
Weights and Bias of Neuron #1



Weights and Bias of Neuron #2



Weights and Bias of Neuron #3

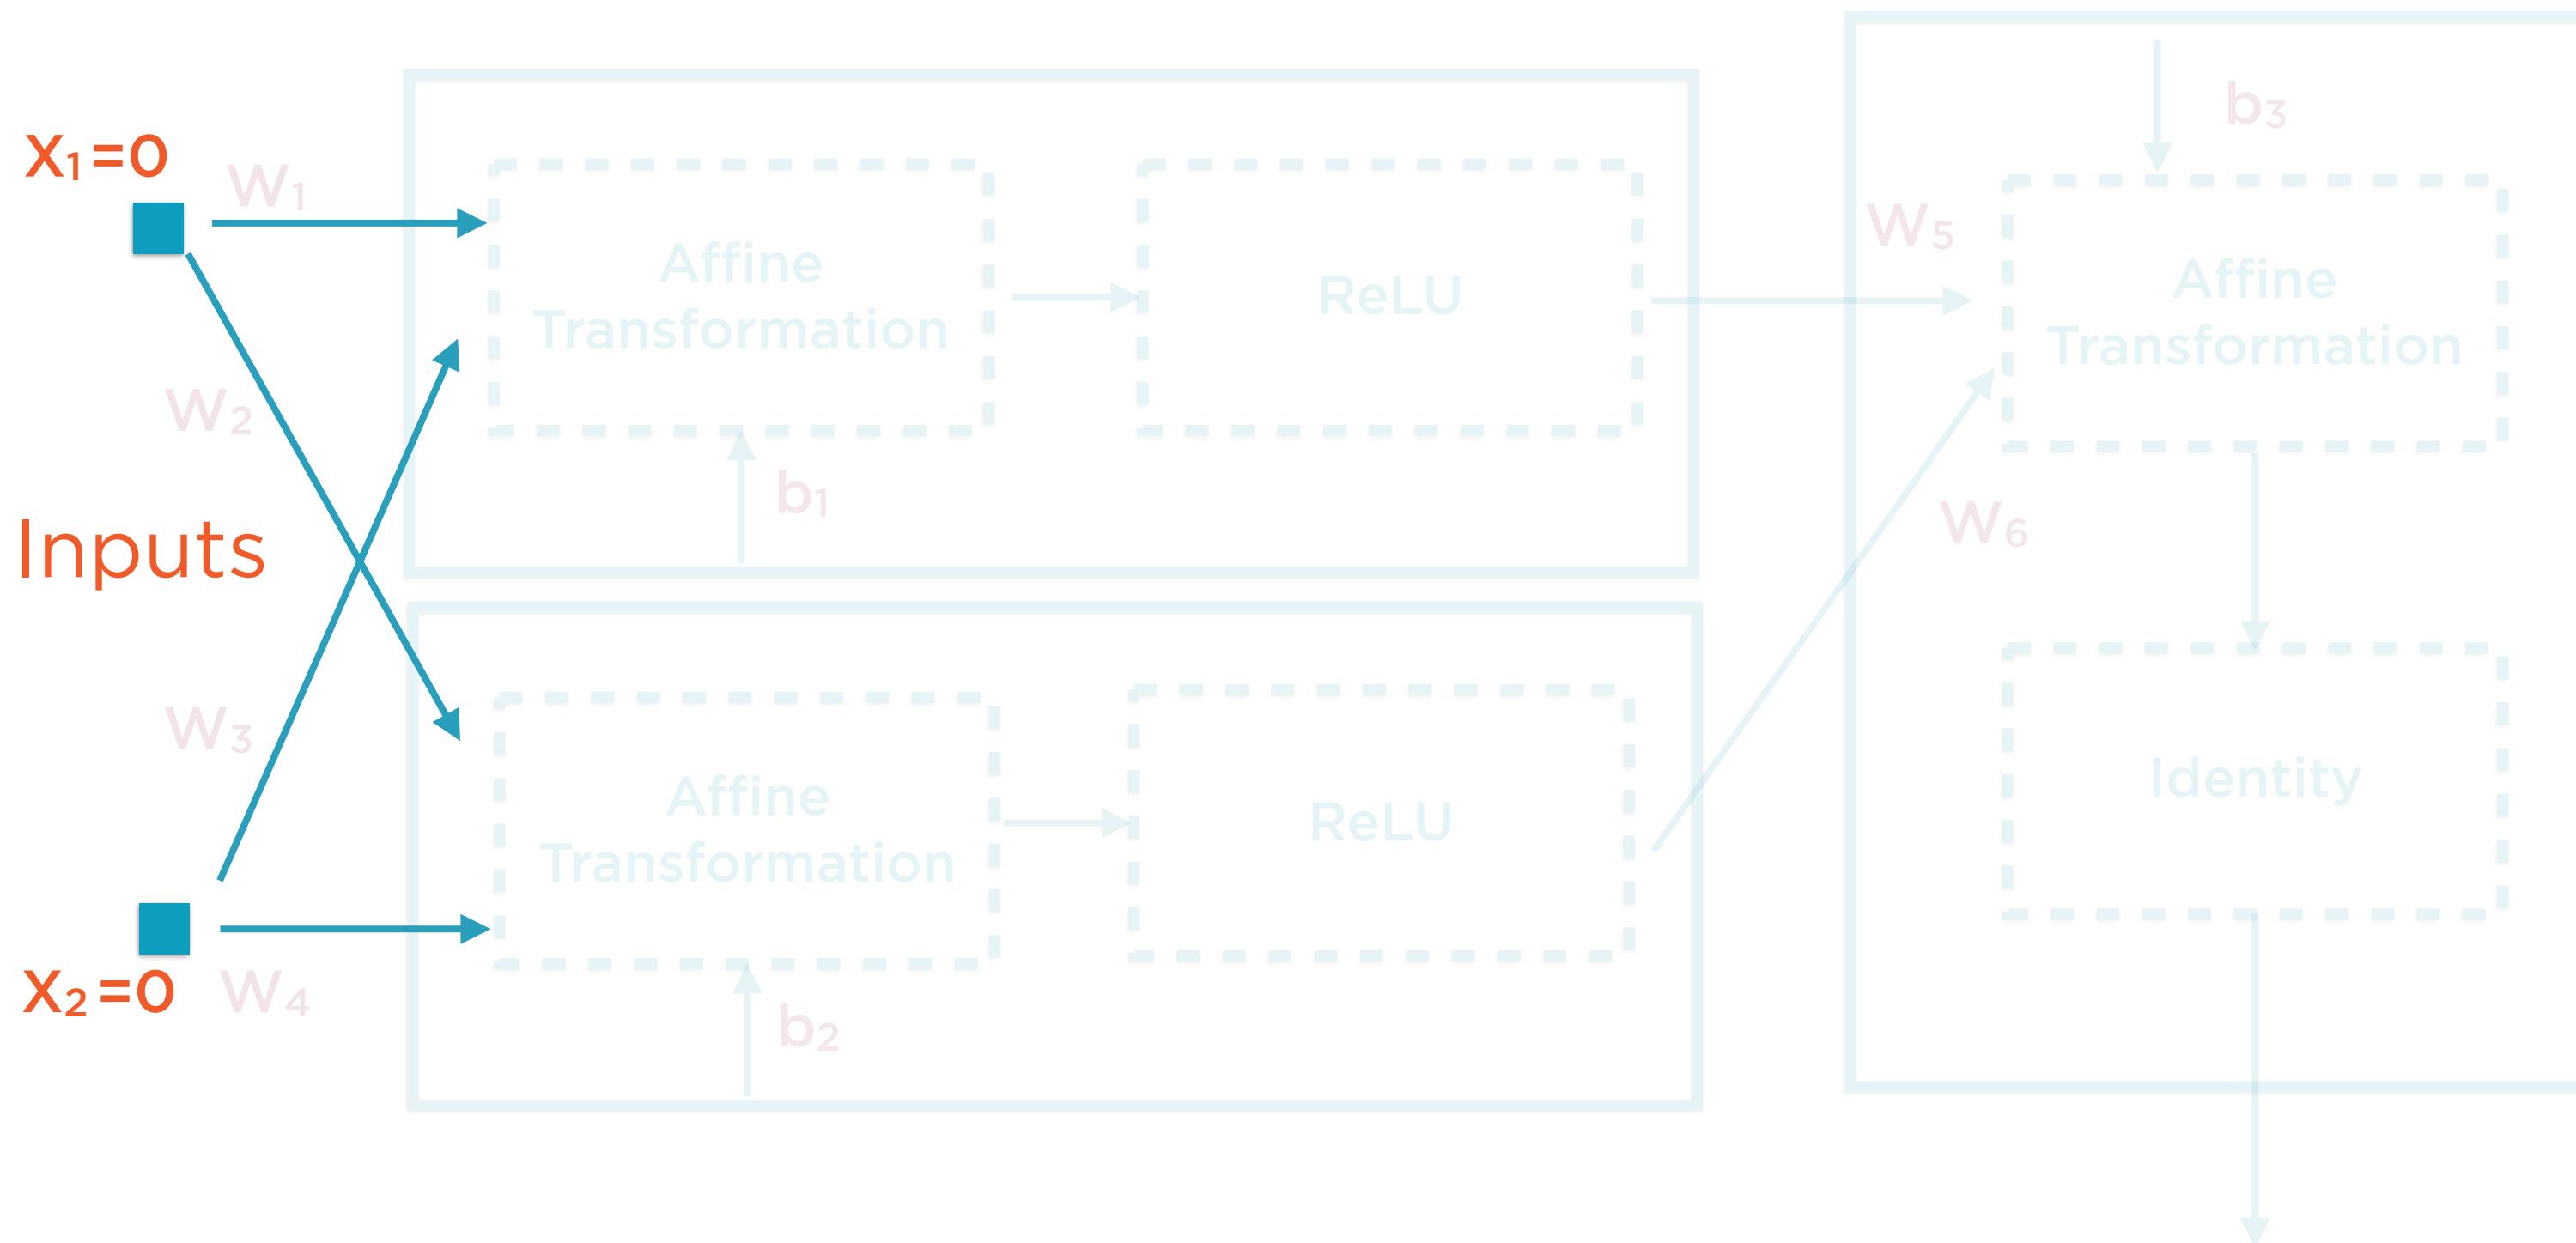


X ₁	X ₂	Y
0	0	0
0	1	1
1	0	1
1	1	0

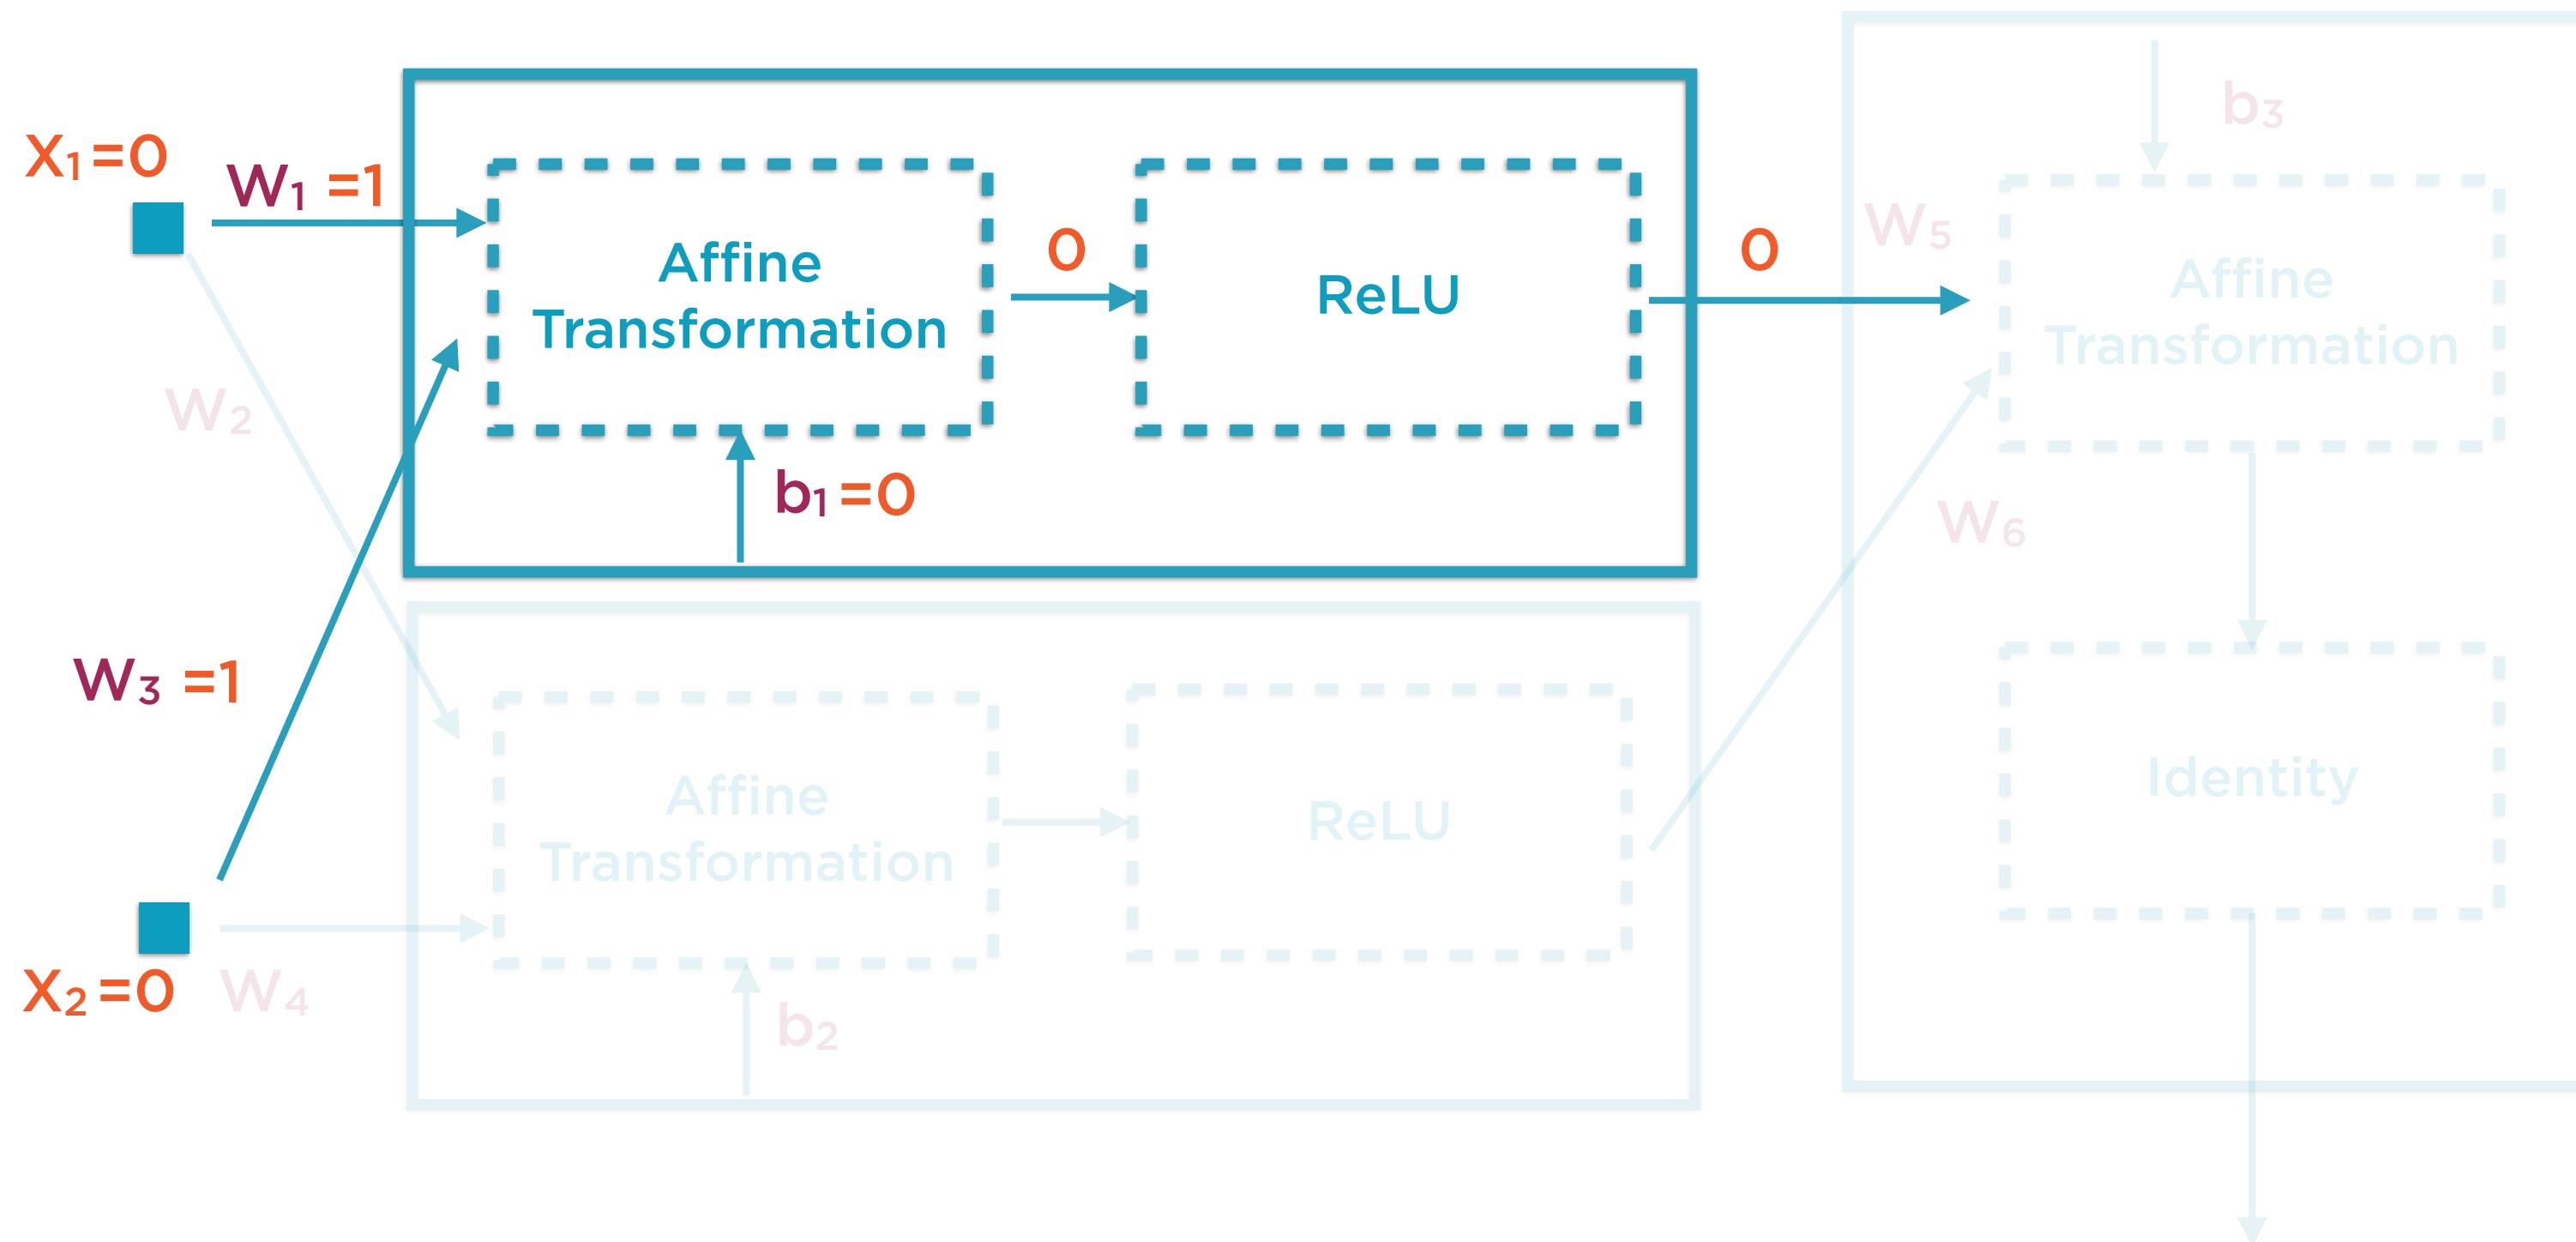
“Learning” XOR

Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

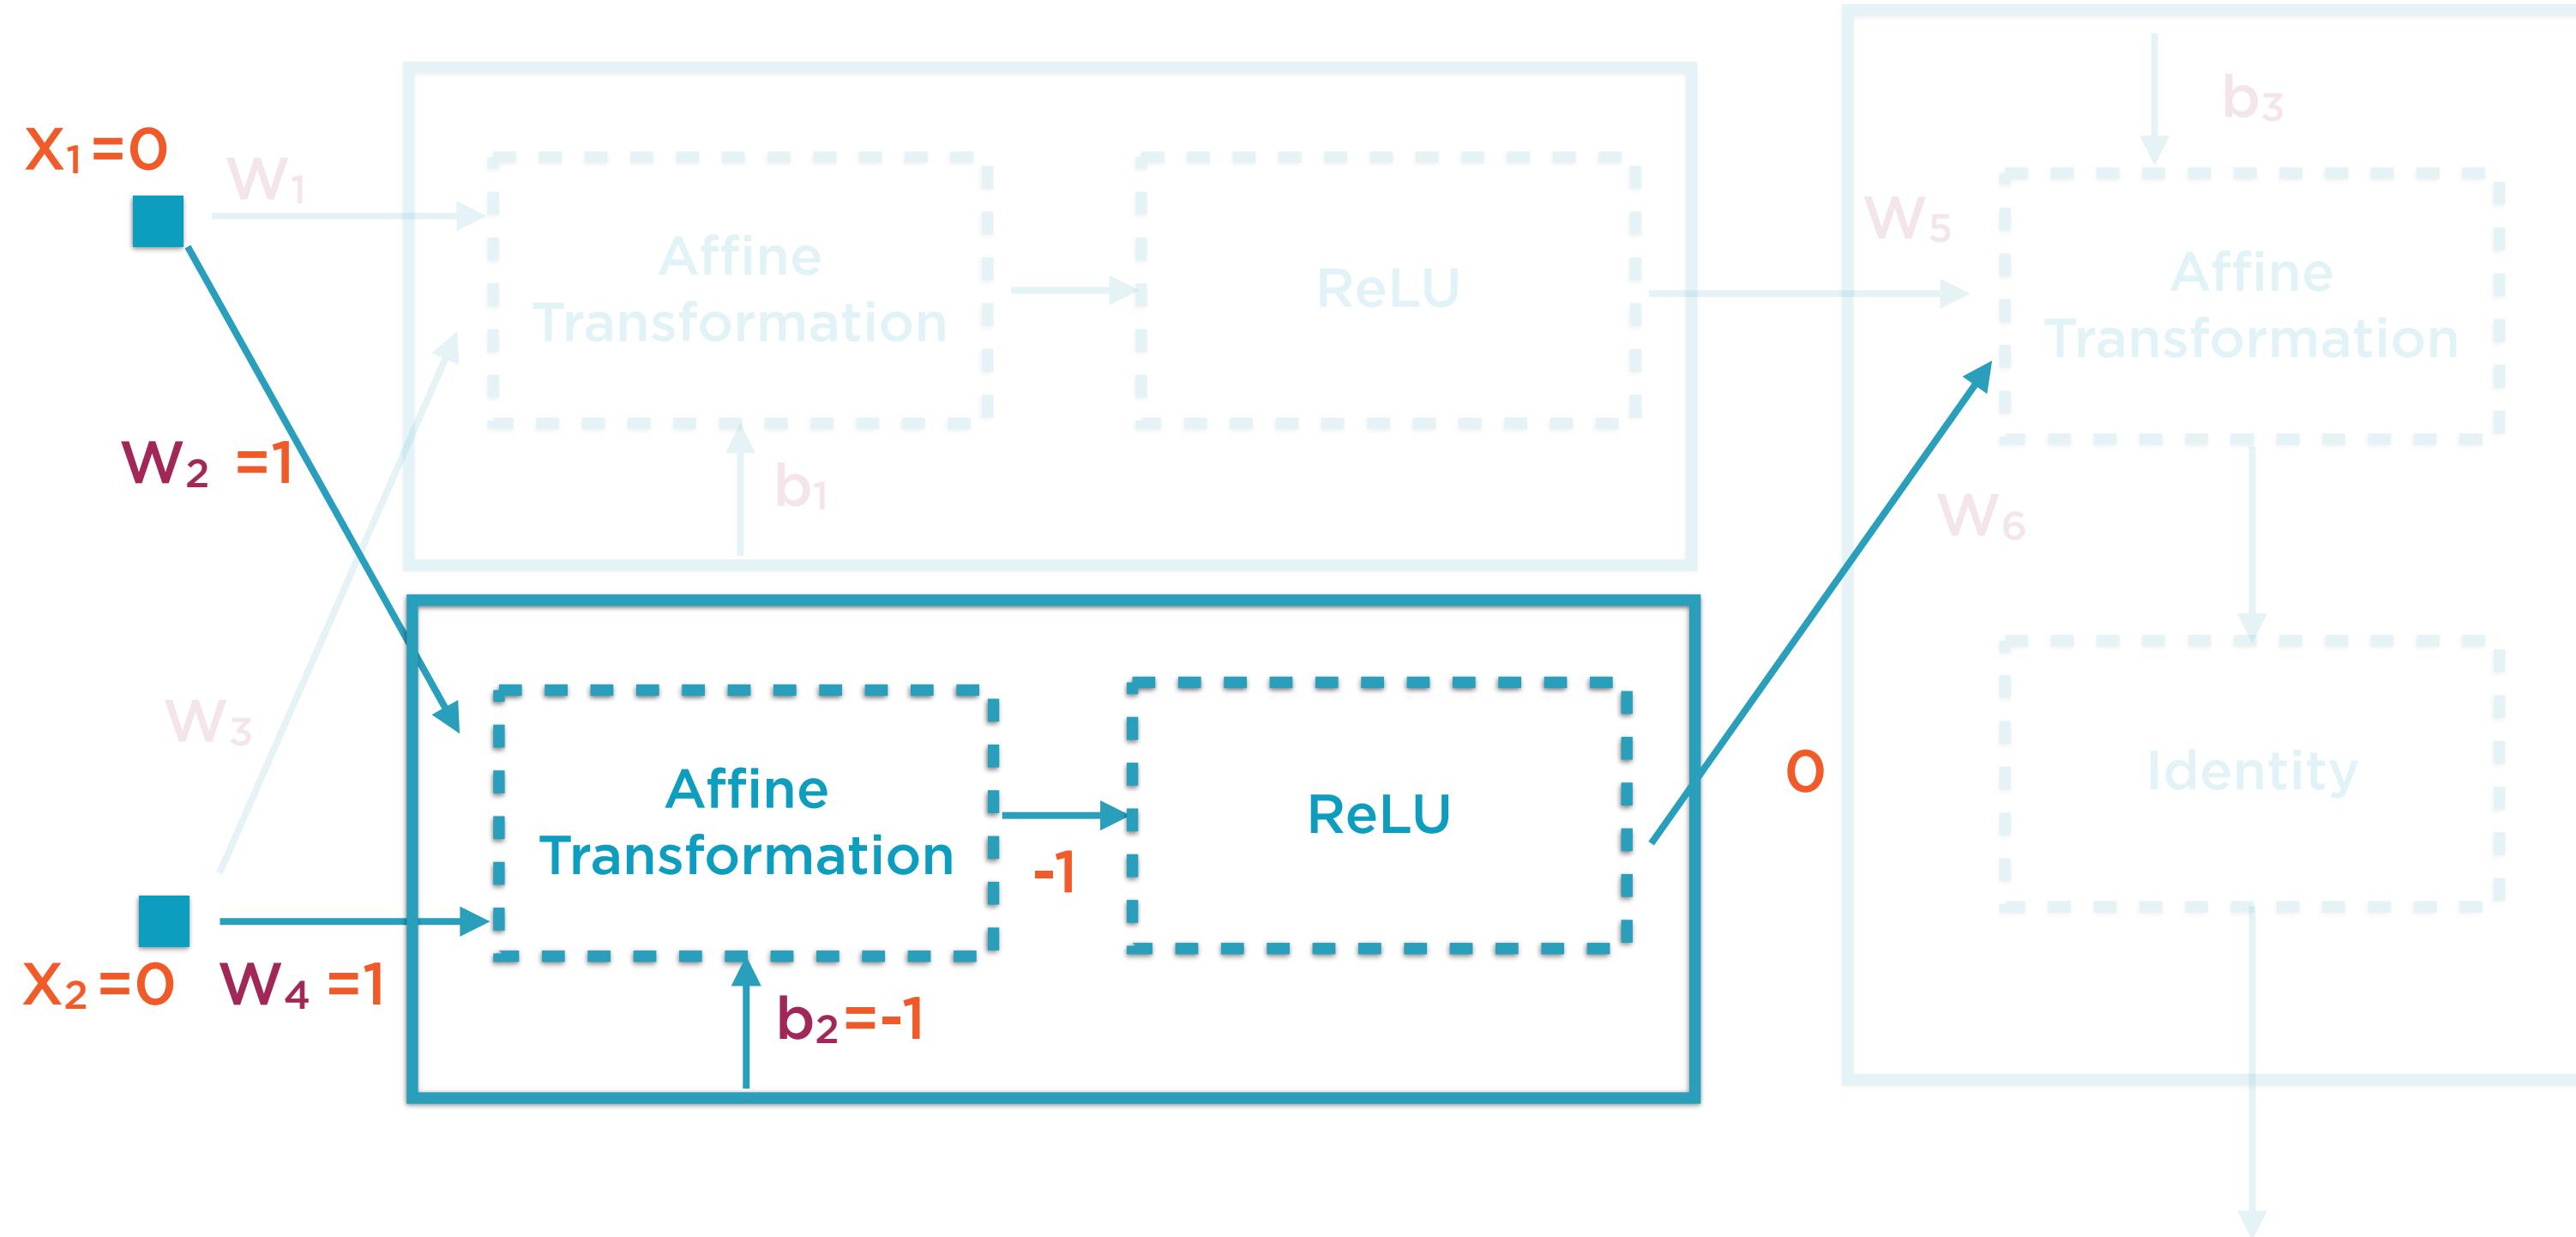
3-Neuron XOR



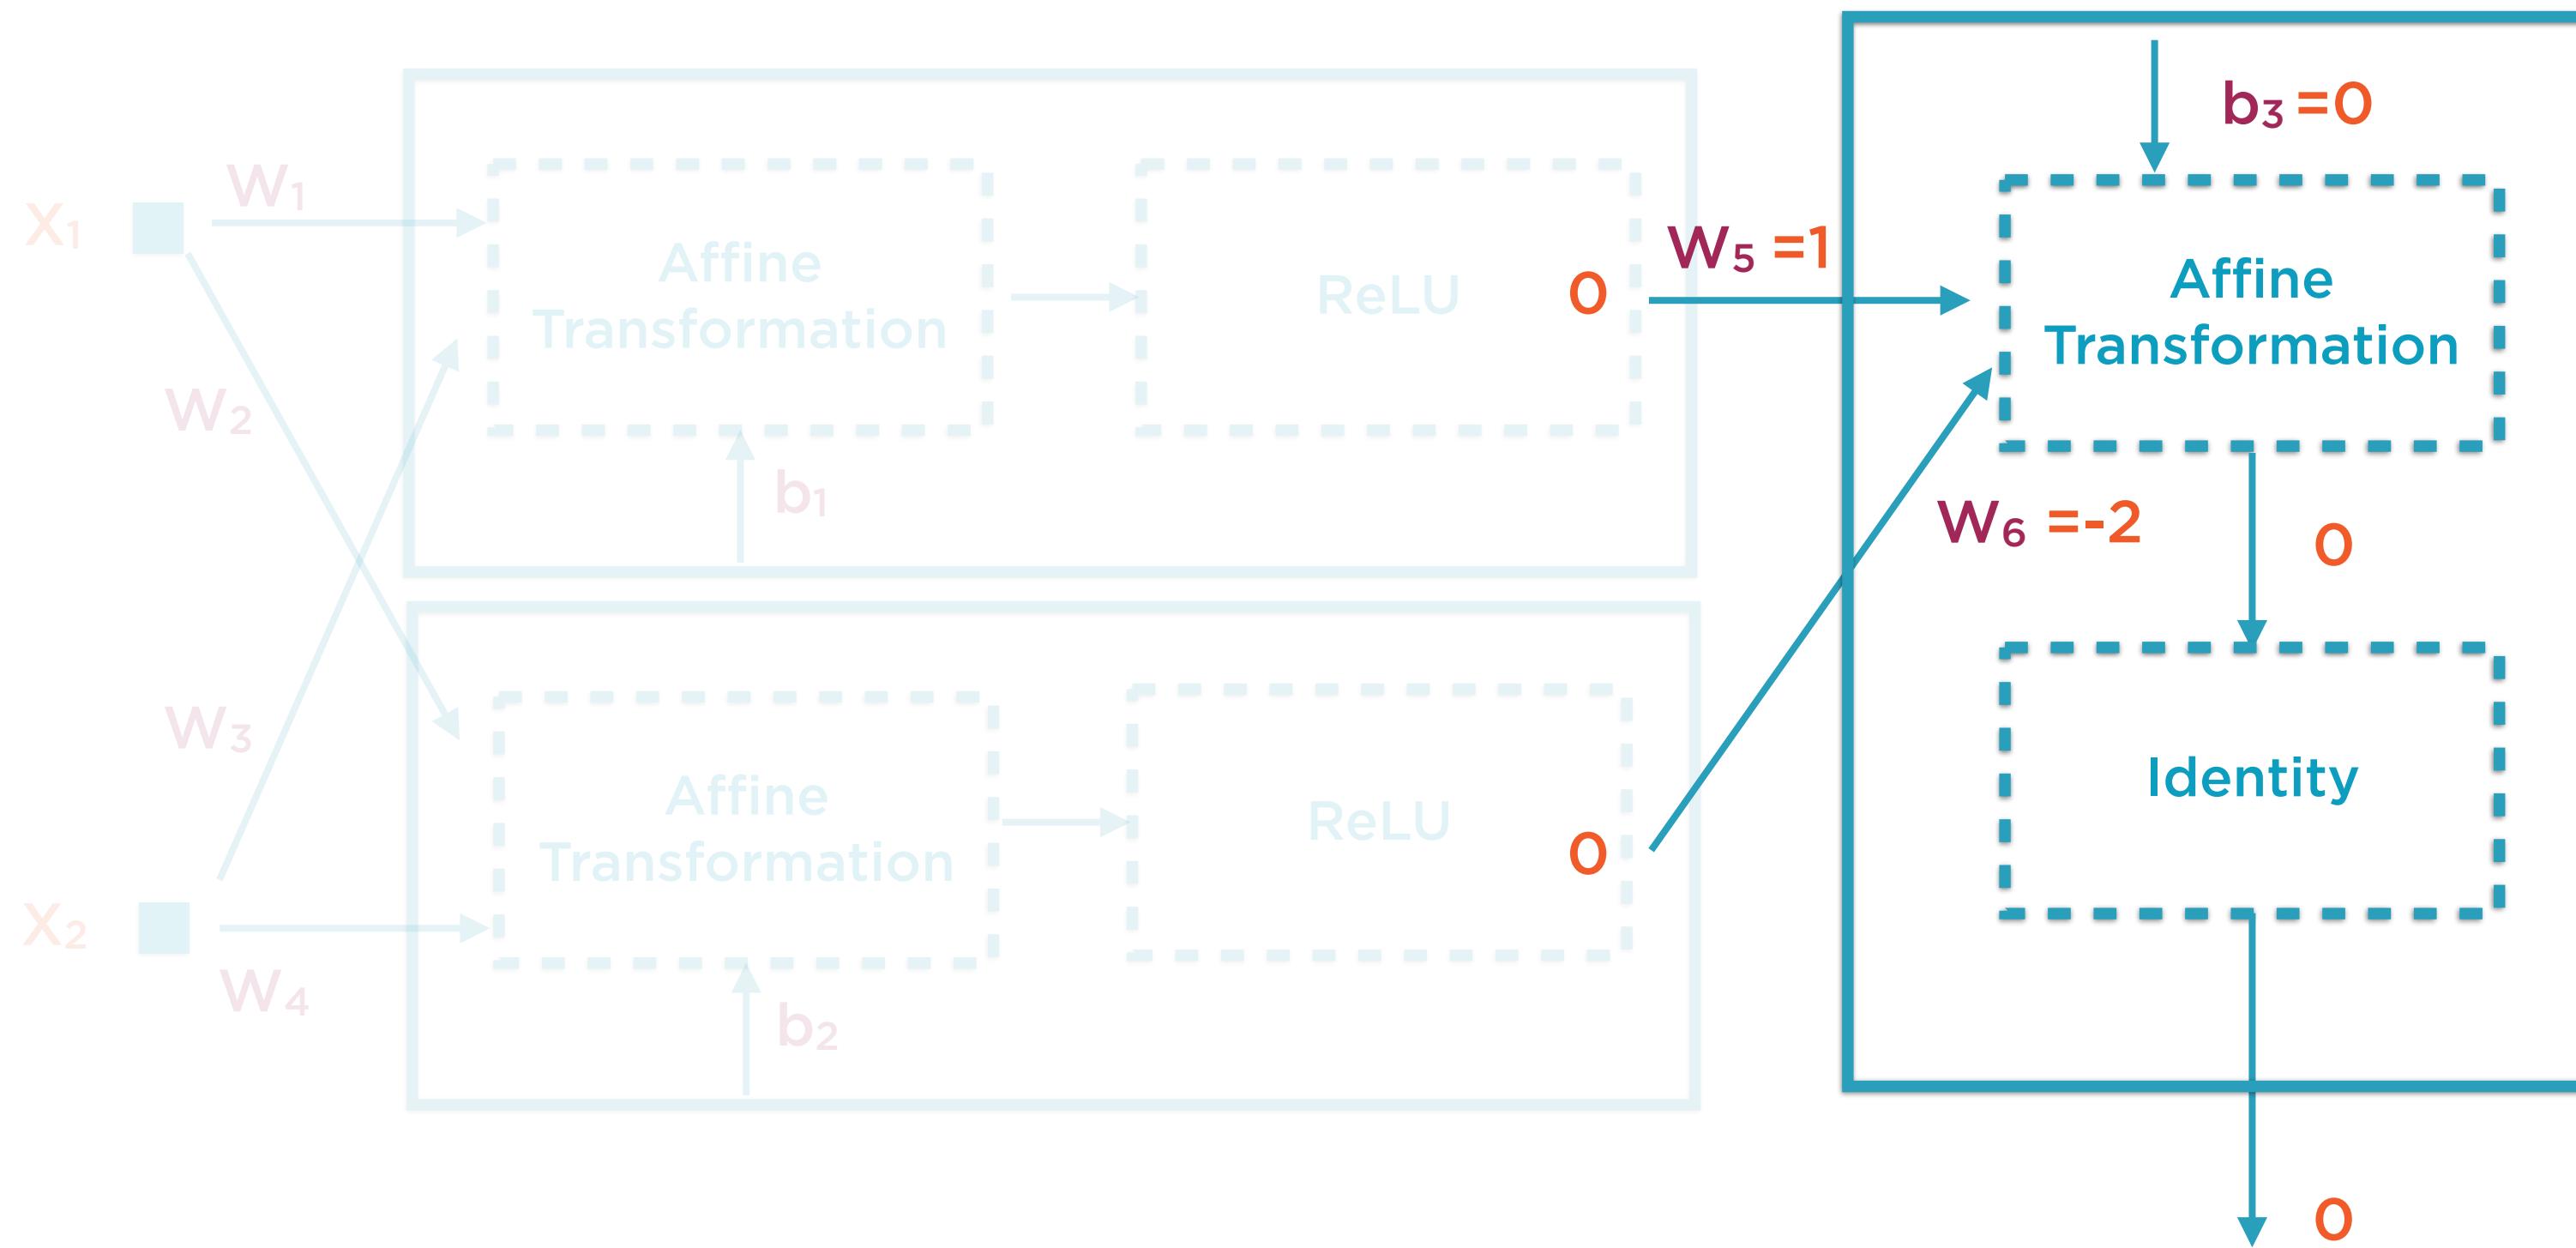
Weights and Bias of Neuron #1



Weights and Bias of Neuron #2



Weights and Bias of Neuron #3

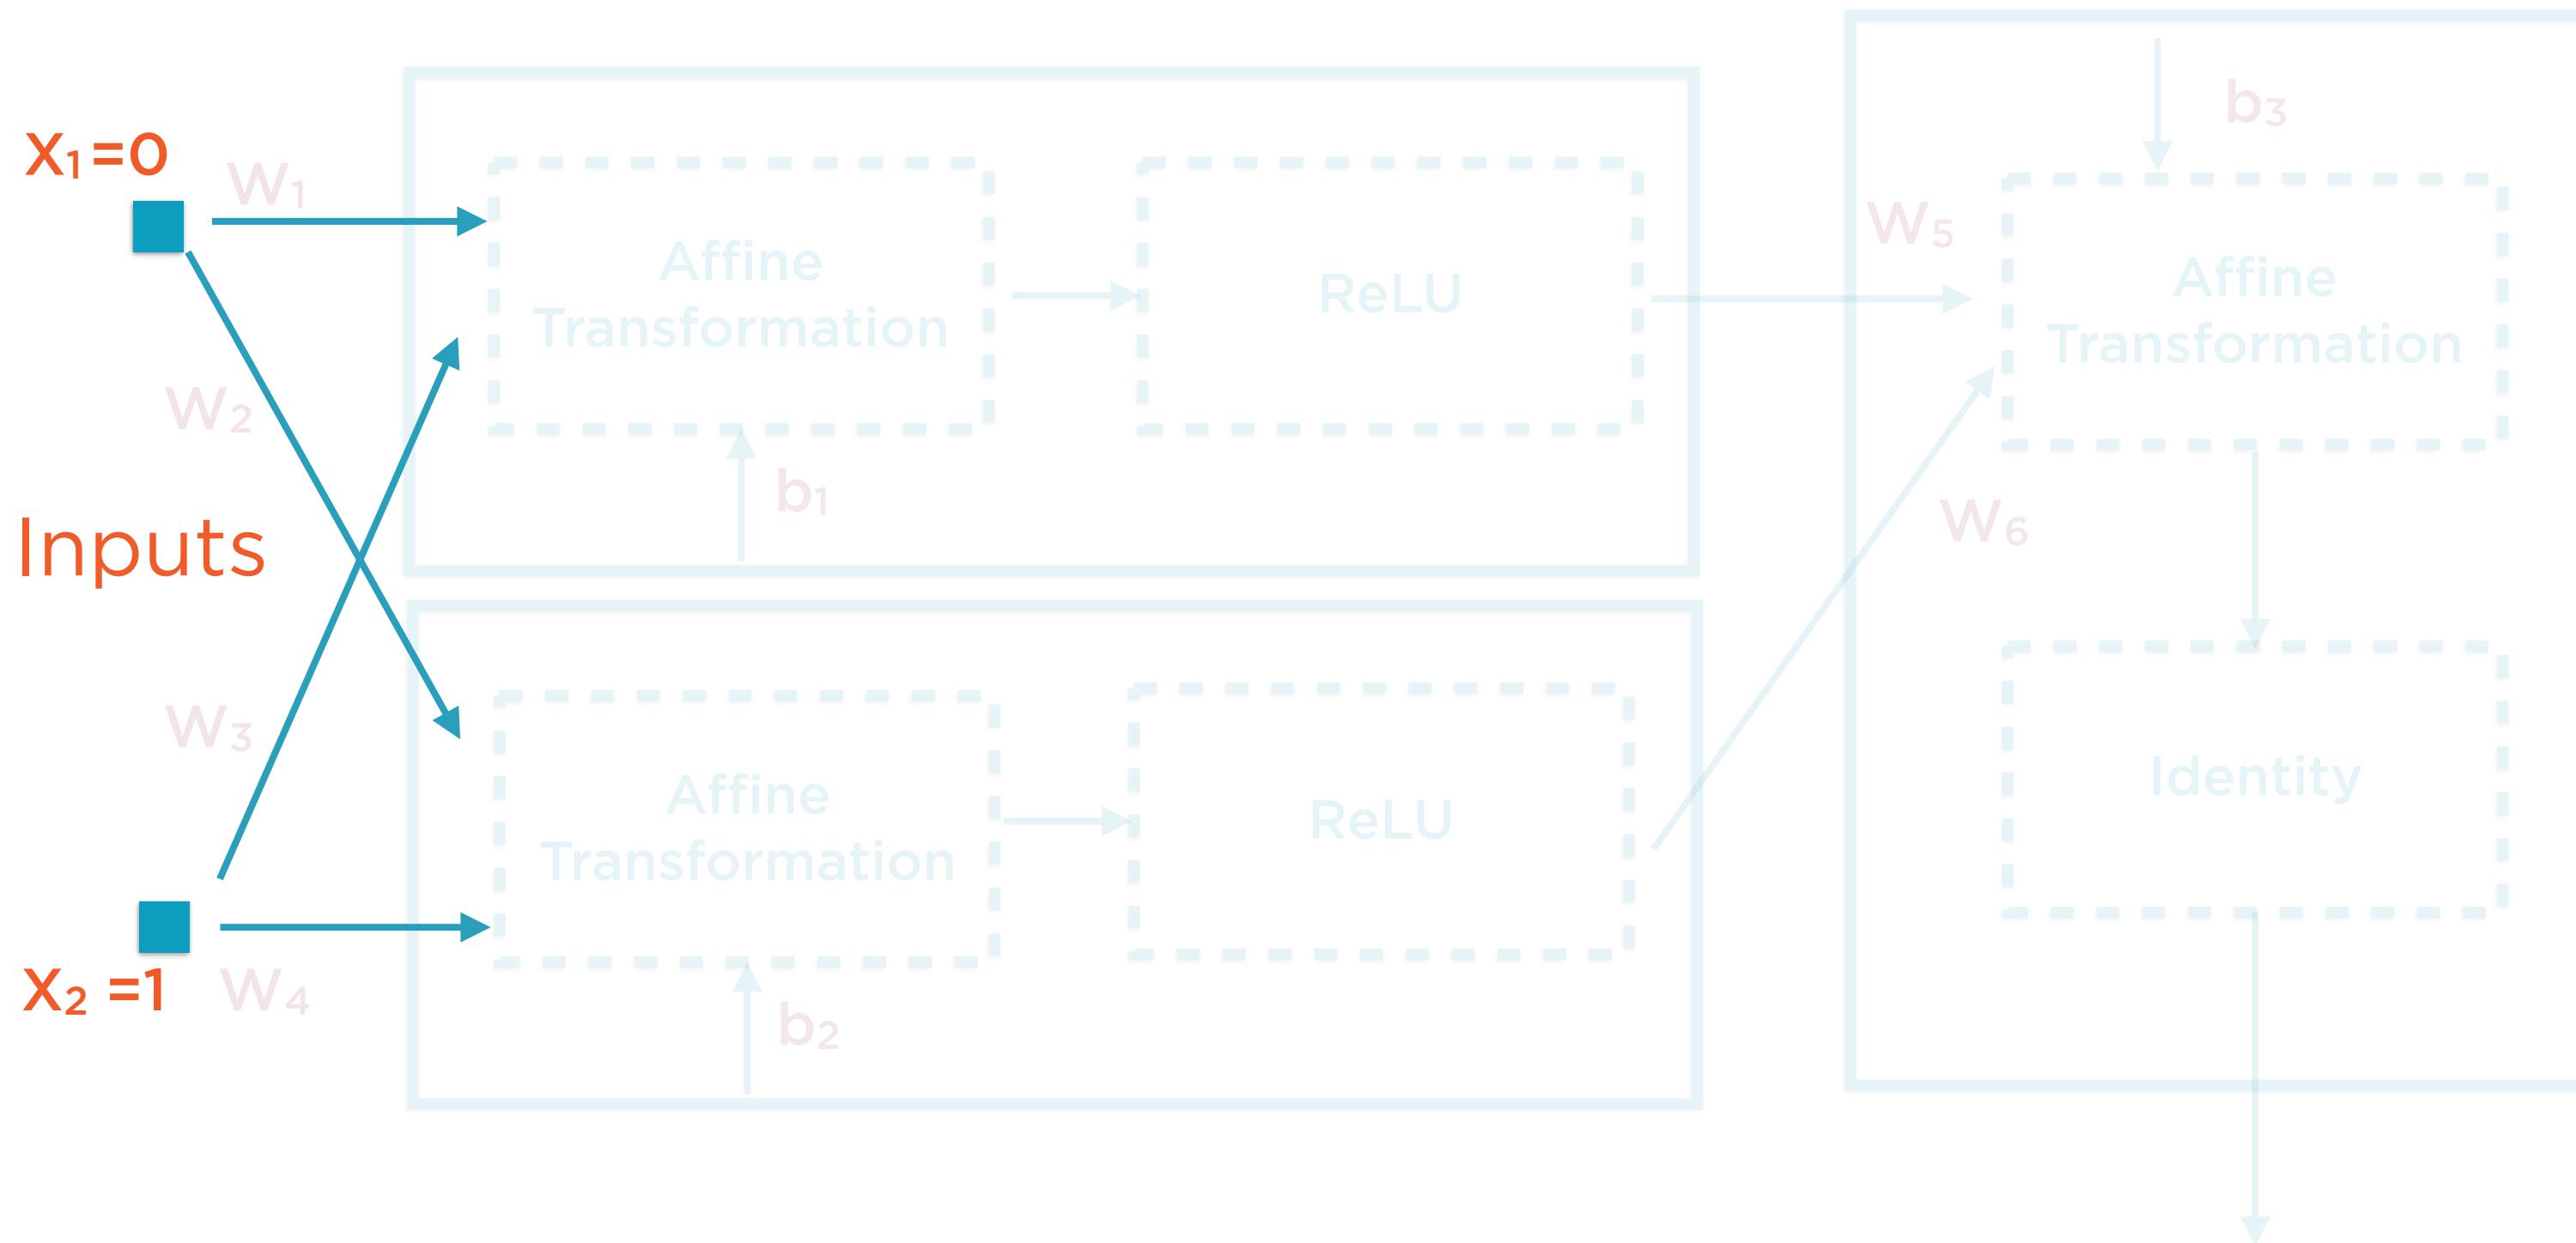


X ₁	X ₂	Y
0	0	0
0	1	1
1	0	1
1	1	0

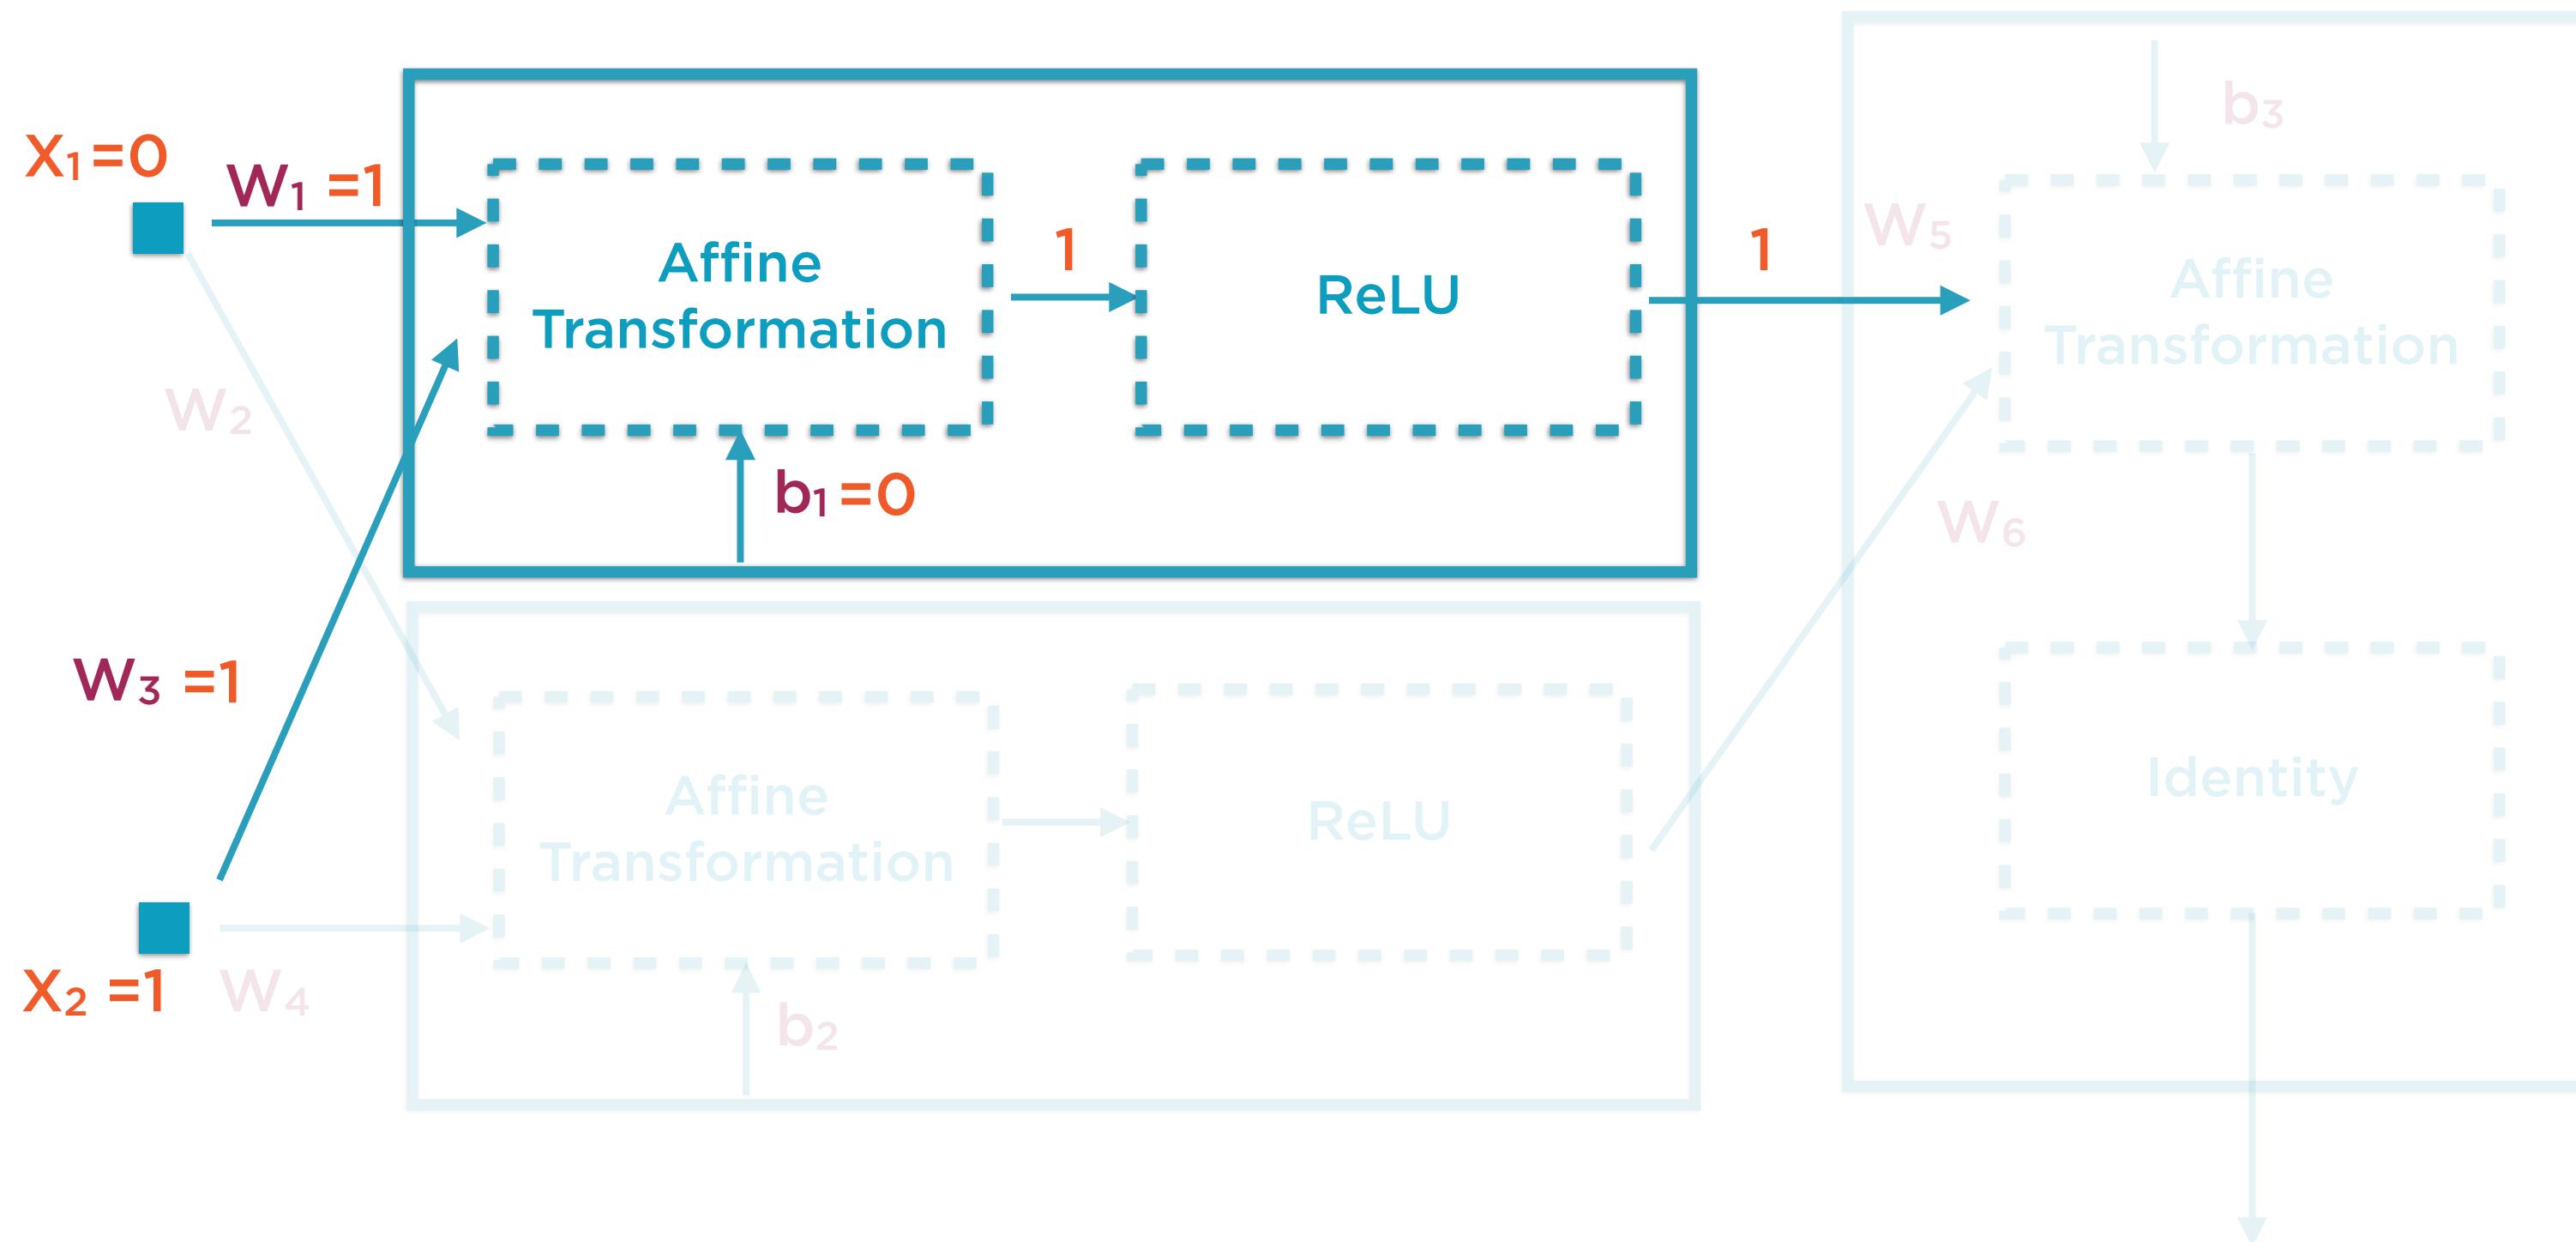
“Learning” XOR

Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

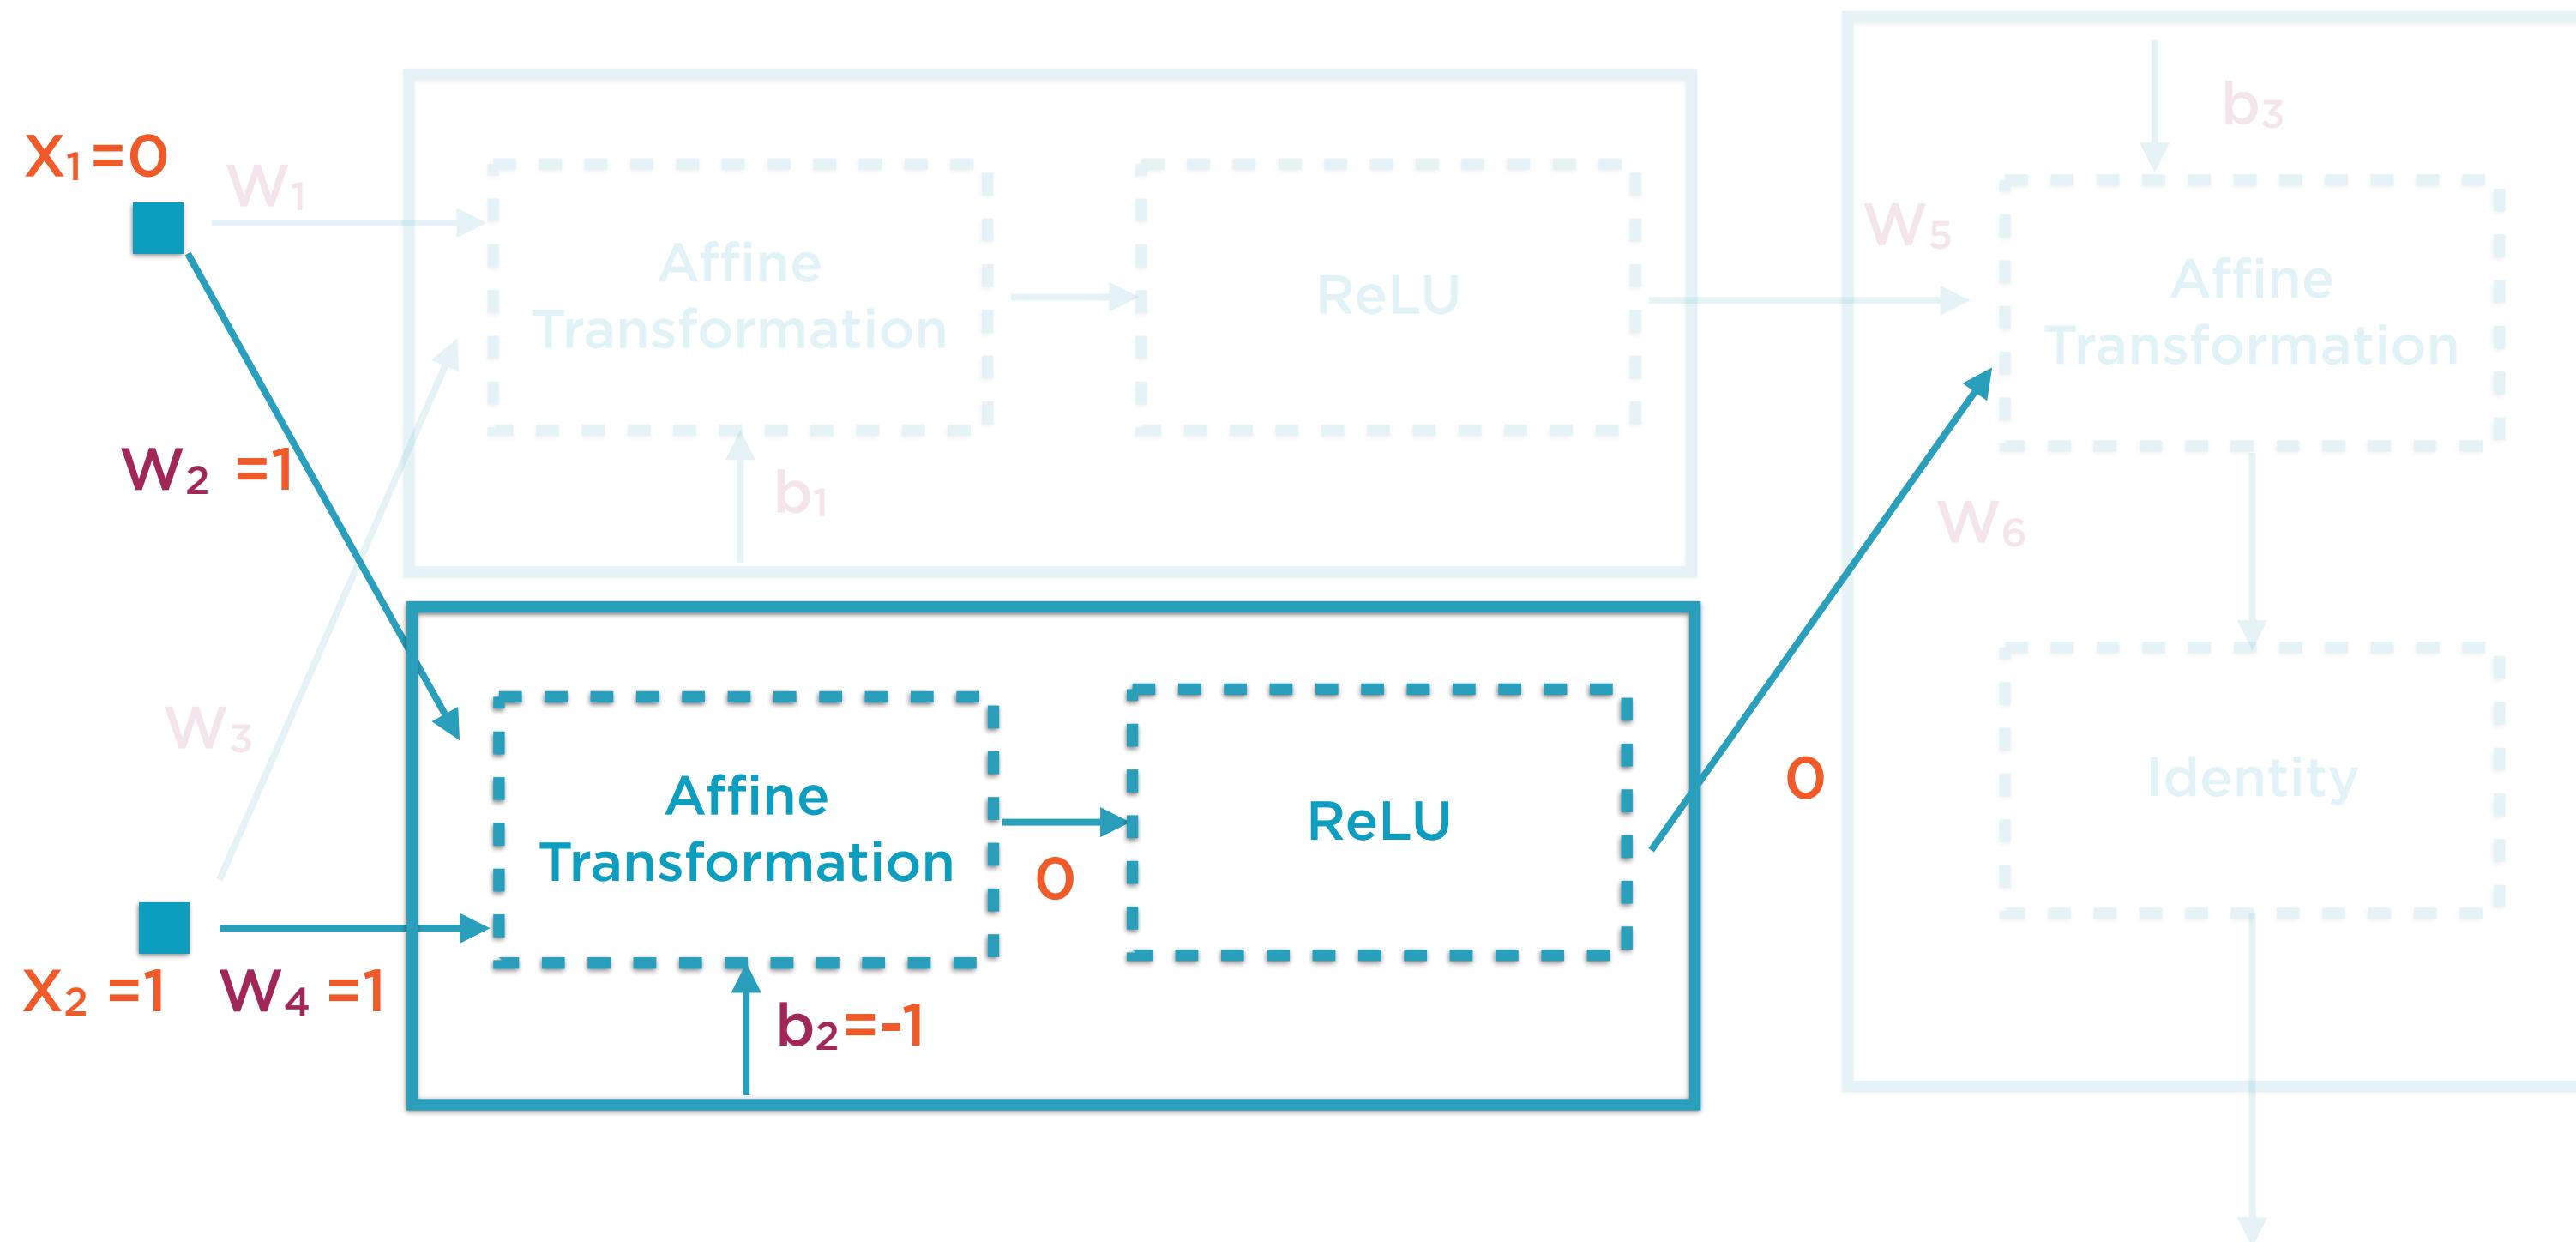
3-Neuron XOR



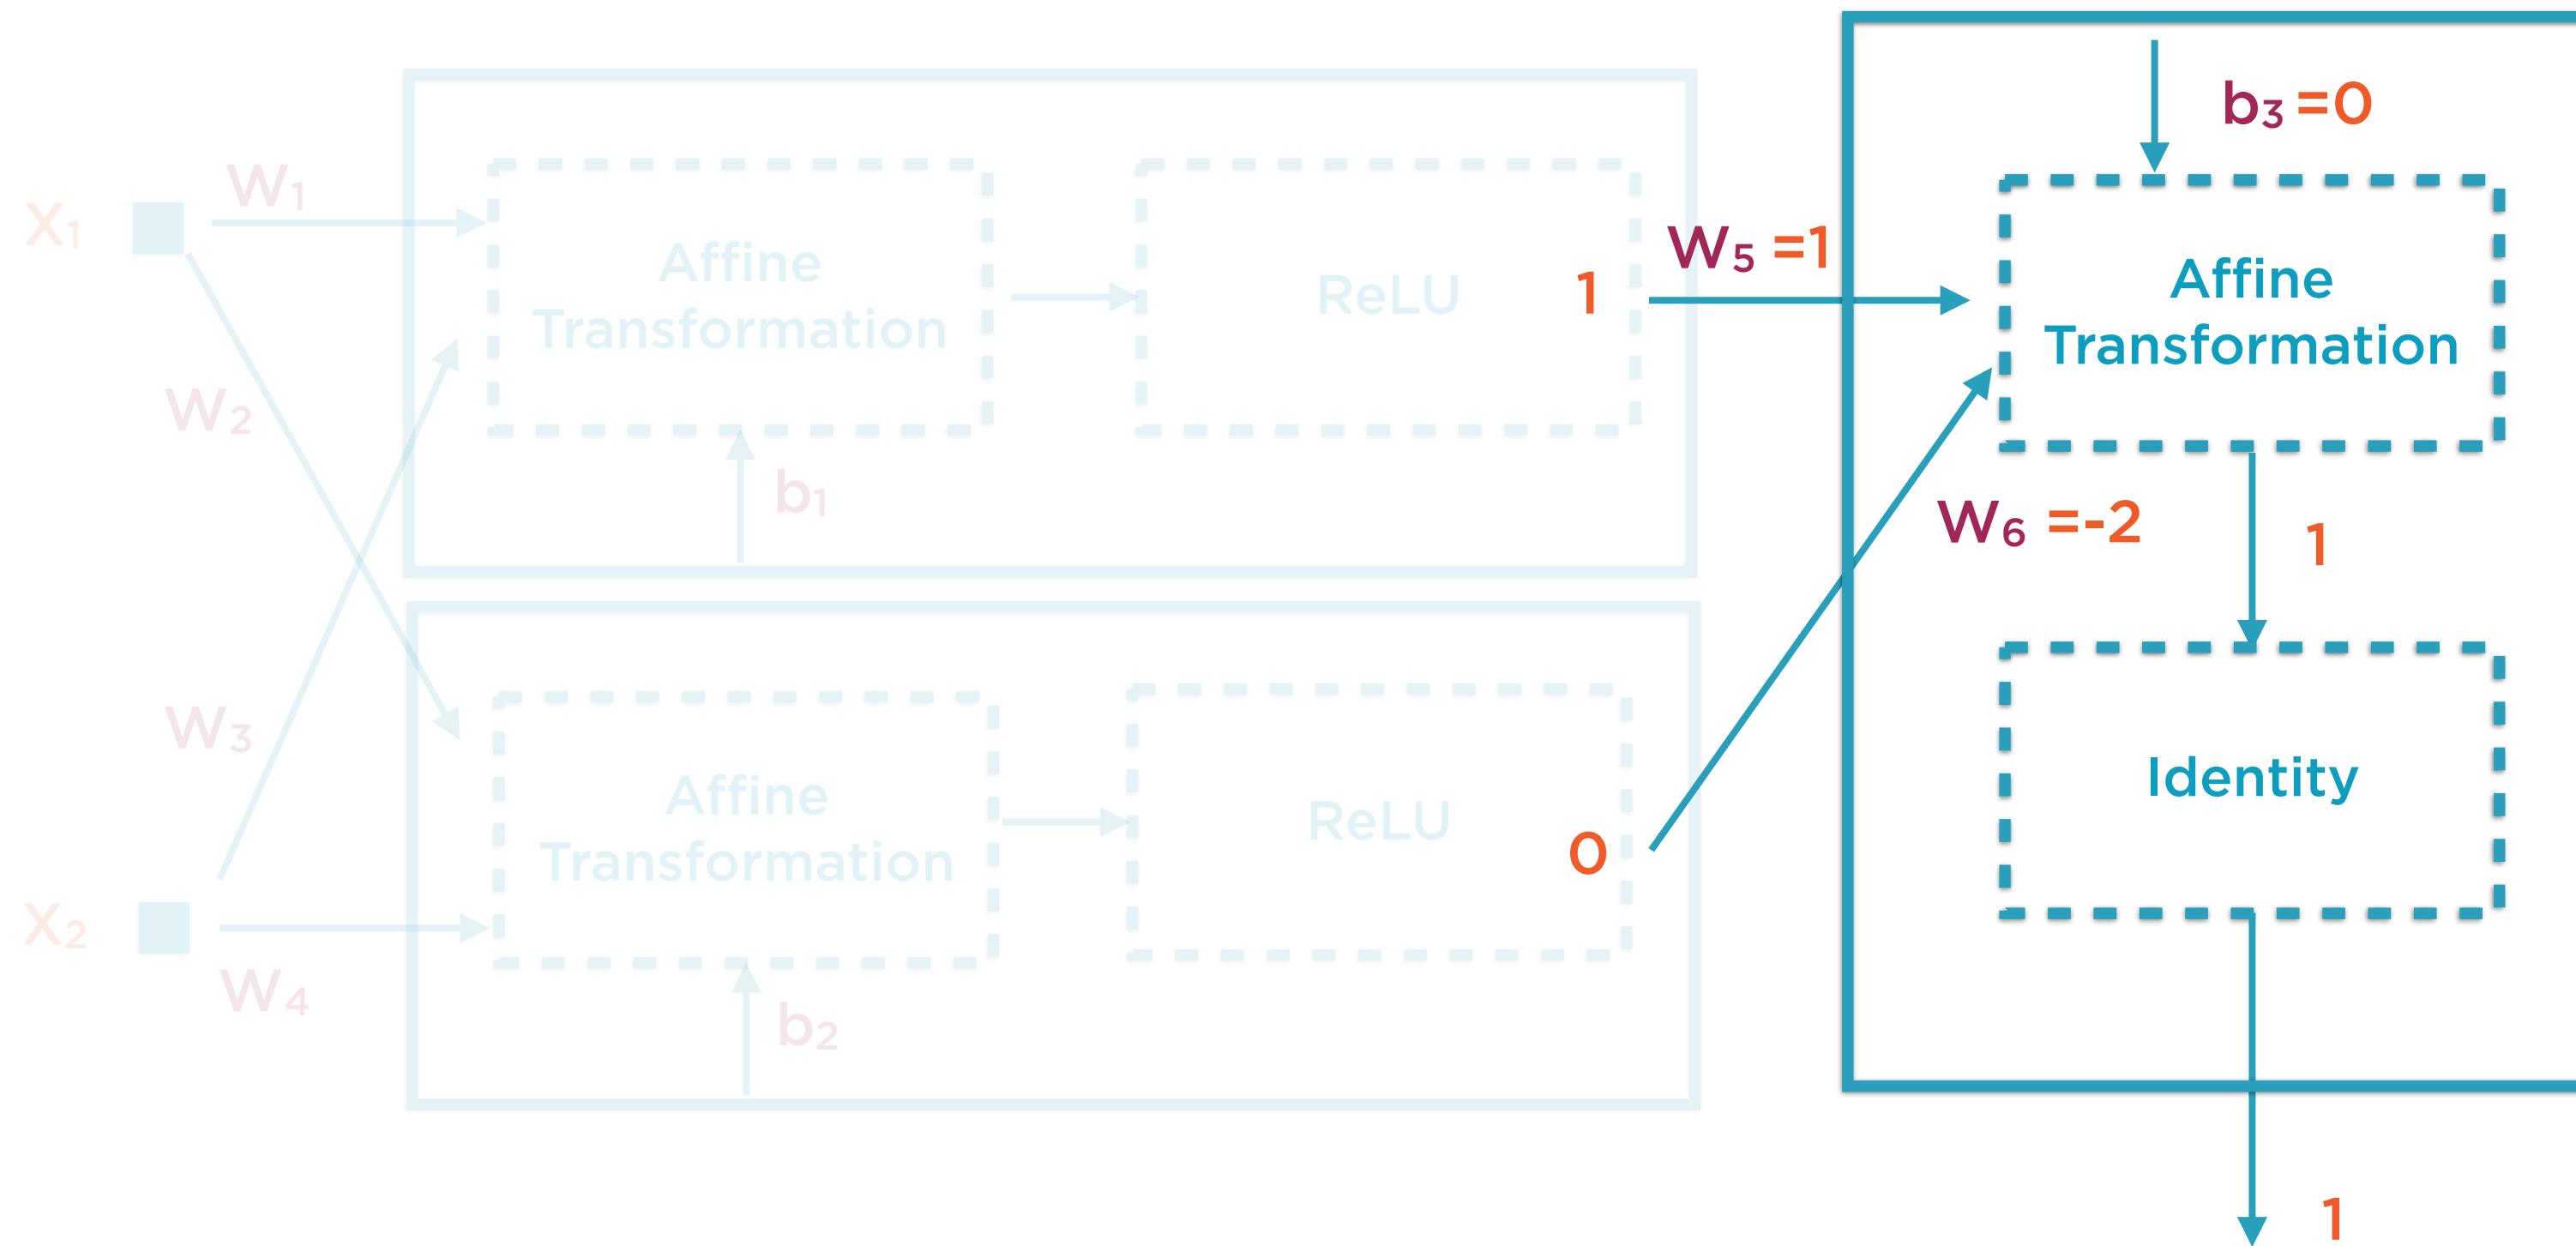
Weights and Bias of Neuron #1



Weights and Bias of Neuron #2



Weights and Bias of Neuron #3

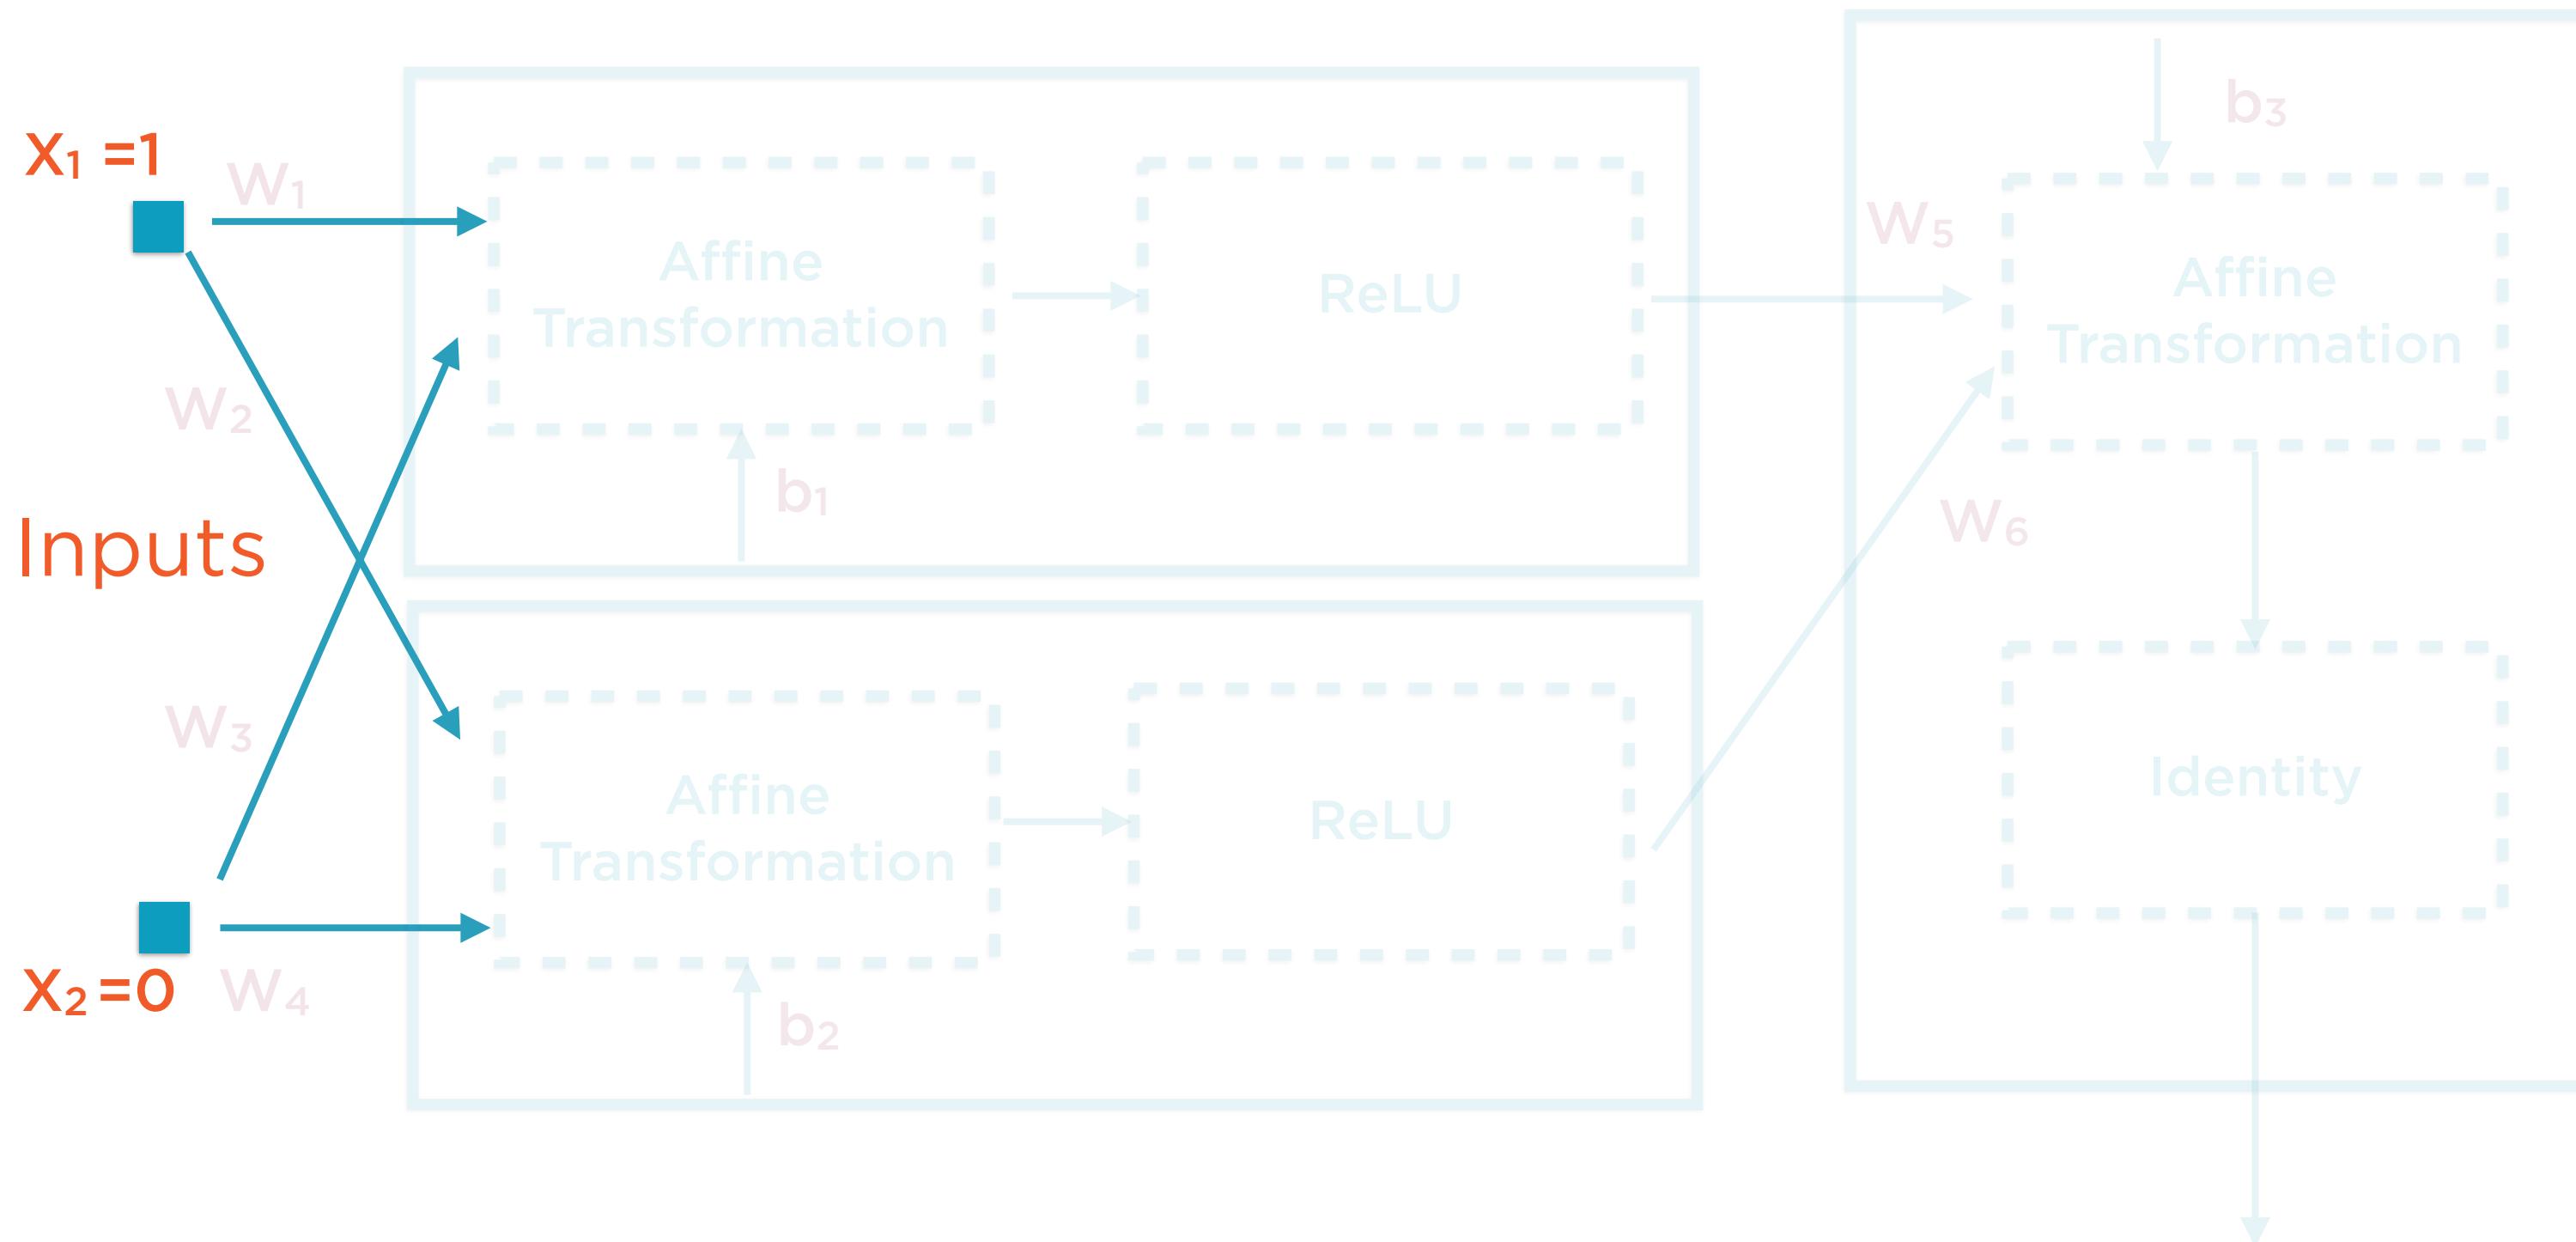


X ₁	X ₂	Y
0	0	0
0	1	1
1	0	1
1	1	0

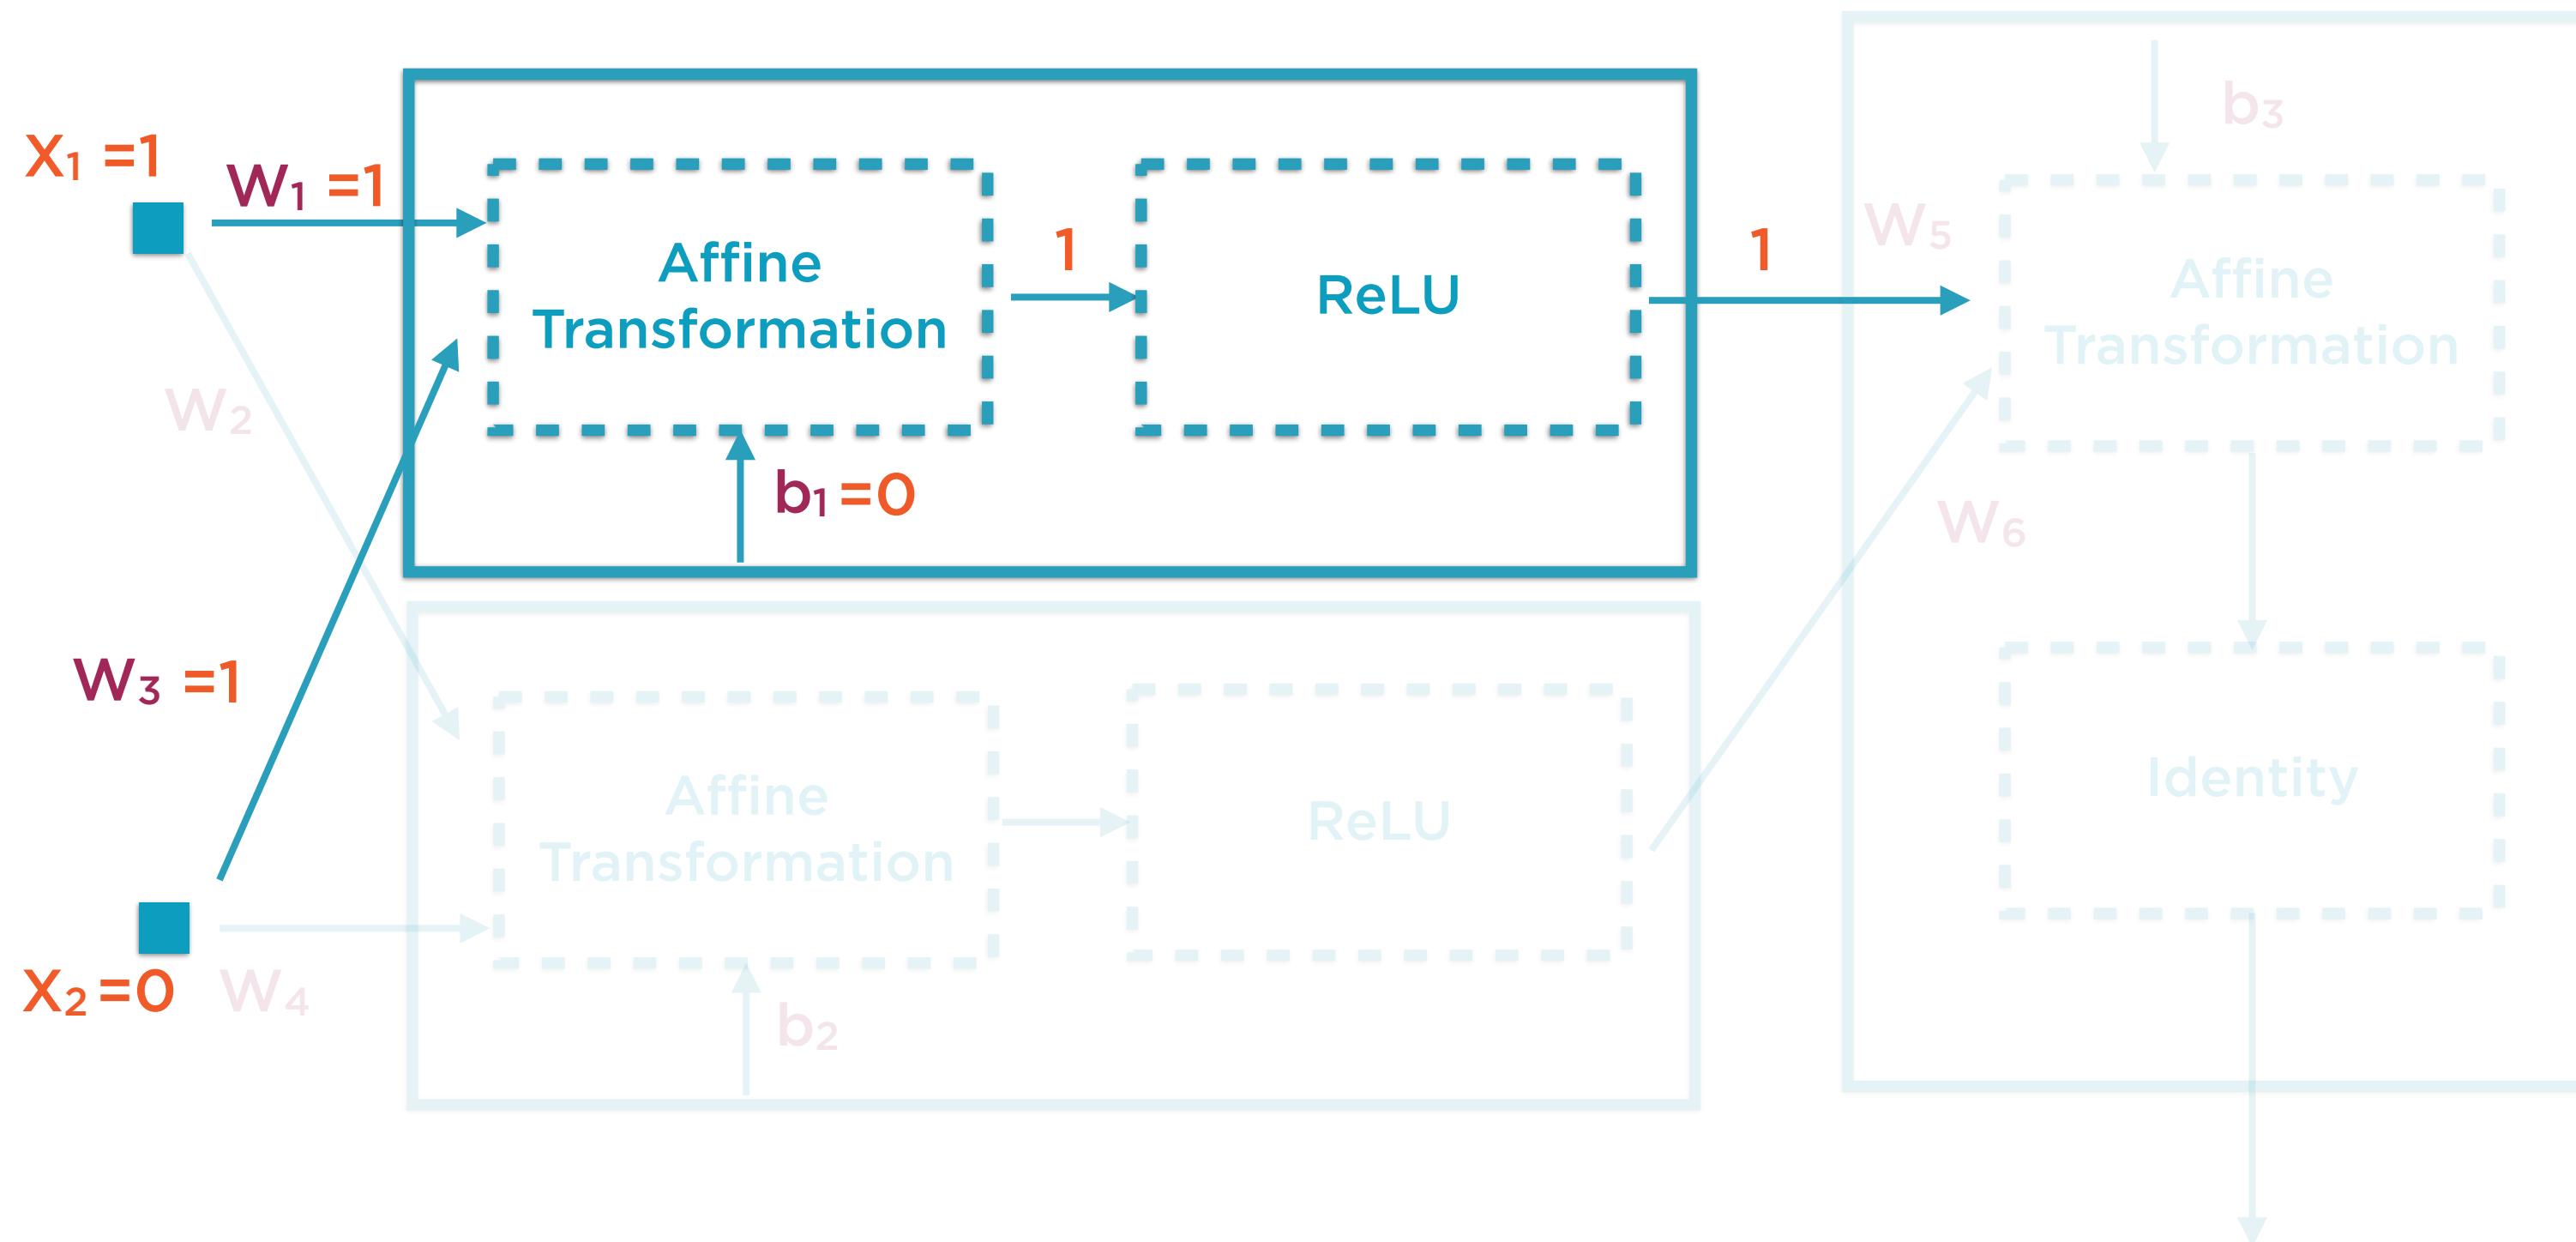
“Learning” XOR

Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

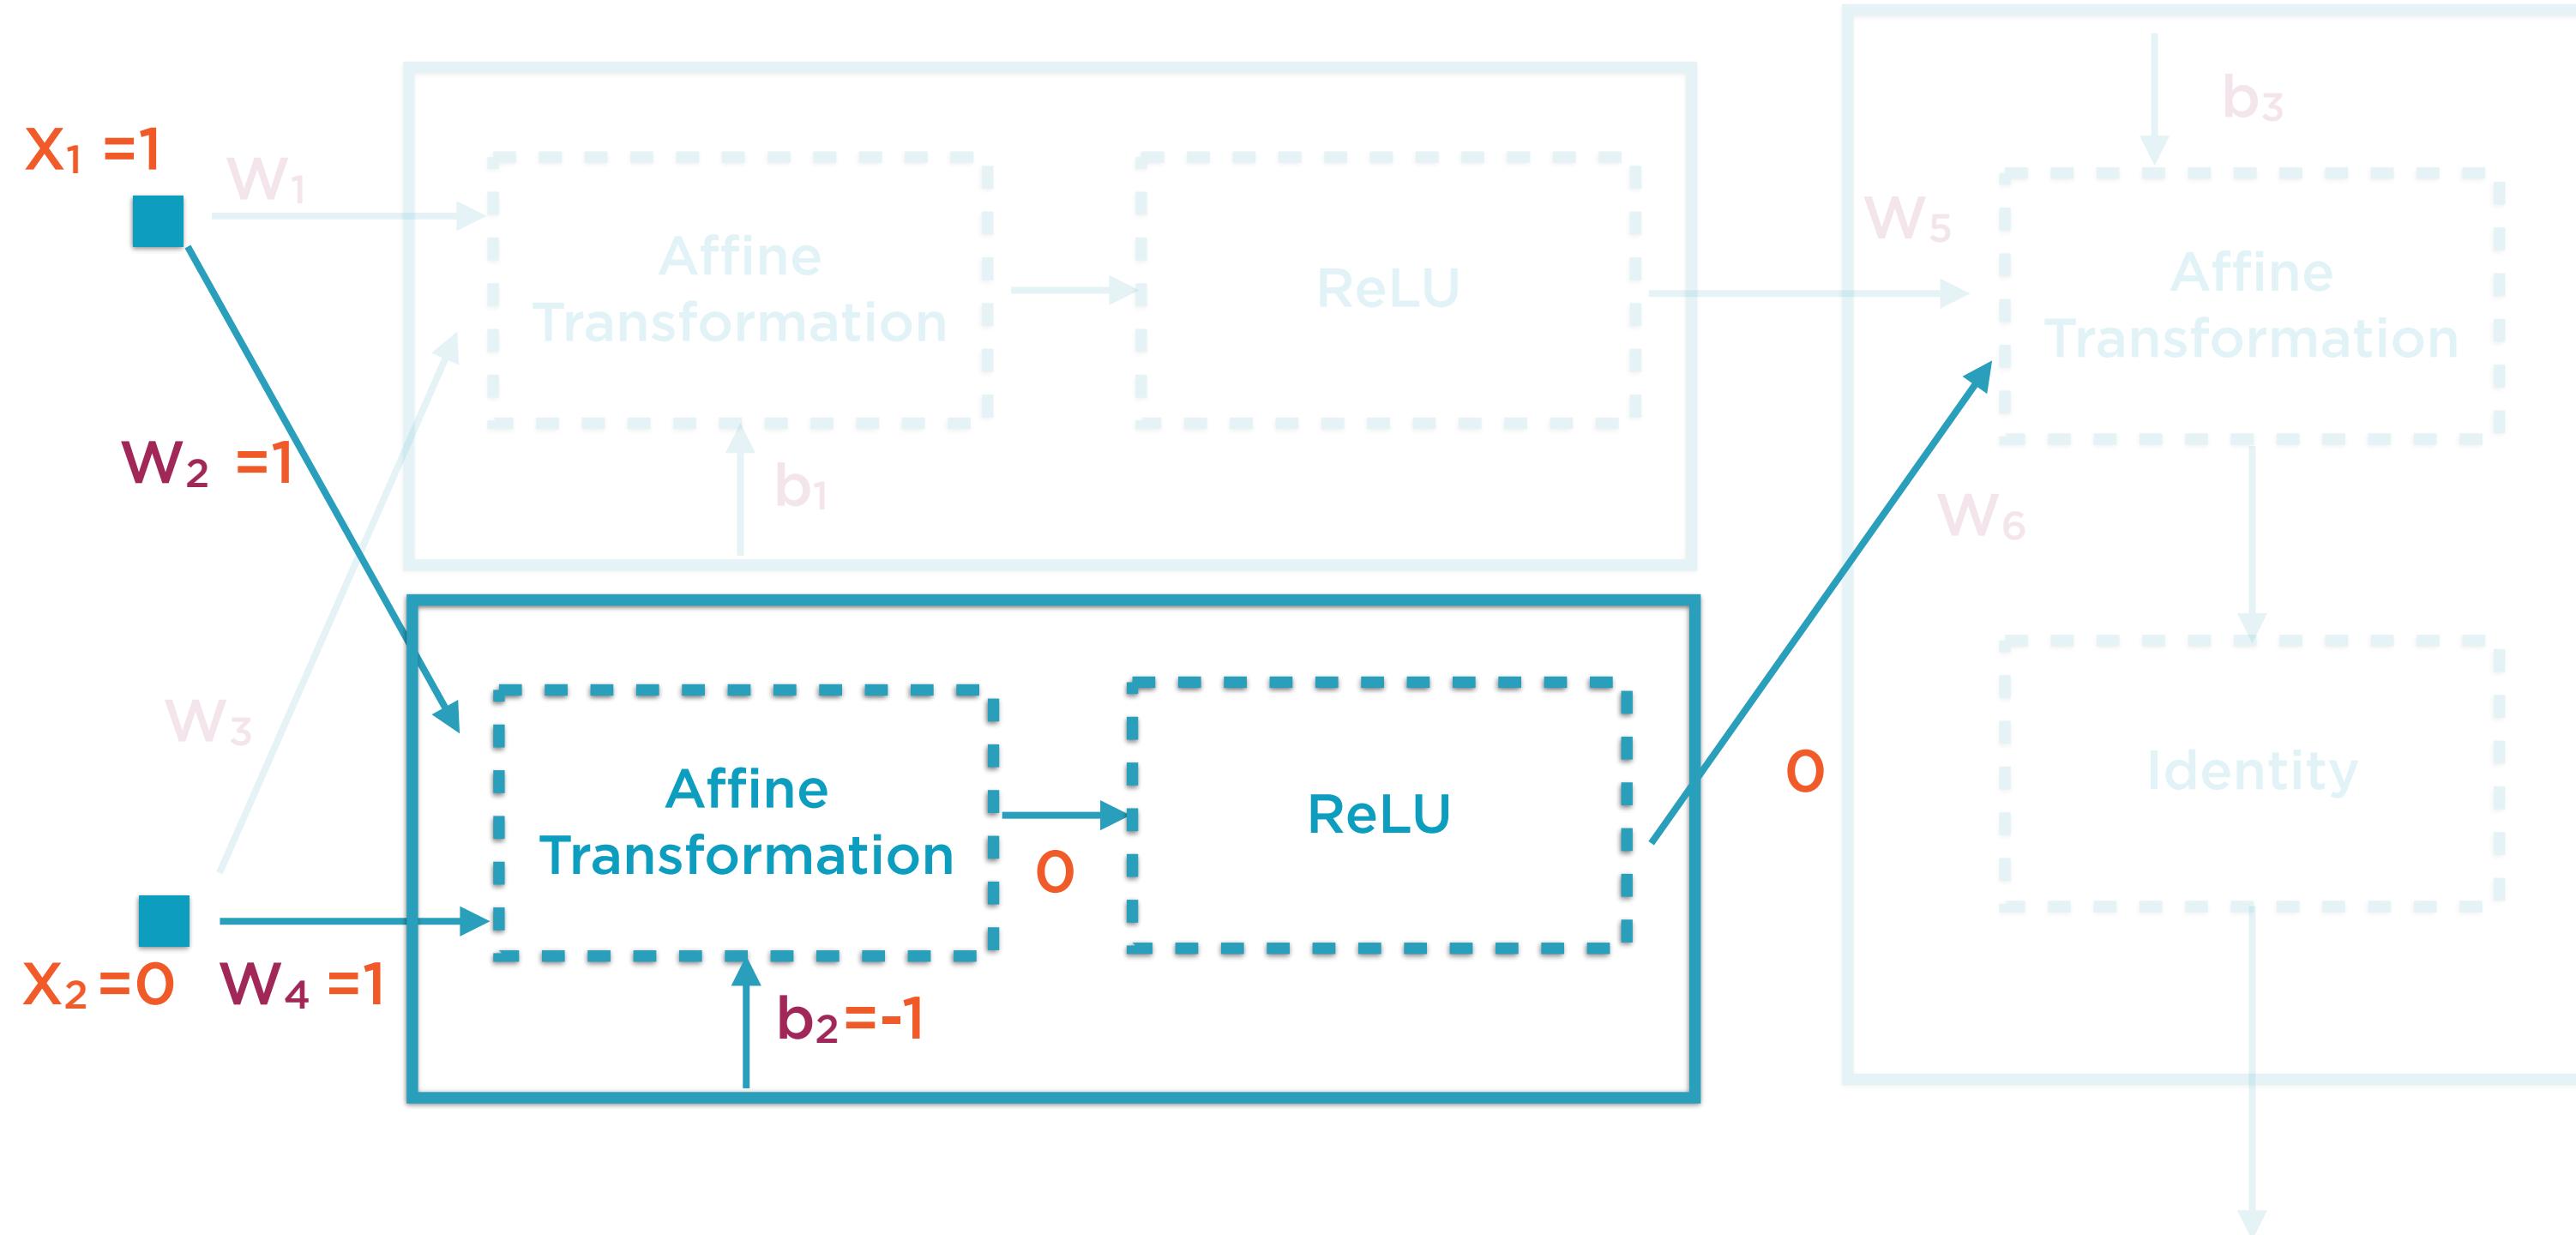
3-Neuron XOR



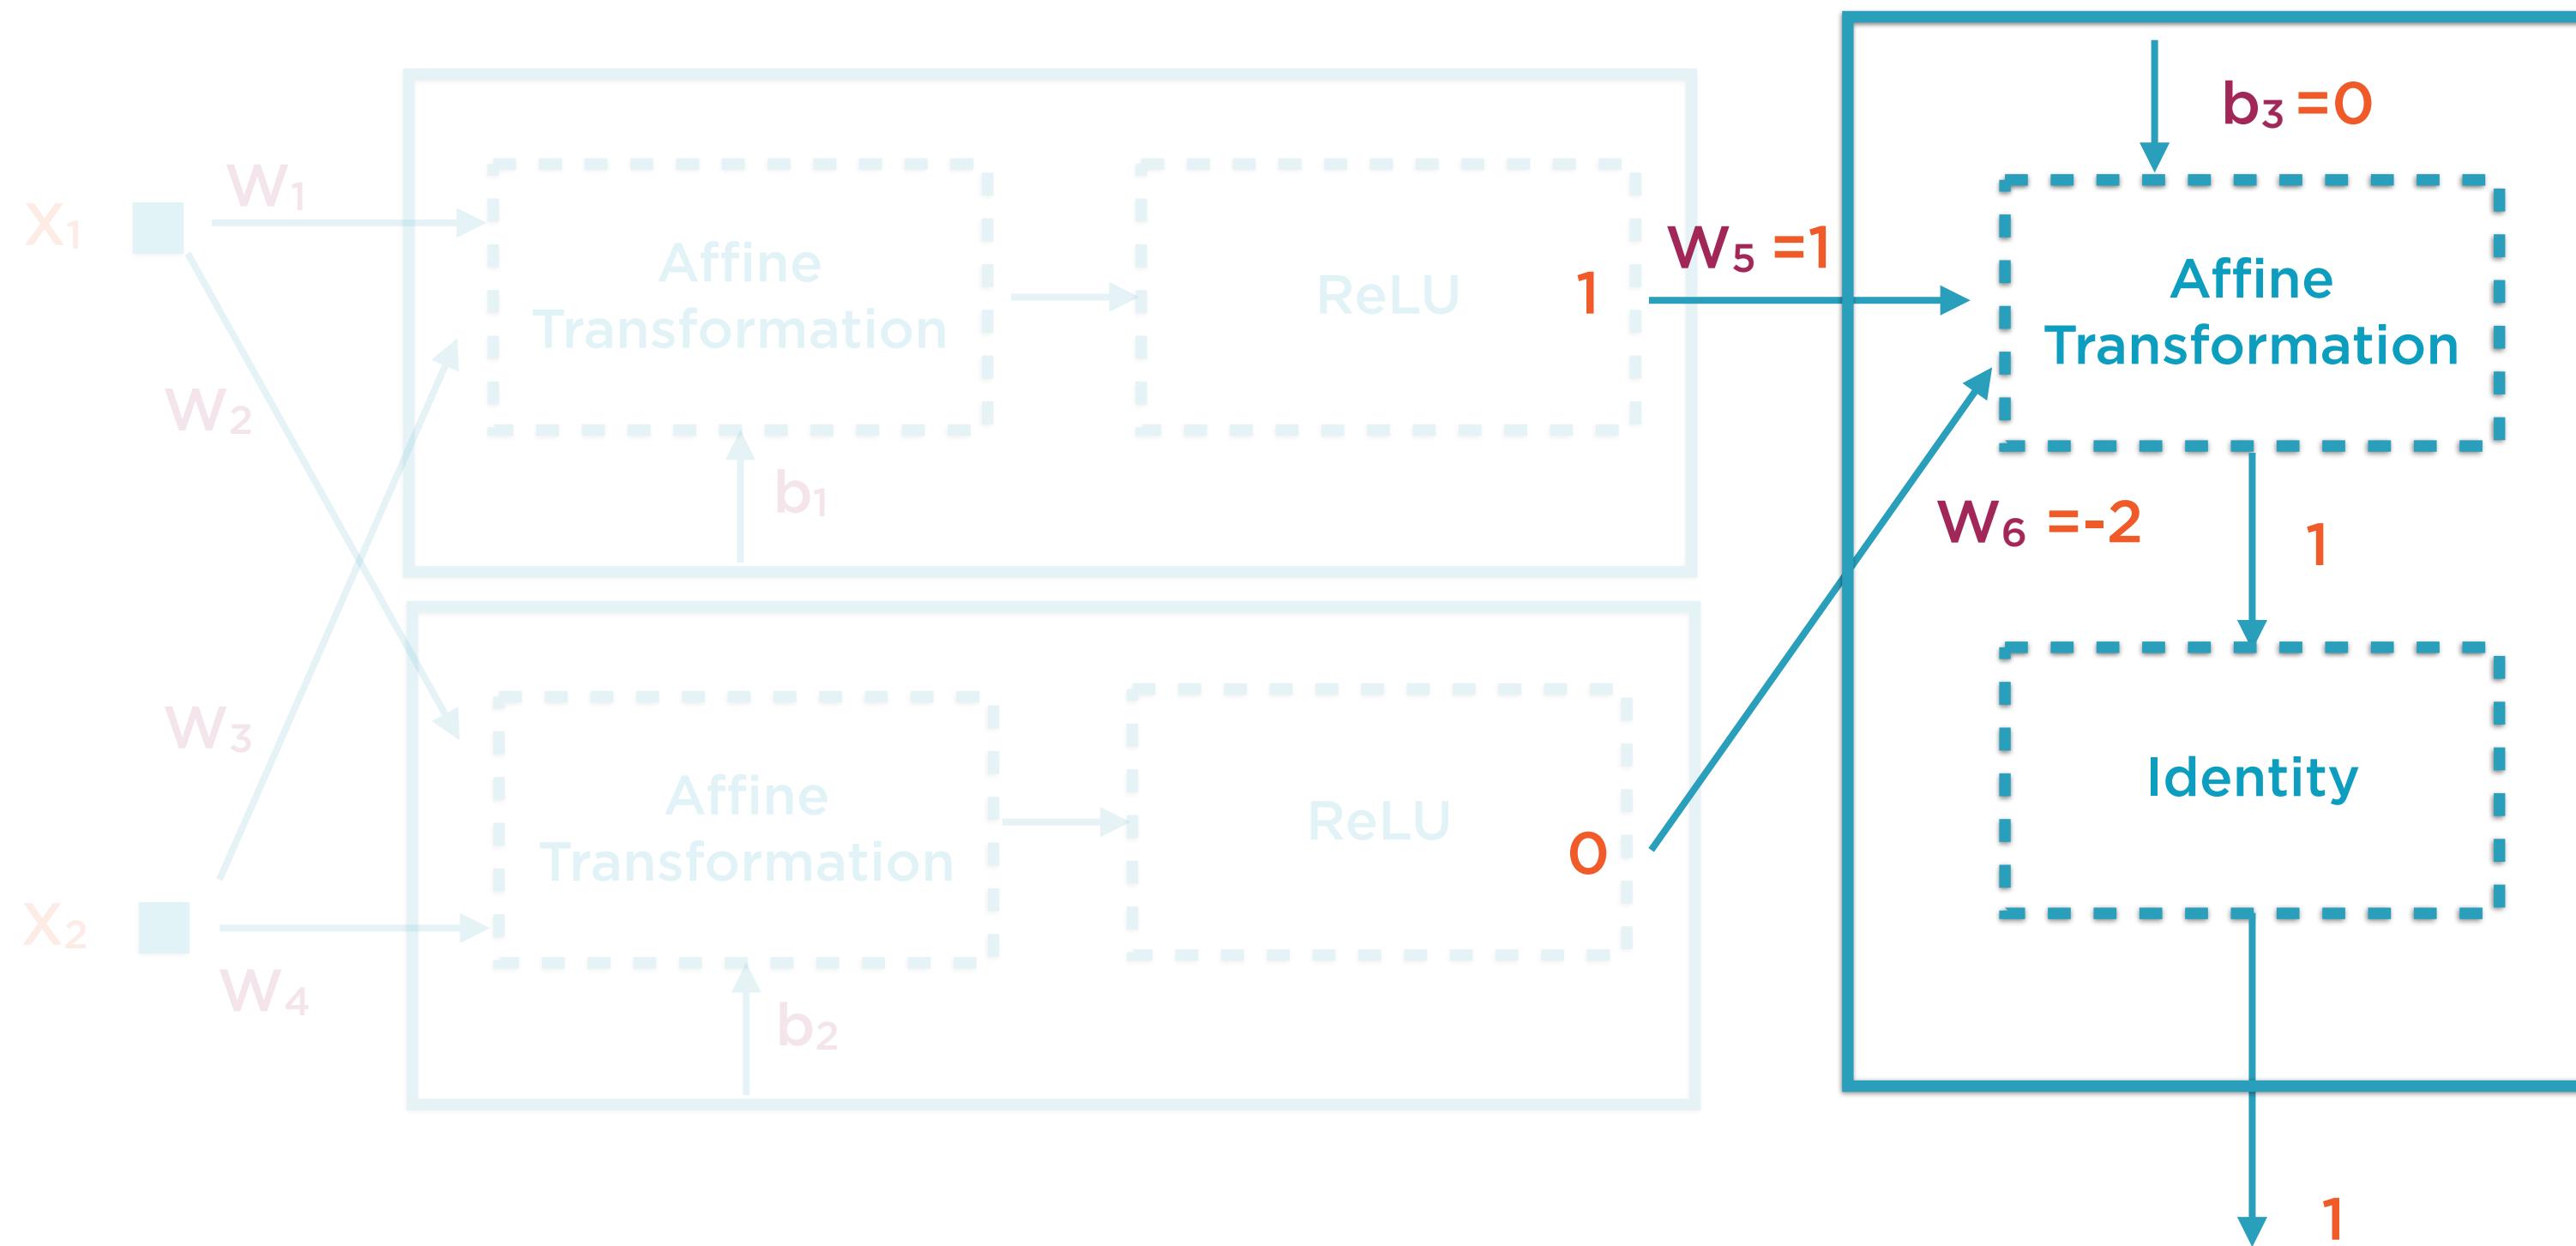
Weights and Bias of Neuron #1



Weights and Bias of Neuron #2



Weights and Bias of Neuron #3

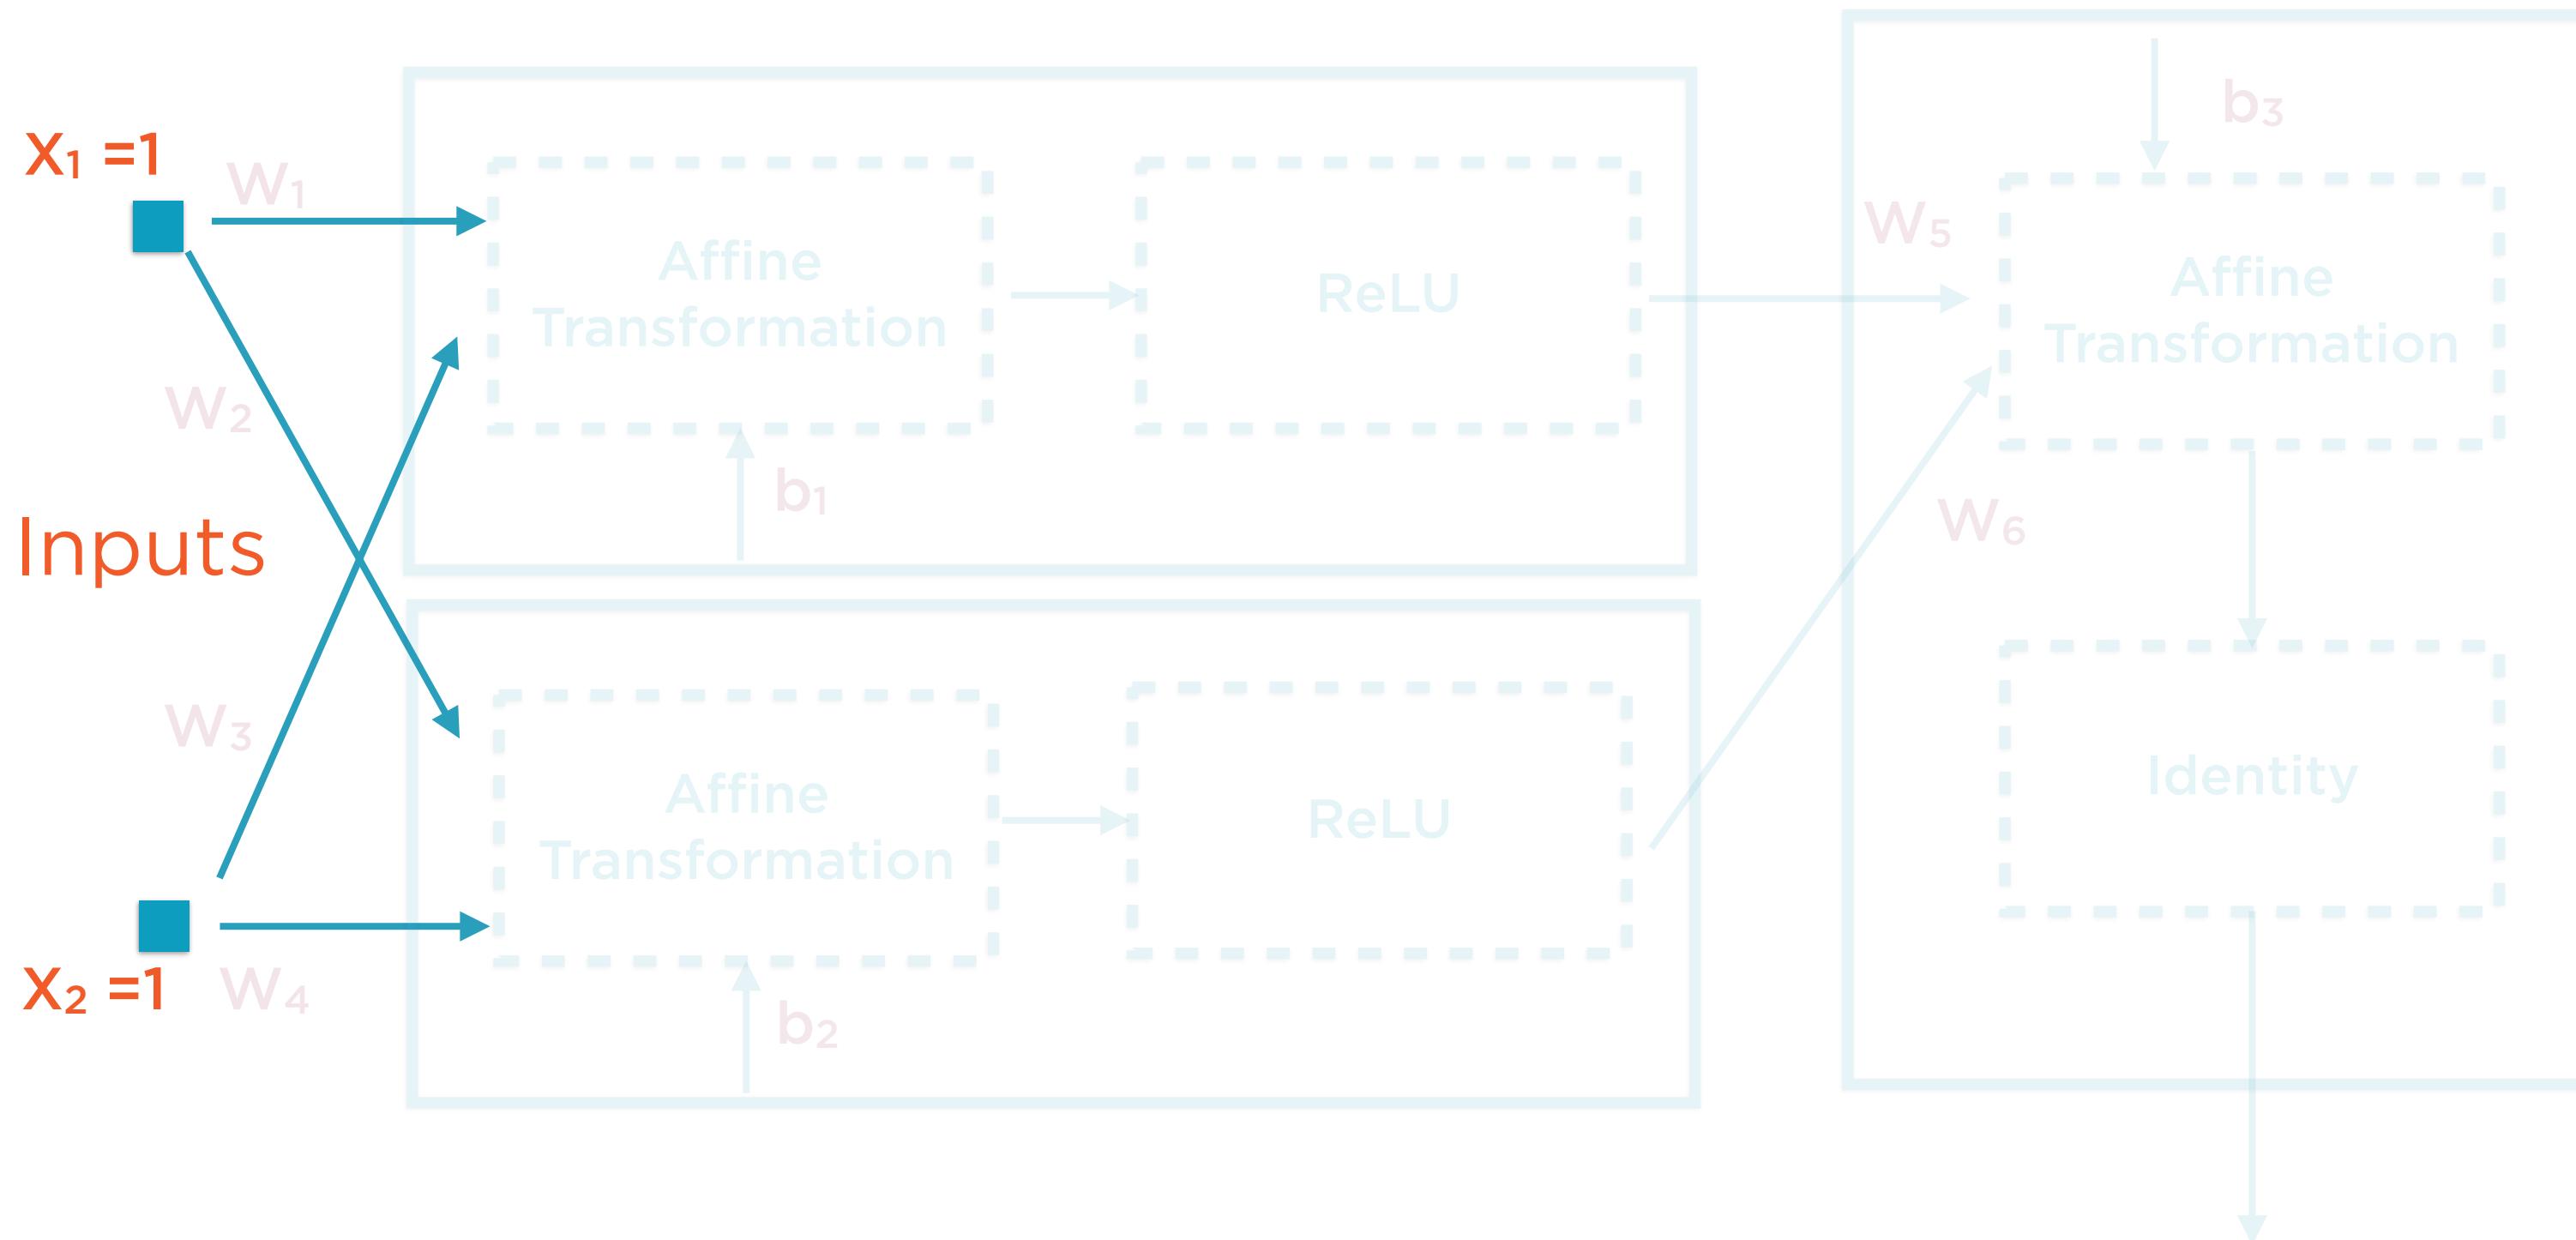


X ₁	X ₂	Y
0	0	0
0	1	1
1	0	1
1	1	0

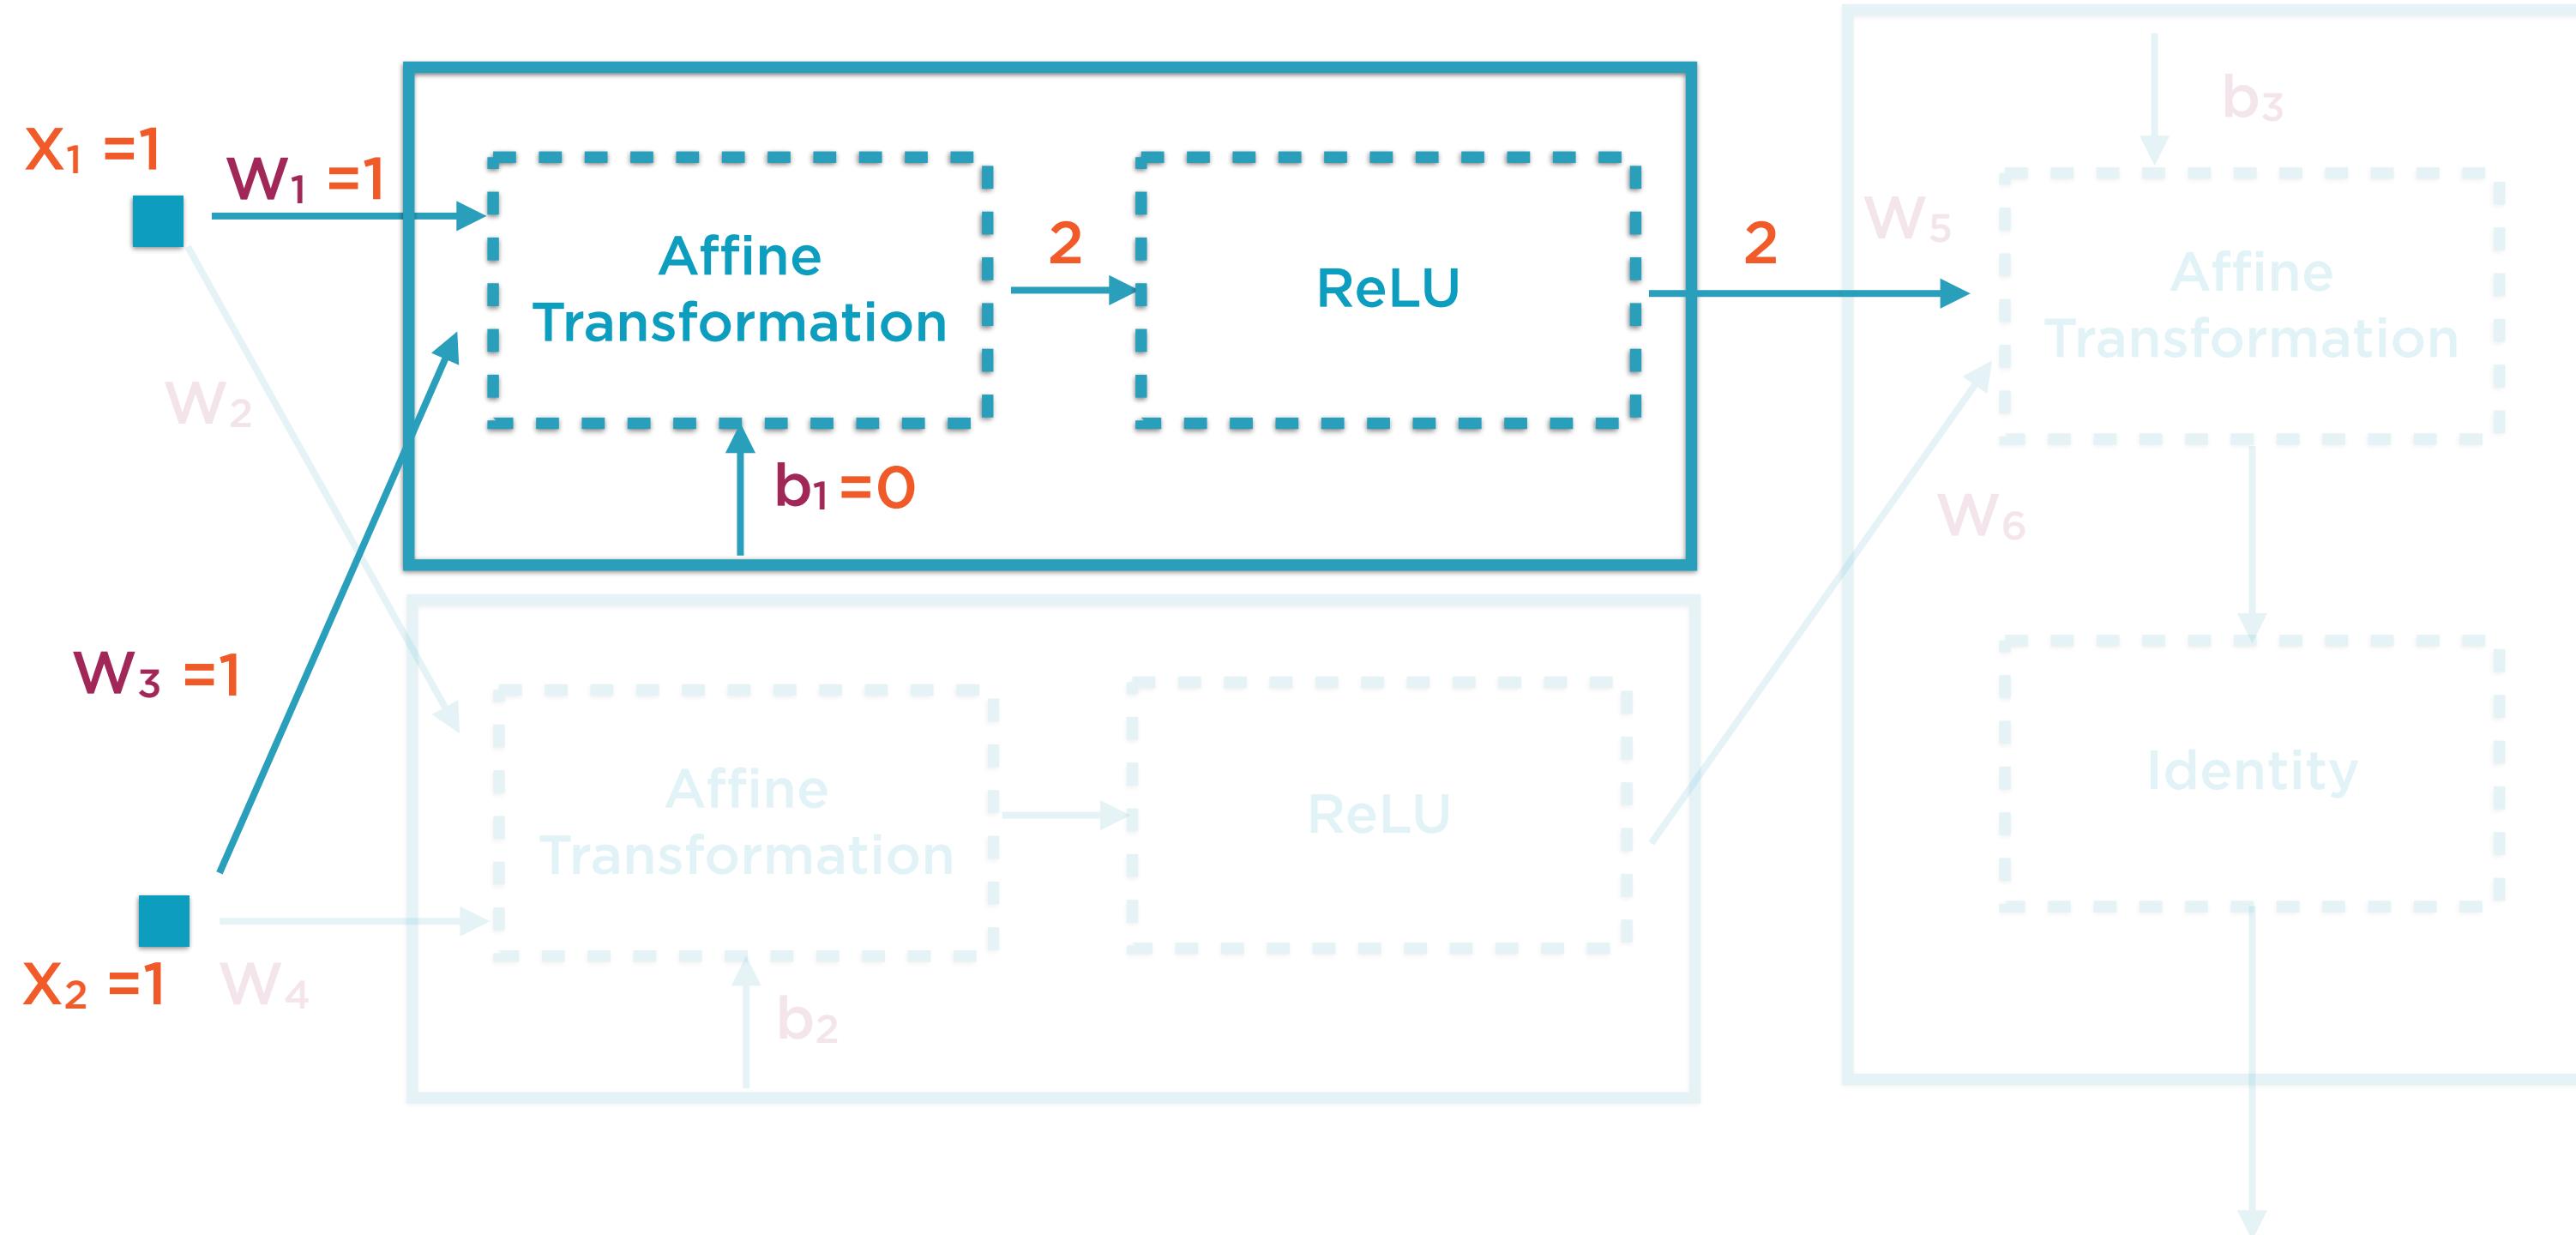
“Learning” XOR

Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

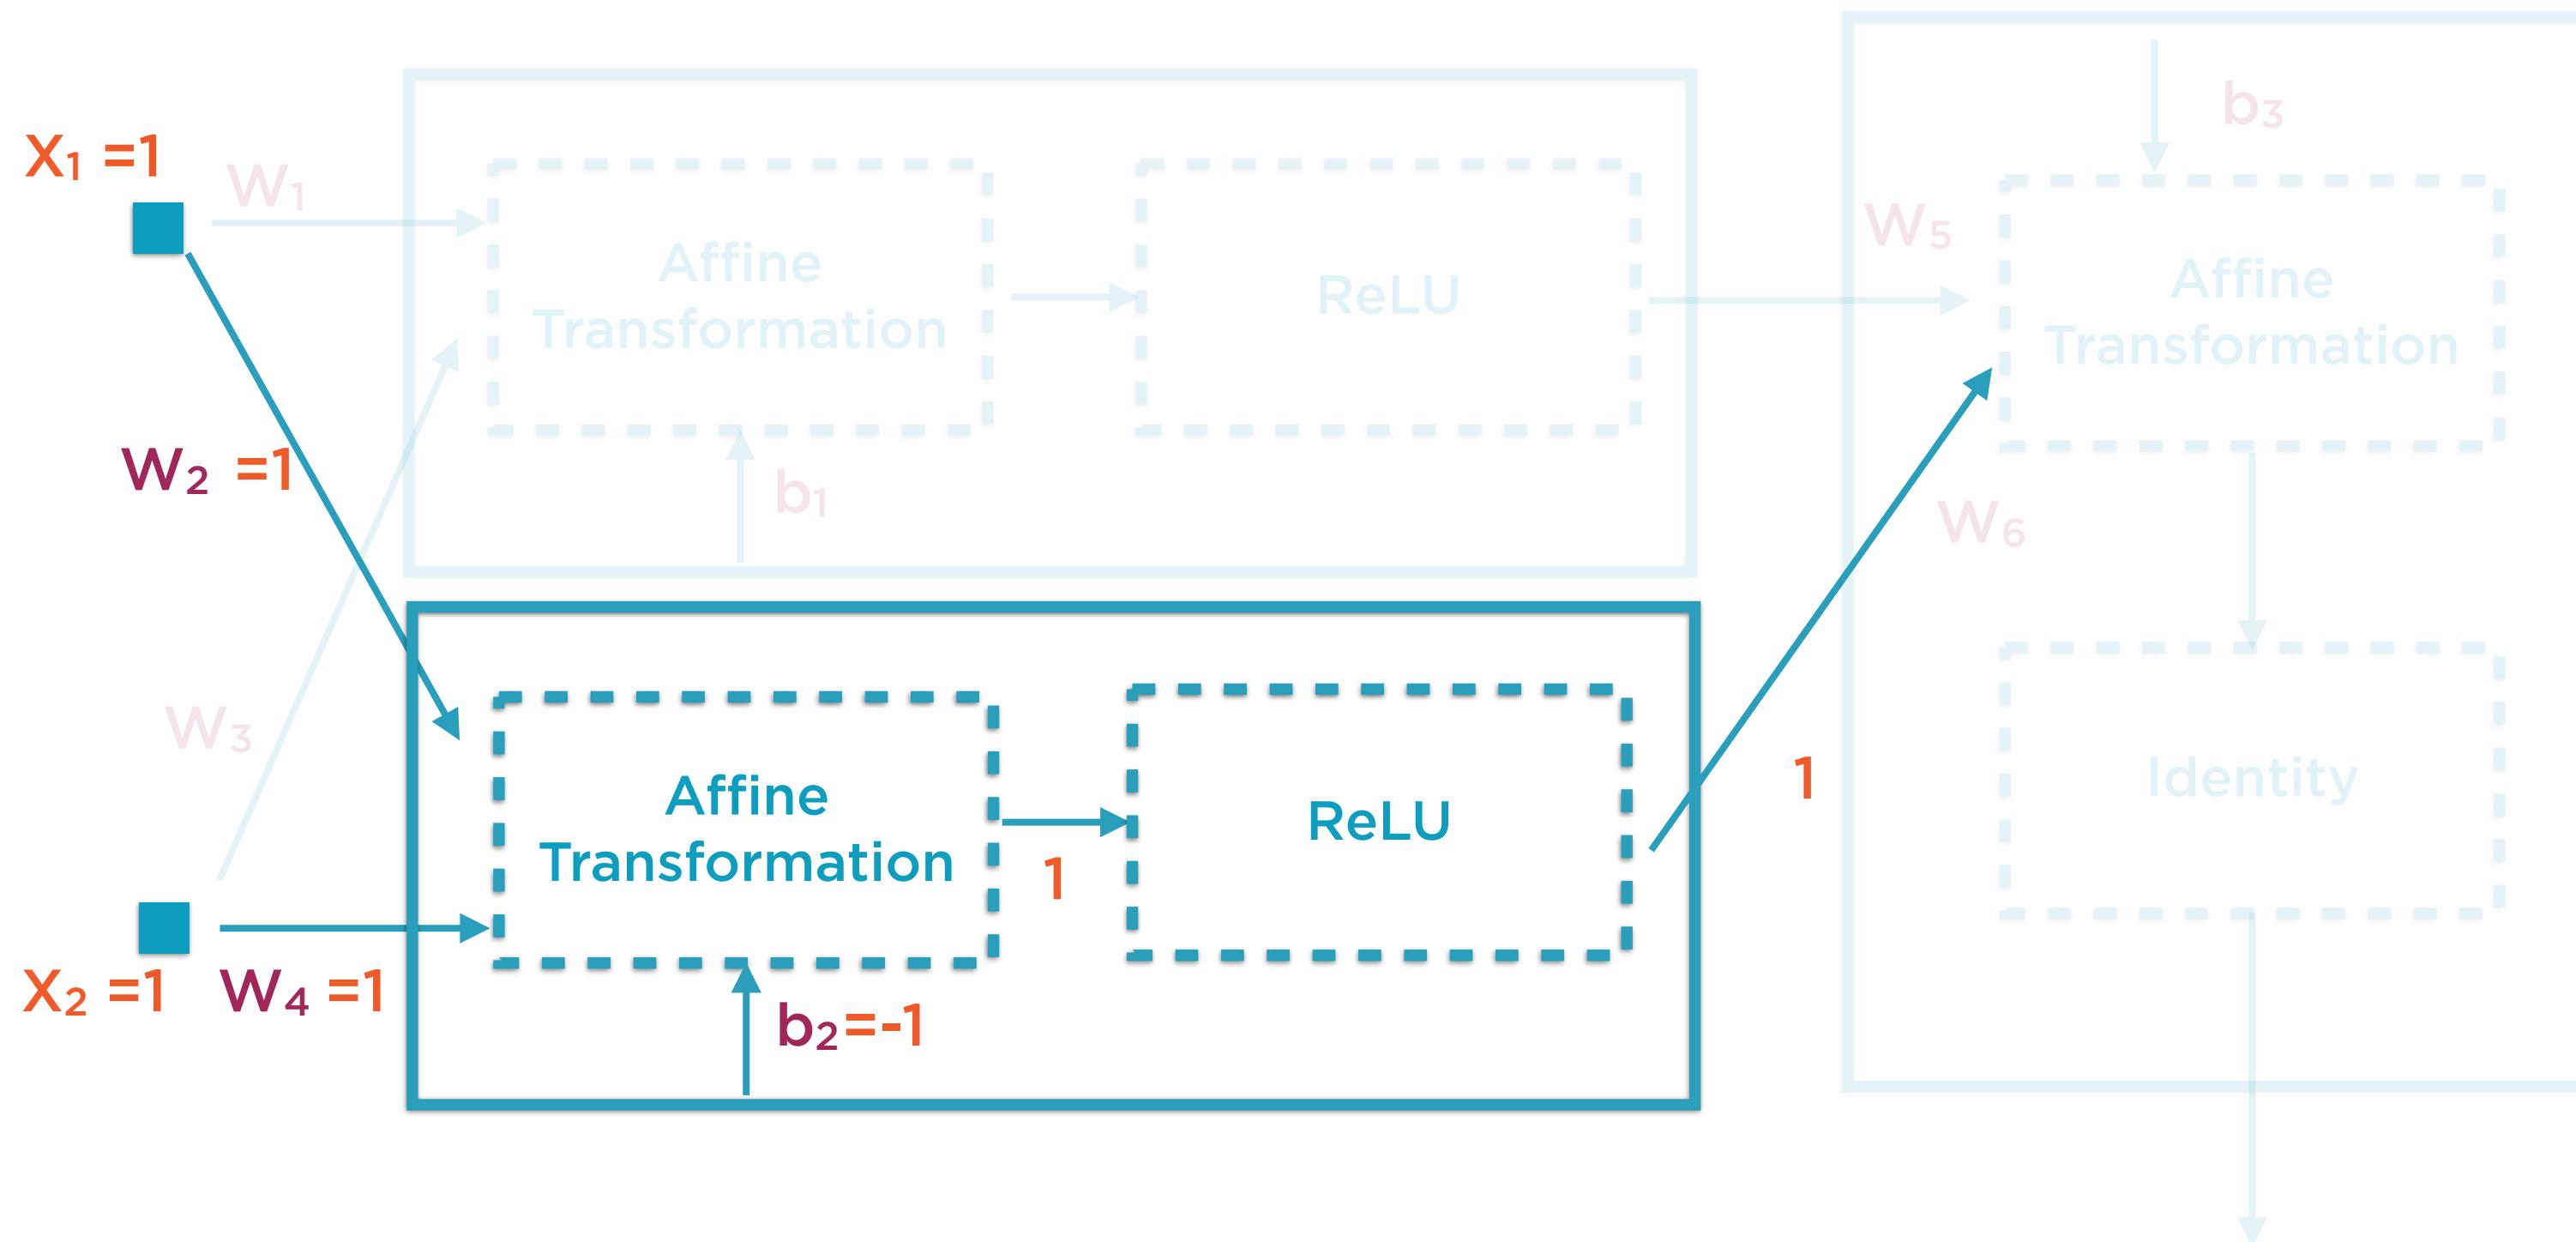
3-Neuron XOR



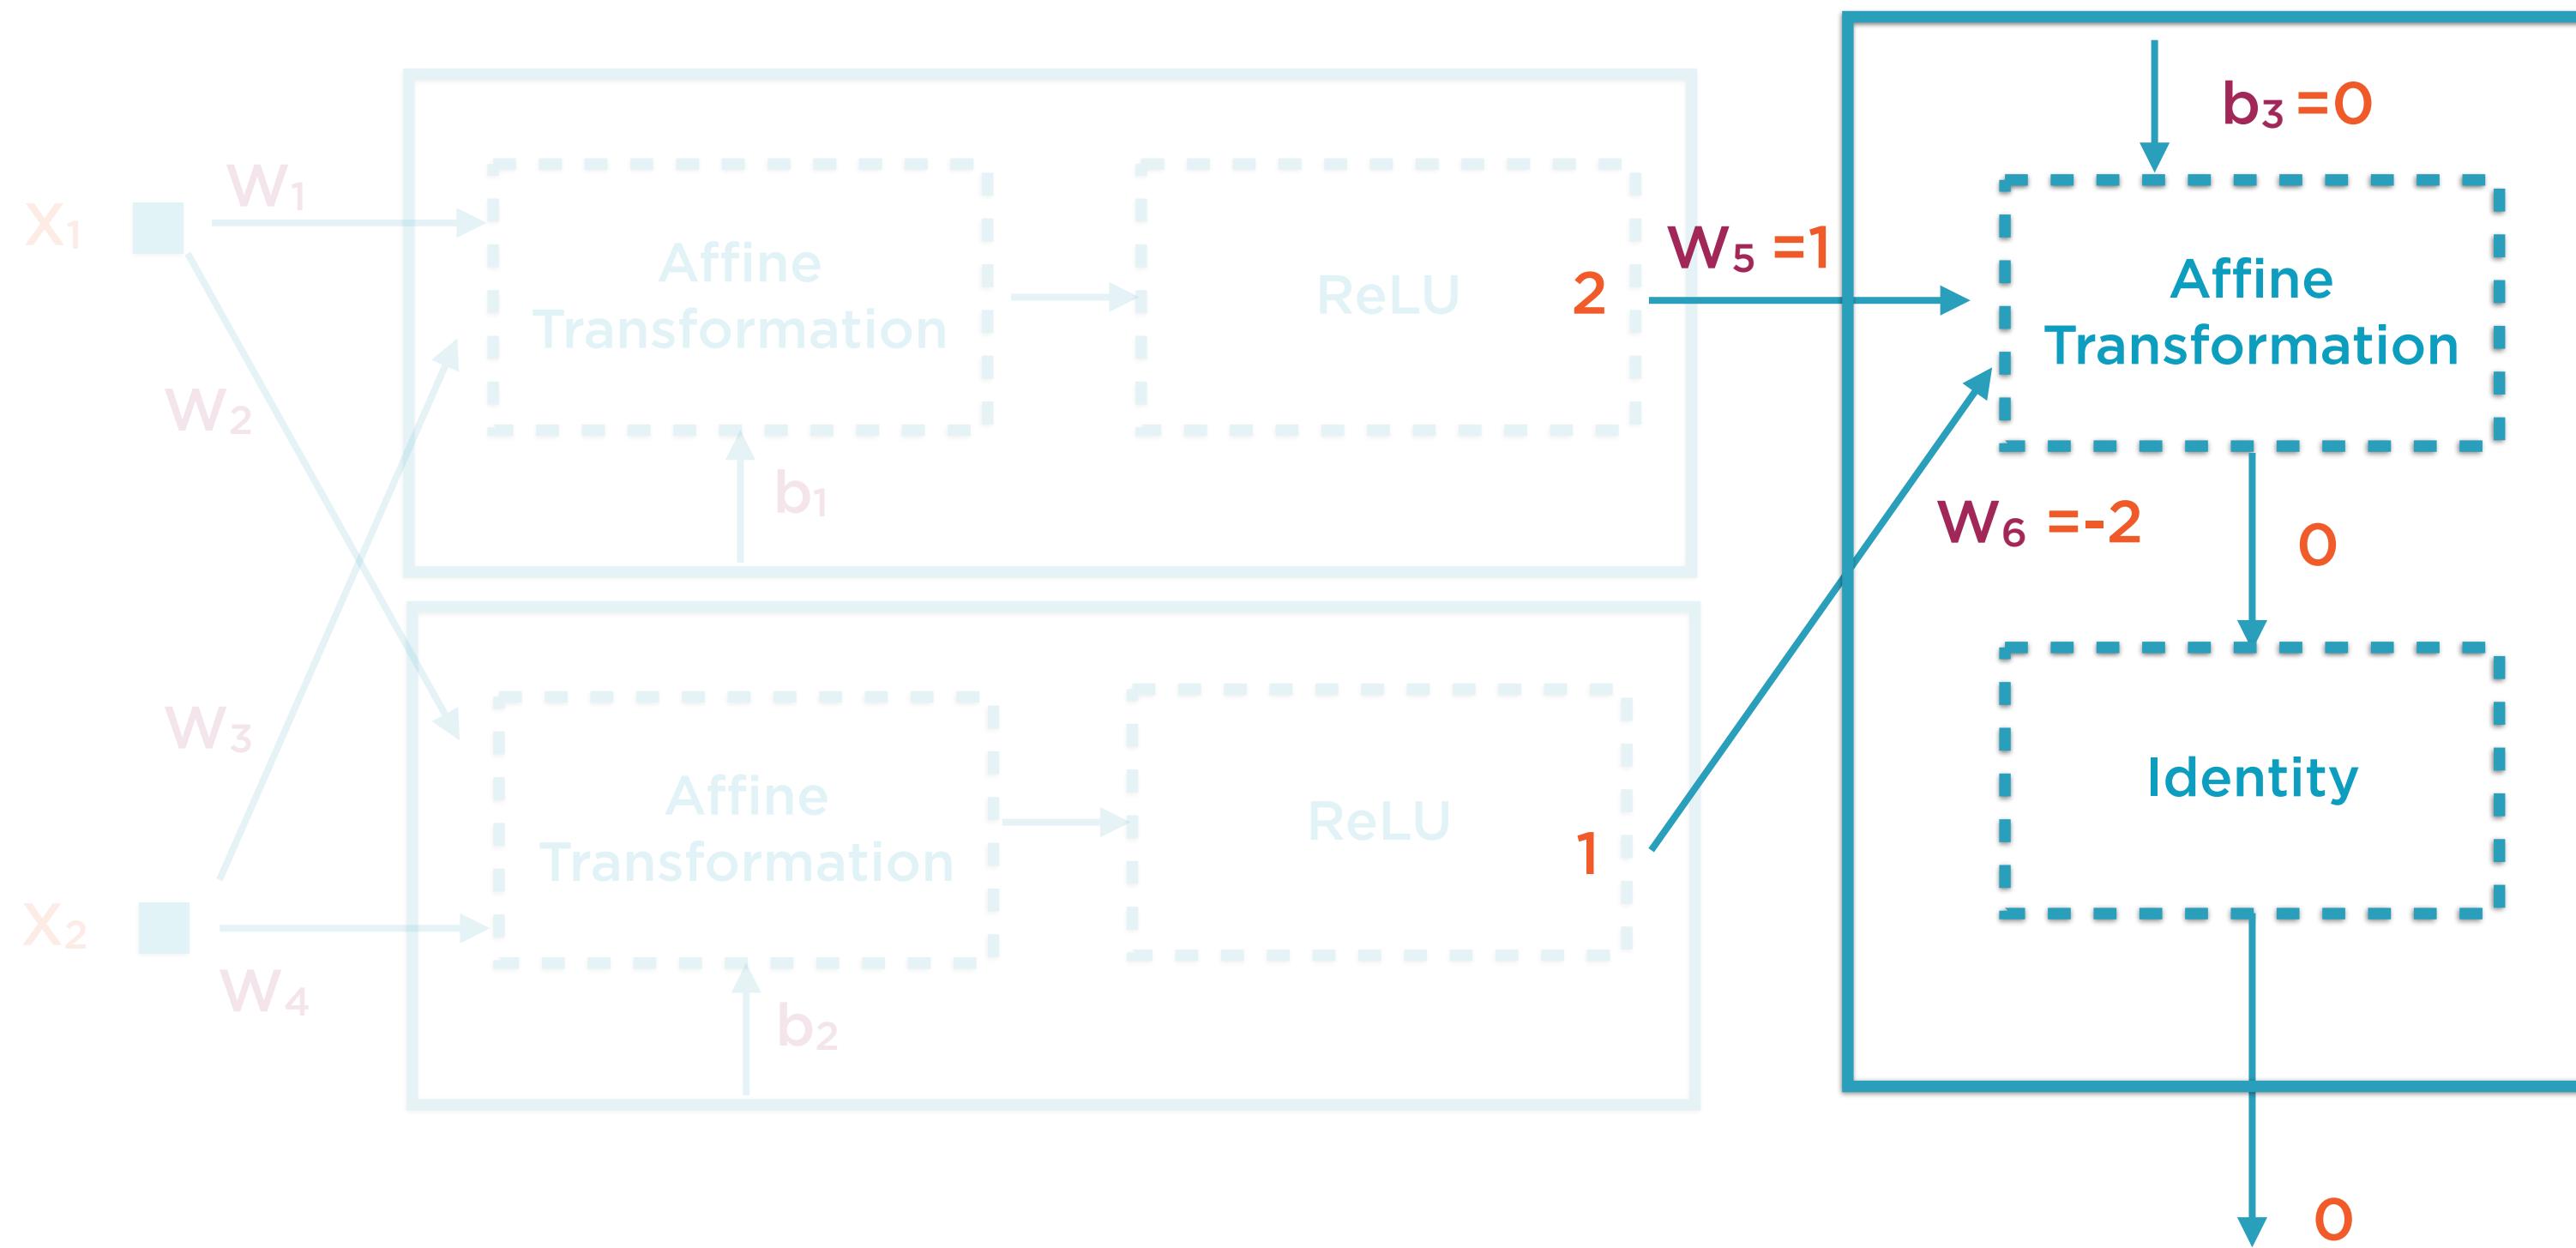
Weights and Bias of Neuron #1



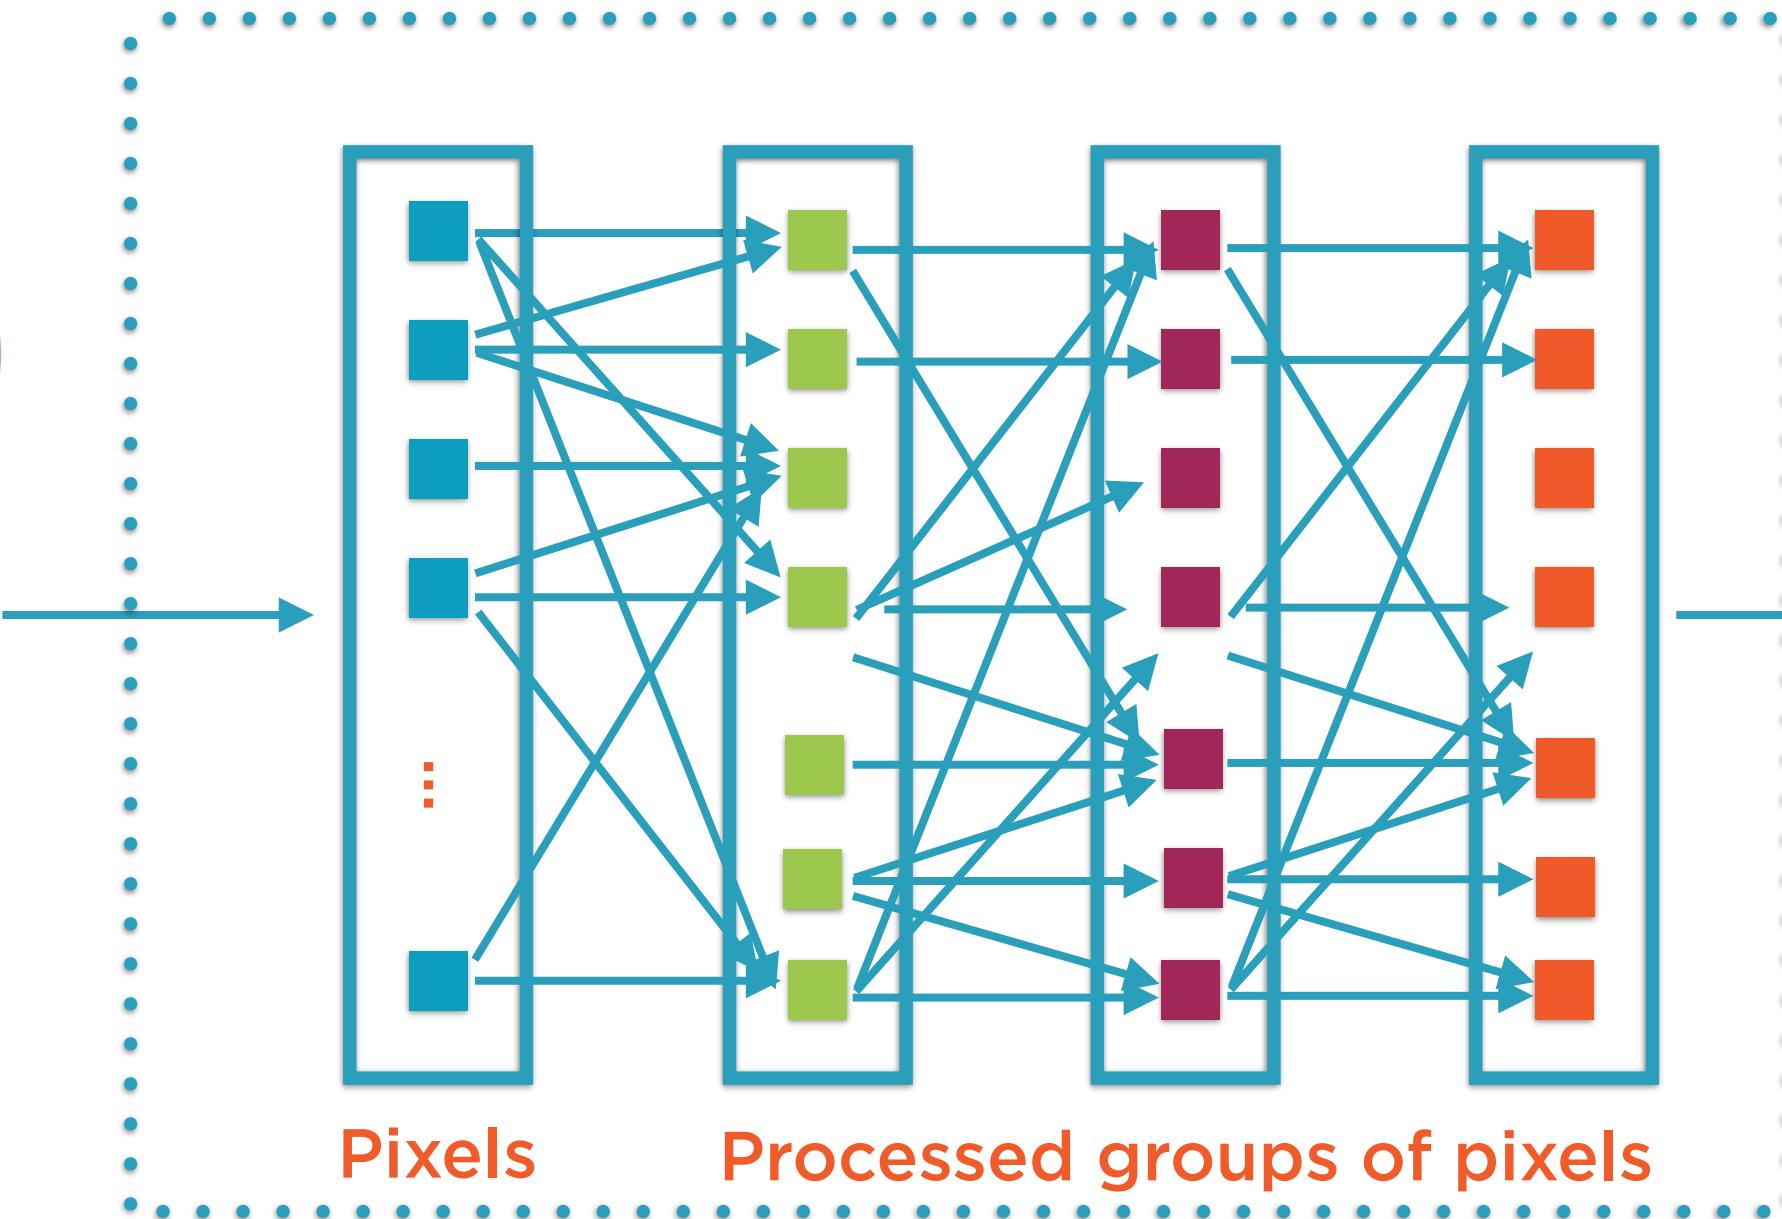
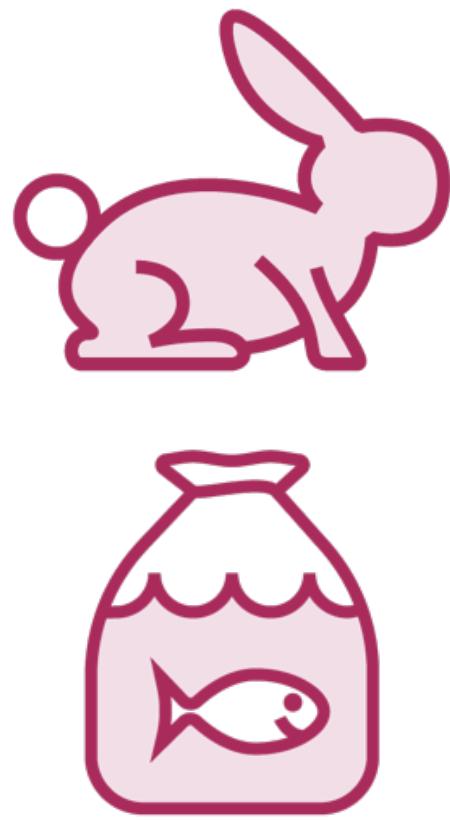
Weights and Bias of Neuron #2



Weights and Bias of Neuron #3



Choice of Activation Function



Corpus of
Images

Neural Network

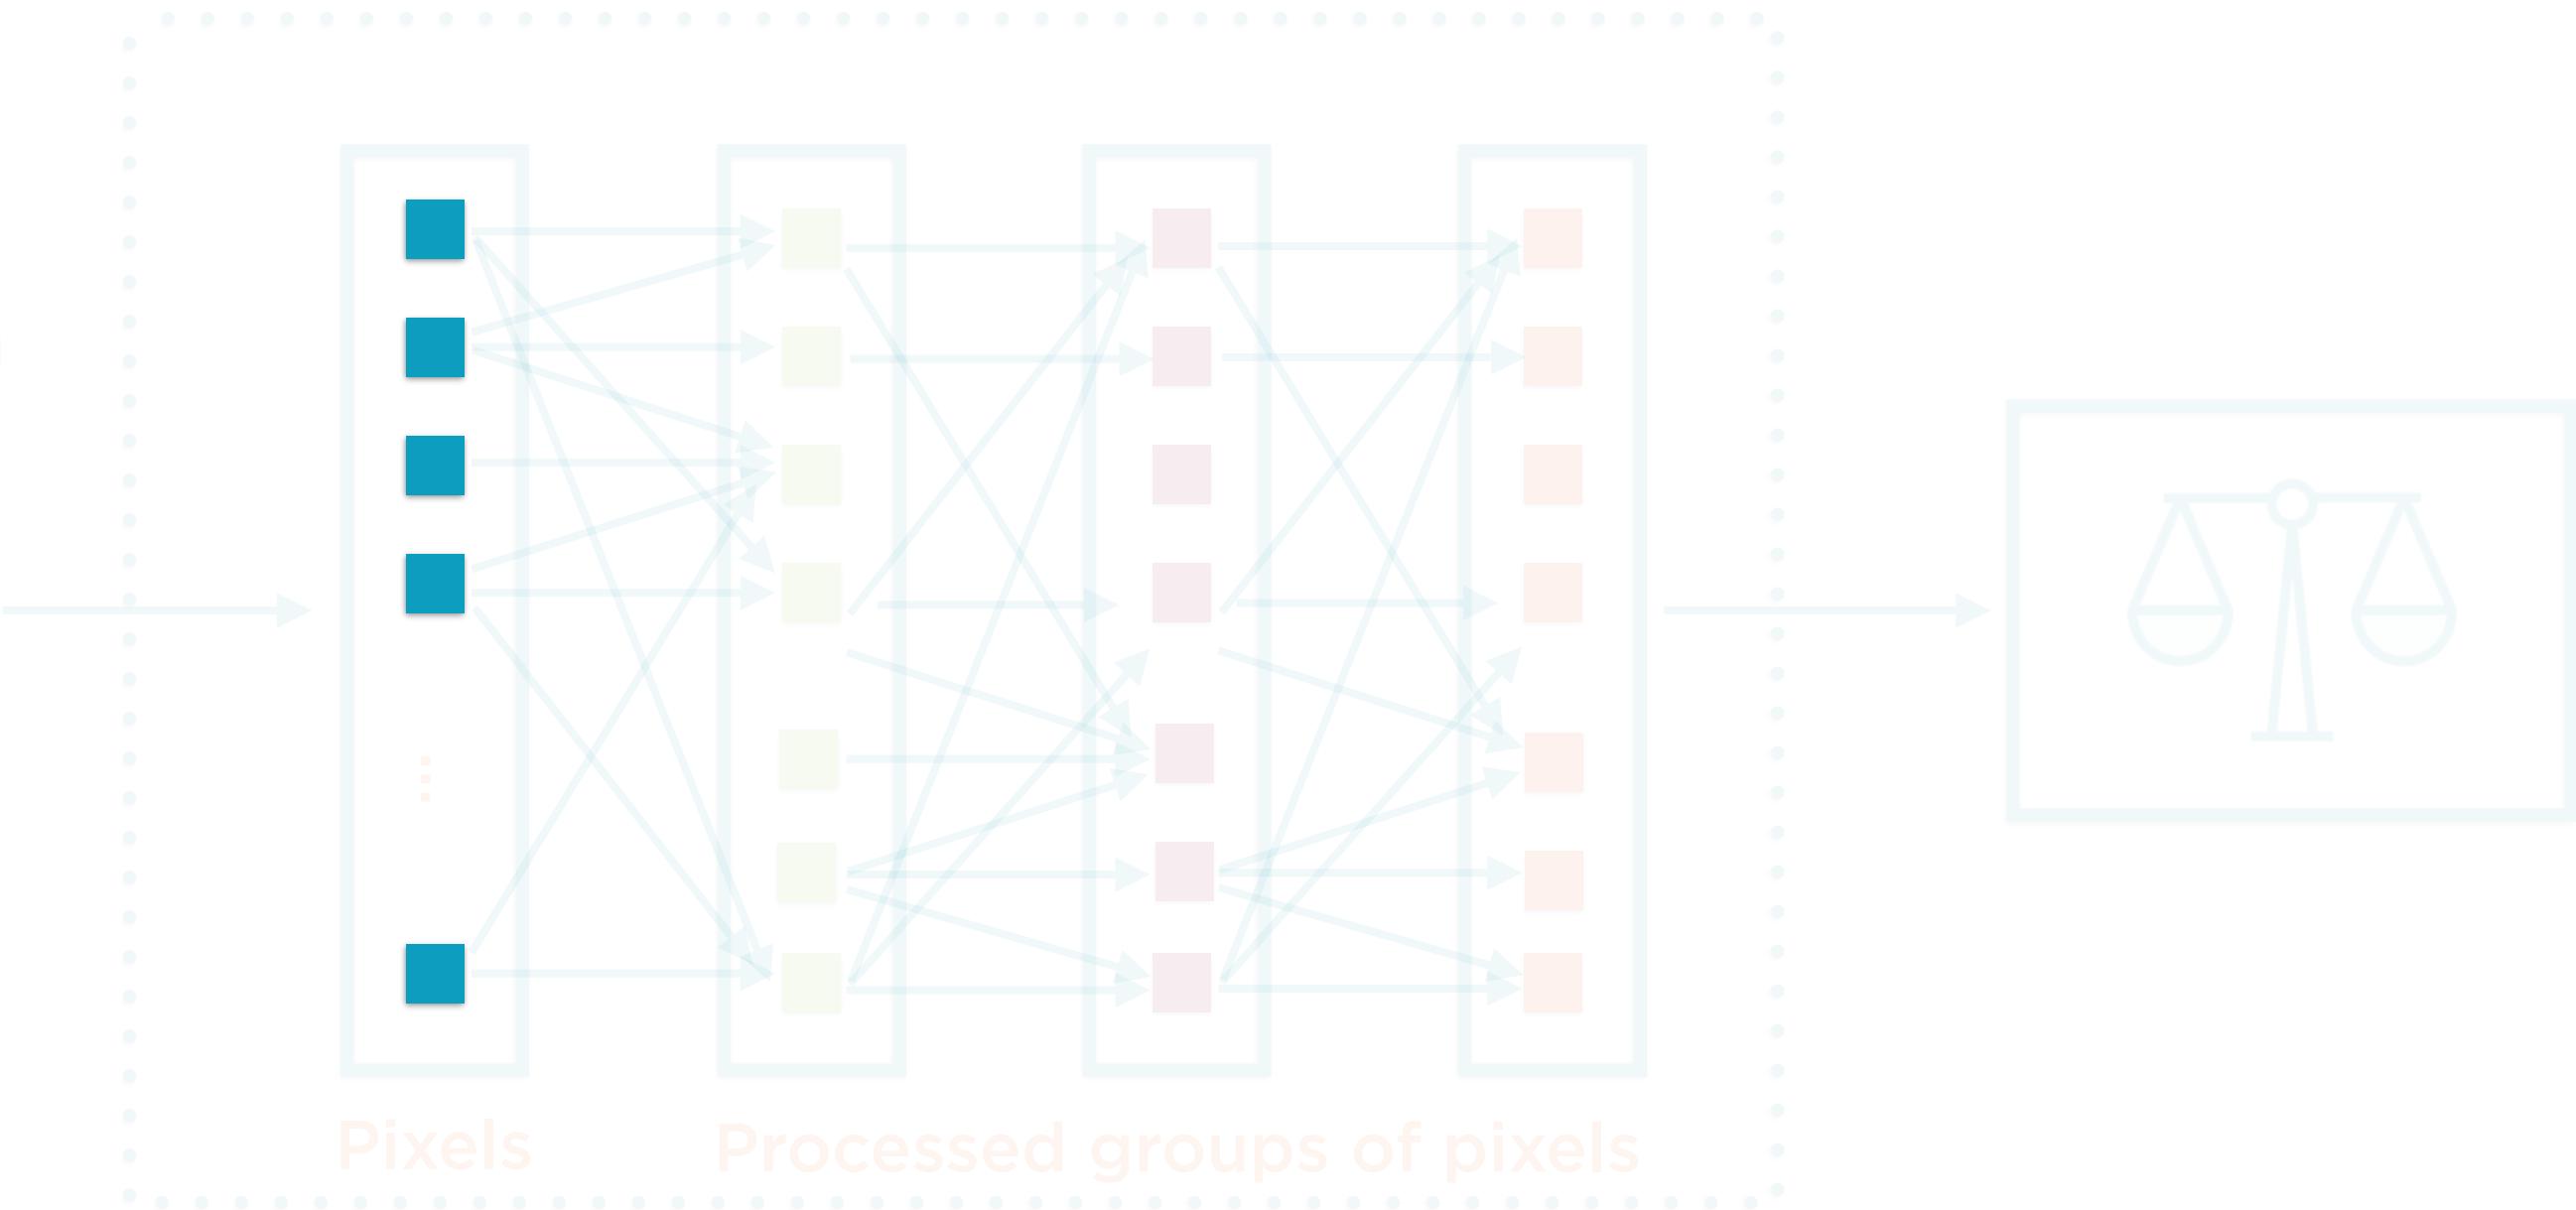
ML-based Classifier



Choice of Activation Function



Corpus of
Images



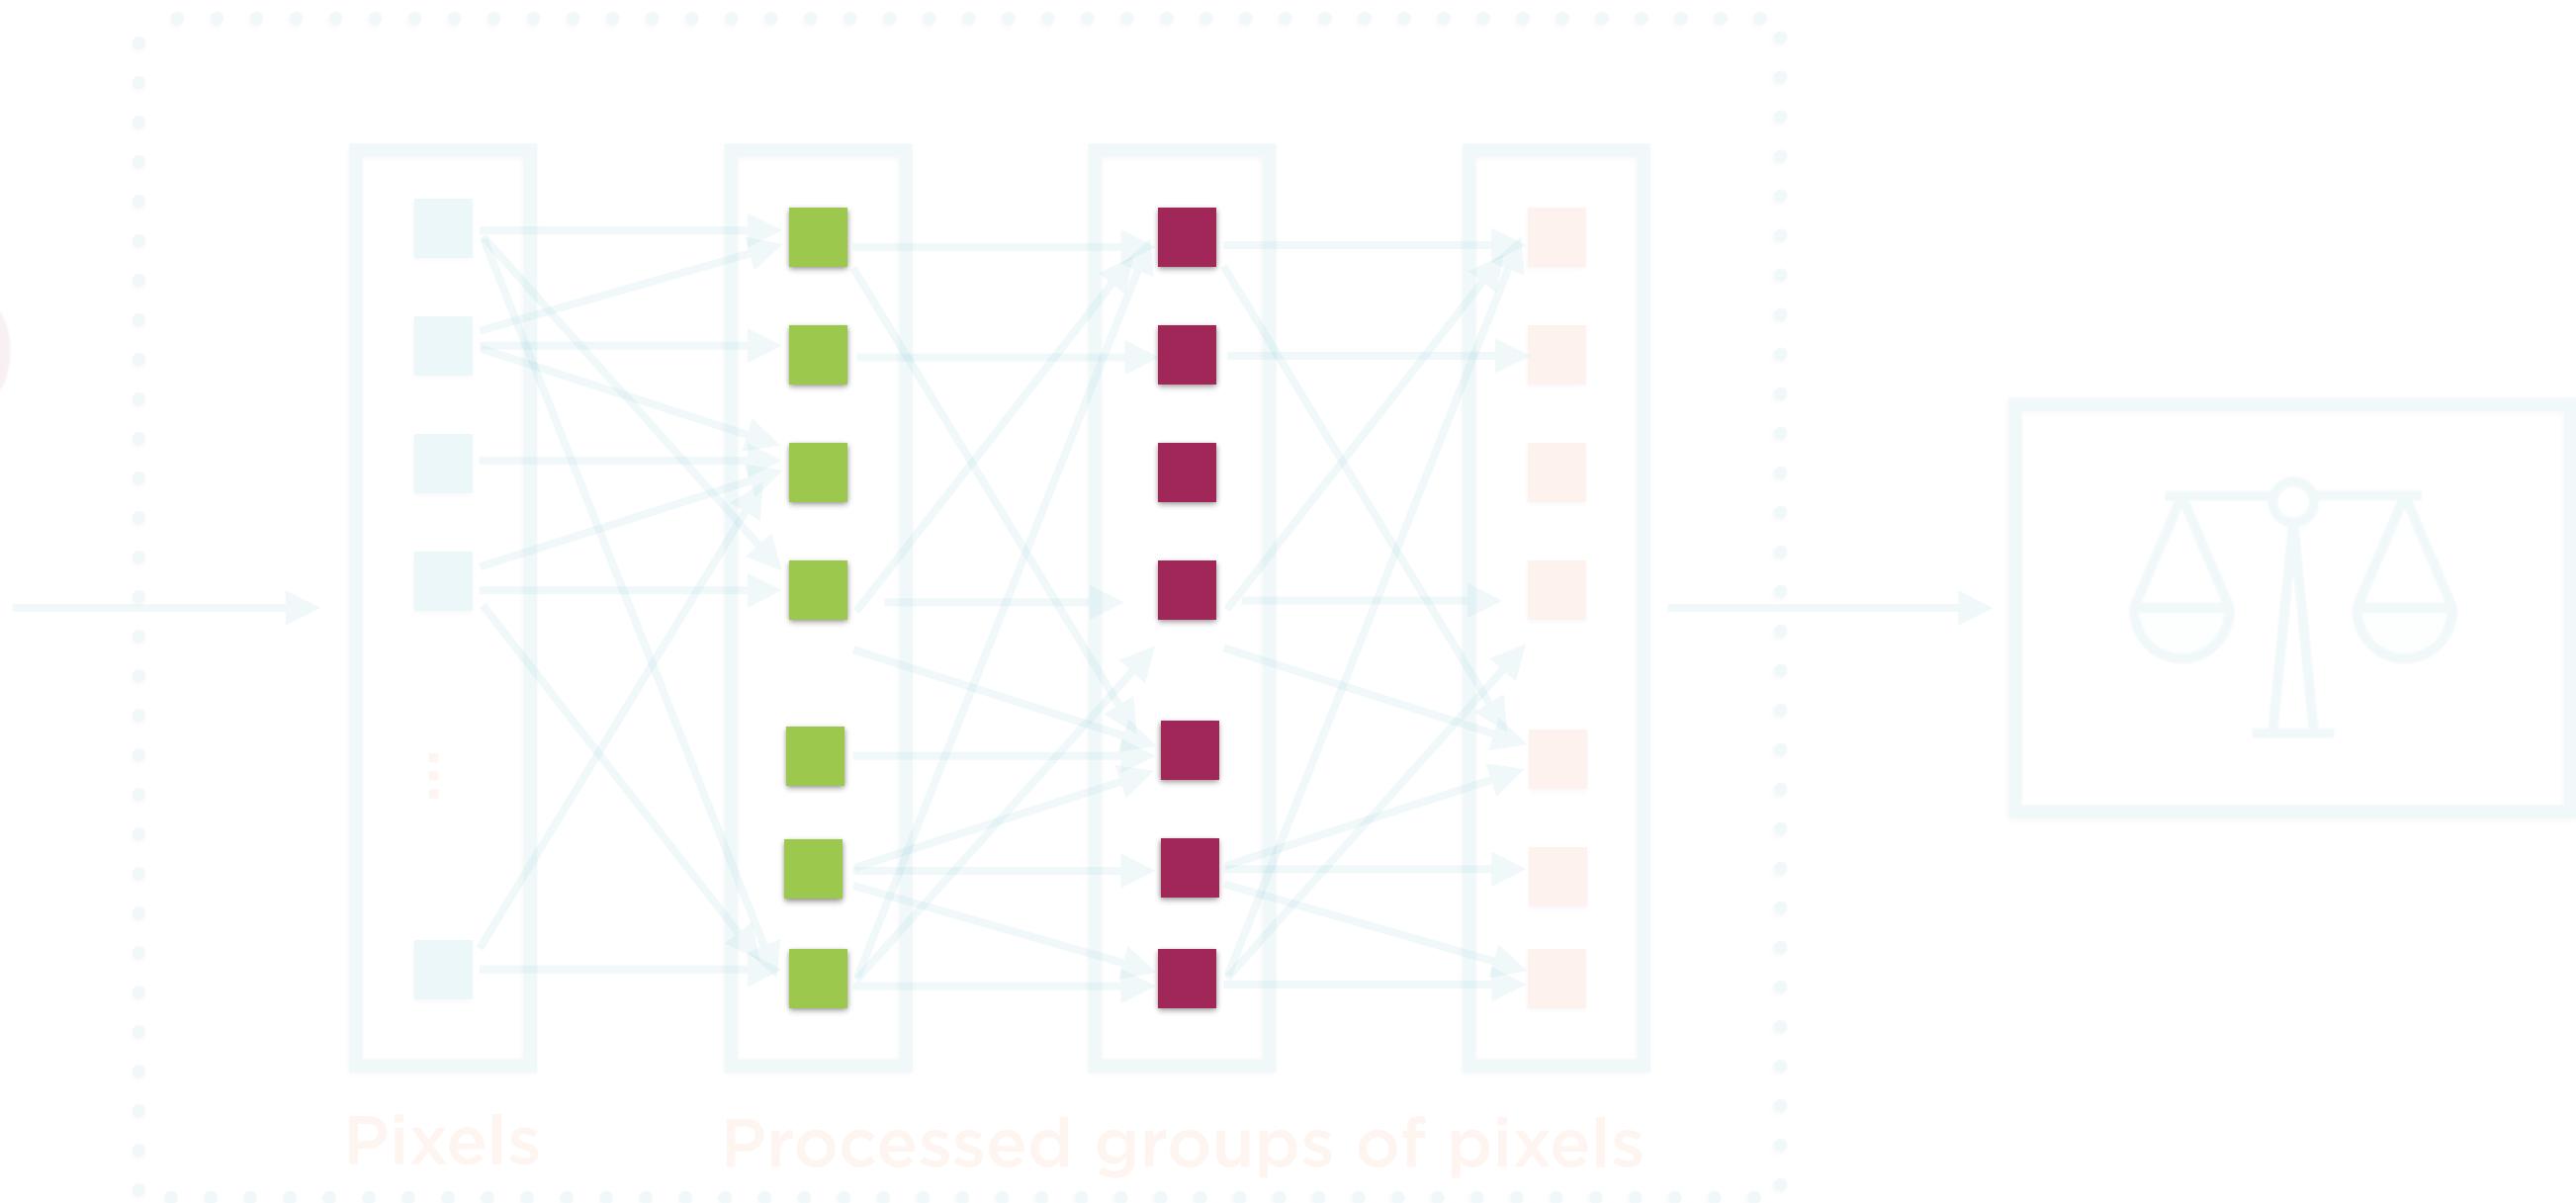
**Input layers use identity
function as activation: $f(x) = x$**

ML-based Classifier

Choice of Activation Function



Corpus of
Images



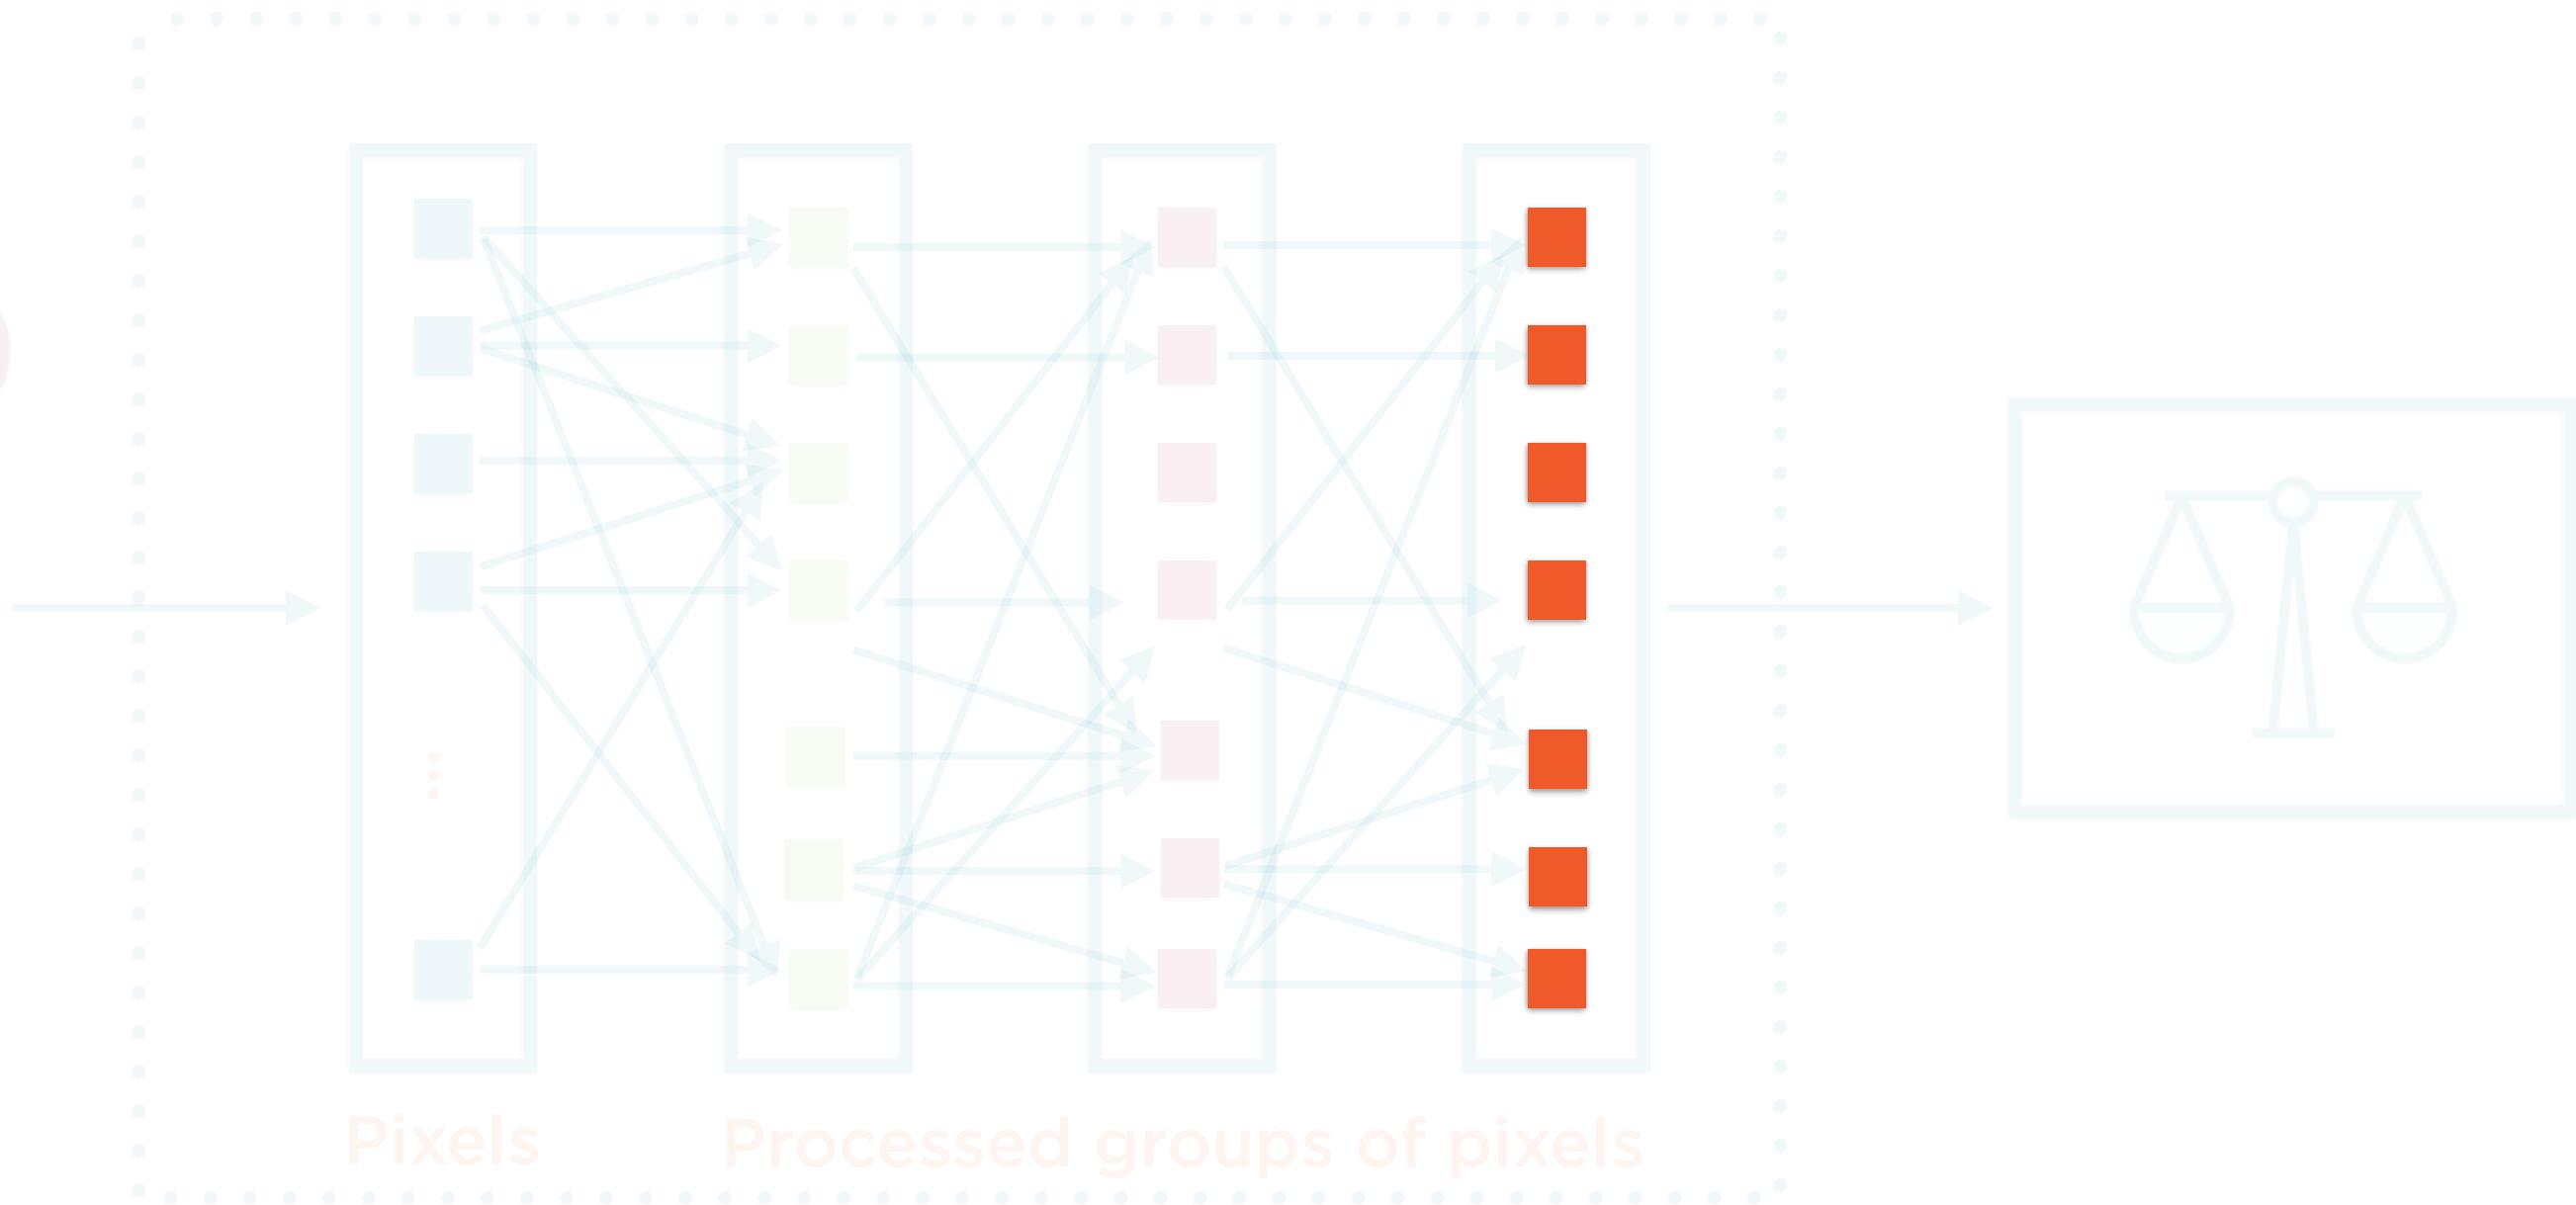
**Inner hidden layers typically
use ReLU as activation function**

ML-based Classifier

Choice of Activation Function



Corpus of
Images



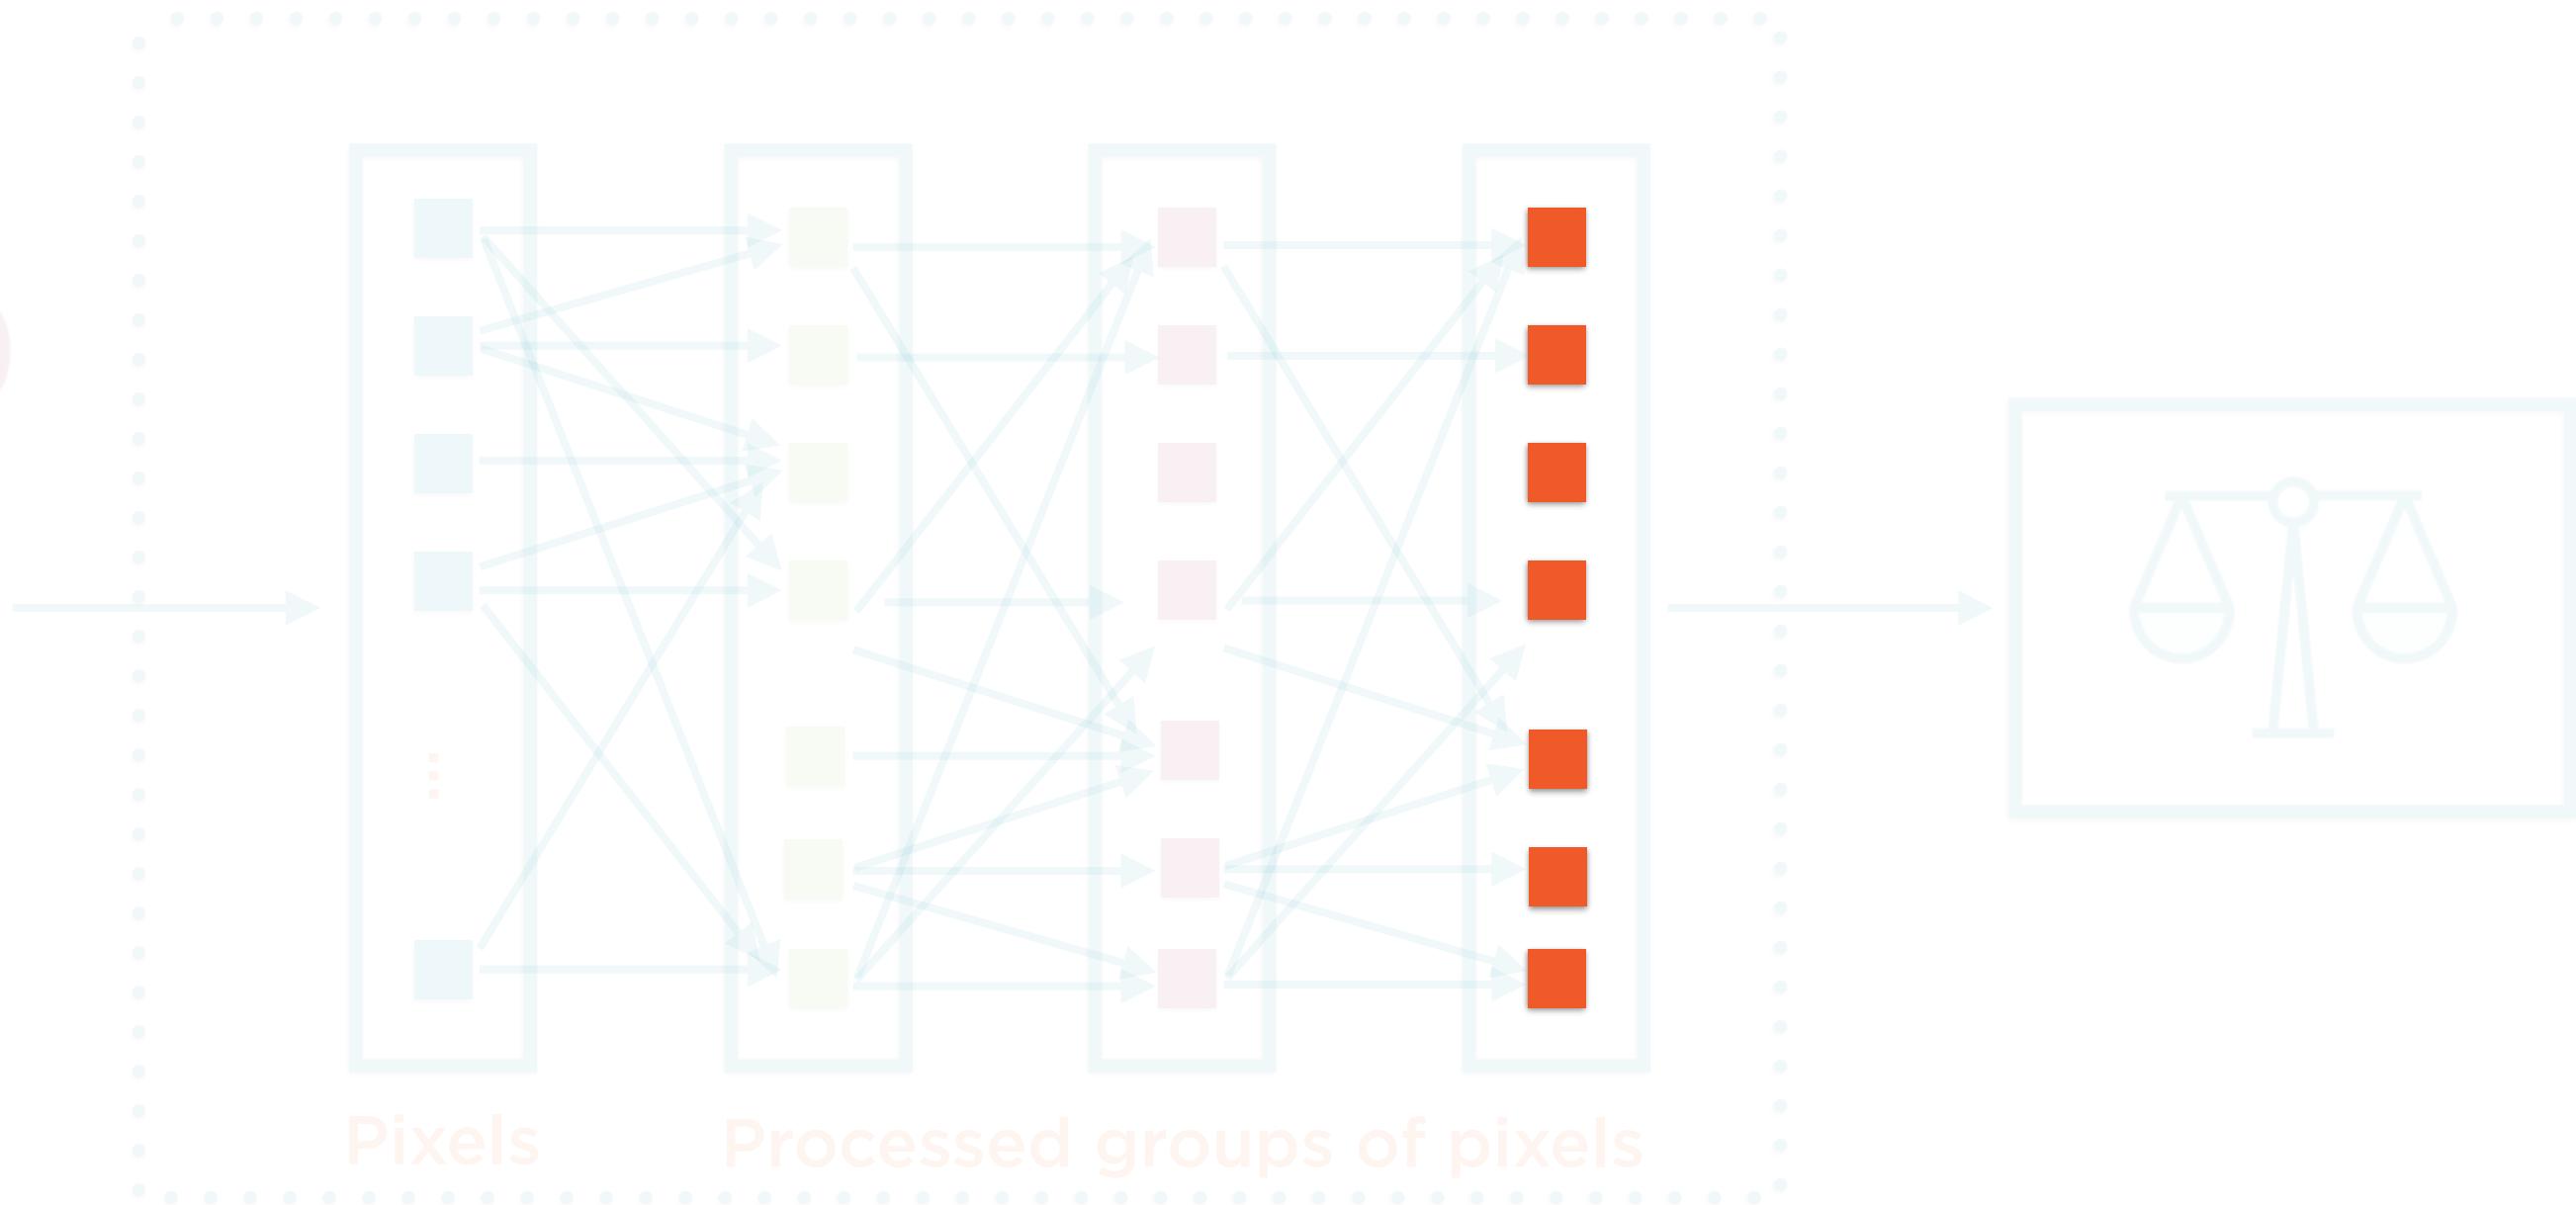
**Output layer in our XOR example
used the identity function**

ML-based Classifier

Choice of Activation Function



Corpus of
Images



**Output layer in classification will
often use SoftMax**

ML-based Classifier



Another very common form of the activation function is the SoftMax

SoftMax(x) outputs a number between 0 and 1

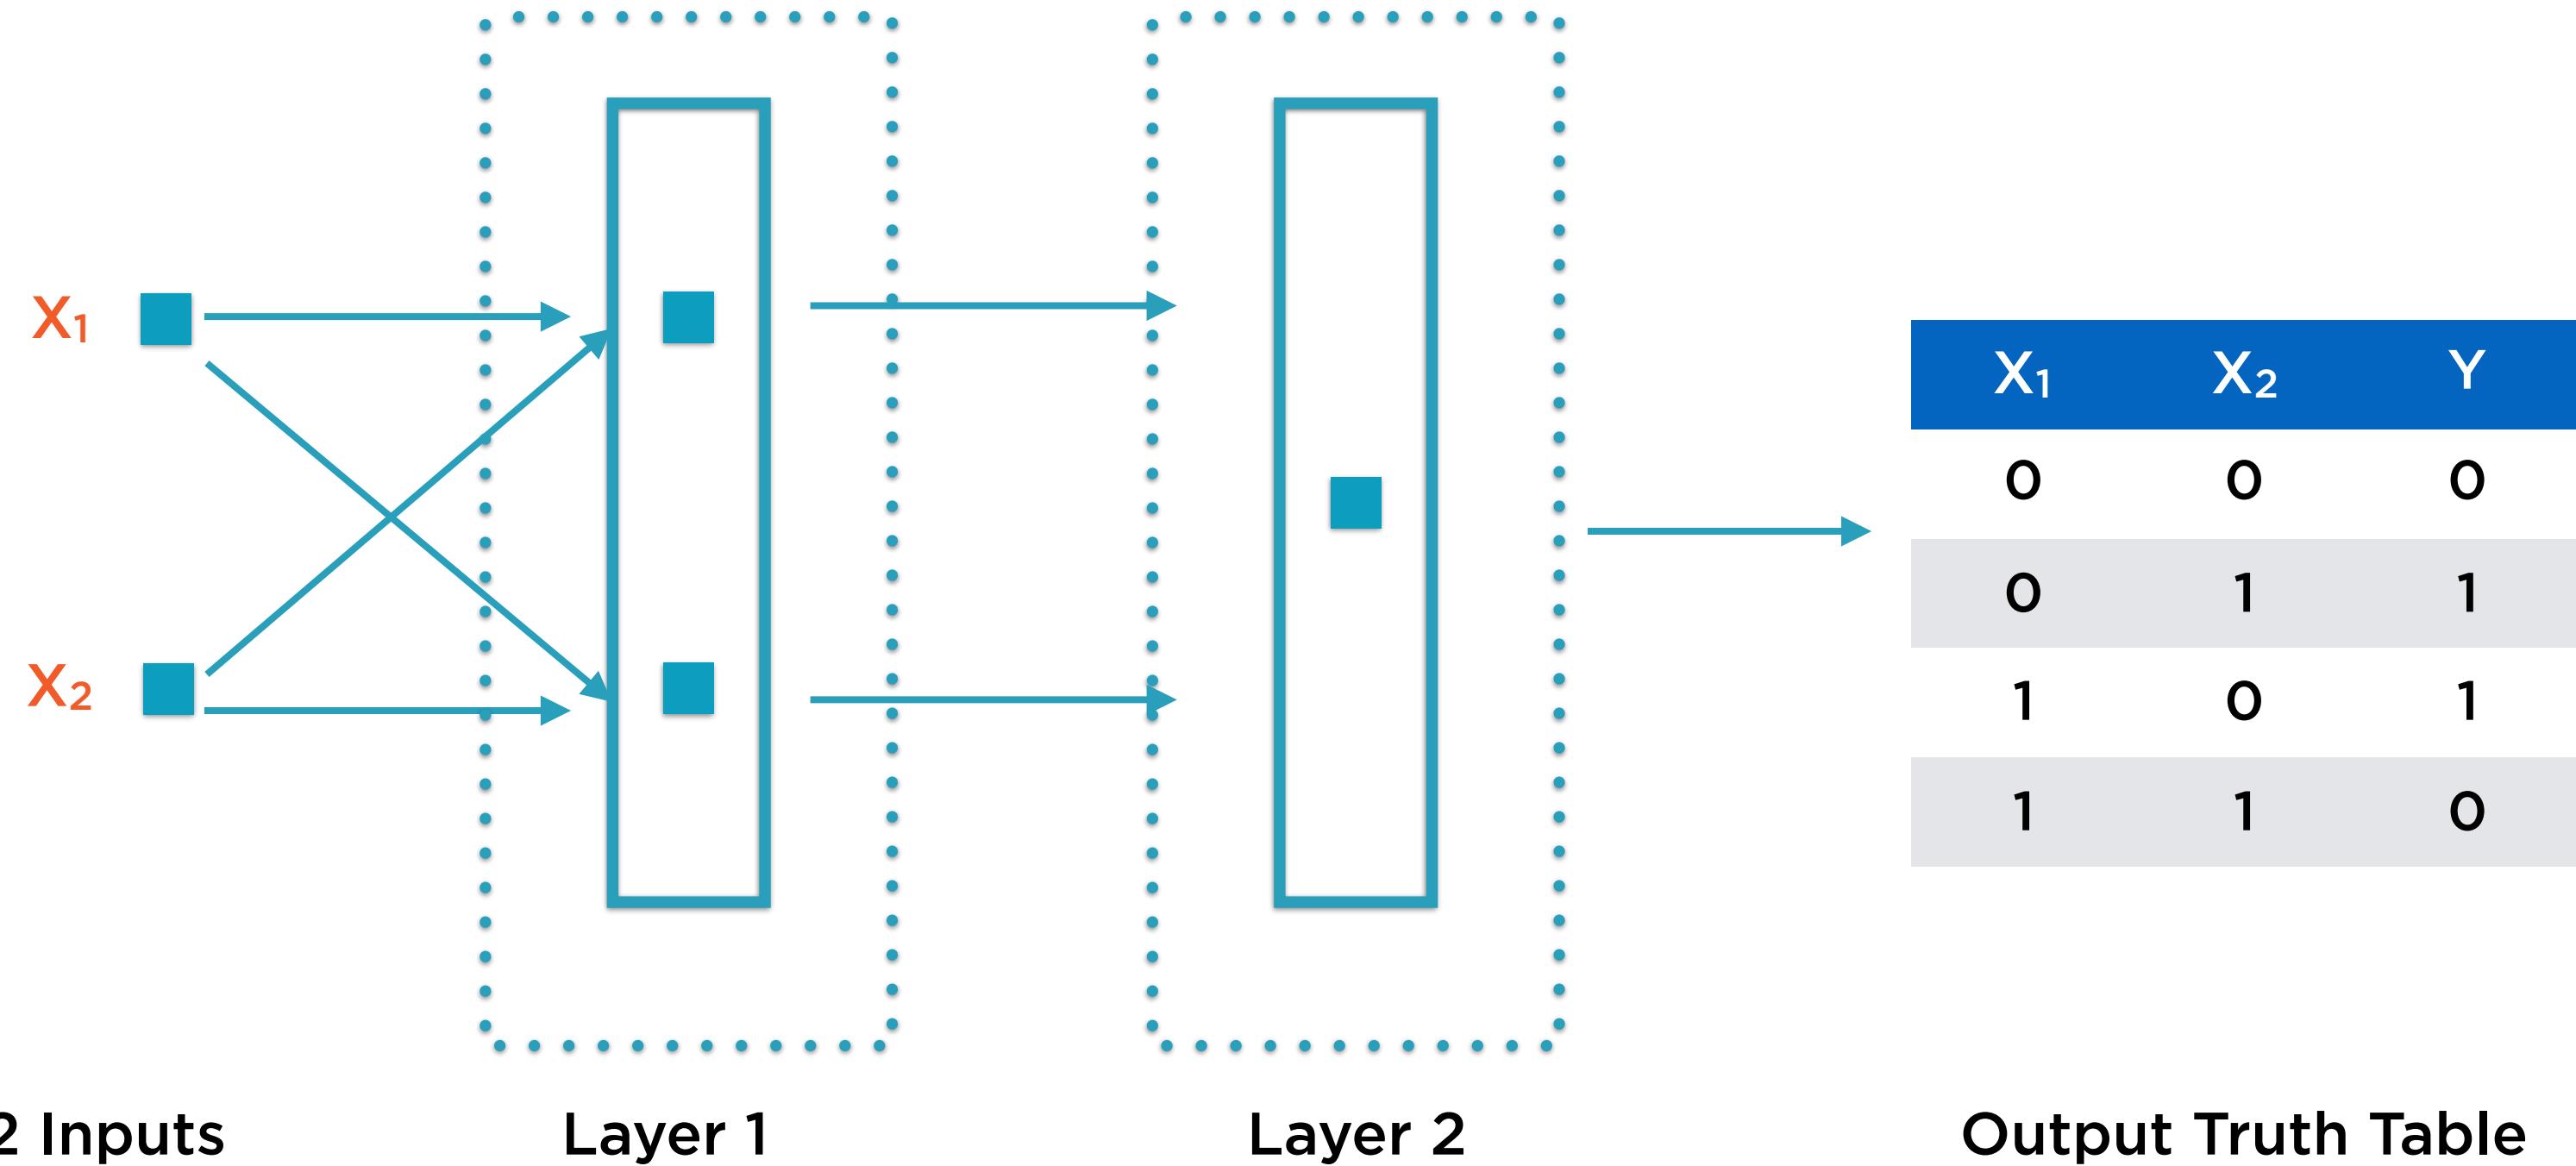
This output can be interpreted as a probability

```
def XOR(x1, x2):  
    if (x1 == x2):  
        return 0  
    return 1
```

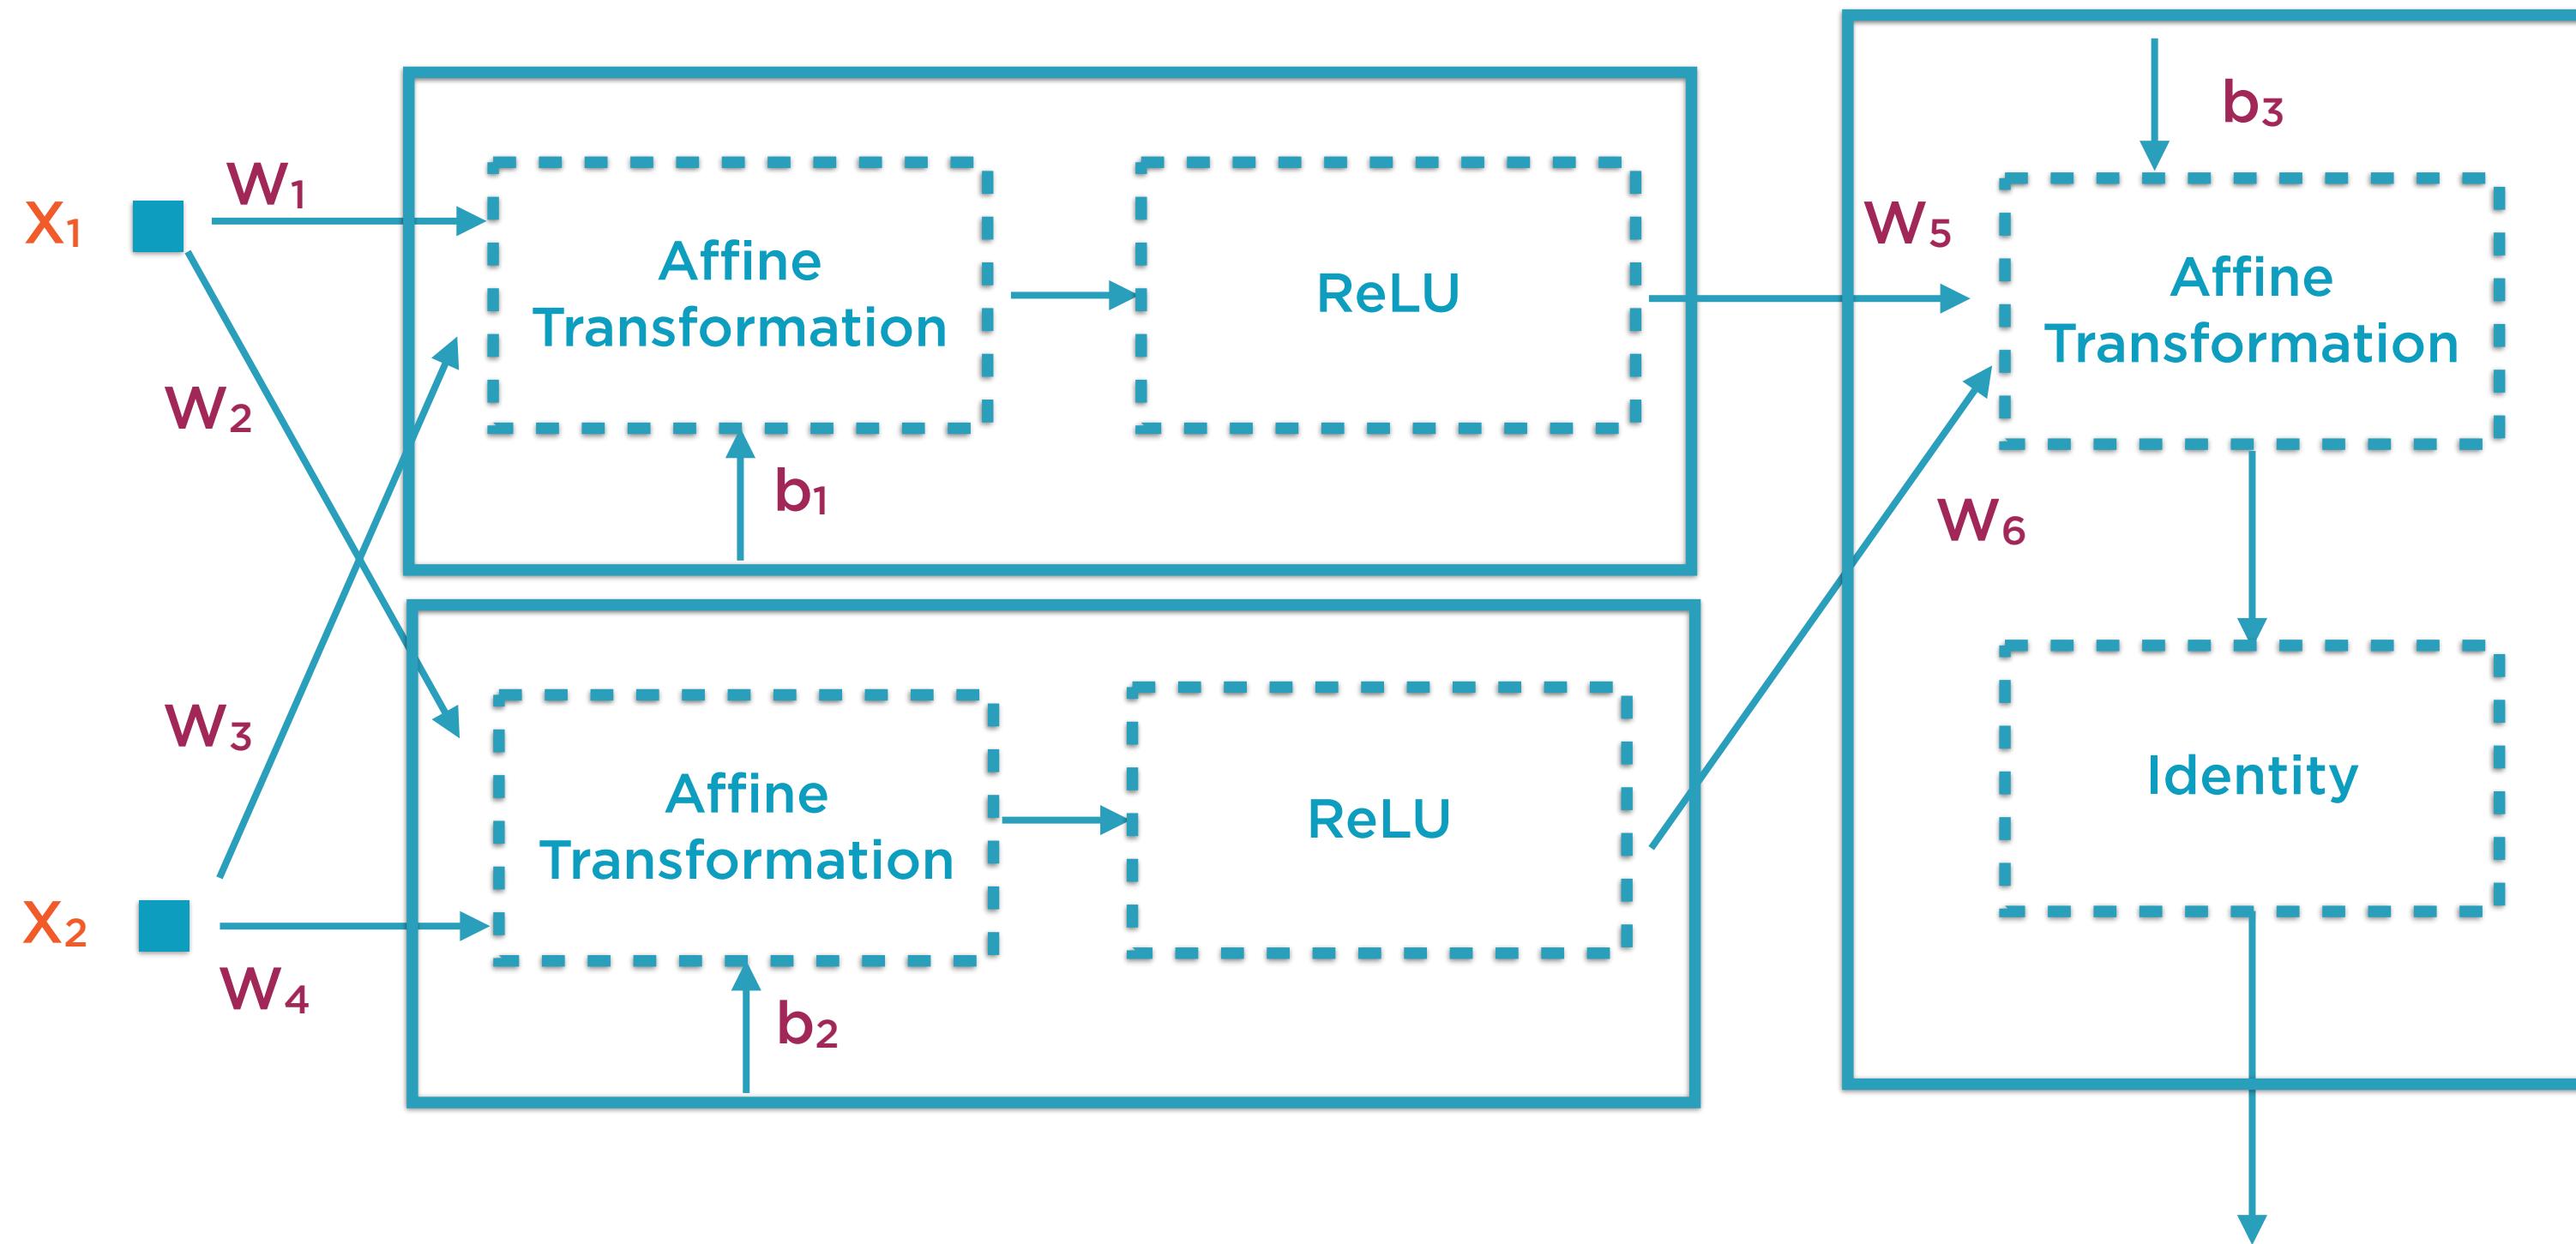
“Learning” XOR

Reverse-engineering XOR requires 3 neurons (arranged in 2 layers) as well as a non-linear activation function

XOR: 3 Neurons, 2 Layers



3-Neuron XOR



```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

“Learning” Arbitrarily Complex Functions

Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything

Summary

A neuron is the smallest entity in a neural network

Linear regression can be learnt by a single neuron

A more complex function such as XOR requires more neurons

Combinations of interconnected neurons can “learn” virtually anything

Training such networks to use the “best” parameter values is vital