

Building Logistic Regression Models Using TensorFlow



Vitthal Srinivasan

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Given causes, predict probability of effects - that's logistic regression

Linear regression and logistic regression are similar, yet quite different

Logistic regression can be used for categorical y-variables

Logistic regression in TensorFlow differs from linear regression in two ways

- Softmax as the activation function**
- cross-entropy as the cost function**

Two Approaches to Deadlines



Start 5 minutes before deadline

Good luck with that



Start 1 year before deadline

Maybe overkill

Neither approach is optimal

Starting a Year in Advance

Probability of meeting the deadline



100%

Probability of getting other important work done

| 0%

Starting Five Minutes in Advance

Probability of meeting the deadline

0%

Probability of getting other important work done

100%

The Goldilocks Solution

Work fast

Start very late and hope
for the best

Work smart

Start as late as possible
to be sure to make it

Work hard

Start very early and do
little else

As usual, the middle path is best

Working Smart

Probability of meeting the deadline



95%

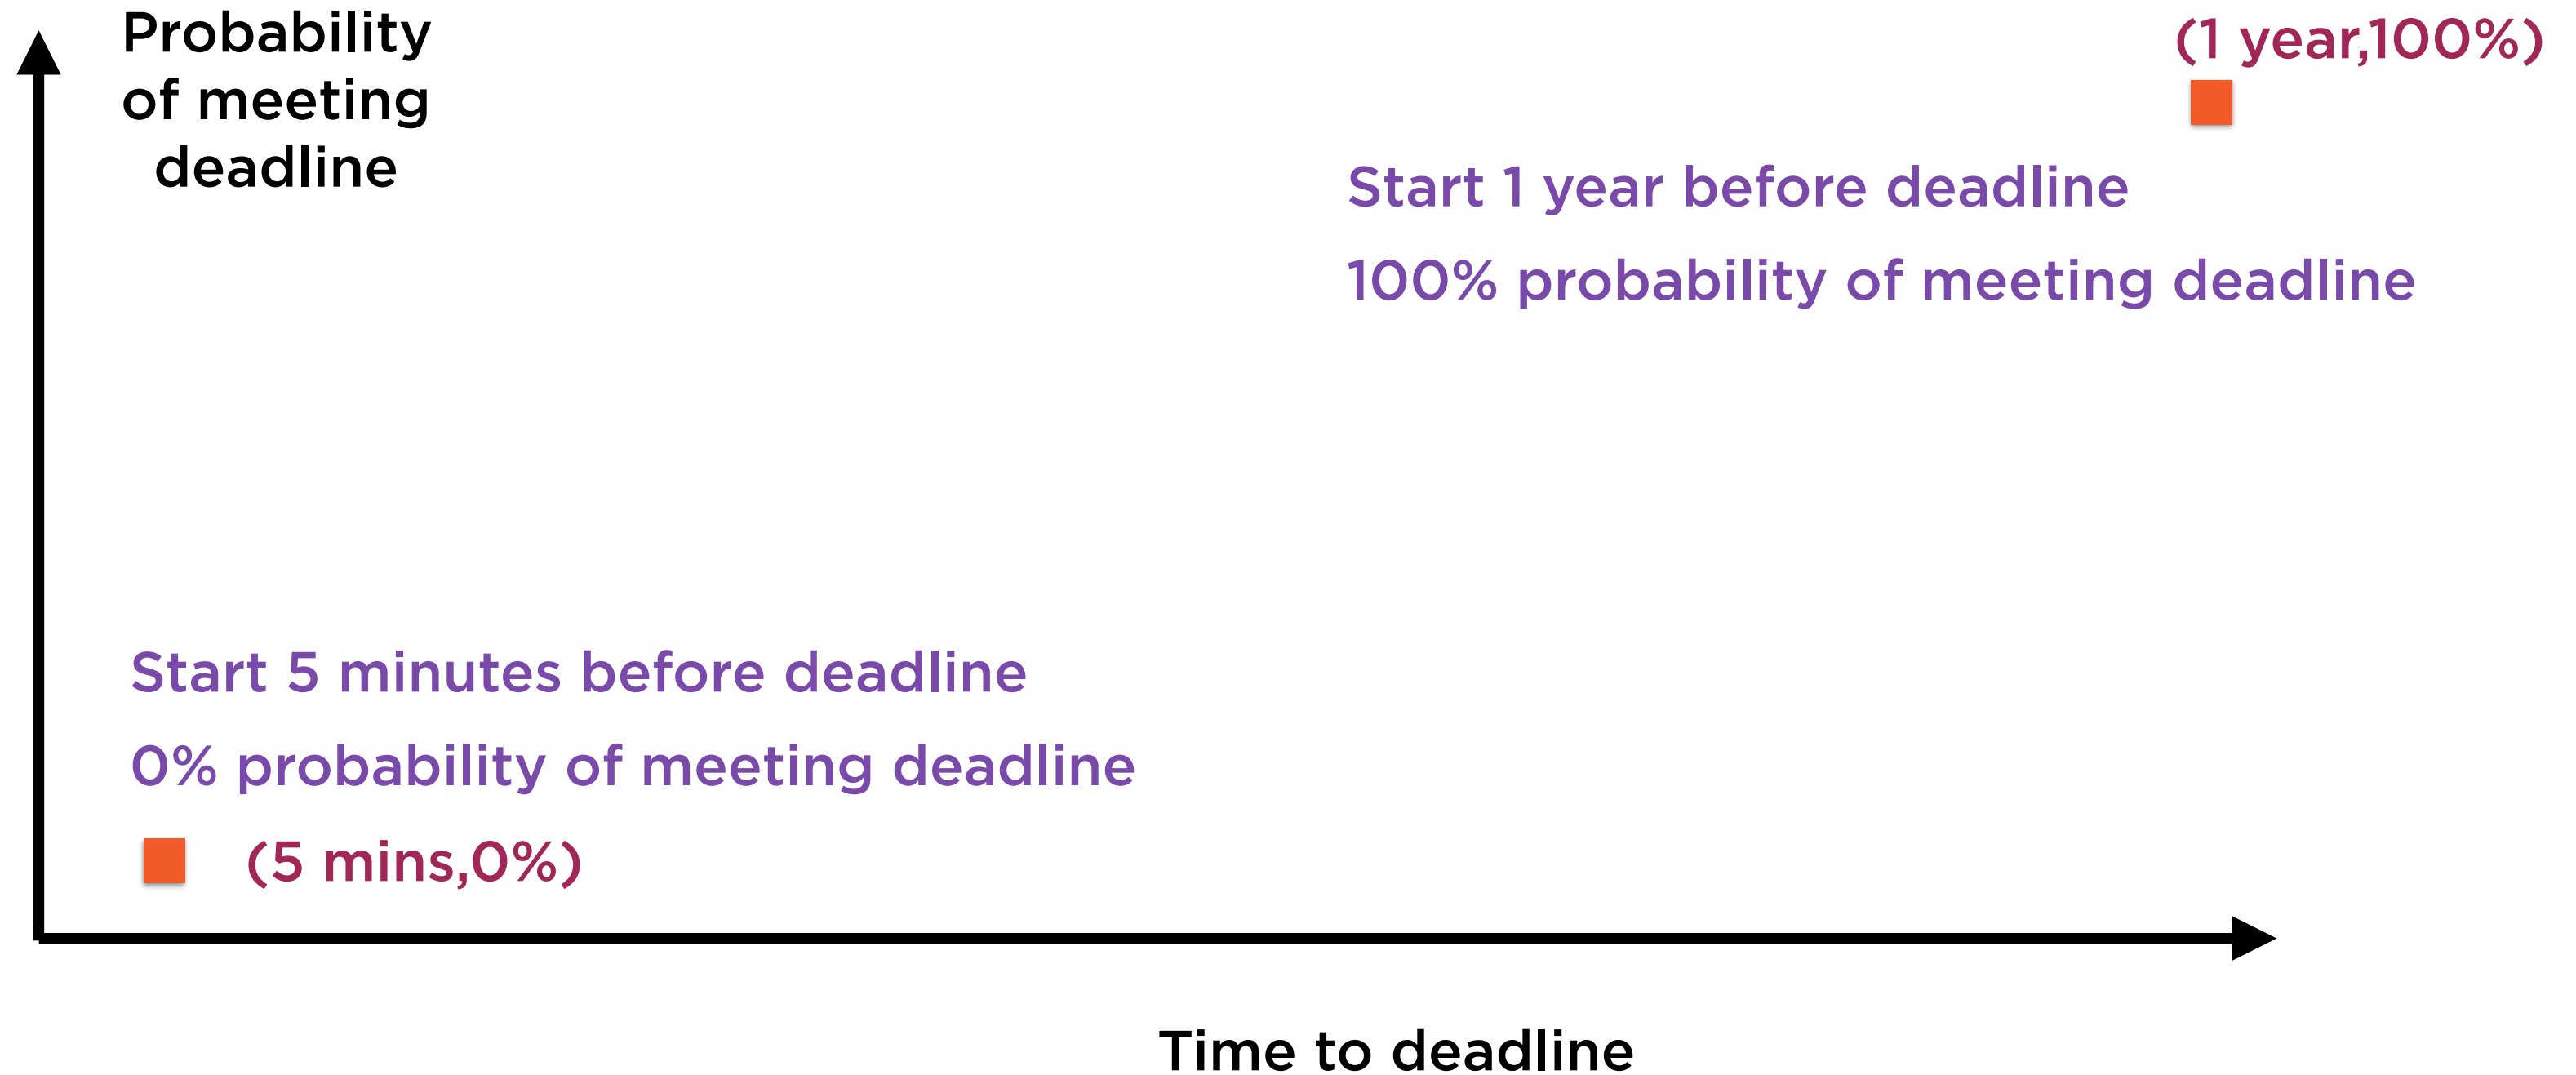


Probability of getting other important work done

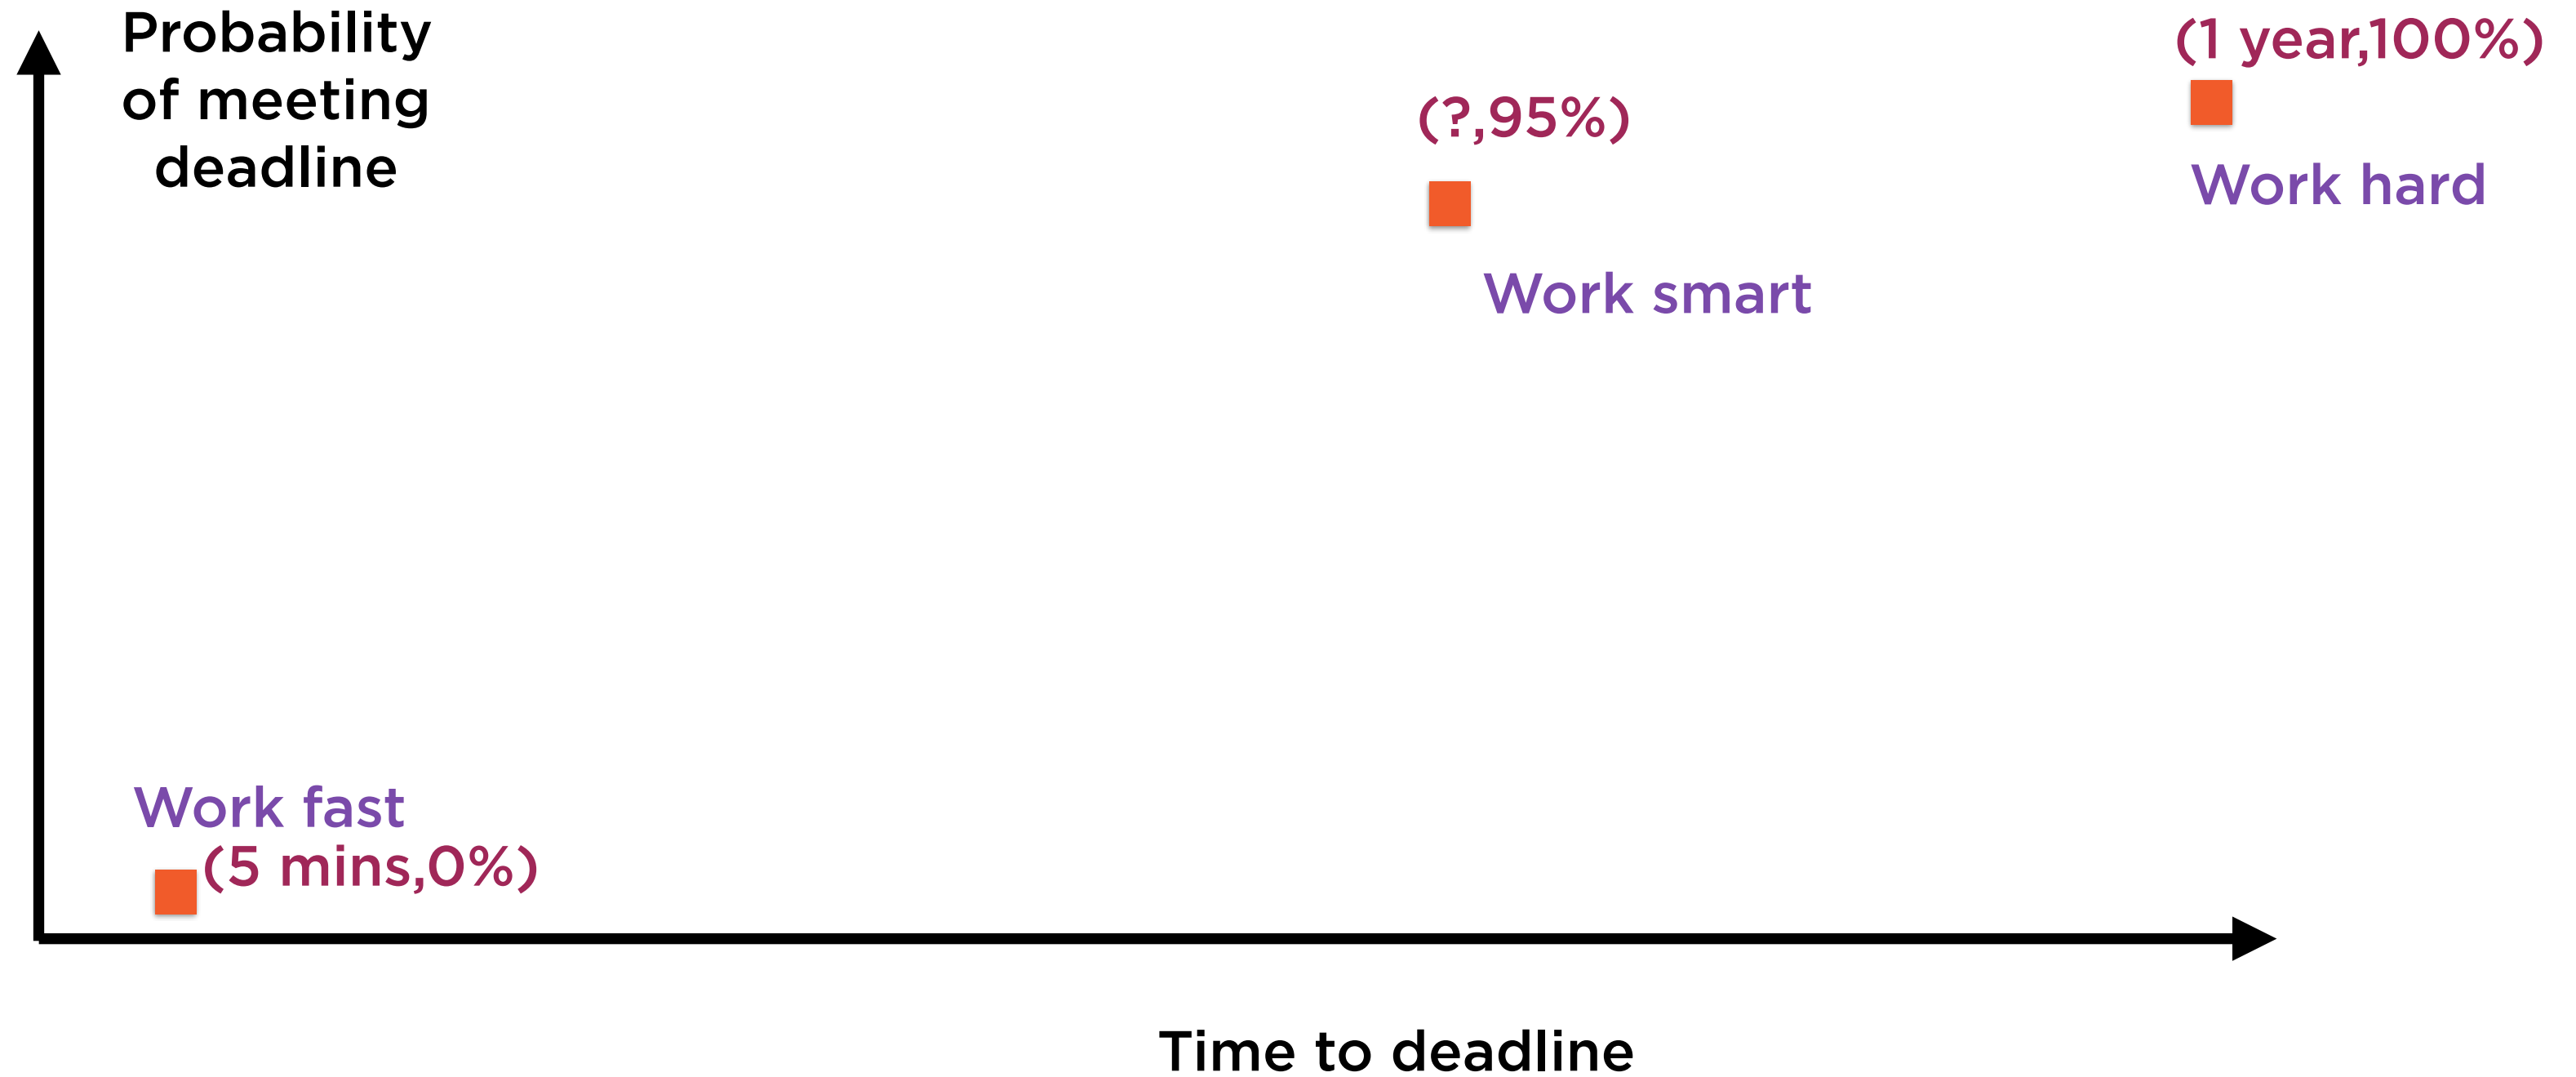


95%


Working Hard, Fast, Smart



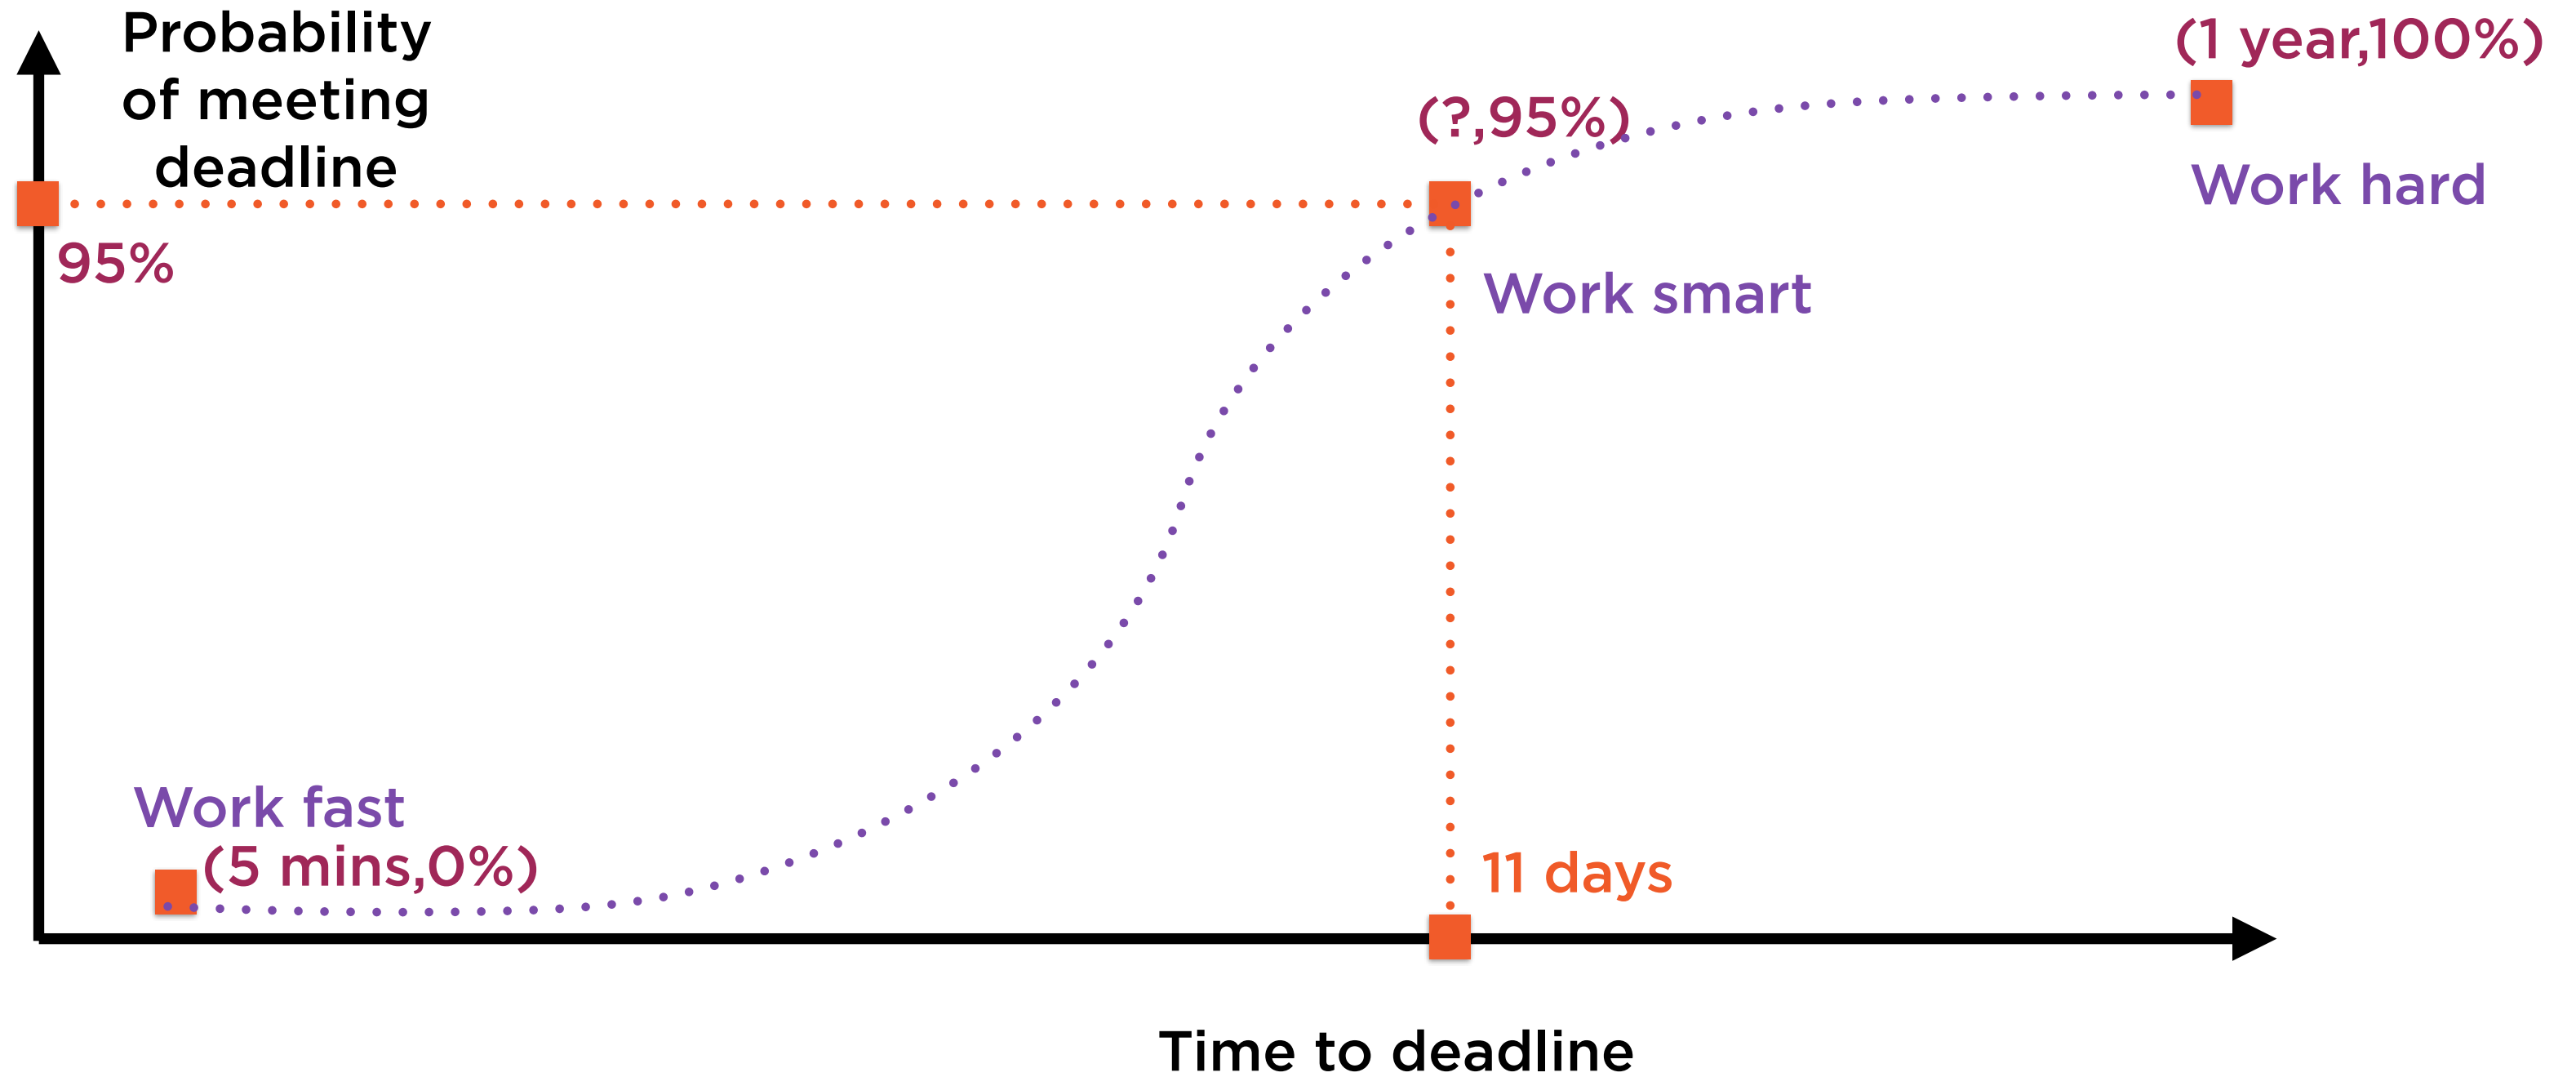
Working Hard, Fast, Smart




Working Hard, Fast, Smart




Working Hard, Fast, Smart



Working Hard, Fast, Smart




Working Hard, Fast, Smart

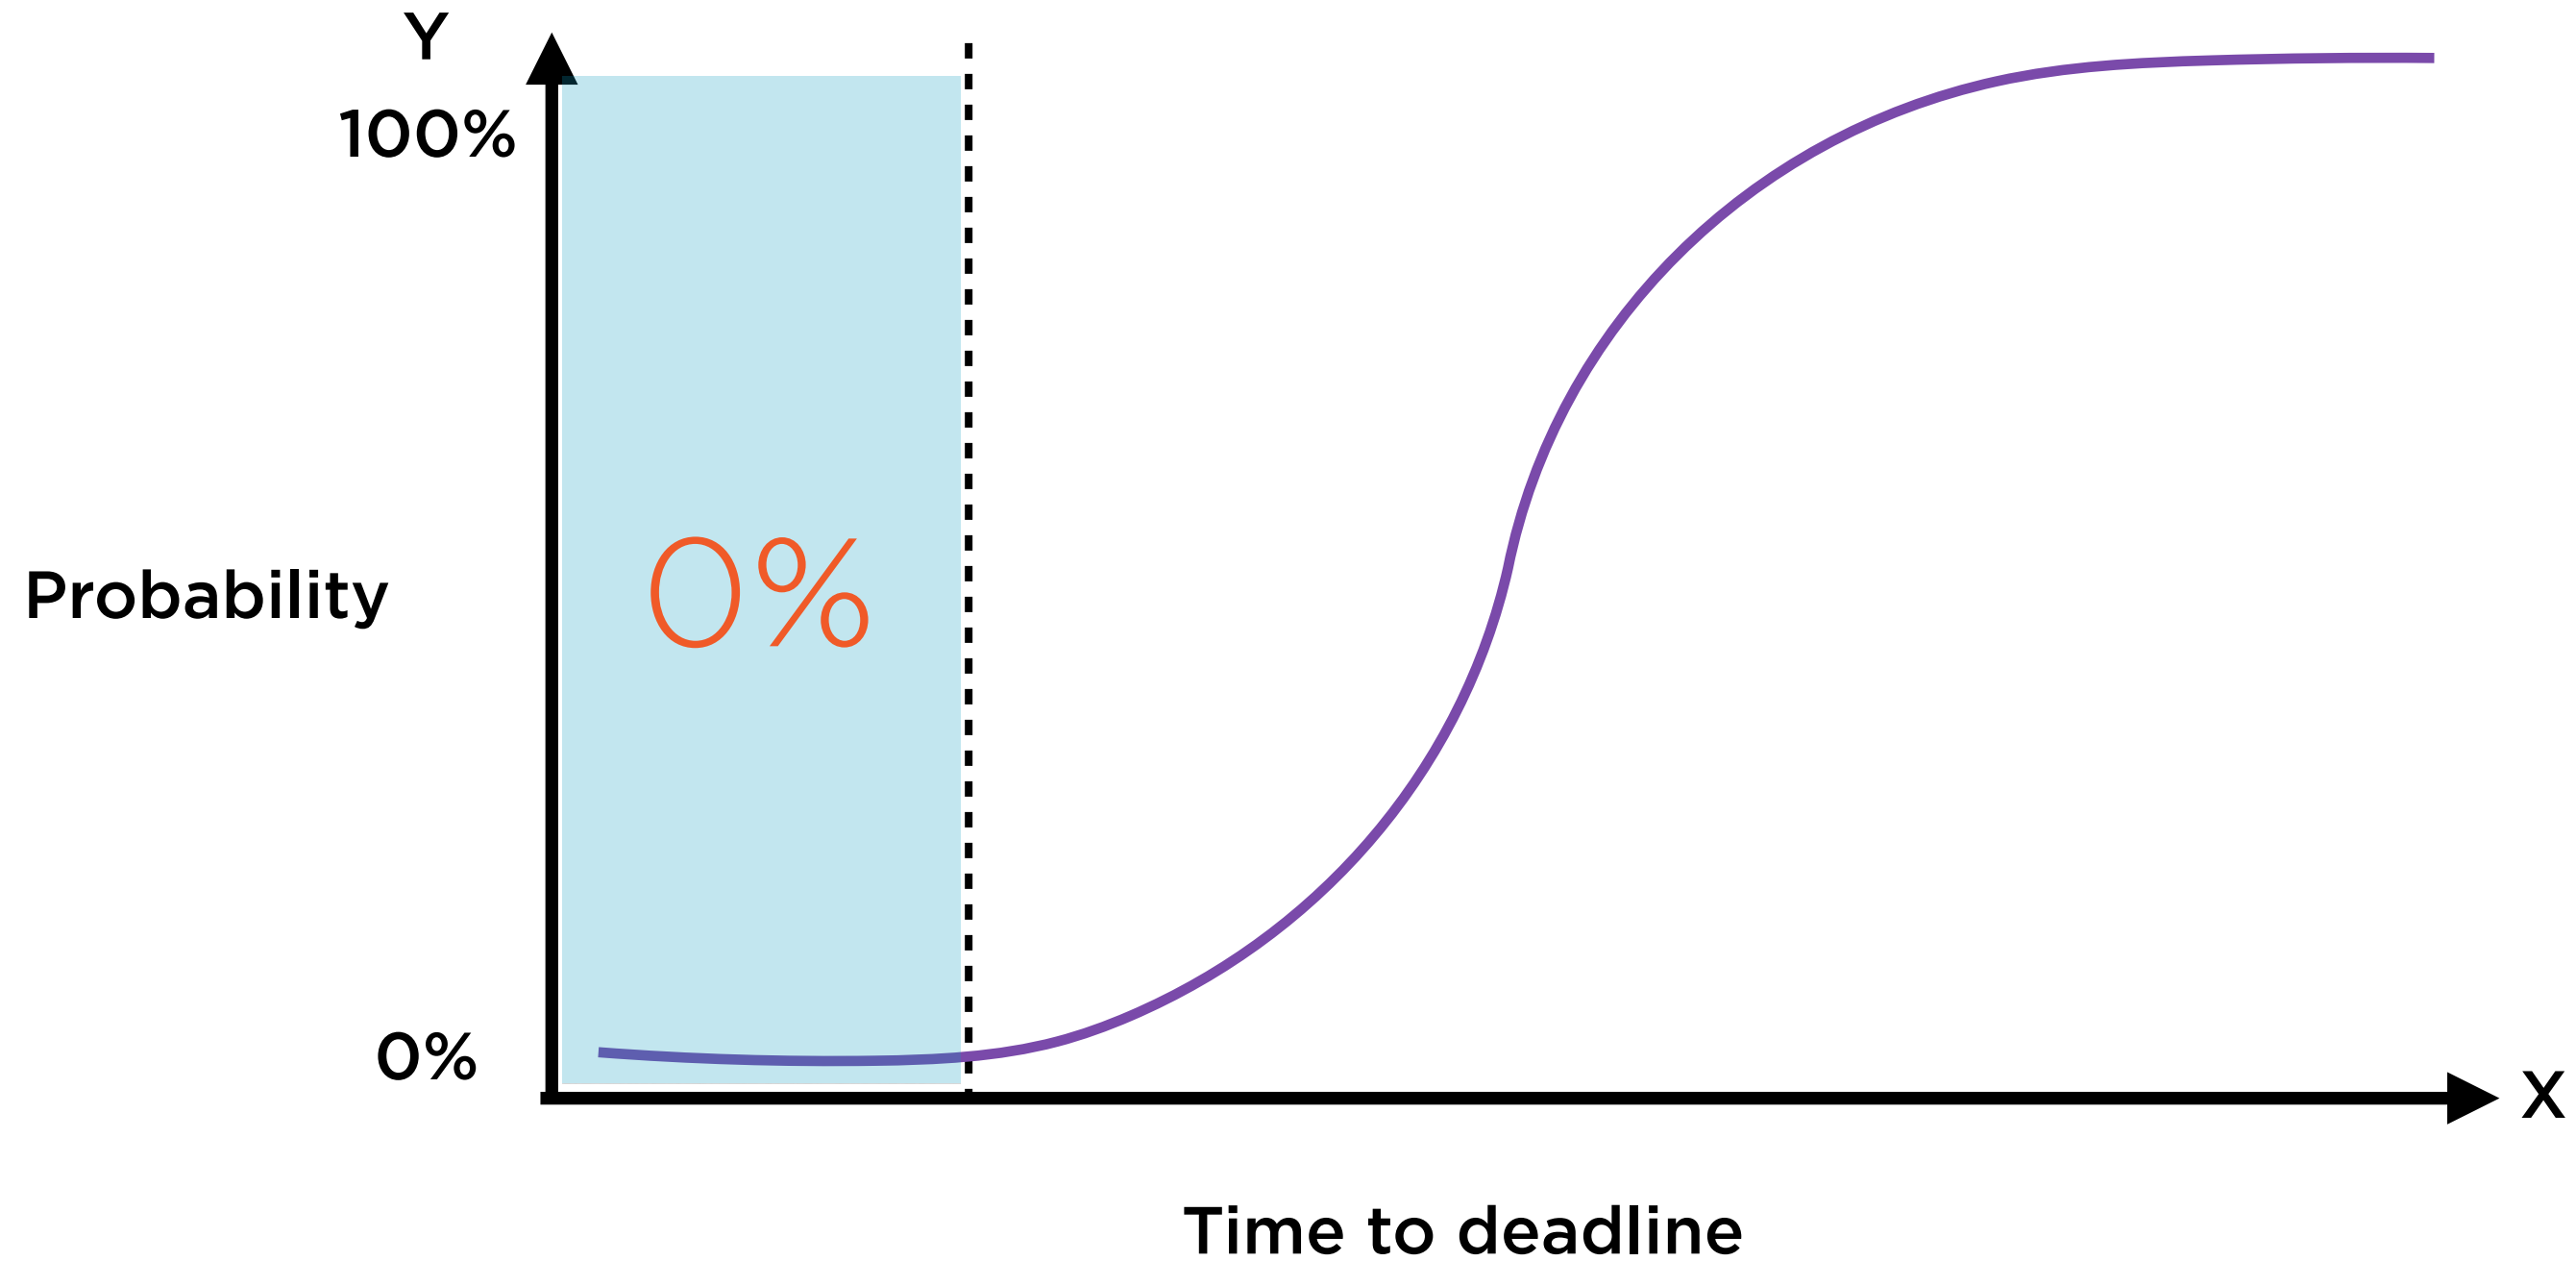


Logistic Regression helps find how probabilities are changed by actions

Working Smart with Logistic Regression




Working Smart with Logistic Regression




Start too late, and you'll definitely miss

Working Smart with Logistic Regression




Start too early, and you'll definitely make it

Working Smart with Logistic Regression



Working smart is knowing when to start




Y-axis: probability of meeting deadline

X-axis: time to deadline

Meeting or missing deadline is binary

Probability curve flattens at ends

- floor of 0
- ceiling of 1



y: hit or miss? (0 or 1?)

x: start time before deadline

$p(y)$: probability of $y = 1$

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$

Logistic regression involves finding the “best fit” such curve

- A is the intercept
- B is the regression coefficient


(e is the constant 2.71828)

S-curves are widely studied, well understood


$$y = \frac{1}{1 + e^{-(A+Bx)}}$$

Logistic regression uses S-curve to estimate probabilities

$$p(y) = \frac{1}{1 + e^{-(A+Bx)}}$$




$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$



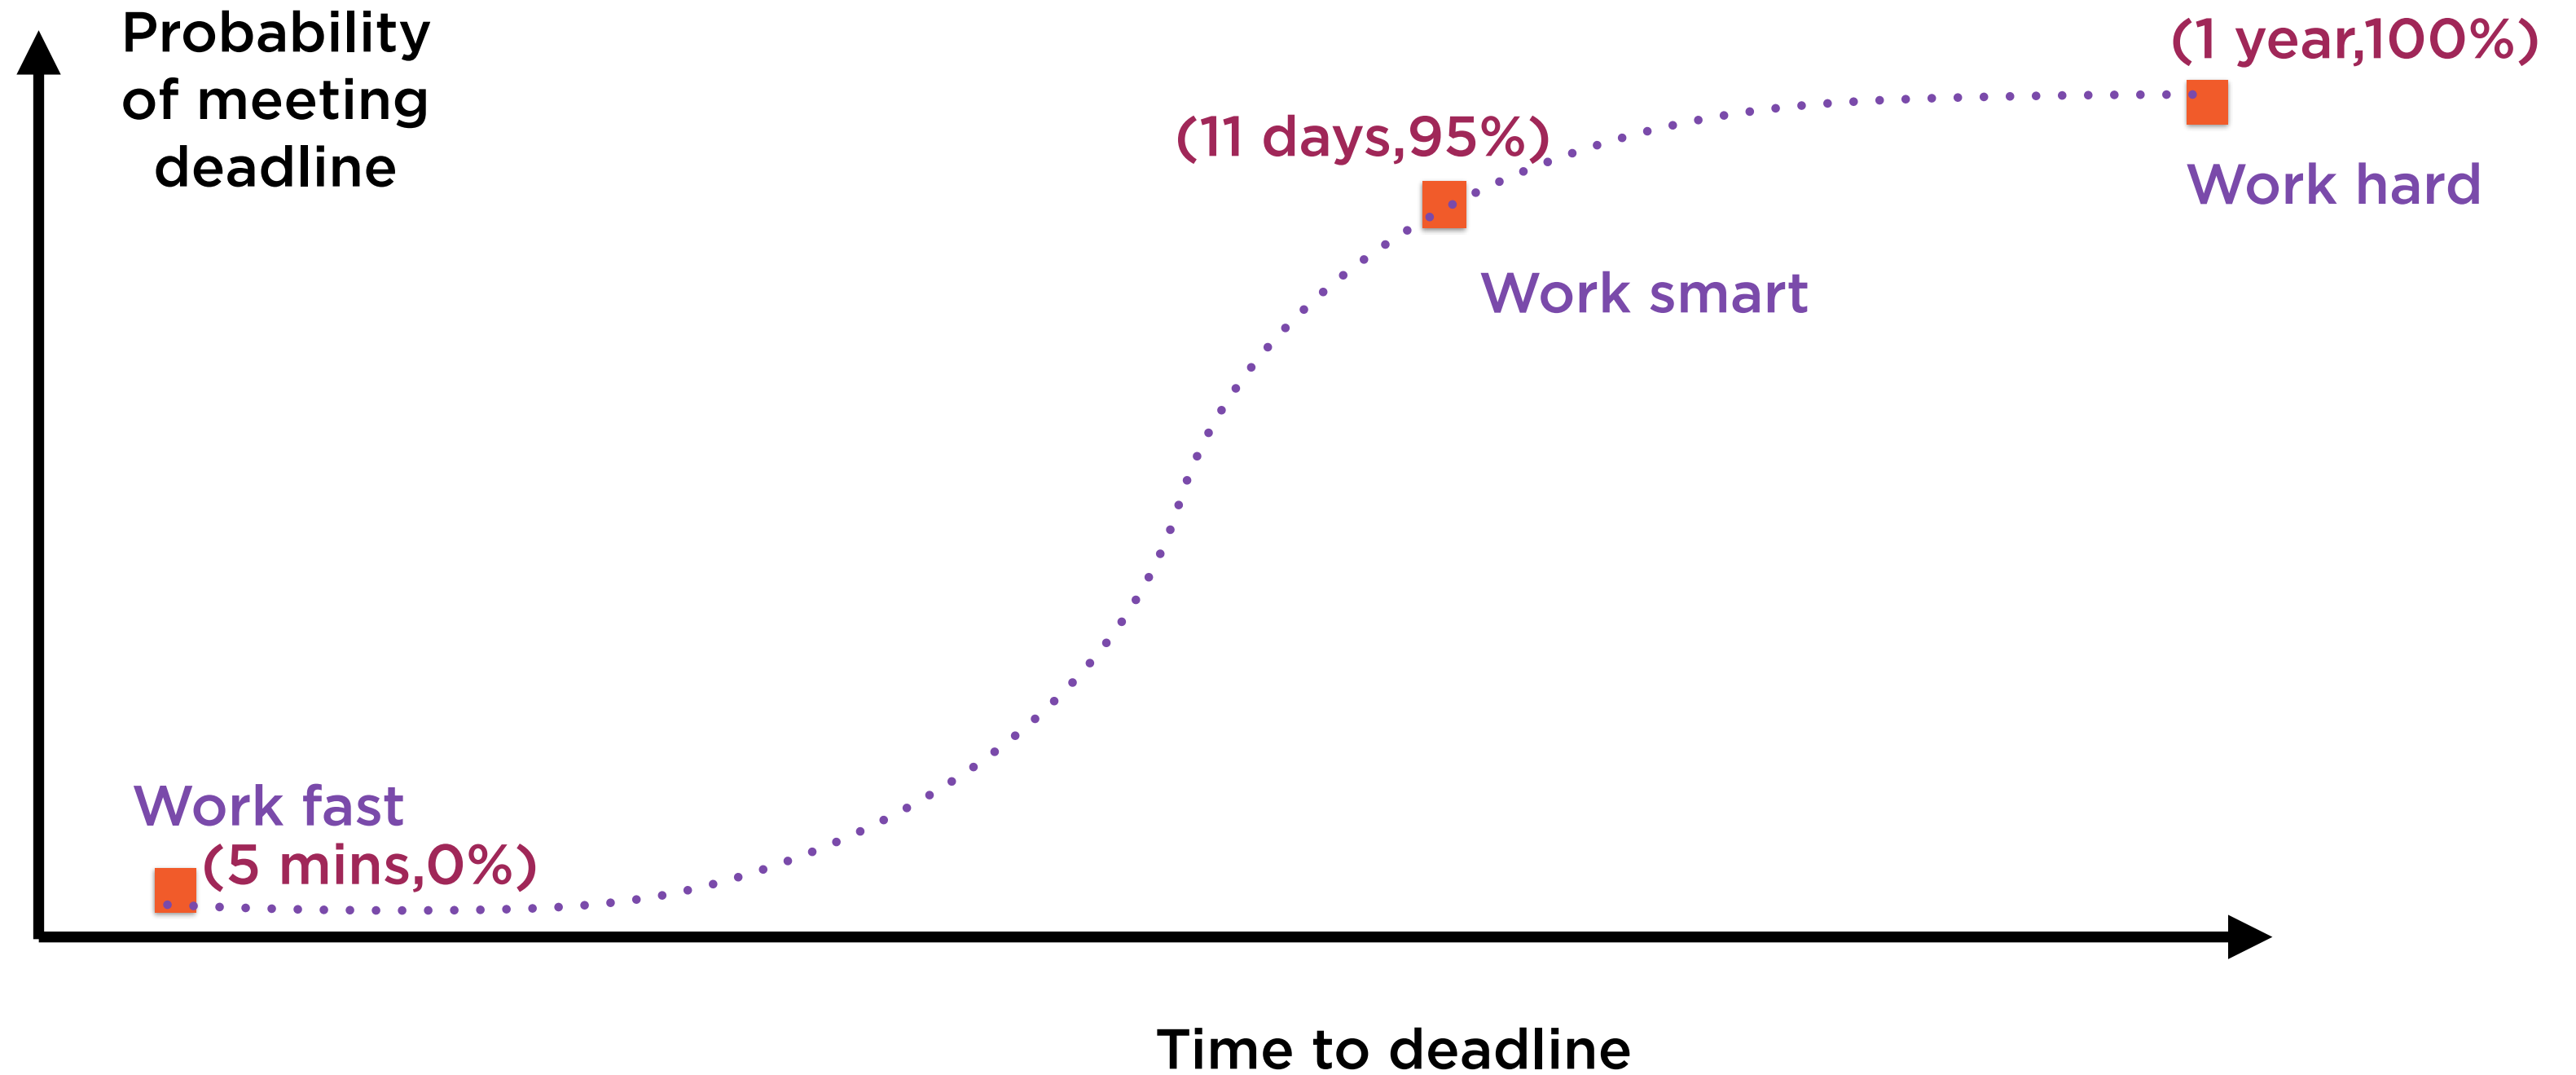
If A and B are **positive**

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$




If A and B are **negative**

Working Hard, Fast, Smart




Working Hard, Fast, Smart




Minimum value of $p(y_i)$

Working Hard, Fast, Smart




Maximum value of $p(y_i)$

Working Hard, Fast, Smart



Between maximum and minimum values of $p(y_i)$

Logistic Regression



$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$

Categorical and Continuous Variables

Continuous

Can take an infinite set of values
(height, weight, income...)

Categorical


Can take a finite set of values
(male/female, day of week...)

Categorical variables that can take just two values are called **binary variables**

Logistic Regression helps estimate
how **probabilities** of **categorical
variables** are influenced by **causes**


Logistic Regression in Classification

Whales: Fish or Mammals



Mammal

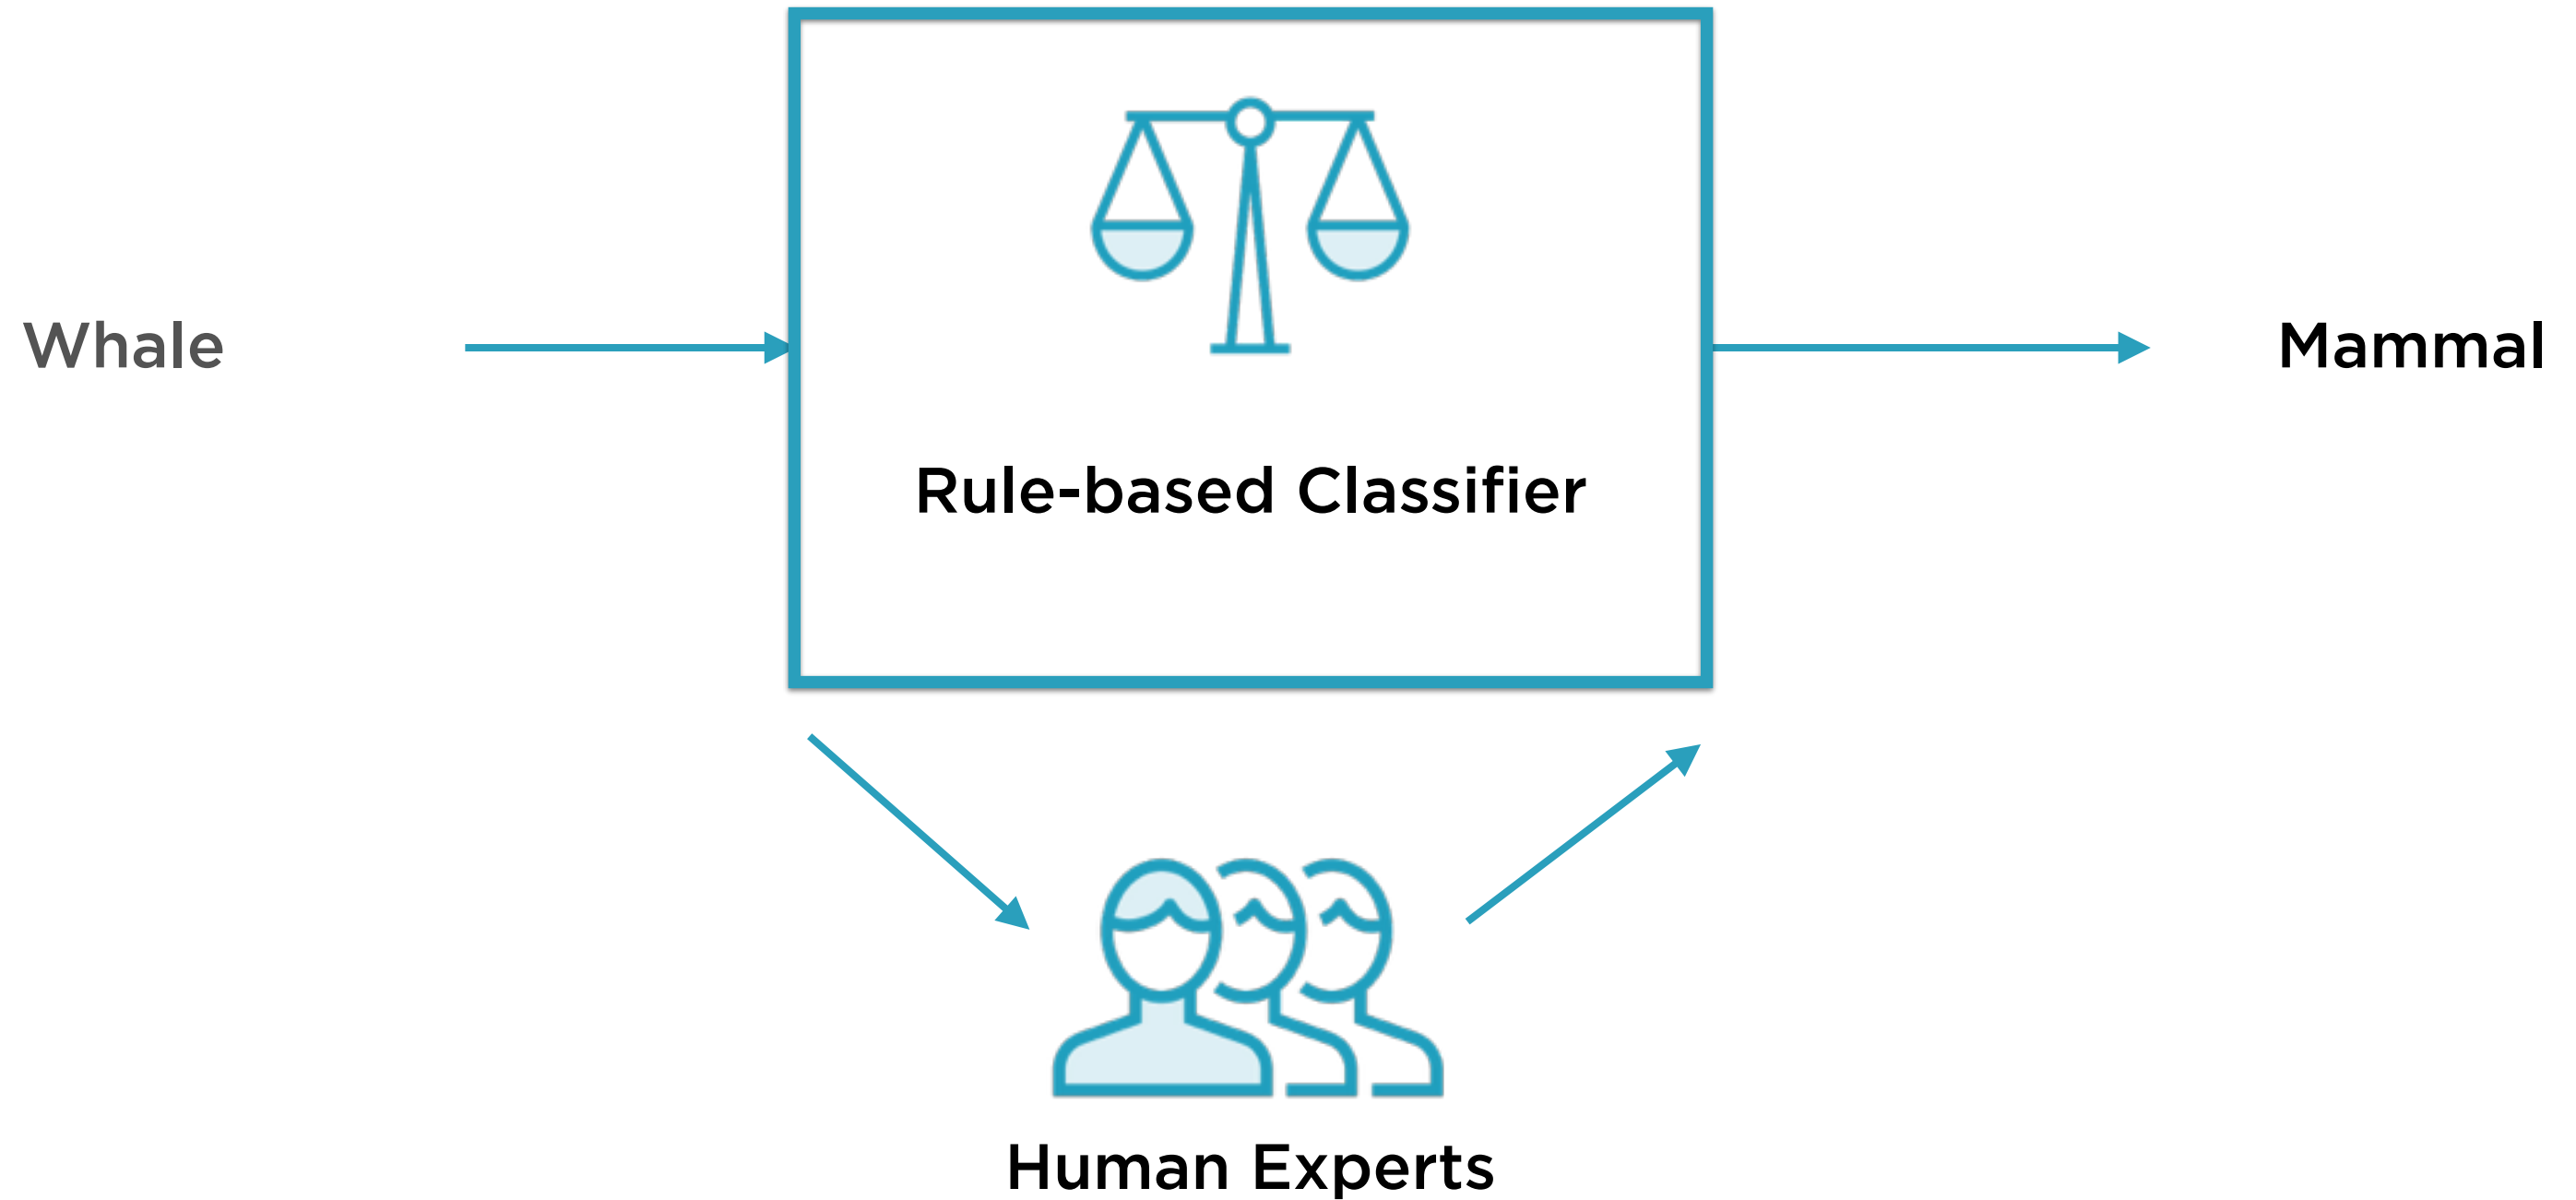
Member of the infraorder
Cetacea



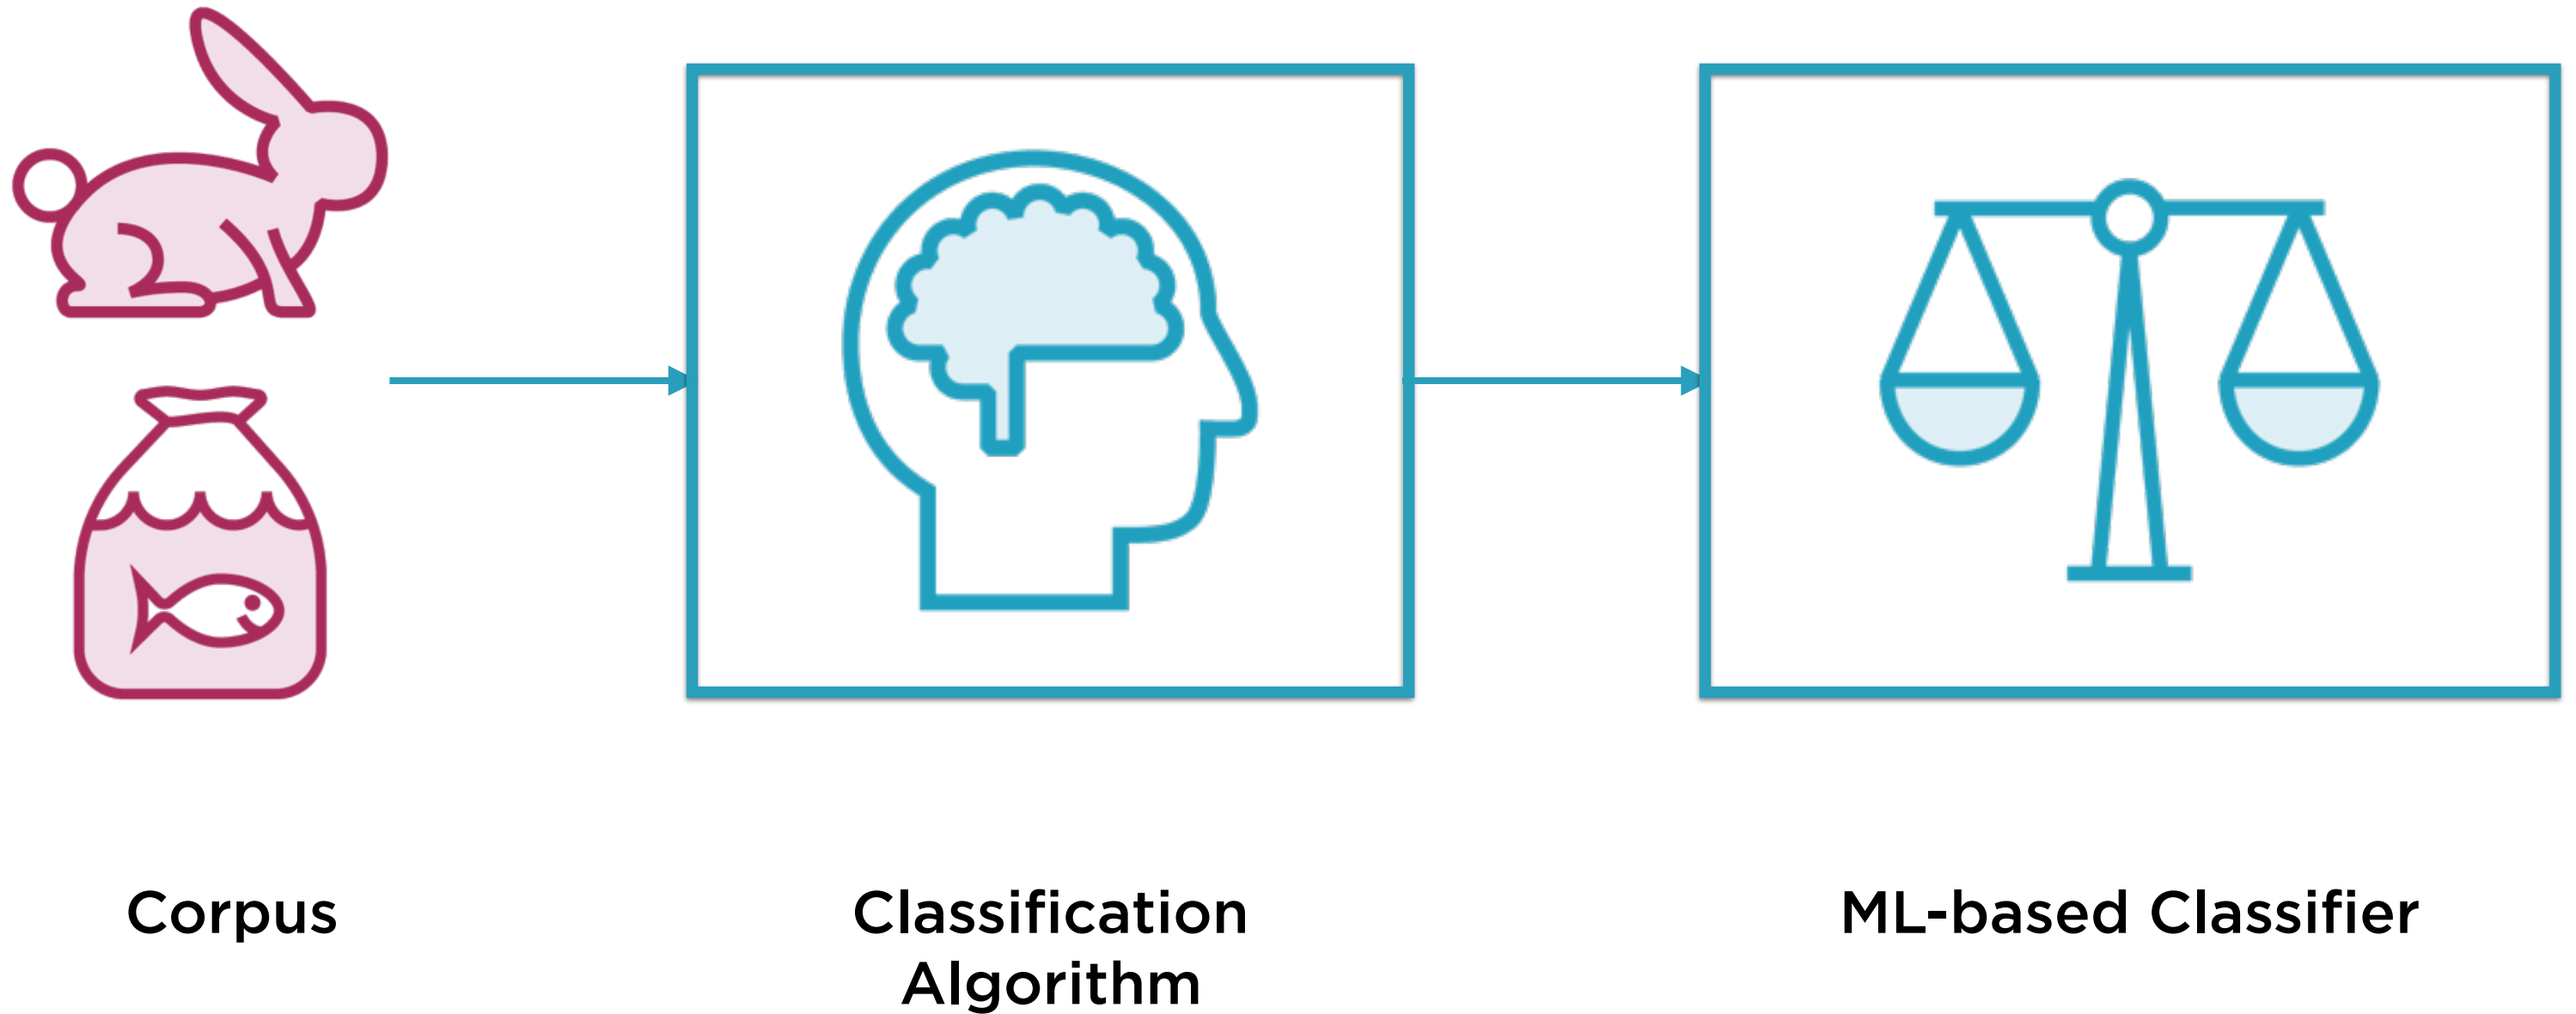
Fish

Looks like a fish, swims like a
fish, moves like a fish

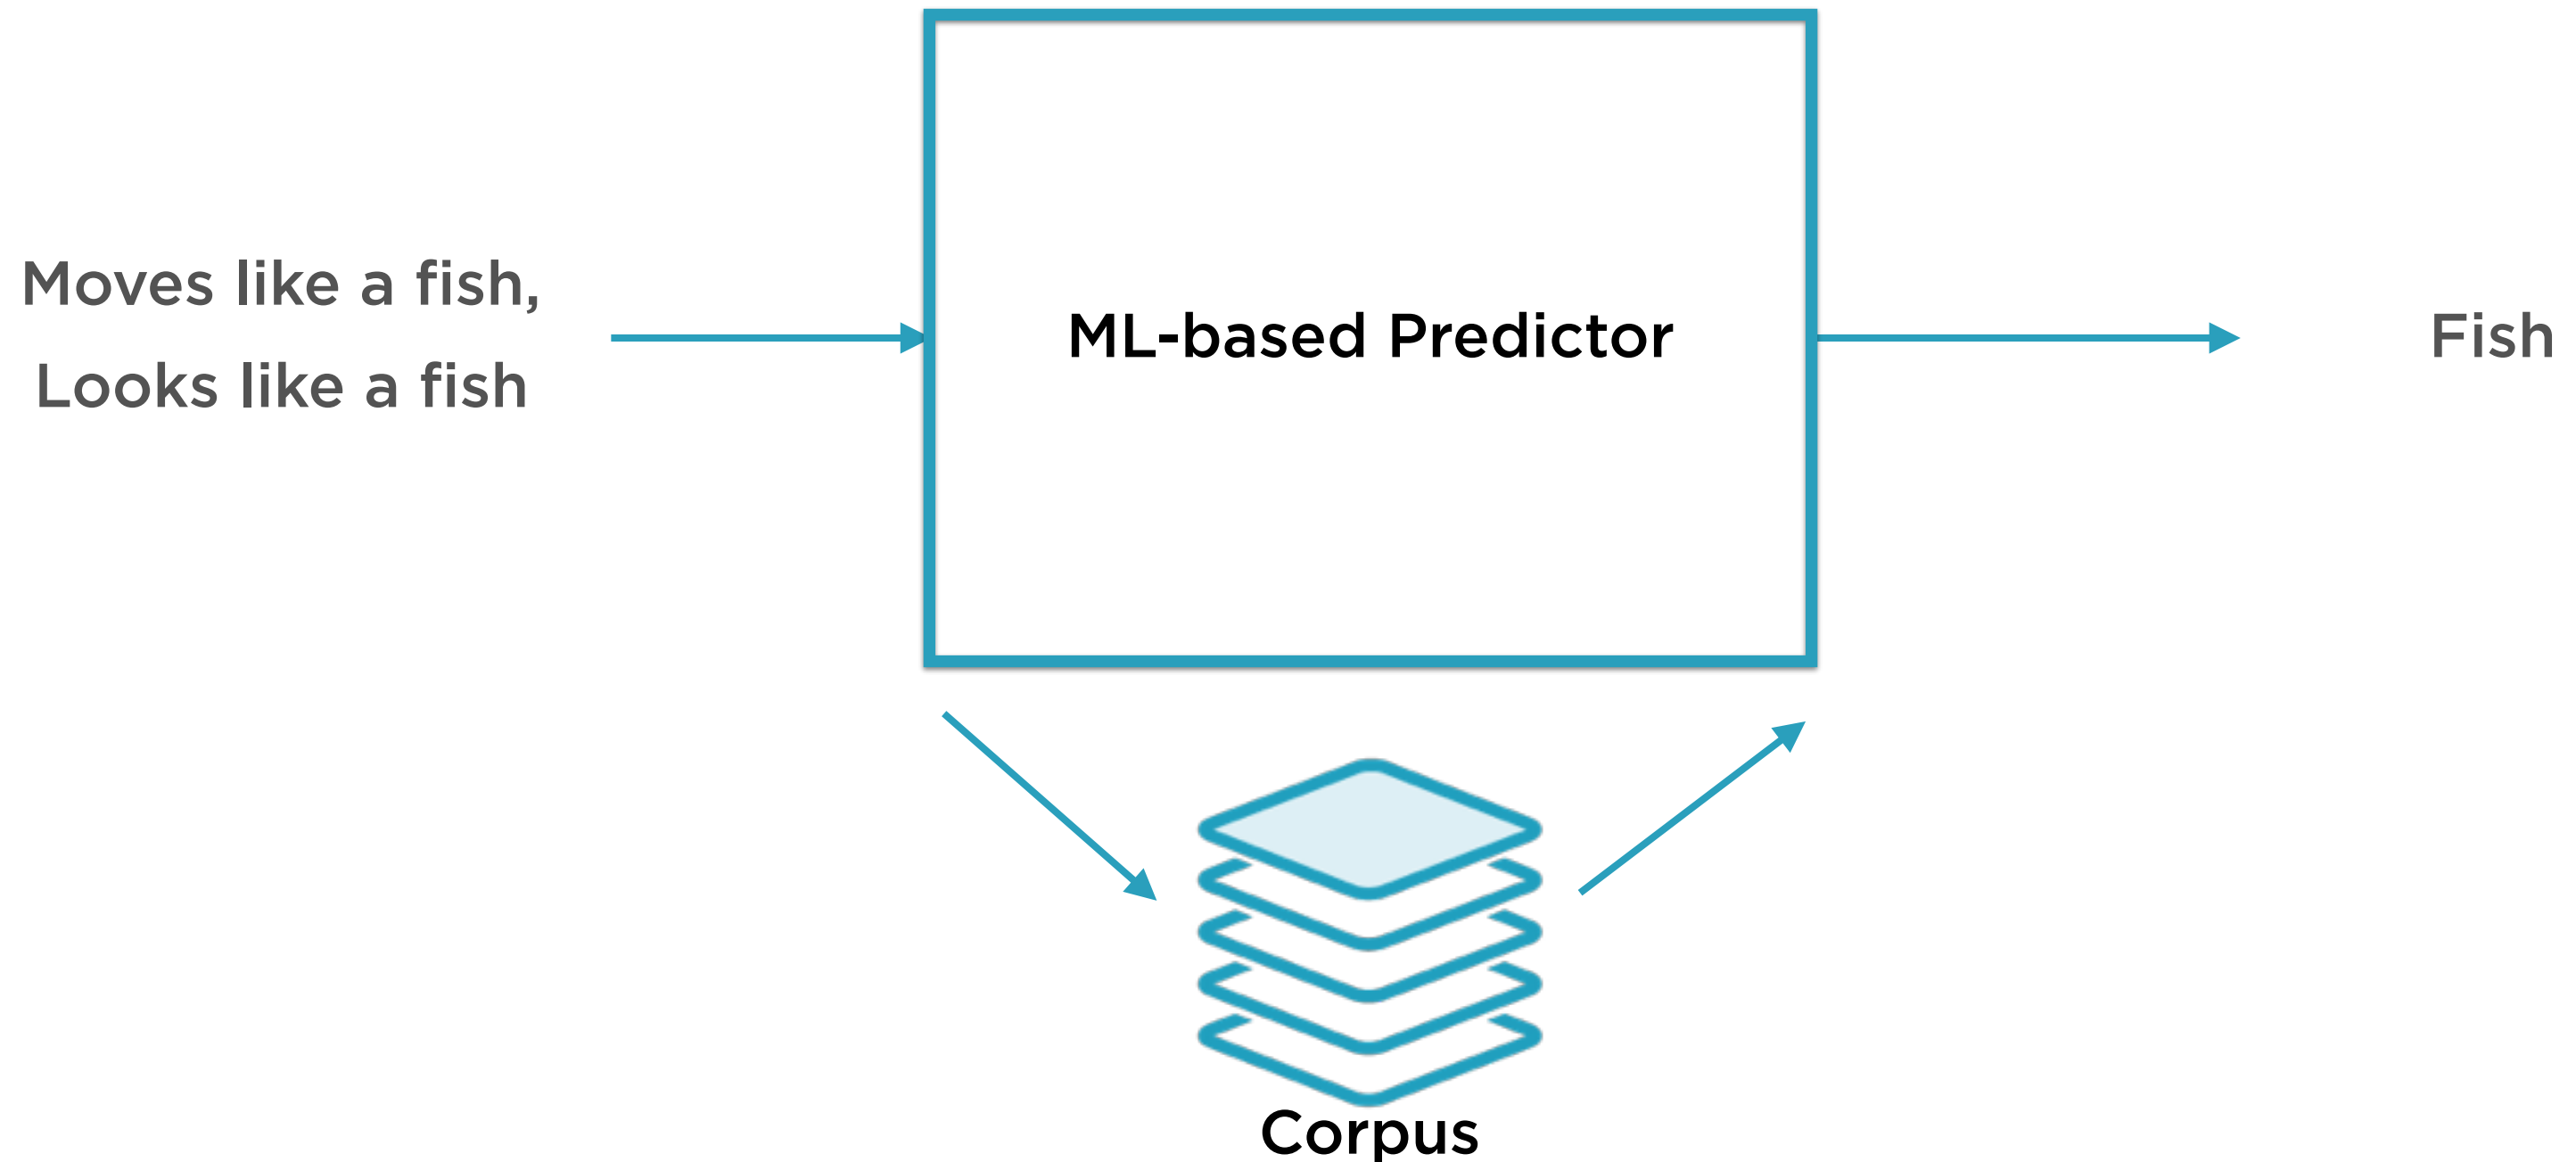
Rule-based Binary Classifier



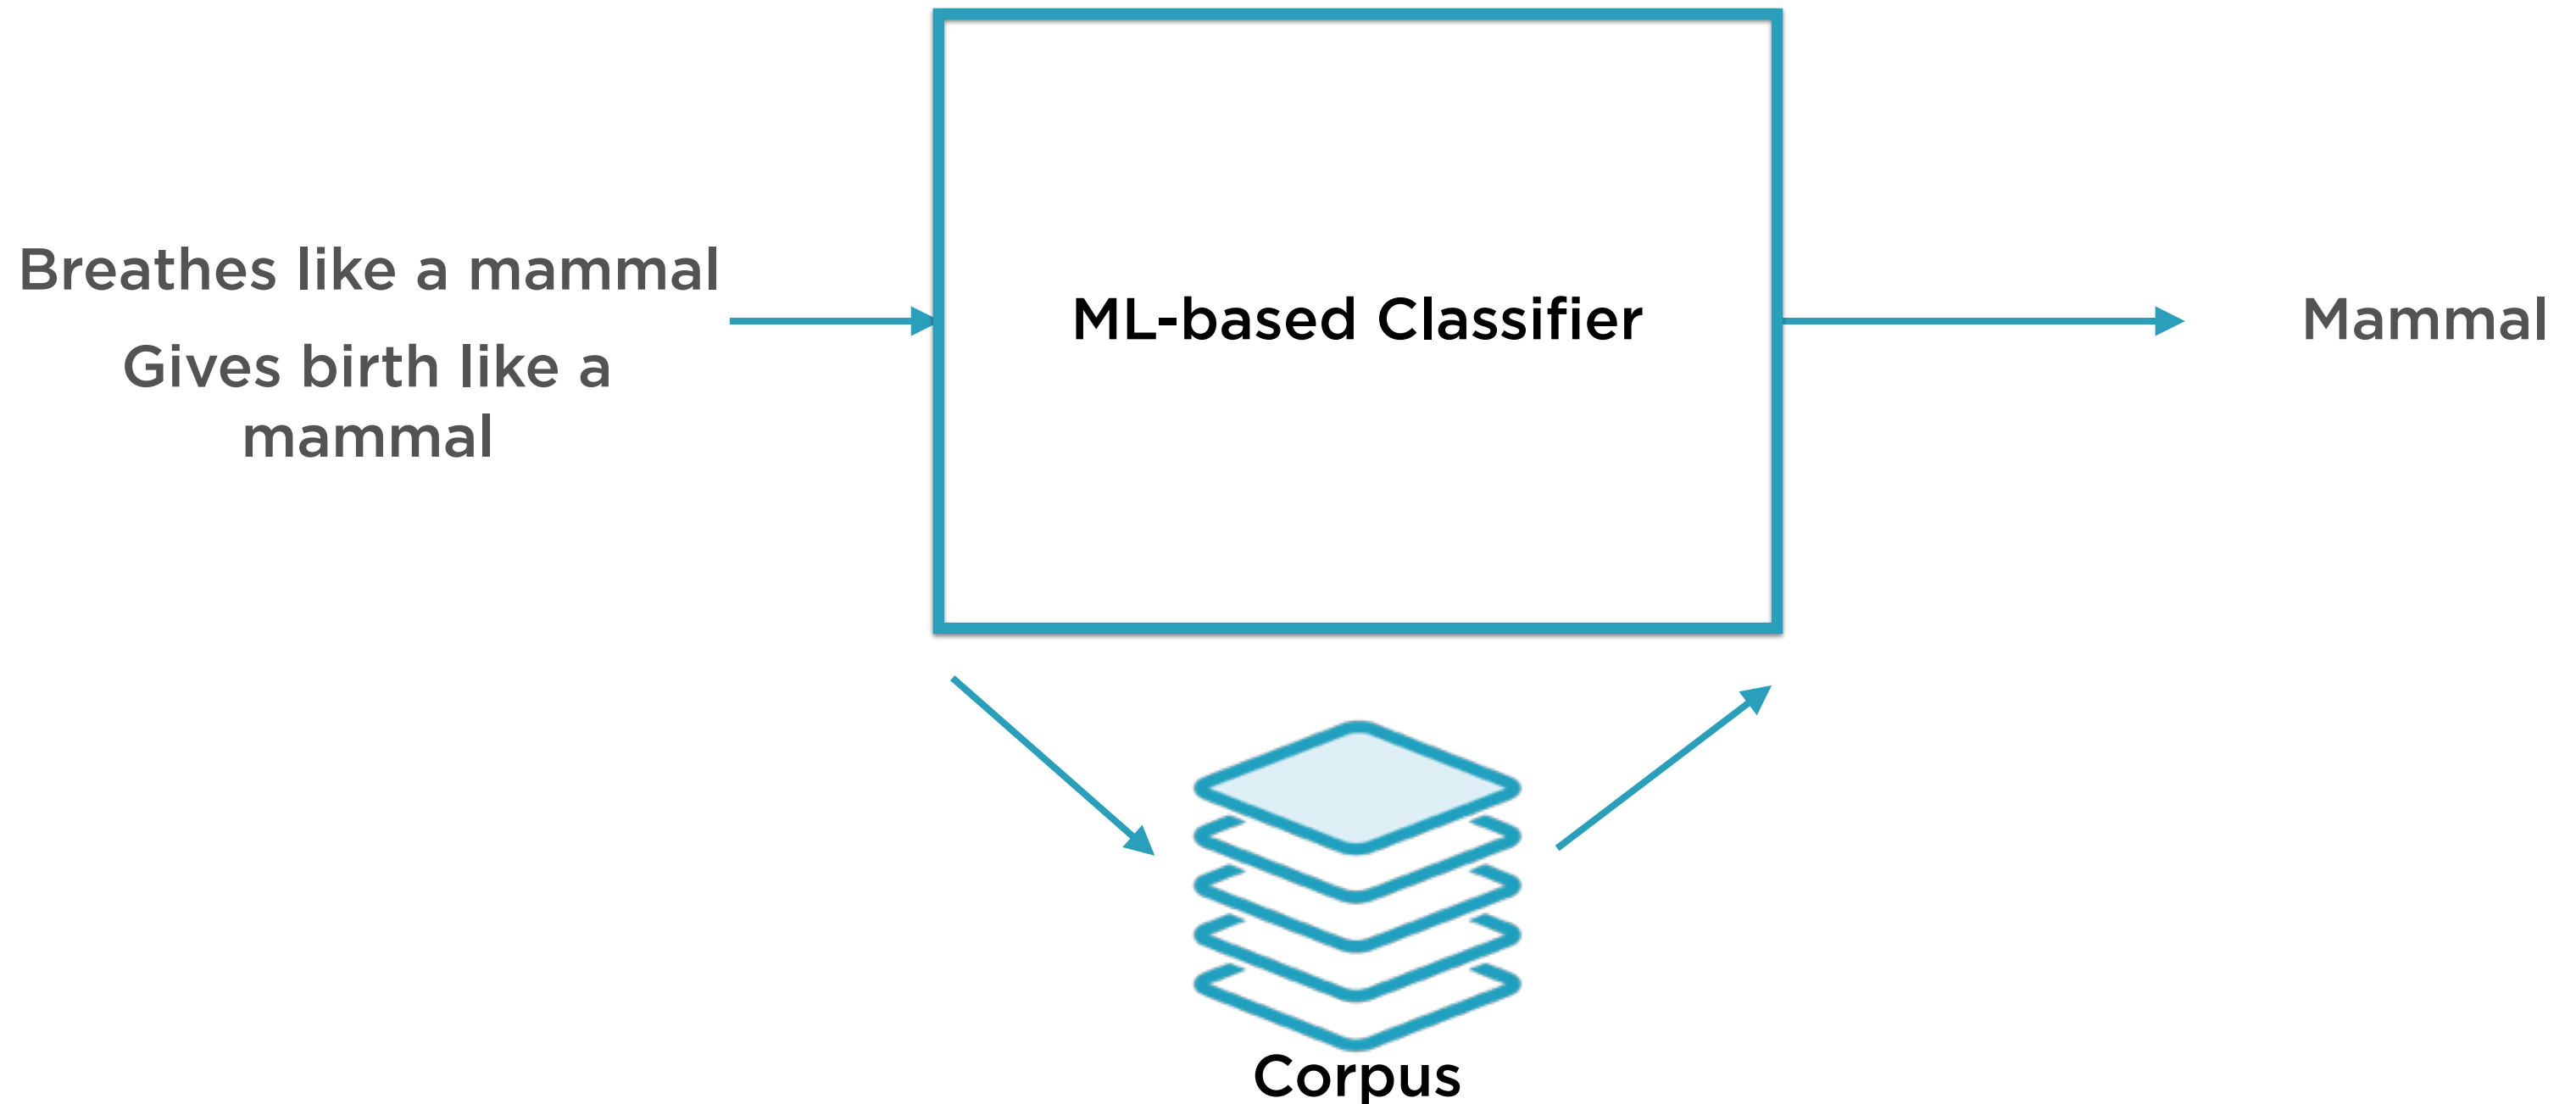
ML-based Binary Classifier



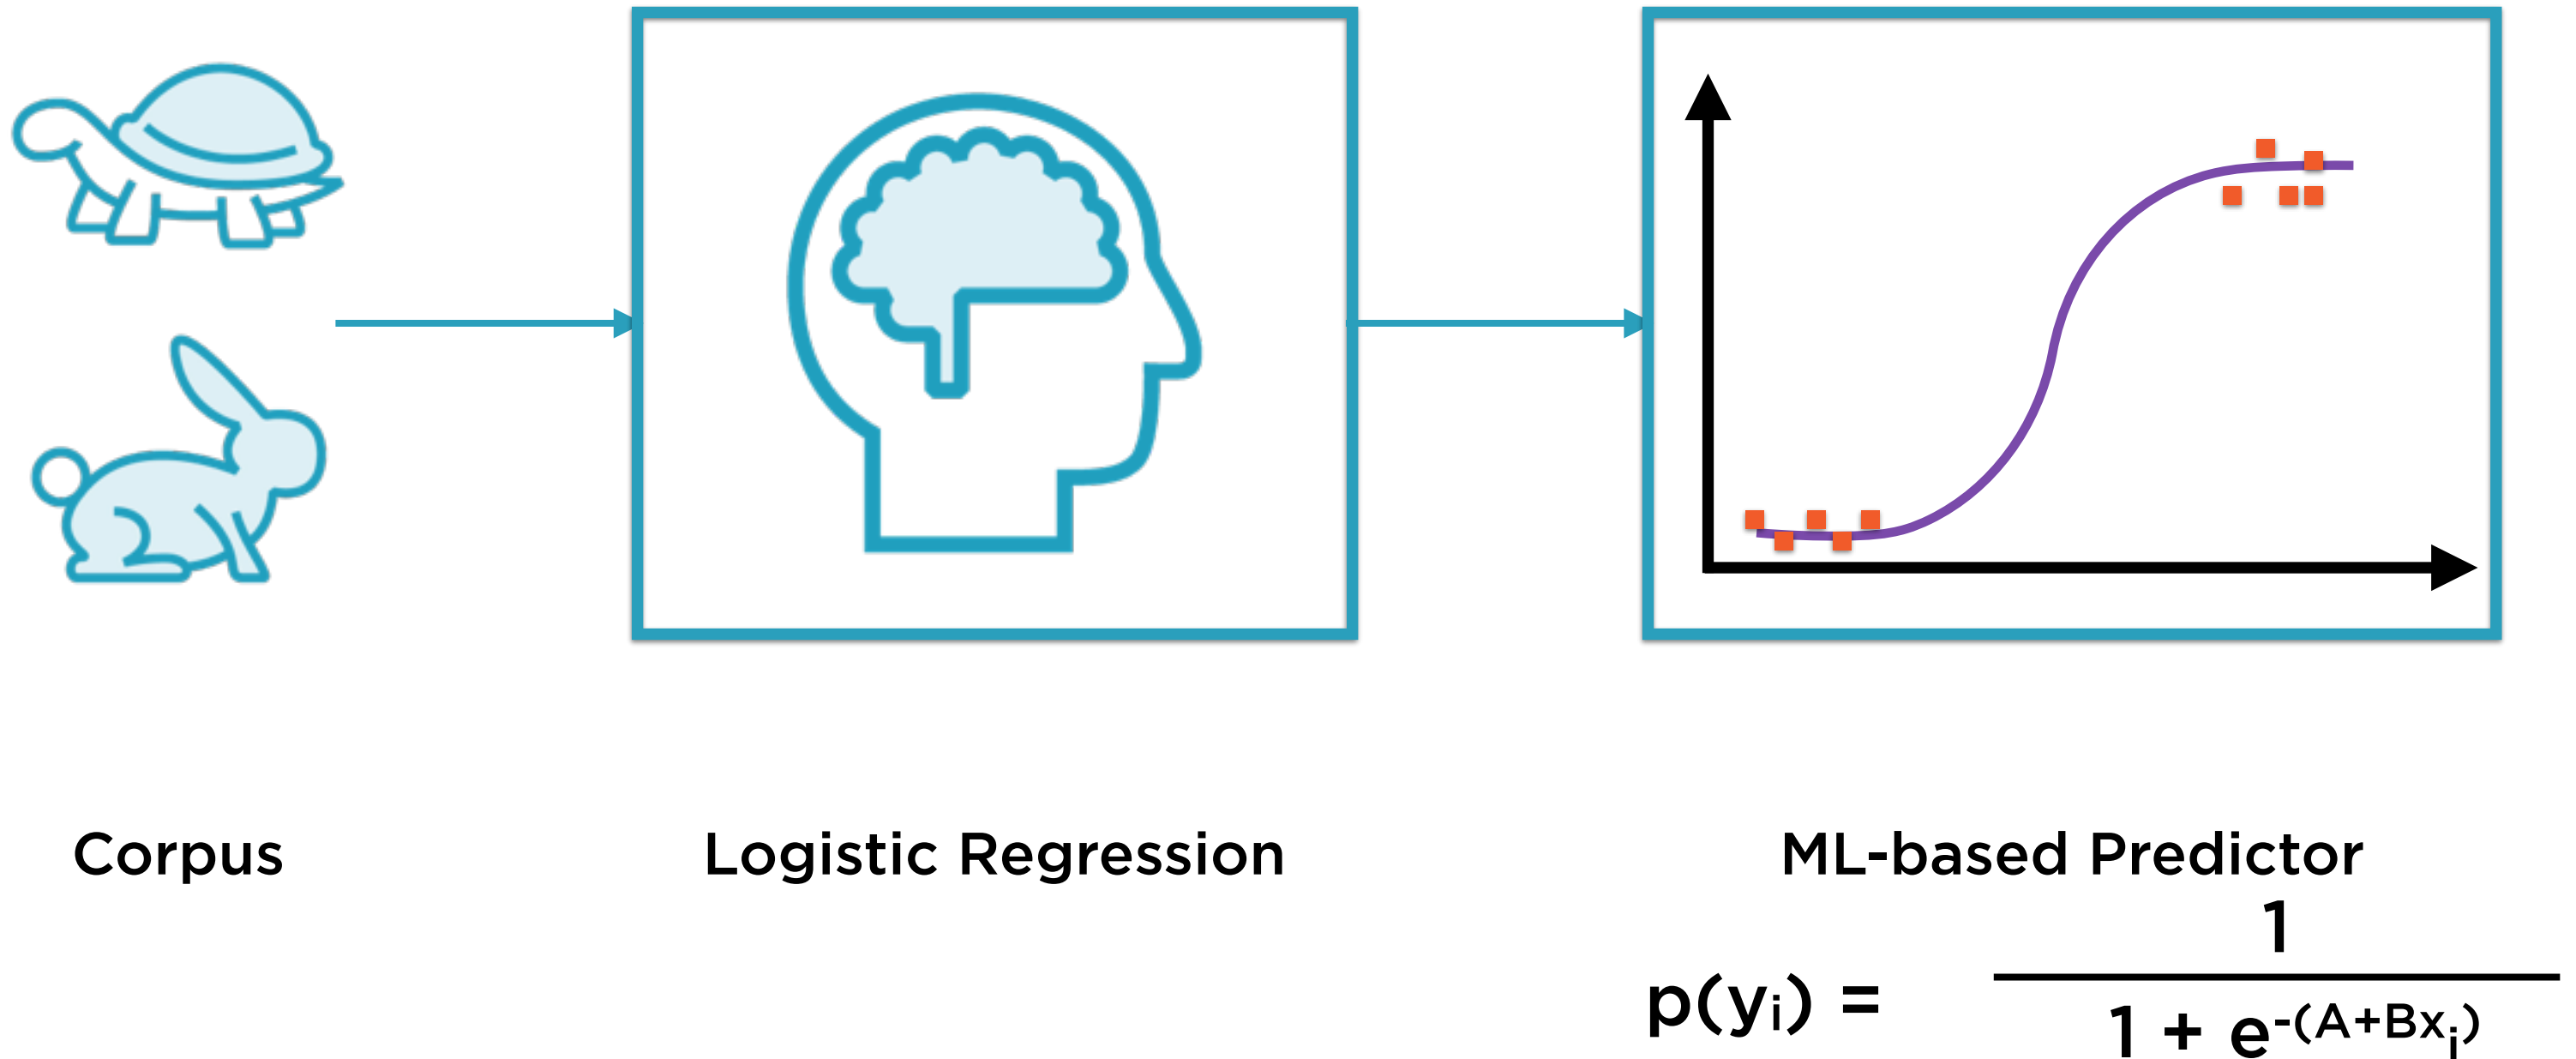
ML-based Binary Classifier



ML-based Binary Classifier




ML-based Predictor



ML-based Predictor

Lives in water,
breathes with
lungs, does not lay
eggs

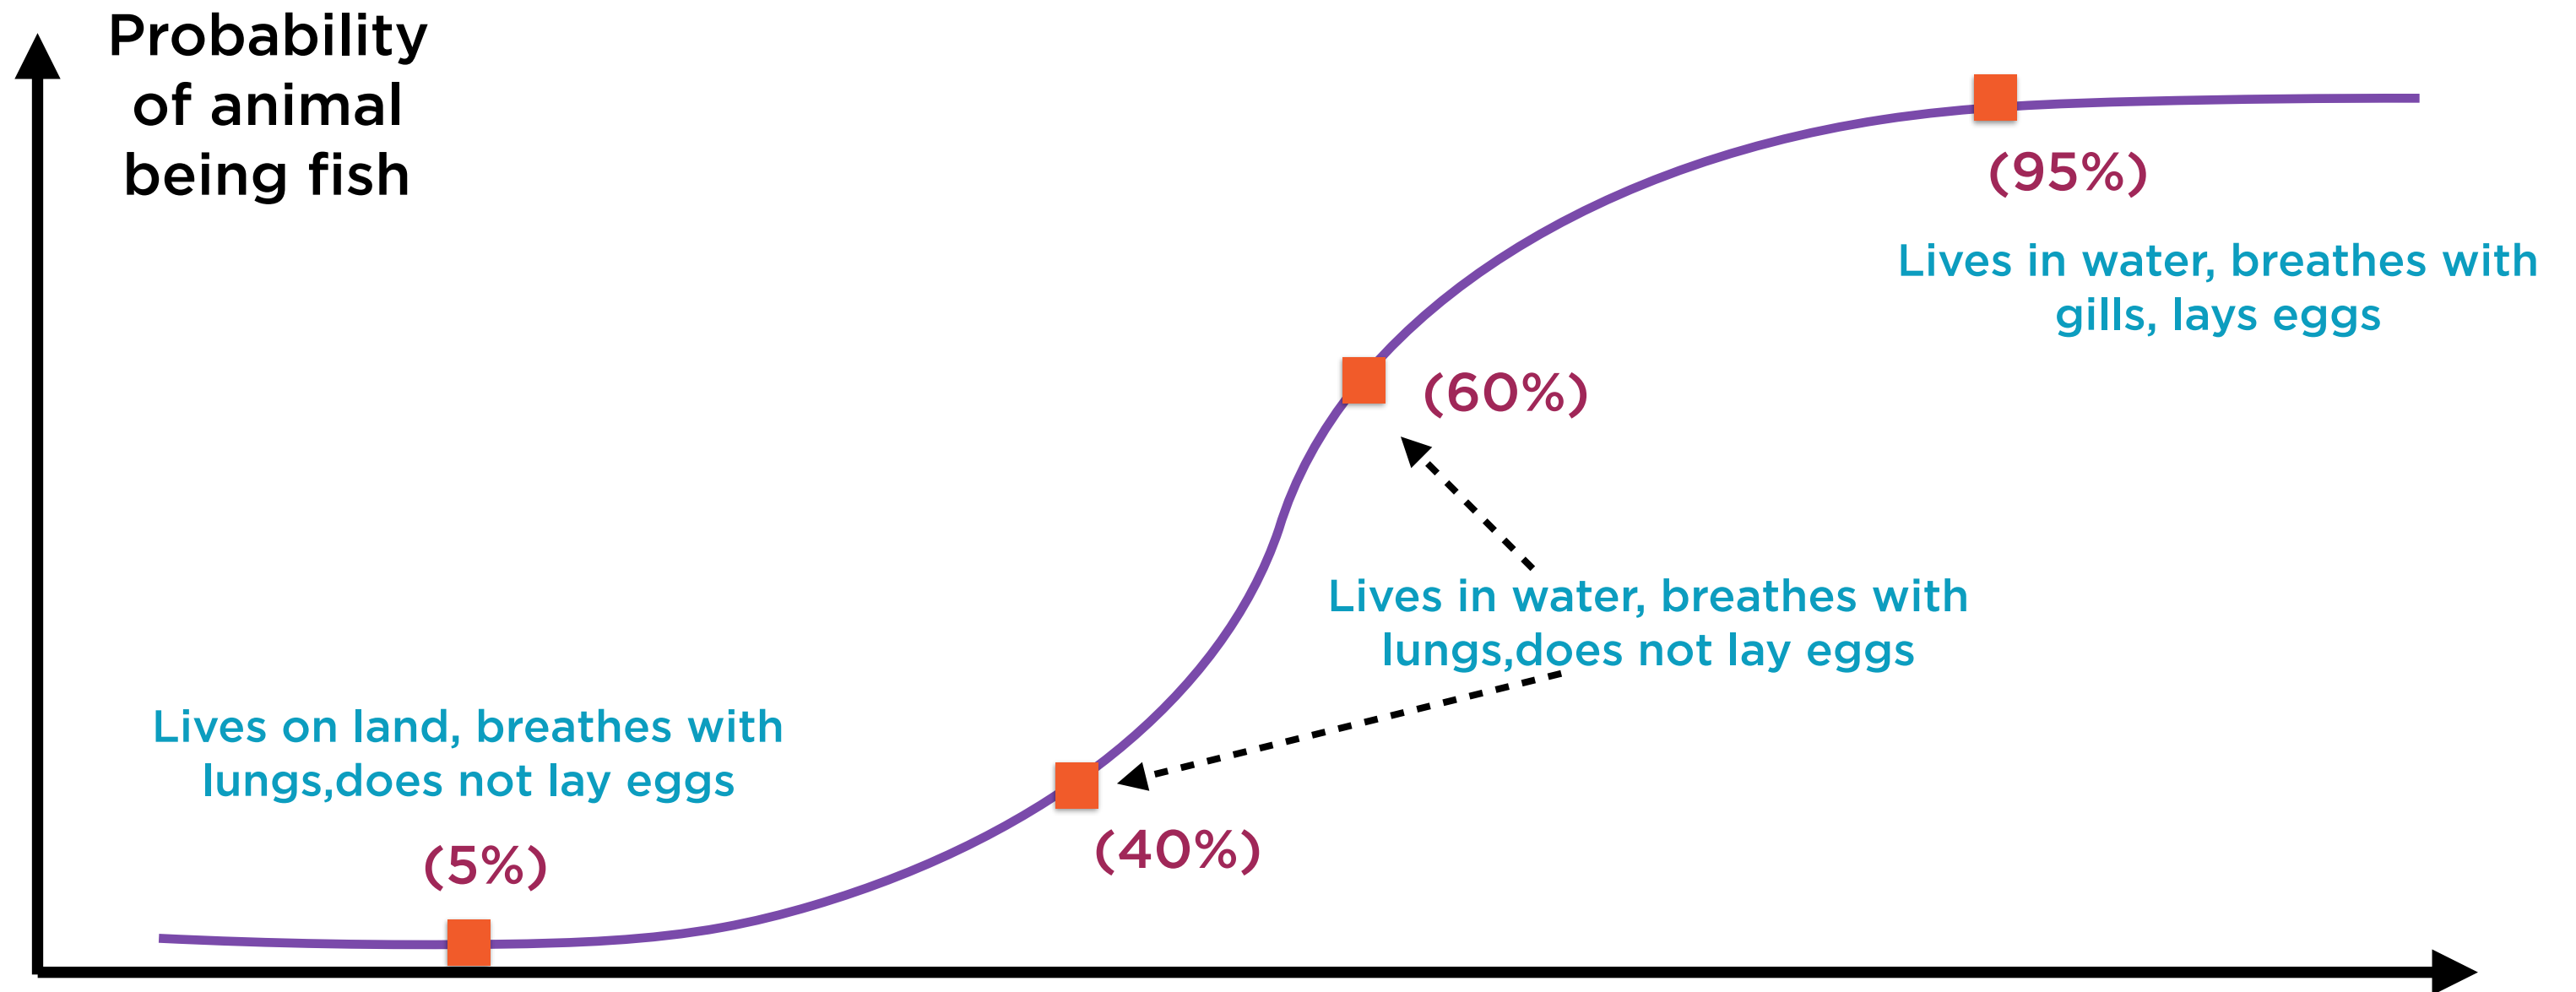


$P(\text{mammal}) = 0.55$



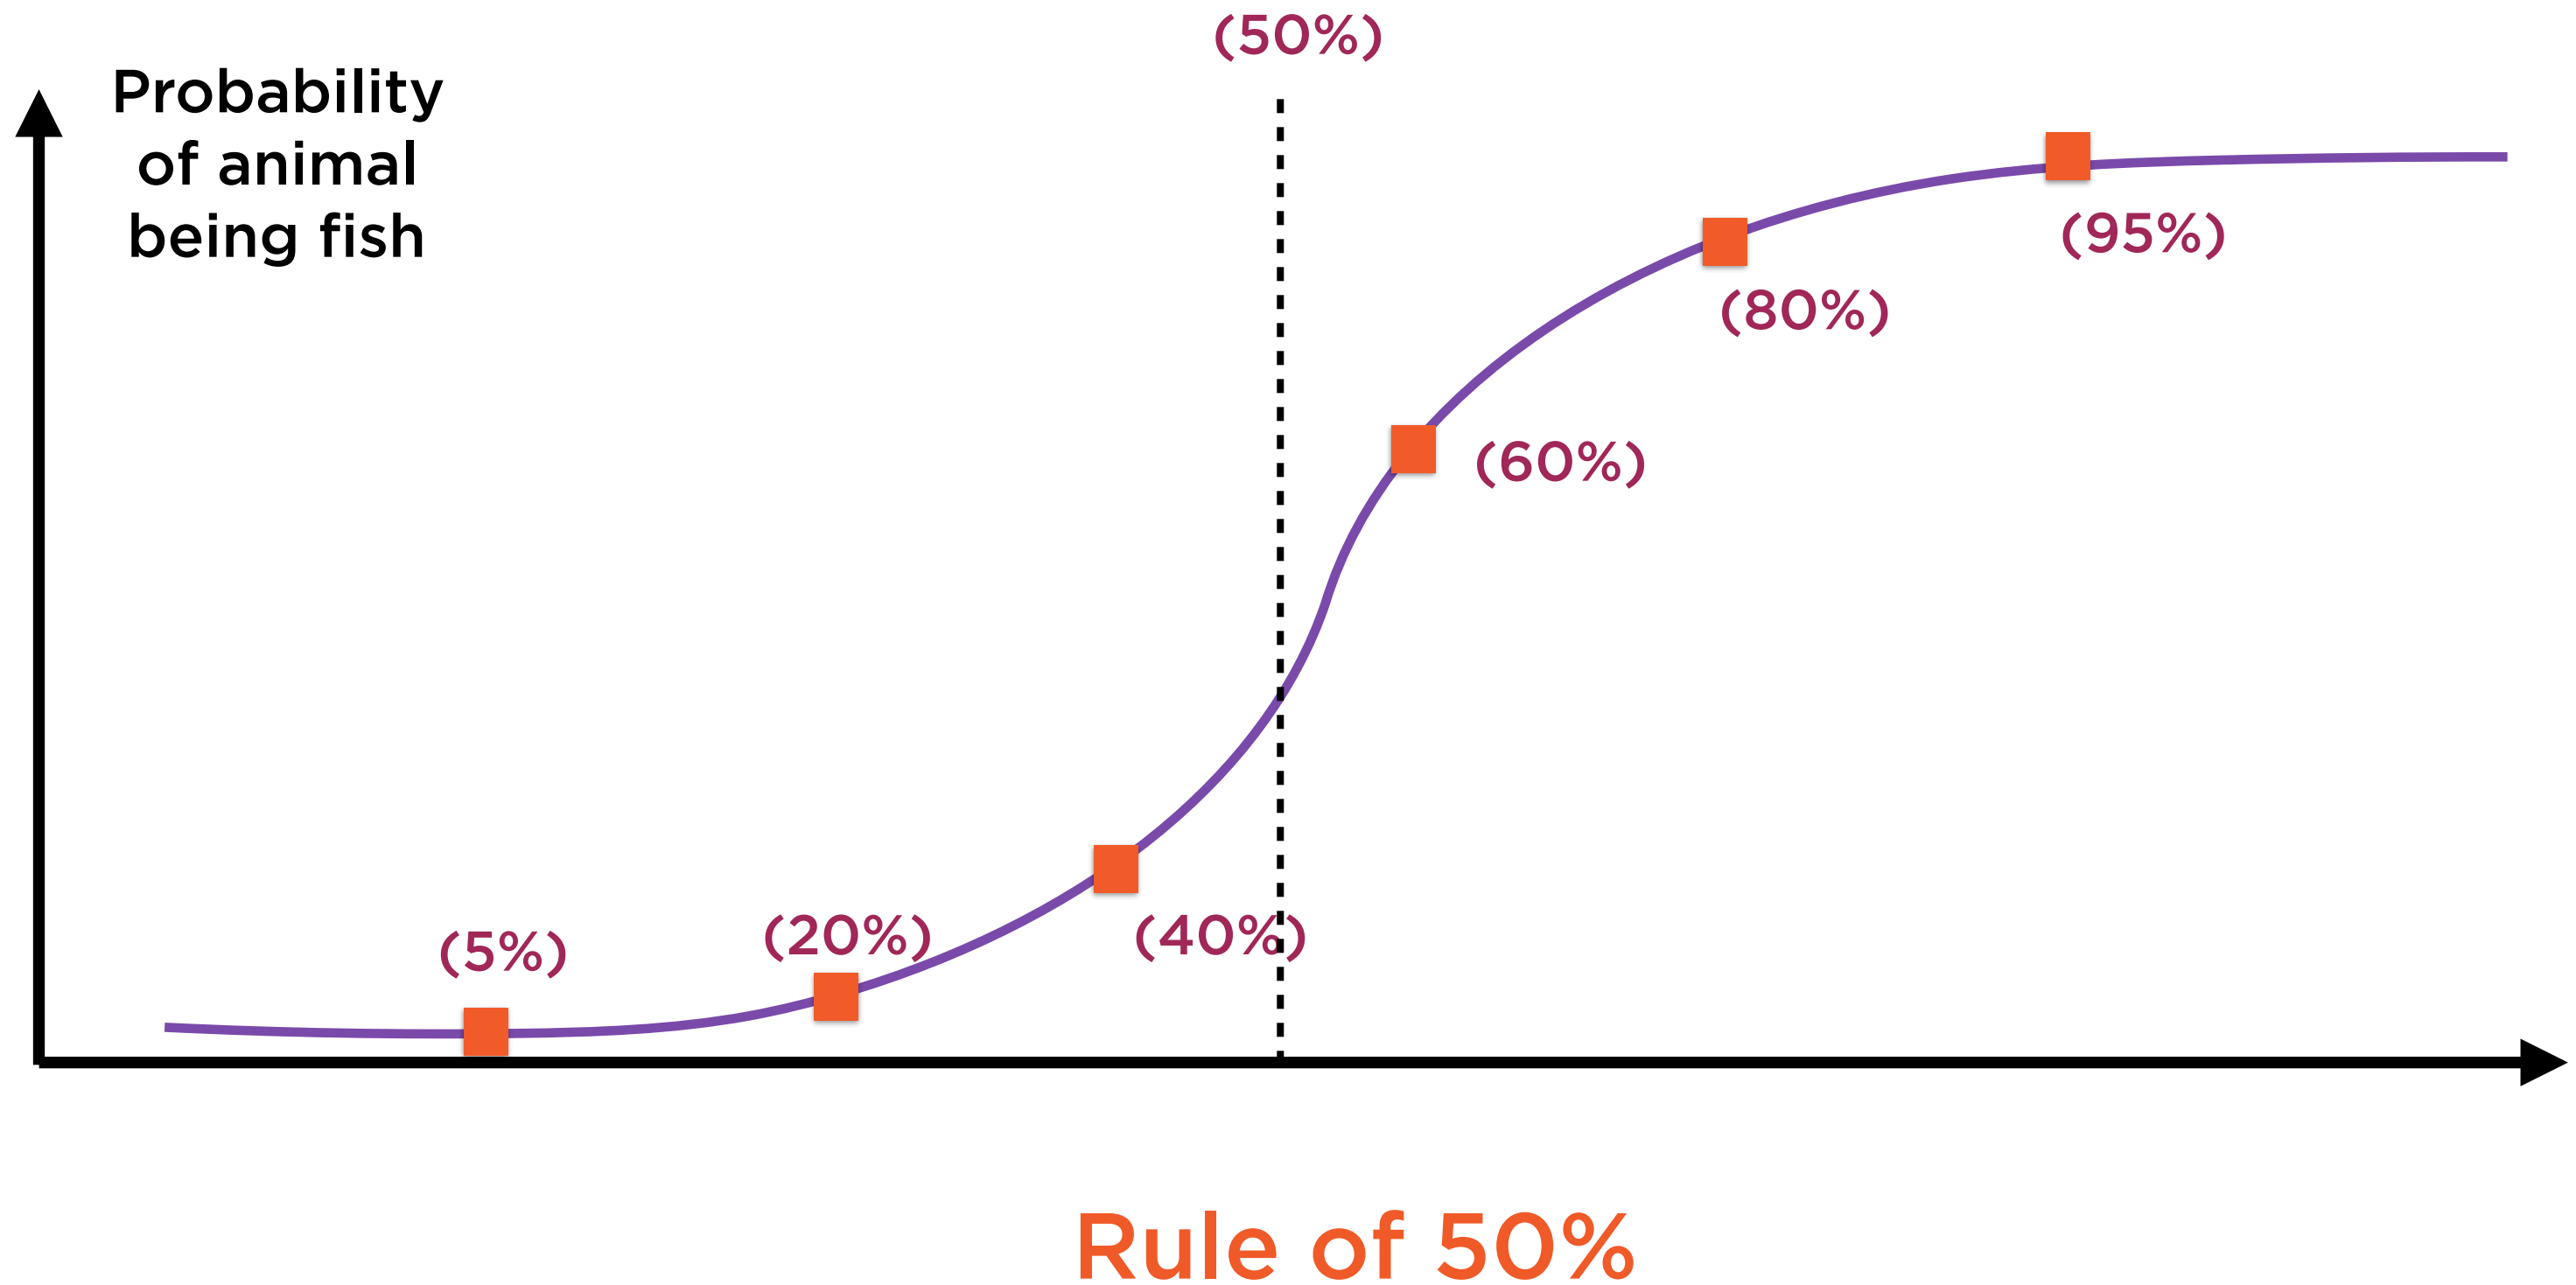
Corpus

Applying Logistic Regression




Whales: Fish or Mammals?

Applying Logistic Regression




Applying Logistic Regression



If probability < 50%, it's a mammal

Applying Logistic Regression



If probability > 50%, it's a fish

Applying Logistic Regression



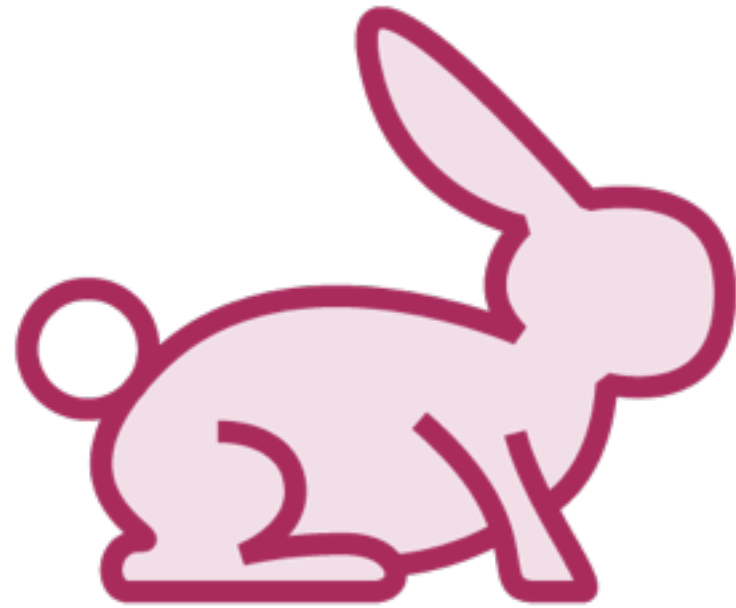
Mammal




Fish

Probability of whales being fish $< 50\%$

Applying Logistic Regression



Mammal



Fish



Probability of whales being fish $> 50\%$

Logistic Regression and Linear Regression

X Causes Y



Cause

Independent variable



Effect

Dependent variable

X Causes Y



Cause

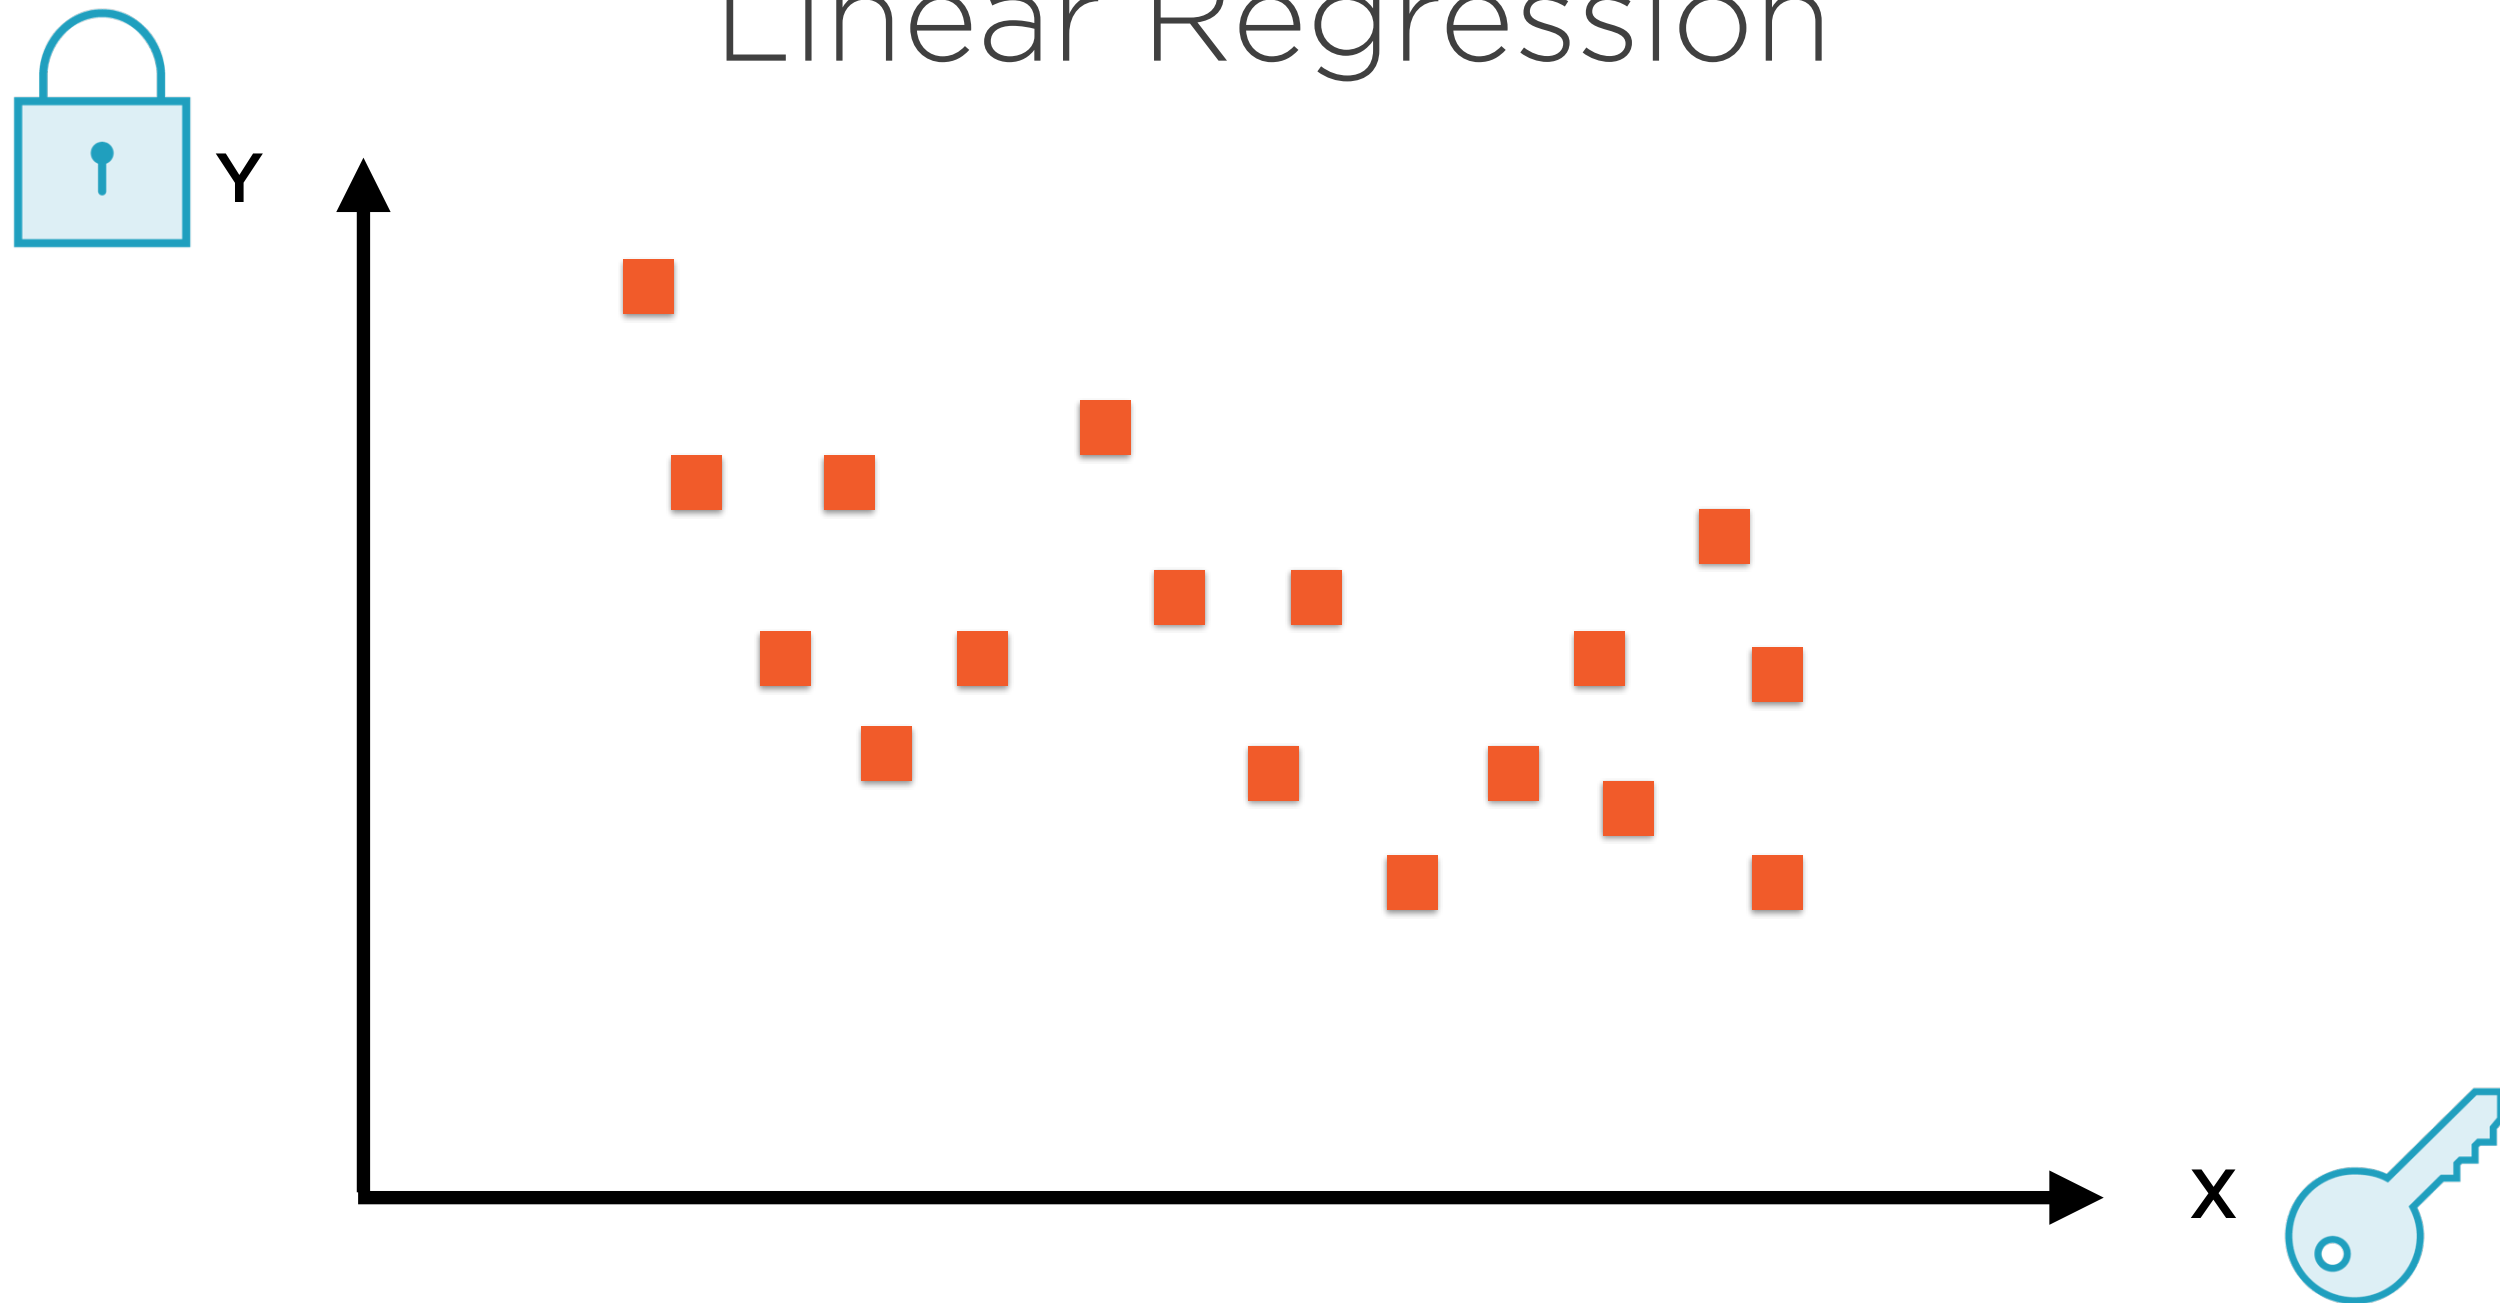
Explanatory variable



Effect


Dependent variable

Linear Regression



Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Linear Regression




Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Logistic Regression



Represent all n points as (x_i, y_i) , where $i = 1$ to n

Logistic Regression




Represent all n points as (x_i, y_i) , where $i = 1$ to n

Similar, yet Different


Linear Regression

Given causes, predict effect



Logistic Regression


Given causes, predict probability of effect



Similar, yet Different


Linear Regression

Effect variable (y) must be continuous



Logistic Regression

Effect variable (y) must be categorical



Similar, yet Different


Linear Regression

Cause variables (x) can be continuous or categorical



Logistic Regression


Cause variables (x) can be continuous or categorical



Similar, yet Different


Linear Regression

Connect the dots with a straight line



Logistic Regression


Connect the dots with an S-curve



Similar, yet Different


Linear Regression

$$y_i = A + Bx_i$$



Logistic Regression

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$



Similar, yet Different

Linear Regression

$$y_i = A + Bx_i$$

Objective of regression is to find A, B
that “best fit” the data

Logistic Regression

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$

Objective of regression is to find A, B
that “best fit” the data

Similar, yet Different

Linear Regression

$$y_i = A + Bx_i$$

Relationship is already linear (by assumption)

Logistic Regression

$$\ln\left(\frac{p(y_i)}{1 - p(y_i)}\right) = A + Bx_i$$

Relationship can be made linear (by log transformation)

Similar, yet Different

Linear Regression

$$y_i = A + Bx_i$$

Solve regression problem using cookie-cutter solvers

Logistic Regression

$$\text{logit}(p) = A + Bx_i$$

$$\text{logit}(p) = \ln\left(\frac{p}{1-p}\right)$$

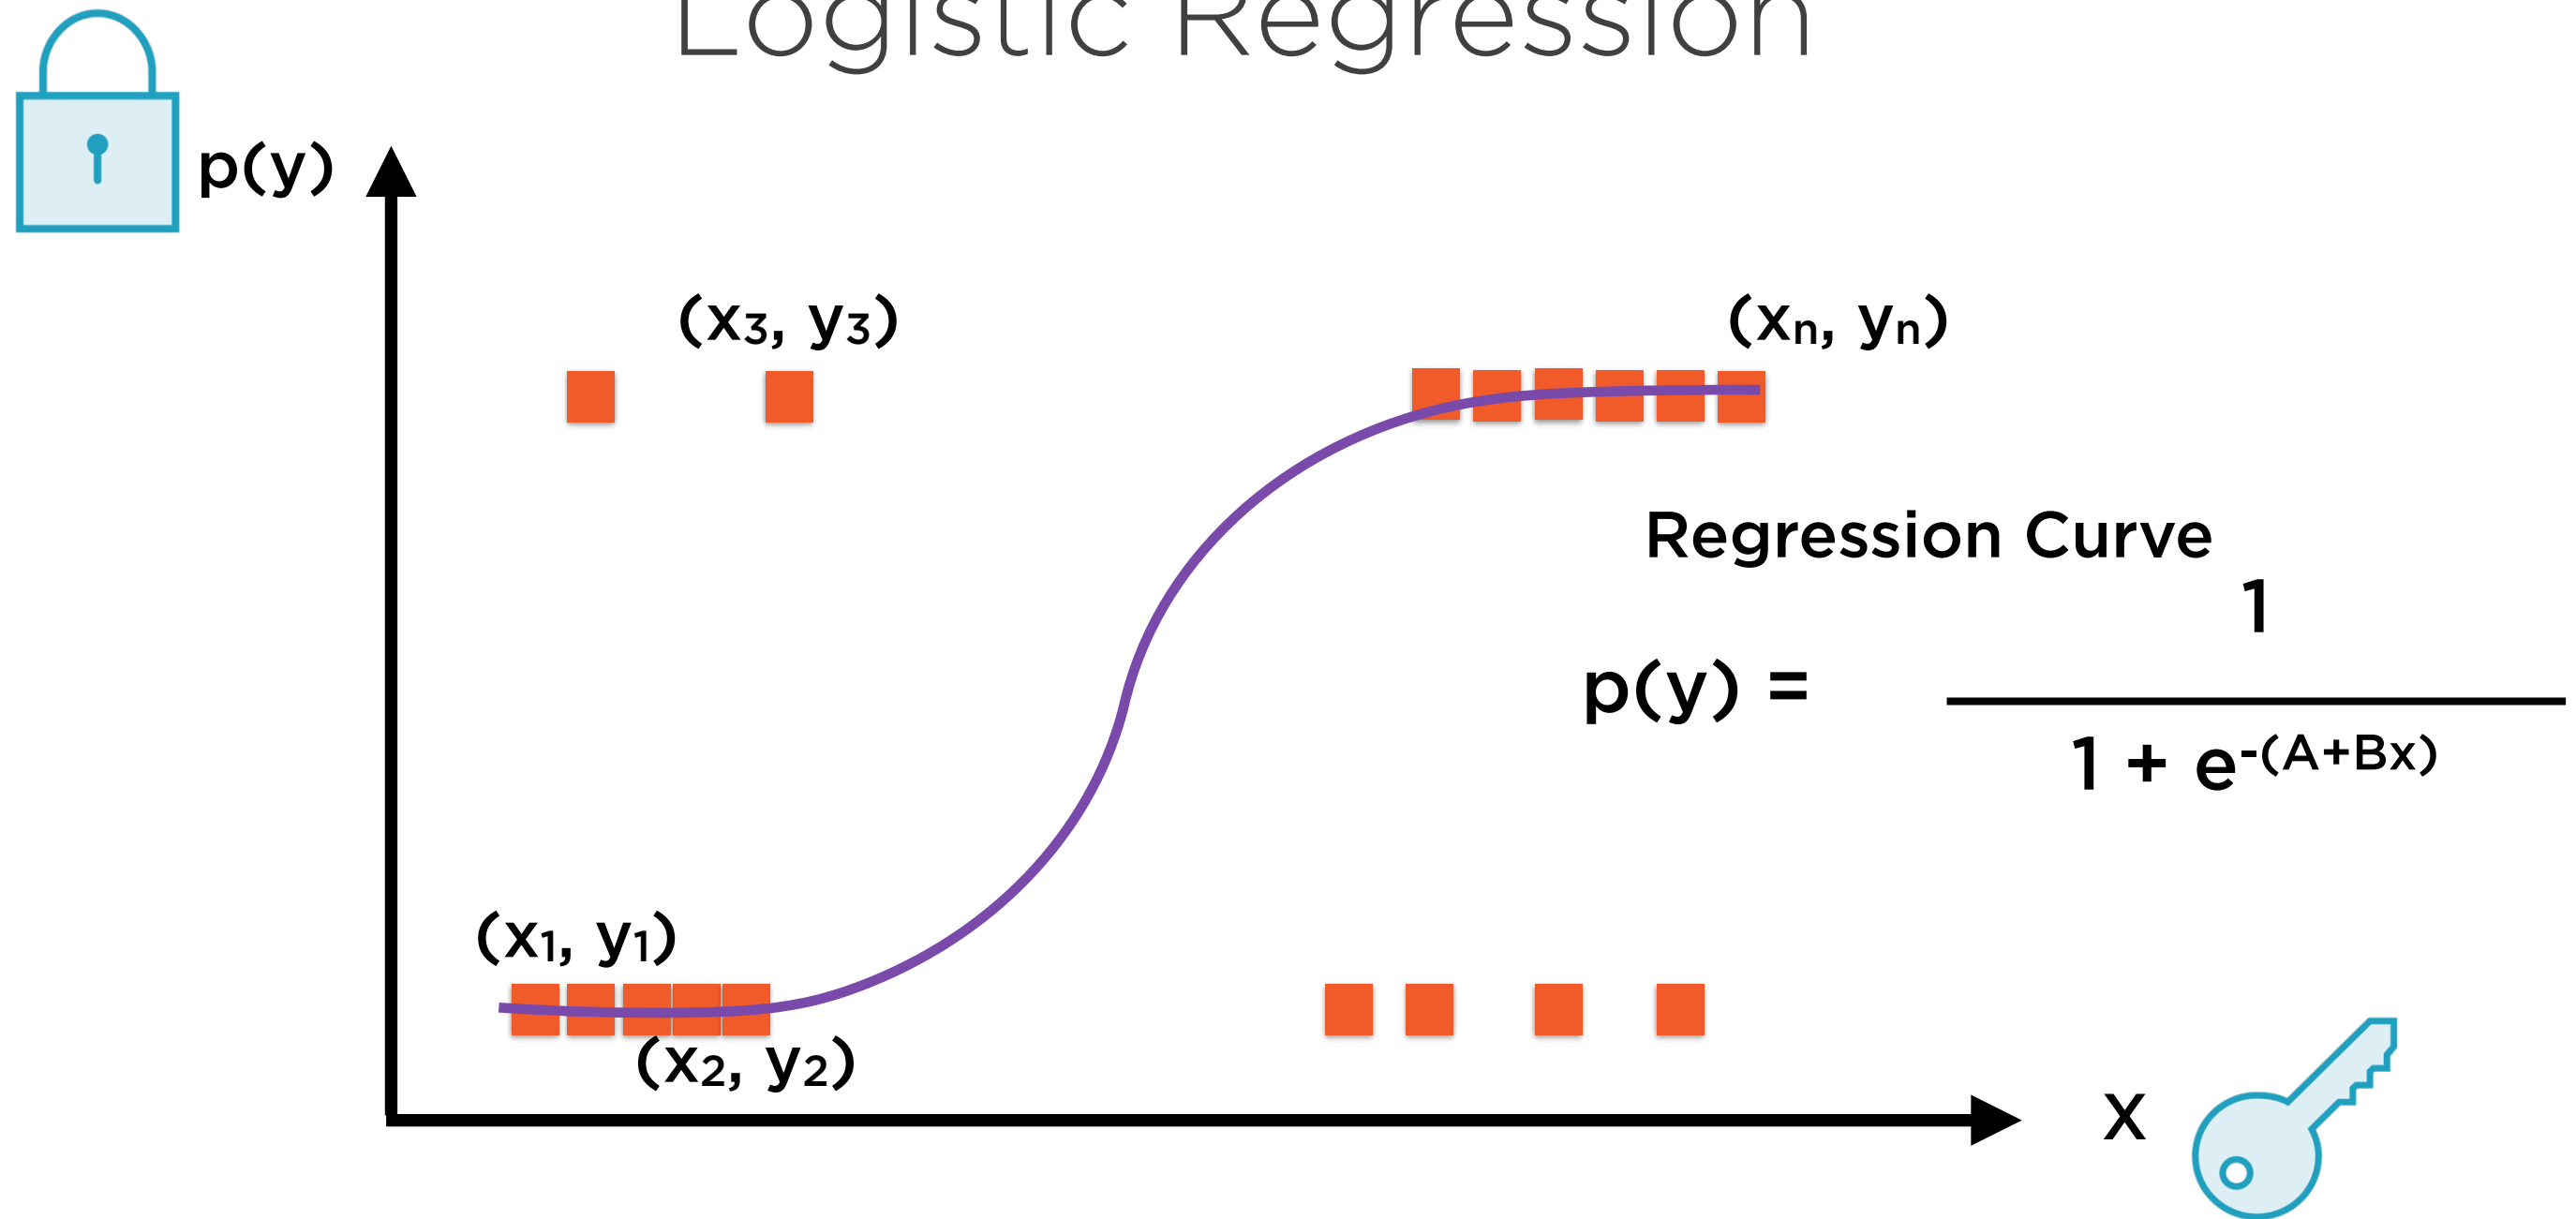
Solve regression problem using cookie-cutter solvers

Logistic Regression



Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Logistic Regression



Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Linear Regression

$$y = A + Bx$$

$$y_1 = A + Bx_1$$

$$y_2 = A + Bx_2$$

$$y_3 = A + Bx_3$$

...

...

$$y_n = A + Bx_n$$

Logistic Regression

$$p(y) = \frac{1}{1 + e^{-(A+Bx)}}$$

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$

$$p(y_1) = \frac{1}{1 + e^{-(A+Bx_1)}}$$

...


$$p(y_n) = \frac{1}{1 + e^{-(A+Bx_n)}}$$

Logistic Regression



Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Logistic Regression



Represent all n points as
 (x_i, y_i) , where $i = 1$ to n

Logistic Regression

Regression Equation:

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$

Solve for A and B that “best fit” the data

Odds from Probabilities

$$\text{Odds}(p) = \frac{p}{1-p}$$

Odds of an Event

$$p = \frac{1}{1 + e^{-(A+Bx)}}$$

$$p = \frac{e^{A + Bx}}{1 + e^{A + Bx}}$$

$$1 - p = 1 - \frac{e^{A + Bx}}{1 + e^{A + Bx}}$$

Odds of an Event

$$1 - p = 1 - \frac{e^{A + Bx}}{1 + e^{A + Bx}}$$

$$1 - p = \frac{1 + e^{A + Bx} - e^{A + Bx}}{1 + e^{A + Bx}}$$

$$1 - p = \frac{1}{1 + e^{A + Bx}}$$

Odds of an Event

$$p = \frac{e^{A + Bx}}{1 + e^{A + Bx}}$$

$$1 - p = \frac{1}{1 + e^{A + Bx}}$$

$$\text{Odds}(p) = \frac{p}{1 - p} = e^{A + Bx}$$

Logit Is Linear

$$\text{Odds}(p) = \frac{p}{1 - p} = e^{A + Bx}$$

$$\text{logit}(p) = A + Bx$$

$\ln(\text{Odds}(p))$ is called the logit function

Logit Is Linear

$$\ln \text{Odds}(p) = \ln(p) - \ln(1-p)$$

$$p = \frac{1}{1 + e^{-(A+Bx)}}$$

$$\text{logit}(p) = \ln \text{Odds}(p) = A + Bx$$

This is a linear function!

Logistic Regression can be solved via **linear regression on the logit function** (log of the odds function)

Logistic Regression in TensorFlow

Logistic Regression



Cause

Changes in S&P 500



Effect

Changes in price of Google Stock

Logistic Regression

**y = Returns on
Google stock
(GOOG)**

**x = Returns
on S&P 500
(S&P500)**


Logistic Regression

$$p(y_i) = \frac{1}{1 + e^{-(A+Bx_i)}}$$


$P(y)$ = Probability of
Google going up in
the current month i

x = Returns on S&P
500 for current
month

Logistic Regression



Set up the Problem



Label GOOG returns as binary (1,0)

Prediction Accuracy

| DATE | ACTUAL | PREDICTED |
|------------|--------|-----------|
| 2005-01-01 | NA | NA |
| 2005-02-01 | 0 | 1 |
| 2005-03-01 | 0 | 0 |
| | | |
| 2017-01-01 | 1 | 1 |
| 2017-02-01 | 1 | 1 |

Compare GOOG's actual labels vs. predicted labels

Linear Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Mean Square Error (MSE)
Quantifying goodness-of-fit

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Affine transformation suffices

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

Linear Regression in TensorFlow

Baseline

Non-TensorFlow implementation

Regular python code

Cost Function

Mean Square Error (MSE)

Quantifying goodness-of-fit

Training

Invoke optimizer in epochs

Batch size for each epoch

Computation Graph

Neural network of 1 neuron

Affine transformation suffices

Optimizer

Gradient Descent optimizers

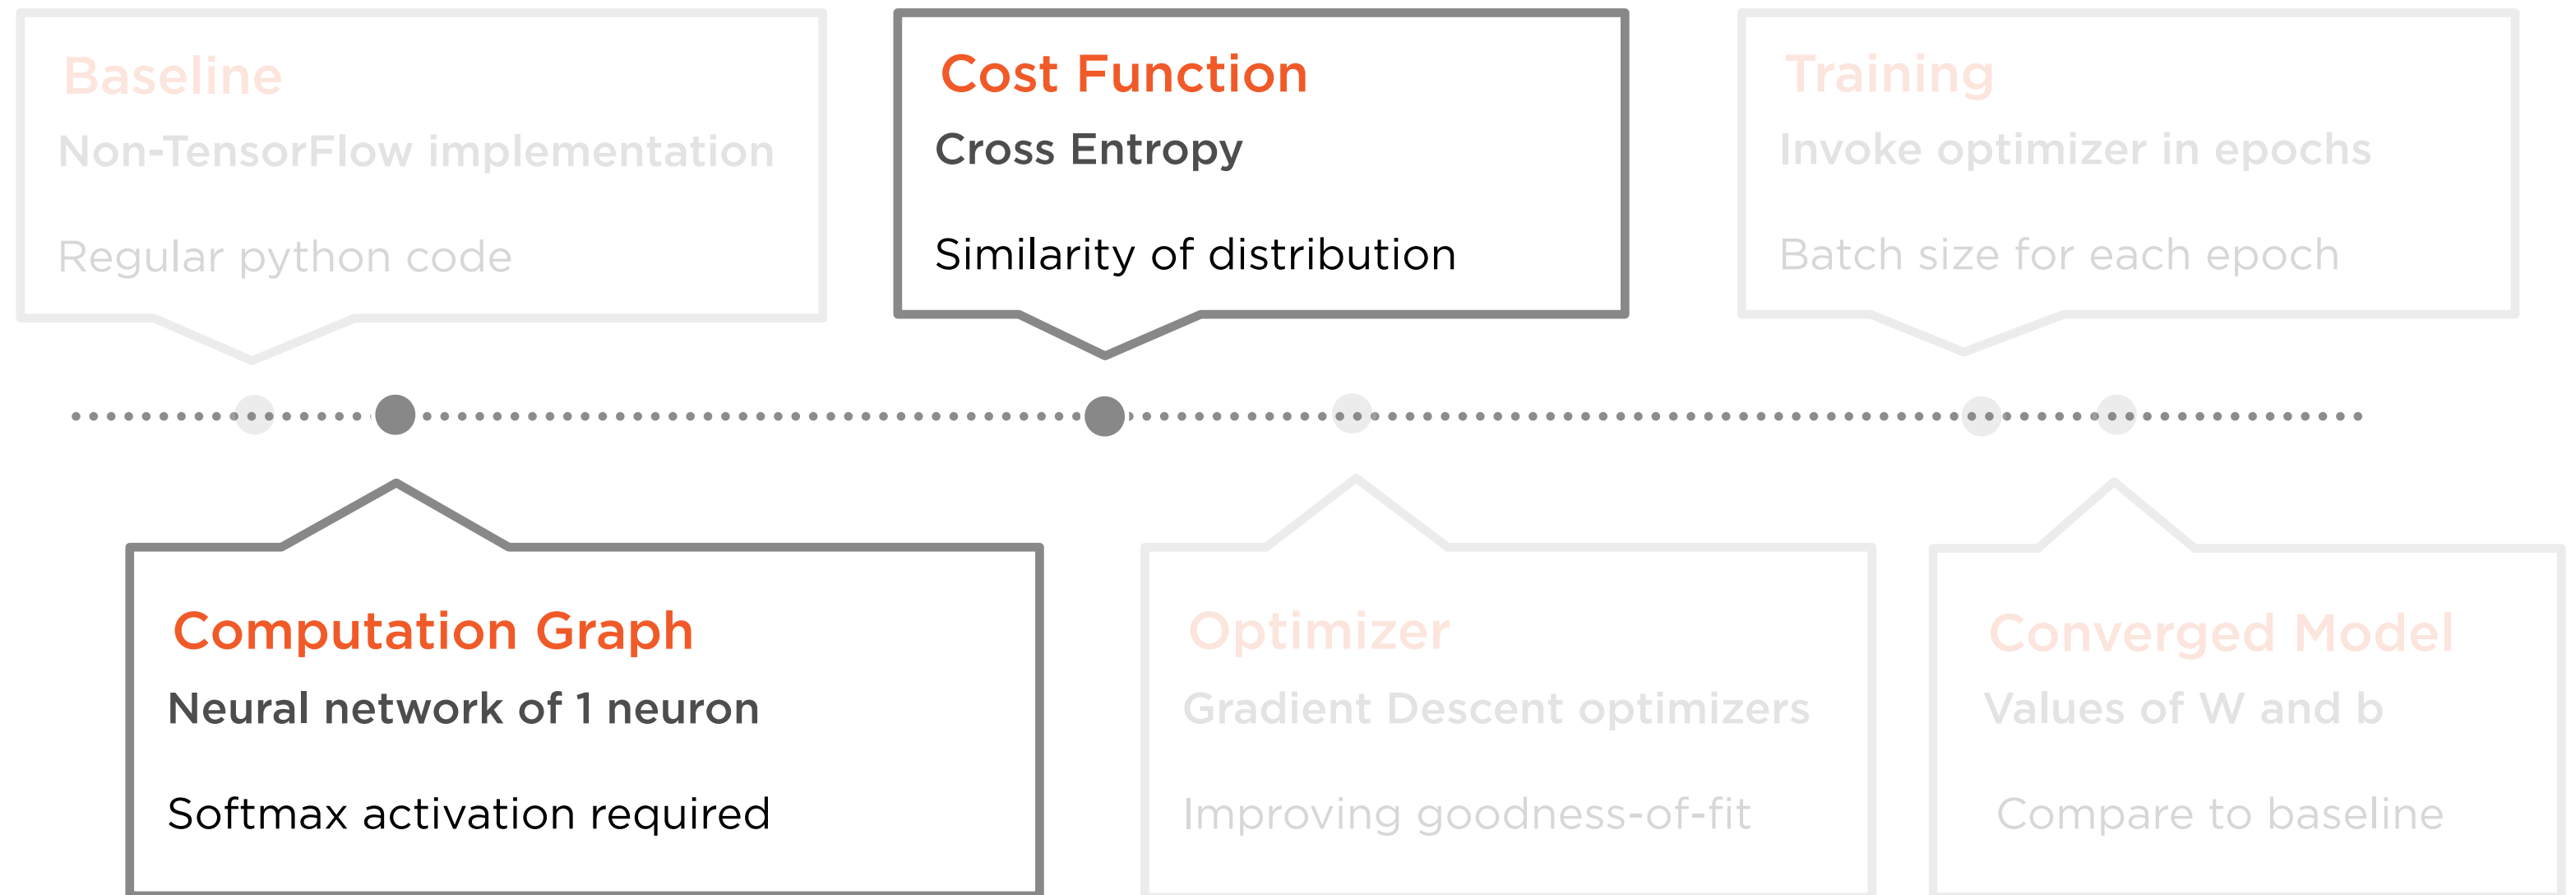
Improving goodness-of-fit

Converged Model

Values of W and b

Compare to baseline

Logistic Regression in TensorFlow



Logistic Regression in TensorFlow

Baseline

Non-TensorFlow implementation

Regular python code

Cost Function

Cross Entropy

Similarity of distribution

Training

Invoke optimizer in epochs

Batch size for each epoch

Computation Graph

Neural network of 1 neuron

Softmax activation required

Optimizer

Gradient Descent optimizers


Improving goodness-of-fit

Converged Model

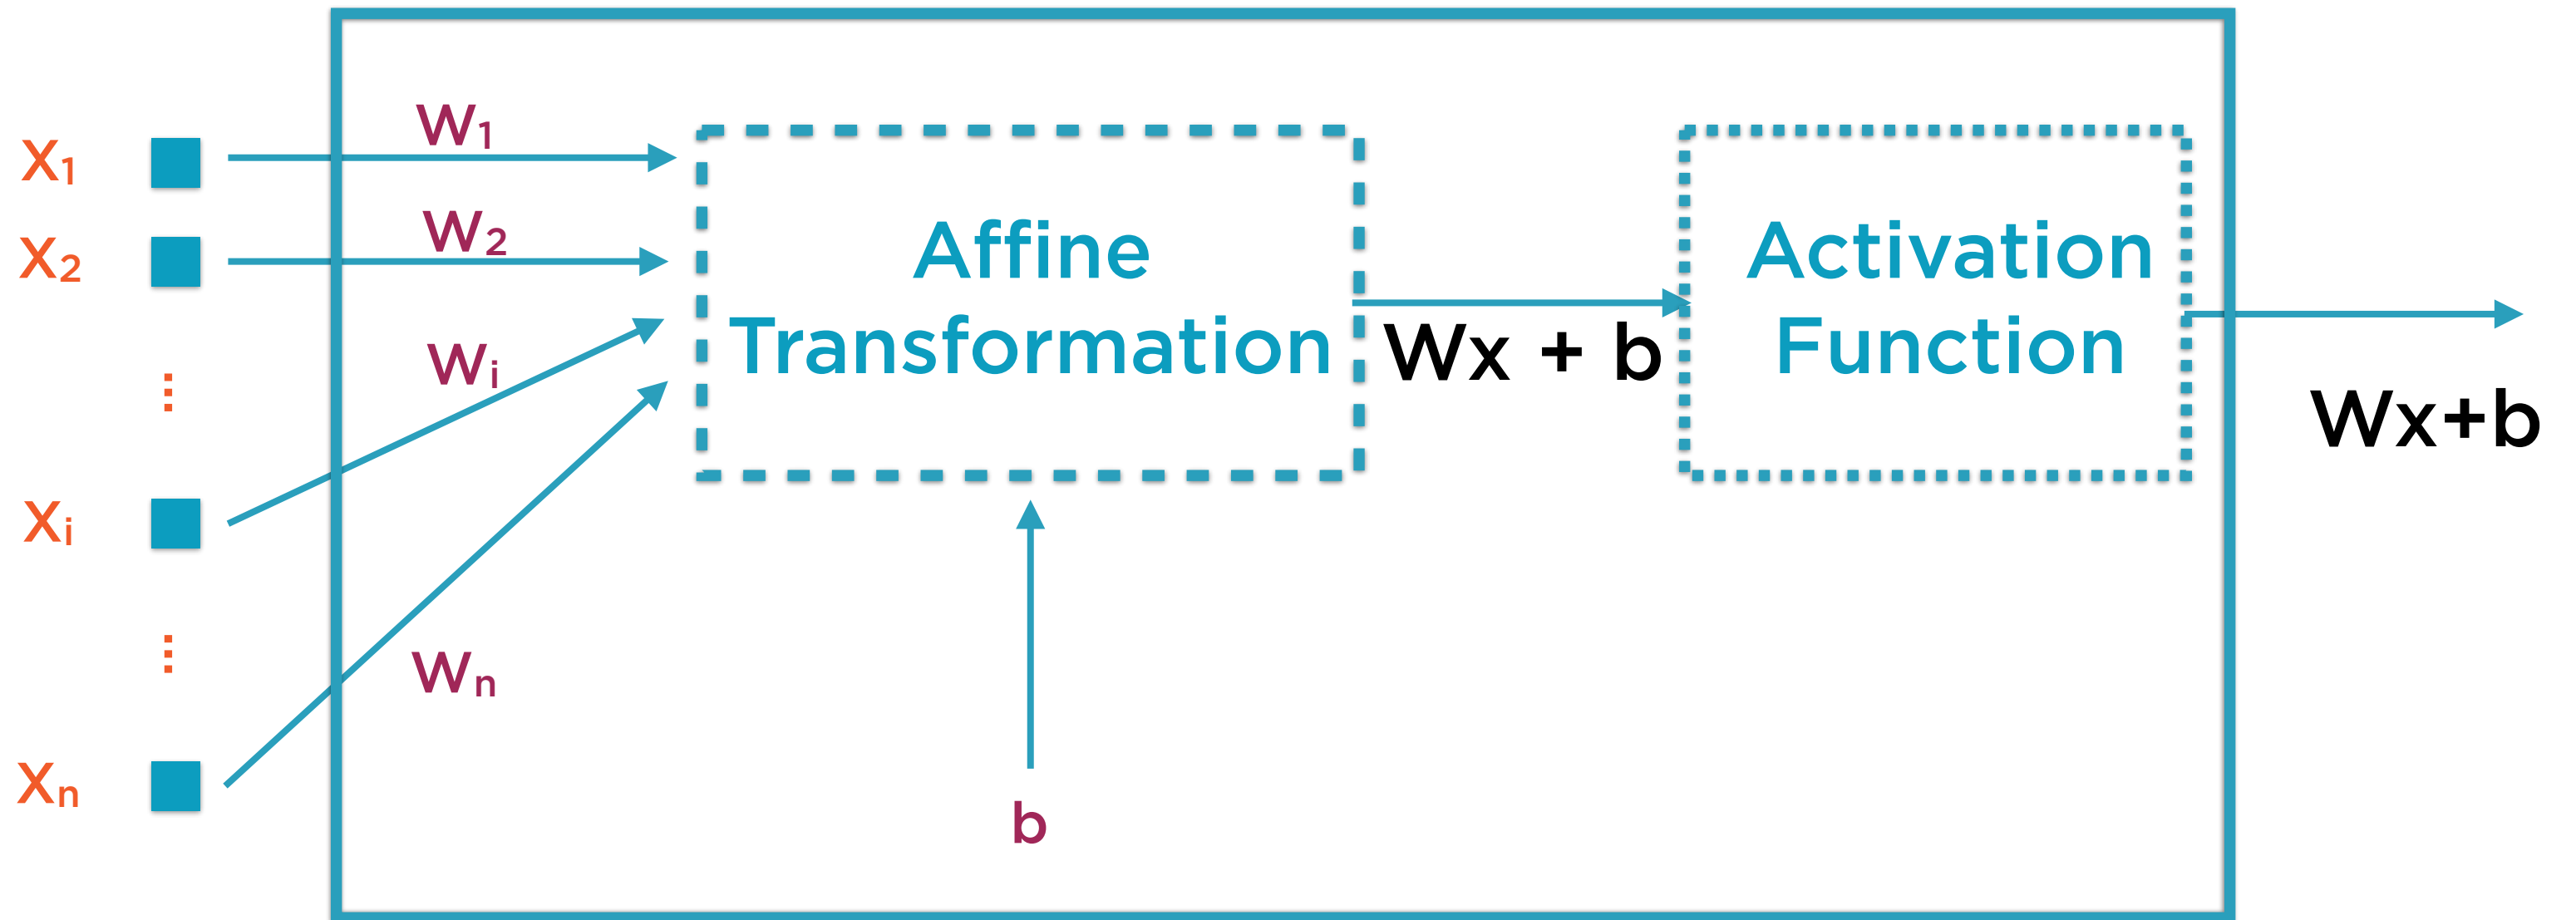
Values of W and b

Compare to baseline

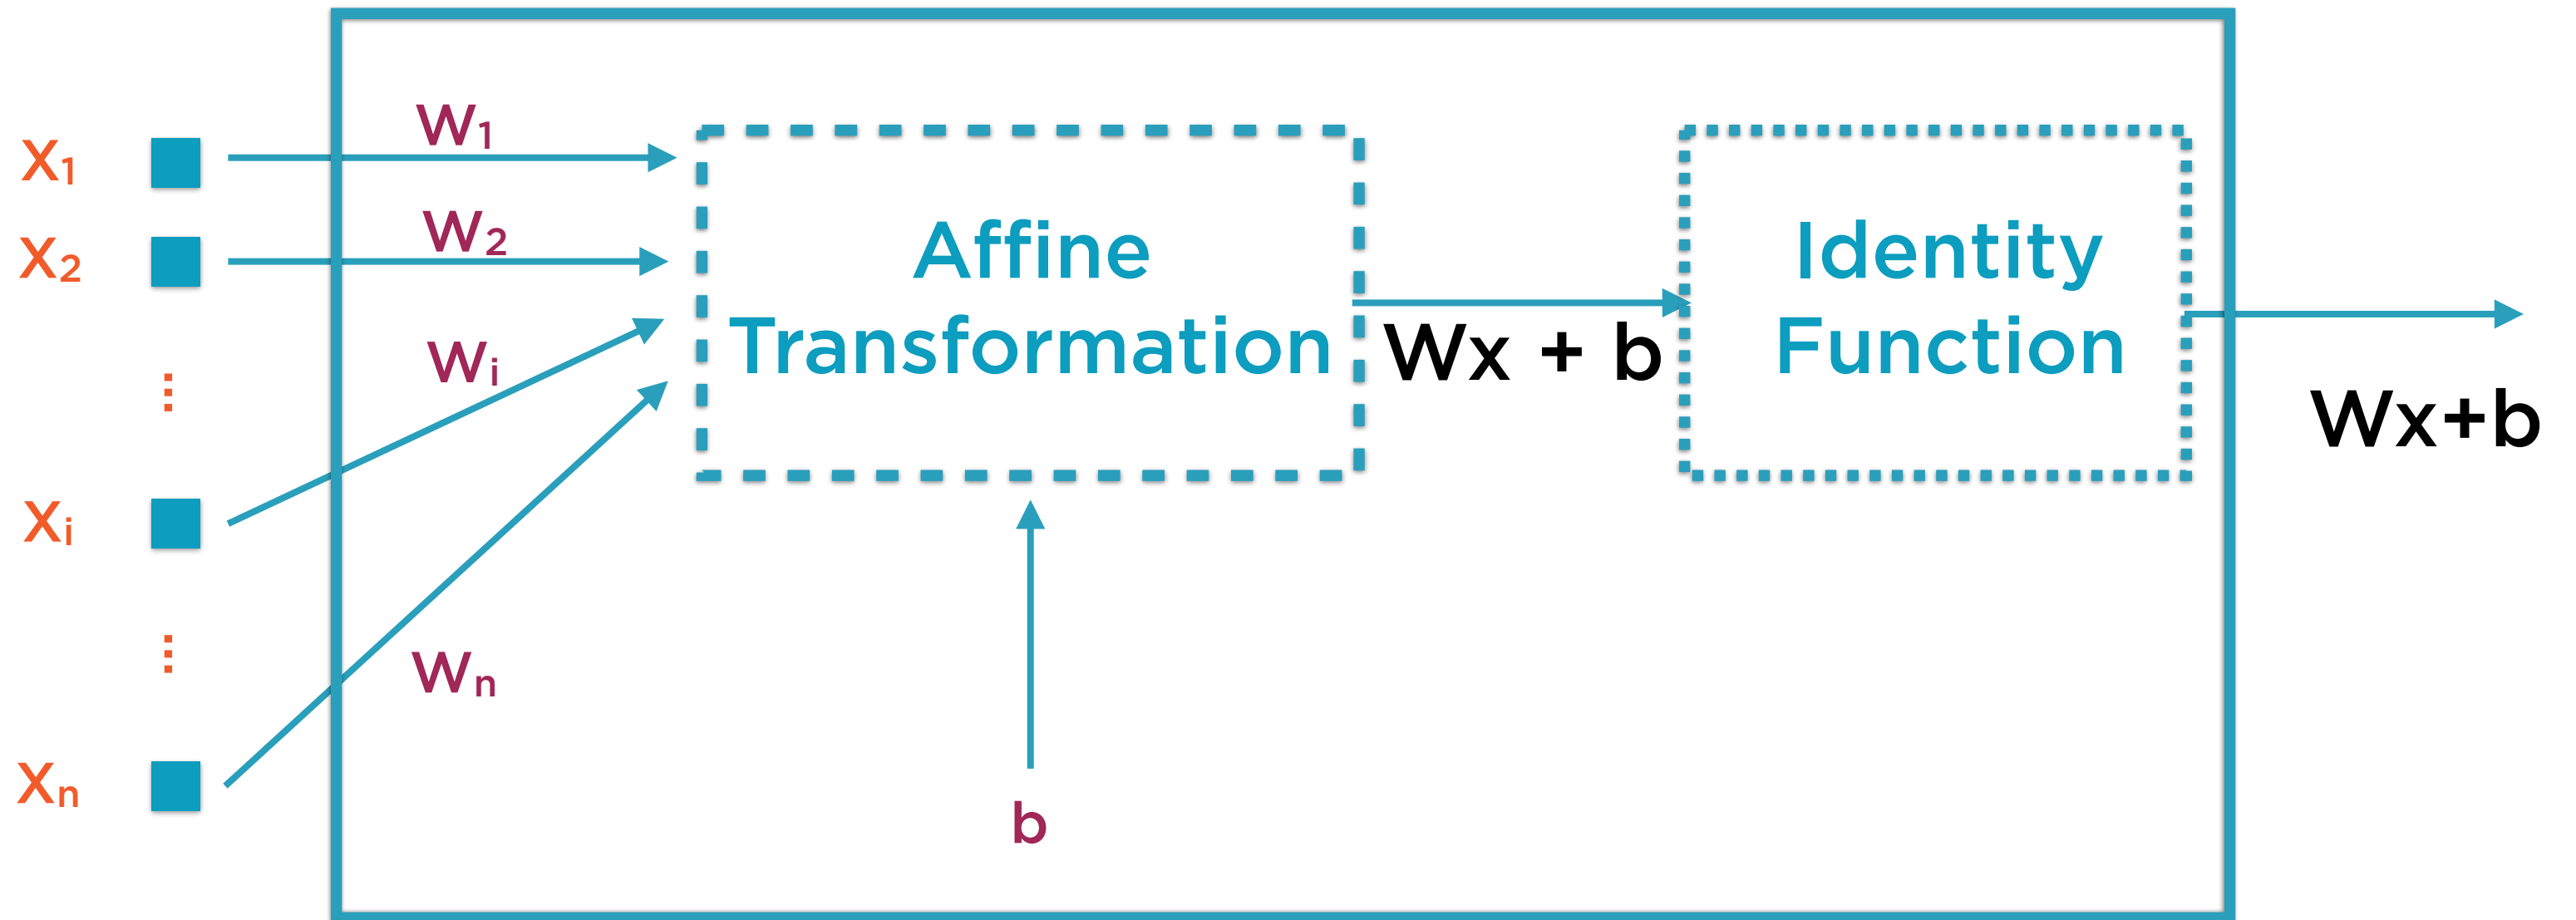
Linear Regression with One Neuron



Linear Regression with One Neuron



Linear Regression with One Neuron



Linear Regression with One Neuron



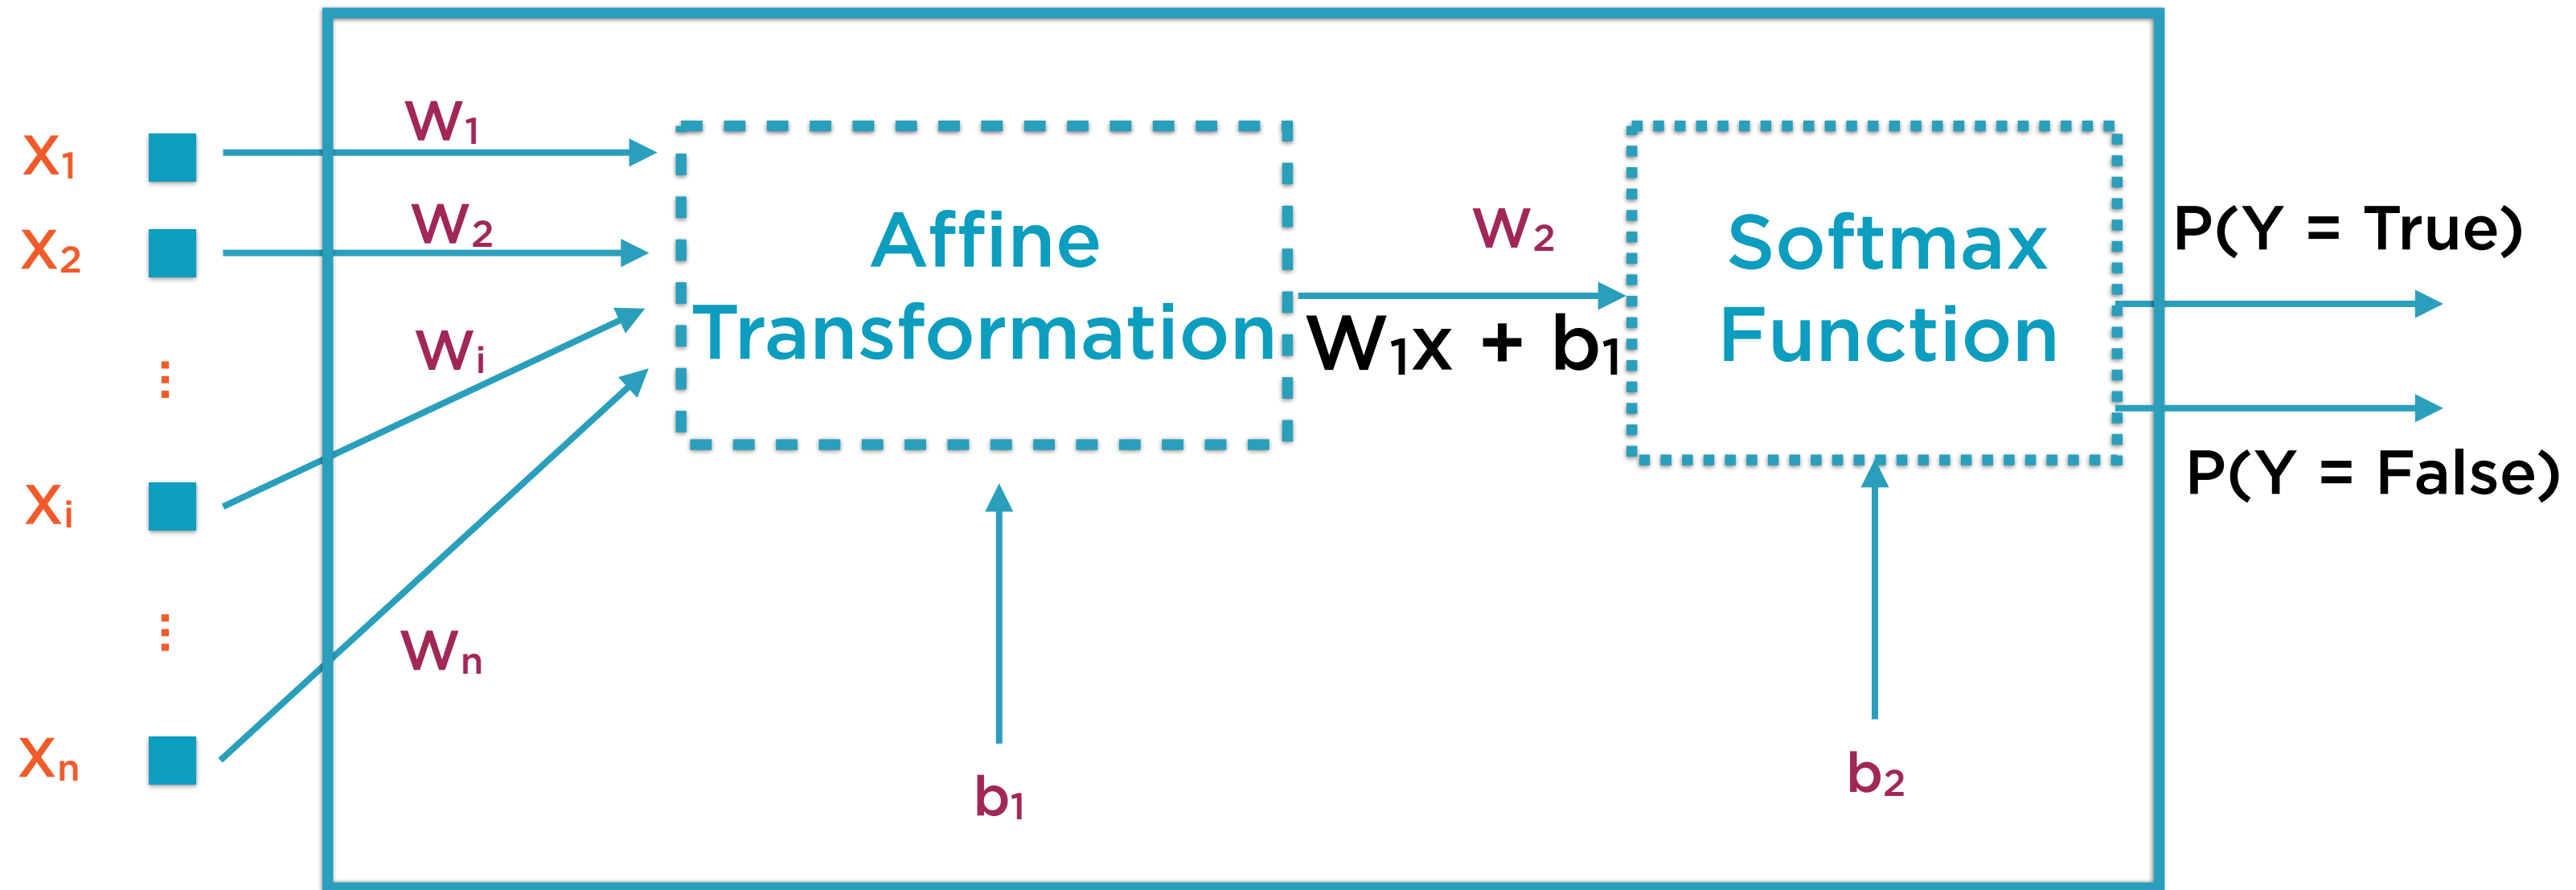
Logistic Regression with One Neuron



Logistic Regression with One Neuron



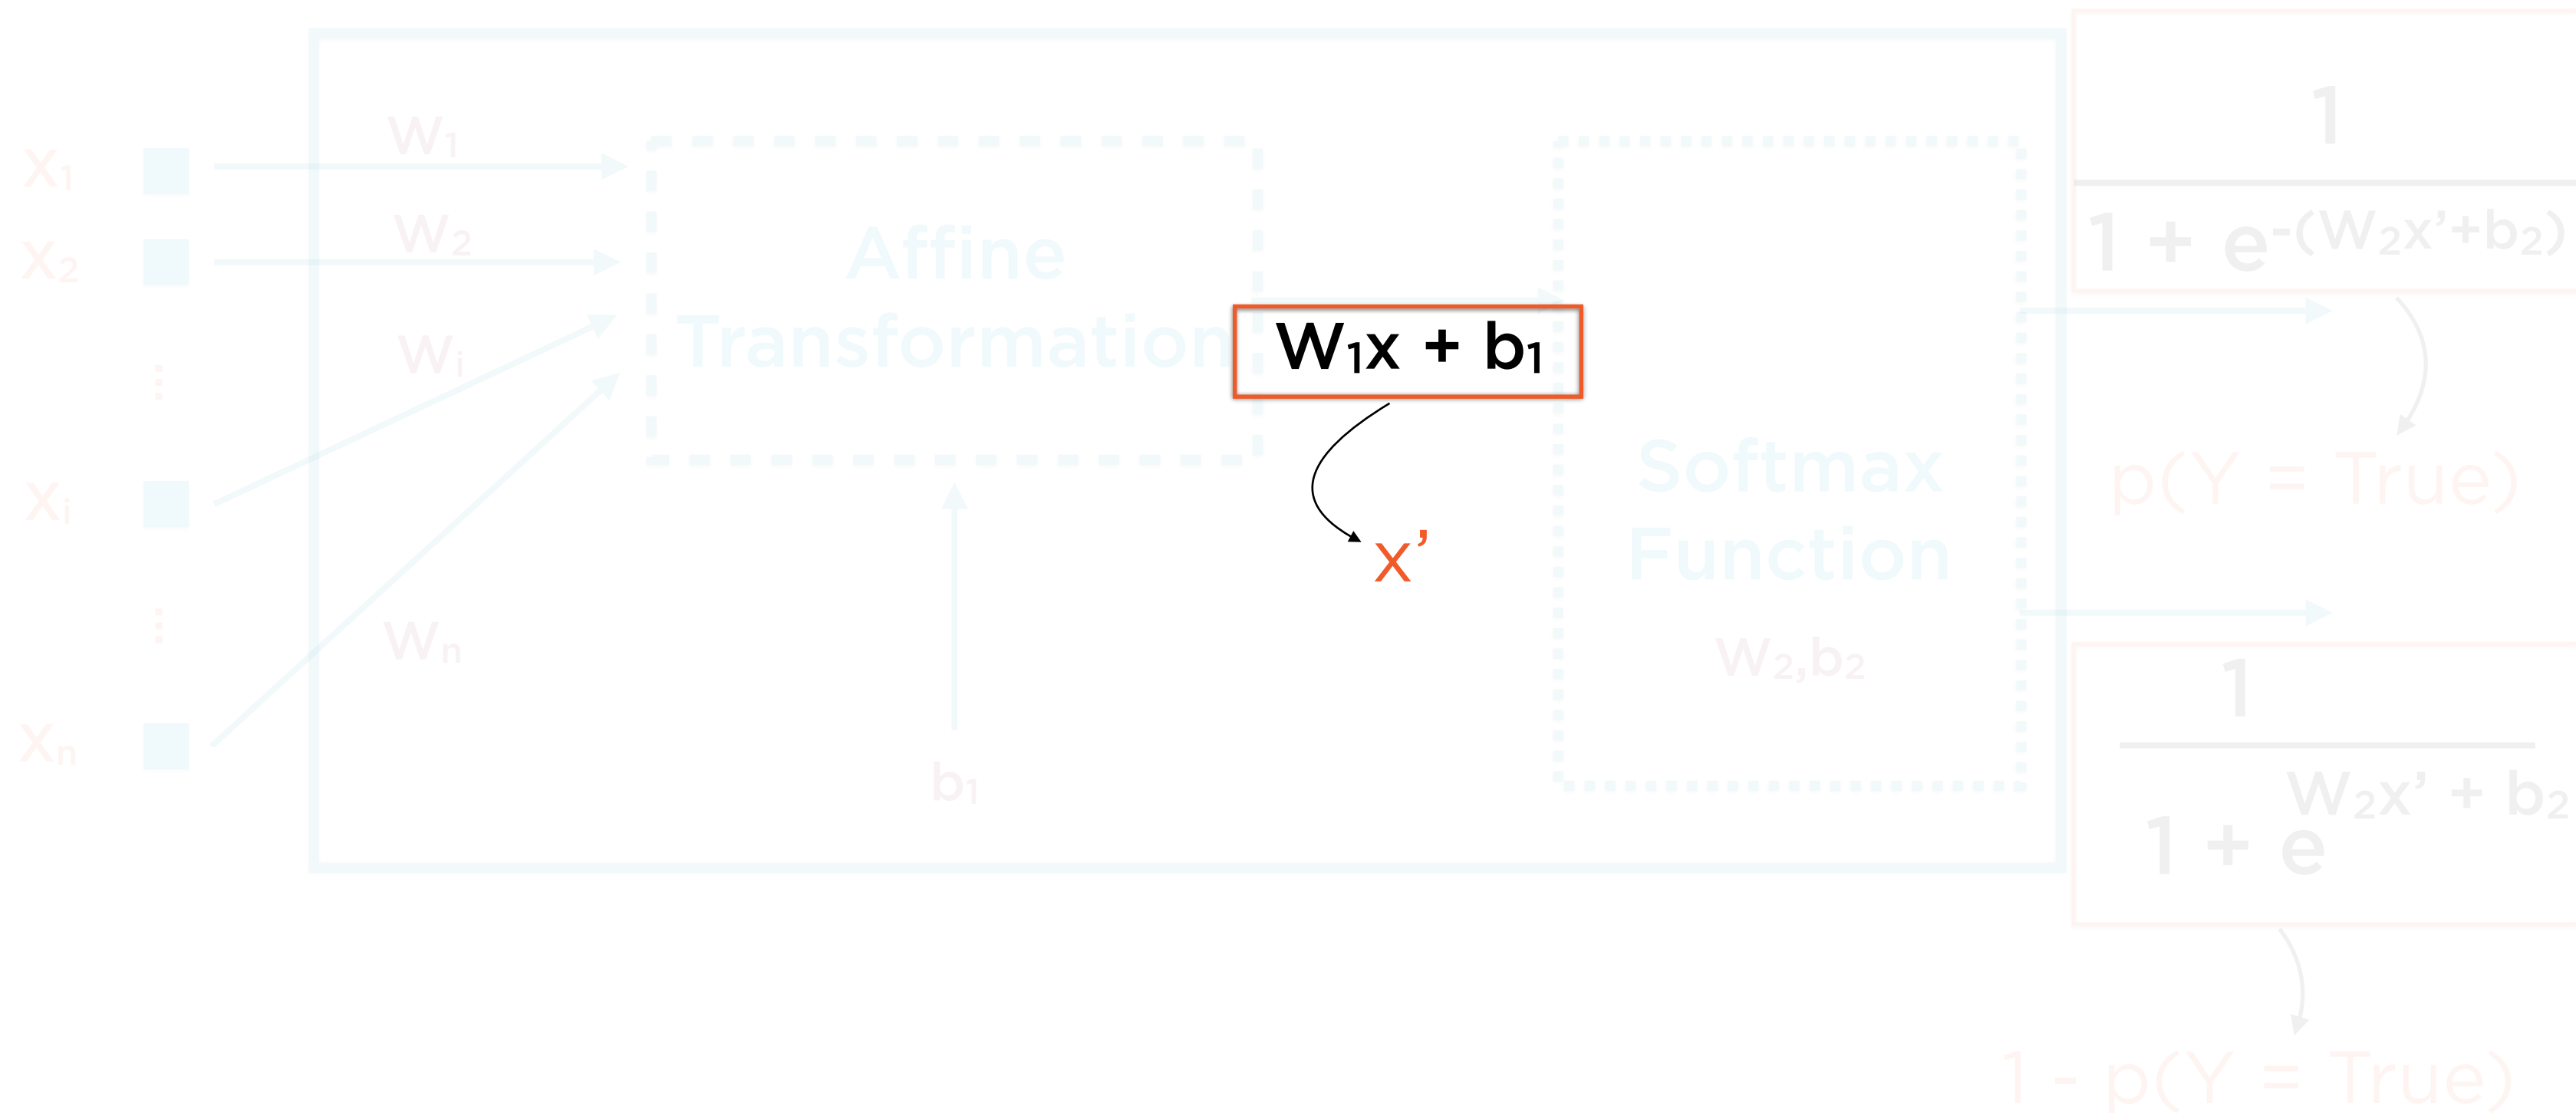
Logistic Regression with One Neuron



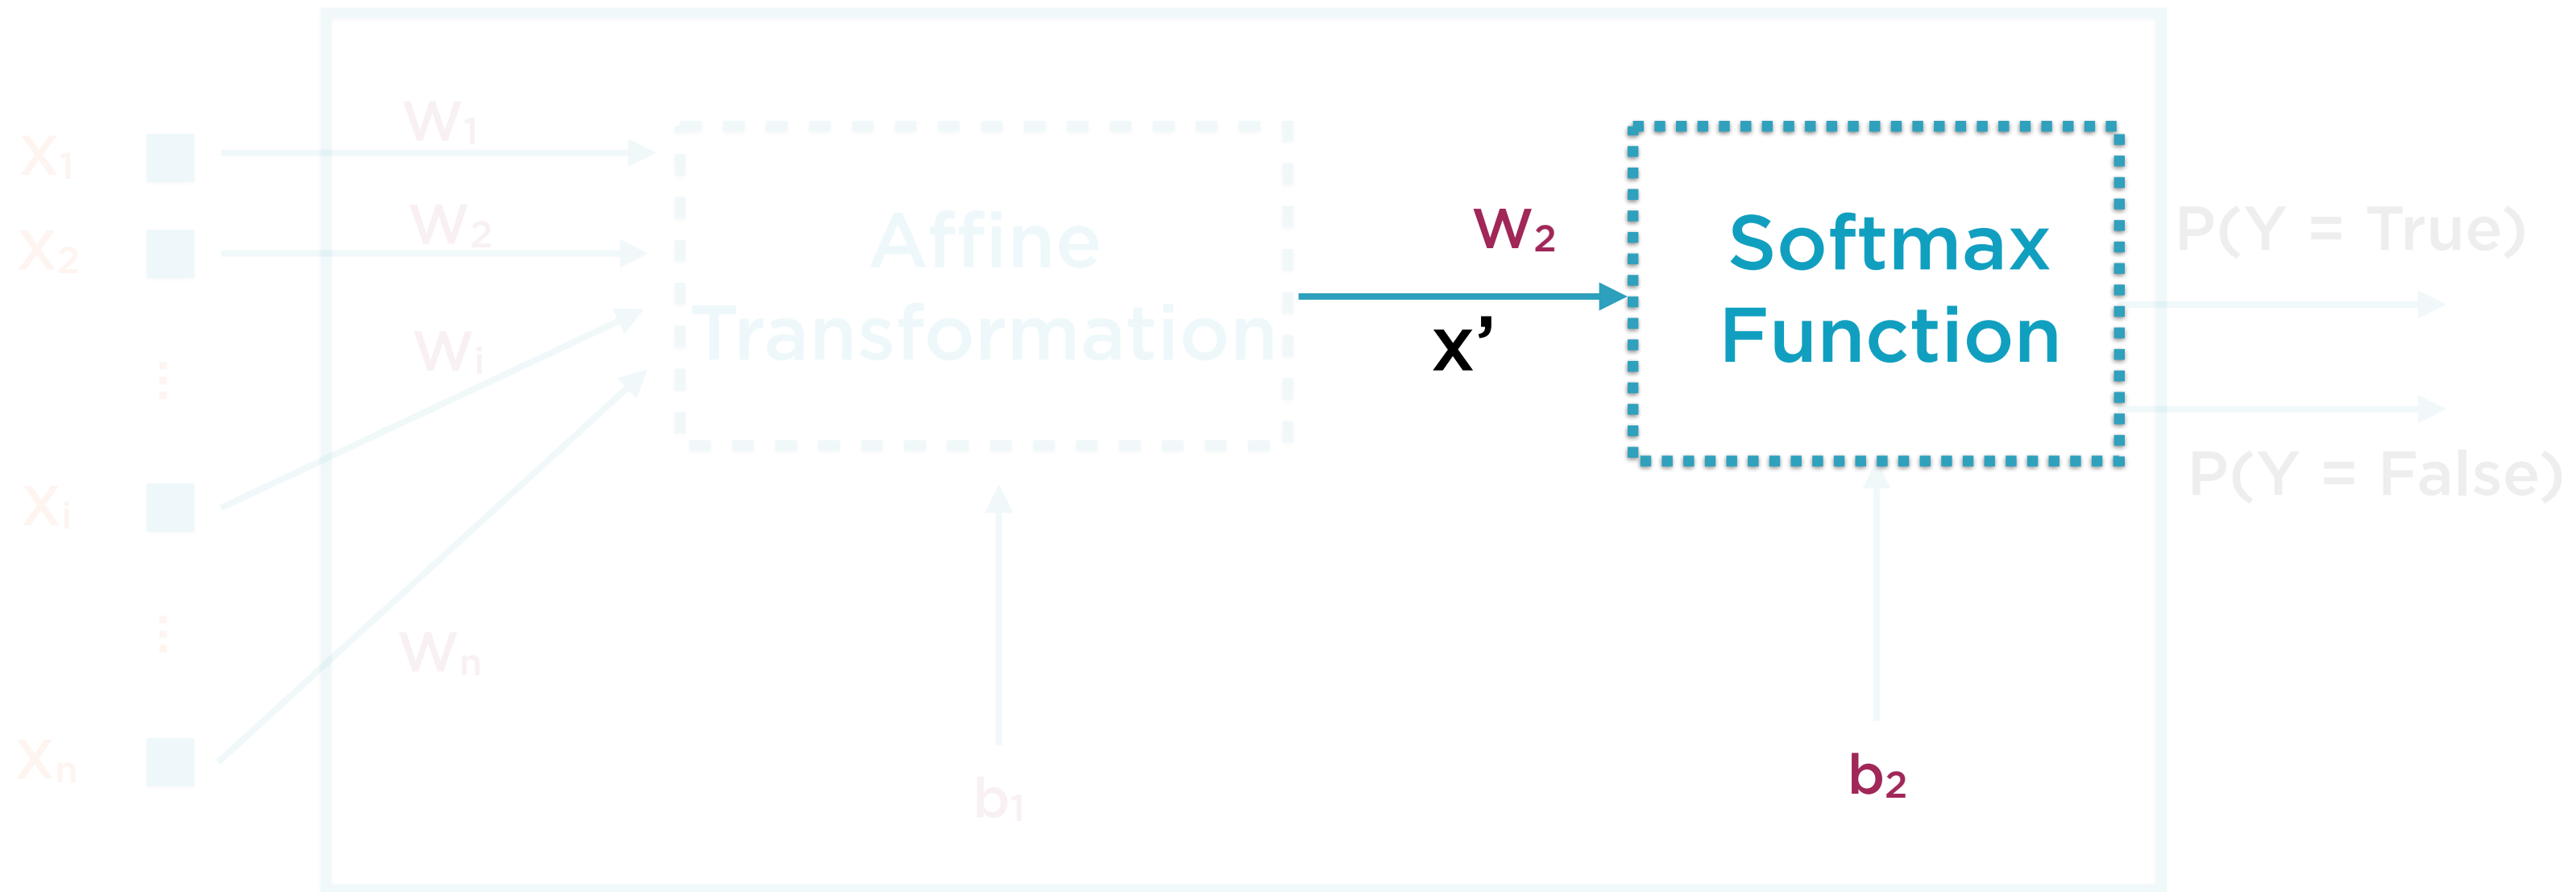
Logistic Regression with One Neuron



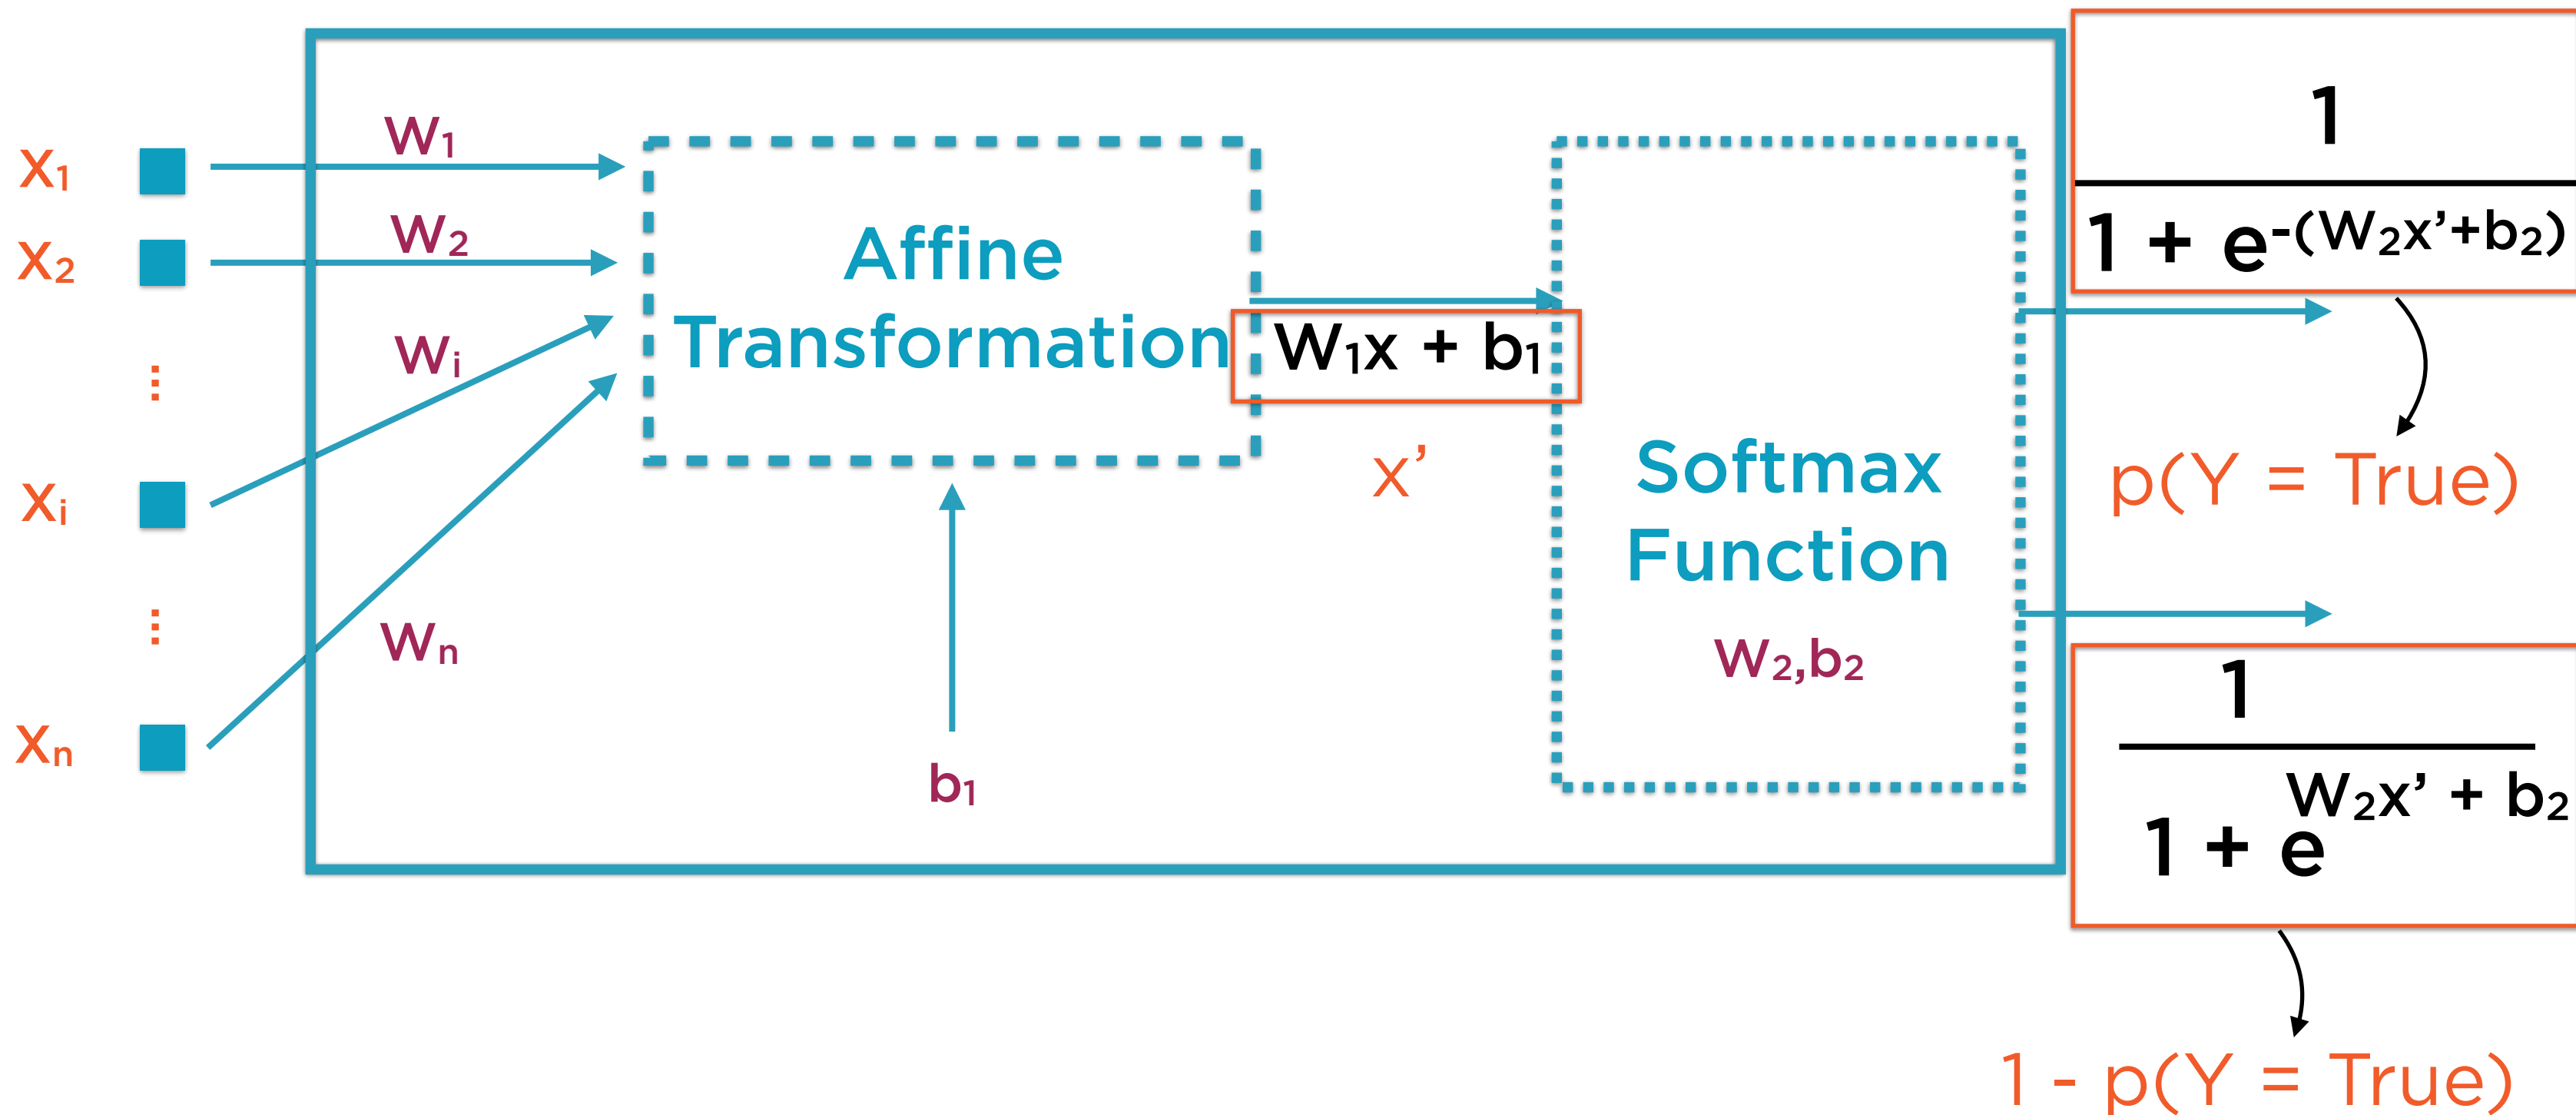
Logistic Regression with One Neuron



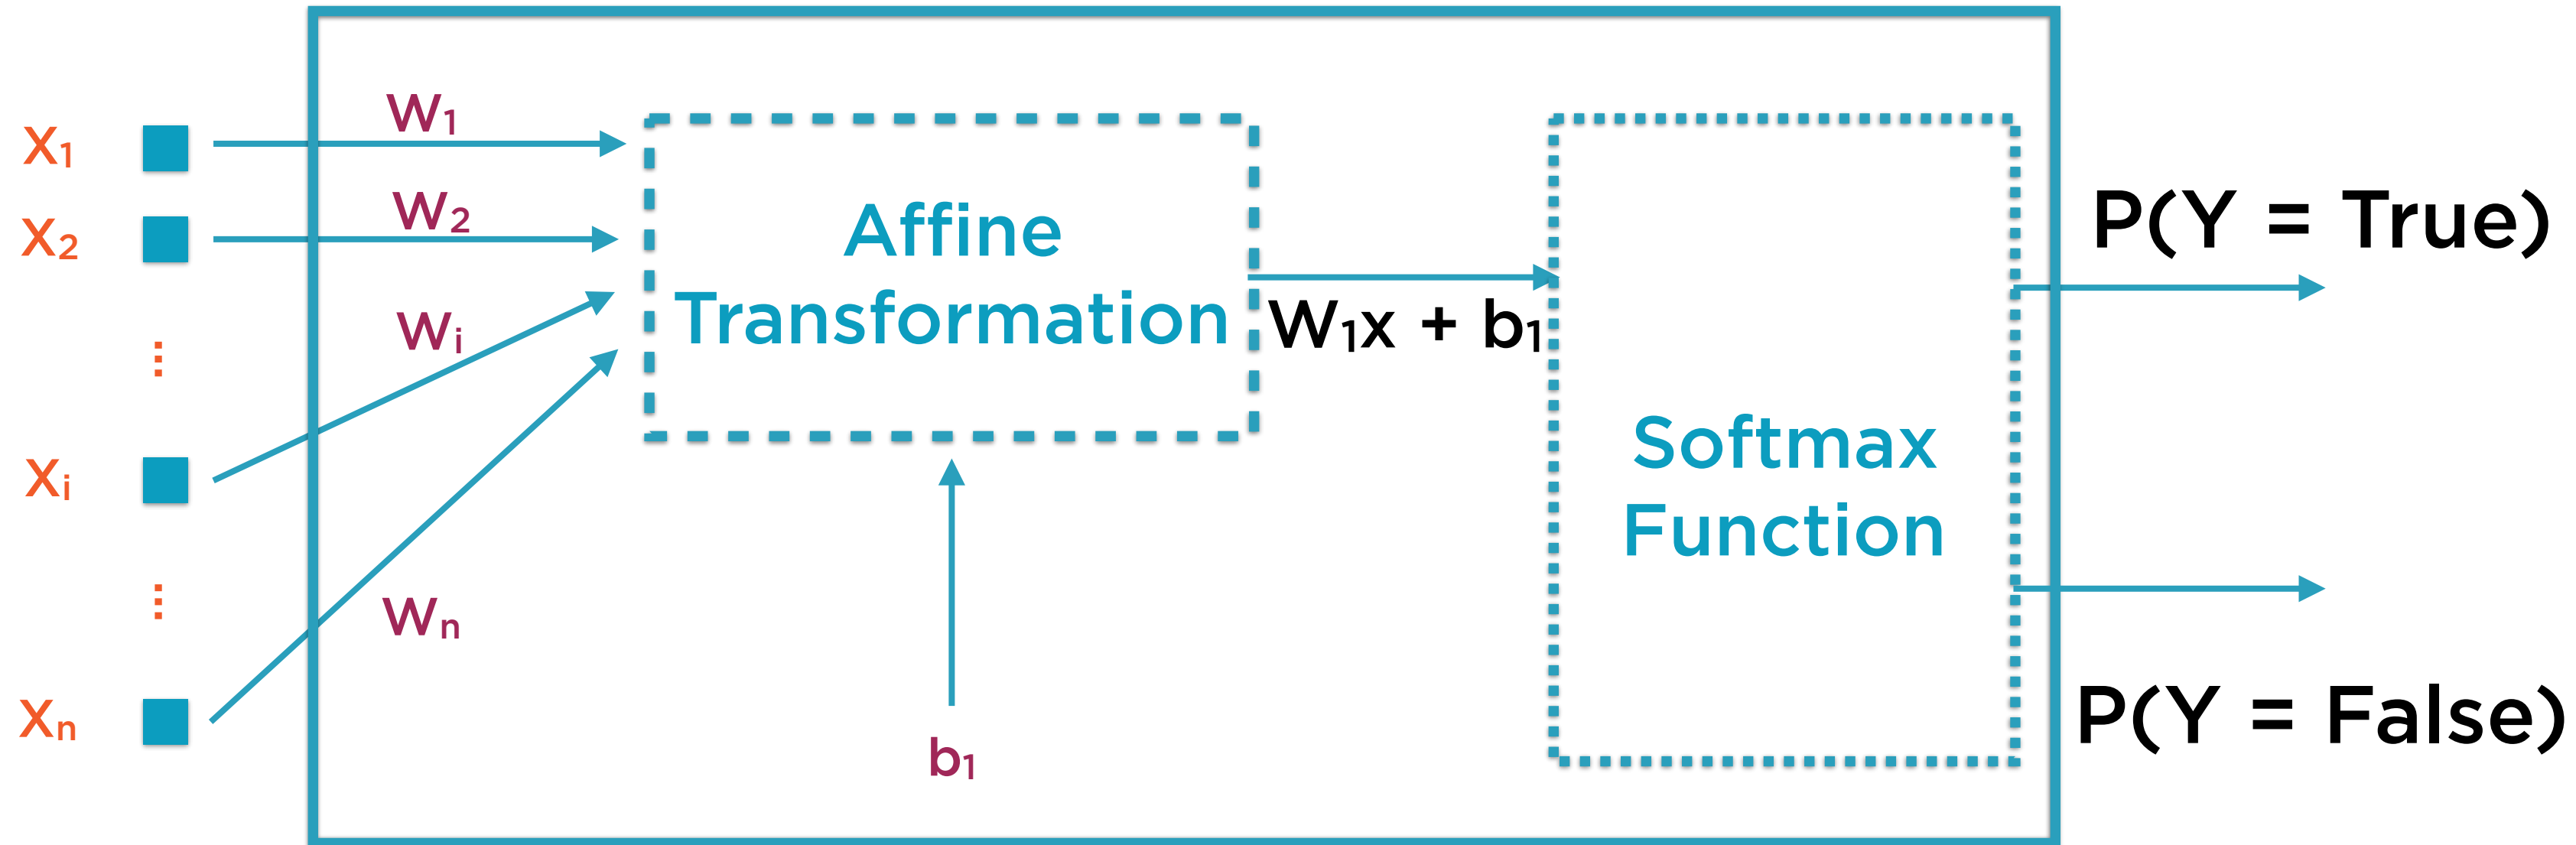
Logistic Regression with One Neuron



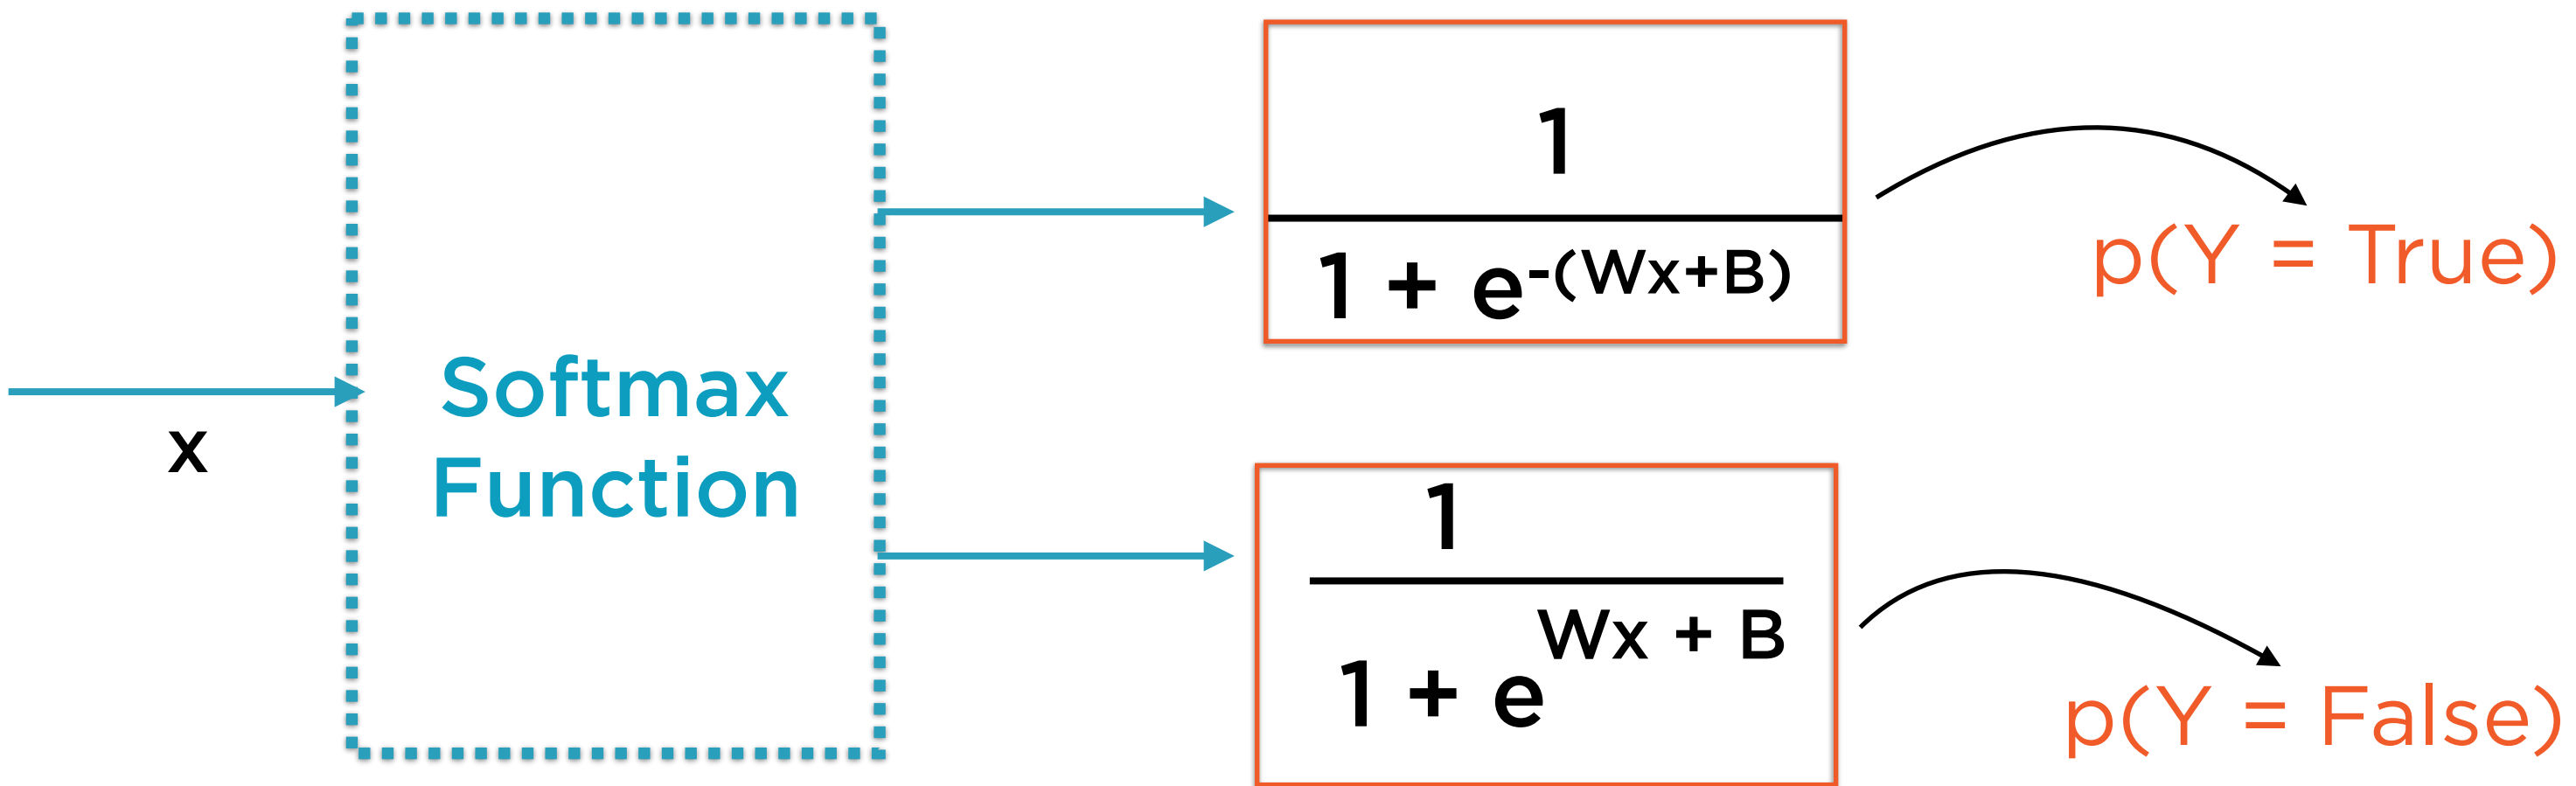
Logistic Regression with One Neuron




Logistic Regression with One Neuron




SoftMax for True/False Classification



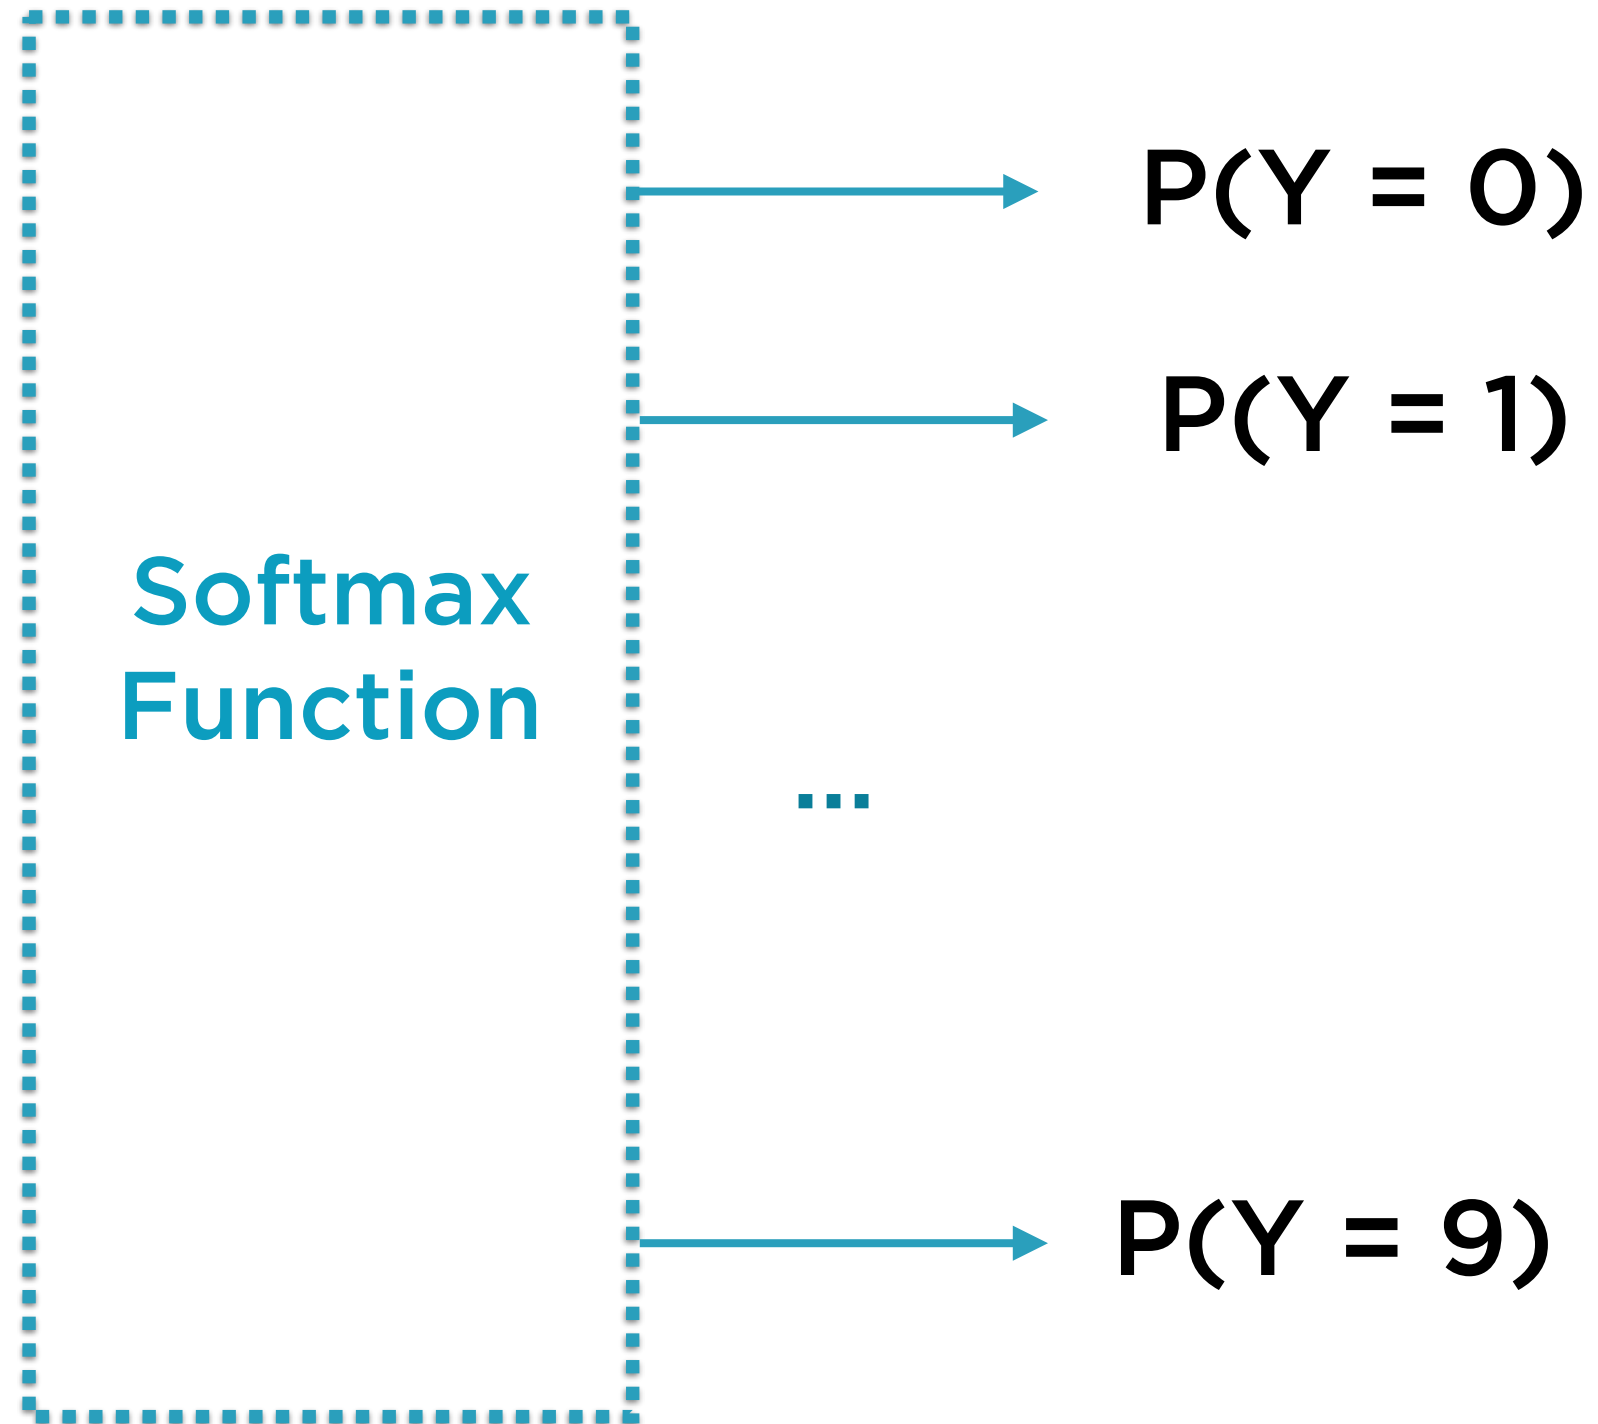
Linear Regression with One Neuron




Logistic Regression with One Neuron



SoftMax for Digit Classification



SoftMax for Digit Classification

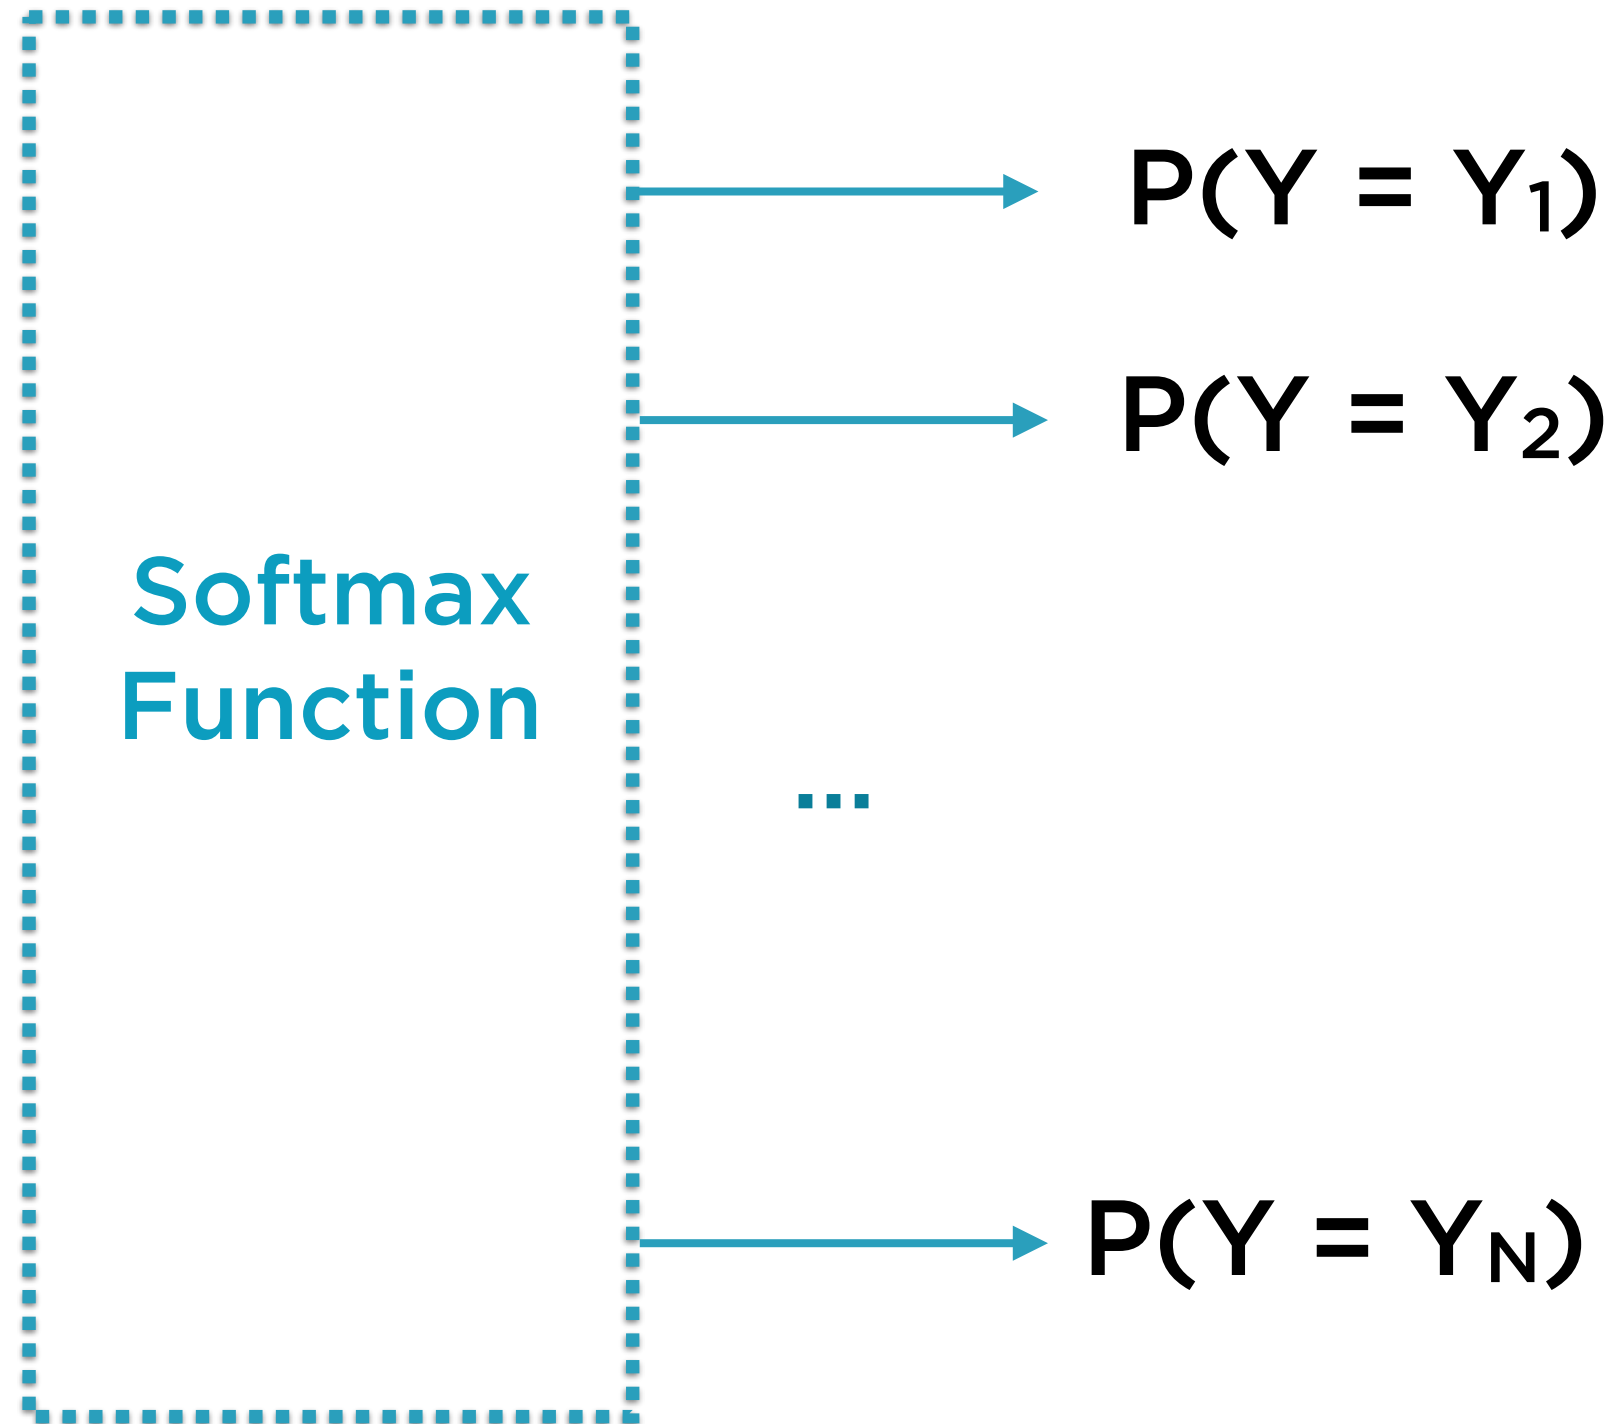


1-dimensional
feature vector


Shape (W) = [1,10]
Shape (b) = [10]

S-Curve

SoftMax N-category Classification



SoftMax N-category Classification




1-dimensional
feature vector

$\text{Shape (W)} = [1, N]$
 $\text{Shape (b)} = [N]$

S-Curve

SoftMax N-category Classification




1-dimensional
feature vector

Shape (W) = $[1, N]$


Shape (b) = $[N]$

S-Curve


SoftMax N-category Classification



SoftMax N-category Classification



SoftMax N-category Classification



**M-dimensional
feature vector**

**Shape (W) = $[M, N]$
Shape (b) = $[N]$**

S-Curve

Logistic Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Cross Entropy
Similarity of distribution

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Softmax activation required


Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline


Logistic Regression in TensorFlow



Linear Regression and MSE




Y




Regression Line:
 $y = A + Bx$

X




The “best fit” line is called the
regression line

Logistic Regression



Set up the Problem



Label GOOG returns as binary (1,0)


Prediction Accuracy

| DATE | ACTUAL | PREDICTED |
|------------|--------|-----------|
| 2005-01-01 | NA | NA |
| 2005-02-01 | 0 | 1 |
| 2005-03-01 | 0 | 0 |
| | | |
| 2017-01-01 | 1 | 1 |
| 2017-02-01 | 1 | 1 |


Compare GOOG's actual labels vs. predicted labels

Intuition: Low Cross Entropy


Y_{actual}



$Y_{\text{predicted}}$




Intuition: Low Cross Entropy



The labels of the two series are in-synch

Intuition: Low Cross Entropy




$-\text{Sum}(Y_{\text{actual}} * \log [Y_{\text{predicted}}])$ will be small


Cross Entropy

Intuition: High Cross Entropy


Y_{actual}



$Y_{\text{predicted}}$




Intuition: High Cross Entropy



The labels of the two series are out-of-synch

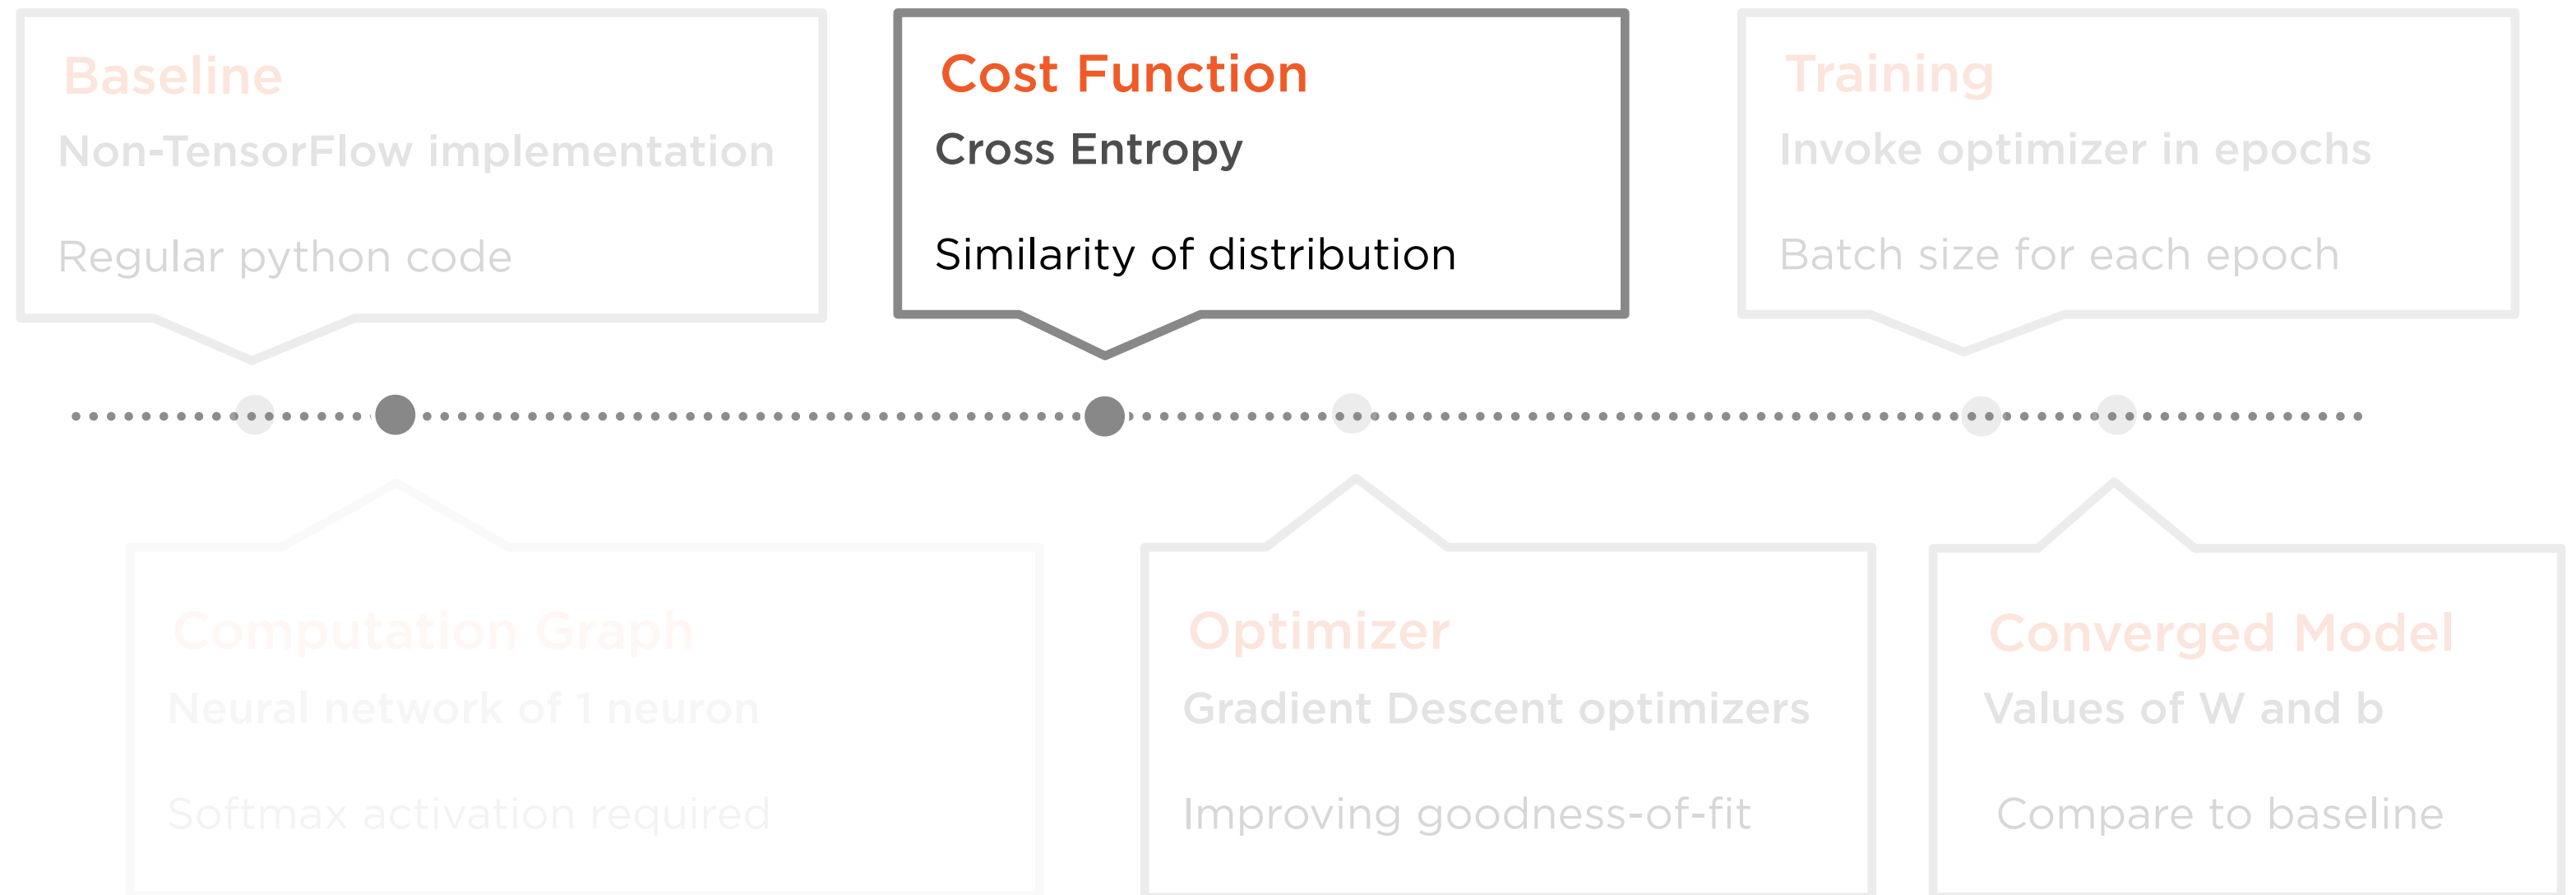
Intuition: High Cross Entropy



$-\text{Sum}(Y_{\text{actual}} * \log [Y_{\text{predicted}}])$ will be large

Cross Entropy

Logistic Regression in TensorFlow



Logistic Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Cross Entropy
Similarity of distribution

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Softmax activation required

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

```
tensorflow.argmax(y, 1)
```

Finding the index of the largest element

Return the index of the largest element of tensor y along dimension k

`tensorflow.argmax(y, 1)`

Tensor



Finding the index of the largest element

Return the index of the largest element of tensor y along dimension k

`tensorflow.argmax(y, 1)`

Dimension



Finding the index of the largest element

Return the index of the largest element of tensor y along dimension k

tf.argmax

Tensor y

Index = 0

1

2

3

4

5

Dimension 0

Dimension 1

| | |
|--|-----|
| | 5 |
| | 15 |
| | 12 |
| | 100 |
| | 74 |
| | 33 |

tf.argmax(y,1)

tf.argmax

Tensor y

Index = 0

1

2

3

4


5

Dimension 0

Dimension 1

| | |
|--|-----|
| | 5 |
| | 15 |
| | 12 |
| | 100 |
| | 74 |
| | 33 |

Return value



tf.argmax(y,1)

tf.argmax

Tensor y

Index = 0

1

2

3

4


5

Dimension 0

Dimension 1

| | |
|--|-----|
| | 5 |
| | 15 |
| | 12 |
| | 100 |
| | 74 |
| | 33 |

Return value



tf.argmax(y,1)

tf.argmax

Tensor y

Index = 0

1

2

3

4

5

Dimension 0

Dimension 1

| | |
|--|-----|
| | 5 |
| | 15 |
| | 12 |
| | 100 |
| | 74 |
| | 33 |

tf.argmax(y, 1)

1

tf.argmax

Tensor y

Index = 0

1

2

3

4

5

Dimension 0

Dimension 1

| | |
|--|-----|
| | 5 |
| | 15 |
| | 12 |
| | 100 |
| | 74 |
| | 33 |

Return value

Largest value

tf.argmax(y,1)

`tf.argmax`

Tensor y

Index = 0

Index = M

Dimension 1

...

Dimension N

| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |
| | | |
| | | |

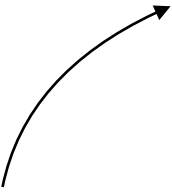
`tf.argmax(y,1)`

```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Actual labels



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Predicted labels

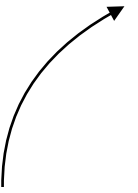
`tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))`



Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

One-hot



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

One-hot Representation

| |
|-------|
| TRUE |
| FALSE |
| FALSE |
| ... |
| |
| TRUE |

Label Vector



| TRUE | FALSE |
|------|-------|
| 1 | 0 |
| 0 | 1 |
| 0 | 1 |
| | |
| | |
| 1 | 0 |

One-hot Label Vector

One-hot $y_{\text{}}$

| |
|-------|
| TRUE |
| FALSE |
| FALSE |
| ... |
| |
| TRUE |

Label Vector



| TRUE | FALSE |
|------|-------|
| 1 | 0 |
| 0 | 1 |
| 0 | 1 |
| | |
| | |
| 1 | 0 |

One-hot Label Vector

$\text{argmax}(y_, 1)$

| 0 | 1 |
|---|---|
| 1 | 0 |
| 0 | 1 |
| 0 | 1 |
| | |
| | |
| 1 | 0 |

One-hot Label Vector



| |
|-----|
| 0 |
| 1 |
| 1 |
| ... |
| |
| 0 |

Index of one-hot element

Predicted labels



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Predicted Probabilities y

| |
|-----------------------|
| P(TRUE) = 0.70 |
| P(TRUE) = 0.44 |
| P(TRUE) = 0.34 |
| ... |
| P(TRUE) = 0.84 |

Probabilities



| P(TRUE) | P(FALSE) |
|----------------|-----------------|
| 0.70 | 0.30 |
| 0.44 | 0.56 |
| 0.34 | 0.66 |
| | |
| 0.84 | 0.16 |

Softmax Output

Predicted Probabilities y

| |
|-----------------------|
| P(TRUE) = 0.70 |
| P(TRUE) = 0.44 |
| P(TRUE) = 0.34 |
| ... |
| P(TRUE) = 0.84 |

Probabilities



| P(TRUE) | P(FALSE) |
|----------------|-----------------|
| 0.70 | 0.30 |
| 0.44 | 0.56 |
| 0.34 | 0.66 |
| | |
| 0.84 | 0.16 |

Softmax Output

Each row
sums to 1

Predicted Probabilities y

| |
|-----------------------|
| P(TRUE) = 0.70 |
| P(TRUE) = 0.44 |
| P(TRUE) = 0.34 |
| ... |
| P(TRUE) = 0.84 |

Probabilities




| P(TRUE) | P(FALSE) |
|-------------|-------------|
| 0.70 | 0.30 |
| 0.44 | 0.56 |
| 0.34 | 0.66 |
| | |
| 0.84 | 0.16 |

Softmax Output

Rule of 50% in Binary Classification



Mammal




Fish

Probability of whales being Fish $< 50\%$

$\text{argmax}(y,1)$

| P(TRUE) | P(FALSE) |
|---------|----------|
| 0.70 | 0.30 |
| 0.44 | 0.56 |
| 0.34 | 0.66 |
| | |
| | |
| 0.84 | 0.16 |

Softmax Output



| |
|-----|
| 0 |
| 1 |
| 1 |
| ... |
| |
| 0 |

$\text{argmax}(y,1)$

One-hot Vectors with Digit Classes

| |
|-----|
| 0 |
| 1 |
| ... |
| ... |
| 9 |

Actual Digits



| | | | |
|---|---|-----|---|
| 0 | 1 | ... | 9 |
| 1 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 |
| | | | |
| | | | |
| 0 | 0 | 0 | 1 |

One-hot Label Vectors

y_: One-hot Vectors with Digit Classes

| |
|-----|
| 0 |
| 1 |
| ... |
| ... |
| 9 |

Actual Digits



| 0 | 1 | ... | 9 |
|---|---|-----|---|
| 1 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 |
| | | | |
| | | | |
| | | | |
| 0 | 0 | 0 | 1 |

One-hot Label Vectors

$\text{argmax}(y_{:,1})$

0

1

...

9

| | | | |
|---|---|---|---|
| 1 | 0 | 0 | 0 |
| 0 | 1 | 0 | 0 |
| | | | |
| | | | |
| | | | |
| 0 | 0 | 0 | 1 |

One-hot Label Vectors



| |
|-----|
| 0 |
| 1 |
| ... |
| ... |
| |
| 9 |

$\text{argmax}(y_{:,1})$

Digit Classification

| $P(X=0)$ | $P(X=1)$ | ... | $P(X=9)$ |
|----------|----------|-----|----------|
| 0.70 | 0.30 | | |
| 0.44 | 0.56 | | |
| | | | |
| | | | |
| | | | |
| | | | 0.66 |

Softmax Output




| |
|-----|
| 0 |
| 1 |
| 1 |
| ... |
| |
| 9 |

$\text{argmax}(y,1)$

y: Predicted Probabilities

| P(X=0) | P(X=1) | ... | P(X=9) |
|--------|--------|-----|--------|
| 0.70 | 0.30 | | |
| 0.44 | 0.56 | | |
| | | | |
| | | | |
| | | | |
| | | | 0.66 |

Softmax Output



| |
|---|
| 0 |
| 1 |
| |
| |
| |
| 9 |

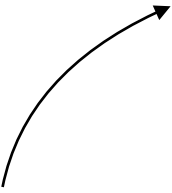
argmax(y,1)

```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Actual labels



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Predicted labels

`tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))`




Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

Tensor of actual labels

Tensor of predicted labels




```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

List of True, False values



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`

True: Correct prediction

False: Incorrect prediction



```
tf.equal(tf.argmax(y_, 1), tf.argmax(y, 1))
```

Two invocations of `tf.argmax`

Once on actual labels `y_`, once on predicted values `y`