

Building Generalized Linear Models Using Estimators



Vitthal Srinivasan

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview


Estimators are cookie-cutter TensorFlow APIs for many standard problems

These high-level APIs reside in `tf.learn` and `tf.contrib.learn`

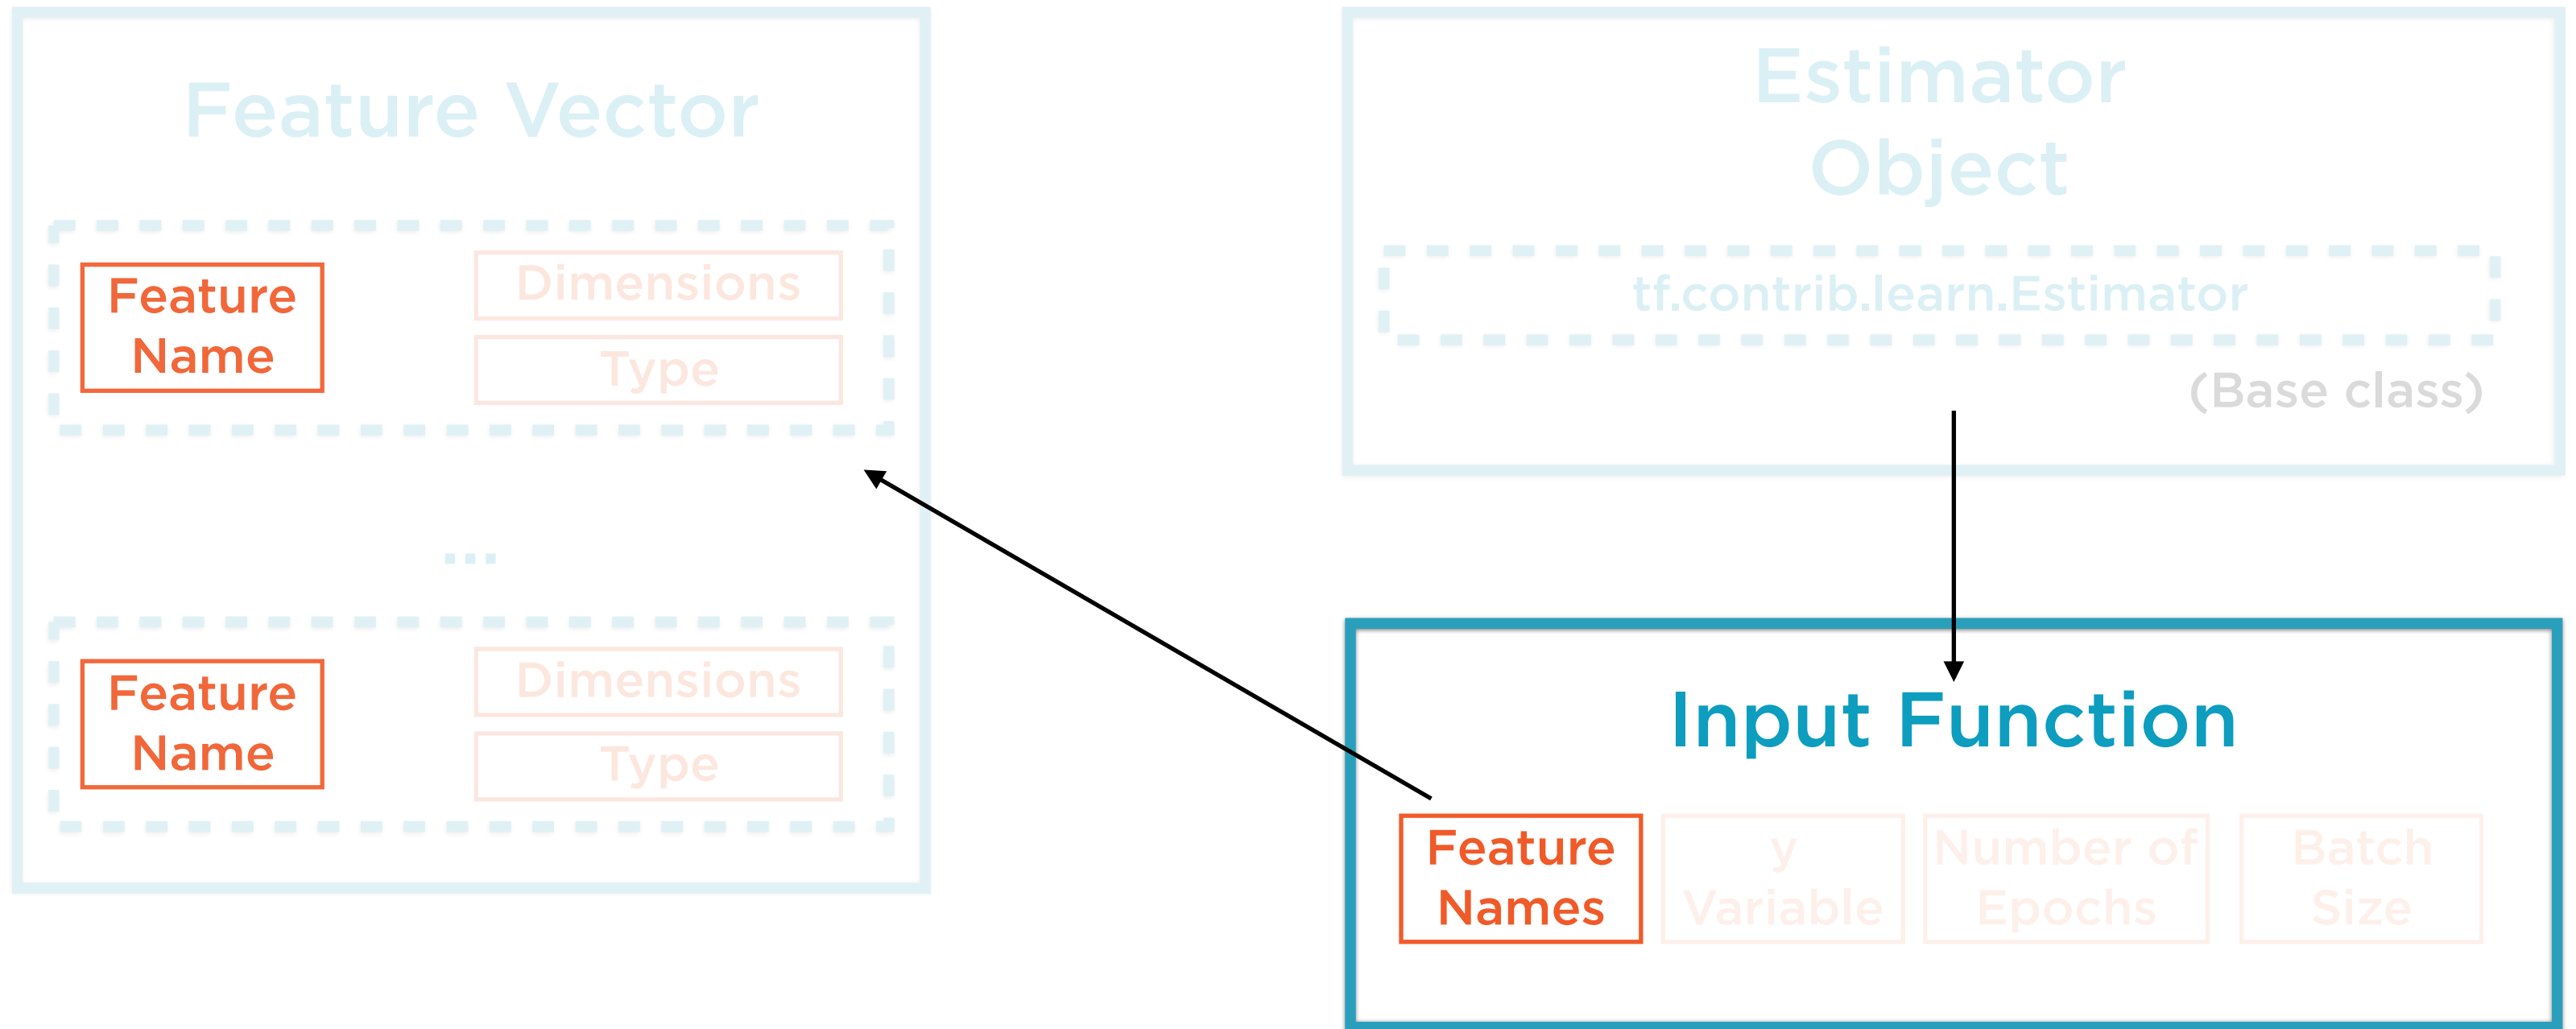
Estimators can be extended by plugging custom models into a base class

That extension relies on composition rather than inheritance

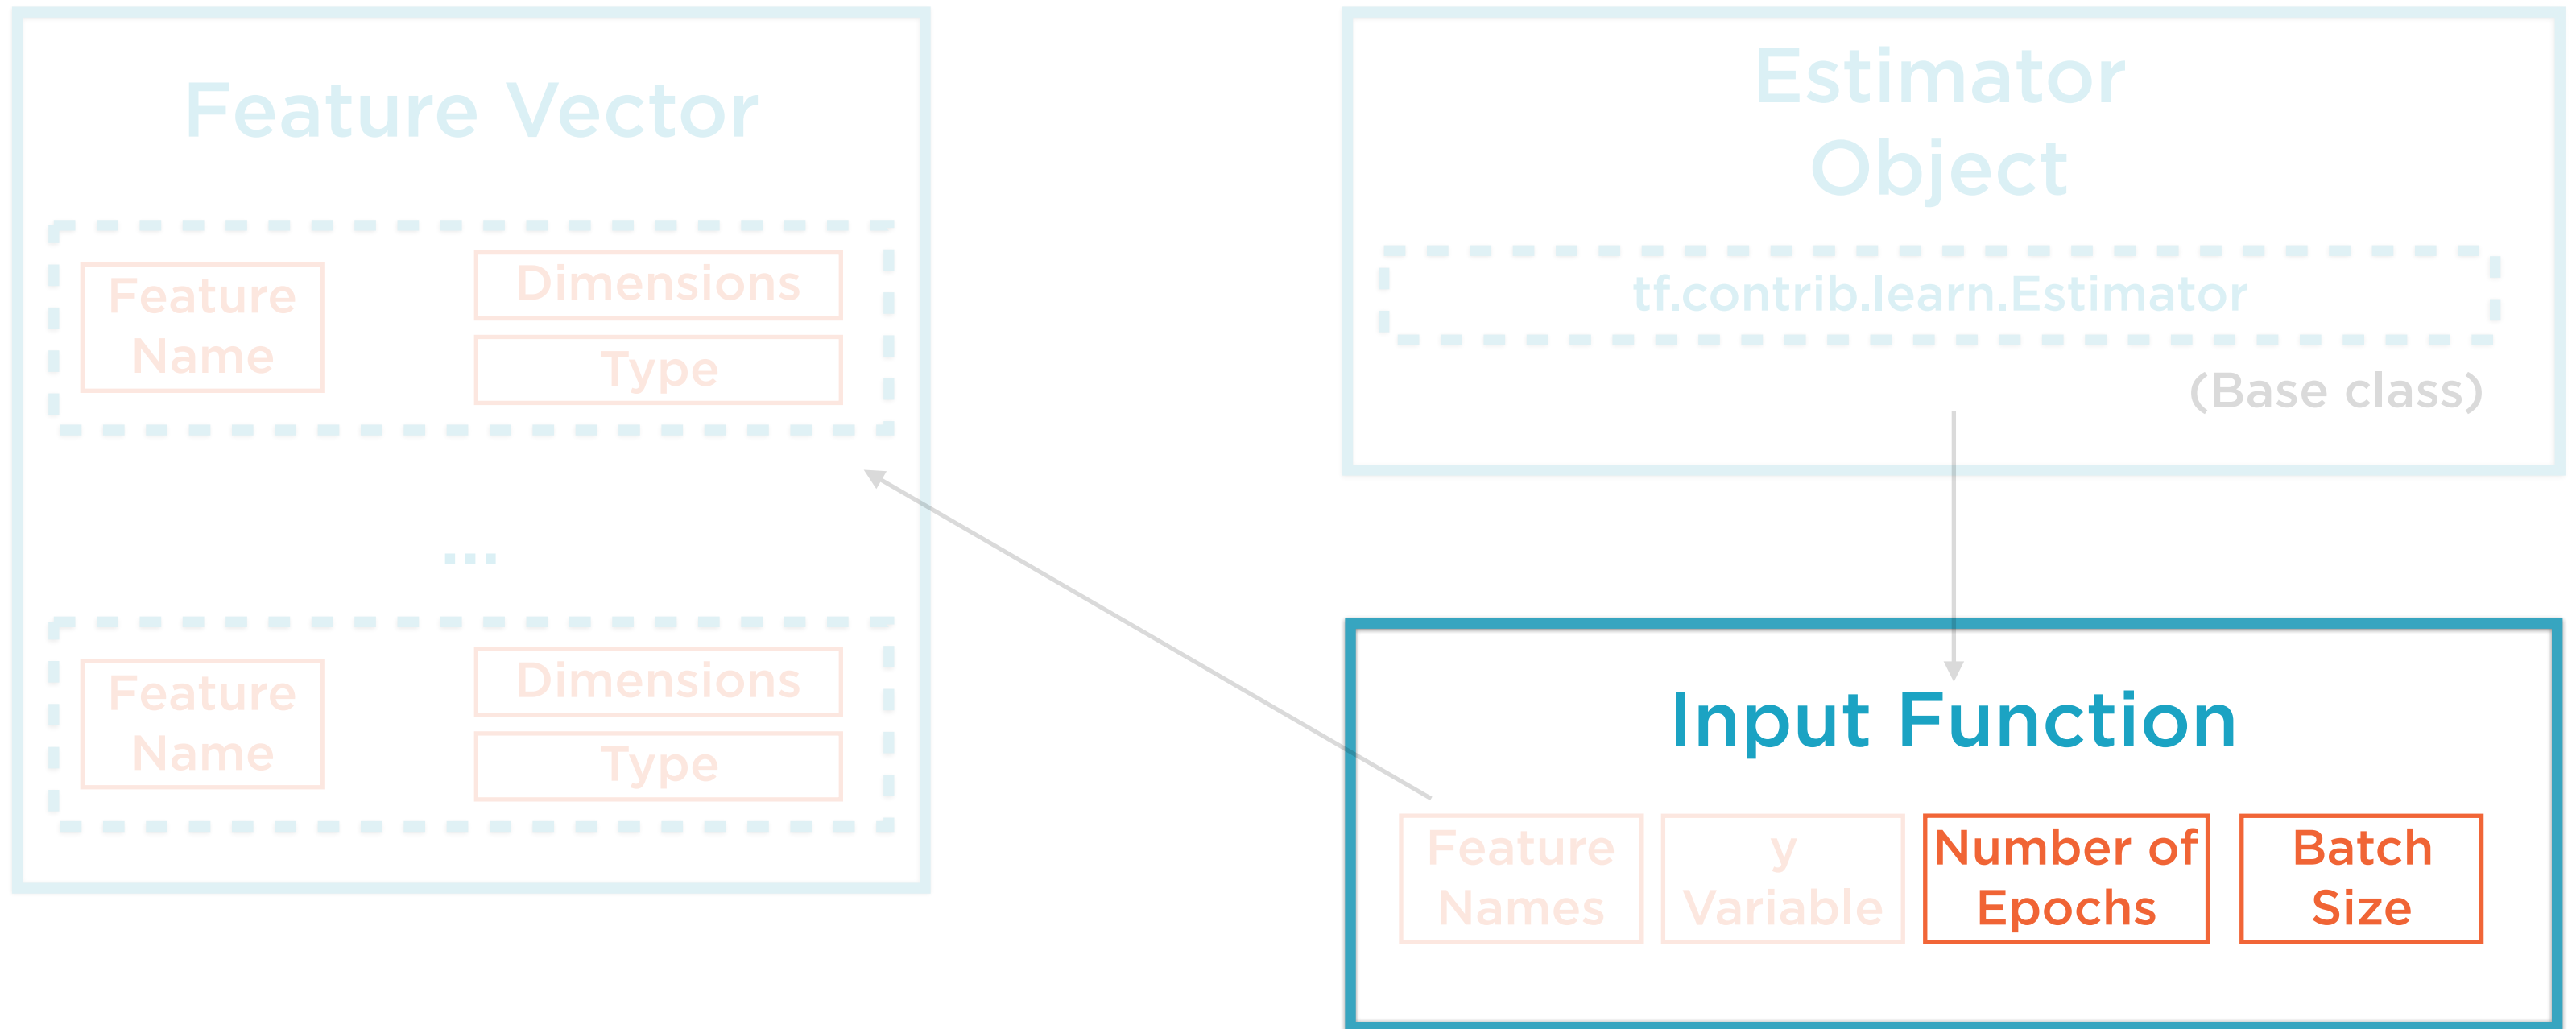
How Estimators Work



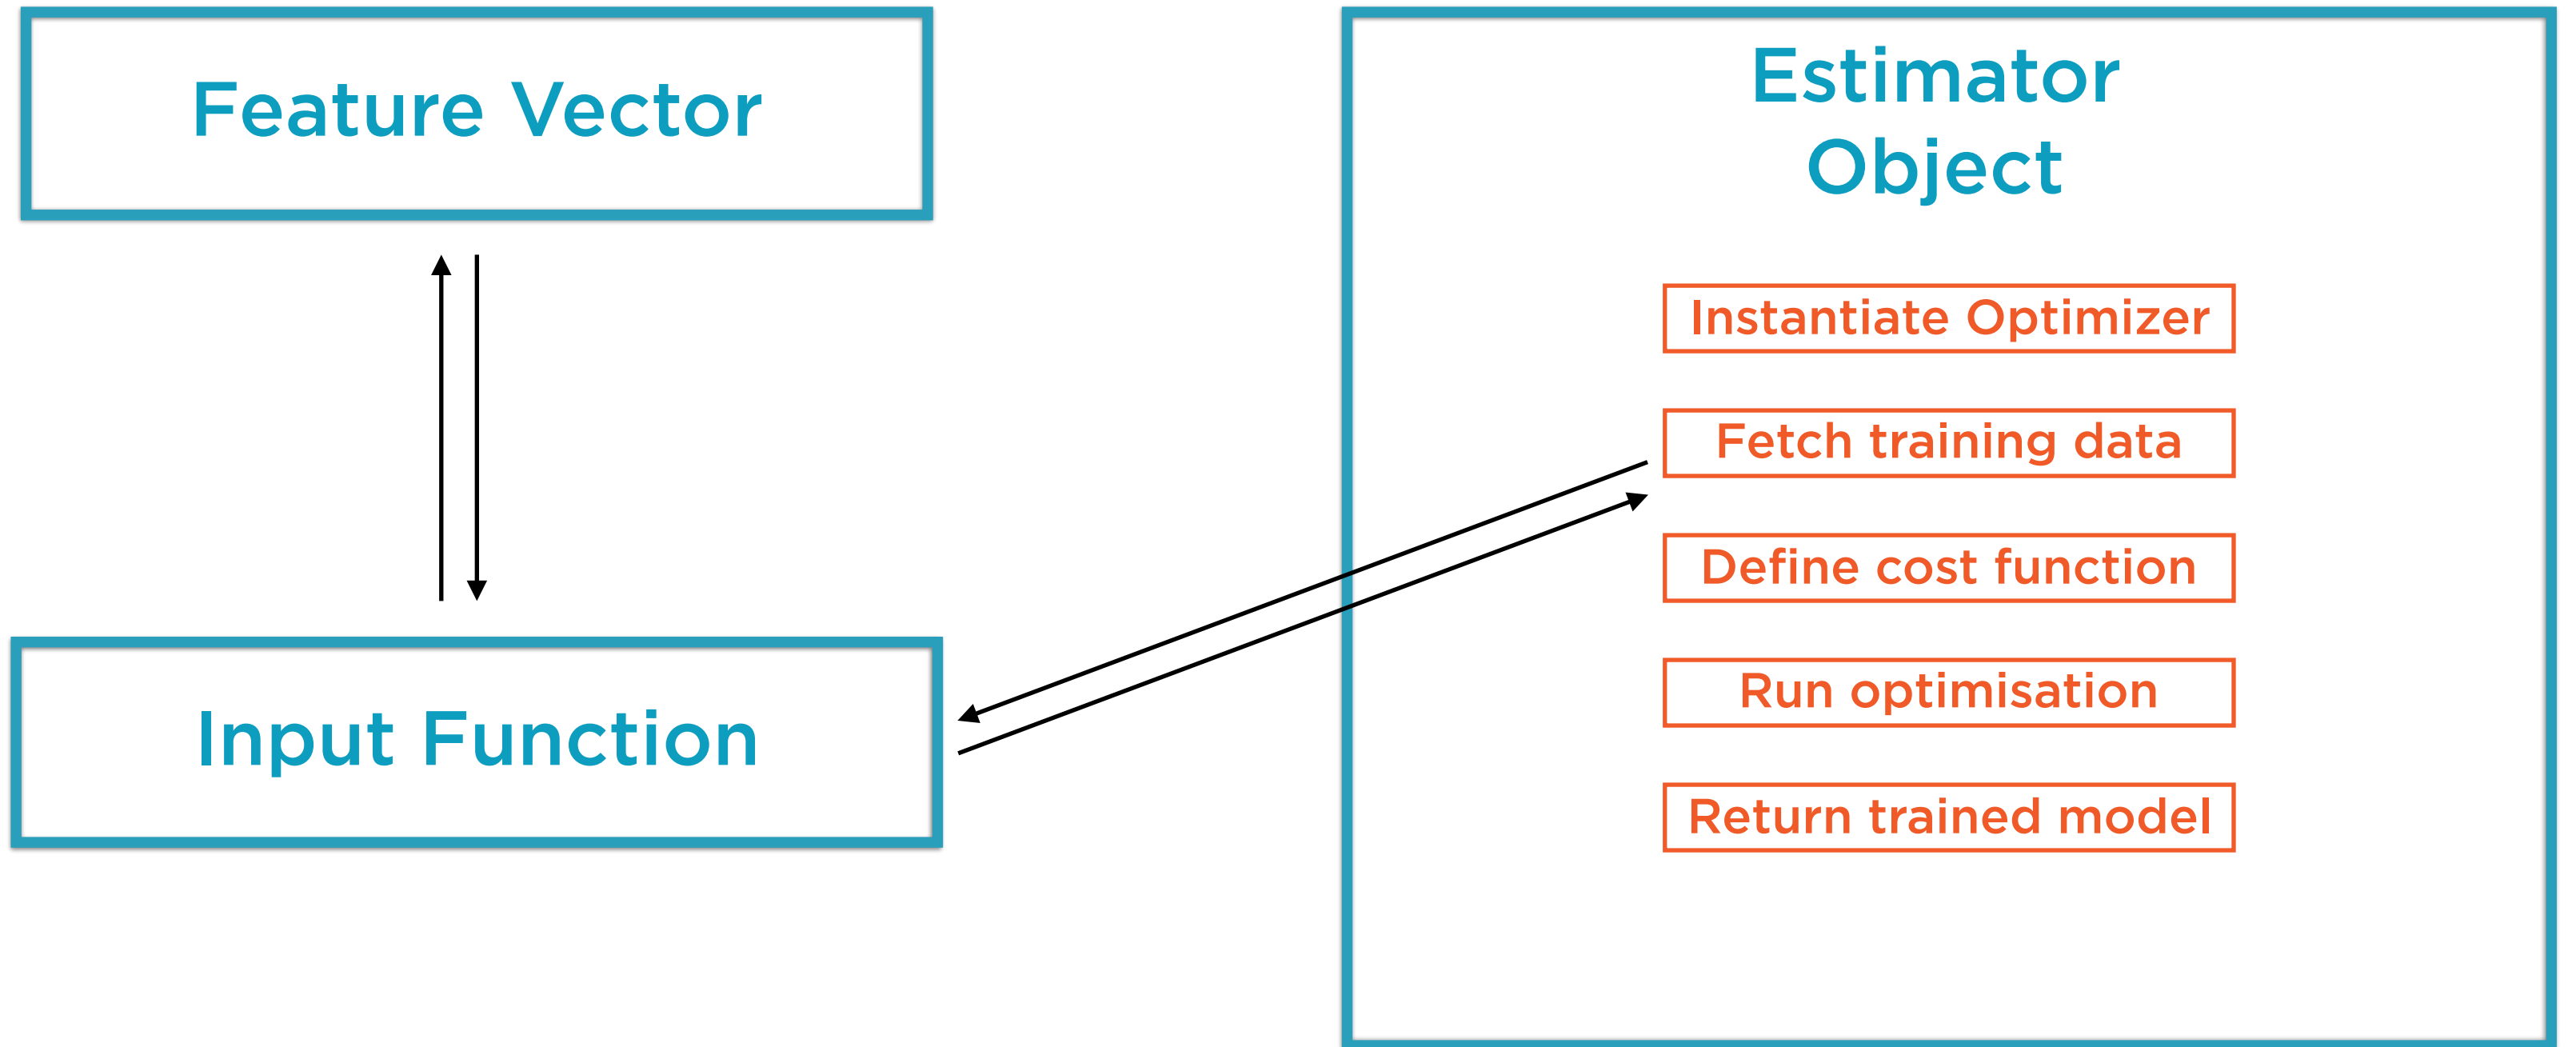
How Estimators Work




How Estimators Work




How Estimators Work



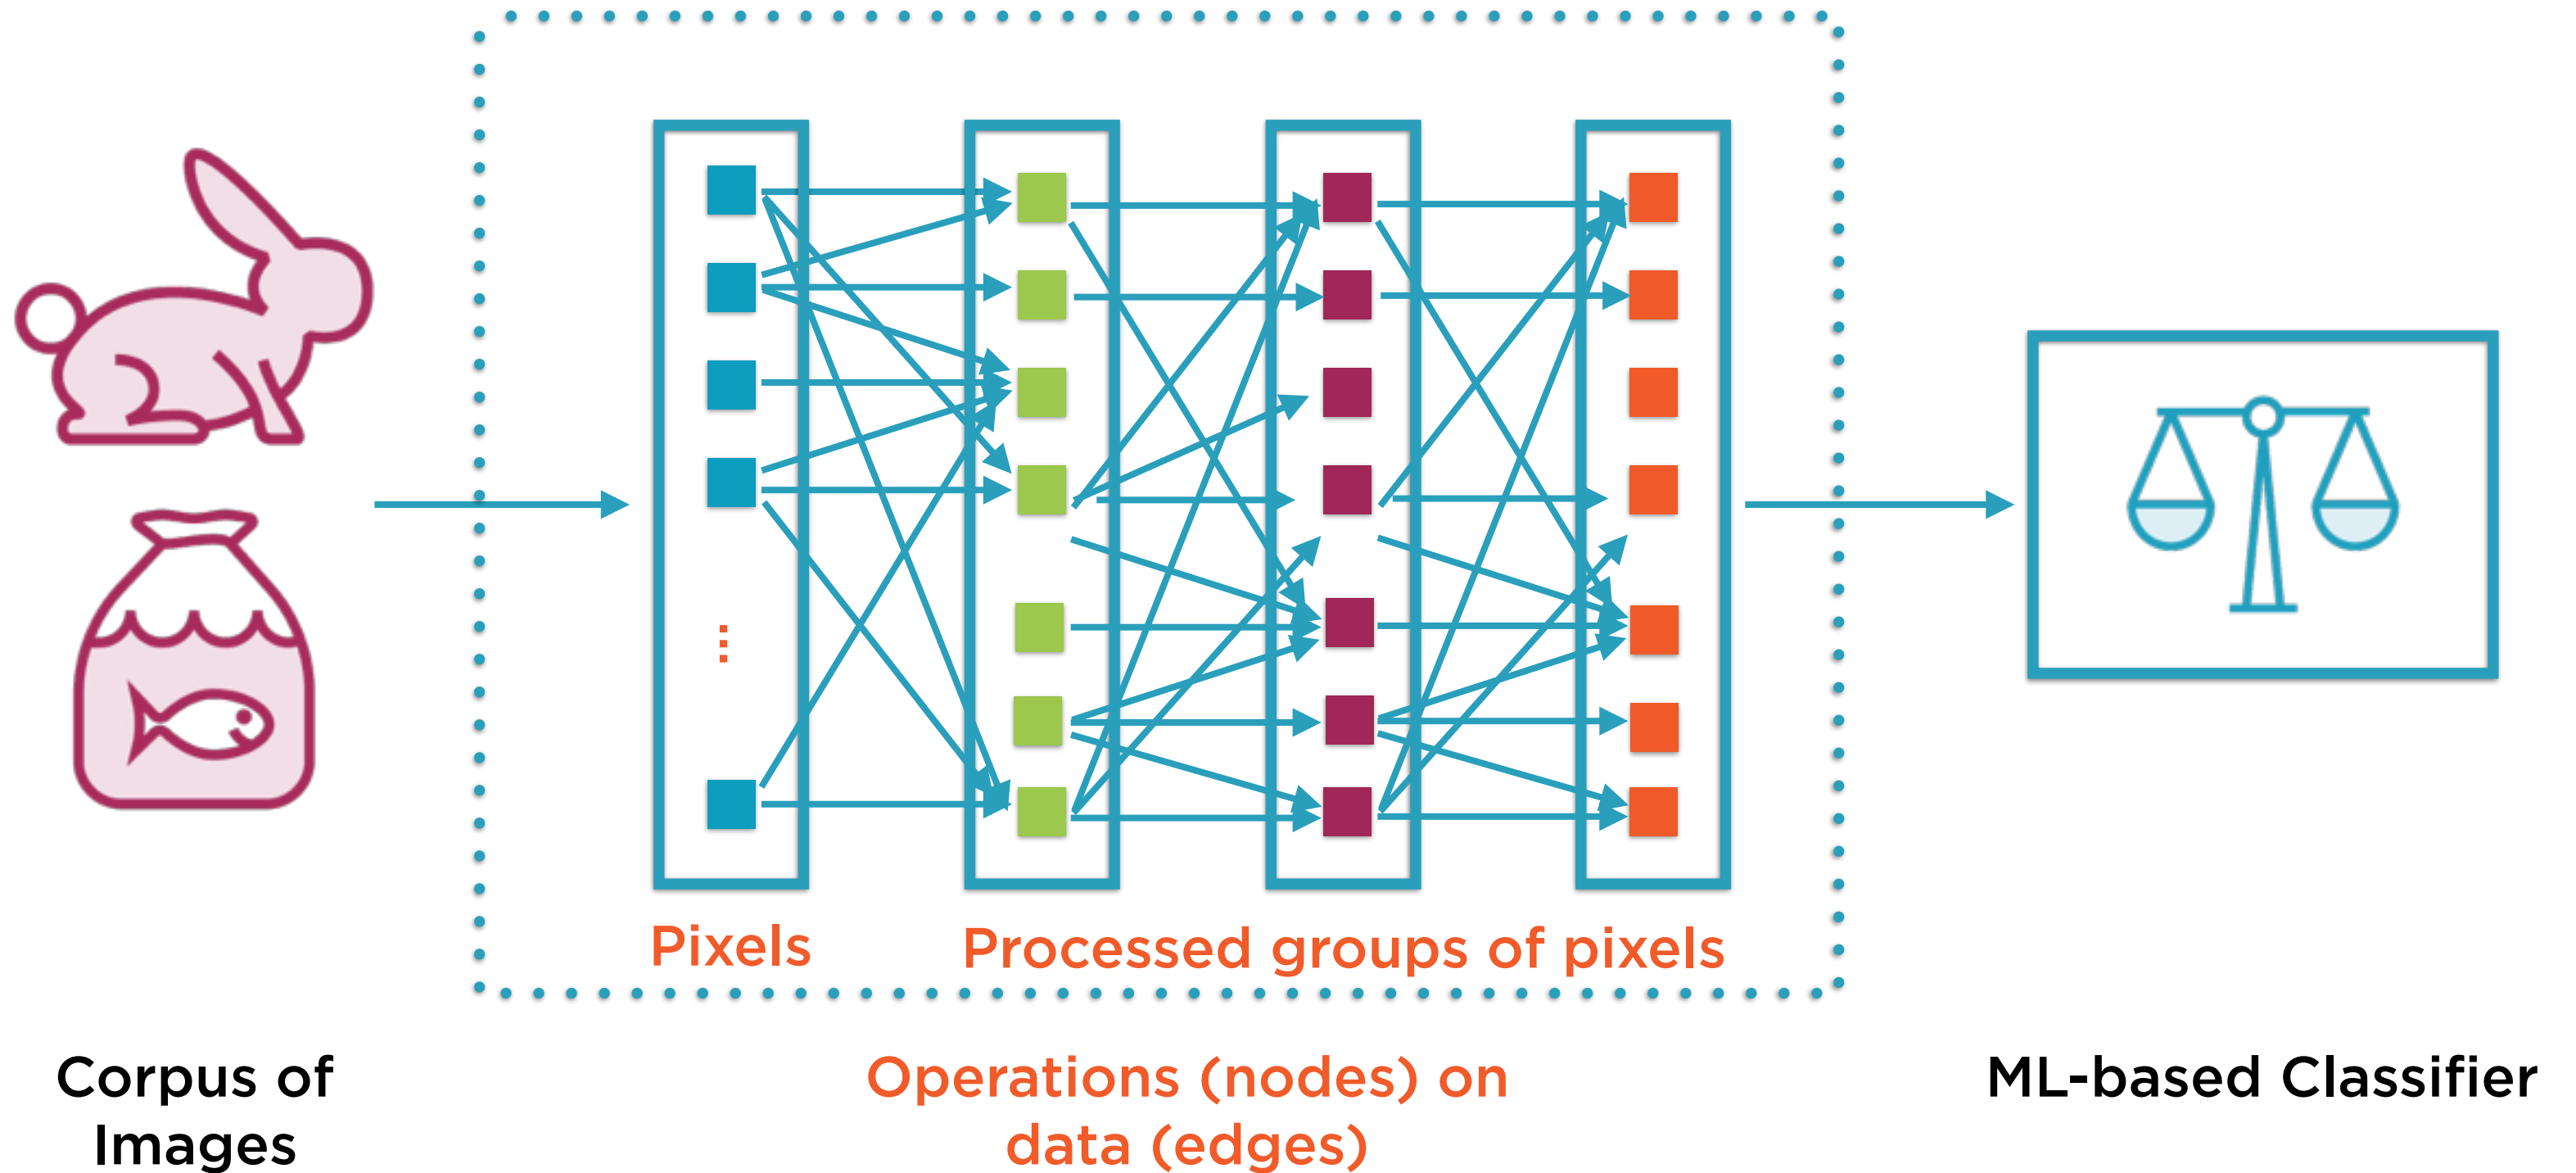
How Estimators Work



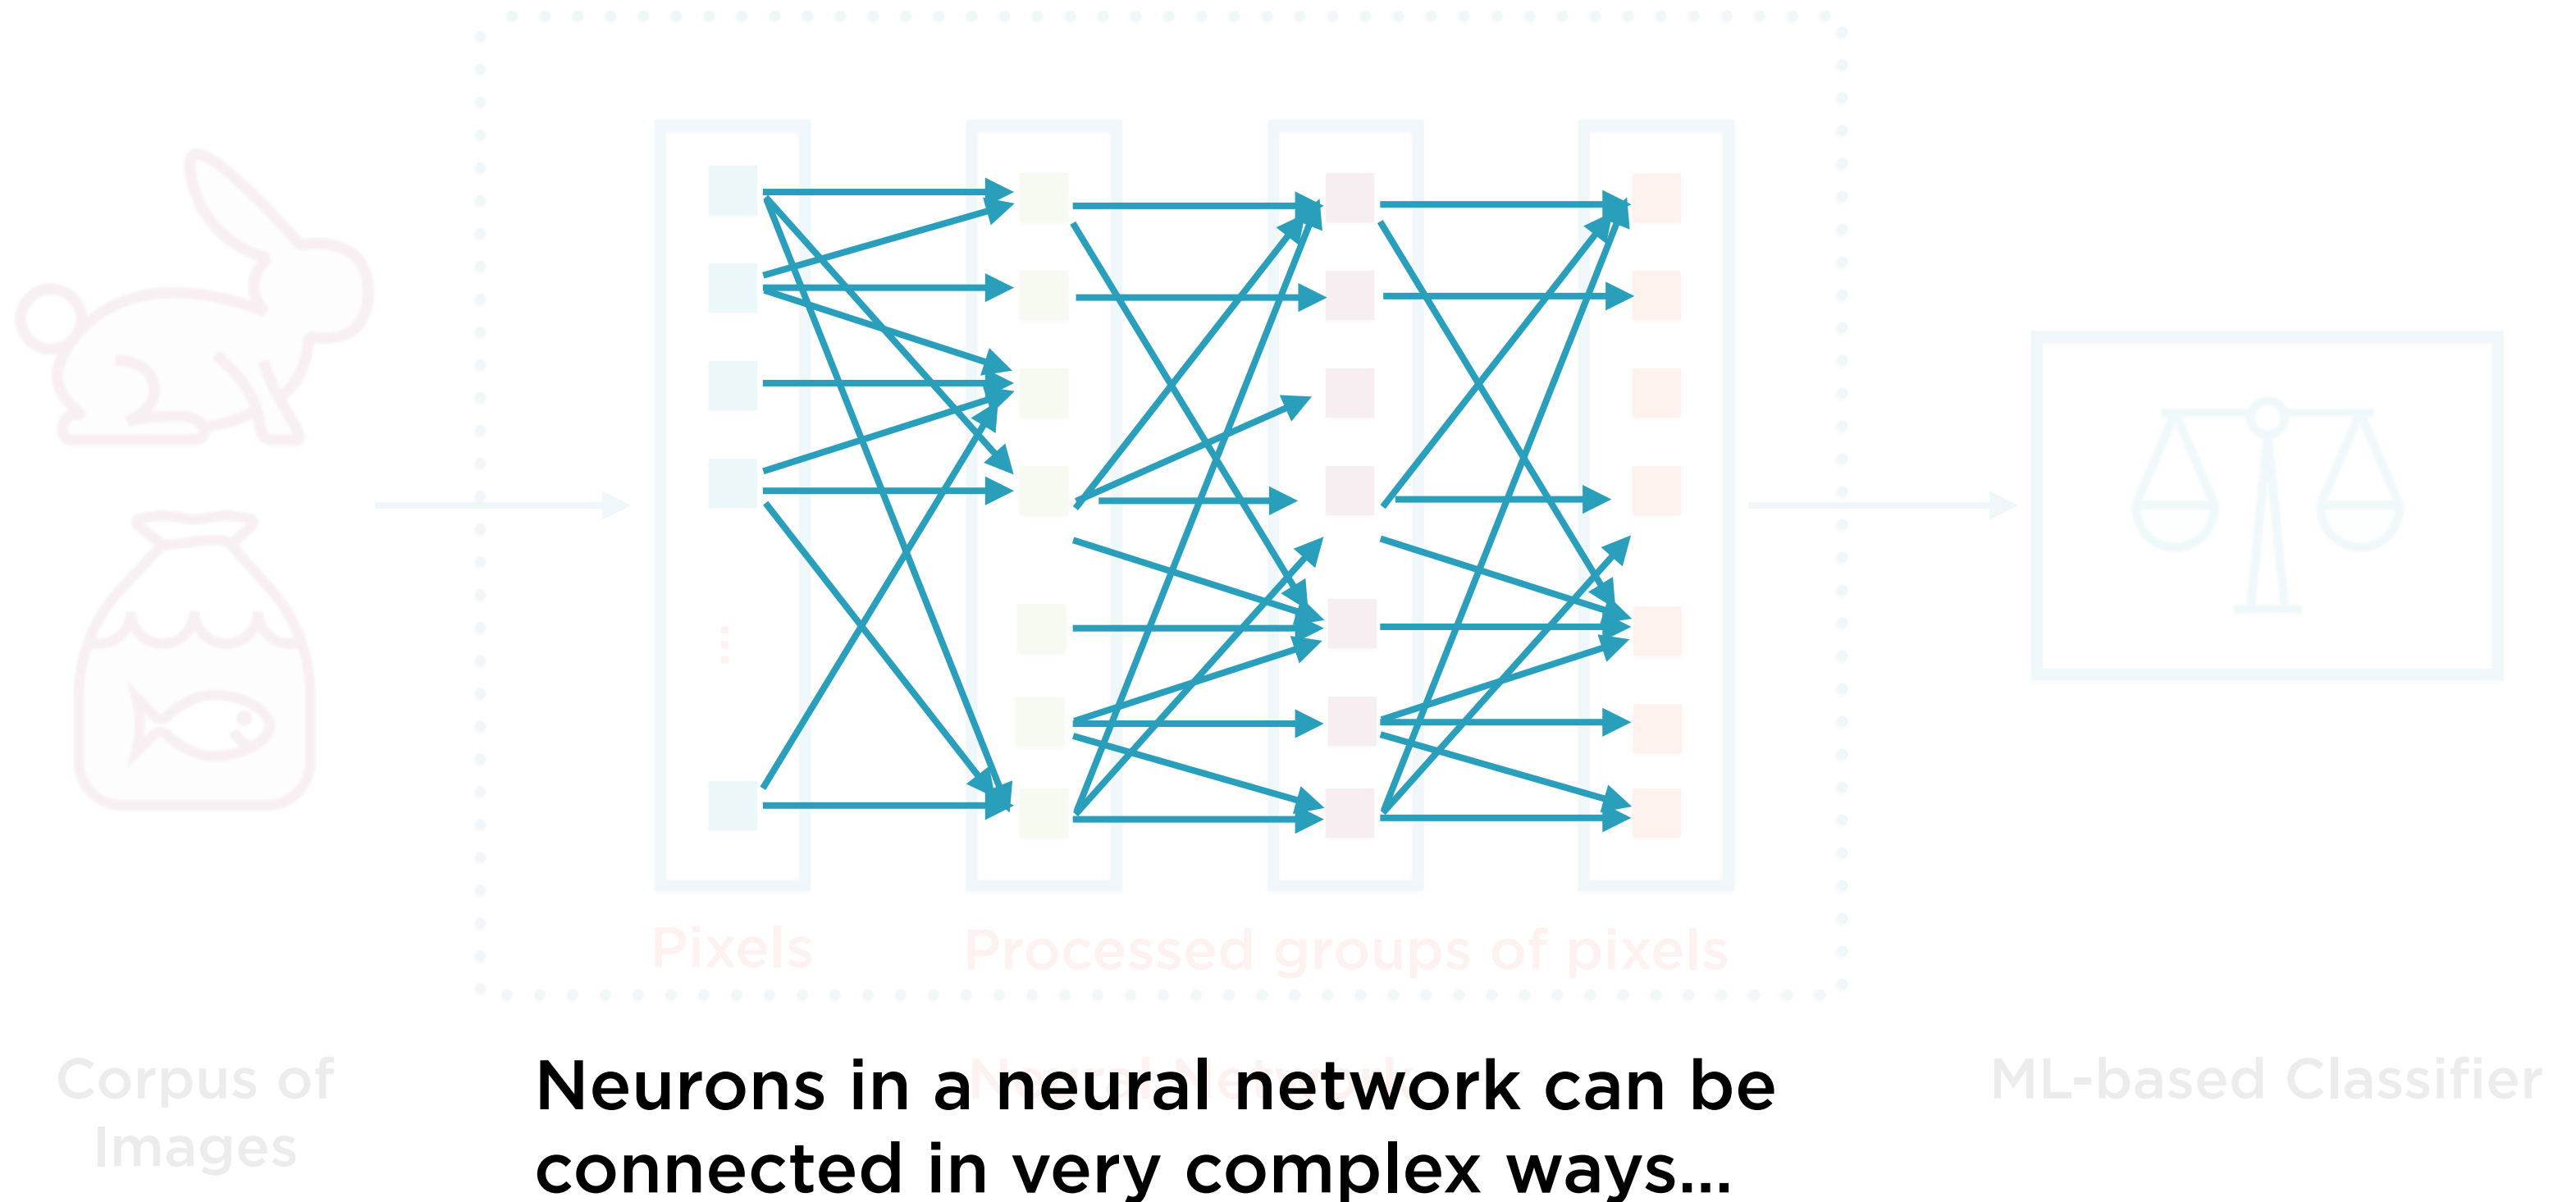
How Estimators Work



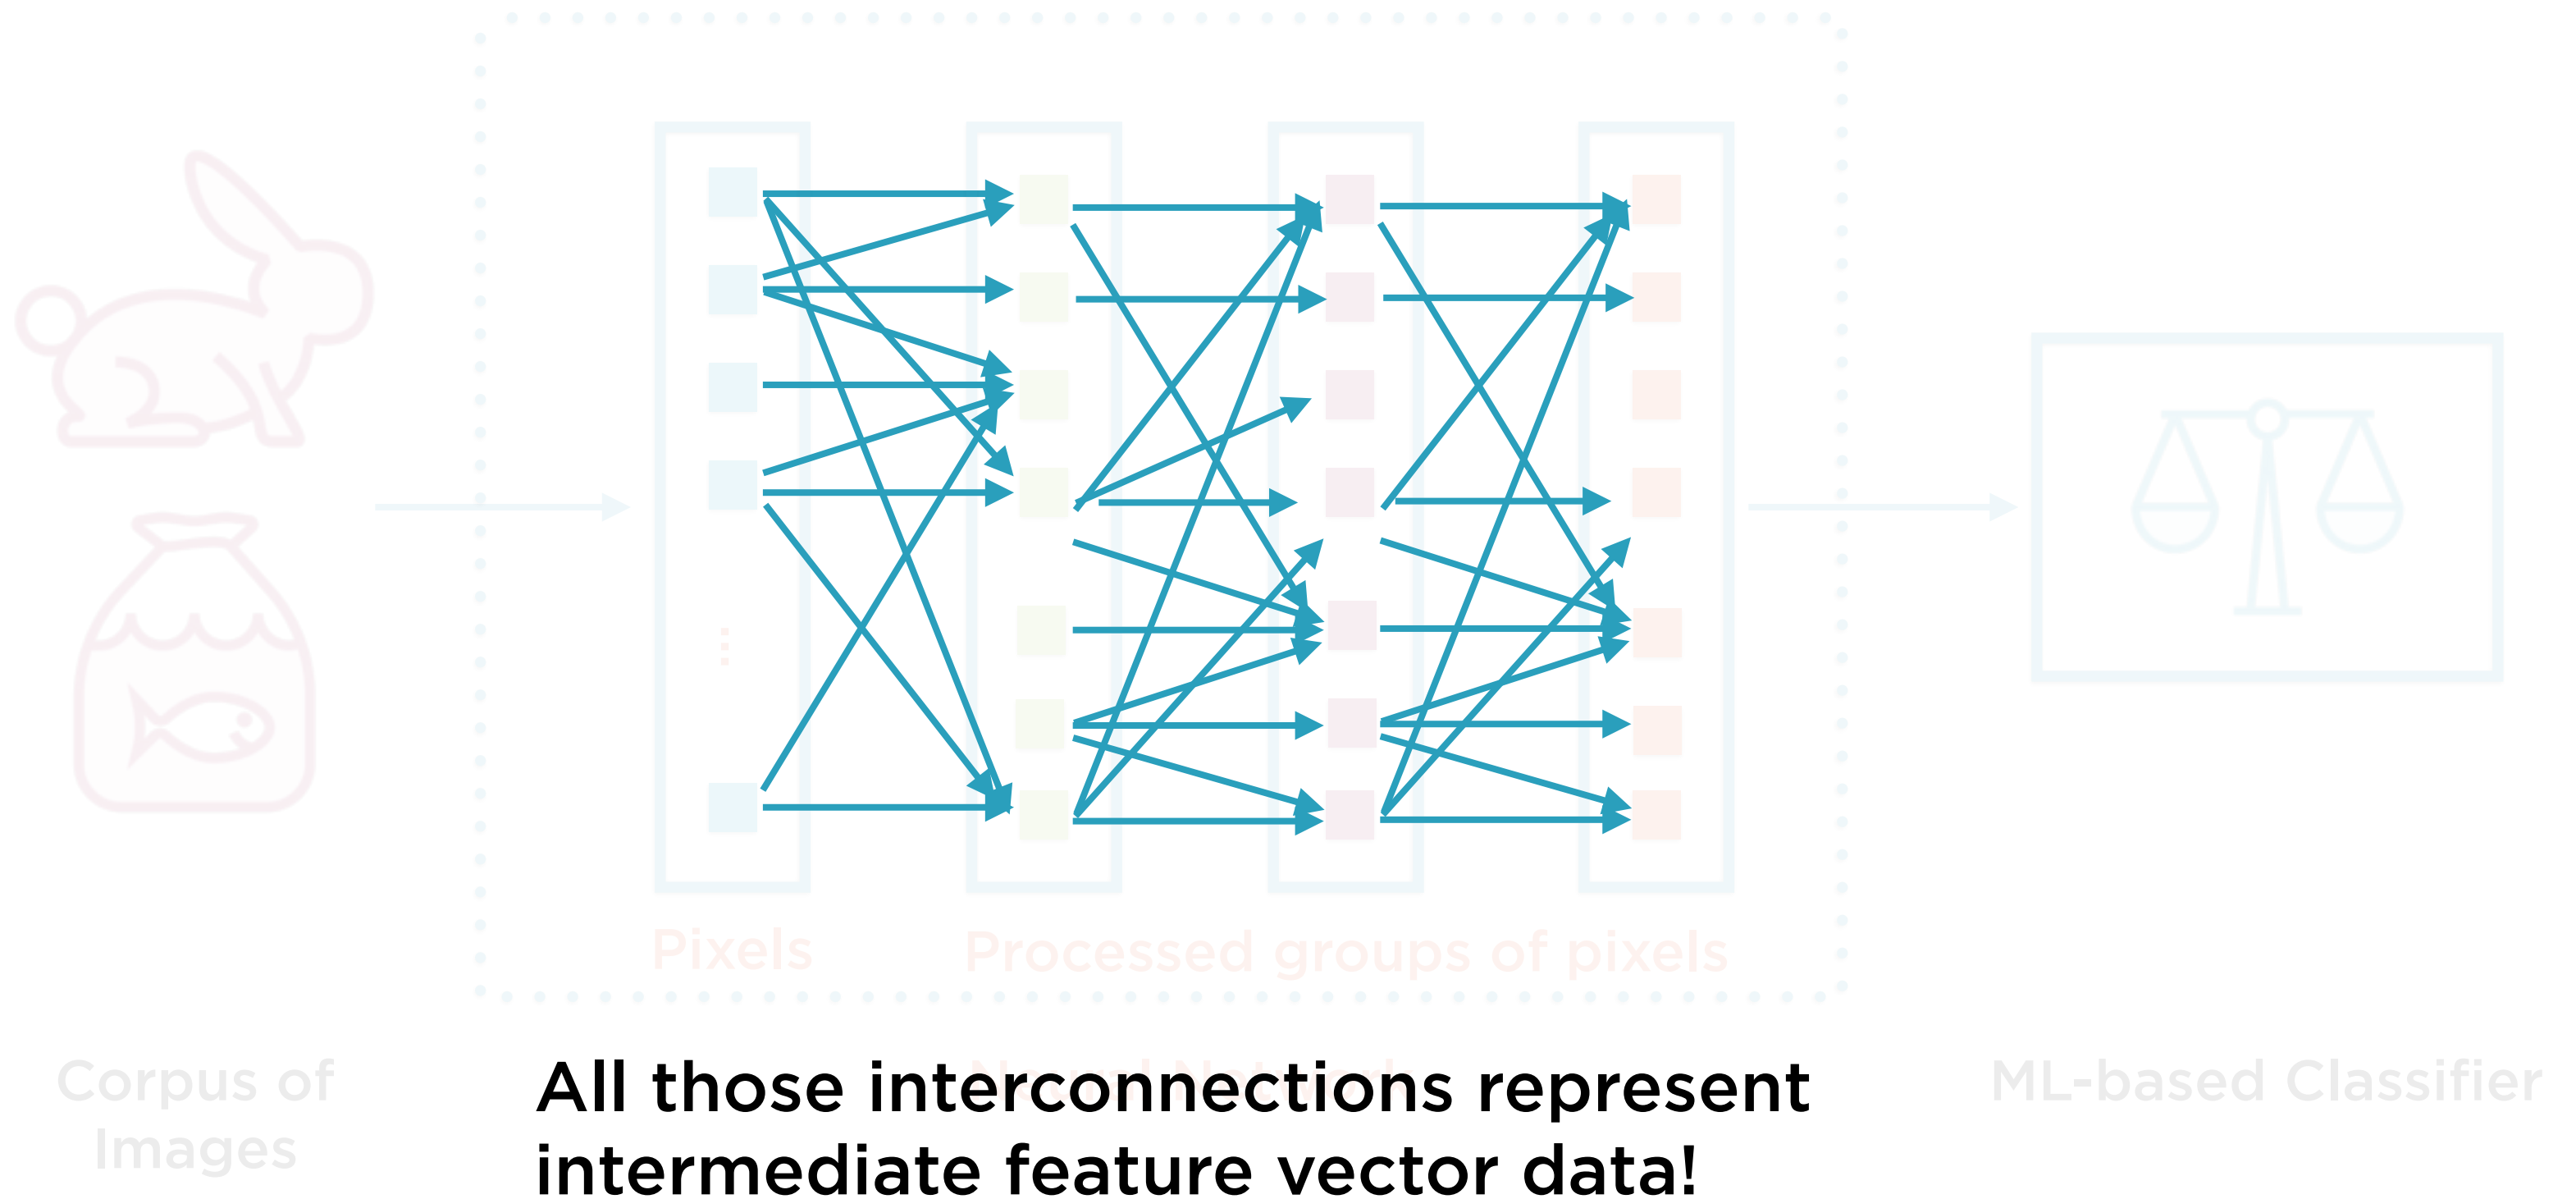
Complex Neural Networks



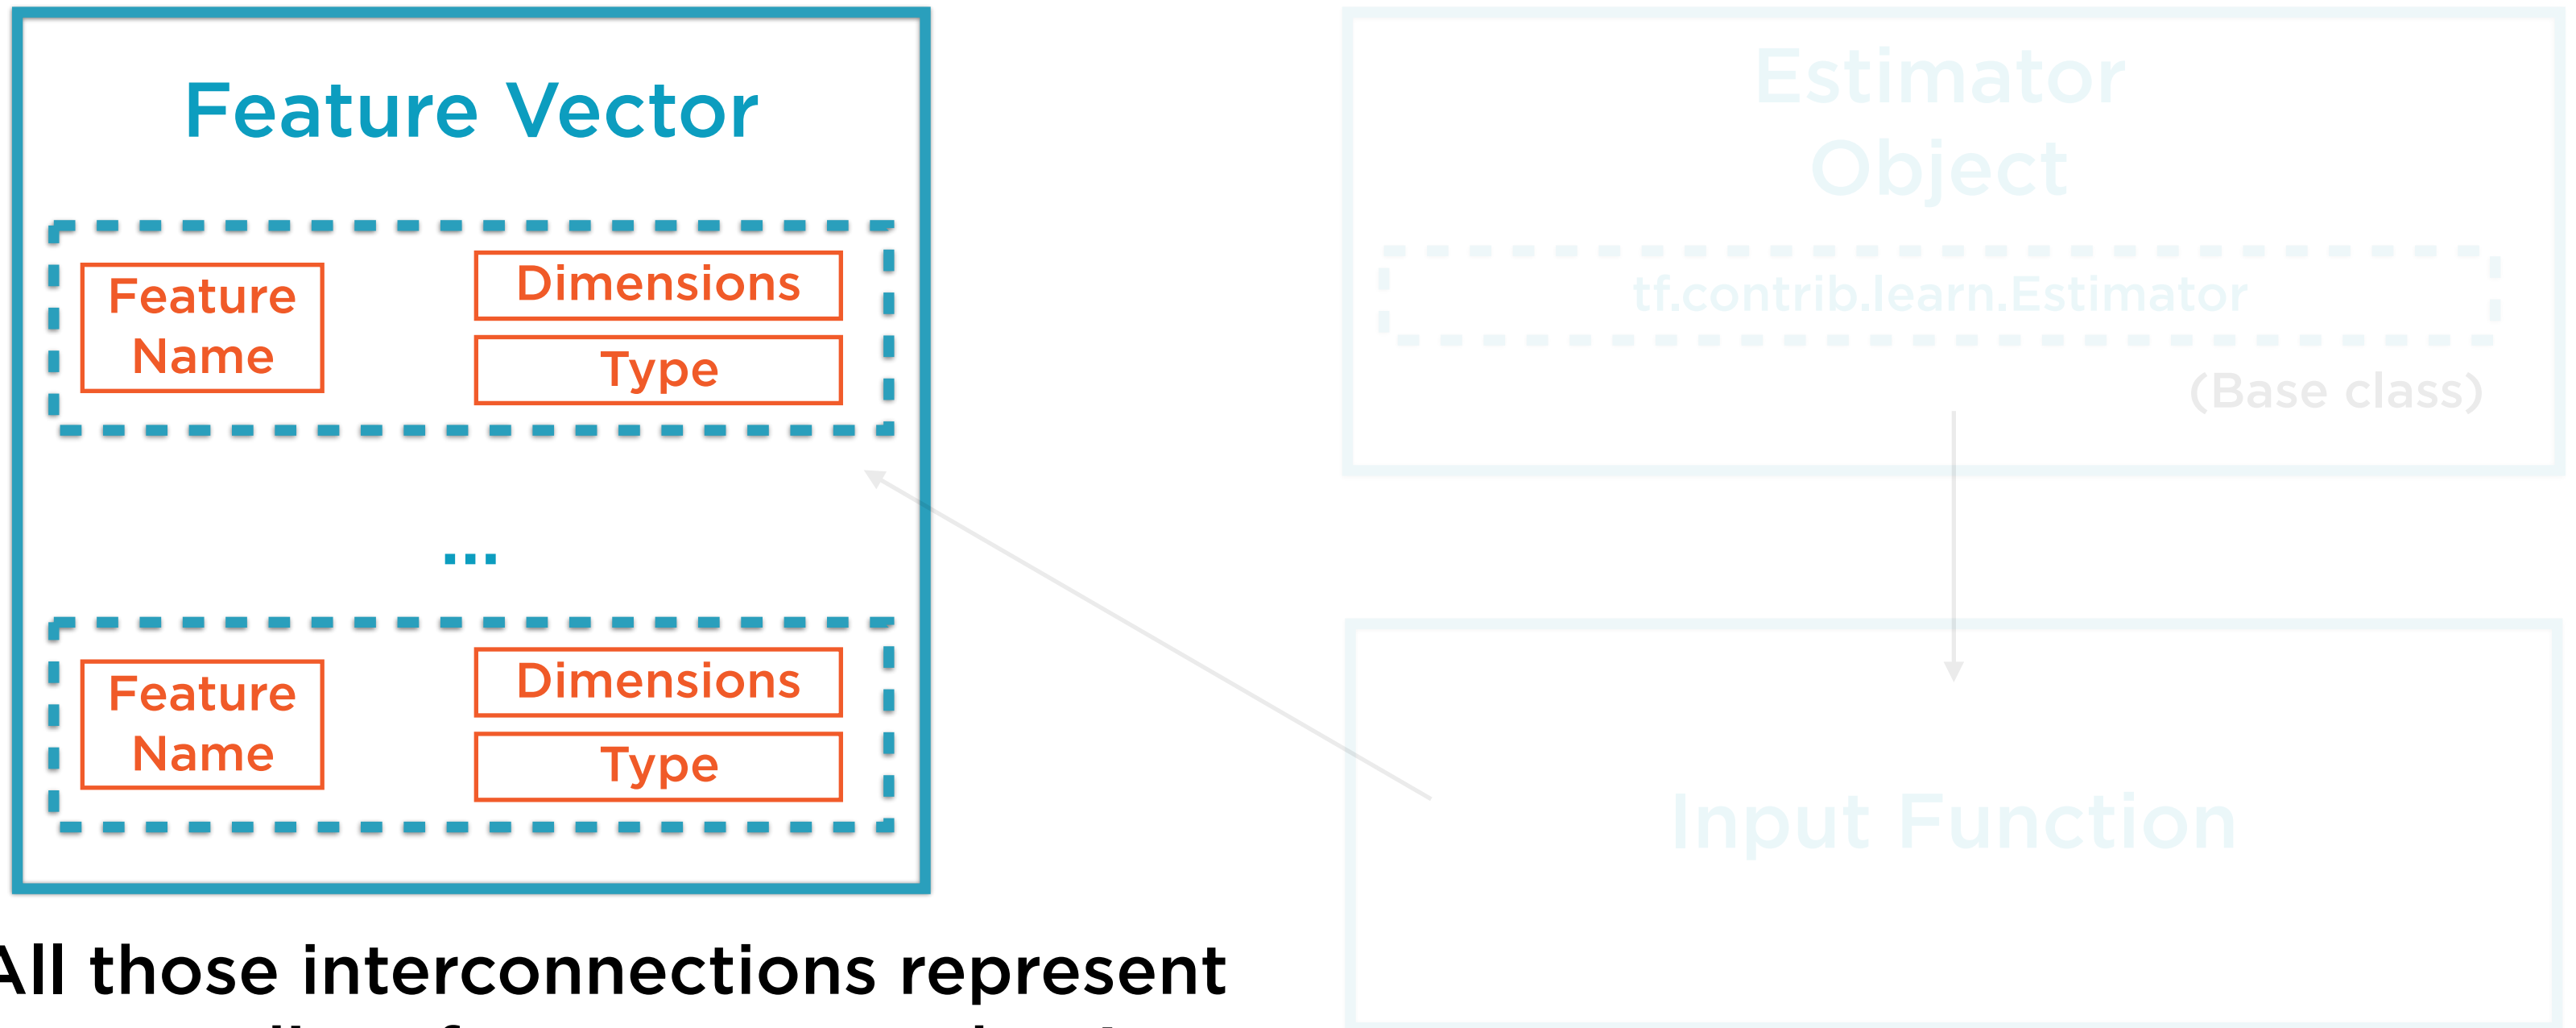
Complex Neural Networks



Complex Neural Networks

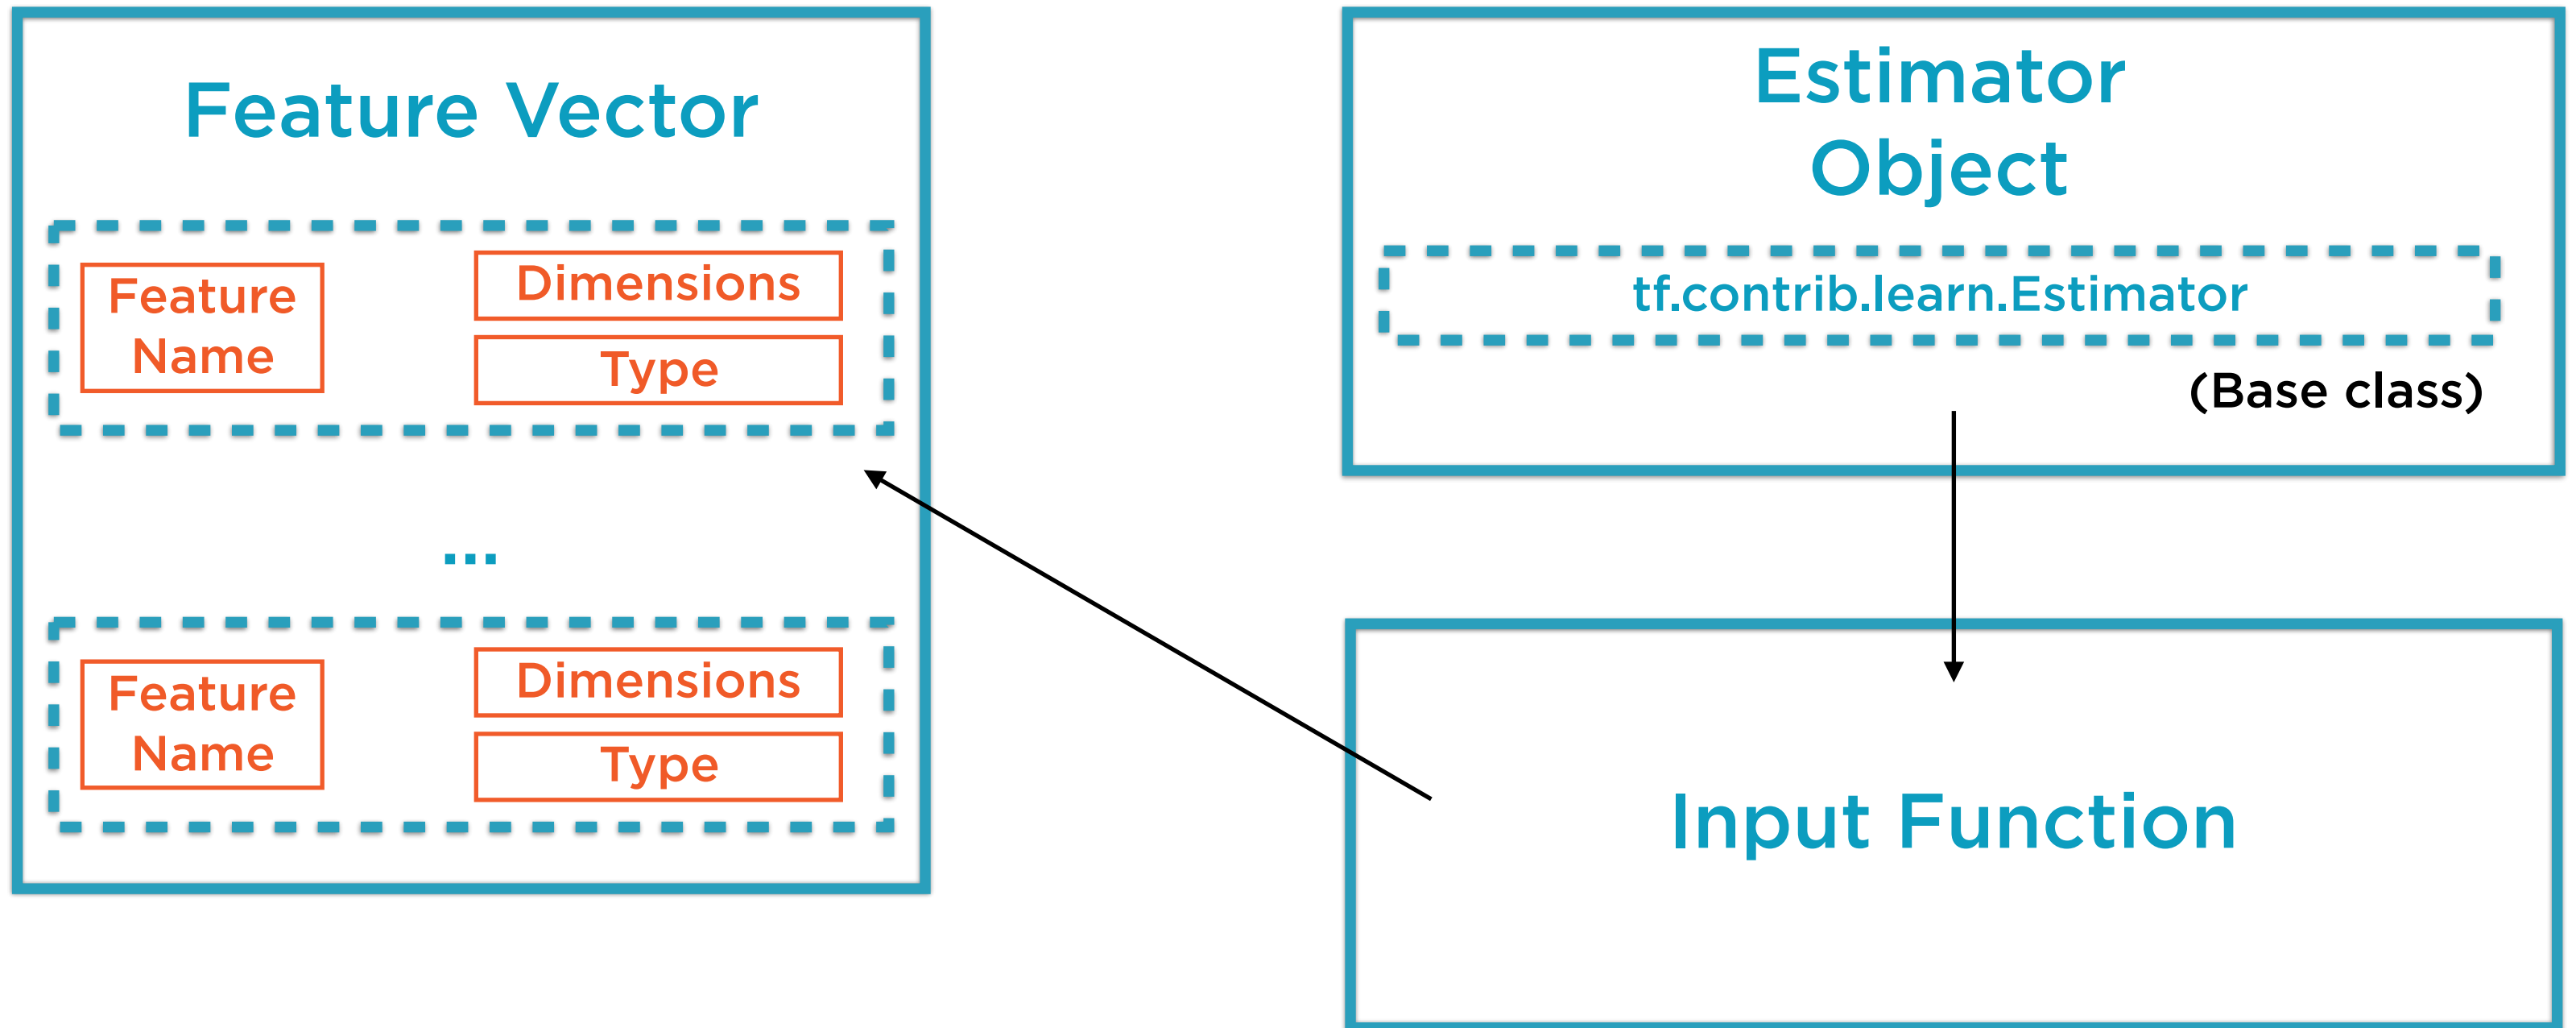


How Estimators Work



All those interconnections represent intermediate feature vector data!

How Estimators Work



Linear Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Mean Square Error (MSE)
Quantifying goodness-of-fit

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Affine transformation suffices

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

Logistic Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Cross Entropy
Similarity of distribution

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Softmax activation required

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

Linear Regression with an Estimator

Baseline

Non-TensorFlow implementation
Regular python code

Input Function

`tf.contrib.learn.io.numpy_input_fn`
Set up X, Y, batch_size, num_epochs

Evaluate

Use trained model
Predict new points (test data)

Instantiate Estimator

`tf.contrib.learn.LinearRegressor`
Abstracts cost and optimizer choices

Fit

Returns trained model
Can re-specify number of training steps

Course Outline

Learning using Neurons

Linear Regression in TensorFlow

Logistic Regression in TensorFlow

Estimators


“Representation” ML-based systems figure out by themselves what features to pay attention to

$$y = Wx + b$$

“Learning” Regression

Regression can be reverse-engineered by a single neuron

Regression: The Simplest Neural Network

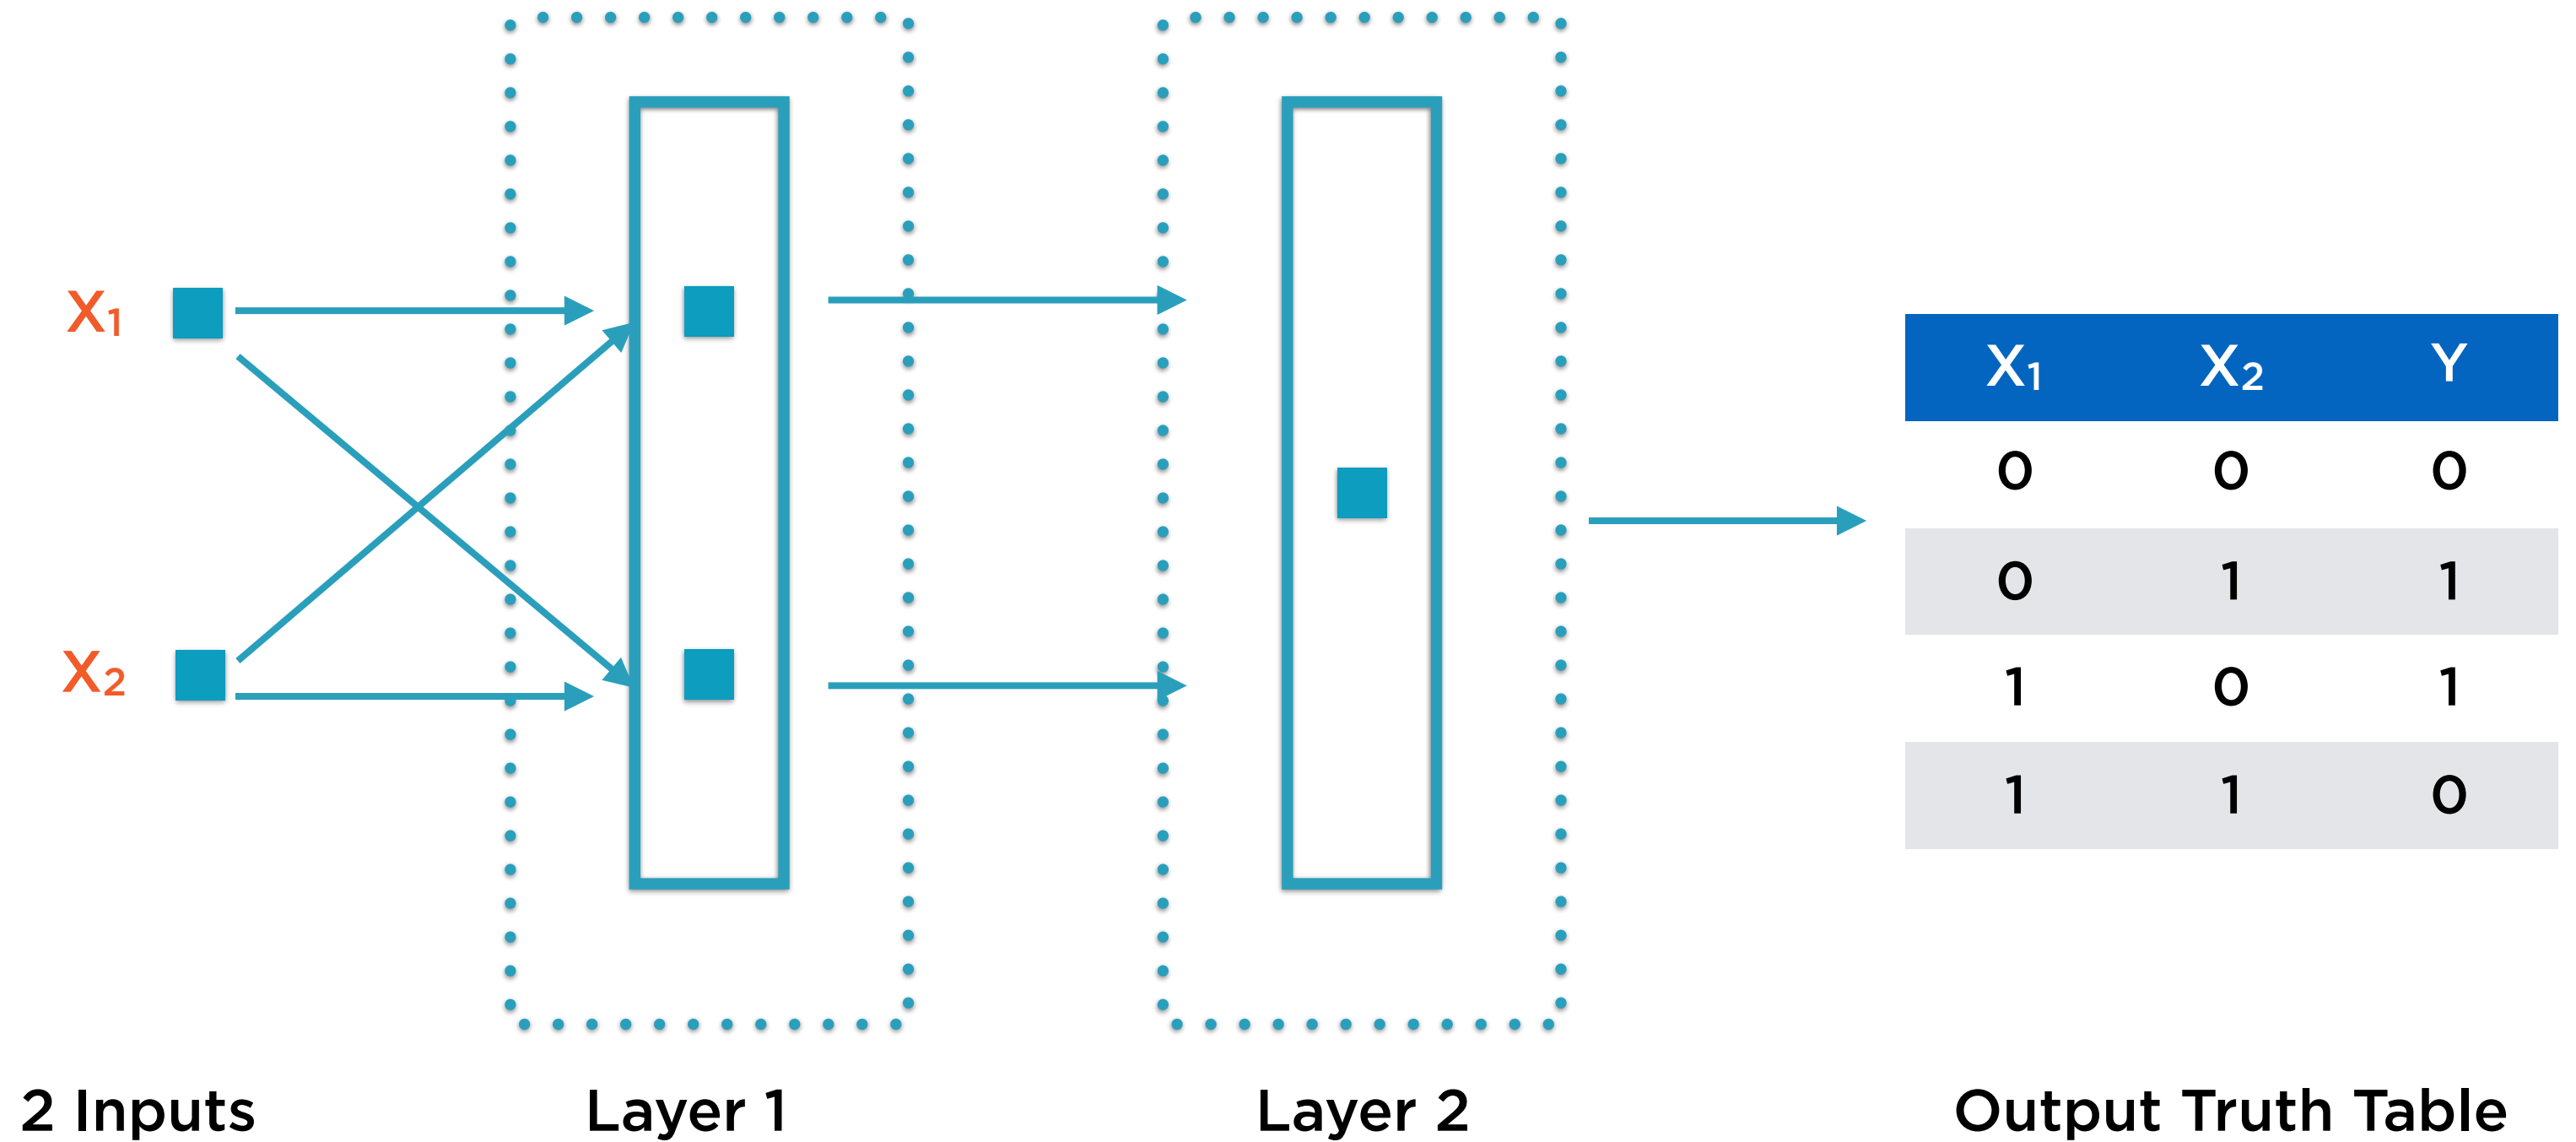


```
def XOR(x1, x2):  
    if (x1 == x2):  
        return 0  
    return 1
```

“Learning” XOR

The XOR function can be reverse-engineered using 3 neurons arranged in 2 layers

XOR: 3 Neurons, 2 Layers

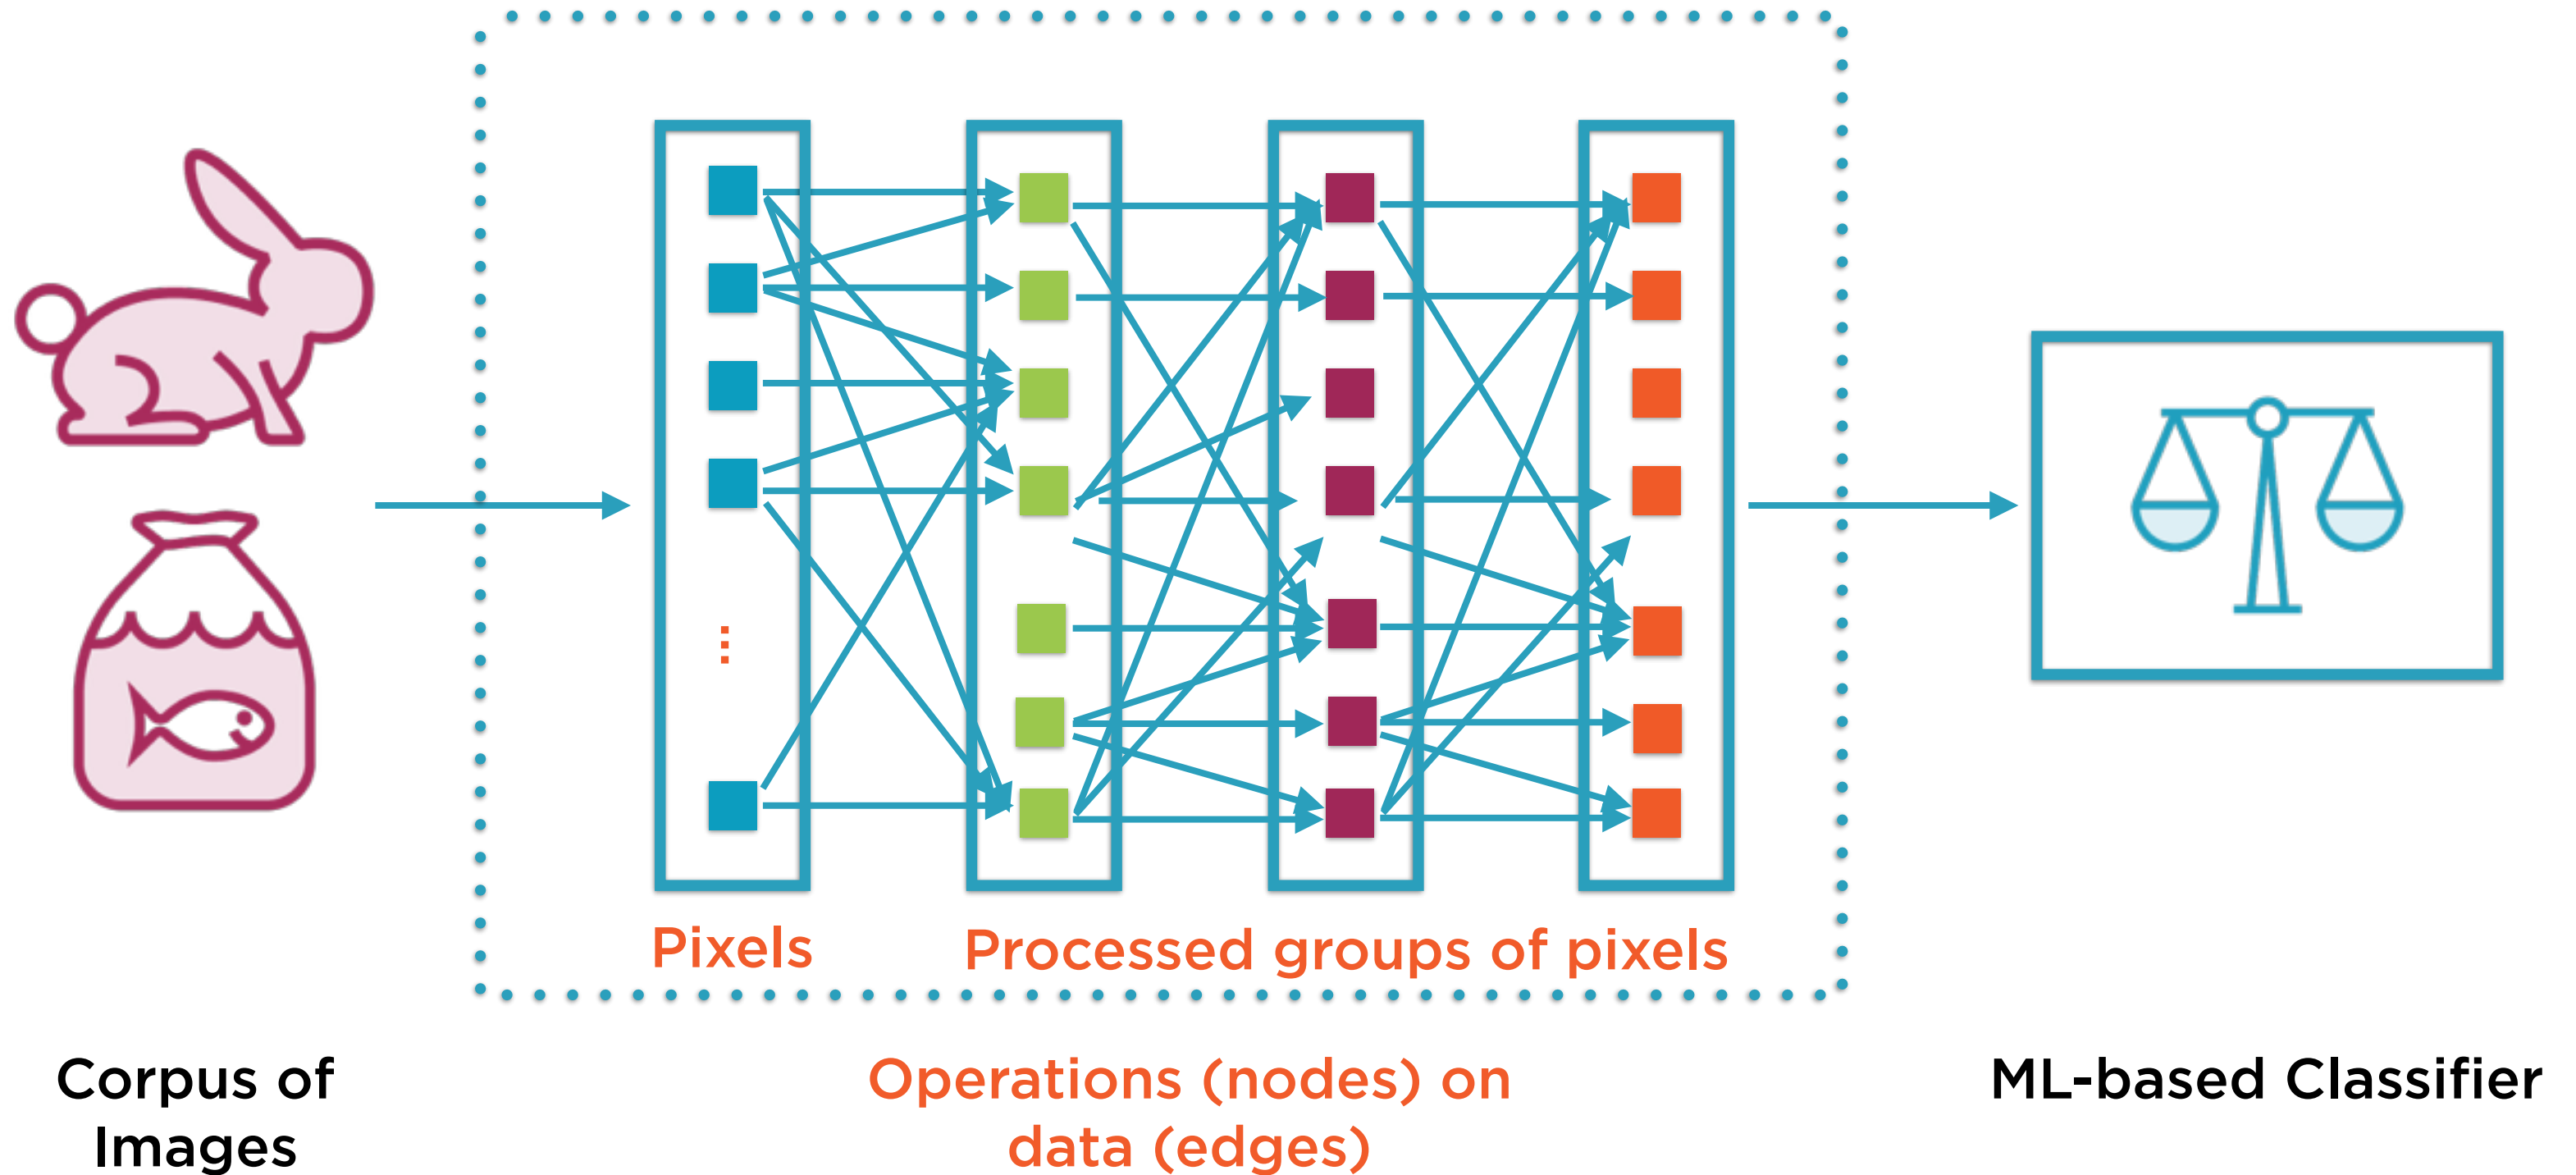


```
def doSomethingReallyComplicated(x1, x2...):  
    ...  
    ...  
    ...  
    return complicatedResult
```

“Learning” Arbitrarily Complex Functions

Adding layers to a neural network can “learn” (reverse-engineer) pretty much anything

Arbitrarily Complex Function



Linear Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Mean Square Error (MSE)
Quantifying goodness-of-fit

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Affine transformation suffices

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

Logistic Regression in TensorFlow

Baseline

Non-TensorFlow implementation
Regular python code

Cost Function

Cross Entropy
Similarity of distribution

Training

Invoke optimizer in epochs
Batch size for each epoch

Computation Graph

Neural network of 1 neuron
Softmax activation required

Optimizer

Gradient Descent optimizers
Improving goodness-of-fit

Converged Model

Values of W and b
Compare to baseline

Logistic Regression Using Estimators
