

Using K-nearest-neighbors for Digit Recognition



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Introduce the MNIST handwritten digit dataset

Understand the K-nearest-neighbors machine learning algorithm

Implement K-nearest-neighbors in TensorFlow to identify handwritten digits from 0 to 9

The MNIST Handwritten Digits Dataset

MNIST Dataset



Handwritten digits database

**Large quantity of handwritten digits
commonly used for training image
processing systems**

MNIST Dataset



Handwritten digits database


**Modified National Institute of Standards
and Technology**

MNIST Dataset



Each digit is in grayscale


MNIST Dataset



**Every image is
standardized to
be of size 28x28**

= 784 pixels

MNIST Dataset




Every pixel holds a **single** value for intensity

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

MNIST Dataset



Every image has an associated label


MNIST Dataset



5



0



4



1

MNIST for machine learning is
the equivalent of the “Hello
World” for programming

The K-nearest-neighbors Algorithm

Types of ML Algorithms



Supervised

Labels associated with the training data is used to correct the algorithm



Unsupervised

The model has to be set up right to learn structure in the data



Supervised Learning

Input variable x and output variable y

Learn the mapping function $y = f(x)$

Approximate the mapping function so
for new values of x we can predict y

Use existing dataset to **correct** our
mapping function approximation



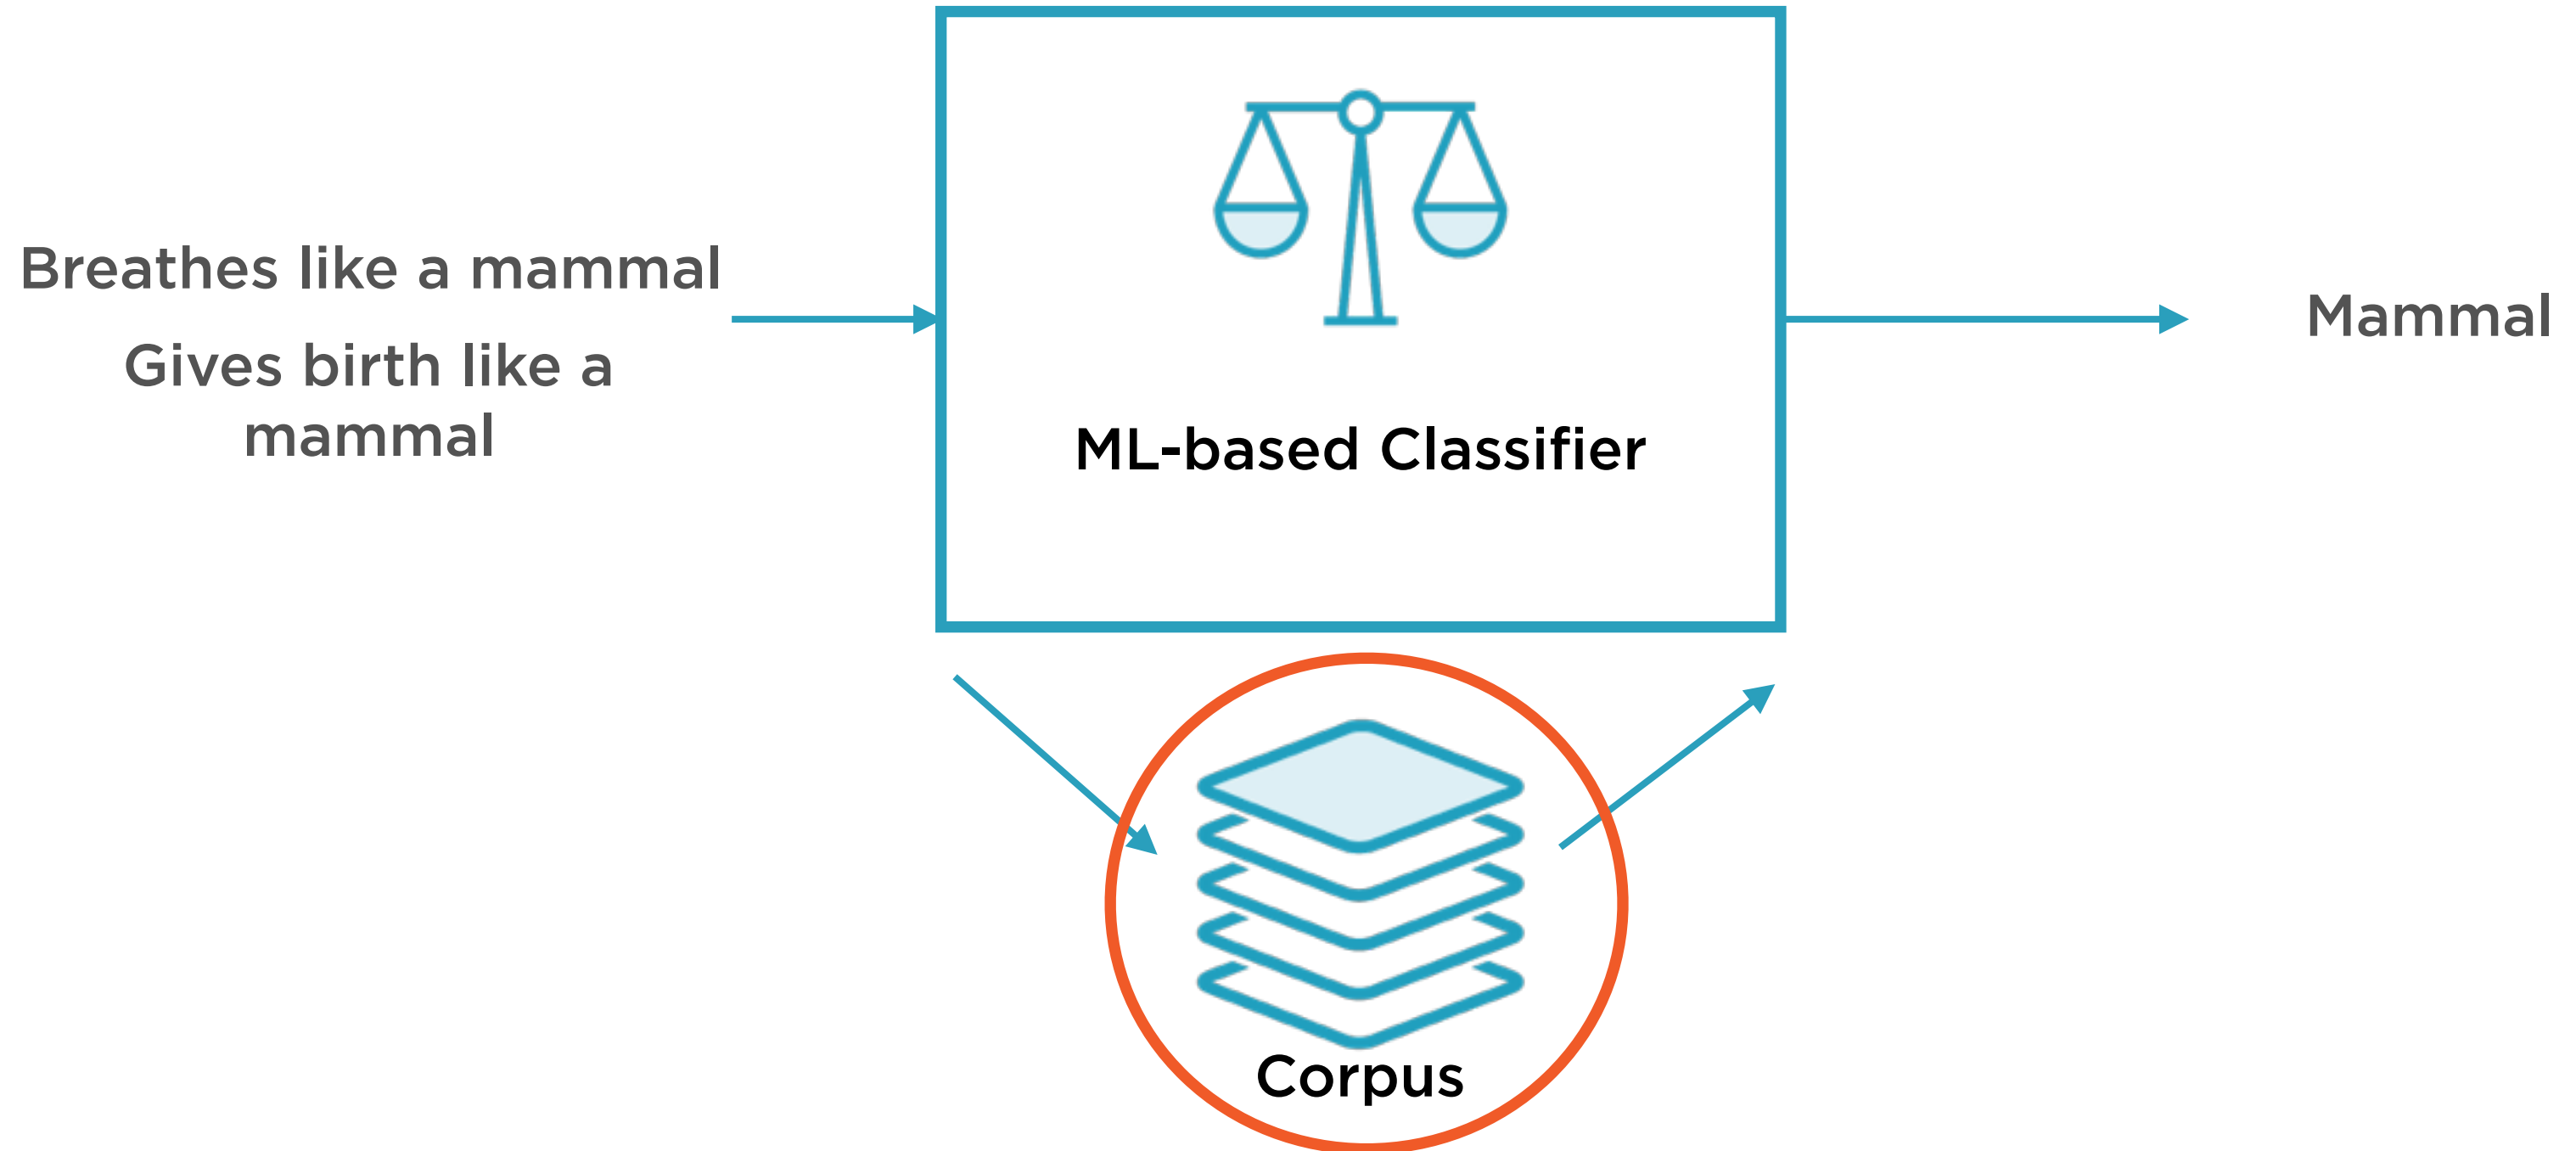
Unsupervised Learning

Only have input data **x** no output data

Model the underlying structure to learn more about data

Algorithms **self discover** the patterns and structure in the data

Training Data




KNN is an supervised learning algorithm which uses training data to find what is **most similar** to the current sample

K-nearest-neighbors



**Uses the entire training
dataset as a model**

K-nearest-neighbors




Rocket



Buildings



Signal




Pig



Shop

**Each element in training
data has a **label****

K-nearest-neighbors




Rocket



Buildings



Signal



Pig




Shop



Predictions for a new sample involves figuring out which element in the training data it is **similar** to

The nearest neighbor

K-nearest-neighbors




Rocket



Buildings




Signal




Pig



Shop



K-nearest-neighbors




Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors




Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors



Rocket



Buildings



Signal




Pig



Shop



K-nearest-neighbors




Rocket



Buildings




Signal



Pig



Shop




K-nearest-neighbors




How do we calculate neighbors of a sample?

K-nearest-neighbors




Distance measures

K-nearest-neighbors




Euclidean distance, Hamming distance, Manhattan distance


K-nearest-neighbors




K-nearest-neighbors



K-nearest-neighbors




Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \text{sqrt}(\text{sum}((x_j - x_{ij})^2))$$

Distance Measures



$$\text{EuclideanDistance}(x, x_i) = \text{sqrt}(\text{sum}((x_j - x_{ij})^2))$$

Distance Measures




$$\text{EuclideanDistance}(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$$

Distance Measures

Euclidean Distance
As the crow flies

Distance Measure



L1 Distance


Distance Measure

L1 distance


Snake distance

City block distance

Manhattan distance




L1 Distance



$$5 - 1 = 4$$

L1 Distance




$$5 - 1 = 4$$

$$4 - 0 = 4$$

L1 Distance

1, 0



5, 4

$$\begin{aligned} 5-1 &= 4 \\ 4-0 &= 4 \\ &= 8 \end{aligned}$$

Demo

Handwritten image recognition using the k-nearest-neighbors ML algorithm

- Use the L1 distance measure to find the nearest neighbor
- Measure the accuracy of the algorithm on the test data

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries

Running the algorithm

Predict labels for all the test data and measure accuracy

Calculating L1 distance

Find the distance between the test digit and all training digits

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries



MNIST Dataset



5



0




4



1

MNIST Dataset



**Every image is
standardized to
be of size 28x28**

= 784 pixels

Representing Labels



Vector

Index

0	0	0	0	1	0	0	0	0	0
0	1	2	3	4	5	6	7	8	9

4

Representing Labels




Vector

Index

0	0	0	0	0	1	0	0	0	0
0	1	2	3	4	5	6	7	8	9

5

Representing Images




= 784 pixels

Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

= 784 pixels


Representing Images



0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0
---	---	---	---	---	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0.2	0.8	0	0.3	0.6	0
---	---	---	---	---	---	-----	-----	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	→
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0.2	0.8	0	0.3	0.6	0	0.2	0.9	0	0.3	0.8	0
---	---	---	---	---	---	-----	-----	---	-----	-----	---	-----	-----	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---


0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images

0	0	0	0	0	0	0	0	0	0.2	0.2	0
---	---	---	---	---	---	----	----	----	----	----	----	---	---	---	-----	-----	---

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Representing Images



= 784 pixels

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries

Calculating L1 distance

Find the distance between the test digit and all training digits

L1 Distance

Training

0.2	0.8	0	0.3	0.6	0
-----	-----	---	-----	-----	---

0.2	0.8	0	0.3	0.6	0
-----	-----	---	-----	-----	---

Test

0	0.6	0.2	0.3	0.3	0.1
---	-----	-----	-----	-----	-----

tf.negative()

0	-0.6	-0.2	-0.3	-0.3	-0.1
---	------	------	------	------	------

tf.add()

0.2	0.2	-0.2	0	0.3	-0.1
-----	-----	------	---	-----	------

L1 Distance

tf.add()

0.2	0.2	-0.2	0	0.3	-0.1
-----	-----	------	---	-----	------

tf.abs()

0.2	0.2	0.2	0	0.3	0.1
-----	-----	-----	---	-----	-----

tf.reduce_sum()

1.0

tf.add()

0	0.2	0.3	0.6	-0.3	-0.1
---	-----	-----	-----	------	------

tf.abs()

0	0.2	0.3	0.6	0.3	0.1
---	-----	-----	-----	-----	-----

tf.reduce_sum()

1.5

tf.add()

-0.2	0.4	-0.2	0	-0.3	-0.1
------	-----	------	---	------	------

tf.abs()

0.2	0.4	0.2	0	0.3	0.1
-----	-----	-----	---	-----	-----

tf.reduce_sum()

1.2

L1 Distance

`tf.reduce_sum()`

1.0

`tf.reduce_sum()`

1.5

`tf.reduce_sum()`

1.2

index = 0

KNN Implemented in TensorFlow

Getting MNIST images

Access the MNIST training and test images in batches using the TensorFlow libraries


Running the algorithm

Predict labels for all the test data and measure accuracy

Calculating L1 distance

Find the distance between the test digit and all training digits

Representing Labels



`np.argmax()`

Summary

Familiar with the MNIST handwritten digit dataset

Understood the logic behind the K-nearest-neighbors algorithm

Implemented K-nearest-neighbors using L1 distance to identify handwritten digits from 0 to 9